

**Object-Oriented Implementation of
Numerical Methods
An Introduction with Java & Smalltalk**

Didier H. Besset
Maintained by S. Ducasse and S. Stinckwich

January 24, 2015

This book is available as a free download from <https://github.com/SquareBracketAssociates/ArchiveOONumericalMethod>.

Copyright © 2001, 2015 by Didier H. Besset.

The contents of this book are protected under Creative Commons Attribution-ShareAlike 3.0 Unported license.
You are free:

to Share — to copy, distribute and transmit the work

to Remix — to adapt the work

Under the following conditions:

Attribution. You must attribute the work in the manner specified by the author or licensor (but not in any way that suggests that they endorse you or your use of the work).

Share Alike. If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.

- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to this web page: creativecommons.org/licenses/by-sa/3.0/
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.



Your fair dealing and other rights are in no way affected by the above.
This is a human-readable summary of the Legal Code (the full license):
creativecommons.org/licenses/by-sa/3.0/legalcode

Published by Square Bracket Associates, Switzerland.
First Edition, 2001

Preface

Si je savais une chose utile à ma nation qui fût ruineuse à une autre,
je ne la proposerais pas à mon prince,
parce que je suis homme avant d'être Français,
parce que je suis nécessairement homme,
et que je ne suis Français que par hasard.¹
Charles de Montesquieux

When I first encountered object-oriented programming I immediately became highly enthusiastic about it, mainly because of my mathematical inclination. After all I learned to use computers as a high-energy physicist. In mathematics, a new, high order concept is always based on previously defined, simpler, concepts. Once a property is demonstrated for a given concept it can be applied to any new concept sharing the same premises as the original one. In object-oriented language, this is called reuse and inheritance. Thus, numerical algorithms using mathematical concepts that can be mapped directly into objects.

This book is intended to be read by object-oriented programmers who need to implement numerical methods in their applications. The algorithms exposed here are mostly fundamental numerical algorithms with a few advanced ones. The purpose of the book is to show that implementing these algorithms in an object-oriented language is feasible and quite easily feasible. We expect readers to be able to implement their own favorite numerical algorithm after seeing the examples discussed in this book.

The scope of the book is limited. It is not a Bible about numerical algorithms. Such Bible-like books already exist and are quoted throughout the chapters. Instead I wanted to illustrate how mapping between mathematical concepts and computer objects. I have limited the book to algorithms, which I have implemented and used in real applications over 12 years of object-oriented programming. Thus, the reader can be certain that the algorithms have been tested in the field.

¹If I knew some trade useful to my country, but which would ruin another, I would not disclose it to my ruler, because I am a man before being French, because I belong to mankind while I am French only by a twist of fate.

Because the intent of the book is showing numerical methods to object-oriented programmers the code presented is commented in depth. Each algorithm is presented with the same organization. First the necessary equations are introduced with short explanation. This book is not one about mathematics so explanations are kept to a minimum. Then the general object-oriented architecture of the algorithm is presented. Finally, this book intending to be a practical one, the code implementation is exposed. First how to use it, for readers who are just interested in using the algorithm. Then, the code implementation is discussed and presented.

The question was which language to use?

I decided to select two languages: Smalltalk and Java. Java is an obvious choice today since the world seems to lean toward Java. In spite of all the current emphasis on Java, I personally prefer Smalltalk for any practical use. But, as the computer industry evolves, and because I am anything but a language fanatic, the reader must be aware that this last statement may no longer be true by the time these lines are read. New computer languages will certainly combine object-oriented features with novel concepts and I am ready to try any serious new contender.

Having stated my preference for Smalltalk does not mean that I despise Java. Java has brought forth interesting innovations and certainly paved a golden way for object-oriented programming. I personally have gained a lot of insight in learning Java.

The code in this book has been written in both languages and is presented in separate sections. My hope is that programmers in Smalltalk and Java will be exposed to the other language by reading this book. If they keep an open mind and do indeed read both code sections, they certainly will gain something from the comparison as I did myself in writing.

As far as possible each algorithm is presented with such an example of use. I did not want to build contrived examples. Instead I have used examples personally encountered in my professional life. Some people may think that some examples are coming from esoteric domains. This is not so. Each example has been selected for its generality. The reader should study each example regardless of the field of application and concentrate on the universal aspects of it.

Acknowledgements

The author wishes to express his thanks to the many people with whom he had interactions about the object-oriented approach — Smalltalk and Java in particular — on the various electronic forums. One special person is Kent Beck whose controversial statements raised hell and started spirited discussions. I also thank Kent for showing me tricks about the Refactoring Browser and eXtreme Programming. I also would like to thank Eric Clayberg for pulling me out of a ditch more than once and for making such fantastic Smalltalk tools.

A special mention goes to Prof. Donald Knuth for being an inspiration for me and many other programmers with his series of books *The Art of Computer Programming*, and for making this wonderful typesetting program \TeX . This present book was typeset with \TeX and \LaTeX .

Furthermore, I would like to give credit to a few people without whom this present book would never have been published. First, Joseph Pelrine who persuaded me that what I was doing was worth sharing with the rest of the object-oriented community.

The author expresses his most sincere thanks to the reviewers who toiled on the early manuscripts of this book. Without their open-mindedness this book would never have made it to a publisher.

Special thanks go to David N. Smith for triggering interesting thoughts about random number generators and to Dr. William Leo for checking the equations.

Another special mention goes to An Lam who did a terrific job at checking the orthodoxy of the Java code.

Finally my immense gratitude is due to Dr. Stéphane Ducasse of the University of Bern who checked the orthodoxy of the Smalltalk code and who did a terrific job of rendering the early manuscript not only readable but entertaining.

Genolier, 11 April 2000

Contents

1	Introduction	1
1.1	Object-oriented paradigm and mathematical objects	2
1.2	Object-oriented concepts in a nutshell	3
1.3	Dealing with numerical data	4
1.3.1	Floating point representation	4
1.3.2	Rounding errors	5
1.3.3	Real example of rounding error	7
1.3.4	Outsmarting rounding errors	8
1.3.5	Wisdom from the past	8
1.4	Finding the numerical precision of a computer	9
1.4.1	Computer numerical precision — General implementation	11
1.4.2	Computer numerical precision — Smalltalk implementation	11
1.4.3	Computer numerical precision — Java implementation	15
1.5	Comparing floating point numbers	21
1.5.1	Comparing floating point numbers — Smalltalk code	22
1.5.2	Comparing floating point numbers — Java code	22
1.6	Speed consideration	23
1.6.1	Smalltalk particular	24
1.7	Conventions	24
1.7.1	Class diagrams	25
1.7.2	Smalltalk code	26
1.7.3	Java code	26
1.8	Road map	27
2	Function evaluation	31
2.1	Function concept	32
2.1.1	Function — Smalltalk implementation	32
2.1.2	Function — Java implementation	33
2.2	Polynomials	36
2.2.1	Mathematical definitions	37
2.2.2	Polynomial — General implementation	38
2.2.3	Polynomial — Smalltalk implementation	39
2.2.4	Polynomial — Java implementation	45
2.3	Error function	53

2.3.1 Mathematical definitions	54
2.3.2 Error function — Smalltalk implementation	55
2.3.3 Error function — Java implementation	57
2.4 Gamma function	59
2.4.1 Mathematical definitions	59
2.4.2 Gamma function — Smalltalk implementation	61
2.4.3 Gamma function — Java implementation	63
2.5 Beta function	66
2.5.1 Mathematical definitions	66
2.5.2 Beta function — Smalltalk implementation	67
2.5.3 Beta function — Java implementation	67
3 Interpolation	69
3.1 General remarks	69
3.2 Lagrange interpolation	75
3.2.1 Lagrange interpolation — Smalltalk implementation	75
3.2.2 Lagrange interpolation — Java implementation	78
3.3 Newton interpolation	83
3.3.1 Newton interpolation — General implementation	84
3.3.2 Newton interpolation — Smalltalk implementation	85
3.3.3 Newton interpolation — Java implementation	86
3.4 Neville interpolation	88
3.4.1 Neville interpolation — General implementation	89
3.4.2 Neville interpolation — Smalltalk implementation	89
3.4.3 Neville interpolation — Java implementation	91
3.5 Bulirsch-Stoer interpolation	94
3.5.1 Bulirsch-Stoer interpolation — General implementation	95
3.5.2 Bulirsch-Stoer interpolation — Smalltalk implementation	95
3.5.3 Bulirsch-Stoer interpolation — Java implementation	96
3.6 Cubic spline interpolation	97
3.6.1 Cubic spline interpolation — General implementation	99
3.6.2 Cubic spline interpolation — Smalltalk implementation	99
3.6.3 Cubic spline interpolation — Java implementation	102
3.7 Which method to choose?	105
4 Iterative algorithms	107
4.1 Successive approximations	107
4.1.1 Iterative process — Smalltalk implementation	112
4.1.2 Iterative process: JAVA implementation	115
4.2 Evaluation with relative precision	120
4.2.1 Relative precision — Smalltalk implementation	121
4.2.2 Relative precision — Java implementation	123
4.3 Examples	125
5 Finding the zero of a function	127
5.1 Introduction	127
5.2 Finding the zeroes of a function — Bisection method	128

5.2.1	Bisection algorithm — General implementation	129
5.2.2	Bisection algorithm — Smalltalk implementation	130
5.2.3	Bisection algorithm — Java implementation	132
5.3	Finding the zero of a function — Newton's method	134
5.3.1	Newton's method — Smalltalk implementation	135
5.3.2	Newton's method — Java implementation	137
5.4	Example of zero-finding — Roots of polynomials	140
5.4.1	Roots of polynomials — Smalltalk implementation	141
5.4.2	Roots of polynomials — Java implementation	142
5.5	Which method to choose	143
6	Integration of functions	145
6.1	Introduction	145
6.2	General framework — Trapeze integration method	146
6.2.1	Trapeze integration — General implementation	149
6.2.2	Trapeze integration — Smalltalk implementation	149
6.2.3	Trapeze integration — Java implementation	151
6.3	Simpson integration algorithm	154
6.3.1	Simpson integration — General implementation	154
6.3.2	Simpson integration — Smalltalk implementation	155
6.3.3	Simpson integration — Java implementation	155
6.4	Romberg integration algorithm	156
6.4.1	Romberg integration — General implementation	157
6.4.2	Romberg integration — Smalltalk implementation	158
6.4.3	Romberg integration — Java implementation	159
6.5	Evaluation of open integrals	161
6.6	Which method to chose?	162
6.6.1	Smalltalk comparison	163
6.6.2	Java comparison	164
7	Series	167
7.1	Introduction	167
7.2	Infinite series	168
7.2.1	Infinite series — Smalltalk implementation	170
7.2.2	Infinite series — Java implementation	171
7.3	Continued fractions	173
7.3.1	Continued fractions — Smalltalk implementation	174
7.3.2	Continued fractions — Java implementation	175
7.4	Incomplete Gamma function	177
7.4.1	Mathematical definitions	177
7.4.2	Incomplete Gamma function — Smalltalk implementation	179
7.4.3	Incomplete Gamma function — Java implementation	182
7.5	Incomplete Beta function	186
7.5.1	Mathematical definitions	187
7.5.2	Incomplete Beta function — Smalltalk implementation	188
7.5.3	Incomplete Beta function — Java implementation	191

8	Linear algebra	195
8.1	Vectors and matrices	195
8.1.1	Vector and matrix — Smalltalk implementation	199
8.1.2	Vector and matrix — Java implementation	210
8.2	Linear equations	230
8.2.1	Linear equations — General implementation	233
8.2.2	Linear equations — Smalltalk implementation	233
8.2.3	Linear equations — Java implementation	238
8.3	LUP decomposition	244
8.3.1	LUP decomposition — General implementation	246
8.3.2	LUP decomposition — Smalltalk implementation	247
8.3.3	LUP decomposition — Java implementation	251
8.4	Computing the determinant of a matrix	259
8.4.1	Computing the determinant of matrix — General implementation	260
8.4.2	Computing the determinant of matrix — Smalltalk implementation	260
8.4.3	Computing the determinant of matrix — Java implementation	261
8.5	Matrix inversion	261
8.5.1	Matrix inversion — Smalltalk implementation	264
8.5.2	Matrix inversion — Java implementation	266
8.5.3	Matrix inversion — Rounding problems	274
8.6	Matrix eigenvalues and eigenvectors of a non-symmetric matrix	275
8.6.1	Finding the largest eigenvalue — General implementation	276
8.6.2	Finding the largest eigenvalue — Smalltalk implementation	277
8.6.3	Finding the largest eigenvalue — Java implementation	279
8.7	Matrix eigenvalues and eigenvectors of a symmetric matrix	282
8.7.1	Jacobi's algorithm — General implementation	286
8.7.2	Jacobi's algorithm — Smalltalk implementation	287
8.7.3	Jacobi's algorithm — Java implementation	291
9	Elements of statistics	297
9.1	Statistical moments	297
9.1.1	Statistical moments — General implementation	299
9.1.2	Statistical moments — Smalltalk implementation	300
9.1.3	Statistical moments — Java implementation	302
9.2	Robust implementation of statistical moments	305
9.2.1	Robust central moments — General implementation	307
9.2.2	Robust central moments — Smalltalk implementation	307
9.2.3	Robust central moments — Java implementation	311
9.3	Histograms	317
9.3.1	Histograms — General implementation	318
9.3.2	Histograms — Smalltalk implementation	319
9.3.3	Histograms — Java implementation	329
9.4	Random number generator	344
9.4.1	Random number generator — Smalltalk implementation	346
9.4.2	Random number generator — Java implementation	352
9.5	Probability distributions	355

9.5.1 Probability distributions — General implementation	357
9.5.2 Probability distributions — Smalltalk implementation	357
9.5.3 Probability distributions — Java implementation	362
9.6 Normal distribution	370
9.6.1 Normal distribution — Smalltalk implementation	370
9.6.2 Normal distribution — Java implementation	373
9.7 Gamma distribution	378
9.7.1 Gamma distribution — Smalltalk implementation	378
9.7.2 Gamma distribution — Java implementation	383
9.8 Experimental distribution	389
9.8.1 Experimental distribution — General implementation	390
9.8.2 Experimental distribution — Smalltalk implementation	390
9.8.3 Experimental distribution — Java implementation	392
10 Statistical analysis	397
10.1 F -test and the Fisher-Snedecor distribution	398
10.1.1 Fisher-Snedecor distribution — Smalltalk implementation	400
10.1.2 Fisher-Snedecor distribution — Java implementation	405
10.2 t -test and the Student distribution	410
10.2.1 Student distribution — Smalltalk implementation	413
10.2.2 Student distribution — Java implementation	417
10.3 χ^2 -test and χ^2 distribution	422
10.3.1 χ^2 distribution — Smalltalk implementation	424
10.3.2 χ^2 distribution — Java implementation	426
10.3.3 Weighted point implementation	428
10.4 χ^2 -test on histograms	434
10.4.1 χ^2 -test on histograms — Smalltalk implementation	435
10.4.2 χ^2 -test on histograms — Java implementation	439
10.5 Definition of estimation	442
10.5.1 Maximum likelihood estimation	443
10.5.2 Least square estimation	443
10.6 Least square fit with linear dependence	445
10.7 Linear regression	446
10.7.1 Linear regression — General implementation	447
10.7.2 Linear regression — Smalltalk implementation	448
10.7.3 Linear regression — Java implementation	451
10.8 Least square fit with polynomials	456
10.8.1 Polynomial least square fits — Smalltalk implementation	459
10.8.2 Polynomial least square fits — Java implementation	462
10.9 Least square fit with non-linear dependence	467
10.9.1 Non-linear fit — General implementation	468
10.9.2 Non-linear fit — Smalltalk implementation	470
10.9.3 Non-linear fit — Java implementation	474
10.10 Maximum likelihood fit of a probability density function	484
10.10.1 Maximum likelihood fit — General implementation	487
10.10.2 Maximum likelihood fit — Smalltalk implementation	487

10.10. Maximum likelihood fit — Java implementation	490
11 Optimization	495
11.1 Introduction	497
11.2 Extended Newton algorithms	498
11.3 Hill climbing algorithms	499
11.3.1 Optimizing — General implementation	499
11.3.2 Common optimizing classes — Smalltalk implementation	500
11.3.3 Common optimizing classes — Java implementation	507
11.4 Optimizing in one dimension	523
11.4.1 Optimizing in one dimension — Smalltalk implementation	523
11.4.2 Optimizing in one dimension — Java implementation	526
11.5 Bracketing the optimum in one dimension	529
11.5.1 Bracketing the optimum — Smalltalk implementation	529
11.5.2 Bracketing the optimum — Java implementation	531
11.6 Powell's algorithm	533
11.6.1 Powell's algorithm — General implementation	534
11.6.2 Powell's algorithm — Smalltalk implementation	535
11.6.3 Powell's algorithm — Java implementation	537
11.7 Simplex algorithm	540
11.7.1 Simplex algorithm — General implementation	542
11.7.2 Simplex algorithm — Smalltalk implementation	542
11.7.3 Simplex algorithm — Java implementation	545
11.8 Genetic algorithm	550
11.8.1 Genetic algorithm — General implementation	553
11.8.2 Genetic algorithm — Smalltalk implementation	555
11.8.3 Genetic algorithm — Java abstract implementation	560
11.8.4 Genetic algorithm — Java implementation with vectors	567
11.9 Multiple strategy approach	574
11.9.1 Multiple strategy approach — General implementation	574
12 Data mining	579
12.1 Data server	580
12.1.1 Data server — Smalltalk implementation	581
12.1.2 Data server — Java implementation	583
12.2 Covariance and covariance matrix	585
12.2.1 Covariance matrix — General implementation	587
12.2.2 Covariance matrix — Smalltalk implementation	587
12.2.3 Covariance matrix — Java implementation	590
12.3 Multidimensional probability distribution	596
12.4 Covariance data reduction	596
12.5 Mahalanobis distance	597
12.5.1 Mahalanobis distance — General implementation	598
12.5.2 Mahalanobis distance — Smalltalk implementation	599
12.5.3 Mahalanobis distance — Java implementation	602
12.6 Cluster analysis	604

12.6.1	Cluster analysis — General implementation	607
12.6.2	Cluster analysis — Smalltalk implementation	608
12.6.3	Cluster analysis — Java implementation	613
12.7	Covariance clusters	622
12.7.1	Covariance clusters — General implementation	622
A	Decimal floating-point simulation	627
B	Smalltalk primer for Java programmers	631
B.1	Syntax in a nutshell	631
B.1.1	Smalltalk expressions	631
B.1.2	Precedence	632
B.1.3	Assignment, equality and identity	632
B.2	Class and methods	633
B.2.1	Instance methods	633
B.2.2	Class methods	635
B.2.3	Block	635
B.3	Iterator methods	636
B.3.1	do:	636
B.3.2	collect:	636
B.3.3	inject:into:	637
B.4	Double dispatching	637
B.5	Multiple dispatching	638
C	Java primer for Smalltalk programmers	641
C.1	Remarks on the syntax	641
C.1.1	Classes	641
C.1.2	Instance variables	642
C.1.3	Method declaration and method calling	642
C.1.4	Objects and non-objects	643
C.1.5	packages	643
C.1.6	Scope qualifiers	644
C.1.7	static qualifier	646
C.2	Abstract class and interface	646
C.3	Exception handling	647
C.4	Collections and related topics	648
D	Additional probability distributions	651
D.1	Beta distribution	651
D.1.1	Beta distribution — Smalltalk implementation	651
D.1.2	Beta distribution — Java implementation	655
D.2	Cauchy distribution	660
D.2.1	Cauchy distribution — Smalltalk implementation	660
D.2.2	Cauchy distribution — Java implementation	663
D.3	Exponential distribution	667
D.3.1	Exponential distribution — Smalltalk implementation	667
D.3.2	Exponential distribution — Java implementation	670

D.4 Fisher-Tippett distribution	674
D.4.1 Fisher-Tippett distribution — Smalltalk implementation	674
D.4.2 Fisher-Tippett distribution — Java implementation	678
D.5 Laplace distribution	682
D.5.1 Laplace distribution — Smalltalk implementation	682
D.5.2 Laplace distribution — Java implementation	685
D.6 Log normal distribution	690
D.6.1 Log normal distribution — Smalltalk implementation	690
D.6.2 Log normal distribution — Java implementation	693
D.7 Triangular distribution	696
D.7.1 Triangular distribution — Smalltalk implementation	697
D.7.2 Triangular distribution — Java implementation	700
D.8 Uniform distribution	704
D.8.1 Uniform distribution — Smalltalk implementation	704
D.8.2 Uniform distribution — Java implementation	706
D.9 Weibull distribution	709
D.9.1 Weibull distribution — Smalltalk implementation	710
D.9.2 Weibull distribution — Java implementation	713
E Accurate accumulation of expectation values	719
E.1 Accurate accumulation of central moments	719
E.2 Accurate accumulation of the covariance	721
F Accompanying CD ROM	723
F.1 Smalltalk code	723
F.1.1 non-ENVY users	723
F.1.2 ENVY users	725
F.2 Java code	726
F.2.1 Visual Age for Java	726
F.2.2 Other Java systems	726
F.3 Utilities	727
F.3.1 Distribution demo	727
F.3.2 File reader	727

List of Figures

1.1 Comparison of achieved precision	7
1.2 A typical class diagram	25
1.3 Book road map	28
2.1 Smalltalk classes related to functions	32
2.2 Java classes related to functions	33
2.3 The error function and the normal distribution	53
3.1 Class diagram for the interpolation classes	70
3.2 Example of interpolation with the Lagrange interpolation polynomial . . .	71
3.3 Comparison between Lagrange interpolation and interpolation with a rational function	72
3.4 Comparison of Lagrange interpolation and cubic spline	73
3.5 Example of misbehaving interpolation	74
4.1 Class diagram for iterative process classes	108
4.2 Successive approximation algorithm	108
4.3 Detailed algorithm for successive approximations	110
4.4 Methods for successive approximations	111
5.1 Class diagram for zero finding classes	128
5.2 The bisection algorithm	129
5.3 Geometrical representation of Newton's zero finding algorithm	135
6.1 Class diagram of integration classes	146
6.2 Geometrical interpretation of the trapeze integration method	147
7.1 Smalltalk class diagram for infinite series and continued fractions	168
7.2 Java class diagram for infinite series and continued fractions	169
7.3 The incomplete gamma function and the gamma distribution	178
7.4 The incomplete beta function and the beta distribution	187
8.1 Linear algebra classes	196
8.2 Comparison of inversion time for non-symmetrical matrices	263

9.1	Classes related to statistics	298
9.2	A typical histogram	317
9.3	Normal distribution for various values of the parameters	371
9.4	Gamma distribution for various values of α	380
10.1	Classes related to estimation	398
10.2	Fisher-Snedecor distribution for a few parameters	400
10.3	Student distribution for a few degrees of freedom	412
10.4	χ^2 distribution for a few degrees of freedom	424
10.5	Example of polynomial fit	457
10.6	Fit results for the fit of figure 10.5	458
10.7	Limitation of polynomial fits	459
10.8	Example of a least square fit	469
10.9	Example of a maximum likelihood fit	486
11.1	Smalltak classes used in optimization	496
11.2	Java classes used in optimization	496
11.3	Local and absolute optima	497
11.4	Operations of the simplex algorithm	542
11.5	Mutation and crossover reproduction of chromosomes	551
11.6	General purpose genetic algorithm	552
11.7	Compared behavior of hill climbing and random based algorithms.	575
12.1	Classes used in data mining	580
12.2	Using the Mahalanobis distance to differentiate between good and fake coins.	599
12.3	Example of cluster algorithm	606
B.1	Triple dispatching	638
B.2	Triple dispatching	639
D.1	Many shapes of the beta distribution	653
D.2	Cauchy distribution for a few parameters	662
D.3	Exponential distribution for a few parameters	669
D.4	Fisher-Tippett distribution for a few parameters	676
D.5	Laplace distribution for a few parameters	684
D.6	Log normal distribution for a few parameters	691
D.7	Weibull distribution for a few parameters	711

List of Tables

1.1 Compared execution speed between C, Smalltalk and Java	23
3.1 Recommended polynomial interpolation algorithms	106
4.1 Algorithms using iterative processes	126
6.1 Comparison between integration algorithms	163
9.1 Public methods for probability density functions	357
9.2 Properties of the Normal distribution	370
9.3 Properties of the gamma distribution	379
10.1 Properties of the Fisher-Snedecor distribution	399
10.2 Covariance test of random number generator	401
10.3 Properties of the Student distribution	411
10.4 Properties of the χ^2 distribution	423
11.1 Optimizing algorithms presented in this book	498
B.1 Sample Smalltalk messages with their Java equivalent	632
B.2 Smalltalk assignment and equalities	633
D.1 Properties of the beta distribution	652
D.2 Properties of the Cauchy distribution	661
D.3 Properties of the exponential distribution	668
D.4 Properties of the Fisher-Tippett distribution	675
D.5 Properties of the Laplace distribution	683
D.6 Properties of the log normal distribution	690
D.7 Properties of the triangular distribution	697
D.8 Properties of the uniform distribution	704
D.9 Properties of the Weibull distribution	710

Chapter 1

Introduction

Science sans conscience n'est que ruine de l'âme.¹
François Rabelais

Teaching numerical methods was a major discipline of computer science at a time computers were only used by a very small amount of professionals such as physicists or operation research technicians. At that time most of the problems solved with the help of a computer were of numerical nature, such as matrix inversion or optimization of a function with many parameters.

With the advent of minicomputers, workstations and foremost, personal computers, the scope of problems solved with a computer shifted from the realm of numerical analysis to that of symbol manipulation. Recently, the main use of a computer has been centered on office automation. Major applications are word processors and database applications.

Today, computers are no longer working stand-alone. Instead they are sharing information with other computers. Large databases are getting commonplace. The wealth of information stored in large databases tends to be ignored, mainly because only few persons knows how to get access to it and an even fewer number know how to extract useful information. Recently people have started to tackle this problem under the buzzword data mining. In truth, data mining is nothing else than good old numerical data analysis performed by high-energy physicists with the help of computers. Of course a few new techniques are been invented recently, but most of the field now consists of rediscovering algorithms used in the past. This past goes back to the day Enrico Fermi used the ENIAC to perform phase shift analysis to determine the nature of nuclear forces.

The interesting point, however, is that, with the advent of data mining, numerical methods are back on the scene of information technologies.

¹Science without consciousness just ruins the soul.

1.1 Object-oriented paradigm and mathematical objects

In the recent years object-oriented programming — OOP for short — has been welcomed for its ability to represent objects from the real world - employees, bank accounts, etc. - inside a computer. Herein resides the formidable leverage of object-oriented programming. It turns out that this way of looking at OOP is somewhat overstated (as these lines are written). Objects manipulated inside an object-oriented program certainly do not behave like their real world counterparts. Computer objects are only models of those of the real world. The UML user guides goes further in stating that a model is a simplification of reality and we should emphasize that it is only that. OOP modeling is so powerful, however, that people tend to forget about it and only think in terms of real world objects.

An area where the behavior of computer objects nearly reproduces that of their real-world counterparts is mathematics. Mathematical objects are organized within hierarchies. For example, natural integers are included in integers (signed integers), which are included in rational numbers, themselves included in real numbers. Mathematical objects use polymorphism in that one operation can be defined on several entities. For example, addition and multiplication are defined for numbers, vectors, matrices, polynomials — as we shall see in this book — and many other mathematical entities. Common properties can be established as an abstract concept — a group e.g. — without the need to specify a concrete implementation. Such concepts can then be used to prove a given property for a concrete case. All this looks very similar to class hierarchies, methods and inheritance.

Because of these similarities OOP offers the possibility to manipulate mathematical objects in such a way that the boundary between real objects and their computer models becomes almost non-existent. This is no surprise since the structure of OOP objects is equivalent to that of mathematical objects². When dealing with numerical evaluations the equivalence between mathematical objects and computer objects is almost perfect. One notable difference remains, however, namely the finite size of the representation for non-integer number in a computer limiting the attainable precision. We shall address this important topic in section 1.3.2.

Most numerical algorithms have been invented long before the wide spread use of computers. Algorithms were designed to speed up human computation and therefore were constructed as to minimize the number of operations to be carried out by the human operator. Minimizing the number of operations is the best thing to do to speed up code execution.

One of the most heralded benefits of object-oriented programming is code reuse, a consequence, in principle, of the hierarchical structure and of inheritance. The last statement is pondered by "in principle" since, to date, code reuse of real world objects is still far from being common place.

For all these reasons, this book tries to convince you that using object-oriented programming for numerical evaluations can exploit the mathematical definitions to maximize code reuse between many different algorithms. Such a high degree of reuse yields very concise code. Not surprisingly, this code is quite efficient and, most impor-

²From the point of view of computer science OOP objects are considered as mathematical objects.

tantly, highly maintainable. Better than an argumentation, we show how to implement some numerical algorithms selected among those which we think are most useful for the areas where object-oriented software is used primarily: finance, medicine and decision support.

1.2 Object-oriented concepts in a nutshell

First let us define what is covered by the adjective object-oriented. Many software vendors are qualifying a piece of software object-oriented as soon as it contains things called objects, even though the behavior of those objects has little to do with object-orientation. For many programmers and most software journalists any software system offering a user interface design tool on which elements can be pasted on a window and linked to some events — even though most of these events are being restricted to user interactions — can be called object-oriented. There are several typical examples of such software, all of them having the prefix Visual in their names³. Visual programming is something entirely different from object-oriented programming.

Object-oriented is something different, not intrinsically linked with the user interface. Recently, object-oriented techniques applied to user interfaces have been widely exposed to the public, hence the confusion. There are 3 properties, which are considered essential for object-oriented software:

1. data encapsulation,
2. class hierarchy and inheritance,
3. polymorphism.

Data encapsulation is the fact that each object hides its internal structure from the rest of the system. Data encapsulation is in fact a misnomer since an object usually chooses to expose some of its data. I prefer to use the expression hiding the implementation, a more precise description of what is usually understood by data encapsulation. Hiding the implementation is a crucial point because an object, once fully tested, is guaranteed to work ever after. It ensures an easy maintainability of applications because the internal implementation of an object can be modified without impacting the application, as long as the public methods are kept identical.

Class hierarchy and inheritance is the keystone implementation of any object-oriented system. A class is a description of all properties of all objects of the same type. These properties can be structural (static) or behavioral (dynamic). Static properties are mostly described with instance variables. Dynamic properties are described by methods. Inheritance is the ability to derive the properties of an object from those of another. The class of the object from which another object is deriving its properties is called the superclass. A powerful technique offered by class hierarchy and inheritance is the overloading of some of the behavior of the superclass.

Polymorphism is the ability to manipulate objects from different classes, not necessarily related by inheritance, through a common set of methods. To take an example

³This is not to say that all products bearing a name with the prefix Visual are not object-oriented.

from this book, polynomials can have the same behavior than signed integers with respect to arithmetic operations: addition, subtraction, multiplication and division.

Most so-called object-oriented development tools (as opposed to languages) usually fail the inheritance and polymorphism requirements.

The code implementation of the algorithms presented in this book is given in two languages: Smalltalk and Java. Both languages are excellent object-oriented languages. I would strongly recommend people reading this book to consult the implementation sections of both languages regardless of their personal taste of language. First, I have made some effort to use of the best feature of each language. Second, each implementation has been made independently. The fact that the code of each implementation is different shows that there is indeed many ways to skin a cat, even when written by the same person. Thus, looking seriously at both implementations can be quite instructive for someone who wants to progress with the object-oriented paradigm.

1.3 Dealing with numerical data

The numerical methods exposed in this book are all applicable to real numbers. As noted earlier the finite representation of numbers within a computer limits the precision of numerical results, thereby causing a departure from the ideal world of mathematics. This section discusses issues related to this limitation.

1.3.1 Floating point representation

Currently mankind is using the decimal system⁴. In this system, however, most rational numbers and all irrational and transcendental numbers escape our way of representation. Numbers such as $1/3$ or π cannot be written in the decimal system other than approximately. One can chose to add more digits to the right of the decimal point to increase the precision of the representation. The true value of the number, however, cannot be represented. Thus, in general, a real number cannot be represented by a finite decimal representation. This kind of limitation has nothing to do with the use of computers. To go around that limitation, mathematicians have invented abstract representations of numbers, which can be manipulated in regular computations. This includes irreducible fractions ($1/3$ e.g.), irrational numbers ($\sqrt{2}$ e.g.), transcendental numbers (π and e the base of natural logarithms e.g.) and normal⁵ infinities ($-\infty$ and $+\infty$).

Like humans, computers are using a representation with a finite number of digits, but the digits are restricted to 0 and 1. Otherwise number representation in a computer can be compared to the way we represent numbers in writing. Compared to humans computers have the notable difference that the number of digits used to represent a number cannot be adjusted during a computation. There is no such thing as adding a few more decimal digits to increase precision. One should note that this is only an

⁴This is of course quite fortuitous. Some civilizations have opted for a different base. The Sumerians have used the base 60 and this habit has survived until now in our time units. The Maya civilization was using the base 20. The reader interested in the history of numbers ought to read the book of Georges Ifrah [Ifrah].

⁵Since Cantor, mathematicians have learned that there are many kinds of infinities. See for example reference [Gullberg].

implementation choice. One could think of designing a computer manipulating numbers with adjustable precision. Of course, some protection should be built in to prevent a number, such as $1/3$, to expand ad infinitum. Probably, such a computer would be much slower. Using digital representation — the word digital being taken in its first sense, that is, a representation with digits — no matter how clever the implementation⁶, most numbers will always escape the possibility of exact representation.

In present day computers, a floating-point number is represented as $m \times r^e$ where the radix r is a fixed number, generally 2. On some machines, however, the radix can be 10 or 16. Thus, each floating-point number is represented in two parts⁷: an integral part called the mantissa m and an exponent e . This way of doing is quite familiar to people using large quantities (astronomers e.g.) or studying the microscopic world (microbiologists e.g.). Of course, the natural radix for people is 10. For example, the average distance from earth to sun expressed in kilometer is written as 1.4959787×10^8 .

In the case of radix 2, the number 18446744073709551616 is represented as 1×2^{64} . Quite a short hand compared to the decimal notation! IEEE standard floating-point numbers use 24 bits for the mantissa (about 8 decimal digits) in single precision; they use 53 bits (about 15 decimal digits) in double precision.

One important property of floating-point number representation is that the relative precision of the representation — that is the ratio between the precision and the number itself — is the same for all numbers except, of course, for the number 0.

1.3.2 Rounding errors

To investigate the problem of rounding let us use our own decimal system limiting ourselves to 15 digits and an exponent. In this system, the number 2^{64} is now written as $184467440737095 \times 10^5$. Let us now perform some elementary arithmetic operations.

First of all, many people are aware of problems occurring with addition or subtraction. Indeed we have:

$$184467440737095 \times 10^5 + 300 = 184467440737095 \times 10^5.$$

More generally, adding or subtracting to 2^{64} any number smaller than 100000 is simply ignored by our representation. This is called a rounding error. This kind of rounding errors have the non-trivial consequence of breaking the associative law of addition. For example,

$$(1 \times 2^{64} + 1 \times 2^{16}) + 1 \times 2^{32} = 184467440780044 \times 10^5,$$

whereas

$$1 \times 2^{64} + (1 \times 2^{16} + 1 \times 2^{32}) = 184467440780045 \times 10^5.$$

In the two last expressions, the operation within the parentheses is performed first and rounded to the precision of our representation, as this is done within the floating point

⁶Symbolic manipulation programs do represent numbers as we do in mathematics. Such programs are not yet suited for quick numerical computation, but research in this area is still open.

⁷This is admittedly a simplification. In practice exponents of floating point numbers are offset to allow negative exponents. This does not change the point being made in this section, however.

arithmetic unit of a microprocessor⁸.

Other type of rounding errors may also occur with factors. Translating the calculation $1 \times 2^{64} \div 1 \times 2^{16} = 1 \times 2^{48}$ into our representation yields:

$$184467440737095 \times 10^5 \div 65536 = 2814744976710655.$$

The result is just off by one since $2^{48} = 2814744976710656$. This seems not to be a big deal since the relative error — that is the ratio between the error and the result — is about $3.6 \times 10^{-16}\%$.

Computing $1 \times 2^{48} - 1 \times 2^{64} \div 1 \times 2^{16}$, however, yields -1 instead of 0 . This time the relative error is 100% or infinite depending of what reference is taken to compute the relative error. Now, imagine that this last expression was used in finding the real (as opposed to complex) solutions of the second order equation:

$$2^{-16}x^2 + 2^{25}x + 2^{64} = 0.$$

The solutions to that equation are:

$$x = \frac{-2^{24} \pm \sqrt{2^{48} - 2^{64} \times 2^{-16}}}{2^{-16}}.$$

Here, the rounding error prevents the square root from being evaluated since $\sqrt{-1}$ cannot be represented as a floating point number. Thus, it has the devastating effect of transforming a result into something, which cannot be computed at all.

This simplistic example shows that rounding errors, however harmless they might seem, can have quite severe consequences. An interested reader can reproduce these results using the Smalltalk class described in appendix A.

In addition to rounding errors of the kind illustrated so far, rounding errors propagate in the computation. Study of error propagation is a wide area going out of the scope of this book. This section was only meant as a reminder that numerical results coming out from a computer must always be taken with a grain of salt. This only good advice to give at this point is to try the algorithm out and compare the changes caused by small variations of the inputs over their expected range. There is no shame in trying things out and you will avoid the ridicule of someone proving that your results are non-sense.

The interested reader will find a wealth of information about floating number representations and their limitations in the book of Knuth [Knuth 2]. The excellent article by David Goldberg — What every computer scientist should know about floating point arithmetic, published in the March 1991 issues of Computing Surveys — is recommend for a quick, but in-depth, survey. This article can be obtained from various WEB sites. Let us conclude this section with a quotation from Donald E. Knuth [Knuth 2].

Floating point arithmetic is by nature inexact, and it is not difficult to misuse it so that the computed answers consist almost entirely of "noise". One of the principal problems of numerical analysis is to determine how accurate the results of certain numerical methods will be.

⁸In modern days microprocessor, a floating point arithmetic unit actually uses more digits than the representation. These extra digits are called guard digits. Such difference is not relevant for our example.




Figure 1.1: Comparison of achieved precision

1.3.3 Real example of rounding error

To illustrate how rounding errors propagate, let us work our way through an example. Let us consider a numerical problem whose solution is known, that is, the solution can be computed exactly.

This numerical problem has one parameter, which measures the complexity of the data. Moreover data can be of two types: general data or special data. Special data have some symmetry properties, which can be exploited by the algorithm. Let us now consider two algorithms A and B able to solve the problem. In general algorithm B is faster than algorithm A.

The precision of each algorithm is determined by computing the deviation of the solution given by the algorithm with the value known theoretically. The precision has been determined for each set of data and for several values of the parameter measuring the complexity of the data.

Figure 1.1 shows the results. The parameter measuring the complexity is laid on the x -axis using a logarithmic scale. The precision is expressed as the negative of the decimal logarithm of the deviation from the known solution. The higher the value the better is the precision. The precision of the floating-point numbers on the machine used in this study corresponds roughly to 16 on the scale of Figure 1.1.

The first observation does not come as a surprise: the precision of each algorithm degrades as the complexity of the problem increases. One can see that when the algorithms can exploit the symmetry properties of the data the precision is better (curves for special data) than for general data. In this case the two algorithms are performing with essentially the same precision. Thus, one can choose the faster algorithm, namely algorithm B. For the general data, however, algorithm B has poorer and poorer precision as the complexity increases. For complexity larger than 50 algorithm B becomes totally

unreliable, to the point of becoming a perfect illustration of Knuth's quotation above. Thus, for general data, one has no choice but to use algorithm A.

Readers who do not like mysteries can go read section 8.5.3 where these algorithms are discussed.

1.3.4 Outsmarting rounding errors

In some instances rounding errors can be significantly reduced if one spends some time reconsidering how to compute the final solution. In this section we like to show an example of such thinking.

Consider the following second order equation, which must be solved when looking for the eigenvalues of a symmetric matrix (c.f. section 8.7):

$$t^2 + 2\alpha t - 1 = 0. \quad (1.1)$$

Without restricting the generality of the argumentation, we shall assume that α is positive. the problem is to find the the root of equation 1.1 having the smallest absolute value. You, reader, should have the answer somewhere in one corner of your brain, left over from high school mathematics:

$$t_{\min} = \sqrt{\alpha^2 + 1} - \alpha. \quad (1.2)$$

Let us now assume that α is very large, so large that adding 1 to α^2 cannot be noticed within the machine precision. Then, the smallest of the solutions of equation 1.1 becomes $t_{\min} \approx 0$, which is of course not true: the left hand side of equation 1.1 evaluates to -1 .

Let us now rewrite equation 1.1 for the variable $x = 1/t$. This gives the following equation:

$$x^2 - 2\alpha x - 1 = 0. \quad (1.3)$$

The smallest of the two solutions of equation 1.1 is the largest of the two solutions of equation 1.3. That is:

$$t_{\min} = \frac{1}{x_{\max}} = \frac{1}{\sqrt{\alpha^2 + 1} + \alpha}. \quad (1.4)$$

Now we have for large α :

$$t_{\min} \approx \frac{1}{2\alpha}. \quad (1.5)$$

This solution has certainly some rounding errors, but much less than the solution of equation 1.2: the left hand side of equation 1.1 evaluates to $\frac{1}{4\alpha^2}$, which goes toward zero for large α , as it should be.

1.3.5 Wisdom from the past

To close the subject of rounding errors, I would like to give the reader a different perspective. There is a big difference between a full control of rounding errors and giving a result with high precision. Granted, high precision computation is required to minimize rounding errors. On the other hand, one only needs to keep the rounding

errors under control to a level up to the precision required for the final results. There is no need to determine a result with non-sensical precision.

To illustrate the point, I am going to use a very old mathematical problem: the determination of the number π . The story is taken from the excellent book of Jan Gullberg, *Mathematics From the Birth of the Numbers* [Gullberg].

Around 300BC, Archimedes devised a simple algorithm to approximate π . For a circle of diameter d , one computes the perimeter p_{in} of a n -sided regular polygon inscribed within the circle and the perimeter p_{out} of a n -sided regular polygon whose sides are tangent to the same circle. We have:

$$\frac{p_{\text{in}}}{d} < \pi < \frac{p_{\text{out}}}{d}. \quad (1.6)$$

By increasing n , one can improve the precision of the determination of π . During the Antiquity and the Middle Age, the computation of the perimeters was a formidable task and an informal competition took place to find who could find the most precise approximation of the number π . In 1424, Jamshid Masud al-Kashi, a Persian scientist, published an approximation of π with 16 decimal digits. The number of sides of the polygons was 3×2^8 . This was quite an achievement, the last of its kind. After that, mathematicians discovered other means of expressing the number π .

In my eyes, however, Jamshid Masud al-Kashi deserves fame and admiration for the note added to his publication that places his result in perspective. He noted that the precision of his determination of the number π was such that,

the error in computing the perimeter of a circle with a radius 600 $\frac{1}{2}$ 000 times that of earth would be less than the thickness of a horse's hair.

The reader should know that the thickness of a horse's hair was a legal unit of measure in ancient Persia corresponding to roughly 0.7 mm. Using present-day knowledge of astronomy, the radius of the circle corresponding to the error quoted by Jamshid Masud al-Kashi is 147 times the distance between the sun and the earth, or about 3 times the radius of the orbit of Pluto, the most distant planet of the solar system.

As Jan Gullberg notes in his book, al-Kashi evidently had a good understanding of the meaninglessness of long chains of decimals. When dealing with numerical precision, you should ask yourself the following question:

Do I really need to know the length of Pluto's orbit to a third of the thickness of a horse's hair?

1.4 Finding the numerical precision of a computer

Object-oriented languages such as Smalltalk and Java give the opportunity to develop an application on one hardware platform and to deploy the application on other platforms running on different operating systems and hardware. It is a well-known fact that the marketing about Java was centered about the concept of Write Once Run Anywhere. What fewer people know is that this concept already existed for Smalltalk 10 years before the advent of Java.

Some numerical algorithms are carried until the estimated precision of the result becomes smaller than a given value, called the desired precision. Since an application can be executing on different hardware, the desired precision is best determined at run time.

The book of Press et al. [Press et al.] shows a clever code determining all the parameters of the floating-point representation of a particular computer. In this book we shall concentrate only on the parameters which are relevant for numerical computations. These parameters correspond to the instance variables of the object responsible to compute them. They are the following:

`radix` the radix of the floating-point representation, that is r .

`machinePrecision` the largest positive number which, when added to 1 yields 1.

`negativeMachinePrecision` the largest positive number which, when subtracted from 1 yields 1.

`smallestNumber` the smallest positive number different from 0.

`largestNumber` the largest positive number which can be represented in the machine.

`defaultNumericalPrecision` the relative precision, which can be expected for a general numerical computation.

`smallNumber` a number, which can be added to some value without noticeably changing the result of the computation.

Computing the radix r is done in two steps. First one computes a number equivalent of the machine precision (c.f. next paragraph) assuming the radix is 2. Then, one keeps adding 1 to this number until the result changes. The number of added ones is the radix.

The machine precision is computed by finding the largest integer n such that:

$$(1 + r^{-n}) - 1 \neq 0 \quad (1.7)$$

This is done with a loop over n . The quantity $\epsilon_+ = r^{-(n+1)}$ is the machine precision.

The negative machine precision is computed by finding the largest integer n such that:

$$(1 - r^{-n}) - 1 \neq 0 \quad (1.8)$$

Computation is made as for the machine precision. The quantity $\epsilon_- = r^{-(n+1)}$ is the negative machine precision. If the floating-point representation uses two-complement to represent negative numbers the machine precision is larger than the negative machine precision.

To compute the smallest and largest number one first compute a number whose mantissa is full. Such a number is obtained by building the expression $f = 1 - r \times \epsilon_-$. The smallest number is then computed by repeatedly dividing this value by the radix until the result produces an underflow. The last value obtained before an underflow occurs is the smallest number. Similarly, the largest number is computed by repeatedly multiplying the value f until an overflow occurs. The last value obtained before an overflow occurs is the largest number.

The variable `defaultNumericalPrecision` contains an estimate of the precision expected for a general numerical computation. For example, one should consider that two numbers a and b are equal if the relative difference between them is less than the default numerical machine precision. This value of the default numerical machine precision has been defined as the square root of the machine precision.

The variable `smallNumber` contains a value, which can be added to some number without noticeably changing the result of the computation. In general an expression of the type $\frac{0}{0}$ is undefined. In some particular case, however, one can define a value based on a limit. For example, the expression $\frac{\sin x}{x}$ is equal to 1 for $x = 0$. For algorithms, where such an undefined expression can occur⁹, adding a small number to the numerator and the denominator can avoid the division by zero exception and can obtain the correct value. This value of the small number has been defined as the square root of the smallest number that can be represented on the machine.

1.4.1 Computer numerical precision — General implementation

The computation of the parameters only needs to be executed once. We have introduced a specific class to hold the variables described earlier and made them available to any object.

Each parameter is computed using lazy initialization within the method bearing the same name as the parameter. Lazy initialization is used while all parameters may not be needed at a given time. Methods in charge of computing the parameters are all prefixed with the word `compute`.

1.4.2 Computer numerical precision — Smalltalk implementation

Listing 1.1 shows the class `DhbFloatingPointMachine` responsible of computing the parameters of the floating-point representation. This class is implemented as a singleton class because the parameters need to be computed once only. For that reason no code optimization was made and priority is given to readability.

The computation of the smallest and largest numbers uses exceptions¹⁰ to detect the underflow and the overflow.

The method `showParameters` can be used to print the values of the parameters onto the Transcript window.

Listing 1.1 Smalltalk code to find the machine precision

Class	DhbFloatingPointMachine
Subclass of	Object
Instance variable names:	<code>defaultNumericalPrecision</code> <code>radix</code> <code>machinePrecision</code> <code>negativeMachinePrecision</code> <code>smallestNumber</code> <code>largestNumber</code> <code>smallNumber</code> <code>largestExponentArgument</code>

⁹Of course, after making sure that the ratio is well defined numerically.

¹⁰The code is using the implementation of Visual Age For Smalltalk .

Class variable names: UniqueInstance

Class methods

new

```
UniqueInstance = nil
  ifTrue: [ UniqueInstance := super new].
^UniqueInstance
```

reset

```
UniqueInstance := nil.
```

Instance methods

computeLargestNumber

```
| zero one floatingRadix fullMantissaNumber |
zero := 0 asFloat.
one := 1 asFloat.
floatingRadix := self radix asFloat.
fullMantissaNumber := one - ( floatingRadix * self
                             negativeMachinePrecision).
largestNumber := fullMantissaNumber.
[ [ fullMantissaNumber := fullMantissaNumber * floatingRadix.
  largestNumber := fullMantissaNumber.
  true] whileTrue: [ ].
] when: ExAll do: [ :signal | signal exitWith: nil].
```

computeMachinePrecision

```
| one zero a b inverseRadix tmp x |
one := 1 asFloat.
zero := 0 asFloat.
inverseRadix := one / self radix asFloat.
machinePrecision := one.
[ tmp := one + machinePrecision.
  tmp - one = zero]
  whileFalse:[ machinePrecision := machinePrecision *
              inverseRadix].
```

computeNegativeMachinePrecision

```
| one zero floatingRadix inverseRadix tmp |
one := 1 asFloat.
```

```

zero := 0 asFloat.
floatingRadix := self radix asFloat.
inverseRadix := one / floatingRadix.
negativeMachinePrecision := one.
[ tmp := one - negativeMachinePrecision.
  tmp - one = zero]
  whileFalse:[ negativeMachinePrecision :=
    negativeMachinePrecision * inverseRadix].

```

computeRadix

```

| one zero a b tmp1 tmp2|
one := 1 asFloat.
zero := 0 asFloat.
a := one.
[ a := a + a.
  tmp1 := a + one.
  tmp2 := tmp1 - a.
  tmp2 - one = zero] whileTrue:[].
b := one.
[ b := b + b.
  tmp1 := a + b.
  radix := ( tmp1 - a) truncated.
  radix = 0 ] whileTrue: [].

```

computeSmallestNumber

```

| zero one floatingRadix inverseRadix fullMantissaNumber |
zero := 0 asFloat.
one := 1 asFloat.
floatingRadix := self radix asFloat.
inverseRadix := one / floatingRadix.
fullMantissaNumber := one - ( floatingRadix * self
  negativeMachinePrecision).
smallestNumber := fullMantissaNumber.
[ [ fullMantissaNumber := fullMantissaNumber * inverseRadix.
  smallestNumber := fullMantissaNumber.
  true] whileTrue: [ ].
] when: ExAll do: [ :signal | signal exitWith: nil].

```

defaultNumericalPrecision

```

defaultNumericalPrecision isNil
  ifTrue: [ defaultNumericalPrecision := self machinePrecision
    sqrt].

```



```
^defaultNumericalPrecision
```

```
largestExponentArgument
```

```
largestExponentArgument isNil
  ifTrue: [ largestExponentArgument := self largestNumber ln].
^largestExponentArgument
```

```
largestNumber
```

```
largestNumber isNil
  ifTrue: [ self computeLargestNumber].
^largestNumber
```

```
machinePrecision
```

```
machinePrecision isNil
  ifTrue: [ self computeMachinePrecision].
^machinePrecision
```

```
negativeMachinePrecision
```

```
negativeMachinePrecision isNil
  ifTrue: [ self computeNegativeMachinePrecision].
^negativeMachinePrecision
```

```
radix
```

```
radix isNil
  ifTrue: [ self computeRadix].
^radix
```

```
showParameters
```

```
Transcript cr; cr;
  nextPutAll: 'Floating-point machine parameters'; cr;
  nextPutAll: '-----';cr;
  nextPutAll: 'Radix: '.
self radix printOn: Transcript.
Transcript cr; nextPutAll: 'Machine precision: '.
self machinePrecision printOn: Transcript.
Transcript cr; nextPutAll: 'Negative machine precision: '.
self negativeMachinePrecision printOn: Transcript.
Transcript cr; nextPutAll: 'Smallest number: '.
```

```

self smallestNumber printOn: Transcript.
Transcript cr; nextPutAll: 'Largest number: '.
self largestNumber printOn: Transcript.

```

```

smallestNumber

```

```

smallestNumber isNil
  ifTrue: [ self computeSmallestNumber].
^smallestNumber

```

```

smallNumber

```

```

smallNumber isNil
  ifTrue: [ smallNumber := self smallestNumber sqrt].
^smallNumber

```

1.4.3 Computer numerical precision — Java implementation

Listing 1.2 shows the Java implementation, which uses static variables and static methods. These methods and variables have been implemented in the class `DhbMath`. This class is responsible of computing the parameters of the floating-point representation, but is also used to implements other mathematical utilities which are too simple to require the creation of a new class. The parameters of the floating-point representation need to be computed once only. Therefore, no code optimization was made; priority is given to readability.

Each parameter is retrieved with a method whose name is constructed with the standard prefix `get` followed with the name of the parameter. Methods in charge of computing the parameters are all prefixed with the word `compute`. The Java virtual machine does not raise an exception when floating underflow occurs. Thus, when computing the smallest number underflow is simply detected by testing that the result is equal to 0. When computing the largest number, the Java virtual machine traps overflows to change the result into infinity. Thus, the overflow is detected when the result becomes infinite.

The method `printParameters` can be used to print the values of the parameters to any output stream. For example, the following code prints the parameters onto the standard output of the console.

```
DhbFloatingPointMachine.printParameters( System.out)
```

Listing 1.2 Java code to find the machine precision

```
package DhbFunctionEvaluation;

import java.io.PrintStream;
/**
 * This class implements additional mathematical functions
 * and determines the parameters of the floating point representation.
 *
 * @author Didier H. Besset
 */
public final class DhbMath
{
    /**
     * Typical meaningful precision for numerical calculations.
     */
    static private double defaultNumericalPrecision = 0;
    /**
     * Typical meaningful small number for numerical calculations.
     */
    static private double smallNumber = 0;
    /**
     * Radix used by floating-point numbers.
     */
    static private int radix = 0;
    /**
     * Largest positive value which, when added to 1.0, yields 0.
     */
    static private double machinePrecision = 0;
    /**
     * Largest positive value which, when subtracted to 1.0, yields 0.
     */
    static private double negativeMachinePrecision = 0;
    /**
     * Smallest number different from zero.
     */
    static private double smallestNumber = 0;
    /**
     * Largest possible number
     */
    static private double largestNumber = 0;
    /**
     * Largest argument for the exponential
     */
    static private double largestExponentialArgument = 0;
    /**
     * Values used to compute human readable scales.
     */
}
```

```

private static final double scales[] = {1.25, 2, 2.5, 4, 5, 7.5, 8, 10};
private static final double semiIntegerScales[] = {2, 2.5, 4, 5, 7.5, 8, 10};
private static final double integerScales[] = {2, 4, 5, 8, 10};

private static void computeLargestNumber()
{
    double floatingRadix = getRadix();
    double fullMantissaNumber = 1.0d -
        floatingRadix * getNegativeMachinePrecision();
    while ( !Double.isInfinite( fullMantissaNumber ) )
    {
        largestNumber = fullMantissaNumber;
        fullMantissaNumber *= floatingRadix;
    }
}

private static void computeMachinePrecision()
{
    double floatingRadix = getRadix();
    double inverseRadix = 1.0d / floatingRadix;
    machinePrecision = 1.0d;
    double tmp = 1.0d + machinePrecision;
    while ( tmp - 1.0d != 0.0d )
    {
        machinePrecision *= inverseRadix;
        tmp = 1.0d + machinePrecision;
    }
}

private static void computeNegativeMachinePrecision()
{
    double floatingRadix = getRadix();
    double inverseRadix = 1.0d / floatingRadix;
    negativeMachinePrecision = 1.0d;
    double tmp = 1.0d - negativeMachinePrecision;
    while ( tmp - 1.0d != 0.0d )
    {
        negativeMachinePrecision *= inverseRadix;
        tmp = 1.0d - negativeMachinePrecision;
    }
}

private static void computeRadix()
{
    double a = 1.0d;
    double tmp1, tmp2;
    do { a += a;
        tmp1 = a + 1.0d;
        tmp2 = tmp1 - a;
    } while ( tmp1 != tmp2 );
}

```

```

        } while ( tmp2 - 1.0d != 0.0d);
double b = 1.0d;
while ( radix == 0)
{
    b += b;
    tmp1 = a + b;
    radix = (int) ( tmp1 - a);
}
}
private static void computeSmallestNumber()
{
    double floatingRadix = getRadix();
    double inverseRadix = 1.0d / floatingRadix;
    double fullMantissaNumber = 1.0d - floatingRadix * getNegativeMachinePrecision();
    while ( fullMantissaNumber != 0.0d )
    {
        smallestNumber = fullMantissaNumber;
        fullMantissaNumber *= inverseRadix;
    }
}
public static double defaultNumericalPrecision()
{
    if ( defaultNumericalPrecision == 0 )
        defaultNumericalPrecision = Math.sqrt( getMachinePrecision());
    return defaultNumericalPrecision;
}
/**
 * @return boolean true if the difference between a and b is
 * less than the default numerical precision
 * @param a double
 * @param b double
 */
public static boolean equal( double a, double b)
{
    return equal( a, b, defaultNumericalPrecision());
}
/**
 * @return boolean true if the relative difference between a and b
 * is less than precision
 * @param a double
 * @param b double
 * @param precision double
 */
public static boolean equal( double a, double b, double precision)
{
    double norm = Math.max( Math.abs(a), Math.abs( b));

```

```

    return norm < precision || Math.abs( a - b ) < precision * norm;
}
public static double getLargestExponentialArgument()
{
    if ( largestExponentialArgument == 0 )
        largestExponentialArgument = Math.log(getLargestNumber());
    return largestExponentialArgument;
}
/**
 * (c) Copyrights Didier BESSET, 1999, all rights reserved.
 */
public static double getLargestNumber()
{
    if ( largestNumber == 0 )
        computeLargestNumber();
    return largestNumber;
}
public static double getMachinePrecision()
{
    if ( machinePrecision == 0 )
        computeMachinePrecision();
    return machinePrecision;
}
public static double getNegativeMachinePrecision()
{
    if ( negativeMachinePrecision == 0 )
        computeNegativeMachinePrecision();
    return negativeMachinePrecision;
}
public static int getRadix()
{
    if ( radix == 0 )
        computeRadix();
    return radix;
}
public static double getSmallestNumber()
{
    if ( smallestNumber == 0 )
        computeSmallestNumber();
    return smallestNumber;
}
public static void printParameters( PrintStream printStream)
{
    printStream.println( "Floating-point machine parameters");
    printStream.println( "-----");
    printStream.println( " ");
}

```

```

        printStream.println( "radix = "+ getRadix());
        printStream.println( "Machine precision = "
                               + getMachinePrecision());
        printStream.println( "Negative machine precision = "
                               + getNegativeMachinePrecision());
        printStream.println( "Smallest number = "+ getSmallestNumber());
        printStream.println( "Largest number = "+ getLargestNumber());
        return;
    }
    public static void reset()
    {
        defaultNumericalPrecision = 0;
        smallNumber = 0;
        radix = 0;
        machinePrecision = 0;
        negativeMachinePrecision = 0;
        smallestNumber = 0;
        largestNumber = 0;
    }
    /**
     * This method returns the specified value rounded to
     * the nearest integer multiple of the specified scale.
     *
     * @param value number to be rounded
     * @param scale defining the rounding scale
     * @return rounded value
     */
    public static double roundTo( double value, double scale)
    {
        return Math.round( value / scale) * scale;
    }
    /**
     * Round the specified value upward to the next scale value.
     * @param the value to be rounded.
     * @param a flag specified whether integer scale are used, otherwise double scale is used.
     * @return a number rounded upward to the next scale value.
     */
    public static double roundToScale( double value, boolean integerValued)
    {
        double[] scaleValues;
        int orderOfMagnitude = (int) Math.floor( Math.log( value) / Math.log( 10.0));
        if ( integerValued )
        {
            orderOfMagnitude = Math.max( 1, orderOfMagnitude);
            if ( orderOfMagnitude == 1)
                scaleValues = integerScales;
        }
    }

```

```

        else if ( orderOfMagnitude == 2)
            scaleValues = semiIntegerScales;
        else
            scaleValues = scales;
    }
    else
        scaleValues = scales;
    double exponent = Math.pow( 10.0, orderOfMagnitude);
    double rValue = value / exponent;
    for ( int n = 0; n < scaleValues.length; n++)
    {
        if ( rValue <= scaleValues[n])
            return scaleValues[n] * exponent;
    }
    return exponent;    // Should never reach here
}
/**
 * (c) Copyrights Didier BESSET, 1999, all rights reserved.
 * @return double
 */
public static double smallNumber()
{
    if ( smallNumber == 0 )
        smallNumber = Math.sqrt( getSmallestNumber());
    return smallNumber;
}
}

```

Note: The two methods equal — one with 2 arguments, one with 3 — are discussed in section 1.5.2.

1.5 Comparing floating point numbers

It is very surprising to see how frequently questions about the lack of equality between two floating-point numbers are posted on the Smalltalk and Java electronic discussion groups. As we have seen in section 1.3.2 one should always expect the result of two different computations that should have yielded the same number from a mathematical standpoint to be different using a finite numerical representation. Somehow the computer courses are not giving enough emphasis about floating-point numbers.

So, how should you check the equality of two floating-point numbers?
The practical answer is: thou shalt not!

As you will see, the algorithms in this book only compare numbers, but never check for equality. If you cannot escape the need for a test of equality, however, the best solution

is to create methods to do this. Since the floating-point representation is keeping a constant relative precision, comparison must be made using relative error. Let a and b be the two numbers to be compared. One should build the following expression:

$$\epsilon = \frac{|a - b|}{\max(|a|, |b|)} \quad (1.9)$$

The two numbers can be considered equal if ϵ is smaller than a given number ϵ_{\max} . If the denominator of the fraction on equation 1.9 is less than ϵ_{\max} , then the two numbers can be considered as being equal. For lack of information on how the numbers a and b have been obtained, one uses for ϵ_{\max} the default numerical precision defined in section 1.4. If one can determine the precision of each number, then the method `relativelyEqual` can be used.

1.5.1 Comparing floating point numbers — Smalltalk code

In Smalltalk this means adding a new method to the class `Number` as shown in Listing 1.3.

Listing 1.3 Comparison of floating point numbers in Smalltalk

```

Class      Number
Subclass of Magnitude

    Instance methods
equalsTo: aNumber

    ^self relativelyEqualsTo: aNumber upTo: DhbFloatingPointMachine
    new defaultNumericalPrecision

relativelyEqualsTo: aNumber upTo: aSmallNumber

    | norm |
    norm := self abs max: aNumber abs.
    ^norm <= DhbFloatingPointMachine new defaultNumericalPrecision
    or: [ (self - aNumber) abs < ( aSmallNumber * norm)]

```

1.5.2 Comparing floating point numbers — Java code

In Java the methods have been added to the class `DhbMath`. The code of that class is shown in listing 1.2.

1.6 Speed consideration

Some people may think that implementing numerical methods for object-oriented languages such as Smalltalk or Java is just a waste of time. Those languages are notoriously slow or so they think.

First of all, things should be put in perspective with other actions performed by the computer. If a computation does not take longer than the time needed to refresh a screen, it does not really matter if the application is interactive. For example, performing a least square fit to a histogram in Smalltalk and Java and drawing the resulting fitted function is usually hardly perceptible to the eye on a personal computer using a 200MHz Pentium. Thus, even though a C version runs 10 times faster, it does not make any difference for the end user. The main difference comes, however, when you need to modify the code. Object-oriented software is well known for its maintainability. As 80% of the code development is spent in maintenance this aspect should first be considered.

Table 1.1 shows measured speed of execution for some of the algorithms exposed in this book. Timing was done on a personal computer equipped with a Pentium II clocked at 200MHz and running Windows NT workstation 4.0. The C code used is the code of [Press et al.] compiled with the C compiler Visual C++ 4.0 from Microsoft Corporation. The time needed to allocate memory for intermediate results was included in the measurement of the C code, otherwise the comparison with object-oriented code would not be fair. The Smalltalk code was run under version 4.0 of Visual Age for Smalltalk from IBM Corporation using the ENVY benchmark tool provided. The Java code was run under version 2.0 of Visual Age for Java from IBM Corporation. Elapsed time were measured by repeating the measured evaluation a sufficient number of time so that the error caused by the CPU clock is less than the last digit shown in the final result.

Table 1.1: Compared execution speed between C, Smalltalk and Java

Operation	Units	C	Smalltalk	Java
Polynomial 10 th degree	msec.	1.1	27.7	9.0
Neville interpolation (20 points)	msec.	0.9	11.0	0.8
LUP matrix inversion (100 × 100)	sec.	3.9	22.9	1.0

One can see that object-oriented code is quite efficient, especially when it comes to complex algorithms: good object-oriented code can actually beat up C code.

My early tests with Java, a couple of years ago, were showing that Java was 5-10 times slower than C. One can see that vendors did a great job in optimizing the generated code and in accelerating the virtual machine. I would like to see the same efforts going in optimizing Smalltalk. The spectacular improvement of Java shows that it is possible. Actually, my early tests made with Visual Smalltalk from Digitalk Inc.¹¹ were 5 times better.

Today admittedly, I would not use Smalltalk to build a structural analysis program, but Java would certainly be a contender. Nevertheless, I have successfully build data

¹¹Unfortunately, the future of Visual Smalltalk now owned by Cincom Inc. is quite uncertain at this time of writing.

mining Smalltalk applications using all the code¹² presented in this book. These applications were not perceived as slow by the end user since most of the computer time was spent drawing the data.

1.6.1 Smalltalk particular

Smalltalk has an interesting property: a division between two integers is by default kept as a fraction. This prevents rounding errors. For example, the multiplication of a matrix of integer numbers with its inverse always yields an exact identity matrix. (c.f. section 8.3 for definitions of these terms).

There is, however, a price to pay for the perfection offered by fractions. When using fractions, the computing time often becomes prohibitive. Resulting fractions are often composed of large integers. This slows down the computing. In the case of matrix inversion, this results in an increase in computing time by several orders of magnitude.

For example, one of my customers was inverting a 301×301 matrix with the code of section 8.3. The numbers used to build the matrix were obtained from a measuring device (using an ADC) and were thus integers. The inversion time was over 2 hours¹³. After converting the matrix components to floating numbers the inversion time became less than 30 seconds!

If you are especially unlucky you may run out of memory when attempting to store a particularly long integer. Thus, it is always a good idea to use floating¹⁴ numbers instead of fractions unless absolute accuracy is your primary goal. My experience has been that using floating numbers speeds up the computation by at least an order of magnitude. In the case of complex computations such as matrix inversion or least square fit this can become prohibitive.

1.7 Conventions

Equations presented in this book are using standard international mathematical notation as described in [Knuth1]. Each section is trying to make a quick derivation of the concepts needed to fully understand the mathematics behind the scene. For readers in a hurry, the equations used by the implementation are flagged as the following sample equation:

$$\ln ab = \ln a + \ln b. \quad (1.10)$$

Main equation⇒

When such an equation is encountered, the reader is sure that the expression is implemented in the code.

In general the code presented in this book adheres to conventions widely used in each language. Having said that, there are a few instances where we have departed from the widely used conventions.

¹²I want to emphasize here that all the code of this book is real code, which I have used personally in real applications.

¹³This particular customer was a very patient person!

¹⁴In most available Smalltalk versions the class Float corresponds to floating numbers with double precision. VisualWorks makes the difference between Float and Double




Figure 1.2: A typical class diagram

1.7.1 Class diagrams

When appropriate a class diagram is shown at the beginning of each chapter. This diagram shows the hierarchy of the classes described in the chapter and eventually the relations with classes of other chapters. The diagrams are drawn using the conventions of the book on design patterns [Gamma et al.].

Figure 1.2 shows a typical class diagram. A rectangular box with 2 or 3 parts represents a class. The top part contains the name of the class in bold face. If the class is an abstract class the name is shown in italic bold face. In figure 1.2 the classes **RelatedClass**, **MySubClass1** and **MySubclass2** are concrete classes; **MyAbstractClass** is an abstract class. The second part of the class box contains a list of the public instance methods. The name of an abstract method is written in italic, for example *abstractMethod3* in the class **MyAbstractClass** of figure 1.2. The third part of the class box, if any, contains the list of all instance variables. If the class does not have any instance variable the class box only consists of 2 parts, for example the class **MySubClass1** of figure 1.2.

A vertical line with a triangle indicates class inheritance. If there are several subclasses the line branches at the triangle, as this is the case in figure 1.2. A horizontal line beginning with a diamond (UML aggregation symbol) indicates the class of an instance variable. For example, figure 1.2 indicates that the instance variable *instanceVariable1* of the class **MyAbstractClass** must be an instance of the class **RelatedClass**. The diamond is black if the instance variable is a collection of instances of the class. A class within a rectangle with rounded corner represents a class already discussed in an earlier chapter; the reference to the chapter is written below the class name. Class **ClassInOtherChapter** in figure 1.2 is such a class. To save space, we have used the Java class names and the Smalltalk method names. It is quite easy to identify methods needing parameters when one uses Smalltalk method names: a colon in the middle or at the end of the method name indicates a parameter. Please refer to appendix B for more details on Smalltalk methods.

1.7.2 Smalltalk code

Most of the Smalltalk systems do not support name spaces. As a consequence, it has become a convention to prefix all class names with 3-letter code identifying the origin of the code. In this book the names of the Smalltalk classes are all prefixed with the author's initials.

There are several ways to store constants needed by all instances of a class. One way is to store the constants in class variables. This requires each class to implement an initialization method, which sets the desired values into the class variables. Another way is to store the constants in a pool dictionary. Here also an initialization method is required. In my opinion pool dictionaries are best used for texts, as they provide a convenient way to change all text from one language to another. Sometimes the creation of a singleton object is used. This is especially useful when the constants are installation specific and, therefore, must be determined at the beginning of the application's execution, such as the precision of the machine (c.f. section 1.4). Finally constants which are not likely to change can be stored in the code. This is acceptable as long as this is done at a unique place. In this book most constants are defined in class methods.

By default a Smalltalk method returns self. For initialization methods, however, we write this return explicitly (`^self`) to ease reading. This adheres to the intention revealing patterns of Kent Beck [Beck].

In [Alpert et al.] it is recommended to use the method name default to implement a singleton class. In this book this convention is not followed. In Smalltalk, however, the normal instance creation method is `new`. Introducing a method default for singleton classes has the effect of departing from this more ancient convention. In fact, requiring the use of default amounts to reveal to the client the details of implementation used by the class. This is in clear contradiction with the principle of hiding the implementation to the external world. Thus, singleton classes in all code presented in this book are obtained by sending the method `new` to the class. A method named default is reserved for the very semantic of the word default: the instance returned by these methods is an instance initialized with some default contents, well specified. Whether or not the instance is a singleton is not the problem of the client application.

1.7.3 Java code

Java forces code to be maintained as files. A file may contain several classes but a class cannot be defined across several files. In this book classes are always presented in a single listing. A note placed at the end of each listing mentioned eventual methods discussed in a subsequent sections.

Java supports name space for class names. The unit of name space is called a package. The names of the packages containing the classes described in this book are all prefixed with the author's initials as recommended by the Java designers. The class names are not prefixed.

The designers of the Java language have chosen to conform to the C syntax as much as possible. In particular, pre- and post-increment operators are available. In the code presented in this book the use of these operators are restricted to indices. If a variable

is used for counting, the assignment with coupled operation is used instead of self incrementation. For example, incrementing a counter is written as:

$$\text{counter} += 1; \quad (1.11)$$

whereas incrementing an index is written as

$$\text{index}++; \quad (1.12)$$

This adheres to the intention revealing patterns of Kent Beck [Beck]. As the two statements are strictly equivalent, the compiler is expected to generate the most efficient byte code.

1.8 Road map

This last section of the introduction describes the road map of the algorithms discussed in the book chapter by chapter. Figure 1.3 shows a schematic view of the major classes discussed in this book together with their dependency relations. In this figure, abstract classes are represented with an ellipse, concrete classes with a rectangle. Dependencies between the classes are represented by lines going from one class to another; the dependent class is always located below. Chapters where the classes are discussed are drawn as grayed rectangles with rounded corners. Hopefully the reader will not be scared by the complexity of the figure. Actually, the figure should be more complex as the classes *Vector* and *Matrix* are used by most objects located in chapters 8 and following. To preserve the readability of figure 1.3 the dependency connections for these two classes have been left out.

Chapter 2 presents a general representation of mathematical functions. Examples are shown. A concrete implementation of polynomial is discussed. Finally three library functions are given: the error function, the gamma function and the beta function.

Chapter 3 discusses interpolation algorithms. A discussion explains when interpolation should be used and which algorithm is more appropriate to which data.

Chapter 4 presents a general framework for iterative process. It also discusses a specialization of the framework to iterative process with a single numerical result. This framework is widely used in the rest of the book.

Chapter 5 discusses two algorithms to find the zeroes of a function: bisection and Newton's zero finding algorithms. Both algorithms use the general framework of chapter 4.

Chapter 6 discusses several algorithms to compute the integral of a function. All algorithms are based on the general framework of chapter 4. This chapter also uses an interpolation technique from chapter 3.

Chapter 7 discusses the specialization of the general framework of chapter 4 to the computation of infinite series and continued fractions. The incomplete gamma function and incomplete beta function are used as concrete examples to illustrate the technique.

Chapter 8 presents a concrete implementation of vector and matrix algebra. It also discusses algorithms to solve systems of linear equations. Algorithms to compute matrix

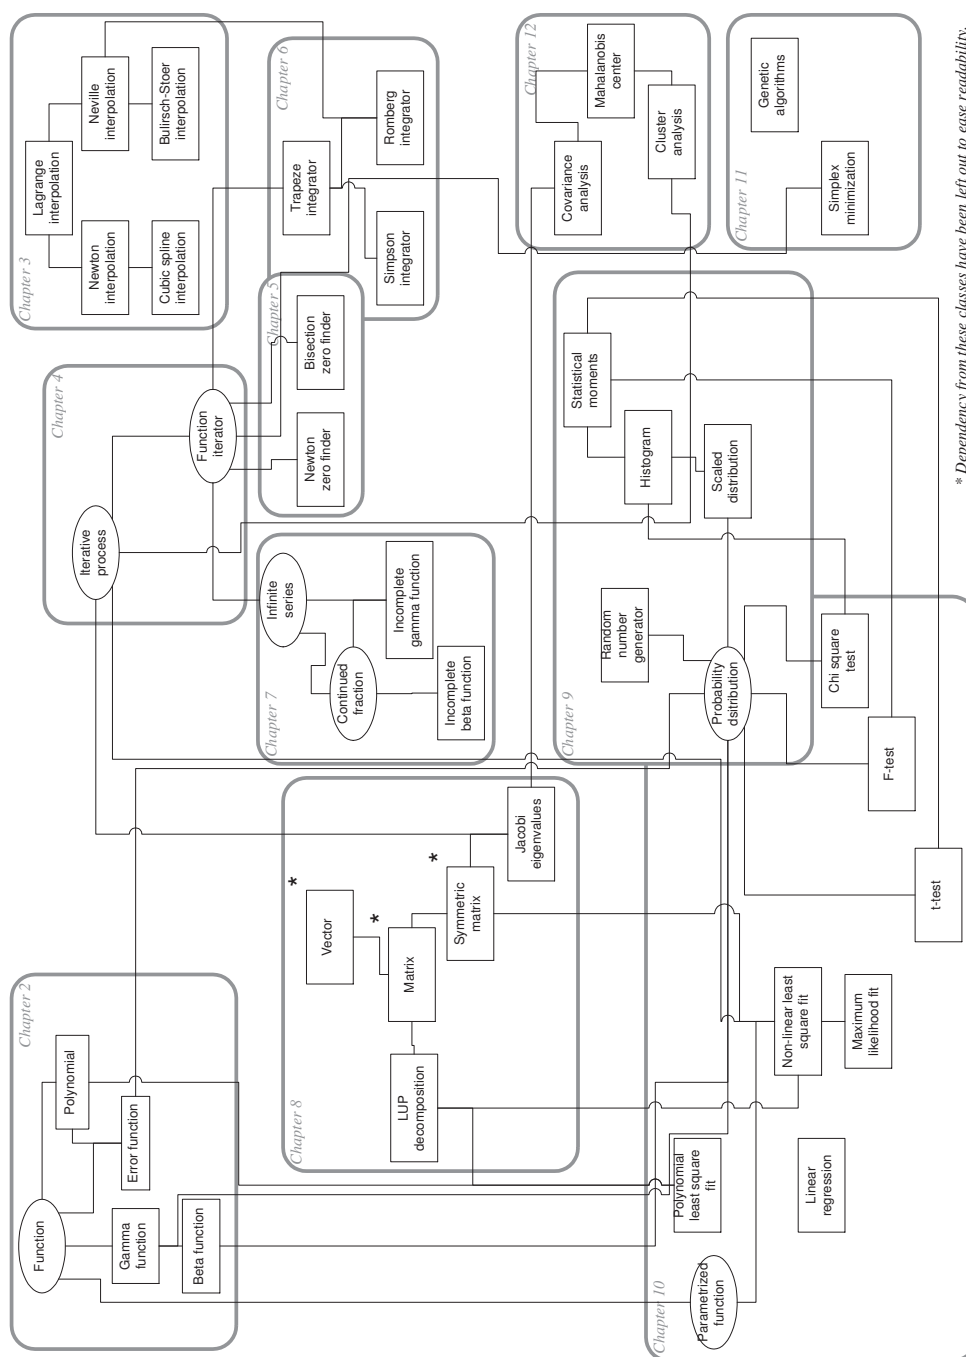


Figure 1.3: Book road map

inversion and the finding of eigenvalues and eigenvectors are exposed. Elements of this chapter are used in other part of this book.

Chapter 9 presents tools to perform statistical analysis. Random number generators are discussed. We give an abstract implementation of a probability distribution with concrete example of the most important distributions. The implementation of other distributions is given in appendix. This chapter uses techniques from chapters 2, 5 and 6.

Chapter 10 discussed the test of hypothesis and estimation. It gives an implementation of the t- and F-tests. It presents a general framework to implement least square fit and maximum likelihood estimation. Concrete implementations of least square fit for linear and polynomial dependence are given. A concrete implementation of the maximum likelihood estimation is given to fit a probability distribution to a histogram. This chapter uses techniques from chapter 2, 4, 8 and 9.

Chapter 11 discusses some techniques used to maximize or minimize a function: classical algorithms (simplex, hill climbing) as well as new ones (genetic algorithms). All these algorithms are using the general framework for iterative process discussed in chapter 4.

Chapter 12 discusses the modern data mining techniques: correlation analysis, cluster analysis and neural networks. A couple of methods invented by the author are also discussed. This chapter uses directly or indirectly techniques from all chapters of this book.

Chapter 2

Function evaluation

Qu'il n'y ait pas de réponse n'excuse pas l'absence de questions.¹
Claude Roy

Many mathematical functions used in numerical computation are defined by an integral, by a recurrence formula or by a series expansion. While such definitions can be useful to a mathematician, they are usually quite complicated to implement on a computer. For one, not every programmer knows how to evaluate an integral numerically². Then, there is the problem of accuracy. Finally, the evaluation of the function as defined mathematically is often too slow to be practical.

Before computers were heavily used, however, people had already found efficient ways of evaluating complicated functions. These methods are usually precise enough and extremely fast. This chapter exposes several functions that are important for statistical analysis. The Handbook of Mathematical Functions by Abramovitz and Stegun [Abramovitz & Stegun] contains a wealth of such function definitions and describes many ways of evaluating them numerically. Most approximations used in this chapter have been taken from this book.

This chapter opens on general considerations on how to implement the concept of function. Then, polynomials are discussed as an example of concrete function implementation. The rest of this chapter introduces three classical functions: the error function, the gamma function and the beta function. We shall use these functions in chapters 9 and 10. Because these functions are fundamental functions used in many areas of mathematics they are implemented as library functions — such as a sine, log or exponential — instead of using the general function formalism described in the first section.

The two language implementations are quite different because the concepts are implemented at a very low level. Figure 2.1 shows the diagram of the Smalltalk classes described in this chapter. Here we have used special notations to indicate that the

¹The absence of answer does not justify the absence of question.

²The good news is that they will if they read the present book (c.f. chapter 6).

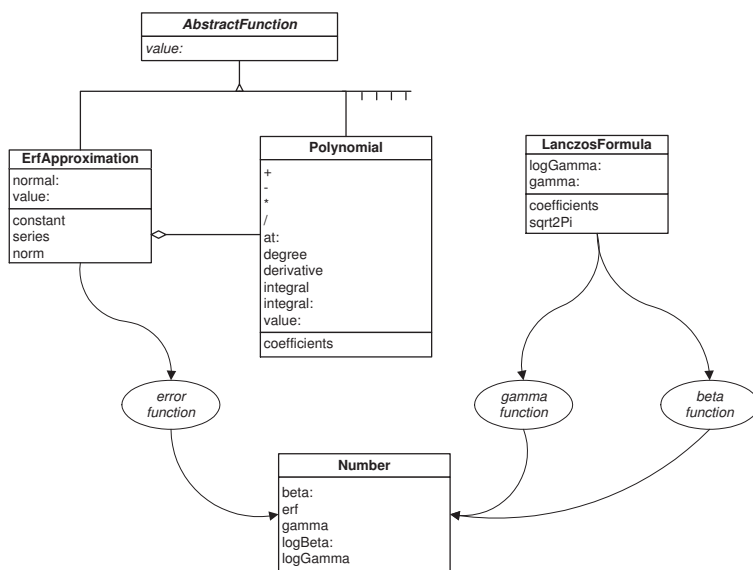


Figure 2.1: Smalltalk classes related to functions

functions are implemented as library functions. The functions are represented by oval and arrows shows which class is used to implement a function for the class **Number**.

Figure 2.2 shows the diagram of the Java classes described in this chapter. The strangeness of the box for the **NormalDistribution** class indicates that only the variables and methods relevant to this chapter are shown here. Both variables and methods are static.

2.1 Function concept

A mathematical function is an object associating a value to a variable. If the variable is a single value one talks about a one variable function. If the variable is an array of values one talks about a multi-variable function. Other types of variables are possible but will not be covered in this book.

We shall assume that the reader is familiar with elementary concepts about functions, namely derivatives and integrals. We shall concentrate mostly on implementation issues.

2.1.1 Function — Smalltalk implementation

A function in Smalltalk can be readily implemented with a block closure. Block closures in Smalltalk are treated like objects; thus, they can be manipulated as any other objects. For example the one variable function defined as:

$$f(x) = \frac{1}{x}, \quad (2.1)$$

Figure 2.1 with the box **AbstractFunction** grayed.

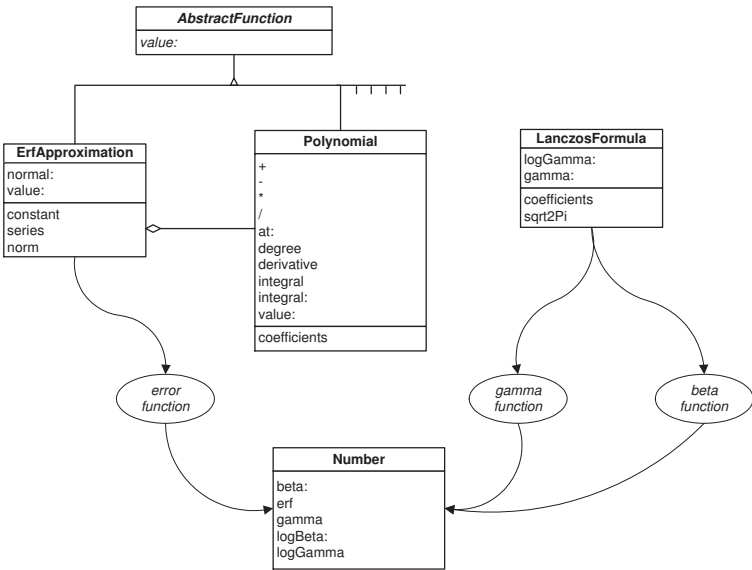


Figure 2.2: Java classes related to functions

can be implemented in Smalltalk as:

$$f := [:x | 1 / x]. \tag{2.2}$$

Evaluation of a block closure is supplied by the method value:. For example, to compute the inverse of 3, one writes:

$$f \text{ value: } 3. \tag{2.3}$$

If the function is more complex a block closure may not be the best solution to implement a function. Instead a class can be created with some instance variables to hold any constants and/or partial results. In order to be able to use functions indifferently implemented as block closures or as classes, one uses polymorphism. Each class implementing a function must implement a method value:. Thus, any object evaluating a function can send the same message selector, namely value:, to the variable holding the function.

To evaluate a multi-variable function, the argument of the method value: is an Array or a vector (c.f. section 8.1). Thus, in Smalltalk multi-variable functions can follow the same polymorphism as for one-variable functions.

2.1.2 Function — Java implementation

The situation in Java is a little more complex. Unlike Smalltalk, Java does not have a concept of block closure. Therefore, code cannot be manipulated as an object. In addition, strong typing brings us to make different declarations for one-variable and multi-variable functions.

Figure 2.2 with the box OneVariableFunction grayed.

One way to keep the same generality as in Smalltalk would be to use a general object declaration and use casting. However, this approach departs from the philosophy of a strongly typed language, which is central to the design of Java. It was therefore not considered.

Java, however, has introduced a very convenient concept, the interface. In this case we need two of them: an interface to declare one variable functions and one to declare multi variable functions. Thus, a class interacting with a function can declare the function - as an instance variable or a method argument - with the required interface.

Listing 2.1 shows the implementation of the Java interface `OneVariableFunction`.

Listing 2.1 Java implementation of the interface `OneVariableFunction`

```
package DhbInterfaces;

/**
 * OneVariableFunction is an interface for mathematical functions of
 * a single variable, that is functions of the form  $f(x)$ .
 *
 * @author Didier H. Besset
 */
public interface OneVariableFunction
{

    /**
     * Returns the value of the function for the specified variable value.
     */
    public double value( double x);
}
```

If the function is simple enough, it can be implemented as an inner class³. For example, the definition of the inverse function as an inner class is:

Code example 2.1

```
new OneVariableFunction() {
    public double value( double x)
    {return 1/x;}}
```

A more complex function must be implemented as a full-fledged class. A function class can be the subclass of any class as long as it implements the interface. There will be numerous examples of such function classes in the rest of this book.

A further example of the use of the `OneVariableFunction` interface is shown in Listing 2.2 showing the implementation of the Java class `FunctionDerivative` using the interface `OneVariableFunction`. This class allows the computation of the derivative

of any function of one variable. The evaluation of the derivative uses the following mathematical approximation:

$$\frac{df(x)}{dx} = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x - \epsilon)}{2\epsilon}. \quad (2.4)$$

This symmetrical definition of the derivative guaranties a higher accuracy than the usual one-sided definition. Of course, evaluating the derivative with this class is slower than coding the derivative explicitly as it requires two computations of the original function for every evaluation. But it is better than nothing, if the function whose derivative is needed is too complicated to be expressed analytically. Last, but not least, this class can be used to check whether the supplied derivatives are indeed correct.

Example of use of the class `FunctionDerivative` is shown in section 5.3.2.

Listing 2.2 Java implementation of a generic derivative evaluation

package DhbFunctionEvaluation;

```
import DhbInterfaces.OneVariableFunction;
/**
 * Evaluate an approximation of the derivative of a given function.
 *
 * @author Didier H. Besset
 */
public final class FunctionDerivative implements OneVariableFunction
{
    /**
     * Function for which the derivative is computed.
     */
    private OneVariableFunction f;
    /**
     * Relative interval variation to compute derivative.
     */
    private double relativePrecision = 0.0001;

    /**
     * Constructor method.
     * @param func the function for which the derivative is computed.
     */
    public FunctionDerivative( OneVariableFunction func)
    {
        this( func, 0.000001);
    }
}
```

³Inner classes are a feature of Java 1.1 and higher versions.

```

/**
 * Constructor method.
 * @param func the function for which the derivative is computed.
 * @param precision the relative step used to compute the derivative.
 */
public FunctionDerivative( OneVariableFunction func, double precision)
{
    f = func;
    relativePrecision = precision;
}
/**
 * Returns the value of the function's derivative
 * for the specified variable value.
 */
public double value( double x)
{
    double x1 = x == 0 ? relativePrecision
                : x * ( 1 + relativePrecision);
    double x2 = 2 * x - x1;
    return (f.value(x1) - f.value(x2)) / (x1 - x2);
}
}

```

2.2 Polynomials

Polynomials are quite important in numerical methods because they are often used in approximating functions. For example, section 2.3 shows how the error function can be approximated with the product of normal distribution times a polynomial.

Polynomials are also useful in approximating functions, which are determined by experimental measurements in the absence of any theory on the nature of the function. For example, the output of a sensor detecting a coin is dependent on the temperature of the coin mechanism. This temperature dependence cannot be predicted theoretically because it is a difficult problem. Instead, one can measure the sensor output at various controlled temperatures. These measurements are used to determine the coefficients of a polynomial reproducing the measured temperature variations. The determination of the coefficients is performed using a polynomial least-square fit (c.f. section 10.8). Using this polynomial the correction for a given temperature can be evaluated for any temperature within the measured range.

The implementations in both languages are discussed in great details. The reader is advised to read carefully both implementation sections as many techniques are introduced at this occasion. Later on those techniques will be mentioned with no further explanations.

2.2.1 Mathematical definitions

A polynomial is a special mathematical function whose value is computed as follows:

$$P(x) = \sum_{k=0}^n a_k x^k. \quad (2.5)$$

n is called the degree of the polynomial. For example, the second order polynomial

$$x^2 - 3x + 2 \quad (2.6)$$

represents a parabola crossing the x -axis at points 1 and 2 and having a minimum at $x = 2/3$. The value of the polynomial at the minimum is $-1/4$.

In equation 2.5 the numbers a_0, \dots, a_n are called the coefficients of the polynomial. Thus, a polynomial can be represented by the array $\{a_0, \dots, a_n\}$. For example, the polynomial of equation 2.6 is represented by the array $\{2, -3, 1\}$.

Evaluating equation 2.5 as such is highly inefficient since one must raise the variable to an integral power at each term. The required number of multiplication is of the order of n^2 . There is of course a better way to evaluate a polynomial. It consists of factoring out x before the evaluation of each term⁴. The following formula shows the resulting expression:

$$(x) = a_0 + x \{a_1 + x [a_2 + x (a_3 + \dots)]\} \quad (2.7)$$

←Main equation

Evaluating the above expression now requires only multiplications. The resulting algorithm is quite straightforward to implement. Expression 2.7 is called Horner's rule because it was first published by W.G. Horner in 1819. 150 years earlier, however, Isaac Newton was already using this method to evaluate polynomials.

In section 5.3 we shall require the derivative of a function. For polynomials this is rather straightforward. The derivative is given by:

$$\frac{dP(x)}{dx} = \sum_{k=1}^n k a_k x^{k-1}. \quad (2.8)$$

Thus, the derivative of a polynomial with n coefficients is another polynomial, with $n - 1$ coefficients⁵ derived from the coefficients of the original polynomial as follows:

$$a'_k = (k + 1) a_{k+1} \quad \text{for } k = 0, \dots, n - 1. \quad (2.9)$$

For example, the derivative of 2.6 is $2x - 3$.

The integral of a polynomial is given by:

$$\int_0^x P(t) dt = \sum_{k=0}^n \frac{a_k}{k+1} x^{k+1}. \quad (2.10)$$

⁴This is actually the first program I ever wrote in my first computer programming class. Back in 1969, the language in fashion was ALGOL.

⁵Notice the change in the range of the summation index in equation 2.8.

Thus, the integral of a polynomial with n coefficients is another polynomial, with $n + 1$ coefficients derived from the coefficients of the original polynomial as follows:

$$\bar{a}_k = \frac{a_{k-1}}{k} \quad \text{for } k = 1, \dots, n + 1. \quad (2.11)$$

For the integral, the coefficient \bar{a}_0 is arbitrary and represents the value of the integral at $x = 0$. For example the integral of 2.6 which has the value -2 at $x = 0$ is the polynomial

$$\frac{x^3}{3} - \frac{3^2}{2} + 2x - 2. \quad (2.12)$$

Conventional arithmetic operations are also defined on polynomials and have the same properties⁶ as for signed integers.

Adding or subtracting two polynomials yields a polynomial whose degree is the maximum of the degrees of the two polynomials. The coefficients of the new polynomial are simply the addition or subtraction of the coefficients of same order.

Multiplying two polynomials yields a polynomial whose degree is the product of the degrees of the two polynomials. If $\{a_0, \dots, a_n\}$ and $\{b_0, \dots, b_n\}$ are the coefficients of two polynomials, the coefficients of the product polynomial are given by:

$$c_k = \sum_{i+j=k} a_i b_j \quad \text{for } k = 0, \dots, n + m. \quad (2.13)$$

In equation 2.13 the coefficients a_k are treated as 0 if $k > n$. Similarly the coefficients b_k are treated as 0 if $k > m$.

Dividing a polynomial by another is akin to integer division with remainder. In other word the following equation:

$$P(x) = Q(x) \cdot T(x) + R(x). \quad (2.14)$$

uniquely defines the two polynomials $Q(x)$, the quotient, and $R(x)$, the remainder, for any two given polynomials $P(x)$ and $T(x)$. The algorithm is similar to the algorithm taught in elementary school for dividing integers [Knuth2].

2.2.2 Polynomial — General implementation

As we have seen a polynomial is uniquely defined by its coefficients. Thus, the creation of a new polynomial instance must have the coefficients given. Our implementation assumes that the first element of the array containing the coefficients is the coefficient of the constant term, the second element the coefficient of the linear term (x), and so on.

The method `value` evaluates the polynomial at the supplied argument. This methods implements equation 2.7.

The methods `derivative` and `integral` return each a new instance of a polynomial. The method `integral`: must have an argument specifying the value of the integral of the

⁶The set of polynomials is a vector space in addition to being a ring.

polynomial at 0. A convenience integral method without argument is equivalent to call the method integral with argument 0.

The implementation of polynomial arithmetic is rarely used in numerical computation though. It is, however, a nice example to illustrate a technique called double dispatching. Double dispatching is described in appendix (c.f. section B.4). The need for double dispatching comes for allowing an operation between object of different nature. In the case of polynomials operations can be defined between two polynomials or between a number and a polynomial. In short, double dispatching allows one to identify the correct method based on the type of the two arguments.

2.2.3 Polynomial — Smalltalk implementation

Being a special case of a function a polynomial must of course implement the behavior of functions as discussed in section 2.1.1. Here is a code example on how to use the class DhbPolynomial.

Figure 2.1 with the box Polynomial grayed.

Code example 2.2

```
| polynomial |
polynomial := DhbPolynomial coefficients: #(2 -3 1).
polynomial value: 1.
```

The code above creates an instance of the class DhbPolynomial by giving the coefficient of the polynomial. In this example the polynomial $x^2 - 3x + 2$. The final line of the code computes the value of the polynomial at $x = 1$.

The next example shows how to manipulate polynomials in symbolic form.

Code example 2.3

```
| pol1 pol2 polynomial polD polI |
pol1 := DhbPolynomial coefficients: #(2 -3 1).
pol2 := DhbPolynomial coefficients: #(-3 7 2 1).
polynomial = pol1 * pol2.
polD := polynomial derivative.
polI := polynomial integral.
```

The first line creates the polynomial of example 2.6. The second line creates the polynomial $x^3 + 2x^2 + 7x - 3$. The third line of the code creates a new polynomial, product of the first two. The last two lines create two polynomials, respectively the derivative and the integral of the polynomial created in the third line.

Listing 2.3 shows the Smalltalk implementation of the class DhbPolynomial.

A beginner may have been tempted to make DhbPolynomial a subclass of Array to spare the need for an instance variable. This is of course quite wrong. An array is a subclass of Collection. Most methods implemented or inherited by Array have nothing to do with the behavior of a polynomial as a mathematical entity.

Thus, a good choice is to make the class DhbPolynomial a subclass of Object. It has a single instance variable, an Array containing the coefficients of the polynomial.

It is always a good idea to implement a method printOn: for each class. This method is used by many system utilities to display an object in readable form, in particular the

debugger and the inspectors. The standard method defined for all objects simply displays the name of the class. Thus, it is hard to decide if two different variables are pointing to the same object. Implementing a method `printOn:` allows displaying parameters particular to each instance so that the instances can easily be identified. It may also be used in quick print on the Transcript and may save you the use on an inspector while debugging. Implementing a method `printOn:` for each class that you create is a good general practice, which can make your life as a Smalltalker much easier.

Working with indices in Smalltalk is somewhat awkward for mathematical formulas because the code is quite verbose. In addition a mathematician using Smalltalk for the first time may be disconcerted with all indices starting at 1 instead of 0. Smalltalk, however, has very powerful iteration methods, which largely compensate for the odd index choice, odd for a mathematician that is. In fact, an experienced Smalltalker seldom uses indices explicitly as Smalltalk provides powerful iterator methods.

The method `value:` uses the Smalltalk iteration method `inject:into:` (c.f. section B.3.3). Using this method requires storing the coefficients in reverse order because the first element fed into the method `inject:into:` corresponds to the coefficient of the largest power of x . It would certainly be quite inefficient to reverse the order of the coefficients at each evaluation. Since this requirement also simplifies the computation of the coefficients of the derivative and of the integral, reversing of the coefficients is done in the creation method to make things transparent.

The methods `derivative` and `integral` return a new instance of the class `DhbPolynomial`. They do not modify the object receiving the message. This is also true for all operations between polynomials. The methods `derivative` and `integral` use the method `collect:` returning a collection of the values returned by the supplied block closure at each argument (c.f. section B.3.2).

The method `at:` allows one to retrieve a given coefficient. To ease readability of the multiplication and division methods, the method `at:` has been defined to allow for indices starting at 0. In addition this method returns zero for any index larger than the polynomial's degree. This allows being lax with the index range. In particular, equation 2.13 can be coded exactly as it is.

The arithmetic operations between polynomials are implemented using double dispatching. This is a general technique widely used in Smalltalk (and all other languages with dynamical typing) consisting of selecting the proper method based on the type of the supplied arguments. Double dispatching is explained in section B.4.

Note: Because Smalltalk is a dynamically typed language, our implementation of polynomial is also valid for polynomials with complex coefficients. This is not the case in Java which requires explicit typing.

Listing 2.3 Smalltalk implementation of the polynomial class

Class	DhbPolynomial
Subclass of	Object
Instance variable names:	coefficients

Class methods

coefficients: anArray

^self new initialize: anArray reverse

Instance methods

* aNumberOrPolynomial

^aNumberOrPolynomial timesPolynomial: self

+ aNumberOrPolynomial

^aNumberOrPolynomial addPolynomial: self

- aNumberOrPolynomial

^aNumberOrPolynomial subtractToPolynomial: self

/ aNumberOrPolynomial

^aNumberOrPolynomial dividingPolynomial: self

addNumber: aNumber

| newCoefficients |

newCoefficients := coefficients reverse.

newCoefficients at: 1 put: newCoefficients first + aNumber.

^self class new: newCoefficients

addPolynomial: aPolynomial

^self class new: ((0 to: (self degree max: aPolynomial degree))

collect: [:n | (aPolynomial at: n) + (self at: n)])

at: anInteger

^anInteger < coefficients size

ifTrue: [coefficients at: (coefficients size - anInteger)]

ifFalse:[0]

coefficients

^coefficients deepCopy

degree

```
^coefficients size - 1
```

derivative

```
| n |
n := coefficients size.
^self class new: ( ( coefficients collect: [ :each | n := n - 1.
each * n]) reverse copyFrom: 2 to: coefficients size)
```

dividingPolynomial: aPolynomial

```
^( self dividingPolynomialWithRemainder: aPolynomial) first
```

dividingPolynomialWithRemainder: aPolynomial

```
| remainderCoefficients quotientCoefficients n m norm
quotientDegree |

n := self degree.
m := aPolynomial degree.
quotientDegree := m - n.
quotientDegree < 0
  ifTrue: [ ^Array with: ( self class new: #(0)) with:
aPolynomial].
quotientCoefficients := Array new: quotientDegree + 1.
remainderCoefficients := ( 0 to: m) collect: [ :k | aPolynomial
at: k].

norm := 1 / coefficients first.
quotientDegree to: 0 by: -1
do: [ :k | | x |
x := ( remainderCoefficients at: n + k + 1) * norm.
quotientCoefficients at: (quotientDegree + 1 - k) put:
x.

(n + k - 1) to: k by: -1
do: [ :j |
remainderCoefficients at: j + 1 put:
( ( remainderCoefficients at: j + 1) - (
x * (self at: j - k)))
].
].

^Array with: ( self class new: quotientCoefficients reverse)
with: ( self class new: ( remainderCoefficients copyFrom:
1 to: n))
```

initialize: anArray

```
coefficients := anArray.
^self
```

integral

```
^self integral: 0
```

integral: aValue

```
| n |
n := coefficients size + 1.
^self class new: ( ( coefficients collect: [ :each | n := n - 1.
each / n]) copyWith: aValue) reverse
```

printOn: aStream

```
| n firstNonZeroCoefficientPrinted |
n := 0.
firstNonZeroCoefficientPrinted := false.
coefficients reverse do:
[ :each |
each = 0
ifFalse:[ firstNonZeroCoefficientPrinted
ifTrue: [ aStream space.
each < 0
ifFalse:[ aStream
nextPut: $+].
aStream space.
]
ifFalse:[ firstNonZeroCoefficientPrinted
:= true].
( each = 1 and: [ n > 0])
ifFalse:[ each printOn: aStream].
n > 0
ifTrue: [ aStream nextPutAll: ' X'.
n > 1
ifTrue: [ aStream
nextPut: $^
n printOn:
aStream.
].
].
].
```

```

    n := n + 1.
  ].

subtractToPolynomial: aPolynomial

  ^self class new: ( ( 0 to: (self degree max: aPolynomial degree))
    collect: [ :n | ( aPolynomial at: n) - ( self at: n)])

timesNumber: aNumber

  ^self class new: ( coefficients reverse collect: [ :each | each *
    aNumber])

timesPolynomial: aPolynomial

  | productCoefficients degree |
  degree := aPolynomial degree + self degree.
  productCoefficients := ( degree to: 0 by: -1)
    collect: [ :n | | sum |
      sum := 0.
      0 to: (degree - n)
        do: [ :k | sum := (self at: k) * (aPolynomial
          at: ( degree - n - k)) + sum].
      sum
    ].
  ^self class new: productCoefficients

value: aNumber

  ^coefficients inject: 0 into: [ :sum :each | sum * aNumber +
    each]

```

Listing 2.4 shows the listing of the methods used by the class Number as part of the double dispatching of the arithmetic operations on polynomials.

Listing 2.4 Method of class Number related to polynomials

Class **Number**
 Subclass of Magnitude

Instance methods
 addPolynomial: aPolynomial

```

^aPolynomial addNumber: self

dividingPolynomial: aPolynomial
    ^aPolynomial timesNumber: (1 / self)

subtractToPolynomial: aPolynomial
    ^aPolynomial addNumber: self negated

timesPolynomial: aPolynomial
    ^aPolynomial timesNumber: self

```

2.2.4 Polynomial — Java implementation

Listing 2.5 shows the Java implementation of the class `PolynomialFunction`. Here is a code example on how to use the class `PolynomialFunction`.

Figure 2.2 with the box Polynomial grayed.

Code example 2.4

```

double[] coefficients = {2, -3, 1};
PolynomialFunction poly = new PolynomialFunction( coefficients);
double p1 = polynomial.value( 1).

```

First the coefficients of the polynomial $x^2 - 3x + 2$ are defined in a variable. The next line creates a new instance of the class `PolynomialFunction` with these coefficients. The final line of the code calculates the value of the polynomial at $x = 1$.

For the same reason explained in the Smalltalk section, the class `PolynomialFunction` is a subclass of `Object`. It implements the interface `OneVariableFunction` (c.f. section 2.1.2) indicating that this class must implement the method `value`. Thus, a polynomial can be used by several other classes defined hereafter in this book.

The Java implementation uses indices since array elements are referred to by indices starting at 0. This makes a direct comparison with the mathematical expression easier than in Smalltalk. Nevertheless, the method `coefficient` was created to access coefficients for any non-negative index. If the index is larger than the polynomial's degree this routine returns 0. Using this method makes the code for adding and multiplying polynomials much simpler. The method `coefficient` corresponds to the method `at:` in the Smalltalk implementation.

Polynomial arithmetic is implemented with special methods since Java does not allow overloading the arithmetic operators. As Java methods and their arguments are typed, a unique method must be written for each argument type. Therefore there is no need for double dispatching⁷.

⁷On the other hand the readability of code using polynomial arithmetic is much better in Smalltalk than in Java.

Finally the method `toString()` is implemented for the same reasons than a method `printOn:` was implemented in Smalltalk.

Listing 2.5 Java implementation of the polynomial class

```
package DhbFunctionEvaluation;

import java.util.Vector;
import java.util.Enumeration;
import DhbInterfaces.OneVariableFunction;
import DhbIterations.NewtonZeroFinder;
import DhbFunctionEvaluation.DhbMath;
/**
 * Mathematical polynomial:
 *  $c[0] + c[1] * x + c[2] * x^2 + \dots$ 
 *
 * @author Didier H. Besset
 */
public class PolynomialFunction implements OneVariableFunction
{
    /**
     * Polynomial coefficients.
     */
    private double[] coefficients;

    /**
     * Constructor method.
     * @param coeffs polynomial coefficients.
     */
    public PolynomialFunction( double[] coeffs)
    {
        coefficients = coeffs;
    }
    /**
     *
     * @param r double    number added to the polynomial.
     * @return DhbFunctionEvaluation.PolynomialFunction
     */
    public PolynomialFunction add( double r)
    {
        int n = coefficients.length;
        double coef[] = new double[n];
        coef[0] = coefficients[0] + r;
        for ( int i = 1; i < n; i++)
            coef[i] = coefficients[i];
    }
}
```

```

    return new PolynomialFunction( coef);
}
/**
 *
 * @param p DhbFunctionEvaluation.PolynomialFunction
 * @return DhbFunctionEvaluation.PolynomialFunction
 */
public PolynomialFunction add( PolynomialFunction p)
{
    int n = Math.max( p.degree(), degree()) + 1;
    double[] coef = new double[n];
    for ( int i = 0; i < n; i++ )
        coef[i] = coefficient(i) + p.coefficient(i);
    return new PolynomialFunction( coef);
}
/**
 * Returns the coefficient value at the desired position
 * @param n int    the position of the coefficient to be returned
 * @return double the coefficient value
 * @version 1.2
 */
public double coefficient( int n)
{
    return n < coefficients.length ? coefficients[n] : 0;
}
/**
 *
 * @param r double    a root of the polynomial (no check made).
 * @return PolynomialFunction the receiver divided by polynomial (x - r).
 */
public PolynomialFunction deflate( double r)
{
    int n = degree();
    double remainder = coefficients[n];
    double[] coef = new double[n];
    for ( int k = n - 1; k >= 0; k--)
    {
        coef[k] = remainder;
        remainder = remainder * r + coefficients[k];
    }
    return new PolynomialFunction( coef);
}
/**
 * Returns degree of this polynomial function
 * @return int degree of this polynomial function
 */

```

```

public int degree()
{
    return coefficients.length - 1;
}
/**
 * Returns the derivative of the receiver.
 * @return PolynomialFunction derivative of the receiver.
 */
public PolynomialFunction derivative()
{
    int n = degree();
    if ( n == 0 )
    {
        double coef[] = {0};
        return new PolynomialFunction( coef);
    }
    double coef[] = new double[n];
    for ( int i = 1; i <= n; i++)
        coef[i-1] = coefficients[i]*i;
    return new PolynomialFunction( coef);
}
/**
 *
 * @param r double
 * @return DhbFunctionEvaluation.PolynomialFunction
 */
public PolynomialFunction divide( double r)
{
    return multiply( 1 / r);
}
/**
 *
 * @param r double
 * @return DhbFunctionEvaluation.PolynomialFunction
 */
public PolynomialFunction divide( PolynomialFunction p)
{
    return divideWithRemainder(p)[0];
}
/**
 *
 * @param r double
 * @return DhbFunctionEvaluation.PolynomialFunction
 */
public PolynomialFunction[] divideWithRemainder( PolynomialFunction p)
{

```

```

PolynomialFunction[] answer = new PolynomialFunction[2];
int m = degree();
int n = p.degree();
if ( m < n )
{
    double[] q = {0};
    answer[0] = new PolynomialFunction( q);
    answer[1] = p;
    return answer;
}
double[] quotient = new double[ m - n + 1];
double[] coef = new double[ m + 1];
for ( int k = 0; k <= m; k++ )
    coef[k] = coefficients[k];
double norm = 1 / p.coefficient( n);
for ( int k = m - n; k >= 0; k-- )
{
    quotient[k] = coef[ n + k] * norm;
    for ( int j = n + k - 1; j >= 0; j-- )
        coef[j] -= quotient[k] * p.coefficient(j-k);
}
double[] remainder = new double[n];
for ( int k = 0; k < n; k++ )
    remainder[k] = coef[k];
answer[0] = new PolynomialFunction( quotient);
answer[1] = new PolynomialFunction( remainder);
return answer;
}
/**
 * Returns the integral of the receiver having the value 0 for X = 0.
 * @return PolynomialFunction integral of the receiver.
 */
public PolynomialFunction integral( )
{
    return integral( 0);
}
/**
 * Returns the integral of the receiver having the specified value for X = 0.
 * @param value double    value of the integral at x=0
 * @return PolynomialFunction integral of the receiver.
 */
public PolynomialFunction integral( double value)
{
    int n = coefficients.length + 1;
    double coef[] = new double[n];
    coef[0] = value;

```

```

        for ( int i = 1; i < n; i++)
            coef[i] = coefficients[i-1]/i;
        return new PolynomialFunction( coef);
    }
    /**
     *
     * @param r double
     * @return DhbFunctionEvaluation.PolynomialFunction
     */
    public PolynomialFunction multiply( double r)
    {
        int n = coefficients.length;
        double coef[] = new double[n];
        for ( int i = 0; i < n; i++)
            coef[i] = coefficients[i] * r;
        return new PolynomialFunction( coef);
    }
    /**
     *
     * @param p DhbFunctionEvaluation.PolynomialFunction
     * @return DhbFunctionEvaluation.PolynomialFunction
     */
    public PolynomialFunction multiply( PolynomialFunction p)
    {
        int n = p.degree() + degree();
        double[] coef = new double[n + 1];
        for ( int i = 0; i <= n; i++)
        {
            coef[i] = 0;
            for ( int k = 0; k <= i; k++)
                coef[i] += p.coefficient(k) * coefficient(i-k);
        }
        return new PolynomialFunction( coef);
    }
    /**
     *
     * @return double[]
     */
    public double[] roots()
    {
        return roots( DhbMath.defaultNumericalPrecision());
    }
    /**
     *
     * @param desiredPrecision double
     * @return double[]

```

```

*/
public double[] roots( double desiredPrecision)
{
    PolynomialFunction dp = derivative();
    double start = 0;
    while ( Math.abs( dp.value( start)) < desiredPrecision )
        start = Math.random();
    PolynomialFunction p = this;
    NewtonZeroFinder rootFinder = new NewtonZeroFinder( this, dp, start);
    rootFinder.setDesiredPrecision( desiredPrecision);
    Vector rootCollection = new Vector( degree());
    while ( true)
    {
        rootFinder.evaluate();
        if ( !rootFinder.hasConverged() )
            break;
        double r = rootFinder.getResult();
        rootCollection.addElement( new Double( r));
        p = p.deflate( r);
        if ( p.degree() == 0 )
            break;
        rootFinder.setFunction( p);
        try { rootFinder.setDerivative( p.derivative());}
        catch ( IllegalArgumentException e) { };
    }
    double[] roots = new double[ rootCollection.size()];
    Enumeration e = rootCollection.elements();
    int n = 0;
    while ( e.hasMoreElements() )
    {
        roots[n++] = ( (Double) e.nextElement()).doubleValue();
    }
    return roots;
}
/**
 *
 * @param p DhbFunctionEvaluation.PolynomialFunction
 * @return DhbFunctionEvaluation.PolynomialFunction
 */
public PolynomialFunction subtract( double r)
{
    return add( -r);
}
/**
 *
 * @return DhbFunctionEvaluation.PolynomialFunction

```

```

    * @param p DhbFunctionEvaluation.PolynomialFunction
    */
    public PolynomialFunction subtract( PolynomialFunction p)
    {
        int n = Math.max( p.degree(), degree()) + 1;
        double[] coef = new double[n];
        for ( int i = 0; i < n; i++ )
            coef[i] = coefficient(i) - p.coefficient(i);
        return new PolynomialFunction( coef);
    }
    /**
    * Returns a string representing the receiver
    */
    public String toString()
    {
        StringBuffer sb = new StringBuffer();
        boolean firstNonZeroCoefficientPrinted = false;
        for ( int n = 0; n < coefficients.length; n++)
        {
            if ( coefficients[n] != 0 )
            {
                if ( firstNonZeroCoefficientPrinted)
                    sb.append( coefficients[n] > 0 ? " + " : " - ");
                else
                    firstNonZeroCoefficientPrinted = true;
                if ( n == 0 || coefficients[n] != 1)
                    sb.append( Double.toString( coefficients[n]) );
                if ( n > 0 )
                    sb.append( " X^"+n);
            }
        }
        return sb.toString();
    }
    /**
    * Returns the value of the polynomial for the specified variable value.
    * @param x double    value at which the polynomial is evaluated
    * @return double polynomial value.
    */
    public double value( double x)
    {
        int n = coefficients.length;
        double answer = coefficients[--n];
        while ( n > 0 )
            answer = answer * x + coefficients[--n];
        return answer;
    }

```




Figure 2.3: The error function and the normal distribution

```
/**
 * Returns the value and the derivative of the polynomial
 * for the specified variable value in an array of two elements
 * @version 1.2
 * @param x double    value at which the polynomial is evaluated
 * @return double[0]  the value of the polynomial
 * @return double[1]  the derivative of the polynomial
 */
public double[] valueAndDerivative( double x)
{
    int n = coefficients.length;
    double[] answer = new double[2];
    answer[0] = coefficients[--n];
    answer[1] = 0;
    while ( n > 0 )
    {
        answer[1] = answer[1] * x + answer[0];
        answer[0] = answer[0] * x + coefficients[--n];
    }
    return answer;
}
}
```

Note: The three methods deflate and roots (roots having two variants) are discussed in section 5.4.2.

2.3 Error function

The error function is the integral of the normal distribution. The error function is used in statistics to evaluate the probability of finding a measurement larger than a given value

when the measurements are distributed according to a normal distribution. Figure 2.3 shows the familiar bell-shaped curve of the probability density function of the normal distribution (dotted line) together with the error function (solid line).

In medical sciences one calls centile the value of the error function expressed in percent. For example, obstetricians look whether the weight at birth of the first born child is located below the 10th centile or above the 90th centile to assess a risk factor for a second pregnancy⁸.

2.3.1 Mathematical definitions

Because it is the integral of the normal distribution, the error function, $\text{erf}(x)$, gives the probability of finding a value lower than x when the values are distributed according to a normal distribution with mean 0 and standard deviation 1. The mean and the standard deviation are explained in section 9.1. This probability is expressed by the following integral⁹:

$$\text{erf}(x) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^x e^{-\frac{t^2}{2}} dt \quad (2.15)$$

The result of the error function lies between 0 and 1.

One could carry out the integral numerically, but there exists several good approximations. The following formula is taken from [Abramovitz & Stegun].

Main equation⇒

$$\text{erf}(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{x^2}{2}} \sum_{i=1}^5 a_i r(x)^i \quad \text{for } x \geq 0. \quad (2.16)$$

where

$$r(x) = \frac{1}{1 - 0.2316419x}. \quad (2.17)$$

and

$$\begin{cases} a_1 = 0.31938153 \\ a_2 = -0.356563782 \\ a_3 = 1.7814779372 \\ a_4 = -1.821255978 \\ a_5 = 1.330274429 \end{cases} \quad (2.18)$$

The error on this formula is better than 7.5×10^{-8} for negative x . To compute the value for positive values, one uses the fact that:

Main equation⇒

$$\text{erf}(x) = 1 - \text{erf}(-x). \quad (2.19)$$

⁸c.f. footnote 10 on page 55

⁹In [Abramovitz & Stegun] and [Press et al.], the error function is defined as:

$$\text{erf}(x) = \frac{2}{\sqrt{\pi}} \int_0^x e^{-\frac{t^2}{2}} dt$$

When dealing with a general Gaussian distribution with average μ and standard deviation σ it is convenient to define a generalized error function as:

$$\text{erf}(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \int_{-\infty}^x e^{-\frac{(x-\mu)^2}{2\sigma^2}} dt. \quad (2.20)$$

A simple change of variable in the integral shows that the generalized error function can be obtained from the error function as:

$$\text{erf}(x; \mu, \sigma) = \text{erf}\left(\frac{x - \mu}{\sigma}\right). \quad (2.21) \quad \Leftarrow \text{Main equation}$$

Thus, one can compute the probability of finding a measurement x within the interval $[\mu - t \cdot \sigma, \mu + t \cdot \sigma]$ when the measurements are distributed according to a Gaussian distribution with average μ and standard deviation σ :

$$\text{Prob}\left(\frac{|x - \mu|}{\sigma} \leq t\right) = 2 \cdot \text{erf}(t) - 1. \quad \text{for } t \geq 0. \quad (2.22)$$

Example

Now we can give the answer to the problem of deciding whether a pregnant woman needs special attention during her second pregnancy. Let the weight at birth of her first child be 2.85 Kg. and let the duration of her first pregnancy be 39 weeks. In this case measurements over a representative sample of all births yielding healthy babies have an average of 3.39 Kg and a standard deviation of 0.44 Kg¹⁰. The probability of having a weight of birth smaller than that of the woman's first child is:

$$\begin{aligned} \text{Prob}(\text{Weight} \leq 2.85 \text{ Kg}) &= \text{erf}\left(\frac{2.85 - 3.39}{0.44}\right), \\ &= 11.2\%. \end{aligned}$$

According to current practice, this second pregnancy does not require special attention.

2.3.2 Error function — Smalltalk implementation

The error function is implemented as a single method for the class Number. Thus, computing the centile of our preceding example is simply coded as:

Code example 2.5

```
| weight average stDev centile |
weight := 2.85.
average := 3.39.
stDev := 0.44.
centile := ( ( weight - average ) / stDev ) erf * 100.
```

Figure 2.1 with the box
ErfApproximation grayed.

¹⁰This is the practice at the department of obstetrics and gynecology of the Chelsea & Westminster Hospital of London. The numbers are reproduced with permission of Prof. P.J. Steer.

If you want to compute the probability for a measurement to lay within 3 standard deviations from its mean, you need to evaluate the following expression using equation 2.22:

Code example 2.6

```
3 errorFunction * 2 - 1
```

If one needs to use the error function as a function, one must use it inside a block closure. In this case one defines a function object as follows:

Code example 2.7

```
| errorFunction |
errorFunction := [ :x | x errorFunction].
```

Listing 2.6 shows the Smalltalk implementation of the error function.

In Smalltalk we are allowed to extend existing classes. Thus, the public method to evaluate the error function is implemented as a method of the base class Number. This method uses the class, DhbErfApproximation, used to store the constants of equation 2.18 and evaluate the formula of equations 2.16 and 2.17. In our case, there is no need to create a separate instance of the class DhbErfApproximation at each time since all instances would actually be exactly identical. Thus, the class DhbErfApproximation is a singleton class. A singleton class is a class, which can only create a single instance [Gamma et al.]. Once the first instance is created, it is kept in a class instance variable. Any subsequent attempt to create an additional instance will return a pointer to the class instance variable holding the first created instance.

One could have implemented all of these methods as class methods to avoid the singleton class. In Smalltalk, however, one tends to reserve class method for behavior needed by the structural definition of the class. So, the use of a singleton class is preferable. A more detailed discussion of this topic can be found in [Alpert et al.].

Listing 2.6 Smalltalk implementation of the Error function

```
Class      Number
Subclass of  Magnitude
```

```
Instance methods
```

```
errorFunction
```

```
^DhbErfApproximation new value: self
```

```
Class      DhbErfApproximation
Subclass of  Object
Instance variable names:  constant series norm
Class variable names:    UniqueInstance
```

```

Class methods

new
  UniqueInstance isNil
    ifTrue: [ UniqueInstance := super new.
              UniqueInstance initialize.
            ].
  ^UniqueInstance

Instance methods

initialize

  constant := 0.2316419.
  norm := 1 / ( Float pi * 2) sqrt.
  series := DhhPolynomial coefficients: #( 0.31938153 -0.356563782
                                           1.781477937 -1.821255978 1.330274429).

normal: aNumber
  ^[ ( aNumber squared * -0.5) exp * norm]
  when: ExAll do: [ :signal | signal exitWith: 0]

value: aNumber
  | t |
  aNumber = 0
    ifTrue: [ ^0.5].
  aNumber > 0
    ifTrue: [ ^1- ( self value: aNumber negated)].
  aNumber < -20
    ifTrue: [ ^0].
  t := 1 / (1 - (constant * aNumber)).
  ^( series value: t) * t * (self normal: aNumber)

```

2.3.3 Error function — Java implementation

Unfortunately Java does not allow for extension of existing classes. Thus, a new class must be created. All methods are implemented as static methods since there is no instance dependency. Static methods in Java have little in common with Smalltalk class methods. In particular there is no inheritance of static methods. Static methods indicate that the code is not instance dependent and are used to generate compiler and run time optimization. Thus, a singleton class is not needed in this case.

Since the error function is almost always used in conjunction with statistical analysis, we have implemented the method as a static method of the class `NormalDistribution` discussed in section 9.6.2. Thus, computing the error function of 2 must be coded as:

Figure 2.2 with the box `NormalDistribution` grayed.

Code example 2.8

```
double x = NormalDistribution.errorFunction(2)
```

If you want to compute the probability for a measurement to lay within 3 standard deviations from its mean, you need to evaluate the following expression:

Code example 2.9

```
double y = 2 * NormalDistribution.errorFunction(3)- 1
```

Listing 2.7 shows the Java implementation of the error function.

The class `NormalDistribution` is tagged as `final` because it does not make sense that this class has any subclass.

Listing 2.7 Java implementation of the Error function (partial listing)

```
package DhbStatistics;

import DhbFunctionEvaluation.PolynomialFunction;
/**
 * Normal distribution, a.k.a. Gaussian distribution.
 *
 */
public final class NormalDistribution extends ProbabilityDensityFunction
{
    private static double baseNorm = Math.sqrt( 2 * Math.PI);
    /**
     * Series to compute the error function.
     */
    private static PolynomialFunction errorFunctionSeries;
    /**
     * Constant needed to compute the argument to the error function series.
     */
    private static double errorFunctionConstant = 0.2316419;

    /**
     * @return error function for the argument.
     * @param x double
     */
    public static double errorFunction ( double x)
    {
        if ( errorFunctionSeries == null )
        {
            double[] coeffs = { 0.31938153,-0.356563782, 1.781477937,-1.821255978, 1.330274429};
            errorFunctionSeries = new PolynomialFunction( coeffs);
        }
        if ( x > 0 )
```

```

return 1 - errorFunction( -x);
double t = 1 / (1 - errorFunctionConstant * x);
return t * errorFunctionSeries.value( t ) * normal( x);
}
/**
 * @return the density probability function for a (0,1) normal distribution.
 * @param x double value for which the probability is evaluated.
 */
static public double normal( double x)
{
return Math.exp( -0.5 * x * x ) / baseNorm;
}
}

```

2.4 Gamma function

The gamma function is used in many mathematical functions. In this book, the gamma function is needed to compute the normalization factor of several probability density functions (c.f. sections 9.7 and 10.3). It is also needed to compute the beta function (c.f. section 2.5).

2.4.1 Mathematical definitions

The Gamma function is defined by the following integral, called Euler's integral¹¹:

$$\Gamma(x) = \int_0^{\infty} t^x e^{-t} dt \quad (2.23)$$

From equation 2.23 a recurrence formula can be derived:

$$\Gamma(x+1) = x \cdot \Gamma(x) \quad (2.24)$$

The value of the Gamma function can be computed for special values of x :

$$\begin{cases} \Gamma(1) = 1 \\ \Gamma(2) = 1 \end{cases} \quad (2.25)$$

From 2.24 and 2.25, the well-known relation between the value of the Gamma function for positive integers and the factorial can be derived:

$$\Gamma(n) = (n-1)! \quad \text{for } n > 0. \quad (2.26)$$

The most precise approximation for the Gamma function is given by a formula discovered by Lanczos [Press et al.]:

¹¹Leonard Euler to be precise as the Euler family produced many mathematicians.

$$\Gamma(x) \approx e^{(x+\frac{5}{2})} \left(x + \frac{5}{2}\right) \frac{\sqrt{2\pi}}{x} \left(c_0 + \sum_{n=1}^6 \frac{c_n}{x+n} + \epsilon\right) \quad (2.27)$$

where

$$\begin{cases} c_0 = & 1.000000000190015 \\ c_1 = & 76.18009172947146 \\ c_2 = & -86.50532032941677 \\ c_3 = & 24.01409824083091 \\ c_4 = & -1.231739572450155 \\ c_5 = & 1.208650973866179 \cdot 10^{-3} \\ c_6 = & -5.395239384953 \cdot 10^{-6} \end{cases} \quad (2.28)$$

This formula approximates $\Gamma(x)$ for $x > 1$ with $\epsilon < 2 \cdot 10^{-10}$. Actually, this remarkable formula can be used to compute the gamma function of any complex number z with $\Re(z) > 1$ to the quoted precision. Combining Lanczos' formula with the recurrence formula 2.24 is sufficient to compute values of the Gamma function for all positive numbers.

For example, $\Gamma(\frac{3}{2}) = \frac{\sqrt{\pi}}{2} = 0.886226925452758$ whereas Lanczos formula yields the value 0.886226925452754, that is, an absolute error of $4 \cdot 10^{-15}$. The corresponding relative precision is almost equal to the floating-point precision of the machine on which this computation was made.

Although this is seldom used, the value of the Gamma function for negative non-integer numbers can be computed using the reflection formula hereafter:

$$\Gamma(x) = \frac{\pi}{\Gamma(1-x) \sin \pi x} \quad (2.29)$$

In summary, the algorithm to compute the Gamma function for any argument goes as follows:

1. If x is a non-positive integer ($x \leq 0$), raise an exception.
2. If x is smaller than or equal to 1 ($x < 1$), use the recurrence formula 2.24.
3. If x is negative ($x < 0$, but non integer), use the reflection formula 2.29.
4. Otherwise use Lanczos' formula 2.27.

One can see from the leading term of Lanczos' formula that the gamma function raises faster than an exponential. Thus, evaluating the gamma function for numbers larger than a few hundreds will exceed the capacity of the floating number representation on most machines. For example, the maximum exponent of a double precision IEEE floating-point number is 1024. Evaluating directly the following expression:

$$\frac{\Gamma(460.5)}{\Gamma(456.3)} \quad (2.30)$$

will fail since $\Gamma(460.5)$ is larger than 10^{1024} . Thus, its evaluation yields a floating-point overflow exception. It is therefore recommended to use the logarithm of the gamma

function whenever it is used in quotients involving large numbers. The expression of equation 2.30 is then evaluated as:

$$\exp [\ln \Gamma (460.5) - \ln \Gamma (456.3)] \quad (2.31)$$

which yield the result $1.497 \cdot 10^{11}$. That result fits comfortably within the floating-point representation.

For similar reasons the leading factors of Lanczos formula are evaluated using logarithms in both implementations.

2.4.2 Gamma function — Smalltalk implementation

Like the error function, the gamma function is implemented as a single method of the class `Number`. Thus, computing the gamma function of 2.5 is simply coded as:

Code example 2.10

```
2.5 gamma
```

To obtain the logarithm of the gamma function, you need to evaluate the following expression:

Code example 2.11

```
2.5 logGamma
```

Listing 11 shows the Smalltalk implementation of the gamma function.

Here, the gamma function is implemented with two methods: one for the class `Integer` and one for the class `Float`. Otherwise, the scheme to define the gamma function is similar to that of the error function. Please refer to section 2.3.2 for detailed explanations.

Since the method `factorial` is already defined for integers in the base classes, the gamma function has been defined using equation 2.26 for integers. An error is generated if one attempts to compute the gamma function for non-positive integers. The class `Number` delegates the computation of Lanczos' formula to a singleton class. This is used by the non-integer subclasses of `Number`: `Float` and `Fraction`.

The execution time to compute the gamma function for floating argument given in Table 1.1 in section 1.6.

Listing 2.8 Smalltalk implementation of the gamma function

Class **Integer**

Subclass of `Number`

Instance methods

`gamma`

`self > 0`

`ifFalse:[^self error: 'Attempt to compute the Gamma function
of a non-positive integer'].`

`^(self - 1) factorial`

Figure 2.1 with the box
`LanczosFormula` grayed.

Class **Number**
 Subclass of Magnitude

Instance methods

gamma

```

^self > 1
  ifTrue: [ ^DhbLanczosFormula new gamma: self]
  ifFalse:[ self < 0
    ifTrue: [ Float pi / ( ( Float pi * self) sin
      * ( 1 - self) gamma)]
    ifFalse:[ ( DhbLanczosFormula new gamma: (
      self + 1)) / self]
  ]

```

logGamma

```

^self > 1
  ifTrue: [ DhbLanczosFormula new logGamma: self]
  ifFalse:[ self > 0
    ifTrue: [ ( DhbLanczosFormula new logGamma: (
      self + 1) ) - self ln ]
    ifFalse: [ ^self error: 'Argument for the log
      gamma function must be positive' ]
  ]

```

Class **DhbLanczosFormula**
 Subclass of Object

Instance variable names: coefficients sqrt2Pi

Class variable names: UniqueInstance

Class methods

new

```

UniqueInstance isNil
  ifTrue: [ UniqueInstance := super new.
    UniqueInstance initialize.
  ].
^UniqueInstance

```

Instance methods

gamma: aNumber

```
^( self leadingFactor: aNumber) exp * ( self series: aNumber) *
  sqrt2Pi / aNumber
```

initialize

```
sqrt2Pi := ( Float pi * 2) sqrt.
coefficients := #( 76.18009172947146 -86.50532032941677
  24.01409824083091 -1.231739572450155 0.1208650973866179e-2
  -0.5395239384953e-5).
^self
```

leadingFactor: aNumber

```
| temp |
temp := aNumber + 5.5.
^( temp ln * ( aNumber + 0.5) - temp)
```

logGamma: aNumber

```
^( self leadingFactor: aNumber) + ( ( self series: aNumber) *
  sqrt2Pi / aNumber) ln
```

series: aNumber

```
| term |
term := aNumber.
^coefficients inject: 1.000000000190015
  into: [ :sum :each | term := term + 1. each /
  term + sum]
```

2.4.3 Gamma function — Java implementation

Like the error function, the gamma function is implemented as a static method in a special class. Thus, computing the gamma function of 2.5 is simply coded as:

Code example 2.12

```
double g = GammaFunction.gamma( 2.5)
```

To obtain the logarithm of the gamma function, you need to evaluate the following expression:

Figure 2.2 with the box GammaFunction grayed.

Code example 2.13

```
double lg = GammaFunction.logGamma( 2.5)
```

Listing 2.9 shows the Java implementation of the Gamma function.

For the Java implementation a specific class containing only static methods was created. This class also implements the beta function discussed in section 2.5.

To prevent the unnecessary computation of a logarithm, the code for the gamma function and the logarithm of the gamma function have duplicated code.

Instead of raising an exception when the argument is a non-positive integer, the function simply returns the system defined value `Double.NaN`, a unique feature of Java to represent a number impossible to compute. This feature is quite handy because `Double.NaN` propagates itself in any subsequent expression without raising an exception. The object recuperating the end result can check for such value without the need to catch an exception.

Strangely enough, Java does not implement an `Integer.NaN`. Thus, the method `factorial` has no other choice than raising an exception when it is called with a negative integer argument.

Listing 2.9 Java implementation of the gamma function

```
package DhbFunctionEvaluation;

/**
 * Gamma function (Euler's integral).
 *
 * @author Didier H. Besset
 */
public final class GammaFunction
{
    static double sqrt2Pi = Math.sqrt( 2 * Math.PI);
    static double[] coefficients = { 76.18009172947146,
                                    -86.50532032941677,
                                    24.01409824083091,
                                    -1.231739572450155,
                                    0.1208650973866179e-2,
                                    -0.5395239384953e-5};

    /**
     * @return double    beta function of the arguments
     * @param x double
     * @param y double
     */
    public static double beta ( double x, double y)
    {
        return Math.exp( logGamma( x) + logGamma( y) - logGamma( x + y));
    }
}
```

```

/**
 * @return long   factorial of n
 * @param n long
 */
public static long factorial ( long n)
{
    return n < 2 ? 1 : n * factorial( n - 1);
}
/**
 * @return double   gamma function
 * @param x double
 */
public static double gamma ( double x)
{
    return x > 1
        ? Math.exp( leadingFactor(x)) * series(x) * sqrt2Pi / x
        : ( x > 0 ? gamma(x + 1) / x
            : Double.NaN);
}
/**
 * @return double
 * @param x double
 */
private static double leadingFactor ( double x)
{
    double temp = x + 5.5;
    return Math.log( temp) * ( x + 0.5) - temp;
}
/**
 * @return double   logarithm of the beta function of the arguments
 * @param x double
 * @param y double
 */
public static double logBeta ( double x, double y)
{
    return logGamma( x) + logGamma( y) - logGamma( x + y);
}
/**
 * @return double   log of the gamma function
 * @param x double
 */
public static double logGamma ( double x)
{
    return x > 1
        ? leadingFactor(x) + Math.log( series(x) * sqrt2Pi / x)
        : ( x > 0 ? logGamma(x + 1) - Math.log( x)

```

```

        : Double.NaN);
    }
    /**
     * @return double    value of the series in Lanczos formula.
     * @param x double
     */
    private static double series( double x)
    {
        double answer = 1.000000000190015;
        double term = x;
        for ( int i = 0; i < 6; i++)
        {
            term += 1;
            answer += coefficients[i] / term;
        }
        return answer;
    }
}

```

Note: Listing 2.9 also contains the method beta needed to compute the beta function discussed in section 2.5.

2.5 Beta function

The beta function is directly related to the gamma function. In this book, the beta function is needed to compute the normalization factor of several probability density functions (c.f. sections 10.1, 10.2 and D.1).

2.5.1 Mathematical definitions

The beta function is defined by the following integral:

$$B(x, y) = \int_0^1 t^{x-1} (1-t)^{y-1} dt \quad (2.32)$$

The beta function is related to the gamma function with the following relation:

$$B(x, y) = \frac{\Gamma(x) \Gamma(y)}{\Gamma(x+y)} \quad (2.33)$$

Thus, computation of the beta function is directly obtained from the gamma function. As evaluating the gamma function might overflow the floating-point exponent (c.f. discussion at the end of section 2.4.1), it is best to evaluate the above formula using the logarithm of the gamma function.

2.5.2 Beta function — Smalltalk implementation

Like the error and gamma functions, the gamma function is implemented as a single method of the class `Number`. Thus, computing the beta function of 2.5 and 5.5 is simply coded as:

Code example 2.14
`2.5 beta: 5.5`

Computing the logarithm of the beta function of 2.5 and 5.5 is simply coded as:

Code example 2.15
`2.5 logBeta: 5.5`

Listing 2.10 shows the implementation of the beta function in Smalltalk.

Listing 2.10 Smalltalk implementation of the beta function

```
Class      Number
Subclass of  Magnitude
```

```
    Instance methods
```

```
beta: aNumber
```

```
    ^( self logBeta: aNumber) exp
```

```
logBeta: aNumber
```

```
    ^self logGamma + aNumber logGamma - ( self + aNumber) logGamma
```

2.5.3 Beta function — Java implementation

The Java implementation of the beta function consists of the two methods `beta` and `logBeta` in the class `GammaFunction` shown in Listing 2.9.

Figure 2.2 with the box `GammaFunction` grayed.

Chapter 3

Interpolation

On ne peut prévoir les choses qu'après qu'elles sont arrivées.¹
Eugène Ionesco

Interpolation is a technique allowing the estimation of a function over the range covered by a set of points at which the function's values are known. These points are called the sample points. Interpolation is useful to compute a function whose evaluation is highly time consuming: with interpolation it suffices to compute the function's values for a small number of well-chosen sample points. Then, evaluation of the function between the sample points can be made with interpolation.

Interpolation can also be used to compute the value of the inverse function, that is finding a value x such that $f(x) = c$ where c is a given number, when the function is known for a few sample points bracketing the sought value. People often overlook this easy and direct computation of the inverse function.

Interpolation is often used interchangeably with extrapolation. This is not correct, however. Extrapolation is the task of estimating a function outside of the range covered by the sample points. If no model exists for the data extrapolation is just gambling. Methods exposed in this chapter should not be used for extrapolation.

Interpolation should not be mistaken with function (or curve) fitting. In the case of interpolation the sample points purely determine the interpolated function. Function fitting allows constraining the fitted function independently from the sample points. As a result fitted functions are more stable than interpolated functions especially when the supplied values are subject to fluctuations coming from rounding or measurement errors. Fitting is discussed in chapter 10.

3.1 General remarks

There are several methods of interpolation. One difference is the type of function used. The other is the particular algorithm used to determine the function. For example, if the

¹One can predict things only after they have occurred.

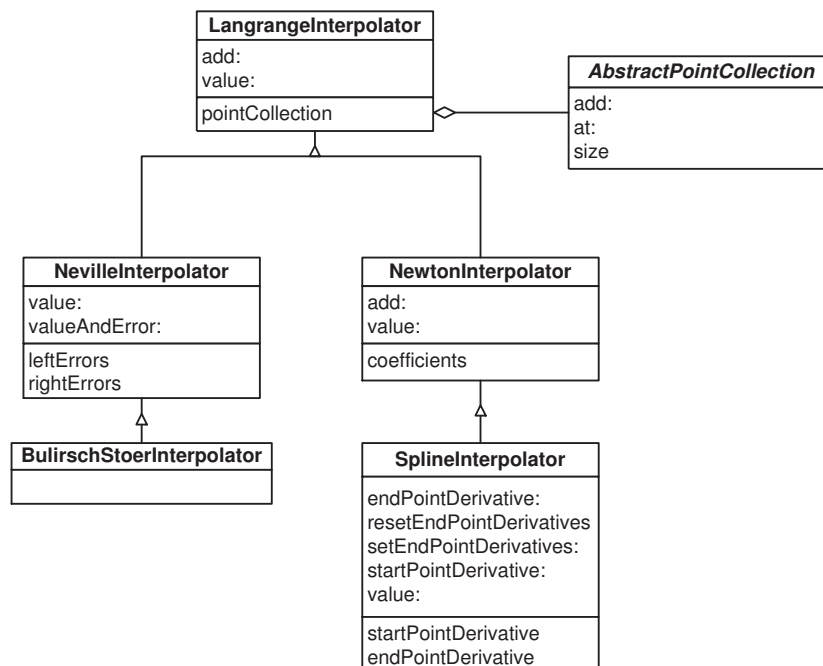


Figure 3.1: Class diagram for the interpolation classes

function is periodic, interpolation can be obtained by computing a sufficient number of coefficients of the Fourier series for that function.

In the absence of any information about the function, polynomial interpolation gives fair results. The function should not have any singularities over the range of interpolation. In addition there should not be any pole in the vicinity of the complex plane near the portion of the real axis corresponding to the range of interpolation. If the function has singularities it is recommended to use rational functions — that is the quotient of two polynomials — instead [Press et al.].

In this chapter we discuss 3 interpolation functions: Lagrange interpolation polynomial, a diagonal rational function (Bulirsch-Stoer interpolation) and cubic spline. Furthermore, we show 3 different implementation of the Lagrange interpolation polynomial: direct implementation of Lagrange's formula, Newton's algorithm and Neville's algorithm. Figure 3.1 shows how the classes corresponding to the different interpolation methods described in this chapter are related to each other.

Definition

The Lagrange interpolation polynomial is the unique polynomial of minimum degree going through the sample points. The degree of the polynomial is equal to the number of supplied points minus one. A diagonal rational function is the quotient of two polynomials where the degree of the polynomial in the numerator is at most equal to that of the denominator. Cubic spline uses piece-wise interpolation with polynomials




Figure 3.2: Example of interpolation with the Lagrange interpolation polynomial

but limits the degree of each polynomial to 3 (hence the adjective cubic).

Examples

Before selecting an interpolation method the user must investigate the validity of the interpolated function over the range of its intended use. Let us illustrate this remark with an example from high-energy physics, that, in addition, will expose the limitation of the methods exposed in this chapter.

Figure 3.1 shows sample points — indicated by crosses — representing correction to the energy measured within a gamma ray detector made of several densely packed crystals. The energy is plotted on a logarithmic scale. The correction is caused by the absorption of energy in the wrapping of each crystal. The sample points were computed using a simulation program², each point requiring several hours of computing time. Interpolation over these points was therefore used to allow a quick computation of the correction at any energy. This is the main point of this example: the determination of each point was expensive in terms of computing time, but the function represented by these points is continuous enough to be interpolated. The simulation program yields results with good precision so that the resulting data are not subjected to fluctuation.

²This program - EGS written by Ralph Nelson of the Stanford Linear Accelerator Center (SLAC) - simulates the absorption of electromagnetic showers inside matter. Besides being used in high-energy physics this program is also used in radiology to dimension detectors of PET scanners and other similar radiology equipment.




Figure 3.3: Comparison between Lagrange interpolation and interpolation with a rational function

The gray thick line in figure 3.1 shows the Lagrange interpolation polynomial obtained from the sample points. It readily shows limitations inherent to the use of interpolation polynomials. The reader can see that for values above 6.5 — corresponding to an energy of 500 MeV — the interpolated function does not reproduce the curve corresponding to the sample points. In fact, above 4.0 — that is, 50 MeV on the scale of figure 3.1 — the correction is expected to be a linear function of the logarithm of the energy.

Figure 3.1 shows a comparison between the Lagrange interpolation polynomial (gray thick line) and interpolation with a rational function (black dotted line) using the same sample points as in figure 3.1. The reader can see that, in the high-energy region (above 4 on the scale of figure 3.1) the rational function does a better job than the Lagrange polynomial. Between the first two points, however, the rational function fails to reproduce the expected behavior.

Figure 3.1 shows a comparison between the Lagrange interpolation polynomial (gray thick line) and cubic spline interpolation (black dotted line) using the same sample points as in figure 3.1. The reader can see that, in the high-energy region (above 4 on the scale of figure 3.1) the cubic spline does a better job than the Lagrange polynomial. In fact, since the dependence is linear over that range, the cubic spline reproduces the theoretical dependence exactly. In the low energy region, however, cubic spline interpolation fails to reproduce the curvature of the theoretical function because of the limitation of the




Figure 3.4: Comparison of Lagrange interpolation and cubic spline




Figure 3.5: Example of misbehaving interpolation

polynomial's degree.

A final example shows a case where interpolation should not be used. Here the sample points represent the dependence of the probability that a coin mechanism accepts a wrong coin as a function of an adjustable threshold. The determination of each point requires 5-10 minutes of computing time. In this case, however, the simulation was based on using experimental data. Contrary to the points of figure 3.1 the points of figure 3.1 are subjected to large fluctuations, because the sample points have been derived from measured data. Thus, interpolation does not work.

As in figure 3.1, the gray thick line is the Lagrange interpolation polynomial and the black dotted line is the cubic spline. Clearly the Lagrange interpolation polynomial is not giving any reasonable interpolation. Cubic spline is not really better as it tries very hard to reproduce the fluctuations of the computed points. In this case, a polynomial fit (c.f. section 10.8) is the best choice: the thin black line shows the result of a fit with a 3rd degree polynomial. Another example of unstable interpolation is given in section 10.8 (figure 10.5).

Three implementations of Lagrange interpolation

Once you have verified that a Lagrange interpolation polynomial can be used to perform reliable interpolation over the sample points, you must choose among 3 algorithms to compute the Lagrange interpolation polynomial: direct Lagrange formula, Newton's

algorithm and Neville's algorithm.

Newton's algorithm stores intermediate values which only depends on the sample points. It is thus recommended, as it is the fastest method to interpolate several values over the same sample points. Newton's algorithm is the method of choice to compute a function from tabulated values.

Neville's algorithm gives an estimate of the numerical error obtained by the interpolation. It can be used when such information is needed. Romberg integration, discussed in section 6.4, uses Neville's method for that reason.

3.2 Lagrange interpolation

Let us assume a set of numbers x_0, \dots, x_n and the corresponding function's values y_0, \dots, y_n . There exist a unique polynomial $P_n(x)$ of degree n such that $P_n(x_i) = y_i$ for all $i = 0, \dots, n$. This polynomial is the Lagrange interpolation polynomial whose expression is given by [Knudth 2]:

$$P_n(x) = \sum_{i=0}^n \frac{\prod_{j \neq i} (x - x_j)}{\prod_{j \neq i} (x_i - x_j)} y_i. \quad (3.1)$$

For example, the Lagrange interpolation polynomial of degree 2 on 3 points is given by:

$$P_2(x) = \frac{(x - x_1)(x - x_2)}{(x_0 - x_1)(x_0 - x_2)} y_0 + \frac{(x - x_0)(x - x_2)}{(x_1 - x_0)(x_1 - x_2)} y_1 + \frac{(x - x_0)(x - x_1)}{(x_2 - x_0)(x_2 - x_1)} y_2 \quad (3.2)$$

The computation of the polynomial occurs in the order of $O(n^2)$ since it involves a double iteration. One can save the evaluation of a few products by rewriting equation 3.1 as:

$$P_n(x) = \prod_{i=0}^n (x - x_i) \sum_{i=0}^n \frac{y_i}{(x - x_i) \prod_{j \neq i} (x_i - x_j)}. \quad (3.3) \quad \Leftarrow \text{Main equation}$$

Of course, equation 3.3 cannot be evaluated at the points defining the interpolation. This is easily solved by returning the defining values as soon as one of the first products becomes zero during the evaluation.

3.2.1 Lagrange interpolation — Smalltalk implementation

The object responsible to implement Lagrange interpolation is defined uniquely by the sample points over which the interpolation is performed. In addition it should behave as a function. In other words it should implement the behavior of a one-variable function as discussed in section 2.1.1. For example linear interpolation behaves as follows:

Code example 3.1

```

| interpolator |
interpolator := DhbLagrangeInterpolator points: ( Array with: 1 @ 2
with: 3 @ 1).

interpolator value: 2.2

```

Figure 3.1 with the box LagrangeInterpolator grayed.

In this example, one creates a new instance of the class `DhbLagrangeInterpolator` by sending the message `points:` to the class `DhbLagrangeInterpolator` with the collection of sample points as argument. The newly created instance is stored in the variable `interpolator`. The next line shows how to compute an interpolated value.

The creation method `points:` takes as argument the collection of sample points. However, it could also accept any object implementing a subset of the methods of the class `Collection` — namely the methods `size`, `at:` and, if we want to be able to add new sample points, `add:`.

One can also spare the creation of an explicit collection object by implementing these collection methods directly in the Lagrange interpolation class. Now, one can also perform interpolation in the following way:

Code example 3.2

```
| interpolator deviation |
interpolator := DhbLagrangeInterpolator new.
1 to: 45 by: 2 do:
    [ :x | interpolator add: x @ (x degreesToRadians sin)].
deviation := (interpolator value: 8) -(8 degreesToRadians sin).
```

The code above creates an instance of the class `DhbLagrangeInterpolator` with an empty collection of sample points. It then adds sample points one by one directly into the interpolator object. Here the sample points are tabulated values of the sine function for odd degree values between 1 and 45 degree. The final line of the code compares the interpolated value with the correct one.

Listing 3.1 shows the full code of the class implementing the interface shown above.

The class `DhbLagrangeInterpolator` is implemented with a single instance variable containing the collection of sample points. Each point contains a pair of values (x_i, y_i) and is implemented with object of the base class `Point` since an instance of `Point` can contain any type of object in its coordinates. There are two creation methods, `points:` and `new`, depending on whether the sample points are supplied as an explicit object or not. Each creation method calls in turn an initialization method, respectively `initialize:` and `initialize`.

The method `points:` takes as argument the collection of the sample points. This object must implement the following methods of the class `Collection`: `size`, `at:` and `add:`. If the class is created with the method `new` an implicit collection object is created with the method `defaultSamplePoints`. This arrangement allows subclasses to select another type of collection if needed. The default collection behavior implemented by the class `DhbLagrangeInterpolator` is minimal, however. If there is a need for more flexible access to the collection of sample points, a proper collection object or a special purpose object should be used.

The interpolation itself is implemented within the single method `value:`. This method is unusually long for object-oriented programming standards. In this case, however, there is no compelling reason to split any portion of the algorithm into a separate method. Moreover, splitting the method would increase the computing time.

A final discussion should be made about the two methods `xPointAt:` and `yPointAt:`. In principle, there is no need for these methods as the value could be grabbed directly from the collection of points. If one needs to change the implementation of the

point collection in a subclass, however, only these two methods need to be modified. Introducing this kind of construct can go a long way in program maintenance.

Listing 3.1 Smalltalk implementation of the Lagrange interpolation

Class **DhbLagrangeInterpolator**

Subclass of Object

Instance variable names: pointCollection

Class methods

new

^super new initialize

points: aCollectionOfPoints

^self new initialize: aCollectionOfPoints

Instance methods

add: aPoint

^pointCollection add: aPoint

defaultSamplePoints

^OrderedCollection new

initialize

^self initialize: self defaultSamplePoints

initialize: aCollectionOfPoints

pointCollection := aCollectionOfPoints.

^self

value: aNumber

| norm dx products answer size |

norm := 1.

size := pointCollection size.

products := Array new: size.

products atAllPut: 1.


```

1 to: size
do: [:n |
    dx := aNumber - ( self xPointAt: n).
    dx = 0
    ifTrue: [ ^( self yPointAt: n)].
    norm := norm * dx.
    1 to: size
    do: [:m |
        m = n
        ifFalse:[ products at: m put: ( ( self
xPointAt: m) - ( self xPointAt: n)) * ( products at: m))].
    ].
answer := 0.
1 to: size do:
    [ :n | answer := ( self yPointAt: n) / ( ( products at: n) *
        ( aNumber - ( self xPointAt: n))) + answer].
^norm * answer

xPointAt: anInteger

^( pointCollection at: anInteger) x

yPointAt: anInteger

^( pointCollection at: anInteger) y

```

Figure 3.1 with the box
LagrangeInterpolator grayed.

3.2.2 Lagrange interpolation — Java implementation

The object responsible to implement Lagrange interpolation is defined uniquely by the sample points over which the interpolation is performed. Unlike in Smalltalk, however, point objects cannot be used to hold the values because the Java class `Point` is only defined for integer values.

Thus, an interface called `PointSeries` was created to define the behavior of the object containing the sample points. This interface is shown in Listing 3.2.

Listing 3.2 Java interface for point series

```

package DhbInterfaces;

/**
 * PointSeries is an interface used by many classes of the package numericalMethods.
 *

```

```

* A PointSeries has the responsibility of handling mathematical
* points in 2-dimensional space.
* It is a BRIDGE to a vector containing the points.
*
* @author Didier H. Besset
*/
public interface PointSeries
{

/**
 * Returns the number of points in the series.
 */
public int size();

/**
 * Returns the x coordinate of the point at the given index.
 * @param index the index of the point.
 * @return x coordinate
 */
public double xValueAt( int index);

/**
 * Returns the y coordinate of the point at the given index.
 * @param index the index of the point.
 * @return y coordinate
 */
public double yValueAt( int index);
}

```

Listing 3.3 shows the class `Curve` implementing the minimum functionality of the interface `PointSeries`. This class is used in a variety of problems where handling geometrical points is required.

Listing 3.3 A possible concrete implementation of the interface `PointSeries`

```

package DhbScientificCurves;

import java.util.Vector;
import DhbInterfaces.PointSeries;

/**
 * A Curve is a series of points. A point is implemented as an array
 * of two doubles. The points are stored in a vector so that points
 * can be added or removed.
 *
 * @author Didier H. Besset
 */
public class Curve implements PointSeries

```

```
{
    /**
     * Vector containing the points.
     */
    protected Vector points;

    /**
     * Constructor method. Initializes the vector.
     */
    public Curve()
    {
        points = new Vector();
    }

    /**
     * Adds a point to the curve defined by its 2-dimensional coordinates.
     * @param x double x-coordinate of the point
     * @param y double y-coordinate of the point
     */
    public void addPoint( double x, double y)
    {
        double point[] = new double[2];
        point[0] = x;
        point[1] = y;
        points.addElement( point);
    }

    /**
     * Removes the point at the specified index.
     * @param int index of the point to remove
     */
    public void removePointAt( int index)
    {
        points.removeElementAt( index);
    }

    /**
     * @return int the number of points in the curve.
     */
    public int size()
    {
        return points.size();
    }

    /**
     * @return double the x coordinate of the point at the given index.
     * @param int index the index of the point.
     */
    public double xValueAt( int index)
```

```

{
    return ((double[]) points.elementAt( index))[0];
}
/**
 * @return double the y coordinate of the point at the given index.
 * @param int index the index of the point.
 */
public double yValueAt( int index)
{
    return ((double[]) points.elementAt( index))[1];
}
}

```

The object implementing Lagrange interpolation should implement the interface `OneVariableFunction` discussed in section 2.1.1. This completes the public interface to the object. Thus an example of linear interpolation behaves as follows:

Code example 3.3

```

DhbScientificCurves.Curve points = new DhbScientificCurves.Curve();
points.addPoint( 1, 2);
points.addPoint( 3, 1);
(new LagrangeInterpolator( points)).value( 2.2)

```

The first line creates the collection of sample points. The next two lines populate the collection with given sample points. The last line creates an instance of the class `LagrangeInterpolator` and evaluates its value at the desired point.

The next example shows how to tabulate a function.

Code example 3.4

```

double radian = Math.PI / 180;
DhbScientificCurves.Curve points = new DhbScientificCurves.Curve();
for( int i = 1; i <= 45; i += 2)
    points.addPoint( i, Math.sin( i * radian));
LagrangeInterpolator interpolator =
    new LagrangeInterpolator( points);
double deviation =interpolator.value( 8) -
    Math.sin( 8 * radian);

```

The code above creates an instance of `Curve`, an object containing the sample points. This object is then populated with the values of the sine function for odd degree values between 1 and 45 degree. Then, an instance of the class `LagrangeInterpolator` is created on the sample points. The last line of the code computes the difference between the interpolated value and the exact value.

Listing 3.4 shows the full code of the class implementing the interface shown above.

The class `LagrangeInterpolator` is implemented with a single instance variable containing the sample points. Each point contains the pair of values (x_i, y_i) defining the function. The object containing the sample points can be any object implementing

the interface `PointSeries`. The single constructor method takes the object containing the sample points as sole argument.

The interpolation itself is implemented within the single method `value()`. This method is unusually long by object-oriented programming standards. In this case, however, there is no compelling reason to split any portion of the algorithm into a separate method. Moreover, splitting the method would increase the computing time.

Listing 3.4 Java implementation of the Lagrange interpolation

```
package DhbInterpolation;

import DhbInterfaces.OneVariableFunction;
import DhbInterfaces.PointSeries;
/**
 * A LagrangeInterpolator can be used to interpolate values between
 * a series of 2-dimensional points. The interpolation function is
 * the Lagrange interpolation polynomial of a degree equal to the
 * number of points in the series minus one.
 *
 * @author Didier H. Besset
 */
public class LagrangeInterpolator implements OneVariableFunction
{
    /**
     * Points containing the values.
     */
    protected PointSeries points;
    /**
     * Constructor method.
     * @param pts the series of points.
     * @see PointSeries
     */
    public LagrangeInterpolator(PointSeries pts)
    {
        points = pts;
    }
    /**
     * Computes the interpolated y value for a given x value.
     * @param aNumber x value.
     * @return interpolated y value.
     */
    public double value( double aNumber)
    {
        double norm = 1.0;
        int size = points.size();
        double products[] = new double[size];
```

```

for ( int i = 0; i < size; i++)
    products[i] = 1;
double dx;
for ( int i = 0; i < size; i++)
{
    dx = aNumber - points.xValueAt( i);
    if ( dx == 0 )
        return points.yValueAt(i);
    norm *= dx;
    for ( int j = 0; j < size; j++)
    {
        if ( i != j)
            products[j] *= points.xValueAt(j)
                          - points.xValueAt(i);
    }
}
double answer = 0.0;
for ( int i = 0; i < size; i++)
    answer += points.yValueAt(i)
             / (products[i] * (aNumber - points.xValueAt(i)));
return norm * answer;
}
}

```

3.3 Newton interpolation

If one must evaluate the Lagrange interpolation polynomial for several values, it is clear that the Lagrange's formula is not efficient. Indeed a portion of the terms in the summation of equation 3.3 depends only on the sample points and does not depend on the value at which the polynomial is evaluated. Thus, one can speed up the evaluation of the polynomial if the invariant parts are computed once and stored.

If one writes the Lagrange interpolation polynomial using a generalized Horner expansion, one obtains the Newton's interpolation formula given by [Knudth 2]:

$$P_n(x) = \alpha_0 + (x - x_0) \cdot [\alpha_1 + (x - x_1) \cdot [\cdots [\alpha_{n-1} + \alpha_n \cdot (x - x_1)]]] \quad (3.4) \quad \Leftarrow \text{Main equation}$$

The coefficients α_i are obtained by evaluating divided differences as follows:

$$\begin{cases} \Delta_i^0 &= y_i \\ \Delta_i^k &= \frac{\Delta_i^{k-1} - \Delta_{i-1}^{k-1}}{x_i - x_{i-k}} \quad \text{for } k = 1, \dots, n \\ \alpha_i &= \Delta_i^i \end{cases} \quad (3.5)$$

Once the coefficients α_i have been obtained, they can be stored in the object and the generalized Horner expansion of equation 3.4 can be used.

The time to evaluate the full Newton's algorithm — that is computing the coefficients and evaluating the generalized Horner expansion — is about twice the time needed to perform a direct Lagrange interpolation. The evaluation of the generalized Horner expansion alone, however, has an execution time of $O(n)$ and is therefore much faster than the evaluation of a direct Lagrange interpolation which goes as $O(n^2)$. Thus, as soon as one needs to interpolate more than 2 points over the same point sample, Newton's algorithm is more efficient than direct Lagrange interpolation.

Note: The implementations of Newton's interpolation algorithm are identical in both languages. Thus, the reader can skip one of the two next subsections without losing anything.

3.3.1 Newton interpolation — General implementation

Figure 3.1 with the box NewtonInterpolator grayed.

The object implementing Newton's interpolation algorithm is best implemented as a subclass of the class DhbLagrangeInterpolator because all methods used to handle the sample points can be reused. This also allows us to keep the interface identical. It has an additional instance variable needed to store the coefficients α_i . Only 4 new methods are needed.

Since the client object can add new sample points at will, one cannot be sure of when it is safe to compute the coefficients. Thus, computing the coefficients is done with lazy initialization. The method `value`: first checks whether the coefficients α_i have been computed. If not, the method `computeCoefficients` is called. Lazy initialization is a technique widely used in object oriented programming whenever some value needs only be computed once.

The generalized Horner expansion is implemented in the method `value`:

If a new sample point is added, the coefficient eventually stored in the object are no longer valid. Thus, the method `add`: first calls the method `resetCoefficients` and then calls the method `add`: of the superclass. The method `resetCoefficients` makes sure that the coefficients will be computed anew at the next evaluation of the interpolation polynomial. The method `resetCoefficients` has been implemented as a separate method so that the reset mechanism can be reused by any subclass.

Another reason to keep the method `resetCoefficients` separate is that it must also be called before doing an interpolation if the sample points have been modified directly by the client application after the last interpolation has been made. An alternative is to implement the Observable/Observer pattern so that resetting of the coefficients happens implicitly using events. However, since modifying the sample points between interpolation should only be a rare occasion when using Newton's algorithm³ our proposed implementation is much simpler.

³If modification of the sample points is not a rare occasion, then Newton's algorithm has no advantage over direct Lagrange interpolation or Neville's algorithm. Those algorithms should be used instead of Newton's algorithm.

3.3.2 Newton interpolation — Smalltalk implementation

Listing 3.5 shows the complete implementation in Smalltalk. The class `NewtonInterpolator` is a subclass of class `LagrangeInterpolator`. The code examples 3.1 and 3.2 can directly be applied to Newton interpolation after replacing the class name `DhbLagrangeInterpolator` with `DhbNewtonInterpolator`.

The generalized Horner expansion is implemented in the method `value:` using explicit indices. One could have used the method `inject:into:` as it was done for Horner's formula when evaluating polynomials. In this case, however, one must still keep track of the index to retrieve the sample point corresponding to each coefficient. Thus, one gains very little in compactness.

Listing 3.5 Smalltalk implementation of the Newton interpolation

```
Class          DhbNewtonInterpolator
Subclass of   DhbLagrangeInterpolator
```

```
Instance variable names: coefficients
```

```
Instance methods
```

```
add: aPoint
```

```
self resetCoefficients.
^super add: aPoint
```

```
computeCoefficients
```

```
  | size k1 kn |
  size := pointCollection size.
  coefficients := ( 1 to: size) collect: [ :n | self yPointAt: n].
  1 to: (size - 1)
    do: [ :n |
      size to: ( n + 1) by: -1
        do: [ :k |
          k1 := k - 1.
          kn := k - n.
          coefficients at: k put: ( (( coefficients at:
                                     k) - ( coefficients at: k1))
                                   / ((self xPointAt: k) -
                                       (self xPointAt: kn))).
        ]
      ].
  ].
```

```
resetCoefficients
```

```
coefficients := nil.
```



```

value: aNumber

  | answer size |
  coefficients isNil
    ifTrue: [ self computeCoefficients].
  size := coefficients size.
  answer := coefficients at: size.
  (size - 1) to: 1 by: -1
    do: [ :n | answer := answer * ( aNumber - (self xPointAt:
                                     n)) + ( coefficients at: n)].
  ^answer

```

3.3.3 Newton interpolation — Java implementation

Listing 3.6 shows the complete implementation in Java. The class `NewtonInterpolator` is a subclass of class `LagrangeInterpolator`. The code examples 3.3 and 3.4 can directly be applied to Newton interpolation after replacing the class name `LagrangeInterpolator` with `NewtonInterpolator`.

Java has an additional argument against implementing a `Observable/Observer` pattern to automatically reset the coefficients: an explicit finalization method is required to detach the link between `Observer` and `Observable` when the interpolator is no longer needed. Our simple implementation does not require finalization.

Listing 3.6 Java implementation of the Newton interpolation

```

package DhbInterpolation;

import DhbInterfaces.PointSeries;
/**
 * A NewtonInterpolator can be used to interpolate values between
 * a series of 2-dimensional points. The interpolation function is
 * the Lagrange interpolation polynomial of a degree equal to
 * the number of points in the series minus one. The coefficients
 * of the polynomial are stored, speeding up interpolation for a
 * series of values.
 *
 * @author Didier H. Besset
 */
public class NewtonInterpolator extends LagrangeInterpolator
{
  /**
   * Polynomial coefficient (modified Horner expansion).
   */
  protected double coefficients[];

```

```

/**
 * Constructor method.
 * @param pts interfaces.PointSeries
 */
public NewtonInterpolator( PointSeries pts) {
    super(pts);
}
/**
 * Computes the coefficients of the interpolation polynomial.
 */
private void computeCoefficients()
{
    int size = points.size();
    int n;
    int k;
    int k1;
    int kn;
    coefficients = new double[size];
    for ( n = 0; n < size; n++)
        coefficients[n] = points.yValueAt(n);
    size -= 1;
    for ( n = 0; n < size; n++)
    {
        for ( k = size; k > n; k--)
        {
            k1 = k - 1;
            kn = k - (n + 1);
            coefficients[k] = ( coefficients[k] - coefficients[k1])
                / ( points.xValueAt(k)
                    - points.xValueAt(kn));
        }
    }
}
/**
 * Forces a new computation of the coefficients. This method must be
 * called whenever the series of points defining the interpolator is
 * modified.
 */
public void resetCoefficients()
{
    coefficients = null;
}
/**
 * Computes the interpolated y value for a given x value.

```

```

* @param aNumber x value.
* @return interpolated y value.
*/
public double value( double aNumber)
{
    if ( coefficients == null)
        computeCoefficients();
    int size = coefficients.length;
    double answer = coefficients[--size];
    while ( --size >= 0 )
        answer = answer * ( aNumber - points.xValueAt(size))
            + coefficients[size];

    return answer;
}
}

```

3.4 Neville interpolation

Neville's algorithm uses a successive approximation approach implemented in practice by calculating divided differences recursively. The idea behind the algorithm is to compute the value of the interpolation's polynomials of all degrees between 0 and n . This algorithm assumes that the sample points have been sorted in increasing order of abscissa.

Let $P_j^i(x)$ be the (partial) Lagrange interpolation polynomials of degree i defined by the sets of values x_j, \dots, x_{j+i} and the corresponding function's values y_j, \dots, y_{j+i} . From equation 3.1 one can derive the following recurrence formula [Press et al.]:

Main equation \Rightarrow

$$P_j^i(x) = \frac{(x - x_{i+j}) P_j^{i-1}(x) + (x_j - x) P_{j+1}^{i-1}(x)}{x_j - x_{i+j}} \quad \text{for } j < i. \quad (3.6)$$

The initial values $P_j^0(x)$ are simply y_j . The value of the final Lagrange's polynomial is $P_0^n(x)$.

Neville's algorithm introduces the differences between the polynomials of various degrees. One defines:

$$\begin{cases} \Delta_{j,i}^{\text{left}}(x) &= P_j^i(x) - P_j^{i-1}(x) \\ \Delta_{j,i}^{\text{right}}(x) &= P_j^i(x) - P_{j+1}^{i-1}(x) \end{cases} \quad (3.7)$$

From the definition above and equation 3.6 one derives a pair of recurrence formulae for the differences:

$$\begin{cases} \Delta_{j,i+1}^{\text{left}}(x) &= \frac{x_i - x}{x_j - x_{i+j+1}} \left(\Delta_{j+1,i}^{\text{left}}(x) - \Delta_{j,i}^{\text{right}}(x) \right) \\ \Delta_{j,i+1}^{\text{right}}(x) &= \frac{x_{i+j+1} - x}{x_j - x_{i+j+1}} \left(\Delta_{j+1,i}^{\text{left}}(x) - \Delta_{j,i}^{\text{right}}(x) \right) \end{cases} \quad (3.8)$$

In practice two arrays of differences — one for left and one for right — are allocated. Computation of each order is made within the same arrays. The differences of the last order can be interpreted as an estimation of the error made in replacing the function by the interpolation's polynomial.

Neville's algorithm is faster than the evaluation of direct Lagrange's interpolation for a small number of points (smaller than about 7⁴). Therefore a simple linear interpolation is best performed using Neville's algorithm. For a large number of points, it becomes significantly slower.

3.4.1 Neville interpolation — General implementation

The object implementing Neville's interpolation's algorithm is best implemented as a subclass of the class `LagrangeInterpolator` since the methods used to handle the sample points can be reused. This also allows us to keep the interface identical.

The new class has two additional instance variables used to store the finite differences $\Delta_{j,i}^{\text{left}}(x)$ and $\Delta_{j,i}^{\text{right}}(x)$ for all j . These instance variables are recycled for all i . Only a few additional methods are needed.

The method `valueAndError:` implementing Neville's algorithm returns an array with two elements: the first element is the interpolated value and the second is the estimated error. The method `value:` calls the former method and returns only the interpolated value.

Unlike other interpolation algorithms, the method `valueAndError:` is broken into smaller methods because the mechanics of computing the finite differences will be reused in the Bulirsch-Stoer algorithm. The method `valueAndError:` begins by calling the method `initializeDifferences:` to populate the arrays containing the finite differences with their initial values. These arrays are created if this is the first time they are used with the current sample points. This prevents unnecessary memory allocation. Then, at each iteration the method `computeDifference:at:order:` computes the differences for the current order.

3.4.2 Neville interpolation — Smalltalk implementation

Listing 3.7 shows the implementation of Neville's algorithm in Smalltalk. The class `DhbNevilleInterpolator` is a subclass of class `DhbLagrangeInterpolator`. The code examples 3.1 and 3.2 can directly be applied to Neville interpolation after replacing the class name `DhbLagrangeInterpolator` with `DhbNevilleInterpolator`. An example of interpolation using the returned estimated error is given in section 6.4.2.

The method `defaultSamplePoints` overrides that of the superclass to return a sorted collection. Thus, each point added to the implicit collection is automatically sorted by increasing abscissa as required by Neville's algorithm.

Figure 3.1 with the box `NevilleInterpolator` grayed.

Listing 3.7 Smalltalk implementation of Neville's algorithm

Class `DhbNevilleInterpolator`

⁴c.f. footnote 8 on page 106

Subclass of `DhbLagrangeInterpolator`

Instance variable names: `leftErrors` `rightErrors`

Instance methods

`computeDifference: aNumber at: anInteger1 order: anInteger2`

```

| leftDist rightDist ratio |
leftDist := ( self xPointAt: anInteger1 ) - aNumber.
rightDist := ( self xPointAt: ( anInteger1 + anInteger2 ) ) -
              aNumber.
ratio := ( ( leftErrors at: ( anInteger1 + 1 ) ) - ( rightErrors
              at: anInteger1 ) ) / ( leftDist - rightDist ).
leftErrors at: anInteger1 put: ratio * leftDist.
rightErrors at: anInteger1 put: ratio * rightDist.

```

`defaultSamplePoints`

```

^SortedCollection sortBlock: [ :a :b | a x < b x ]

```

`initializeDifferences: aNumber`

```

| size nearestIndex dist minDist |
size := pointCollection size.
leftErrors size = size
  ifFalse: [ leftErrors := Array new: size.
             rightErrors := Array new: size.
             ].
minDist := ( ( self xPointAt: 1 ) - aNumber ) abs.
nearestIndex := 1.
leftErrors at: 1 put: ( self yPointAt: 1 ).
rightErrors at: 1 put: leftErrors first.
2 to: size do:
  [ :n |
    dist := ( ( self xPointAt: n ) - aNumber ) abs.
    dist < minDist
      ifTrue: [ dist = 0
                ifTrue: [ ^n negated ].
                nearestIndex := n.
                minDist := dist.
              ].
    leftErrors at: n put: ( self yPointAt: n ).
    rightErrors at: n put: ( leftErrors at: n ).
  ].
^nearestIndex

```

value: aNumber

^(self valueAndError: aNumber) first

valueAndError: aNumber

```

| size nearestIndex answer error |
nearestIndex := self initializeDifferences: aNumber.
nearestIndex < 0
    ifTrue: [ ^Array with: ( self yPointAt: nearestIndex negated)
                        with: 0].
answer := leftErrors at: nearestIndex.
nearestIndex := nearestIndex - 1.
size := pointCollection size.
1 to: ( size - 1) do:
    [ :m |
        1 to: ( size - m) do:
            [ :n | self computeDifference: aNumber at: n order: m].
        size - m > ( 2 * nearestIndex)
            ifTrue: [ error := leftErrors at: ( nearestIndex + 1)
                    ]
            ifFalse:[ error := rightErrors at: ( nearestIndex).
                    nearestIndex := nearestIndex - 1.
                    ].
        answer := answer + error.
    ].
^Array with: answer with: error abs

```

3.4.3 Neville interpolation — Java implementation

Listing 3.8 shows the implementation of Neville's algorithm in Java. The class `NevilleInterpolator` is a subclass of class `LagrangeInterpolator`. The code examples 3.3 and 3.4 can directly be applied to Neville interpolation after replacing the class name `LagrangeInterpolator` with `NevilleInterpolator`. An example of interpolation using the returned estimated error is given in section 6.4.3.

Note: Since there is no sorted collection in Java, this class supposes that the calling application supplies the point sorted by ascending value of the abscissa.

Listing 3.8 Java implementation of Neville interpolation

```

package DhbInterpolation;

/**
 * NevilleInterpolator
 *
 * @author Didier H. Besset
 */
public class NevilleInterpolator extends LagrangeInterpolator
{
    protected double[] leftErrors = null;
    protected double[] rightErrors = null;
    /**
     * Constructor method.
     * @param pts DhbInterfaces.PointSeries contains the points sampling
     *           the function to interpolate.
     * @exception java.lang.IllegalArgumentException points are not sorted
     *           in increasing x values.
     */
    public NevilleInterpolator(DhbInterfaces.PointSeries pts)
    {
        super(pts);
        for ( int i = 1; i < pts.size(); i++ )
        {
            if ( pts.xValueAt( i - 1 ) >= pts.xValueAt( i ) )
                throw new IllegalArgumentException
                    ( "Points must be sorted in increasing x value");
        }
    }
    /**
     * @param m int    order of the difference
     * @param n int    index of difference
     * @param x double  argument
     */
    protected void computeNextDifference( int m, int n, double x)
    {
        double leftDist = points.xValueAt(n) - x;
        double rightDist = points.xValueAt(n + m + 1) - x;
        double ratio = ( leftErrors[n+1] - rightErrors[n])
                       / ( leftDist - rightDist);
        leftErrors[n] = ratio * leftDist;
        rightErrors[n] = ratio * rightDist;
    }
    /**
     * @return int
     */
    private int initializeDifferences( double x)

```

```

{
    int size = points.size();
    if ( leftErrors == null || leftErrors.length != size )
    {
        leftErrors = new double[ size];
        rightErrors = new double[ size];
    }
    double minDist = Math.abs( x - points.xValueAt(0));
    if ( minDist == 0 )
        return -1;
    int nearestIndex = 0;
    leftErrors[0] = points.yValueAt(0);
    rightErrors[0] = leftErrors[0];
    for ( int n = 1; n < size; n++)
    {
        double dist = Math.abs( x - points.xValueAt(n));
        if ( dist < minDist )
        {
            if ( dist == 0)
                return -n-1;
            minDist = dist;
            nearestIndex = n;
        }
        leftErrors[n] = points.yValueAt(n);
        rightErrors[n] = leftErrors[n];
    }
    return nearestIndex;
}
/**
 * @return double
 * @param aNumber double
 */
public double value( double aNumber)
{
    return valueAndError( aNumber)[0];
}
/**
 * @return double[]  an array with 2 elements:
 *      [0] interpolated value, [1] estimated error
 * @param x double
 */
public double[] valueAndError( double x)
{
    double[] answer = new double[2];
    int nearestIndex = initializeDifferences( x);
    if ( nearestIndex < 0 )

```



```

{
    answer[0] = points.yValueAt(-1-nearestIndex);
    answer[1] = 0;
    return answer;
}
int size = points.size();
answer[0] = leftErrors[ nearestIndex--];
double leftDist, rightDist, ratio;
for ( int m = 0; m < size - 1; m++)
{
    for ( int n = 0; n < size - 1 - m; n++)
    {
        computeNextDifference( m, n, x);
    }
    answer[1] = ( size - m > 2 * ( nearestIndex + 1 ) )
        ? leftErrors[ nearestIndex + 1]
        : rightErrors[ nearestIndex--];
    answer[0] += answer[1];
}
return answer;
}
}

```

3.5 Bulirsch-Stoer interpolation

If the function to interpolate is known to have poles⁵ in the vicinity of the real axis over the range of the sample points a polynomial cannot do a good interpolation job [Press et al.].

In this case it is better to use rational function, that is a quotient of two polynomials as defined hereafter:

$$R(x) = \frac{P(x)}{Q(x)} \quad (3.9)$$

The coefficients of both polynomials are only defined up to a common factor. Thus, if p is the degree of polynomial $P(x)$ and q is the degree of polynomial $Q(x)$, we must have the relation $p + q + 1 = n$ where n is the number of sample points. This of course is not enough to restrict the variety of possible rational functions.

Bulirsch and Stoer have proposed an algorithm for a rational function where $p = \lfloor \frac{n-1}{2} \rfloor$. This means that q is either equal to p if the number of sample points is odd or equal to $p + 1$ if the number of sample points is even. Such a rational function is called a diagonal rational function. This restriction, of course, limits the type of function shapes that can be interpolated.

The Bulirsch-Stoer algorithm is constructed like Neville's algorithm: finite differences are constructed until all points have been taken into account.

⁵That is, a singularity in the complex plane.

Let $R_j^i(x)$ be the (partial) diagonal rational functions of order i defined by the sets of values x_j, \dots, x_{j+i} and the corresponding function's values y_j, \dots, y_{j+i} . As in the case of Neville's algorithm, one can establish a recurrence formula between functions of successive orders. We have [Press et al.]:

$$R_j^i(x) = R_{j+1}^{i-1}(x) + \frac{R_{j+1}^{i-1}(x) - R_j^{i-1}(x)}{\frac{x-x_j}{x-x_{i+j}} \left(1 - \frac{R_{j+1}^{i-1}(x) - R_j^{i-1}(x)}{R_{j+1}^{i-1}(x) - R_{j+1}^{i-2}(x)}\right)} \quad \text{for } j < i. \quad (3.10)$$

The initial values $R_j^0(x)$ are simply y_j . The final rational function is $R_0^n(x)$.

Like in Neville's algorithm one introduces the differences between the functions of various orders. One defines:

$$\begin{cases} \Delta_{j,i}^{\text{left}}(x) &= R_j^i(x) - R_j^{i-1}(x) \\ \Delta_{j,i}^{\text{right}}(x) &= R_j^i(x) - R_{j+1}^{i-1}(x) \end{cases} \quad (3.11)$$

From the definition above and equation 3.10 one derives a pair of recurrence formulae for the differences:

$$\begin{cases} \Delta_{j,i+1}^{\text{left}}(x) &= \frac{\frac{x-x_j}{x-x_{i+j+1}} \Delta_{j,i}^{\text{right}}(x) [\Delta_{j+1,i}^{\text{left}}(x) - \Delta_{j,i}^{\text{right}}(x)]}{\frac{x-x_j}{x-x_{i+j+1}} \Delta_{j,i}^{\text{right}}(x) - \Delta_{j+1,i}^{\text{left}}(x)} \\ \Delta_{j,i}^{\text{right}}(x) &= \frac{\frac{x-x_j}{x-x_{i+j+1}} \Delta_{j+1,i}^{\text{left}}(x) [\Delta_{j+1,i}^{\text{left}}(x) - \Delta_{j,i}^{\text{right}}(x)]}{\frac{x-x_j}{x-x_{i+j+1}} \Delta_{j,i}^{\text{right}}(x) - \Delta_{j+1,i}^{\text{left}}(x)} \end{cases} \quad (3.12)$$

Like for Neville's algorithm, two arrays of differences — one for left and one for right — are allocated. Computation of each order is made within the same arrays. The differences of the last order can be interpreted as an estimation of the error made in replacing the function by the interpolating rational function. Given the many similarities with Neville's algorithm many methods of that algorithm can be reused.

3.5.1 Bulirsch-Stoer interpolation — General implementation

The object implementing Bulirsch-Stoer interpolation's algorithm is best implemented as a subclass of the class `DhbNevilleInterpolator` since the methods used to manage the computation of the finite differences can be reused. The public interface is identical.

Only a single method — the one responsible for the evaluation of the finite differences at each order — must be implemented. All other methods of Neville's interpolation can be reused.

This shows the great power of object-oriented approach. Code written in procedural language cannot be reused that easily. In [Press et al.] the two codes implementing Neville's and Bulirsch-Stoer interpolation are of comparable length; not surprisingly they also have much in common.

3.5.2 Bulirsch-Stoer interpolation — Smalltalk implementation

Listing 3.9 shows the implementation of Bulirsch-Stoer interpolation in Smalltalk. The class `DhbBulirschStoerInterpolator` is a subclass of class `DhbNevilleInterpolator`. The

Figure 3.1 with the box `BulirschStoerInterpolator` grayed.


```

* Constructor method.
* @param pts DhbInterfaces.PointSeries
*/
public BulirschStoerInterpolator(DhbInterfaces.PointSeries pts) {
    super(pts);
}
/**
* @param m int
* @param n int
* @param x double
*/
protected void computeNextDifference( int m, int n, double x)
{
    double ratio = ( points.xValueAt(n) - x) * rightErrors[n]
                  / ( points.xValueAt(n + m + 1) - x);
    double diff = ( leftErrors[n+1] - rightErrors[n])
                  / ( ratio - leftErrors[n+1]);
    if( Double.isNaN( diff) )
    {
        diff = 0;
    }
    rightErrors[n] = leftErrors[n+1] * diff;
    leftErrors[n] = ratio * diff;
}
}

```

3.6 Cubic spline interpolation

The Lagrange interpolation polynomial is defined globally over the set of given points and respective function's values. As we have seen in figure 3.1 and to a lesser degree in figure 3.1 Lagrange's interpolation polynomial can have large fluctuations between two adjacent points because the degree of the interpolating polynomial is not constrained.

One practical method for interpolating a set of function's value with a polynomial of constrained degree is to use cubic splines. A cubic spline is a 3rd order polynomial constrained in its derivatives at the end points. A unique cubic spline is defined for each interval between two adjacent points. The interpolated function is required to be continuous up to the second derivative at each of the points.

Before the advent of computers, people were drawing smooth curves by sticking nails at the location of computed points and placing flat bands of metal between the nails. The bands were then used as rulers to draw the desired curve. These bands of metal were called splines and this is where the name of the interpolation algorithm comes from. The elasticity property of the splines correspond to the continuity property of the cubic spline function.

The algorithm exposed hereafter assumes that the sample points have been sorted in increasing order of abscissa.

To derive the expression for the cubic spline, one first assumes that the second derivatives of the splines, y_i'' , are known at each point. Then one writes the cubic spline between x_{i-1} and x_i in the following symmetric form:

Main equation \Rightarrow

$$P_i(x) = y_{i-1}A_i(x) + y_iB_i(x) + y_{i-1}''C_i(x) + y_i''D_i(x), \quad (3.13)$$

where

$$\begin{cases} A_i(x) &= \frac{x_i - x}{x_i - x_{i-1}}, \\ B_i(x) &= \frac{x - x_{i-1}}{x_i - x_{i-1}}. \end{cases} \quad (3.14)$$

Using the definition above, the first two terms in equation 3.13 represents the linear interpolation between the two points x_{i-1} and x_i . Thus, the last two terms of must vanish at x_{i-1} and x_i . In addition we must have by definition:

$$\begin{cases} \left. \frac{d^2 P_i(x)}{dx^2} \right|_{x=x_{i-1}} &= y_{i-1}'', \\ \left. \frac{d^2 P_i(x)}{dx^2} \right|_{x=x_i} &= y_i''. \end{cases} \quad (3.15)$$

One can rewrite the first equation in 3.15 as a differential equation for the function C_i as a function of A_i . Similarly, the second equation is rewritten as a differential equation for the function D_i as a function of B_i . This yields:

$$\begin{cases} C_i(x) &= \frac{A_i(x) [A_i(x)^2 - 1]}{6} (x_i - x_{i-1})^2, \\ D_i(x) &= \frac{B_i(x) [B_i(x)^2 - 1]}{6} (x_i - x_{i-1})^2, \end{cases} \quad (3.16)$$

Finally, one must use the fact that the first derivatives of each spline must be equal at each end points of the interval, that is:

$$\frac{dP_i(x)}{dx} = \frac{dP_{i+1}(x)}{dx}. \quad (3.17)$$

This yields the following equations for the second derivatives y_i'' :

$$\frac{x_{i+1} - x_i}{6} y_{i+1}'' + \frac{x_{i+1} - x_{i-1}}{6} y_i'' + \frac{x_i - x_{i-1}}{6} y_{i-1}'' = \frac{y_{i+1} - y_i}{x_{i+1} - x_i} - \frac{y_i - y_{i-1}}{x_i - x_{i-1}}. \quad (3.18)$$

There are $n - 1$ equations for the n unknowns y_i'' . We are thus missing two equations. There are two ways of defining two additional equations to obtain a unique solution.

- The first method is the so-called natural cubic spline for which one sets $y_0'' = y_n'' = 0$. This means that the spline is flat at the end points.
- The second method is called constrained cubic spline. In this case the first derivatives of the function at x_0 and x_n , y_0' and y_n' , are set to given values.

In the case of constrained cubic spline, one obtains two additional equations by evaluating the derivatives of equation 3.13 at x_0 and x_n :

$$\begin{cases} \frac{3A_1(x)^2-1}{6} (x_1 - x_0) y_0'' - \frac{3B_1(x)^2-1}{6} (x_1 - x_0) y_1'' &= y_0' - \frac{y_1 - y_0}{x_1 - x_0}, \\ \frac{3A_n(x)^2-1}{6} (x_n - x_{n-1}) y_n'' - \frac{3B_n(x)^2-1}{6} (x_n - x_{n-1}) y_{n-1}'' &= y_n' - \frac{y_n - y_{n-1}}{x_n - x_{n-1}}. \end{cases} \quad (3.19)$$

The choice between natural or constrained spline can be made independently at each end point.

One solves the system of equations 3.18, and possibly 3.19, using direct Gaussian elimination and back substitution (c.f. section 8.2). Because the corresponding matrix is tridiagonal, each pivoting step only involves one operation. Thus, resorting to a general algorithm for solving a system of linear equations is not necessary.

3.6.1 Cubic spline interpolation — General implementation

In both languages the object implementing cubic spline interpolation is a subclass of the Newton interpolator. The reader might be surprised by this choice since, mathematically speaking, these two objects do not have anything in common.

However, from the behavioral point of view, they are quite similar. Like for Newton interpolation, cubic spline interpolation first needs to compute a series of coefficients, namely the second derivatives, which only depends on the sample points. This calculation only needs to be performed once. Then the evaluation of the function can be done using equations 3.13, 3.14 and 3.16. Finally, as for the Newton interpolator, any modification of the points requires a new computation of the coefficients. The behavior can be reused from the class `NewtonInterpolator`.

The second derivatives needed by the algorithm are stored in the variable used to store the coefficients of Newton's algorithm.

The class `SplineInterpolator` has two additional instance variables needed to store the end point derivatives y_0' and y_n' . Corresponding methods needed to set or reset these values are implemented. If the value of y_0' or y_n' is changed then the coefficients must be reset.

Natural or constrained cubic spline is flagged independently at each point by testing if the corresponding end-point derivative has been supplied or not. The second derivatives are computed using lazy initialization by the method `computeSecondDerivatives`.

Figure 3.1 with the box `SplineInterpolator` grayed.

3.6.2 Cubic spline interpolation — Smalltalk implementation

Listing 3.11 shows the implementation of cubic spline interpolation in Smalltalk. The class `DhbSplineInterpolator` is a subclass of class `DhbNewtonInterpolator`. The code examples 3.1 and 3.2 can directly be applied to cubic spline interpolation after replacing the class name `DhbLagrangeInterpolator` with `DhbSplineInterpolator`.

If the end-point derivative is nil the corresponding end-point is treated as a natural spline.

The method `defaultSamplePoints` overrides that of the superclass to create a sorted collection. Thus, as each point is added to the implicit collection, the collection of

sample points remains in increasing order of abscissa as required by the cubic spline algorithm.

Listing 3.11 Smalltalk implementation of cubic spline interpolation

Class **DhbSplineInterpolator**
 Subclass of **DhbNewtonInterpolator**
 Instance variable names: startPointDerivative endPointDerivative

Instance methods

computeSecondDerivatives

```

| size u w s dx inv2dx |
size := pointCollection size.
coefficients := Array new: size.
u := Array new: size - 1.
startPointDerivative isNil
  ifTrue:
    [coefficients at: 1 put: 0.
     u at: 1 put: 0]
  ifFalse:
    [coefficients at: 1 put: -1 / 2.
     s := 1 / (( self xPointAt: 2) x - ( self xPointAt: 1) x).
     u at: 1
       put: 3 * s
           * (s * (( self yPointAt: size) - ( self
               yPointAt: size - 1))
             - startPointDerivative)].
2 to: size - 1
do:
  [:n |
    dx := (self xPointAt: n) - (self xPointAt: ( n - 1)).
    inv2dx := 1 / (( self xPointAt: n + 1) - (self xPointAt:
        n - 1)).
    s := dx * inv2dx.
    w := 1 / (s * (coefficients at: n - 1) + 2).
    coefficients at: n put: (s - 1) * w.
    u at: n
      put: ((( self yPointAt: n + 1) - ( self yPointAt:
          n))
        / (( self xPointAt: n + 1) - ( self xPointAt:
          n))
        - ((( self yPointAt: n) - ( self
          yPointAt: n - 1)) / dx)) * 6
        * inv2dx - ((u at: n - 1) * s))

```

```

        * w].
endPointDerivative isNil
  ifTrue: [coefficients at: size put: 0]
  ifFalse:
    [w := 1 / 2.
     s := 1 / ((self xPointAt: size) - (self xPointAt: (size
       - 1))).
     u at: 1
       put: 3 * s * (endPointDerivative
         - (s * (self yPointAt: size) - (self
           yPointAt: size - 1))).
     coefficients at: size
       put: s - (w * (u at: size - 1) / ((coefficients at:
         size - 1) * w + 1))].
size - 1 to: 1
  by: -1
  do:
    [:n |
     coefficients at: n
       put: (coefficients at: n) * (coefficients at: n + 1)
         + (u at: n)]

defaultSamplePoints
  ^SortedCollection sortBlock: [ :a :b | a x < b x]

endPointDerivative: aNumber
  endPointDerivative := aNumber.
  self resetCoefficients.

resetEndPointDerivatives
  self setEndPointDerivatives: (Array new: 2).

setEndPointDerivatives: anArray
  startPointDerivative := anArray at: 1.
  endPointDerivative := anArray at: 2.
  self resetCoefficients.

startPointDerivative: aNumber
  startPointDerivative := aNumber.
  self resetCoefficients.

```



```

value: aNumber

| answer n1 n2 n step a b |
coefficients isNil ifTrue: [self computeSecondDerivatives].
n2 := pointCollection size.
n1 := 1.
[n2 - n1 > 1] whileTrue:
    [n := (n1 + n2) // 2.
     (self xPointAt: n) > aNumber ifTrue: [n2 := n] ifFalse:
        [n1 := n]].
step := (self xPointAt: n2) - (self xPointAt: n1).
a := ((self xPointAt: n2) - aNumber) / step.
b := (aNumber - (self xPointAt: n1)) / step.
^a * (self yPointAt: n1) + (b * (self yPointAt: n2))
  + ((a * (a squared - 1) * (coefficients at: n1)
      + (b * (b squared - 1) * (coefficients at: n2))) *
      step squared
      / 6)

```

3.6.3 Cubic spline interpolation — Java implementation

Listing 3.12 shows the implementation of cubic spline interpolation in Java. The class `SplineInterpolator` is a subclass of class `NewtonInterpolator`. The code examples 3.3 and 3.4 can directly be applied to cubic spline interpolation after replacing the class name `LagrangeInterpolator` with `SplineInterpolator`.

The fact that an end-point derivative has not been supplied is flagged with the special Java value `Double.NaN`.

Note: Since there is no sorted collection in Java, this class supposes that the calling application supplies the point sorted by ascending value of the abscissa.

Listing 3.12 Java implementation of cubic spline interpolation

```

package DhbInterpolation;

import DhbInterfaces.OneVariableFunction;
import DhbInterfaces.PointSeries;
/**
 * A SplineInterpolator can be used to interpolate values between
 * a series of 2-dimensional points. The interpolation function is
 * a cubic spline with first derivatives defined at the end points.
 * If the first derivatives are not defined for the end points,
 * a so-called natural spline is used with second derivatives at

```

```

* the end points set to zero.
*
* @author Didier H. Besset
*/
public class SplineInterpolator extends NewtonInterpolator
    implements DhbInterfaces.OneVariableFunction
{
    /**
     * First derivative at first point.
     */
    private double startPointDerivative = Double.NaN;
    /**
     * First derivative at last point.
     */
    private double endPointDerivative = Double.NaN;
    /**
     * This method creates a new instance of a spline interpolator over
     * a given set of points. The points must be sorted by strictly
     * increasing x values.
     * @param pts DhbInterfaces.PointSeries contains the points sampling
     *           the function to interpolate.
     * @exception java.lang.IllegalArgumentException points are not sorted
     *           in increasing x values.
     */
    public SplineInterpolator ( PointSeries pts)
        throws IllegalArgumentException
    {
        super( pts);
        for ( int i = 1; i < pts.size(); i++ )
        {
            if ( pts.xValueAt( i - 1 ) >= pts.xValueAt( i ) )
                throw new IllegalArgumentException
                    ( "Points must be sorted in increasing x value");
        }
    }
    private void computeSecondDerivatives( )
    {
        int n = points.size();
        double w, s;
        double[] u = new double[ n - 1];
        coefficients = new double[ n];
        if ( Double.isNaN( startPointDerivative ) )
            coefficients[0] = u[0] = 0;
        else
        {
            coefficients[0] = -0.5;

```

```

        u[0] = 3.0 / ( points.xValueAt( 1) - points.xValueAt( 0))
                * ( ( points.yValueAt( 1) - points.yValueAt( 0))
                  / ( points.xValueAt( 1) - points.xValueAt( 0))
                  - startPointDerivative);
    }
    for ( int i = 1; i < n - 1; i ++ )
    {
        double invStep2 = 1 / ( points.xValueAt( i + 1)
                                - points.xValueAt( i - 1));
        s = ( points.xValueAt( i) - points.xValueAt( i - 1)) * invStep2;
        w = 1 / ( s * coefficients[ i - 1] + 2);
        coefficients[ i] = ( s - 1) * w;
        u[i] = ( 6 * invStep2 * (
            ( points.yValueAt( i + 1) - points.yValueAt( i))
            / ( points.xValueAt( i + 1) - points.xValueAt( i))
            - ( points.yValueAt( i) - points.yValueAt( i - 1))
            / ( points.xValueAt( i) - points.xValueAt( i - 1))
            ) - s * u[ i - 1]) * w;
    }
    if ( Double.isNaN( endPointDerivative) )
        w = s = 0;
    else
    {
        w = -0.5;
        s = 3.0 / ( points.xValueAt( n - 1) - points.xValueAt( n - 2))
            * ( endPointDerivative - ( points.yValueAt( n - 1)
                                      - points.yValueAt( n - 2))
              / ( points.xValueAt( n - 1) -
                  points.xValueAt( n - 2)));
    }
    coefficients[ n - 1] = ( s - w * u[ n - 2])
                          / ( w * coefficients[ n - 2] + 1);
    for ( int i = n - 2; i >= 0; i--)
        coefficients[i] = coefficients[i] * coefficients[i + 1] + u[i];
    return;
}
/**
 * Computes the interpolated y value for a given x value.
 * @param aNumber x value.
 * @return interpolated y value.
 */
public double value(double x)
{
    if ( coefficients == null )
        computeSecondDerivatives();
    int n1 = 0;

```

```

int n2 = points.size() - 1;
while ( n2 - n1 > 1 )
{
    int n = (n1 + n2) / 2;
    if ( points.xValueAt( n) > x )
        n2 = n;
    else
        n1 = n;
}
double step = points.xValueAt( n2) - points.xValueAt( n1);
double a = ( points.xValueAt( n2) - x) / step;
double b = ( x - points.xValueAt( n1)) / step;
return a * points.yValueAt( n1) + b * points.yValueAt( n2)
        + ( a * ( a * a - 1) * coefficients[n1]
            + b * ( b * b - 1) * coefficients[n2])
            * step * step / 6;
}
}

```

3.7 Which method to choose?

At this point some reader might experience some difficulty in choosing among the many interpolation algorithms discussed in this book. There are indeed many ways to skin a cat. Selecting a method depends on what the user intends to do with the data.

First of all, the reader should be reminded that Lagrange interpolation, Newton interpolation and Neville's algorithm are different algorithms computing the values of the same function, namely the Lagrange interpolation polynomial. In other words, the interpolated value resulting from each 3 algorithms is the same (up to rounding errors of course).

The Lagrange interpolation polynomial can be subject to strong variations (if not wild in some cases, figure 3.1 for example) if the sampling points are not smooth enough. A cubic spline may depart from the desired function if the derivatives on the end points are not constrained to proper values. A rational function can do a good job in cases where polynomials have problems. To conclude, let me give you some rules of thumb to select the best interpolation method based on my personal experience.

If the function to interpolate is not smooth enough, which maybe the case when not enough sampling points are available, a cubic spline is preferable to the Lagrange interpolation polynomial. Cubic splines are traditionally used in curve drawing programs. Once the second derivatives have been computed, evaluation time is of the order of $O(n)$. You must keep in your mind the limitation⁶ imposed on the curvature when using a 3rd order polynomial.

⁶The curvature of a cubic spline is somewhat limited. What happens is that the curvature and the slope (first derivative) are strongly coupled. As a consequence a cubic spline gives a smooth approximation to the interpolated points.

If the Lagrange interpolation polynomial is used to quickly evaluate a tabulated⁷ function, Newton interpolation is the algorithm of choice. Like for cubic spline interpolation, the evaluation time is of the order of $O(n)$ once the coefficients have been computed.

Neville's algorithm is the only choice if an estimate of error is needed in addition to the interpolated value. The evaluation time of the algorithm is of the order of $O(n^2)$.

Lagrange interpolation can be used for occasional interpolation or when the values over which interpolation is made are changing at each interpolation. The evaluation time of the algorithm is of the order of $O(n^2)$. Lagrange interpolation is slightly slower than Neville's algorithm as soon as the number of points is larger than 3⁸. However, Neville's algorithm needs to allocate more memory. Depending on the operating system and the amount of available memory the exact place where Lagrange interpolation becomes slower than Neville's algorithm is likely to change.

If the function is smooth but a Lagrange polynomial is not reproducing the function in a proper way, a rational function can be tried using Bulirsch-Stoer interpolation.

Table 3.1 shows a summary of the discussion. If you are in doubt, I recommend

Table 3.1: Recommended polynomial interpolation algorithms

Feature	Recommended algorithm
Error estimate desired	Neville
Couple of sample points	Lagrange
Medium to large number of sample points	Neville
Many evaluations on fixed sample	Newton
Keep curvature under constraint	Cubic spline
Function hard to reproduce	Bulirsch-Stoer

that you make a test first for accuracy and then for speed of execution. Drawing a graph such as in the figures presented in this chapter is quite helpful to get a proper feeling about the possibility offered by various interpolation algorithms on a given set of sample points. If neither Lagrange interpolation nor Bulirsch-Stoer nor cubic spline is doing a good job at interpolating the sample points, you should consider using curve fitting (c.f. chapter 10) with an ad-hoc function.

⁷A tabulated function is a function, which has been computed at a finite number of its argument.

⁸Such a number is strongly dependent on the operating system and virtual machine. Thus, the reader should check this number him/herself.

Chapter 4

Iterative algorithms

Cent fois sur le métier remettez votre ouvrage.
Nicolas Boileau

When a mathematical function cannot be approximated with a clever expression, such as Lanczos formula introduced in the chapter 2.4.1, one must resort to compute that function using the integral, the recurrence formula or the series expansion. All these algorithms have one central feature in common: the repetition of the same computation until some convergence criteria is met. Such repetitive computation is called iteration.

Figure 4.1 shows the class diagram of the classes discussed in this chapter. This chapter first discusses the implementation of a general-purpose iterative process. Then, we describe a generalization for the finding of a numerical result. Other chapters discuss examples of sub-classing of these classes to implement specific algorithms.

Iteration is used to find the solution of a wide variety of problems other than just function evaluation. Finding the location where a function is zero, reached a maximum or a minimum is another example. Some data mining algorithms also use iteration to find a solution (c.f. section 12.6).

4.1 Successive approximations

A general-purpose iterative process can be decomposed in three main steps:

- a set-up phase
- an iteration phase until the result is acceptable
- a clean-up phase

These steps are translated schematically into the flow diagram shown in Figure 4.2.

The set-up phase allows determining constant parameters used by the subsequent computations. Often a first estimation of the solution is defined at this time. In any case an object representing the approximate solution is constructed. Depending on the

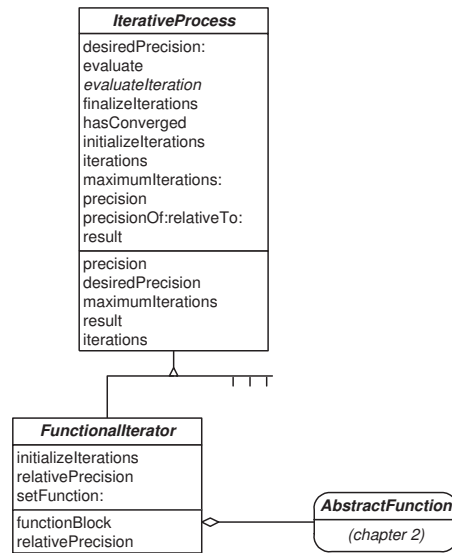


Figure 4.1: Class diagram for iterative process classes




Figure 4.2: Successive approximation algorithm

complexity of the problem a class will explicitly represent the solution object. Otherwise the solution shall be described by a few instance variables of simple types (numbers and arrays).

After the set-up phase the iterative process proper is started. A transformation is applied to the solution object to obtain a new object. This process is repeated unless the solution object resulting from the last transformation can be considered close enough to the sought solution.

During the clean-up phase resources used by the iterative process must be release. In some cases additional results may be derived before leaving the algorithm. Let us now explicit each of the three stages of the algorithm.

The step computing or choosing an initial object is strongly dependent on the nature of the problem to be solved. In some methods, a good estimate of the solution can be computed from the data. In others using randomly generated objects yields good results. Finally, one can also ask the application's user for directions. In many cases this step is also used to initialize parameters needed by the algorithm.

The step computing the next object contains the essence of the algorithm. In general a new object is generated based on the history of the algorithm.

The step deciding whether or not an object is sufficiently close to the sought solution is more general. If the algorithm is capable of estimating the precision of the solution — that is, how close the current object is located from the exact solution — one can decide to stop the algorithm by comparing the precision to a desired value. This is not always the case, however. Some algorithms, genetic algorithms for example, do not have a criterion for stopping.

Whether or not a well-defined stopping criterion exists, the algorithm must be prevented from taking an arbitrary large amount of time. Thus, the object implementing an iterative process ought to keep track of the number of iterations and interrupt the algorithm if the number of iterations becomes larger than a given number.

Design

Now we can add some details to the algorithm. The new details are shown in figure 4.1 with the box IterativeProcess grayed. Figure 4.3. This schema allows us to determine the structure of a general object implementing the iterative process. It will be implemented as an abstract class. An abstract class is a class with does not have object instances. A object implementing a specific algorithm is an instance of a particular subclass of the abstract class.

The gray boxes in figure 4.3 represent the methods, which must be implemented explicitly by the subclass. The abstract class calls them. However, the exact implementation of these methods is not defined at this stage. Such methods are called hook methods.

Using this architecture the abstract class is able to implement the iterative process without any deep knowledge of the algorithm. Algorithm specific methods are implemented by the subclass of the abstract class.

Let us call IterativeProcess the class of the abstract object. The class IterativeProcess needs the following instance variables.

iterations keeps track of the number of iterations, that is the number of successive approximations,

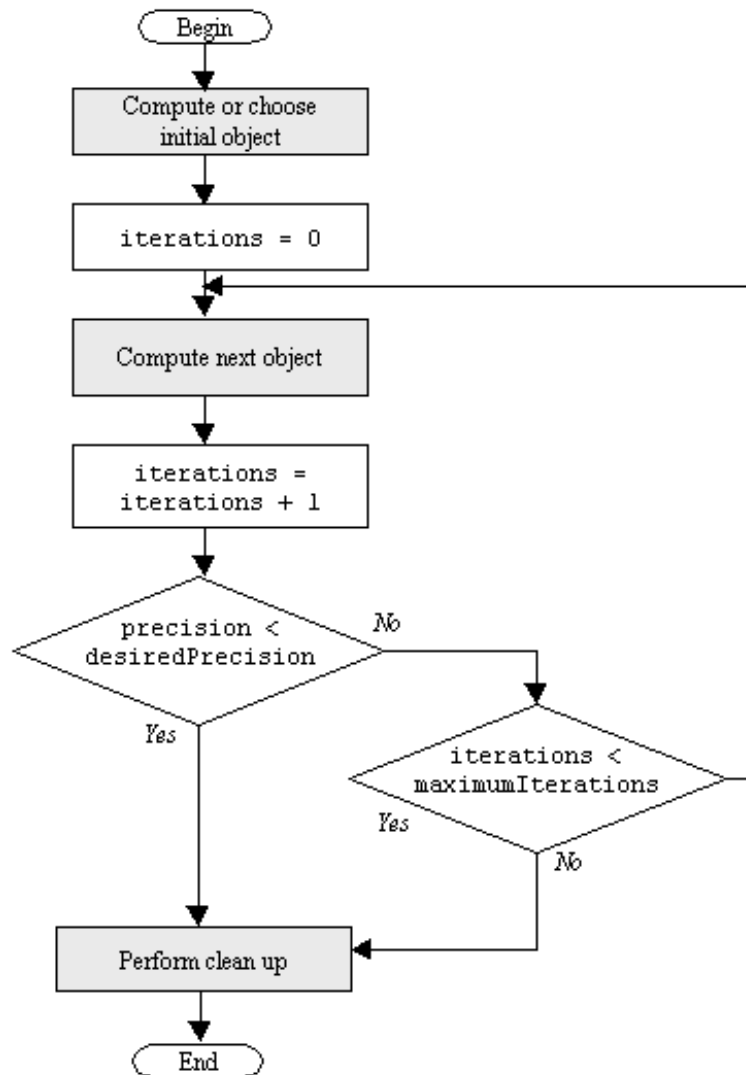


Figure 4.3: Detailed algorithm for successive approximations

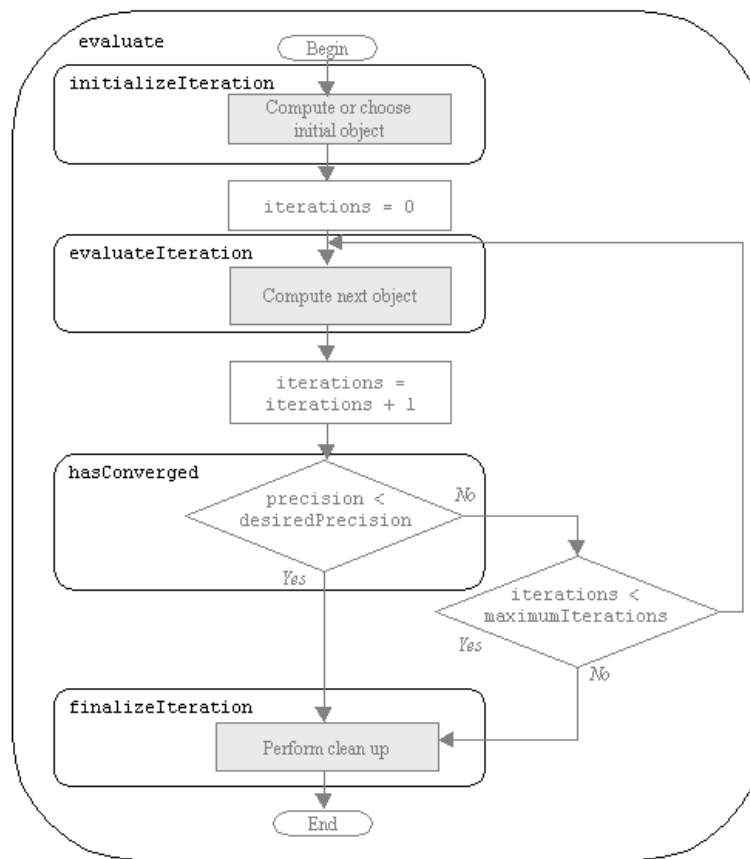


Figure 4.4: Methods for successive approximations

`maximumIterations` maximum number of allowed iterations,

`desiredPrecision` the precision to attain, that is, how close to the solution should the solution object be when the algorithm is terminated,

`precision` the precision achieved by the process. Its value is updated after each iteration and it is used to decide when to stop.

The methods of the class `IterativeProcess` are shown in figure 4.4 in correspondence with the general execution flow shown in figure 4.3. The two methods `initializeIterations` and `finalizeIterations` should be implemented by the subclass but the abstract class provides a default behavior: doing nothing. The method `evaluateIteration` must be implemented by each subclass.

Since the precision of the last iteration is kept in an instance variable, the method `hasConverged` can be called at any time after evaluation, thus providing a way for client classes to check whether the evaluation has converged or not.

4.1.1 Iterative process — Smalltalk implementation

Even though we are dealing for the moment with an abstract class we are able to present a scenario of use illustrating the public interface to the class. Here is how a basic utilization of an iterative process object would look like.

Code example 4.1

```

| iterativeProcess result |
iterativeProcess := <a subclass of DhbIterativeProcess> new.
result := iterativeProcess evaluate.
iterativeProcess hasConverged
ifFalse:[ <special case processing>].

```

The first statement creates an object to handle the iterative process. The second one performs the process and retrieves the result, whatever it is. The final statement checks for convergence.

To give the user a possibility to have more control, one can extend the public interface of the object to allow defining the parameters of the iterative process: the desired precision and the maximum number of iterations. In addition, the user may want to know the precision of the attained result and the number of iterations needed to obtain the result. The following code sample shows an example of use for all public methods defined for an iterative process. The precision of the attained result and the number of iterations are printed on the transcript window.

Code example 4.2

```

| iterativeProcess result precision |
iterativeProcess := <a subclass of DhbIterativeProcess> new.
iterativeProcess desiredPrecision: 1.0e-6; maximumIterations: 25.
result := iterativeProcess evaluate.
iterativeProcess hasConverged
ifTrue: [ Transcript nextPutAll: 'Result obtained after '.
iterativeProcess iteration printOn: Transcript.
Transcript nextPutAll: 'iterations. Attained precision is '.
iterativeProcess precision printOn: Transcript.
]
ifFalse:[ Transcript nextPutAll: 'Process did not converge' ].
Transcript cr.

```

Listing 4.1 shows the Smalltalk implementation of the iterative process.

In the Smalltalk implementation, the class `IterativeProcess` has one instance variable in addition to the ones described in the preceding section. This variable, called `result`, is used to keep the solution object of the process. The method `result` allows direct access to it. Thus, all subclasses can use this instance variable as a placeholder to store any type of result. As a convenience the method `evaluate` also returns the instance variable `result`.

Default values for the desired precision and the maximum number of iterations are kept in class methods for easy editing. The method `initialize` loads these default values

for each newly created instance. The default precision is set to the machine precision discussed in section 1.3.2.

The methods used to modify the desired precision (`desiredPrecision:`) and the maximum number of iterations (`maximumIterations:`) check the value to prevent illegal definitions, which could prevent the algorithm from terminating.

Since there is no explicit declaration of abstract class and abstract methods in Smalltalk¹ the three methods `initializeIterations`, `evaluateIteration` and `finalizeIterations`, are implemented with a reasonable default behavior. The methods `initializeIterations` and `finalizeIterations` do nothing. The method `evaluateIteration` calls the method `subclassResponsibility`, which raises an exception when called. Using this technique is the Smalltalk way of creating an abstract method.

Listing 4.1 Smalltalk implementation of an iterative process

```

Class                DhbIterativeProcess
Subclass of         Object
Instance variable names:  precision desiredPrecision maximumIterations result iterations

Class methods
defaultMaximumIterations
    ^50

defaultPrecision
    ^DhbFloatingPointMachine new defaultNumericalPrecision

new
    ^super new initialize

Instance methods
desiredPrecision: aNumber
    aNumber > 0
    ifFalse: [ ^self error: 'Illegal precision: ', aNumber
                printString].
    desiredPrecision := aNumber.

evaluate

```

¹An abstract class is a class containing at least an abstract method; an abstract method contains the single conventional statement:
`self subclassResponsibility`

```

iterations := 0.
self initializeIterations.
[iterations := iterations + 1.
precision := self evaluateIteration.
self hasConverged or: [iterations >= maximumIterations]]
    whileFalse: [].
self finalizeIterations.
^self result

evaluateIteration

^self subclassResponsibility

finalizeIterations

hasConverged

^precision <= desiredPrecision

initialize

desiredPrecision := self class defaultPrecision.
maximumIterations := self class defaultMaximumIterations.
^self

initializeIterations

iterations

^iterations

maximumIterations: anInteger

( anInteger isInteger and: [ anInteger > 1])
    ifFalse: [ ^self error: 'Invalid maximum number of iteration:
        ', anInteger printString].
maximumIterations := anInteger.

precision

^precision

```

precisionOf: aNumber1 relativeTo: aNumber2

```
^aNumber2 > DhhFloatingPointMachine new defaultNumericalPrecision  
  ifTrue: [ aNumber1 / aNumber2]  
  ifFalse:[ aNumber1]
```

result

```
^result
```

Note: The method precisionOf:relativeTo: implements the computation of the relative precision. This is discussed in section 4.2.1.

4.1.2 Iterative process: JAVA implementation

Even though we are dealing for the moment with an abstract class we are able to present a scenario of use illustrating the public interface to the class. Here is how a basic utilization of an iterative process object would look like.

Code example 4.3 <a subclass of IterativeProcess> iterativeProcess = new
 <a subclass of IterativeProcess>();
 iterativeProcess.evaluate();
 if (iterativeProcess.hasConverged) {
 <retrieve the result>;
 } else {
 < special case processing >;
 }

The first statement creates an object to handle the iterative process. The second one performs the process. The final statement checks for convergence.

To give the user a possibility to have more control, one can extend the public interface of the object to allow defining the parameters of the iterative process: the desired precision and the maximum number of iterations. In addition, the user may want to know the precision of the attained result and the number of iterations needed to obtain the result. The following code sample shows an example of use for all public methods defined for an iterative process. The precision of the attained result and the number of iterations are printed on console output

```

Code example 4.4      | iterativeProcess result |
    <a subclass of IterativeProcess> iterativeProcess = new <a subclass of
    IterativeProcess>();
    iterativeProcess. setMaximumIterations (25);
    iterativeProcess. setDesiredPrecision (1.0e-6);
    iterativeProcess.evaluate();
    if ( iterativeProcess.hasConverged)
    System.out.println( "Result obtained after "
    +iterativeProcess.getIterations()+" iterations. Attained precision is "+iterativeProcess.getPrecision());
    else
    System.out.println( "Process did not converge");

```

Listing 4.2 shows the Java implementation of the iterative process.

The instance variables are kept private since no subclass should be able to modify them. Otherwise the efficacy of the algorithm could be disturbed.

In the case of a general-purpose algorithm the nature of the result is not known. It would be possible to declare a variable of type Object to hold a solution object of any type. However, this would force each subclass and each client of the subclass to use an explicit cast to obtain the result into the proper class. In general cast ought to be avoided because they beat the purpose of using strong typing, that is checking code consistency at compile time. Thus, unlike the Smalltalk implementation, no instance variable can be used to hold the result.

For the same reason the method evaluate cannot return the result. It is thus declared as void. Only the subclass knows the exact type of the result. A method getResult must be implemented by each subclass to allow the client class to retrieve the result.

The method evaluateIteration is declared as an abstract method to make sure that any subclass implements it. The methods initializeIterations and finalizeIterations are supplied with a default behavior, which is to do nothing.

Default values for the desired precision and the maximum number of iterations are assigned statically. Any subclass needing to change these values must use the corresponding setting methods in its constructor method.

Otherwise the Java implementation is identical to the Smalltalk one.

Listing 4.2 Java implementation of an iterative process

```

package DhbIterations;

import DhbFunctionEvaluation.DhbMath;
/**
 * An iterative process is a general structure managing iterations.
 *
 * @author Didier H. Besset
 */
public abstract class IterativeProcess
{

```



```

/**
 * Number of iterations performed.
 */
private int iterations;
/**
 * Maximum allowed number of iterations.
 */
private int maximumIterations = 50;
/**
 * Desired precision.
 */
private double desiredPrecision = DhbMath.defaultNumericalPrecision();
/**
 * Achieved precision.
 */
private double precision;

/**
 * Generic constructor.
 */
public IterativeProcess() {
}
/**
 * Performs the iterative process.
 * Note: this method does not return anything because Java does not
 * allow mixing double, int, or objects
 */
public void evaluate()
{
    iterations = 0;
    initializeIterations();
    while ( iterations++ < maximumIterations )
    {
        precision = evaluateIteration();
        if ( hasConverged() )
            break;
    }
    finalizeIterations();
}
/**
 * Evaluate the result of the current interation.
 * @return the estimated precision of the result.
 */
abstract public double evaluateIteration();
/**

```

```
* Perform eventual clean-up operations
* (mustbe implement by subclass when needed).
*/
public void finalizeIterations ( )
{
}
/**
* Returns the desired precision.
*/
public double getDesiredPrecision( )
{
    return desiredPrecision;
}
/**
* Returns the number of iterations performed.
*/
public int getIterations()
{
    return iterations;
}
/**
* Returns the maximum allowed number of iterations.
*/
public int getMaximumIterations( )
{
    return maximumIterations;
}
/**
* Returns the attained precision.
*/
public double getPrecision()
{
    return precision;
}
/**
* Check to see if the result has been attained.
* @return boolean
*/
public boolean hasConverged()
{
    return precision < desiredPrecision;
}
/**
* Initializes internal parameters to start the iterative process.
*/
public void initializeIterations()
```

```

{
}
/**
 * @return double
 * @param epsilon double
 * @param x double
 */
public double relativePrecision( double epsilon, double x)
{
    return x > DhbMath.defaultNumericalPrecision()
        ? epsilon / x: epsilon;
}
/**
 * Defines the desired precision.
 */
public void setDesiredPrecision( double prec )
    throws IllegalArgumentException
{
    if ( prec <= 0 )
        throw new IllegalArgumentException
            ( "Non-positive precision: "+prec);
    desiredPrecision = prec;
}
/**
 * Defines the maximum allowed number of iterations.
 */
public void setMaximumIterations( int maxIter)
    throws IllegalArgumentException
{
    if ( maxIter < 1 )
        throw new IllegalArgumentException
            ( "Non-positive maximum iteration: "+maxIter);
    maximumIterations = maxIter;
}
}

```

Note: The method `relativePrecision` implements the computation of the relative precision. This is discussed in section 4.2.2.

4.2 Evaluation with relative precision

Figure 4.1 with the box `FunctionalIterator` grayed.

So far we have made no assumption about the nature of the solution searched by an iterative process. In this section we want to discuss the case when the solution is a numerical value.

As discussed in section 1.3.2 a floating-point number is a representation with constant relative precision. It is thus meaningless to use absolute precision to determine the convergence of an algorithm. The precision of an algorithm resulting in a numerical value ought to be determined relatively.

One way to do it is to have the method `evaluateIteration` returning a relative precision instead of an absolute number. Relative precision, however, can only be evaluated if the final result is different from zero. If the result is zero, the only possibility is to check for absolute precision. Of course, in practice one does not check for equality with zero. The computation of a relative precision is carried only if the absolute value of the result is larger than the desired precision.

The reasoning behind the computation of the relative error is quite general. Thus, a general-purpose class `FunctionalIterator` has been created to implement a method computing the relative precision from an absolute precision and a numerical result. In addition, since all subclasses of `FunctionalIterator` use a function a general method to handle the definition of that function is also supplied.

4.2.1 Relative precision — Smalltalk implementation

In this case the public interface is extended with a creation method taking as argument the function on which the process operates. The code example of section 4.1.1 then becomes:

```

Code example 4.5      | iterativeProcess result |
iterativeProcess := <a subclass of DhbFunctionalIterator> function: ( DhbPolynomial
coefficients: #(1 2 3).
result := iterativeProcess evaluate.
iterativeProcess hasConverged
ifFalse:[ <special case processing>].

```

In this example the function on which the process will operate is the polynomial $3x^2 + 2x + 1$ (c.f. section 2.2).

Listing 4.3 shows the implementation of the abstract class `DhbFunctionalIterator` in Smalltalk.

This class has one instance variable `functionBlock` to store the function. A single class method allows creating a new instance while defining the function.

As we have seen in section 2.1.1, a function can be any object responding to the message `value:`. This allows supplying any block of Smalltalk code as argument to the constructor method. However, the user can also supply a class implementing the computation of the function with a method with selector `value:`. For example, an instance of the class `DhbPolynomial` discussed in section 2.2.3 can be used.

The instance method `setFunction:` is used to set the instance variable `functionBlock`. In order to prevent a client class from sending the wrong object, the method first checks whether the supplied object responds to the message `value:`. This is one way of ensuring that the arguments passed to a method conform to the expected protocol. This way of doing is only shown as an example, however. It is not recommended in practice. The responsibility of supplying the correct arguments to a Smalltalk method is usually the responsibility of the client class.

The method `initializeIterations` first checks whether a function block has been defined. Then it calls the method `computeInitialValues`. This method is a hook method, which a subclass must implement to compute the value of the result at the beginning of the iterative process.

The computation of relative precision is implemented at two levels. One general method, `precisionOf:relativeTo:`, implemented by the superclass allows the computation of the relative precision relative to any value. Any iterative process can use this method. The method `relativePrecision` implements the computation of the precision relative to the current result.

Listing 4.3 Smalltalk implementation of the class `DhbFunctionalIterator`

```

Class      DhbFunctionalIterator
Subclass of DhbIterativeProcess
Instance variable names:  functionBlock relativePrecision

```

Class methods

```
function: aBlock
```

```
^self new setFunction: aBlock; yourself
```

Instance methods

initializeIterations

```
functionBlock isNil ifTrue: [self error: 'No function supplied'].  
self computeInitialValues
```

relativePrecision: aNumber

```
^self precisionOf: aNumber relativeTo: result abs
```

setFunction: aBlock

```
( aBlock respondsTo: #value:)  
  ifFalse:[ self error: 'Function block must implement the  
                    method value:'].  
functionBlock := aBlock.
```

4.2.2 Relative precision — Java implementation

In this case the public interface is extended with a creation method taking as argument the function on which the process operates. The code example of section 4.1.2 then becomes:

```

Code example 4.6      double[]coefficients = 1, 2, 3;
                     <a subclass of FunctionalIterator> iterativeProcess = new <a subclass of
                     FunctionalIterator>(new Polynomial(coefficients));
                     iterativeProcess.evaluate();
                     if ( iterativeProcess.hasConverged) {
                     double result = iterativeProcess.getResult();
                     } else {
                     < special case processing >;
                     }

```

In this example, the function on which the process will operate is the polynomial $3x^2 + 2x + 1$ (c.f. section 2.2). In the case of a numerical result the type of the result is known; thus, we can retrieve the result explicitly.

Listing 4.4 shows the implementation of the abstract class `FunctionalIterator` in Java.

In Java the abstract class has two instances variables: one to hold the function as in Smalltalk and one to hold the result now that the type of the result is defined (c.f. the remark in section 4.1.2). An access method for the result is implemented.

The computation of relative precision is implemented at two levels. One general method, `relativePrecision` with two arguments, implemented by the superclass allows the computation of the relative precision relative to any value. Any iterative process can use this method. The method `relativePrecision` with one argument implements the computation of the precision relative to the current result.

Otherwise the implementation is similar to that of Smalltalk.

Listing 4.4 Java implementation of the class `FunctionalIterator`

```

package DhbIterations;

import DhbInterfaces.OneVariableFunction;
/**
 * Iterative process based on a one-variable function,
 * having a single numerical result.
 *
 * @author Didier H. Besset
 */
public abstract class FunctionalIterator extends IterativeProcess
{
    /**
     * Best approximation of the zero.
     */
    protected double result = Double.NaN;
    /**
     * Function for which the zero will be found.
     */
    protected OneVariableFunction f;
}

```

```

* Generic constructor.
* @param func OneVariableFunction
* @param start double
*/
public FunctionalIterator(OneVariableFunction func)
{
    setFunction( func);
}
/**
* Returns the result (assuming convergence has been attained).
*/
public double getResult( )
{
    return result;
}
/**
* @return double
* @param epsilon double
*/
public double relativePrecision( double epsilon)
{
    return relativePrecision( epsilon, Math.abs( result));
}
/**
* @param func DhhInterfaces.OneVariableFunction
*/
public void setFunction( OneVariableFunction func)
{
    f = func;
}
/**
* @param x double
*/
public void setInitialValue( double x)
{
    result = x;
}
}

```

4.3 Examples

As we have dealt with abstract classes, this chapter did not give concrete examples of use. By consulting the rest of this book the reader will find numerous examples of subclasses of the two classes described in this chapter. Table 4.1 lists the sections where each algorithm using the iterative process framework is discussed.

Table 4.1: Algorithms using iterative processes

Algorithm or class of algorithm	Superclass	Chapter or section
Zero finding	Function iterator	Chapter 5
Integration	Function iterator	Chapter 6
Infinite series and continued fractions	Function iterator	Chapter 7
Matrix eigenvalues	Iterative process	Section 8.6
Non-linear least square fit	Iterative process	Section 10.9
Maximum likelihood fit	Iterative process	Section 10.10
Function minimization	Function iterator	Chapter 11
Cluster analysis	Iterative process	Section 12.6

Chapter 5

Finding the zero of a function

Le zéro, collier du néant.¹
Jean Cocteau

The zeroes of a function are the values of the function's variable for which the value of the function is zero. Mathematically, given the function $f(x)$, z is a zero of when $f(z) = 0$. This kind of problem is can be extended to the general problem of computing the value of the inverse function, that is finding a value x such that $f(x) = c$ where c is a given number. The inverse function is noted as $f^{-1}(x)$. Thus, one wants to find the value of $f^{-1}(c)$ for any c . The problem can be transformed into the problem of finding the zero of the function $\tilde{f}(x) = f(x) - c$.

The problem of finding the values at which a function takes a maximum or minimum value is called searching for the extremes of a function. This problem can be transformed into a zero-finding problem if the derivative of the function can be easily computed. The extremes are the zeroes of the function's derivative.

Figure 5.1 shows the class diagram of the classes discussed in this chapter.

5.1 Introduction

Let us begin with a concrete example.

Often an experimental result is obtained by measuring the same quantity several times. In scientific publications, such a result is published with two numbers: the average and the standard deviation of the measurements. This is true for medical publication as well. As we have already discussed in section 2.3.1, obstetricians prefer to think in terms of risk and prefer to use centiles instead of average and standard deviation. Assuming that the measurements were distributed according to a normal distribution (c.f. section 9.6), the 90th centile is the solution to the following equation:

$$\operatorname{erf}(x) = 0.9 \tag{5.1}$$

¹The zero, a necklace for emptiness.

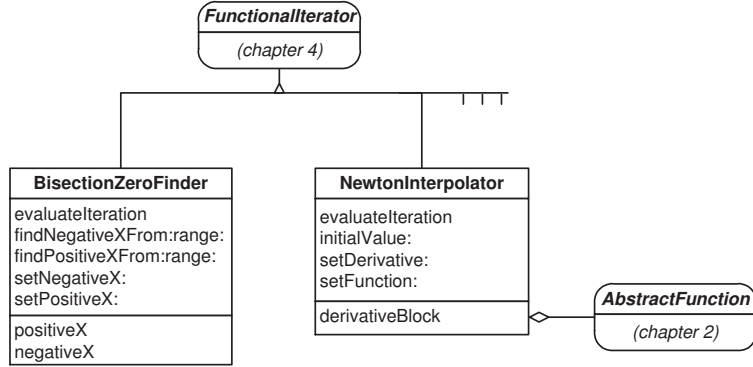


Figure 5.1: Class diagram for zero finding classes

That is, we need to find the zero of the function $f(x) = \text{erf}(x) - 0.9$. The answer is $x = 1.28$ with a precision of two decimals. Thus, if μ and σ are respectively the average and standard deviation of a published measurement, the 90th centile is given by $\mu + 1.28 \cdot \sigma$. Using equation 2.19 the 10th centile is given by $\mu - 1.28 \cdot \sigma$.

5.2 Finding the zeroes of a function — Bisection method

Figure 5.1 with the box BisectionZeroFinder grayed.

Let assume that one knows two values of x for which the function takes values of opposite sign. Let us call x_{pos} the value such that $f(x_{\text{pos}}) > 0$ and x_{neg} the value such that $f(x_{\text{neg}}) < 0$. If the function is continuous between x_{pos} and x_{neg} , there exists at least one zero of the function in the interval $[x_{\text{pos}}, x_{\text{neg}}]$. This is illustrated in figure 5.2. If the function f is not continuous over the interval where the sign of the function changes, then the presence of a zero cannot be guaranteed². The continuity requirement is essential for the application of the bisection algorithm.

The values x_{pos} and x_{neg} are the initial values of the bisection algorithm. The algorithm goes as follows:

1. Compute $x = \frac{x_{\text{pos}} + x_{\text{neg}}}{2}$.
2. If $f(x) > 0$, set $x_{\text{pos}} = x$ and goto step 4.
3. Otherwise set $x_{\text{neg}} = x$.
4. If $|x_{\text{pos}} - x_{\text{neg}}| > \epsilon$ go back to step 1. ϵ is the desired precision of the solution.

The first couple of steps of the bisection algorithm are represented geometrically on figure 5.2. Given the two initial values, x_{pos} and x_{neg} , the first iteration of the algorithm replaces x_{pos} with x_1 . The next step replaces x_{neg} with x_2 .

²The inverse function is such an example. It changes sign over 0 but has no zeroes for any finite x




Figure 5.2: The bisection algorithm

For a given pair of initial values, x_{pos} and x_{neg} , the number of iterations required to attain a precision ϵ is given by:

$$n = \left\lceil \log_2 \frac{|x_{\text{pos}}, x_{\text{neg}}|}{\epsilon} \right\rceil. \quad (5.2)$$

For example if the distance between the two initial values is 1 the number of iterations required to attain a precision of 10^{-8} is 30. It shows that the bisection algorithm is rather slow.

Knowledge of the initial values, x_{pos} and x_{neg} , is essential for starting the algorithm. Methods to define them must be supplied. Two convenience methods are supplied to sample the function randomly over a given range to find each initial value. The random number generator is discussed in section 9.4.

The bisection algorithm is a concrete implementation of an iterative process. In this case, the method `evaluateIteration` of figure 4.4 implements steps 2, 3 and 4. The precision at each iteration is $|x_{\text{pos}} - x_{\text{neg}}|$ since the zero of the function is always inside the interval defined by x_{pos} and x_{neg} .

5.2.1 Bisection algorithm — General implementation

The class of the object implementing the bisection algorithm is a subclass of the abstract class `FunctionalIterator`. The class `BisectionZeroFinder` needs the following additional instance variables.

`positiveX` x_{pos} and

`negativeX` x_{neg}

The bisection algorithm proper is implemented only within the method `evaluateIteration`. Other necessary methods have already been implemented in the iterative process class.

5.2.2 Bisection algorithm — Smalltalk implementation

Finding the zero of a function is performed by creating an instance of the class `DhbBisectionZeroFinder` and giving the function as the argument of the creation method as explained in section 4.2.1. For example the following code finds the solution of equation 5.1.

Code example 5.1

```
| zeroFinder result |
zeroFinder:= DhbBisectionZeroFinder function: [ :x | x errorFunction - 0.9].
zeroFinder setPositiveX: 10; setNegativeX: 0.
result := zeroFinder evaluate. zeroFinder
hasConverged

ifFalse:[ <special case processing>].
```

The second line creates the object responsible to find the zero. The third line defines the initial values, x_{pos} and x_{neg} . The fourth line performs the algorithm and stores the result if the algorithm has converged. The last two lines check for convergence and take corrective action if the algorithm did not converge.

Listing 5.1 shows the implementation of the bisection zero finding algorithm in Smalltalk.

The class `DhbBisectionZeroFinder` is a subclass of the class `DhbFunctionalIterator`. As one can see only a few methods need to be implemented. Most of them pertain to the definition of the initial interval. In particular, convenience methods are supplied to find a positive and negative function value over a given interval.

The methods defining the initial values, x_{pos} and x_{neg} , are `setPositiveX:` and `setNegativeX:` respectively. An error is generated in each method if the function's value does not have the proper sign. The convenience methods to find random starting values are respectively `findPositiveXFrom:range:` and `findNegativeXFrom:range:`. The method `computeInitialValues` does not compute the initial values. Instead it makes sure that x_{pos} and x_{neg} have been properly defined.

Listing 5.1 Smalltalk implementation of the bisection algorithm

Class	DhbBisectionZeroFinder
Subclass of	<code>DhbFunctionalIterator</code>
Instance variable names:	<code>positiveX</code> <code>negativeX</code>

Instance methods

`computeInitialValues`

`positiveX isNil`

`ifTrue: [self error: 'No positive value supplied'].`

`negativeX isNil`

`ifTrue: [self error: 'No negative value supplied'].`

evaluateIteration

```

result := ( positiveX + negativeX ) * 0.5.
( functionBlock value: result ) > 0
  ifTrue: [ positiveX := result ]
  ifFalse: [ negativeX := result ].
^self relativePrecision: ( positiveX - negativeX ) abs

```

findNegativeXFrom: aNumber1 range: aNumber2

```

| n |
n := 0.
[ negativeX := Number random * aNumber2 + aNumber1.
  ( functionBlock value: negativeX ) < 0
  ] whileFalse: [ n := n + 0.1.
    n > maximumIterations
    ifTrue: [ self error: 'Unable to find a
      negative function value' ].
  ].

```

findPositiveXFrom: aNumber1 range: aNumber2

```

| n |
n := 0.
[ positiveX := Number random * aNumber2 + aNumber1.
  ( functionBlock value: positiveX ) > 0
  ] whileFalse: [ n := n + 1.
    n > maximumIterations
    ifTrue: [ self error: 'Unable to find a
      positive function value' ].
  ].

```

setNegativeX: aNumber

```

( functionBlock value: aNumber ) < 0
  ifFalse: [ self error: 'Function is not negative at x = ',
    aNumber printString ].
negativeX := aNumber.

```

setPositiveX: aNumber

```

( functionBlock value: aNumber ) > 0
  ifFalse: [ self error: 'Function is not positive at x = ',
    aNumber printString ].
positiveX := aNumber.

```

5.2.3 Bisection algorithm — Java implementation

Finding the zero of a function is performed by creating an instance of the class `BisectionZeroFinder` and giving the function as the argument of the constructor method as explained in section 4.2.2. For example the following code finds the solution of equation 5.1.

Code example 5.2

```
BisectionZeroFinder zeroFinder = new BisectionZeroFinder( new
    OneVariableFunction() { public double value( double x)
        { return NormalDistribution.errorFunction(x)-0.9;}});
try { zeroFinder.setNegativeX( 0);
    zeroFinder.setPositiveX ( 5);
    } catch( IllegalArgumentException e) { return};
zeroFinder.evaluate();
double result = zeroFinder.getResult();
```

The first line creates the object responsible to find the zero. It uses an inner class to define the logarithm function. The second line defines the initial values, x_{pos} and x_{neg} within an exception catching construct (try...catch). The line before last performs the algorithm. The last line retrieves the result assuming the algorithm has converged.

Listing 5.2 shows the implementation of the bisection zero finding algorithm in Java.

The class `BisectionZeroFinder` is a subclass of the class `FunctionalIterator`. As one can see only a few methods need to be implemented. Most of them pertain to the definition of the initial interval.

The methods defining the initial value, x_{pos} and x_{neg} , are respectively `setPositiveX` and `setNegativeX`. Both of them throw an illegal argument exception if the function's value does not have the proper sign. The convenience methods to find random starting values are respectively `findPositiveX` and `findNegativeX`.

Listing 5.2 Java implementation of a generic derivative evaluation

```
package DhbIterations;

import DhbInterfaces.OneVariableFunction;
/**
 * Zero finding by bisection.
 *
 * @author Didier H. Besset
 */
public class BisectionZeroFinder extends FunctionalIterator
{
    /**
     * Value at which the function's value is negative.
     */
    private double xNeg;
    /**
     * Value at which the function's value is positive.
```

```

    */
    private double xPos;
/**
 * @param func DhbInterfaces.OneVariableFunction
 */
public BisectionZeroFinder(DhbInterfaces.OneVariableFunction func) {
    super(func);
}
/**
 * @param func DhbInterfaces.OneVariableFunction
 * @param x1 location at which the function yields a negative value
 * @param x2 location at which the function yields a positive value
 */
public BisectionZeroFinder( OneVariableFunction func, double x1, double x2)
    throws IllegalArgumentException
{
    this(func);
    setNegativeX( x1);
    setPositiveX( x2);
}
/**
 * @return double
 */
public double evaluateIteration()
{
    result = ( xPos + xNeg) * 0.5;
    if ( f.value(result) > 0 )
        xPos = result;
    else
        xNeg = result;
    return relativePrecision( Math.abs( xPos - xNeg));
}
/**
 * @param x double
 * @exception java.lang.IllegalArgumentException
 *         if the function's value is not negative
 */
public void setNegativeX( double x) throws IllegalArgumentException
{
    if ( f.value( x) > 0 )
        throw new IllegalArgumentException( "("+x+
            ") is positive instead of negative");
    xNeg = x;
}
/**
 * (c) Copyrights Didier BESSET, 1999, all rights reserved.

```



```

* @param x double
* @exception java.lang.IllegalArgumentException
*         if the function's value is not positive
*/
public void setPositiveX( double x) throws IllegalArgumentException
{
    if ( f.value( x) < 0 )
        throw new IllegalArgumentException( "f("+x+
            ") is negative instead of positive");
    xPos = x;
}
}

```

Figure 5.1 with the box
NewtonZeroFinder grayed.

5.3 Finding the zero of a function — Newton's method

Isaac Newton has designed an algorithm working by successive approximations[Bass]. Given a value x_0 chosen in the vicinity of the desired zero, the following series:

Main equation⇒

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}, \quad (5.3)$$

where $f'(x)$ is the first derivative of $f(x)$, converges toward a zero of the function . This algorithm is sometimes called Newton-Ralphson[Press et al.].

Figure 5.3 shows the geometrical interpretation of the series. $f'(x)$ is the slope of the tangent to the curve of the function $f(x)$ at the point x_n . The equation of this tangent is thus given by:

$$y = (x - x_n) \cdot f'(x_n) + f(x_n) \quad (5.4)$$

One can then see that x_{n+1} is the point where the tangent to the curve at the point x_n crosses the x -axis. The algorithm can be started at any point where the function's derivative is non- zero.

The technique used in Newton's algorithm is a general technique often used in approximations. The function is replaced by a linear approximation³, that is a straight line going through the point defined by the preceding value and its function's value. The slope of the straight line is given by the first derivative of the function. The procedure is repeated until the variation between the new value and the preceding one is sufficiently small. We shall see other examples of this technique in the remainder of this book (c.f. sections 10.9, 10.10 and 11.1).

From equation 5.3, one can see that the series may not converge if $f'(x)$ becomes zero. If the derivative of the function is zero in the vicinity of the zero, the bisection algorithm gives better results. Otherwise Newton's algorithm is highly efficient. It usually requires 5-10 times less iteration than the bisection algorithm. This largely compensates for the additional time spent in computing the derivative.

³Mathematically, this corresponds to estimate the function using the first two terms of its Taylor series.




Figure 5.3: Geometrical representation of Newton's zero finding algorithm

The class implementing Newton's algorithm belongs to a subclass of the functional iterator described in section 4.2. An additional instance variable is needed to store the function's derivative.

5.3.1 Newton's method — Smalltalk implementation

Listing 5.3 shows the complete implementation in Smalltalk. The class `DhbNewtonZeroFinder` is a subclass of the class `DhbFunctionalIterator` described in section 4.2.1. For example the following code finds the solution of equation 5.1.

Code example 5.3

```
| zeroFinder result |
zeroFinder := DhbNewtonZeroFinder
    function: [ :x | x errorFunction - 0.9 ]
    derivative: [ :x | DhbErfApproximation new normal: x ].
zeroFinder initialValue: 1.
result := zeroFinder evaluate.
zeroFinder hasConverged
    ifFalse: [ <special case processing> ].
```

The second line creates the object responsible to find the zero supplying the function and the derivative⁴. The third line defines the starting value. The fourth line performs the algorithm and stores the result if the algorithm has converged. The last two lines check for convergence and take corrective action if the algorithm did not converge.

The method `computeInitialValues` is somewhat complex. First, it checks whether the user supplied an initial value. If not, it is assigned to 0. Then the method checks

⁴As we have seen in section 2.3, the normal distribution is the derivative of the error function.

whether the user supplied a derivative. If not a default derivative function is supplied as a block closure by the method `defaultDerivativeBlock`. The supplied block closure implements the formula of equation 2.4 (on page 35). If a derivative is supplied, it is compared to the result of the derivative supplied by default. This may save a lot of trouble if the user made an error in coding the derivative. Not supplying a derivative has some negative effect on the speed and limits the precision of the final result. The method `initializeIterations` also checks whether the derivative is nearly zero for the initial value. If that is the case, the initial value is changed with a random walk algorithm. If no value can be found such that the derivative is non-zero an error is generated.

If the function is changed, the supplied derivative must be suppressed. Thus, the method `setFunction:` must also force a redefinition of the derivative. A method allows defining the initial value. A creation method defining the function and derivative is also supplied for convenience.

Like for the bisection, the algorithm itself is coded within the method `evaluateIteration`. Other methods needed by the algorithm have been already implemented in the superclasses.

Listing 5.3 Smalltalk implementation of Newton's zero-finding method

```

Class                DhbNewtonZeroFinder
Subclass of          DhbFunctionalIterator

Instance variable names:  derivativeBlock

Class methods

function: aBlock1 derivative: aBlock2

^(self new) setFunction: aBlock1; setDerivative: aBlock2;
               yourself

Instance methods

computeInitialValues

| n |
result isNil
  ifTrue: [ result := 0].
derivativeBlock isNil
  ifTrue: [ derivativeBlock := self defaultDerivativeBlock].
n := 0.
[ (derivativeBlock value: result) equalsTo: 0]
  whileTrue: [ n := n + 1.
              n > maximumIterations
                ifTrue: [ self error: 'Function''s derivative
                                seems to be zero everywhere'].
              result := Number random + result].

```

defaultDerivativeBlock

```
^[ :x | 5000 * ( ( functionBlock value: (x + 0.0001)) - (
    functionBlock value: (x - 0.0001)))]
```

evaluateIteration

```
| delta |
delta := ( functionBlock value: result) / ( derivativeBlock
    value: result).
result := result - delta.
^self relativePrecision: delta abs
```

initialValue: aNumber

```
result := aNumber.
```

setDerivative: aBlock

```
| x |
( aBlock respondsTo: #value:)
  ifFalse:[ self error: 'Derivative block must implement the
    method value:'].
x := result ifNil: [ Number random] ifNot: [ :base | base +
    Number random].
( ( aBlock value: x) relativelyEqualsTo: (self
    defaultDerivativeBlock value: x) upTo: 0.0001)
  ifFalse:[ self error: 'Supplied derivative is not correct'].
derivativeBlock := aBlock.
```

setFunction: aBlock

```
super setFunction: aBlock.
derivativeBlock := nil.
```

5.3.2 Newton's method — Java implementation

The class is a subclass of the class `FunctionalIterator` described in section 4.3. For example the following code finds the solution of equation 5.1.

```
NewtonZeroFinder zeroFinder = new NewtonZeroFinder(
    new OneVariableFunction() { public double value( double x)
        { return NormalDistribution.errorFunction(x) - 0.9; } });
zeroFinder.setDerivative(new OneVariableFunction()
```

```

    { public double value( double x){ return NormalDistribution.normal(x);}});
zeroFinder. setStartingValue( 1.0); zeroFinder.evaluate();
double result = zeroFinder.getResult();

```

The first line creates the object responsible to find the zero. It uses an inner class to define the error function. The second line defines the derivative also defined as an inner class. The third line defines the starting value. The line before last performs the algorithm. The last line retrieves the result assuming the algorithm has converged.

Listing 5.4 shows the complete implementation in Java.

The strong typing of Java saves us the need for checking supplied parameters. The supplied function and derivative must implement the `OneVariableFunction` interface (c.f. section 2.1.2). If the derivative function is not supplied, an object constructed with the class `FunctionDerivative` (c.f. section 2.1.2) is used instead.

Since the superclass of `NewtonZeroFinder` is an abstract class, two constructor methods are supplied. One defines the function and the initial value, the second defines also the derivative.

Otherwise the code is similar to that of the Smalltalk implementation.

Listing 5.4 Java implementation of Newton's zero-finding method

```

package DhbIterations;

import DhbFunctionEvaluation.DhbMath;
import DhbFunctionEvaluation.FunctionDerivative;
import DhbInterfaces.OneVariableFunction;
/**
 * Finds the zeroes of a function using Newton approximation.
 * Note: the zero of a function is the value at which the function's
 * value is zero.
 *
 * @author Didier H. Besset
 */
public class NewtonZeroFinder extends FunctionalIterator
{
    /**
     * Derivative of the function for which the zero will be found.
     */
    private OneVariableFunction df;
    /**
     * Constructor method.
     * @param func the function for which the zero will be found.
     * @param start the initial value for the search.
     */
    public NewtonZeroFinder( OneVariableFunction func, double start)
    {
        super( func);
    }

```

```

    setStartingValue( start);
}
/**
 * Constructor method.
 * @param func the function for which the zero will be found.
 * @param dFunc derivative of func.
 * @param start the initial value for the search.
 */
public NewtonZeroFinder( OneVariableFunction func,
                        OneVariableFunction dFunc, double start)
                        throws IllegalArgumentException
{
    this( func, start);
    setDerivative( dFunc);
}
/**
 * Evaluate the result of the current iteration.
 * @return the estimated precision of the result.
 */
public double evaluateIteration()
{
    double delta = f.value( result) / df.value( result);
    result -= delta;
    return relativePrecision( Math.abs( delta));
}
/**
 * Initializes internal parameters to start the iterative process.
 * Assigns default derivative if necessary.
 */
public void initializeIterations()
{
    if ( df == null)
        df = new FunctionDerivative( f);
    if ( Double.isNaN( result) )
        result = 0;
    int n = 0;
    while ( DhbMath.equal( df.value( result), 0) )
    {
        if ( ++n > getMaximumIterations() )
            break;
        result += Math.random();
    }
}
/**
 * (c) Copyrights Didier BESSET, 1999, all rights reserved.
 * @param dFunc DhbInterfaces.OneVariableFunction

```

```

* @exception java.lang.IllegalArgumentException
*         if the derivative is not accurate.
*/
public void setDerivative( OneVariableFunction dFunc)
                        throws IllegalArgumentException
{
    df = new FunctionDerivative( f);
    if ( !DhbMath.equal( df.value( result), dFunc.value( result), 0.001) )
        throw new IllegalArgumentException
            ( "Supplied derative function is inaccurate");
    df = dFunc;
}
/**
 * (c) Copyrights Didier BESSET, 1999, all rights reserved.
 */
public void setFunction( OneVariableFunction func)
{
    super.setFunction( func);
    df = null;
}
/**
 * Defines the initial value for the search.
 */
public void setStartingValue( double start)
{
    result = start;
}
}

```

5.4 Example of zero-finding — Roots of polynomials

The zeroes of a polynomial function are called the roots of the polynomial. A polynomial of degree n has at most n real roots. Some⁵ of them maybe complex, but are not covered in this book.

If x_0 is a root of the polynomial $P(x)$, then $P(x)$ can be exactly divided by the polynomial $x - x_0$. In other words there exists a polynomial $P_1(x)$ such that:

$$P(x) = (x - x_0) \cdot P_1(x) \quad (5.5)$$

Equation 5.5 also shows that all roots of $P_1(x)$ are also roots of $P(x)$. Thus, one can carry the search of the roots using recurrence. In practice a loop is more efficient⁶. The

⁵If the degree of the polynomial is odd, there is always at least one non-complex root. Polynomials of even degree may have only complex roots and no real roots.

⁶The overhead comes from allocating the structures needed by the method in each call.

process is repeated at most n times and will be interrupted if a zero finding step does not converge.

One could use the division algorithm of section 2.2.1 to find $P_1(x)$. In this case, however, the inner loop of the division algorithm — that is, the loop over the coefficients of the dividing polynomial — is not needed since the dividing polynomial has only two terms. In fact, one does not need to express $x - x_0$ at all as a polynomial. To carry the division one uses a specialized algorithm taking the root as the only argument. This specialized division algorithm is called deflation [Press et al.].

Polynomials are very smooth so Newton's algorithm is quite efficient for finding the first root. To ensure the best accuracy for the deflation it is recommended to find the root of smallest absolute value first. This works without additional effort since our implementation of Newton's algorithm uses 0 at the starting point by default. At each step the convergence of the zero-finder is checked. If a root could not be found the process must be stopped. Otherwise, the root finding loop is terminated when the degree of the deflated polynomial becomes zero.

5.4.1 Roots of polynomials — Smalltalk implementation

Roots of a polynomial can be obtained as an `OrderedCollection`. For example, the following code sample retrieves the roots of the polynomial $x^3 - 2x^2 - 13x - 10$:

Code example 5.4

```
(DhbPolynomial coefficients: #(-10 -13 -2 1)) roots
```

The methods needed to get the roots are shown in Listing 5.5.

The deflation algorithm is implemented in the method `deflateAt:` using the iterator method `collect:` (c.f. section B.3.2). An instance variable is keeping track of the remainder of the division within the block closure used by the method `collect:`.

The roots are kept in an `OrderedCollection` object constructed in the method `roots:`. The size of the `OrderedCollection` is initialize to the maximum expected number of real roots. Since some of the roots may be complex, we are storing the roots in an `OrderedCollection`, instead of an `Array`, so that the number of found real roots can easily be obtained. This method takes as argument the desired precision used in the zero finding algorithm. A method `root` uses the default numerical machine precision as discussed in section 1.4.

Listing 5.5 Smalltalk implementation of finding the roots of a polynomial

Class	DhbPolynomial
Subclass of	Object
Instance variable names:	coefficients

Instance methods

deflatedAt: aNumber


```

| remainder next newCoefficients|
remainder := 0.
newCoefficients := coefficients collect:
    [ :each |
      next := remainder.
      remainder := remainder * aNumber + each.
      next].
^self class new: ( newCoefficients copyFrom: 2 to:
                  newCoefficients size) reverse

roots

^self roots: DhbFloatingPointMachine new
              defaultNumericalPrecision

roots: aNumber

| pol roots x rootFinder |
rootFinder := DhbNewtonZeroFinder new.
rootFinder desiredPrecision: aNumber.
pol := self class new: ( coefficients reverse collect: [ :each |
                                                         each asFloat]).
roots := OrderedCollection new: self degree.
[ rootFinder setFunction: pol; setDerivative: pol derivative.
  x := rootFinder evaluate.
  rootFinder hasConverged
  ] whileTrue: [ roots add: x.
                pol := pol deflatedAt: x.
                pol degree > 0
                ifFalse: [ ^roots].
                ].
^roots

```

5.4.2 Roots of polynomials — Java implementation

Roots of a polynomial can be obtained as an array. For example, the following code sample retrieves the roots of the polynomial $x^3 - 2x^2 - 13x - 10$:

Code example 5.5

```

double[] coef = { -10, -13, -2, 1 };
PolynomialFunction p = new PolynomialFunction( coef);
double[] roots = p.roots();

```

Finding the roots of a polynomial in Java is implemented in three methods — deflate and roots (two flavor of this one) — of listing 2.5 (on page 46). The implementation follows that of Smalltalk.

Since the number of roots in a polynomial is not known at the beginning of the method, roots are kept in a Java Vector object. A Vector object is an array, which can grow its size. The manipulation of Vector objects, however, differs significantly from that of an array. First, the elements of a Vector object are instances of class Object. Thus, using the elements always requires casting into the proper class. The roots are primitive types (double) so they must first be cast into the corresponding wrapper class Double. Looping over the elements of a Vector object must be made using an enumeration instead of the conventional for loop.

5.5 Which method to choose

There are other zero-finding techniques: regula falsi, Brent [Press et al.]. For each of these methods, however, a specialist of numerical methods can design a function causing that particular method to fail.

In practice the bisection algorithm is quite slow as can be seen from equation 5.2. Newton's algorithm is faster for most functions you will encounter. For example, it takes 5 iterations to find the zero of the logarithm function with Newton's algorithm to a precision of $3 \cdot 10^{-9}$ whereas the bisection algorithm requires 29 to reach a similar precision. On the other hand bisection is rock solid and will always converge over an interval where the function has no singularity. Thus, it can be used as a recovery when Newton's algorithm fails.

My own experience is that Newton's algorithm is quite robust and very fast. It should suffice in most cases. As we have seen Newton's algorithm will fail if it encounters a value for which the derivative of the function is very small. In this case, the algorithm jumps far away from the solution. For these cases, the chances are that the bisection algorithm will find the solution if there is any. Thus, combining Newton's algorithm with bisection is the best strategy if you need to design a foolproof algorithm.

Implementing an object combining both algorithms is left as an exercise to the reader. Here is a quick outline of the strategy to adopt. Newton's algorithm must be modified to keep track of values for which the function takes negative values and positive values — that is the values x_{pos} and x_{neg} — making sure that the value $|x_{\text{pos}} - x_{\text{neg}}|$ never increases. Then, at each step, one must check that the computed change does not cause the solution to jump outside of the interval defined by x_{pos} and x_{neg} . If that is the case, Newton's algorithm must be interrupted for one step using the bisection algorithm.

Chapter 6

Integration of functions

Les petits ruisseaux font les grandes rivières¹
French proverb

Many functions are defined by an integral. For example, the three functions discussed in the last 3 sections of chapter 2 were all defined by an integral. When no other method is available the only way to compute such function is to evaluate the integral. Integrals are also useful in probability theory to compute the probability of obtaining a value over a given interval. This aspect will be discussed in chapter 9. Finally integrals come up in the computation of surfaces and of many physical quantities related to energy and power. For example, the power contained in an electromagnetic signal is proportional to the integral of the square of the signal's amplitude.

The French proverb quoted at the beginning of this chapter is here to remind people that an integral is defined formally as the infinite sum of infinitesimal quantities.

6.1 Introduction

Let us begin with a concrete example. This time we shall take a problem from physics 101.

When light is transmitted through a narrow slit, it is diffracted. The intensity of the light transmitted at an angle ϑ , $I(\vartheta)$, is given by:

$$I(\vartheta) = \frac{\sin^2 \vartheta}{\vartheta^2} \quad (6.1)$$

If one wants to compute the fraction of light which is transmitted within the first diffraction peak, one must compute the expression:

$$I(\vartheta) = \frac{1}{\pi} \int_{-\pi}^{\pi} \frac{\sin^2 \vartheta}{\vartheta^2} d\vartheta. \quad (6.2)$$

¹Small streams build great rivers.

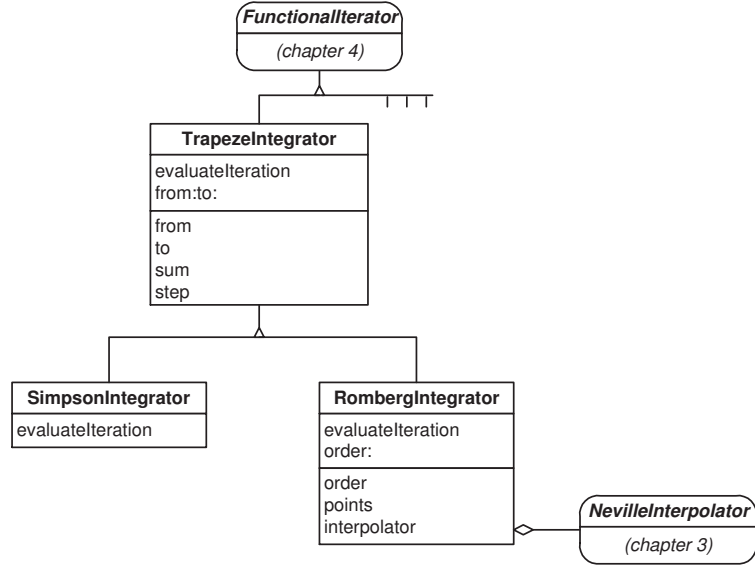


Figure 6.1: Class diagram of integration classes

The division by π is there because the integral of $I(\vartheta)$ from $-\infty$ to $+\infty$ is equal to π . No closed form exists for the integral of equation 6.2: it must be computed numerically. This answer is 90.3%.

In this chapter we introduce 3 integration algorithms. Figure 6.1 shows the corresponding class diagram. The first one, trapeze integration, is only introduced for the sake of defining a common framework for the next two algorithms: Simpson and Romberg integration. In general, the reader should use Romberg's algorithm. It is fast and very precise. There are, however, some instances where Simpson's algorithm can be faster if high accuracy is not required.

6.2 General framework — Trapeze integration method

Let us state it at the beginning. One should not use the trapeze integration algorithm in practice. The interest of this algorithm is to define a general framework for numerical integration. All subclasses of the class responsible for implementing the trapeze integration algorithm will reuse most the mechanisms described in this section.

The trapeze numerical integration method takes its origin in the series expansion of an integral. This series expansion is expressed by the Euler-Maclaurin formula shown hereafter [Bass]:

$$\int_a^b f(x) dx = \frac{b-a}{2} [f(a) + f(b)] - \sum_n \frac{(b-a)^2}{(2n)!} B_{2n} \left[\frac{d^{2n-1} f(b)}{dx^{2n-1}} - \frac{d^{2n-1} f(a)}{dx^{2n-1}} \right], \quad (6.3)$$

where the numbers B_{2n} are the Bernouilli numbers.




Figure 6.2: Geometrical interpretation of the trapeze integration method

The next observation is that, if the interval of integration is small enough, the series in the second term of equation 6.3 would yield a contribution negligible compared to that of the first term. Thus, we can write:

$$\int_a^b f(x) dx \approx \frac{b-a}{2} [f(a) + f(b)], \quad (6.4)$$

if $b-a$ is sufficiently small. The approximation of equation 6.4 represents the area of a trapeze whose summits are the circled points in Figure 6.2. Finally, one must remember the additive formula between integrals:

$$\int_a^b f(x) dx = \int_a^c f(x) dx + \int_c^b f(x) dx, \quad (6.5)$$

for any c . We shall use this property by choosing a c located between a and b .

The resulting strategy is a divide-and-conquer strategy. The integration interval is divided until one can be sure that the second term of equation 6.3 becomes indeed negligible. As one would like to re-use the points at which the function has been evaluated during the course of the algorithm, the integration interval is halved at each iteration. The first few steps are outlined in figure 6.2. An estimation of the integral is obtained by summing the areas of the trapezes corresponding to each partition.

Let $x_0^{(n)}, \dots, x_{2^n}^{(n)}$ be the partition of the interval at iteration n . Let $\epsilon^{(n)}$ be the length of each interval between these points. We thus have:

$$\begin{cases} \epsilon^{(n)} &= \frac{b-a}{2^n} \\ x_0^{(n)} &= a \\ x_i^{(n)} &= a + i\epsilon^{(n)} \quad \text{for } i = 1, \dots, 2^n. \end{cases} \quad (6.6)$$

The corresponding estimation for the integral is:

$$I^{(n)} = \epsilon^{(n)} \left[f(a) + f(b) + 2 \sum_{i=1}^{2^n-1} f(x_i^{(n)}) \right]. \quad (6.7)$$

To compute the next estimation, it suffices to compute the value of the function at the even values of the partition because the odd values were already computed before. One can derive the following recurrence relation:

Main equation \Rightarrow

$$I^{(n+1)} = \frac{I^{(n)}}{2} + \epsilon^{(n)} \sum_{i=1}^{2^n-1} f(x_{2i-1}^{(n)}), \quad (6.8)$$

with the initial condition:

$$I^{(0)} = \frac{b-a}{2} [f(a) + f(b)]. \quad (6.9)$$

Note that the sum on the right-hand side of equation 6.8 represents the sum of the function's values at the new points of the partition.

End game strategy

The final question is when should the algorithm be stopped? A real honest answer is we do not know. The magnitude of the series in equation 6.3 is difficult to estimate as the Bernoulli numbers become very large with increasing n . An experimental way is to watch for the change of the integral estimate. In other words the absolute value of the last variation, $|I^{(n)} - I^{(n+1)}|$, is considered a good estimate of the precision. This kind of heuristic works for most functions.

At each iteration the number of function evaluation doubles. This means that the time spent in the algorithm grows exponentially with the number of iterations. Thus, the default maximum number of iteration must be kept quite low compared to that of the other methods.

In practice, however, trapeze integration converges quite slowly and should not be used. Why bother implementing it then? It turns out that the more elaborate methods, Simpson and Romberg integration, require the computation of the same sums needed by the trapeze integration. Thus, the trapeze integration is introduced to be the superclass of the other better integration methods.

One must keep in mind, however, that the magnitude of the series in equation 6.3 can become large for any function whose derivatives of high orders have singularities over the interval of integration. The convergence of the algorithm can be seriously compromised for such functions. This remark is true for the other algorithms described in this chapter. For example, none of the algorithms is able to give a precise estimate of the beta function using equation 2.32 with $x > 1$ and $y < 1$ (c.f. section 2.5) because, for these values, the derivatives of the function to integrate have a singularity at $t = 1$.

Another problem can come up if the function is nearly zeroes at regular intervals. For example, evaluating the integral of the function $f(x) = \frac{\sin(2^m x)}{x}$ from $-\pi$ to π for a moderate value of m . In this case, the terms $I^{(0)}$ to $I^{(m)}$ will have a null contribution.

This would cause the algorithm to stop prematurely. Such special function behavior is of course quite rare. Nevertheless the reader must be aware of the limitations of the algorithm. This remark is valid for all algorithms exposed in this chapter.

6.2.1 Trapeze integration — General implementation

The class implementing trapeze integration is a subclass of the functional iterator discussed in section 4.2. Two instance variables are needed to define the integration interval. Additional instance variables must keep track of the partition of the interval and of the successive estimations. Consequently, the class has the following instance variables.

`from` contains the lower limit of the integration's interval, i.e. a .

`to` contains the upper limit of the integration's interval, i.e. b .

`step` contains the size of the interval's partition, i.e. $\epsilon^{(n)}$.

`sum` contains the intermediate sums, i.e. $I^{(n)}$.

Although trapeze integration is not a practical algorithm, we give an example of coding for both language implementations. The reason is that the public interface used by trapeze integration is the same for all integration classes.

The example shown in the next two sections is the integration of the inverse function. In mathematics the natural logarithm of x , $\ln x$, is defined as the integral from 1 to x of the inverse function. Of course, using numerical integration is a very inefficient way of computing a logarithm. This example, however, allows the reader to investigate the accuracy (since the exact solution is known) and performances of all algorithms presented in this chapter. The interested reader should try the example for various setting of the desired precision and look at the number of iterations needed for a desired precision. She can also verify how accurate is the estimated precision.

6.2.2 Trapeze integration — Smalltalk implementation

Listing 6.1 shows the Smalltalk implementation of the trapeze integration method. In Smalltalk the code for the computation of the integral defining the natural logarithm is as follows:

Code example 6.1

```

| integrator ln2 ln3 |
integrator := DhbTrapezeIntegrator function: [:x | 1.0 / x] from: 1 to: 2.
ln2 := integrator evaluate.
integrator from: 1 to: 3.
ln3 := integrator evaluate.
```

The line after the declaration creates a new instance of the class `DhbTrapezeIntegrator` for the inverse function. The limits of the integration interval are set from 1 to 2 at

Figure 6.1 with the box `TrapezeIntegrator` grayed.

creation time. The third line retrieves the value of the integral. The fourth line changes the integration interval and the last line retrieves the value of the integral over the new integration interval.

The class `DhbTrapezeIntegrator` is a subclass of the class `AbstractFunctionIterator` defined in section 4.2.1. The default creation class method `new` has been overloaded to prevent creating an object without initialized instance variables. The proper creation class method defines the function and the integration interval.

The method `from:to:` allows changing the integration interval for a new computation with the same function.

Note that the initialization of the iterations (method `computeInitialValues`, c.f. section 4.1.1) also corresponds to the first iteration of the algorithm. The method `highOrderSum` computes the sum of the right-hand side of equation 6.8.

Listing 6.1 Smalltalk implementation of trapeze integration

```

Class                                DhbTrapezeIntegrator
Subclass of                          DhbFunctionalIterator
Instance variable names:  from to sum step

Class methods
defaultMaximumIterations
^13

new
^self error: 'Method new:from:to: must be used'

new: aBlock from: aNumber1 to: aNumber2
^super new initialize: aBlock from: aNumber1 to: aNumber2

Instance methods
computeInitialValues

step := to - from.
sum := ( ( functionBlock value: from) + ( functionBlock value:
                                         to)) * step /2.

result := sum.

evaluateIteration

```


The first line creates a new instance of the class `TrapezeIntegrator` for the inverse function defined as an inner class. The limits of the integration interval are set from 1 to 2. The second line retrieves the value of the integral. The third line changes the integration interval and the last line retrieves the value of the integral over the new integration interval. Required checks for convergence are not included in this example.

The class `TrapezeIntegrator` is a subclass of the class `FunctionalIterator` defined in section 4.2.2. A single constructor method creates a new instance with a function and the integration interval. The method `setInterval` allows changing the integration interval for a new computation with the same function.

Note that the initialization of the iterations (method `initializeIterations`, c.f. section 4.1.2) also corresponds to the first iteration of the algorithm. The method `highOrderSum` computes the sum of the right-hand side of equation 6.8.

Listing 6.2 Java implementation of trapeze integration

```
package DhbIterations;

import DhbInterfaces.OneVariableFunction;
/**
 * Trapeze integration method
 *
 * @author Didier H. Besset
 */
public class TrapezeIntegrator extends FunctionalIterator
{
    /**
     * Low integral bound.
     */
    private double from;
    /**
     * High integral bound.
     */
    private double to;
    /**
     * Sum
     */
    protected double sum;
    /**
     * Interval partition.
     */
    private double step;
    /**
     * Constructor
     * @param func DhbInterfaces.OneVariableFunction
     * @param from double
     * @param to double
     */
}
```

```
*/
public TrapezeIntegrator(OneVariableFunction f, double from, double to)
{
    super( f);
    setInterval( from, to);
    setMaximumIterations( 13);
}
public double evaluateIteration()
{
    double oldResult = result;
    result = highOrderSum();
    return relativePrecision( Math.abs( result - oldResult));
}
/**
 * @return double
 */
protected double highOrderSum()
{
    double x = from + 0.5 * step;
    double newSum = 0;
    while ( x < to )
    {
        newSum += f.value(x);
        x += step;
    }
    sum = ( step * newSum + sum) * 0.5;
    step *= 0.5;
    return sum;
}
public void initializeIterations()
{
    step = to - from;
    sum = ( f.value(from) + f.value(to)) * step * 0.5;
    result = sum;
}
/**
 * Defines integration interval.
 * @param double a low integral bound.
 * @param double b high integral bound.
 */
public void setInterval( double a, double b)
{
    from = a;
    to = b;
}
}
```

Figure 6.1 with the box
SimpsonIntegrator grayed.

6.3 Simpson integration algorithm

Simpson integration algorithm consists in replacing the function to integrate by a second order Lagrange interpolation polynomial[Bass] (c.f. section 3.2). One can then carry the integral analytically. Let $f(x)$ be the function to integrate. For a given interval of integration, $[a, b]$, the function is evaluated at the extremities of the interval and at its middle point $c = \frac{a+b}{2}$. As defined in equation 3.1 the second order Lagrange interpolation polynomial is then given by:

$$P_2(x) = \frac{2}{(b-a)^2} [(x-b)(x-c)f(a) + (x-c)(x-a)f(b) + (x-a)(x-b)f(c)]. \quad (6.10)$$

The integral of the polynomial over the interpolation interval is given by:

$$\int_a^b P_2(x) dx = \frac{b-a}{6} [f(a) + f(b) + 4f(c)]. \quad (6.11)$$

As for trapeze integration, the interval is partitioned into small intervals. Let us assume that the interval has been divided into subintervals. By repeating equation 6.11 over each subinterval, we obtain:

$$\int_a^b P_2(x) dx = \frac{\epsilon^{(n)}}{3} \left[f(a) + f(b) + 2 \sum_{i=1}^{2^{n-1}} f(x_{2i-1}^{(n)}) + 4 \sum_{i=0}^{2^{n-1}} f(x_{2i}^{(n)}) \right]. \quad (6.12)$$

Equation 6.12 uses the notations introduced in section 6.2. Except for the first iteration, the right-hand side of equation 6.12 can be computed from the quantities $I^{(n)}$ defined in equation 6.7. Thus, we have:

Main equation \Rightarrow

$$\int_a^b P_2(x) dx = \frac{1}{3} [4I^{(n)} - I^{(n-1)}] \quad \text{for } n > 1 \quad (6.13)$$

This can be checked by verifying that $I^{(n-1)}$ is equal to the first sum of equation 6.12 times and that $I^{(n)}$ is equal to the addition of the two sums of equation 6.12 times $\epsilon^{(n)}$. As advertised in section 6.2 we can re-use the major parts of the trapeze algorithm: computation of the sums and partition of the integration interval.

Like in the case of the trapeze algorithm, the precision of the algorithm is estimated by looking at the differences between the estimation obtained previously and the current estimation. At the first iteration only one function point is computed. This can cause the process to stop prematurely if the function is nearly zero at the middle of the integration interval. Thus, a protection must be built in to prevent the algorithm from stopping at the first iteration.

6.3.1 Simpson integration — General implementation

The class implementing Simpson algorithm is a subclass of the class implementing trapeze integration. The method `evaluateIteration` is the only method needing change.

The number of iterations is checked to prevent returning after the first iteration.

The public interface is the same as that of the superclass. Thus, all the examples shown in sections 6.2.2 and 6.2.3 can be used for Simpson algorithm by just changing the name of the class.

6.3.2 Simpson integration — Smalltalk implementation

Listing 6.3 shows the complete implementation in Smalltalk.

The class `DhbSimpsonIntegrator` is a subclass of the class `DhbTrapezeIntegrator` defined in section 6.2.2.

Listing 6.3 Smalltalk implementation of the Simpson integration algorithm

```

Class      DhbSimpsonIntegrator
Subclass of DhbTrapezeIntegrator

    Instance methods

evaluateIteration

    | oldResult oldSum |
    iterations < 2
        ifTrue: [ self higherOrderSum.
            ^1
        ].
    oldResult := result.
    oldSum := sum.
    result := (self higherOrderSum * 4 - oldSum) / 3.
    ^self relativePrecision: ( result - oldResult) abs

```

6.3.3 Simpson integration — Java implementation

Listing 6.4 shows the complete implementation in Java.

Except for the need of an explicit constructor method, the Java implementation is exactly similar to that in Smalltalk.

Listing 6.4 Java implementation of the Simpson integration algorithm

```

package DhbIterations;

/**
 * Simpson integration method
 *
 * @author Didier H. Besset

```

```

*/
public class SimpsonIntegrator extends TrapezeIntegrator
{
/**
 * SimpsonIntegrator constructor.
 * @param f DhbInterfaces.OneVariableFunction
 * @param from double
 * @param to double
 */
public SimpsonIntegrator(DhbInterfaces.OneVariableFunction f,
                        double from, double to)
{
    super(f, from, to);
}
/**
 * @return double
 */
public double evaluateIteration()
{
    if ( getIterations() < 2 )
    {
        highOrderSum();
        return getDesiredPrecision( );
    }
    double oldResult = result;
    double oldSum = sum;
    result = ( 4 * highOrderSum() - oldSum ) / 3.0;
    return relativePrecision( Math.abs( result - oldResult ));
}
}

```

6.4 Romberg integration algorithm

Figure 6.1 with the box RombergIntegrator grayed.

If one goes back to equation 6.3 one notices that the second term is of the order of the square of the integration interval. Romberg's algorithm uses this fact to postulate that $I^{(n)}$ is a smooth function of the square of the size of interval's partition $\epsilon^{(n)}$. Romberg's algorithm introduces the parameter k where $k - 1$ is the degree of the interpolation's polynomial². The result of the integral is estimated by extrapolating the series $I^{(n-k)}, \dots, I^{(n)}$ at the value $\epsilon^{(n)} = 0$. Since we have:

$$\epsilon^{(n)} = \frac{\epsilon^{(n-1)}}{2}, \quad (6.14)$$

²In other words, k is the number of points over which the interpolation is performed (c.f. section 3.2).

one only needs to interpolate over successive powers of $1/4$, starting with 1: 1, $1/4$, $1/16$, $1/256$, etc. In this case, extrapolation is safe because the value at which extrapolation is made is very close to the end of the interval defined by the sample points and actually becomes closer and closer with every iteration.

Thus, Romberg's algorithm requires at least k iterations. The good news is that this algorithm converges very quickly and, in general, only a few iterations are needed after the five initial ones. A polynomial of 4^{th} degree — that is $k = 5$ — is generally sufficient [Press et al.].

Extrapolation is performed using Neville's algorithm (c.f. section 3.4) because it computes an estimate of the error on the interpolated value. That error estimate can then be used as the estimate of the error on the final result.

If $k = 1$ Romberg's algorithm is equivalent to trapeze integration. If $k = 2$, the interpolation polynomial is given by:

$$P_1(x) = y_1 + \frac{x - x_1}{x_2 - x_1} (y_2 - y_1). \quad (6.15)$$

At the n^{th} iteration we have: $y_1 = I^{(n-1)}$, $y_2 = I^{(n)}$ and $x_2 = x_1/4$. Thus, the interpolated value at 0 is:

$$P_1(0) = I^{(n-1)} + \frac{4}{3} [I^{(n)} - I^{(n-1)}] = \frac{1}{3} [4I^{(n)} - I^{(n-1)}] \quad (6.16)$$

Thus, for $k = 2$ Romberg's algorithm is equivalent to Simpson's algorithm. For higher order, however, Romberg's algorithm is much more efficient than Simpson method as soon as precision is required (c.f. a comparison of the result in section 6.6).

Using interpolation on the successive results of an iterative process to obtain the final result is a general procedure known as Richardson's deferred approach to the limit [Press et al.]. This technique can be used whenever the estimated error can be expressed as a function of a suitable parameter depending on the iterations. The choice of the parameter, however, is critical. For example, if one had interpolated over the size of the interval's partition instead of its square, the method would not converge as well³.

6.4.1 Romberg integration — General implementation

The class implementing Romberg's algorithm needs the following additional instance variables:

`order` the order of the interpolation, i.e. k ,

`interpolator` an instance of Neville's interpolation class,

`points` an `OrderedCollection` containing the most recent sum estimates, i.e. $I^{(n-k)}, \dots, I^{(n)}$.

³It converges at the same speed as Simpson's algorithm. This can be verified by running the comparison programs after changing the factor 0.25 used to compute the abscissa of the next point into 0.5.

The method `evaluateIteration` (c.f. section 4.1) contains the entire algorithm. At each iteration the collection of point receives a new point with an abscissa equal to the quarter of that of the last point and an ordinate equal to the next sum estimate . If not enough points are available, the method returns a precision such that the iterative process will continue. Otherwise, the extrapolation is performed. After the result of the extrapolation has been obtained the oldest point is removed. In other words, the collection of points is used as a last-in-last-out list with a constant number of elements equal to the order of the interpolation. Of the two values returned by Neville's interpolation (c.f. section 3.4), the interpolated value is stored in the result and the error estimate is returned as the precision for the other.

6.4.2 Romberg integration — Smalltalk implementation

Listing 6.5 shows the Smalltalk implementation of Romberg's algorithm. The class `DhbRombergIntegrator` is a subclass of the class `DhbTrapezeIntegrator` defined in section 6.2.2.

The class method `defaultOrder` defines the default order to 5. This method is used in the method `initialize` so that each newly created instance is created with the default interpolation order. The method `order:` allows changing the default order if needed.

The sample points defining the interpolation are stored in an `OrderedCollection`. This collection is created in the method `computeInitialValues`. Since the number of points will never exceed the order of the interpolation the maximum size is preset when the collection is created. The method `computeInitialValues` also creates the object in charge of performing the interpolation and it stores the first point in the collection of sample points.

Listing 6.5 Smalltalk implementation of Romberg integration

```

Class                DhbRombergIntegrator
Subclass of          DhbTrapezeIntegrator
Instance variable names:  order points interpolator

    Class methods
defaultOrder
    ^5

    Instance methods
computeInitialValues
    super computeInitialValues.
    points := OrderedCollection new: order.
    interpolator := DhbNevilleInterpolator points: points.
    points add: 1 @ sum.
```

```
evaluateIteration
```

```
    | interpolation |
    points addLast: (points last x * 0.25) @ self higherOrderSum.
    points size < order
        ifTrue: [ ^1].
    interpolation := interpolator valueAndError: 0.
    points removeFirst.
    result := interpolation at: 1.
    ^self relativePrecision: ( interpolation at: 2) abs
```

```
initialize
```

```
    order := self class defaultOrder.
    ^super initialize
```

```
order: anInteger
```

```
    anInteger < 2
        ifTrue: [ self error: 'Order for Romberg integration must be
                                larger than 1'].
    order := anInteger.
```

6.4.3 Romberg integration — Java implementation

Listing 6.6 shows the Java implementation of Romberg's algorithm. The class `DhbRombergIntegrator` is a subclass of the class `TrapezeIntegrator` defined in section 6.2.3.

The default order is assigned in the instance variable so that each newly created instance is created with the default interpolation order (5). The method `setOrder` allows changing the default order when needed.

The sample points defining the interpolation are stored in a structure implementing the `PointSeries` interface as discussed in section 3.2.2. The method `initializeIterations` creates this structure. It also creates the `interpolatorobject` in charge of performing the interpolation and it stores the first point in the collection of sample points.

Listing 6.6 Java implementation of Romberg integration

```
package DhbIterations;

import DhbInterfaces.OneVariableFunction;
import DhbInterpolation.NevilleInterpolator;
import DhbScientificCurves.Curve;
/**
```

```

* Romberg integration method
*
* @author Didier H. Besset
*/
public class RombergIntegrator extends TrapezeIntegrator
{
    /**
     * Order of the interpolation.
     */
    private int order = 5;
    /**
     * Structure containing the last estimations.
     */
    private Curve estimates;
    /**
     * Neville interpolator.
     */
    private NevilleInterpolator interpolator;
    /**
     * RombergIntegrator constructor.
     * @param func DhbInterfaces.OneVariableFunction
     * @param from double
     * @param to double
     */
    public RombergIntegrator(DhbInterfaces.OneVariableFunction func,
                             double from, double to)
    {
        super(func, from, to);
    }
    /**
     * @return double
     */
    public double evaluateIteration()
    {
        estimates.addPoint( estimates.xValueAt(estimates.size() - 1) * 0.25,
                           highOrderSum());
        if ( estimates.size() < order )
            return 1;
        double[] interpolation = interpolator.valueAndError( 0);
        estimates.removePointAt( 0);
        result = interpolation[0];
        return relativePrecision( Math.abs( interpolation[1]));
    }
    public void initializeIterations()
    {
        super.initializeIterations();
    }

```

```

    estimates = new Curve();
    interpolator = new NevilleInterpolator( estimates);
    estimates.addPoint( 1, sum);
}
/**
 * @param n int
 */
public void setOrder( int n)
{
    order = n;
}
}

```

6.5 Evaluation of open integrals

An open integral is an integral for which the function to integrate cannot be evaluated at the boundaries of the integration interval. This is the case when one of the limit is infinite or when the function to integrate has a singularity at one of the limits. If the function to integrate has one or more singularity in the middle of the integration interval, the case can be reduced to that of having a singularity at the limits using the additive formula between integrals 6.5. Generalization of the trapeze algorithm and the corresponding adaptation of Romberg's algorithm can be found in [Press et al.].

Bag of tricks

My experience is that using a suitable change of variable can often remove the problem. In particular, integrals whose integration interval extends to infinity can be rewritten as integrals over a finite interval. We give a few examples below.

For an integral starting from minus infinity, a change of variable $t = \frac{1}{x}$ can be used as follows:

$$\int_{-\infty}^a f(x) dx = \int_{\frac{1}{a}}^0 f\left(\frac{1}{t}\right) \frac{dt}{t^2} \quad \text{for } a < 0. \quad (6.17)$$

For such integral to be defined, the function must vanish at minus infinity faster than x^2 . This means that:

$$\lim_{t \rightarrow 0} \frac{1}{t^2} f\left(\frac{1}{t}\right) = 0. \quad (6.18)$$

If $a > 0$, the integration must be evaluated in two steps, for example one over the interval $]-\infty, -1]$ using the change of variable of equation 6.17 and one over the interval $[-1, a]$ using the original function.

For integral ending at positive infinity the same change of variable can also be made. However, if the interval of integration is positive, the change of variable $t = e^{-x}$ can be more efficient. In this case, one makes the following transformation:

$$\int_a^{+\infty} f(x) dx = \int_0^{e^{-a}} f(-\ln t) \frac{dt}{t} \quad \text{for } a > 0. \quad (6.19)$$

For this integral to be defined one must have:

$$\lim_{t \rightarrow 0} \frac{1}{t} f(\ln t) = 0. \quad (6.20)$$

By breaking up the interval of integration is several pieces one can chose a change of variable most appropriate for each piece.

6.6 Which method to chose?

An example comparing the results of the three algorithms is given in section 6.6.1 for Smalltalk and section 6.6.2 for Java. The function to integrate is the inverse function in both cases. The integration interval is from 1 to 2 so that the value of the result is known, namely $\ln 2$. Integration is carried for various values of the desired precision. The reader can then compare the attained precision (both predicted and real) and the number of iterations⁴ required for each algorithm. Let us recall that the number of required function evaluations grows exponentially with the number of iterations.

The results clearly show that the trapeze algorithm is ruled out as a practical method. As advertised in section 6.2 it is not converging quickly toward a solution.

Romberg's algorithm is the clear winner. At given precision, it requires the least number of iterations. This is the algorithm of choice in most cases.

Simspon's algorithm may be useful if the required precision is not too high and if the time to evaluate the function is small compared to the interpolation. In such cases Simspon's algorithm can be faster than Romberg's algorithm.

Sections 6.6.1 and 6.6.2 gives some sample code the reader can use to investigate the various integration algorithms. The results of the code execution are shown in table 6.1. The columns of this table are:

ϵ_{\max} the desired precision,

n the number of required iterations; let us recall that the corresponding number of function's evaluations is 2^{n+1} ,

$\tilde{\epsilon}$ the estimated precision of the result,

ϵ the effective precision of the result, that is the absolute value of the difference between the result of the integration algorithm and the true result.

⁴Note that the number of iterations printed in the examples is one less than the real number of iterations because the first iteration is performed in the set-up phase of the iterative process.

Table 6.1: Comparison between integration algorithms

ϵ_{\max}	Trapeze algorithm			Simpson algorithm			Romberg algorithm		
	n	$\tilde{\epsilon}$	ϵ	n	$\tilde{\epsilon}$	ϵ	n	$\tilde{\epsilon}$	ϵ
10^{-5}	8	$4.1 \cdot 10^{-6}$	$9.5 \cdot 10^{-7}$	4	$9.9 \cdot 10^{-6}$	$4.7 \cdot 10^{-7}$	4	$1.7 \cdot 10^{-9}$	$1.4 \cdot 10^{-9}$
10^{-7}	11	$6.4 \cdot 10^{-8}$	$1.5 \cdot 10^{-8}$	6	$4.0 \cdot 10^{-8}$	$1.9 \cdot 10^{-9}$	4	$1.7 \cdot 10^{-9}$	$1.4 \cdot 10^{-9}$
10^{-9}	15	$2.5 \cdot 10^{-10}$	$5.8 \cdot 10^{-11}$	8	$1.5 \cdot 10^{-10}$	$7.3 \cdot 10^{-12}$	5	$1.4 \cdot 10^{-11}$	$3.7 \cdot 10^{-12}$
10^{-11}	18	$3.9 \cdot 10^{-12}$	$9.0 \cdot 10^{-13}$	9	$9.8 \cdot 10^{-12}$	$5.7 \cdot 10^{-13}$	6	$7.6 \cdot 10^{-14}$	$5.7 \cdot 10^{-15}$
10^{-13}	21	$4.8 \cdot 10^{-14}$	$2.8 \cdot 10^{-14}$	11	$3.8 \cdot 10^{-14}$	$1.9 \cdot 10^{-15}$	6	$7.6 \cdot 10^{-14}$	$5.7 \cdot 10^{-15}$

6.6.1 Smalltalk comparison

The script of Listing 6.7 can be executed as such in any Smalltalk window.

The function to integrate is specified as a block closure as discussed in section 2.1.1.

Listing 6.7 Smalltalk comparison script for integration algorithms

```

| a b integrators |
a := 1.0.
b := 2.0.
integrators := Array with: ( DhbTrapezeIntegrator new: [ :x | 1.0 / x] from: a to: b)
with: ( DhbSimpsonIntegrator new: [ :x | 1.0 / x] from: a to: b)
with: ( DhbRombergIntegrator new: [ :x | 1.0 / x] from: a to: b).
#(1.0e-5 1.0e-7 1.0e-9 1.0e-11 1.0e-13) do: [ :precision |
Transcript cr; cr; nextPutAll: '==> Precision: '.
precision printOn: Transcript.
integrators do: [ :integrator |
Transcript cr; nextPutAll: '***** ', integrator class name, ':'; cr.
integrator desiredPrecision: precision.
Transcript nextPutAll: 'Integral of 1/x from '.
a printOn: Transcript.
Transcript nextPutAll: ' to '.
b printOn: Transcript.
Transcript nextPutAll: ' = '.
integrator evaluate printOn: Transcript.
Transcript nextPutAll: ' +- '.
integrator precision printOn: Transcript.
Transcript cr; nextPutAll: ' ( '.
integrator iterations printOn: Transcript.
Transcript nextPutAll: ' iterations, true error = '.
( integrator result - 2 ln) printOn: Transcript.
Transcript nextPutAll: ')'; cr.
]]

```

6.6.2 Java comparison

The method shown in Listing 6.8 can be compiled as a static method in any class. To execute the comparison test this method should be called from a main method with the desired precision. The function to integrate is implemented as an inner class as discussed in section 2.1.2.

Listing 6.8 Java comparison method for integration algorithms

```
public static void executeAtPrecision( double precision)
{
    double a = 1;
    double b = 2;
    System.out.println(" ");
    System.out.println("====> desired precision = " + precision);
    System.out.println("***** Trapeze method:");
    try {
        TrapezeIntegrator integratorT = new TrapezeIntegrator(
            new OneVariableFunction()
                { public double value( double x){ return 1/x; }}, a, b);
        integratorT.setDesiredPrecision( precision);
        integratorT.setMaximumIterations( 25);
        integratorT.evaluate();
        System.out.print("Integral of ln(x) from "+a+" to "+b);
        System.out.println(" is "+integratorT.getResult()+" +/- "
            + integratorT.getPrecision());
        System.out.print(" ( "+integratorT.getIterations()+" iterations");
        System.out.println(", error = "+(Math.log( b)
            - integratorT.getResult())+"");
    } catch ( IllegalArgumentException e)
    { System.out.println("Illegal precision specified in integration");};
    System.out.println("***** Simpson method:");
    try {
        SimpsonIntegrator integratorS = new SimpsonIntegrator(
            new OneVariableFunction()
                { public double value( double x){ return 1/x; }}, a, b);
        integratorS.setDesiredPrecision( precision);
        integratorS.evaluate();
        System.out.print("Integral of ln(x) from "+a+" to "+b);
        System.out.println(" is "+integratorS.getResult()+" +/- "
            +integratorS.getPrecision());
        System.out.print(" ( "+integratorS.getIterations()+" iterations");
        System.out.println(", error = "+(Math.log( b)
            - integratorS.getResult())+"");
    } catch ( IllegalArgumentException e)
    { System.out.println("Illegal precision specified in integration");};
}
```

```
System.out.println("***** Romberg method:");
try {
    RombergIntegrator integratorR = new RombergIntegrator(
        new OneVariableFunction()
        { public double value( double x){ return 1/x;}} , a, b);
    integratorR.setDesiredPrecision( precision);
    integratorR.evaluate();
    System.out.print("Integral of ln(x) from "+a+" to "+b);
    System.out.println(" is "+integratorR.getResult()+" +/- "
        +integratorR.getPrecision());
    System.out.print(" ( "+integratorR.getIterations()+" iterations");
    System.out.println(", error = "+(Math.log( b)
        - integratorR.getResult()+""));
} catch ( IllegalArgumentException e)
{ System.out.println("Illegal precision specified in integration");};
}
```

Chapter 7

Series

On ne peut pas partir de l'infini, on peut y aller.¹
Jules Lachelier

Whole families of functions are defined with infinite series expansion or a continued fraction. Before the advent of mechanical calculators, a person could earn a Ph.D. in mathematics by publishing tabulated values of a function evaluated by its series expansion or continued fraction. Some people developed a talent to perform such tasks.

Some reported stories make the task of evaluating series sound like a real business. A German nobleman discovered that one of its peasants had a talent for numbers. He then housed him in his mansion and put him to work on the calculation of a table of logarithms. The table was published under the nobleman's name[Ifrah].

Nowadays we do not need to look for some talented peasant, but we still publish the result of computations made by other than ourselves. Overall computers are better treated than peasants were, though. . .

7.1 Introduction

It will not come as a surprise to the reader that the computation of infinite series is made on a computer by computing a sufficient but finite number of terms. The same is true for continued fractions. Thus, the computation of infinite series and continued fractions uses the iterative process framework described in chapter 4. In this case the iteration consists of computing successive terms.

The present chapter begins by exposing a general framework on how to compute infinite series and continued fractions. Then, we show two examples of application of this framework by implementing two functions, which are very important to compute probabilities: the incomplete gamma function and the incomplete beta function.

For illustrative purposes, the implementation in Smalltalk is using a different architecture from the one used by the Java implementation. It should be noted that each implementation could have been implemented in the other language. Figure 7.1 shows the class diagram

¹One cannot start at infinity; one can reach it, however.

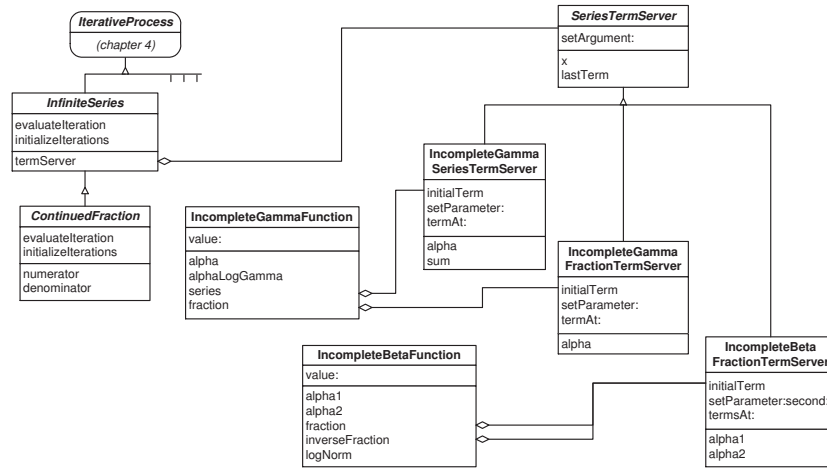


Figure 7.1: Smalltalk class diagram for infinite series and continued fractions

of the Smalltalk implementation. Figure 7.2 shows the class diagram of the Java implementation.

The Smalltalk implementation uses two general-purpose classes to implement an infinite series and a continued fraction respectively. Each class then use a Strategy pattern class [Gamma et al.] to compute each term of the expansion.

The Java implementation uses two abstract classes to implement an infinite series and a continued fraction respectively. Each concrete implementation necessitates the creation of a concrete subclass.

In spite of the difference in architecture, the reader can verify on each class diagram that the number of classes needed for a concrete implementation is the same in each case.

An interesting exercise for the reader is to implement the architecture presented in Java in Smalltalk and vice versa.

7.2 Infinite series

Many functions are defined with an infinite series, that is a sum of an infinite number of terms. The most well known example is the series for the exponential function:

$$e^x = \sum_{n=0}^{\infty} \frac{x^n}{n!}. \quad (7.1)$$

For such a series to be defined, the terms of the series must become very small as the index increases. If that is the case an infinite series may be used to evaluate a function, for which no closed expression exists. For this to be practical, however, the series should converge quickly so that only a few terms are needed. For example, computing




Figure 7.2: Java class diagram for infinite series and continued fractions

the exponential of 6 to the precision of an IEEE 32 bit floating number requires nearly 40 terms. This is clearly not an efficient way to compute the exponential.

Discussing the convergence of a series is outside of the scope of this book. Let us just state that in general numerical convergence of a series is much harder to achieve than mathematical convergence. In other words the fact that a series is defined mathematically does not ensure that it can be evaluated numerically.

A special remark pertains to alternating series. In an alternating series the signs of consecutive terms are opposite. Trigonometric functions have such a series expansion. Alternating series have very bad numerical convergence properties: if the terms are large rounding errors might suppress the convergence altogether. If one cannot compute the function in another way it is best to compute the terms of an alternating series in pairs to avoid rounding errors.

In practice, a series must be tested for quick numerical convergence prior to its implementation. As for rounding errors the safest way is to do this experimentally, that is, print out the terms of the series for a few representative² values of the variable. In the rest of this chapter we shall assume that this essential step has been made.

To evaluate an infinite series, one carries the summation until the last added term becomes smaller than the desired precision. This kind of logic is quite similar to that of an iterative process. Thus, the object used to compute an infinite series belongs to a subclass of the iterative process classe discussed in chapter 4.

²By representative, I mean either values, which are covering the domain over which the function will be evaluated, or values, which are suspected to give convergence problems.

Figure 7.1 with the boxes InfiniteSeries and SeriesTermServer grayed.

7.2.1 Infinite series — Smalltalk implementation

Listing 7.1 shows a general Smalltalk implementation of a class evaluating an infinite series. The class being abstract, we do not give examples here. Concrete examples are given in section 7.4.2.

The Smalltalk implementation uses a Strategy pattern. The class `DhbInfiniteSeries` is a subclass of the class `DhbIterativeProcess`, discussed in section 4.1.1. This class does not implement the algorithm needed to compute the terms of the series directly. It delegates this responsibility to an object stored in the instance variable `termServer`. Two hook methods, `initialTerm` and `termAt:` are used to obtain the terms of the series from the term server object.

The method `evaluateIteration` uses the method `precisionOf:relativeTo:` to return a relative precision as discussed in section 4.2.1.

To implement a specific series, an object of the class `DhbInfiniteSeries` is instantiated with a specific term server. A concrete example will be shown in section 7.4.2.

Because of its generic nature, the class `DhbInfiniteSeries` does not implement the function behavior described in section 2.1.1 (method `value:`). It is the responsibility of each object combining an infinite series with a specific term server to implement the function behavior. An example is given in section 7.4.2.

Listing 7.1 Smalltalk implementation of an infinite series

```

Class                DhbInfiniteSeries
Subclass of          DhbIterativeProcess
Instance variable names: termServer

    Class methods
server: aTermServer

^self new initialize: aTermServer

    Instance methods
evaluateIteration

| delta |
delta := termServer termAt: iterations.
result := result + delta.
^self precisionOf: delta abs relativeTo: result abs

initialize: aTermServer

termServer := aTermServer.
^self

```

```
initializeIterations
```

```
result := termServer initialTerm
```

The computation of the terms of the series is delegated to an object instantiated from a server class. The abstract server class is called `DhbInfiniteSeriesTermServer`. It is responsible to compute the terms at each iteration. This class receives the argument of the function defined by the series, which is kept in the instance variable `x`. The instance variable `lastTerm` is provided to keep the last computed term since the next term can often be computed from the previous one. The code of this abstract class is shown in Listing 44.

Listing 7.2 Smalltalk implementation of a term server

Class	DhbSeriesTermServer
Subclass of	Object
Instance variable names:	x lastTerm

Instance methods

setArgument: aNumber

x := aNumber asFloat.

7.2.2 Infinite series — Java implementation

Listing 7.3 shows the implementation of an infinite series in Java. The class being abstract, we do not give examples here. Concrete examples are given in section 7.4.3.

The Java implementation uses an abstract class, `DhbInfiniteSeries` — a subclass of the class `IterativeProcess` described in section 4.1.2 — to evaluate the series. Abstract methods define the common interface needed to retrieve the terms of the series. As a series is a function, it must implement the `OneVariableFunction` interface. Specific series are implemented as a subclass of the class `DhbInfiniteSeries`, each subclass implementing the methods needed to compute each term.

The argument of the series and each last term are kept into protected instance variables for efficiency purposes. Thus, subclasses can have direct access to these variables without the need to call an accessor method.

The method `evaluateIteration` uses the method `relativePrecision` to return a relative precision as discussed in section 4.2.2.

Listing 7.3 Java implementation of an infinite series

Figure 7.2 with the box `InfiniteSeries` grayed.

```
package DhbIterations;

/**
 * InfiniteSeries
 *
 * @author Didier H. Besset
 */
public abstract class InfiniteSeries extends IterativeProcess
{
    /**
     * Best approximation of the sum.
     */
    private double result;
    /**
     * Series argument.
     */
    protected double x;
    /**
     * Value of the last term.
     */
    protected double lastTerm;

    /**
     * Computes the n-th term of the series and stores it in lastTerm.
     * @param n int
     */
    protected abstract void computeTermAt ( int n);
    public double evaluateIteration()
    {
        computeTermAt( getIterations());
        result += lastTerm;
        return relativePrecision( Math.abs( lastTerm), Math.abs( result));
    }
    /**
     * @return double
     */
    public double getResult ( )
    {
        return result;
    }
    /**
     * Set the initial value for the sum.
     */
    public void initializeIterations()
    {
        result = initialValue();
    }
}
```

```

}
/**
 * @return double    the 0-th term of the series
 */
protected abstract double initialValue ( );
/**
 * @param r double   the value of the series argument.
 */
public void setArgument ( double r)
{
    x = r;
    return;
}
}

```

7.3 Continued fractions

A continued fraction is an infinite series of cascading fractions of the following form:

$$f(x) = b_0 + \frac{a_1}{b_1 + \frac{a_2}{b_2 + \frac{a_3}{b_3 + \frac{a_4}{b_4 + \dots}}}} \quad (7.2)$$

In general, both sets of coefficients a_0, \dots and b_0, \dots depend on the function's argument x . This dependence is implicit in equation 7.2 to keep the notation simple. Since the above expression is quite awkward to read - besides being a printer's nightmare - one usually uses a linear notation as follows [Abramovitz & Stegun], [Press et al.]:

$$f(x) = b_0 + \frac{a_1}{b_1 +} \frac{a_2}{b_2 +} \frac{a_3}{b_3 +} \frac{a_4}{b_4 +} \dots \quad (7.3)$$

The problem in evaluating such a fraction is that, a priori, one must begin the evaluation from the last term. Fortunately, methods allowing the evaluation from the beginning of the fractions have been around since the seventeen's century. A detailed discussion of several methods is given in [Press et al.]. In this book, we shall only discuss the modified Lentz' method which has the advantage to work for a large class of fractions.

Implementing the other methods discussed in [Press et al.] is left as an exercise to the reader. The corresponding classes can be subclassed from the classes found in this chapter.

In 1976, Lentz proposed the following two auxiliary series:

$$\begin{cases} C_0 &= b_0, \\ D_0 &= 0, \\ C_n &= \frac{a_n}{C_{n-1}} + b_n \quad \text{for } n > 0, \\ D_n &= \frac{1}{a_n D_{n-1} + b_n} \quad \text{for } n > 0. \end{cases} \quad (7.4)$$

These two series are used to construct the series:

$$\begin{cases} f_0 &= C_0, \\ f_n &= f_{n-1}C_nD_n. \end{cases} \quad (7.5)$$

One can prove by induction that this series converges toward the continued fraction as n gets large.

In general continued fractions have excellent convergence properties. Some care, however, must be given when one of the auxiliary terms C_n or $1/D_n$ become nearly zero³. To avoid rounding errors, Thompson and Barnett, in 1986, proposed a modification of the Lentz method in which any value of the coefficients smaller than a small floor value is adjusted to the floor value [Press et al.]. The floor value is chosen to be the machine precision of the floating-point representation (instance variable `smallNumber` described in section 1.4).

In terms of architecture, the implementation of a continued fraction is similar to that of the infinite series.

7.3.1 Continued fractions — Smalltalk implementation

Figure 7.1 with the box ContinuedFraction grayed.

Listing 7.4 shows the implementation of a continued fraction in Smalltalk.

The class `DhbContinuedFraction` is built as a subclass of the class `DhbInfiniteSeries`. Thus, it uses also the Strategy pattern.

The method `limitedSmallValue:` implements the prescription of Thompson and Barnett.

Listing 7.4 Smalltalk implementation of a continued fraction

```

Class                DhbContinuedFraction
Subclass of          DhbInfiniteSeries

Instance variable names:  numerator denominator

    Instance methods

evaluateIteration

    | terms delta |
    terms := termServer termsAt: iterations.
    denominator := 1 / ( self limitedSmallValue: ( (terms at: 1) *
                                                    denominator + (terms at: 2))).
    numerator := self limitedSmallValue: ( (terms at: 1) / numerator
                                           + (terms at: 2)).

    delta := numerator * denominator.
    result := result * delta.
    ^ ( delta - 1) abs

```

³That is, a value which is zero within the precision of the numerical representation.

```
initializeIterations
```

```
    numerator := self limitedSmallValue: termServer initialTerm.
    denominator := 0.
    result := numerator
```

```
limitedSmallValue: aNumber
```

```
    ^aNumber abs < DhbFloatingPointMachine new smallNumber
      ifTrue: [ DhbFloatingPointMachine new smallNumber]
      ifFalse:[ aNumber]
```

7.3.2 Continued fractions — Java implementation

Listing 7.5 shows the implementation of a continued fraction in Java.

As for infinite series, an abstract class is in charge of implementing the modified Lentz method. As the two classes have nothing in common, there is no reason for the class `ContinuedFraction` to be a subclass of `InfiniteSeries`. It is a subclass of the class `IterativeProcess` instead.

The method `limitedSmallValue` implements the prescription of Thompson and Barnett.

Figure 7.2 with the box `ContinuedFraction` grayed.

Listing 7.5 Java implementation of an infinite series

```
package DhbIterations;

import DhbFunctionEvaluation.DhbMath;
/**
 * Continued fraction
 *
 * @author Didier H. Besset
 */
public abstract class ContinuedFraction extends IterativeProcess
{
    /**
     * Best approximation of the fraction.
     */
    private double result;
    /**
     * Fraction's argument.
     */
    protected double x;
    /**
     * Fraction's accumulated numerator.
     */
```

```

private double numerator;
/**
 * Fraction's accumulated denominator.
 */
private double denominator;
/**
 * Fraction's next factors.
 */
protected double[] factors = new double[2];
/**
 * Compute the pair numerator/denominator for iteration n.
 * @param n int
 */
protected abstract void computeFactorsAt(int n);
/**
 * @return double
 */
public double evaluateIteration()
{
    computeFactorsAt( getIterations());
    denominator = 1 / limitedSmallValue( factors[0] * denominator
                                         + factors[1]);
    numerator = limitedSmallValue( factors[0] / numerator + factors[1]);
    double delta = numerator * denominator;
    result *= delta;
    return Math.abs( delta - 1);
}
/**
 * @return double
 */
public double getResult ( )
{
    return result;
}
public void initializeIterations()
{
    numerator = limitedSmallValue( initialValue());
    denominator = 0;
    result = numerator;
    return;
}
/**
 * @return double
 */
protected abstract double initialValue();
/**

```

```

* Protection against small factors.
* @return double
* @param r double
*/
private double limitedSmallValue ( double r)
{
    return Math.abs( r) < DhbMath.smallNumber()
        ? DhbMath.smallNumber() : r;
}
/**
* @param r double    the value of the series argument.
*/
public void setArgument ( double r)
{
    x = r;
    return;
}
}

```

7.4 Incomplete Gamma function

The incomplete gamma function is the integral of a gamma distribution. It is used in statistics to evaluate the probability of finding a measurement larger than a given value when the measurements are distributed according to a gamma distribution. In particular, the incomplete gamma function is used to compute the confidence level of χ^2 values when assessing the validity of a parametric fit. Several examples of use of this function will be introduced in chapters 9 and 10.

Figure 7.3 shows the incomplete gamma function (solid line) and its corresponding probability density function (dotted line) for $\alpha = 2.5$.

The gamma distribution is discussed in section 9.7. The χ^2 confidence level is discussed in section 10.3. General χ^2 fits are discussed in section 10.9.

7.4.1 Mathematical definitions

The incomplete gamma function is defined by the following integral:

$$\Gamma(x, \alpha) = \frac{1}{\Gamma(\alpha)} \int_0^x t^{\alpha-1} e^{-t} dt. \quad (7.6)$$

Thus, the value of the incomplete gamma function lies between 0 and 1. The function has one parameter α . The incomplete gamma function is the distribution function of a gamma probability density function with parameters α and 1 (c.f. section 9.7 for a description of the gamma distribution and its parameters). This integral can be expressed




Figure 7.3: The incomplete gamma function and the gamma distribution

as the following infinite series [Abramovitz & Stegun]:

$$\Gamma(x, \alpha) = \frac{e^{-x} x^\alpha}{\Gamma(\alpha)} \sum_{n=0}^{\infty} \frac{\Gamma(\alpha)}{\Gamma(\alpha + 1 + n)} x^n. \quad (7.7)$$

Written in this form we can see that each term of the series can be computed from the previous one. Using the recurrence formula for the gamma function — equation 2.24 in section 2.4.1 — we have:

Main equation⇒

$$\begin{cases} a_0 &= \frac{1}{\alpha}, \\ a_n &= \frac{x}{\alpha + 1 + n} a_{n-1}. \end{cases} \quad (7.8)$$

The series in equation 7.7 converges well for $x < \alpha + 1$.

The incomplete gamma function can also be written as [Abramovitz & Stegun]:

$$\Gamma(x, \alpha) = \frac{e^{-x} x^\alpha}{\Gamma(\alpha)} \frac{1}{F(x - \alpha + 1, \alpha)}, \quad (7.9)$$

where $F(x, \alpha)$ is the continued fraction:

$$F(x, \alpha) = x + \frac{1(\alpha - 1)}{x + 2} + \frac{2(\alpha - 2)}{x + 4} + \frac{3(\alpha - 3)}{x + 6} + \dots \quad (7.10)$$

Using the notation introduced in equation 7.3 in section 7.3 the terms of the continued fraction are given by the following expressions:

Main equation⇒

$$\begin{cases} b_n &= x - \alpha + 2n \quad \text{for } n = 0, 1, 2, \dots \\ a_n &= n(\alpha - n) \quad \text{for } n = 1, 2, \dots \end{cases} \quad (7.11)$$

It turns out that the continued fraction in equation 7.9 converges for $x > \alpha + 1$ [Press et al.], that is, exactly where the series expansion of equation 7.7 did not converge

very well. Thus, the incomplete gamma function can be computed using one of the two methods depending on the range of the argument.

The reader will notice that equations 7.7 and 7.9 have a common factor. The denominator of that factor can be evaluated in advance in logarithmic form to avoid floating-point overflow (c.f. discussion in section 2.4.1). For each function evaluation the entire factor is computed in logarithmic form to reduce rounding errors. Then it is combined with the value of the series or the continued fraction to compute the final result.

To avoid a floating-point error when evaluating the common factor, the value of the incomplete gamma function at $x = 0$ — which is of course 0 — must be returned separately.

7.4.2 Incomplete Gamma function — Smalltalk implementation

Three classes are needed to implement the incomplete gamma function in Smalltalk. The class `DhbIncompleteGamaFunction` is in charge of computing the function itself. This is the object, which responds to the method `value:` to provide a function-like behavior to the object. It is shown in Listing 7.6 and has the following instance variables.

Figure 7.1 with the boxes `IncompleteGammaFunction`, `IncompleteGammaSeriesTermServer` and `IncompleteGammaFractionTermSe` grayed.

`alpha` contains the function's parameter, i.e. α ,

`alphaLogGamma` used to cache the value of $\Gamma(\alpha)$ for efficiency purposes,

`series` contains the infinite series associated to the function,

`fraction` contains the continued fraction associated to the function.

The instance variables `series` and `fraction` are assigned using lazy initialization.

Depending on the range of the argument, the class delegates the rest of the computing to either a series or a continued fraction. In each case, a term server class provides the computation of the terms. They are shown in listings 7.7 and 7.8.

Listing 7.6 Smalltalk implementation of the incomplete gamma function

```

Class          DhbIncompleteGammaFunction
Subclass of    Object
Instance variable names:  alpha alphaLogGamma series fraction

```

Class methods

`shape:` aNumber

`^super new initialize:` aNumber

Instance methods

`evaluateFraction:` aNumber

```

fraction isNil
  ifTrue:
    [fraction := DhbIncompleteGammaFractionTermServer new.
     fraction setParameter: alpha].
fraction setArgument: aNumber.
^(DhbContinuedFraction server: fraction)
  desiredPrecision: DhbFloatingPointMachine new
    defaultNumericalPrecision;
  evaluate

evaluateSeries: aNumber

series isNil
  ifTrue: [ series := DhbIncompleteGammaSeriesTermServer new.
    series setParameter: alpha.
  ].
series setArgument: aNumber.
^(DhbInfiniteSeries server: series)
  desiredPrecision: DhbFloatingPointMachine new
    defaultNumericalPrecision;
  evaluate

initialize: aNumber

alpha := aNumber asFloat.
alphaLogGamma := alpha logGamma.
^self

value: aNumber

| x norm |
aNumber = 0
  ifTrue: [ ^0].
x := aNumber asFloat.
norm := [ ( x ln * alpha - x - alphaLogGamma) exp] when: ExAll
  do: [ :signal | signal exitWith: nil].

norm isNil
  ifTrue: [ ^1].
^x - 1 < alpha
  ifTrue: [ ( self evaluateSeries: x) * norm]
  ifFalse:[ 1 - ( norm / ( self evaluateFraction: x))]
```

Listing 7.7 shows the implementation of the term server for the series expansion. It needs two instance variables: one to store the parameter α ; one to store the sum accumulated

in the denominator of equation 7.8. The two lines of equation 7.8 are implemented respectively by the methods `initialTerm` (for $n = 0$) and `termAt:` (for $n \geq 1$).

Listing 7.7 Smalltalk implementation of the series term server for the incomplete gamma function

```

Class                DhbIncompleteGammaSeriesTermServer
Subclass of          DhbSeriesTermServer

Instance variable names:  alpha sum

    Instance methods

initialTerm

    lastTerm := 1 / alpha.
    sum := alpha.
    ^lastTerm

setParameter: aNumber

    alpha := aNumber asFloat

termAt: anInteger

    sum := sum + 1.
    lastTerm := lastTerm * x / sum.
    ^lastTerm

```

Listing 7.8 shows the implementation of the term server for the continued fraction. It needs one instance variable to store the parameter α . Equation 7.11 is implemented by the methods `initialTerm` (for $n = 0$) and `termsAt:` (for $n \geq 1$).

Listing 7.8 Smalltalk implementation of the fraction term server for the incomplete gamma function

```

Class                DhbIncompleteGammaFractionTermServer
Subclass of          DhbSeriesTermServer

Instance variable names:  alpha

    Instance methods

initialTerm

```



```
lastTerm := x - alpha + 1.
^lastTerm
```

```
setParameter: aNumber
```

```
alpha := aNumber asFloat
```

```
termsAt: anInteger
```

```
lastTerm := lastTerm + 2.
^Array with: (alpha - anInteger) * anInteger with: lastTerm
```

An example of use of the incomplete gamma function can be found in section 9.7.1.

7.4.3 Incomplete Gamma function — Java implementation

Figure 7.2 with the boxes IncompleteGammaFunction, IncompleteGammaFunctionSeries and IncompleteGammaFunctionFraction grayed.

In spite of the difference of architecture with Smalltalk, three classes are also needed in the Java implementation.

Listing 7.9 shows the implementation of the class IncompleteGammaFunction implementing the incomplete gamma function proper. It is constructed as the corresponding Smalltalk class.

Listing 7.9 Java implementation of the incomplete gamma function

```
package DhbIterations;

import DhbFunctionEvaluation.DhbMath;
import DhbFunctionEvaluation.GammaFunction;

/**
 * IncompleteGamma function
 *
 * @author Didier H. Besset
 */
public class IncompleteGammaFunction implements DhbInterfaces.OneVariableFunction
{
    /**
     * Function parameter.
     */
    private double alpha;
    /**
     * Constant to be computed once only.
     */
    private double alphaLogGamma;
```

```

/**
 * Infinite series.
 */
private IncompleteGammaFunctionSeries series;
/**
 * Continued fraction.
 */
private IncompleteGammaFunctionFraction fraction;

/**
 * Constructor method.
 */
public IncompleteGammaFunction ( double a)
{
    alpha = a;
    alphaLogGamma = GammaFunction.logGamma( alpha);
}
/**
 * @return double
 * @param x double
 */
private double evaluateFraction ( double x)
{
    if ( fraction == null )
    {
        fraction = new IncompleteGammaFunctionFraction( alpha);
        fraction.setDesiredPrecision(
            DhbMath.defaultNumericalPrecision());
    }
    fraction.setArgument( x);
    fraction.evaluate();
    return fraction.getResult();
}
/**
 * @return double    evaluate the series of the incomplete gamma function.
 * @param x double
 */
private double evaluateSeries ( double x)
{
    if ( series == null )
    {
        series = new IncompleteGammaFunctionSeries( alpha);
        series.setDesiredPrecision(
            DhbMath.defaultNumericalPrecision());
    }
    series.setArgument( x);

```

```

    series.evaluate();
    return series.getResult();
}
/**
 * Returns the value of the function for the specified variable value.
 */
public double value(double x)
{
    if ( x == 0 )
        return 0;
    double norm = Math.exp( Math.log(x) * alpha - x - alphaLogGamma);
    return x - 1 < alpha
        ? evaluateSeries( x) * norm
        : 1 - norm / evaluateFraction( x);
}
}

```

Listing 7.10 shows the implementation of the term server used by the infinite series defining the incomplete gamma function in Java. It is implemented as a subclass of the class `InfiniteSeries` defined in section 7.2.1. The class `IncompleteGammaFunctionSeries` needs two instance variables: One to store the parameter α ; one to store the sum accumulated in the denominator of equation 7.8. The two lines of equation 7.8 are implemented respectively by the methods `initialValue` (for $n = 0$) and `computeTermAt`: (for $n \geq 1$).

Listing 7.10 Java implementation of the infinite series term server for the incomplete gamma function

```

package DhblIterations;

/**
 * Series for the incompleteGamma function
 *
 * @author Didier H. Besset
 */
public class IncompleteGammaFunctionSeries extends InfiniteSeries
{
    /**
     * Series parameter.
     */
    private double alpha;
    /**
     * Auxiliary sum.
     */
    private double sum;
}

```

```

/**
 * Constructor method
 * @param a double series parameter
 */
public IncompleteGammaFunctionSeries ( double a)
{
    alpha = a;
}
/**
 * Computes the n-th term of the series and stores it in lastTerm.
 * @param n int
 */
protected void computeTermAt(int n)
{
    sum += 1;
    lastTerm *= x / sum;
    return;
}
/**
 * initializes the series and return the 0-th term.
 */
protected double initialValue()
{
    lastTerm = 1 / alpha;
    sum = alpha;
    return lastTerm;
}
}

```

Listing 7.11 shows the implementation of the term server used by the continued fraction defining the incomplete gamma function in Java. It is implemented as a subclass of the class `ContinuedFraction` defined in section 7.3.2. The class `IncompleteGammaFunctionFraction` needs one instance variable to store the parameter x . Equation 7.11 is implemented by the methods `initialValue` (for $n = 0$) and `computeFactorsAt`: (for $n \geq 1$).

Listing 7.11 Java implementation of the fraction term server for the incomplete gamma function

```

package DhIterations;

/**
 * Continued fraction for the incompleteGamma function
 *
 * @author Didier H. Besset

```

```

*/
public class IncompleteGammaFunctionFraction extends ContinuedFraction
{
    /**
     * Series parameter.
     */
    private double alpha;
    /**
     * Auxiliary sum.
     */
    private double sum;

    /**
     * Constructor method.
     * @param a double
     */
    public IncompleteGammaFunctionFraction ( double a)
    {
        alpha = a;
    }
    /**
     * Compute the pair numerator/denominator for iteration n.
     * @param n int
     */
    protected void computeFactorsAt(int n)
    {
        sum += 2;
        factors[0] = ( alpha - n) * n;
        factors[1] = sum;
        return;
    }
    protected double initialValue()
    {
        sum = x - alpha + 1;
        return sum;
    }
}

```

An example of use of the incomplete gamma function can be found in section 9.7.2.

7.5 Incomplete Beta function

The incomplete beta function is the integral of a beta distribution. It used in statistics to evaluate the probability of finding a measurement larger than a given value when the measurements are distributed according to a beta distribution. It is also used to compute




Figure 7.4: The incomplete beta function and the beta distribution

the confidence level of the Student distribution (t -test) and of the Fisher-Snedecor distribution (F -test). The beta distribution is discussed in section D.1. The t -test is discussed in section 10.2. The F -test is discussed in section 10.1.

Figure 7.4 shows the incomplete beta function (solid line) and its corresponding probability density function (dotted line) for $\alpha_1 = 4.5$ and $\alpha_2 = 2.5$.

7.5.1 Mathematical definitions

The incomplete beta function is defined over the interval $[0, 1]$ by the following integral:

$$B(x; \alpha_1, \alpha_2) = \frac{1}{B(\alpha_1, \alpha_2)} \int_0^x t^{\alpha_1-1} (1-t)^{\alpha_2-1} dt, \quad (7.12)$$

where $B(\alpha_1, \alpha_2)$ is the beta function defined in section 2.5. The function has two parameters α_1 and α_2 . By definition, the value of the incomplete beta function is comprised between 0 and 1.

None of the series expansions of this integral have good numerical convergence. There is, however, a continued fraction development which converges over a sufficient range [Abramovitz & Stegun]:

$$B(x; \alpha_1, \alpha_2) = \frac{x^{\alpha_1-1} (1-x)^{\alpha_2-1}}{\alpha_1 B(\alpha_1, \alpha_2)} \frac{1}{F(x; \alpha_1, \alpha_2)}, \quad (7.13)$$

where

$$F(x; \alpha_1, \alpha_2) = 1 + \frac{a_1}{1+} \frac{a_2}{1+} \frac{a_3}{1+} \dots \quad (7.14)$$

Using the notation introduced in section 7.3 we have:

⇐Main equation

$$\begin{cases} b_n &= 1 \text{ for } n = 0, 1, 2, \dots \\ a_{2n} &= \frac{n(\alpha_2 - n)x}{(\alpha_1 + 2n)(\alpha_1 + 2n - 1)} \text{ for } n = 1, 2, \dots \\ a_{2n+1} &= \frac{(\alpha_1 + n)(\alpha_1 + \alpha_2 + n)x}{(\alpha_1 + 2n)(\alpha_1 + 2n - 1)} \text{ for } n = 1, 2, \dots \end{cases} \quad (7.15)$$

The continued fraction in equation 7.13 converges rapidly for $x > \frac{\alpha_1+1}{\alpha_1+\alpha_2+2}$ [Press et al.]. To compute the incomplete beta function over the complementary range, one uses the following symmetry property of the function:

$$B(x; \alpha_1, \alpha_2) = 1 - B(1 - x; \alpha_2, \alpha_1) \quad (7.16)$$

Since $1 - x < \frac{\alpha_2+1}{\alpha_1+\alpha_2+2}$ if $x < \frac{\alpha_1+1}{\alpha_1+\alpha_2+2}$, we can now compute the function over the entire range.

To avoid a floating-point error when evaluating the leading factor of equation 7.13, the values of the incomplete beta function at $x = 0$ — which is 0 — and at $x = 1$ — which is 1 — must be returned separately.

7.5.2 Incomplete Beta function — Smalltalk implementation

Figure 7.1 with the boxes IncompleteBetaFunction and IncompleteBetaFractionTermServer grayed.

Listing 7.12 shows the implementation of the incomplete beta function in Smalltalk. Two classes are needed to implement the incomplete beta function. The class DhbIncompleteBetaFunction is in charge of computing the function itself. This class has the following instance variables.

alpha1 contains the first function's parameter, i.e. α_1 ,

alpha2 contains the second function's parameter, i.e. α_2 ,

logNorm used to cache the value of $\ln B(\alpha_1, \alpha_2)$ for efficiency purposes,

fraction contains the continued fraction associated to the function $B(x; \alpha_1, \alpha_2)$,

inverseFraction contains the continued fraction associated to the function $B(1 - x; \alpha_2, \alpha_1)$.

Depending on the range of the argument, the class delegates the rest of the computing to a continued fraction using the original parameters or the reversed parameters if the symmetry relation must be used. A term server class allows the computing of the terms. Its code is shown in listing 7.13. The two instance variables - fraction and inverseFraction -, contain an instance of the term server, one for each permutation of the parameters, thus preventing the unnecessary creation of new instances of the term server at each evaluation. These instance variables are assigned using lazy initialization.

Listing 7.12 Smalltalk implementation of the incomplete beta function

Class	DhbIncompleteBetaFunction
Subclass of	Object

Instance variable names: alpha1 alpha2 fraction inverseFraction logNorm

Class methods

shape: aNumber1 shape: aNumber2

^super new initialize: aNumber1 shape: aNumber2

Instance methods

evaluateFraction: aNumber

fraction isNil

ifTrue:

[fraction := DhbIncompleteBetaFractionTermServer new.
fraction setParameter: alpha1 second: alpha2].

fraction setArgument: aNumber.

^(DhbContinuedFraction server: fraction)

desiredPrecision: DhbFloatingPointMachine new
defaultNumericalPrecision;

evaluate

evaluateInverseFraction: aNumber

inverseFraction isNil

ifTrue:

[inverseFraction := DhbIncompleteBetaFractionTermServer
new.

inverseFraction setParameter: alpha2 second: alpha1].

inverseFraction setArgument: (1 - aNumber).

^(DhbContinuedFraction server: inverseFraction)

desiredPrecision: DhbFloatingPointMachine new
defaultNumericalPrecision;

evaluate

initialize: aNumber1 shape: aNumber2

alpha1 := aNumber1.

alpha2 := aNumber2.

logNorm := (alpha1 + alpha2) logGamma - alpha1 logGamma - alpha2
logGamma.

^self

value: aNumber


```

| norm |
aNumber = 0
ifTrue: [ ^0].
aNumber = 1
ifTrue: [ ^1].
norm := ( aNumber ln * alpha1 + ( ( 1 - aNumber) ln * alpha2) +
          logNorm) exp.
^( alpha1 + alpha2 + 2) * aNumber < ( alpha1 + 1)
ifTrue: [ norm / ( ( self evaluateFraction: aNumber) *
                  alpha1)]
ifFalse:[ 1 - ( norm / ( ( self evaluateInverseFraction:
                          aNumber) * alpha2))]

```

Listing 7.13 shows the implementation of the term server. It needs two instance variables to store the parameters α_1 and α_2 . Equation 7.15 is implemented by the methods `initialTerm` (for $n = 0$) and `termsAt:` (for $n \geq 1$).

Listing 7.13 Smalltalk implementation of the term server for the incomplete beta function

```

Class                DhbIncompleteBetaFractionTermServer
Subclass of          DhbSeriesTermServer
Instance variable names:  alpha1 alpha2

```

```

    Instance methods
initialTerm
    ^1

setParameter: aNumber1 second: aNumber2

    alpha1 := aNumber1.
    alpha2 := aNumber2

termsAt: anInteger

    | n n2 |
    n := anInteger // 2.
    n2 := 2 * n.
    ^Array with: ( n2 < anInteger
    ifTrue:
        [x negated * (alpha1 + n) * (alpha1 + alpha2 + n)
          / ((alpha1 + n2) * (alpha1 + 1 + n2))]

```

```

ifFalse: [x * n * (alpha2 - n) / ((alpha1 + n2) * (alpha1 - 1
                                + n2))]
with: 1

```

An example of use of the incomplete beta function can be found in sections 10.1 and 10.2.

7.5.3 Incomplete Beta function — Java implementation

The Java implementation of the incomplete beta function needs two classes.

Listing 7.14 shows the implementation of the class `IncompleteBetaFunction` implementing the incomplete beta function proper. It is constructed as the corresponding Smalltalk class.

Figure 7.2 with the boxes `IncompleteBetaFunction` and `IncompleteBetaFunctionFraction` grayed.

Listing 7.14 Java implementation of the incomplete beta function

```

package DhbIterations;

import DhbFunctionEvaluation.DhbMath;
import DhbFunctionEvaluation.GammaFunction;
import DhbInterfaces.OneVariableFunction;

/**
 * Incomplete Beta function
 *
 * @author Didier H. Besset
 */
public class IncompleteBetaFunction implements OneVariableFunction
{
    /**
     * Function parameters.
     */
    private double alpha1;
    private double alpha2;
    /**
     * Constant to be computed once only.
     */
    private double logNorm;
    /**
     * Continued fractions.
     */
    private IncompleteBetaFunctionFraction fraction;
    private IncompleteBetaFunctionFraction inverseFraction;

```

```

/**
 * Constructor method.
 * @param a1 double
 * @param a2 double
 */
public IncompleteBetaFunction ( double a1, double a2)
{
    alpha1 = a1;
    alpha2 = a2;
    logNorm = GammaFunction.logGamma( alpha1 + alpha2)
              - GammaFunction.logGamma( alpha1)
              - GammaFunction.logGamma( alpha2);
}
/**
 * @return double
 * @param x double
 */
private double evaluateFraction ( double x)
{
    if ( fraction == null )
    {
        fraction = new IncompleteBetaFunctionFraction( alpha1, alpha2);
        fraction.setDesiredPrecision( DhbMath.defaultNumericalPrecision());
    }
    fraction.setArgument( x);
    fraction.evaluate();
    return fraction.getResult();
}
/**
 * @return double
 * @param x double
 */
private double evaluateInverseFraction ( double x)
{
    if ( fraction == null )
    {
        fraction = new IncompleteBetaFunctionFraction( alpha2, alpha1);
        fraction.setDesiredPrecision( DhbMath.defaultNumericalPrecision());
    }
    fraction.setArgument( x);
    fraction.evaluate();
    return fraction.getResult();
}
public double value(double x)
{
    if ( x == 0 )

```

```

    return 0;
    if ( x == 1 )
        return 1;
    double norm = Math.exp( alpha1 * Math.log(x)
        + alpha2 * Math.log(1 - x) + logNorm);
    return ( alpha1 + alpha2 + 2) * x < ( alpha1 + 1)
        ? norm / ( evaluateFraction( x) * alpha1)
        : 1 - norm / ( evaluateInverseFraction(1 - x)
            * alpha2);
}
}

```

Listing 7.15 shows the implementation of the term server used by the continued fraction defining the incomplete beta function in Java. It needs two instance variables to store the parameters α_1 and α_2 . Equation 7.15 is implemented by the methods `initialTerm` (for $n = 0$) and `computeFactorsAt`: (for $n \geq 1$).

Listing 7.15 Java implementation of the continued fraction for the incomplete beta function

```

package DhbIterations;

/**
 * Incomplete Beta function fraction
 *
 * @author Didier H. Besset
 */
public class IncompleteBetaFunctionFraction extends ContinuedFraction
{
    /**
     * Fraction's parameters.
     */
    private double alpha1;
    private double alpha2;

    /**
     * Constructor method.
     * @param a1 double
     * @param a2 double
     */
    public IncompleteBetaFunctionFraction ( double a1, double a2)
    {
        alpha1 = a1;
        alpha2 = a2;
    }
}

```

```
/**
 * Compute the pair numerator/denominator for iteration n.
 * @param n int
 */
protected void computeFactorsAt(int n)
{
    int m = n / 2;
    int m2 = 2 * m;
    factors[0] = m2 == n
        ? x * m * (alpha2 - m)
          / ((alpha1 + m2) * (alpha1 + m2 - 1))
        : -x * (alpha1 + m) * (alpha1 + alpha2 + m)
          / ((alpha1 + m2) * (alpha1 + m2 + 1));
    return;
}
protected double initialValue()
{
    factors[1] = 1;
    return 1;
}
}
```

Chapter 8

Linear algebra

On ne trouve pas l'espace, il faut toujours le construire.¹
Gaston Bachelard

Linear algebra concerns itself with the manipulation of vectors and matrices. The concepts of linear algebra are not difficult and linear algebra is usually taught in the first year of university. Solving systems of linear equations are even taught in high school. Of course, one must get used to the book keeping of the indices. The concise notation introduced in linear algebra for vector and matrix operations allows expressing difficult problems in a few short equations. This notation can be directly adapted to object oriented programming.

Figure 8.1 shows the classes described in this chapter. Like chapter 2, this chapter discusses some fundamental concepts and operations that shall be used throughout the rest of the book. It might appear austere to many readers because, unlike the preceding chapters, it does not contains concrete examples. However, the reader will find example of use of linear algebra in nearly all remaining chapters of this book.

The chapter begins with a reminder of operations defined on vectors and matrices. Then, two methods for solving systems of linear equations are discussed. This leads to the important concept of matrix inversion. Finally the chapter closes of the problem of finding eigenvalues and eigenvectors.

8.1 Vectors and matrices

Linear algebra concerns itself with vectors in multidimensional spaces and the properties of operations on these vectors. It is a remarkable fact that such properties can be studied without explicit specification of the space dimension².

A vector is an object in a multidimensional space. It is represented by its components measured on a reference system. A reference system is a series of vectors from which

¹Space is not to be found; it must always be constructed.

²In fact, most mathematical properties discussed in this chapter are valid for space with an infinite number of dimensions (Hilbert spaces).

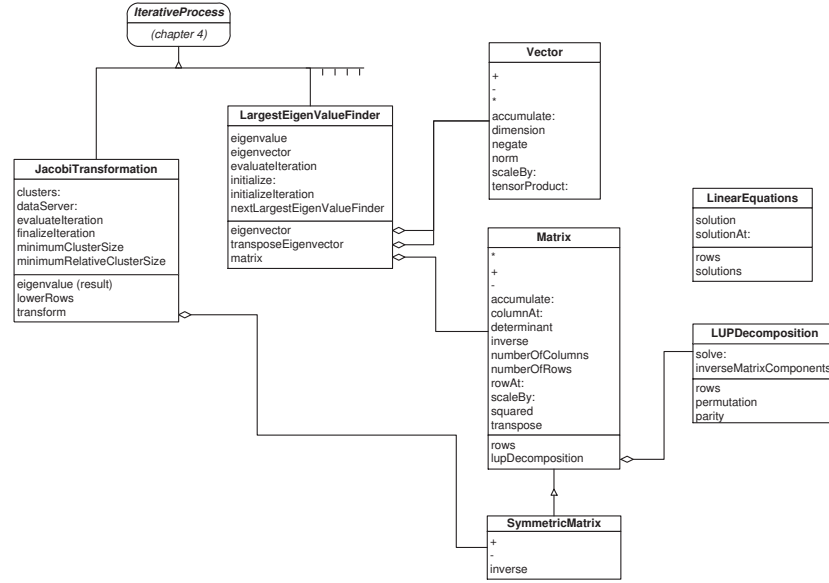


Figure 8.1: Linear algebra classes

the entire space can be generated. A commonly used mathematical notation for a vector is a lower case bold letter, \mathbf{v} for example. If the set of vectors $\mathbf{u}_1, \dots, \mathbf{u}_n$ is a reference system for a space with n dimension, then any vector of the space can be written as:

$$\mathbf{v} = v_1 \mathbf{u}_1 + \dots + v_n \mathbf{u}_n, \quad (8.1)$$

where v_1, \dots, v_n are real numbers in the case of a real space or complex numbers in a complex space. The numbers v_1, \dots, v_n are called the components of the vector.

A matrix is a linear operator over vectors from one space to vectors in another space not necessarily of the same dimension. This means that the application of a matrix on a vector is another vector. To explain what linear means, we must quickly introduce some notation.

A matrix is commonly represented with an upper case bold letter, \mathbf{M} for example. The application of the matrix \mathbf{M} on the vector \mathbf{v} is denoted by $\mathbf{M} \cdot \mathbf{v}$. The fact that a matrix is a linear operator means that

$$\mathbf{M} \cdot (\alpha \mathbf{u} + \beta \mathbf{v}) = \alpha \mathbf{M} \cdot \mathbf{u} + \beta \mathbf{M} \cdot \mathbf{v}, \quad (8.2)$$

for any matrix \mathbf{M} , any vectors \mathbf{u} and \mathbf{v} , any numbers α and β .

Matrices are usually represented using a table of numbers. In general, the number of rows and the number of columns are not the same. A square matrix is a matrix having the same number of rows and columns. A square matrix maps a vector onto another vector of the same space.

Vectors and matrices have an infinite number of representations depending on the choice of reference system. Some properties of matrices are independent from the

reference system. Very often the reference system is not specified explicitly. For example, the vector \mathbf{v} of equation 8.1 is represented by the array of numbers $(v_1 v_2 \cdots v_n)$ where n is the dimension of the vector. Writing the components of a vector within parentheses is customary. Similarly a matrix is represented with a table of numbers called the components of the matrix; the table is also enclosed within parentheses. For example, the n by m matrix \mathbf{A} is represented by:

$$\mathbf{A} = \begin{pmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{pmatrix}. \quad (8.3)$$

The components can be real or complex numbers. In this book we shall deal only with vectors and matrices having real components.

For simplicity a matrix can also be written with matrix components. That is, the n by m matrix \mathbf{A} can be written in the following form:

$$\mathbf{A} = \begin{pmatrix} \mathbf{B} & \mathbf{C} \\ \mathbf{D} & \mathbf{E} \end{pmatrix}, \quad (8.4)$$

where \mathbf{B} , \mathbf{C} , \mathbf{D} and \mathbf{E} are all matrices of lower dimensions. Let \mathbf{B} be a p by q matrix. Then, \mathbf{C} is a p by $m - q$ matrix, \mathbf{D} is a $n - p$ by q matrix and \mathbf{E} is a $n - p$ by $m - q$ matrix.

Using this notation one can carry all conventional matrix operations using formulas similar to those written with the ordinary number components. There is, however, one notable exception: the order of the products must be preserved since matrix multiplication is not commutative. For example, the product between two matrices expressed as matrix components can be carried out as:

$$\begin{pmatrix} \mathbf{B}_1 & \mathbf{C}_1 \\ \mathbf{D}_1 & \mathbf{E}_1 \end{pmatrix} \cdot \begin{pmatrix} \mathbf{B}_2 & \mathbf{C}_2 \\ \mathbf{D}_2 & \mathbf{E}_2 \end{pmatrix} = \begin{pmatrix} \mathbf{B}_1 \cdot \mathbf{B}_2 + \mathbf{C}_1 \cdot \mathbf{D}_2 & \mathbf{B}_1 \cdot \mathbf{C}_2 + \mathbf{C}_1 \cdot \mathbf{E}_2 \\ \mathbf{D}_1 \cdot \mathbf{B}_2 + \mathbf{E}_1 \cdot \mathbf{D}_2 & \mathbf{D}_1 \cdot \mathbf{C}_2 + \mathbf{E}_1 \cdot \mathbf{E}_2 \end{pmatrix}, \quad (8.5)$$

In equation 8.5 the dimension of the respective matrix components must permit the corresponding product. For example the number of rows of the matrices \mathbf{B}_1 and \mathbf{C}_1 must be equal to the number of columns of the matrices \mathbf{B}_2 and \mathbf{C}_2 respectively.

Common operations defined on vectors and matrices are summarized below. In each equation, the left-hand side shows the vector notation and the right-hand side shows the expression for coordinates and components.

The sum of two vectors of dimension n is a vector of dimension n :

$$\mathbf{w} = \mathbf{u} + \mathbf{v} \quad w_i = u_i + v_i \quad \text{for } i = 1, \dots, n \quad (8.6)$$

The product of a vector of dimension n by a number α is a vector of dimension n :

$$\mathbf{w} = \alpha \mathbf{v} \quad w_i = \alpha v_i \quad \text{for } i = 1, \dots, n \quad (8.7)$$

The scalar product of two vectors of dimension n is a number:

$$s = \mathbf{u} \cdot \mathbf{v} \quad s = \sum_{i=1}^n u_i v_i \quad (8.8)$$

The norm of a vector is denoted $|\mathbf{v}|$. The norm is the square root of the scalar product with itself.

$$|\mathbf{v}| = \sqrt{\mathbf{v} \cdot \mathbf{v}} \quad |\mathbf{v}| = \sqrt{\sum_{i=1}^n v_i v_i} \quad (8.9)$$

The tensor product of two vectors of respective dimensions n and m is an n by m matrix:

$$\mathbf{T} = \mathbf{u} \otimes \mathbf{v} \quad T_{ij} = u_i v_j \quad \begin{array}{l} \text{for } i = 1, \dots, n \\ \text{and } j = 1, \dots, m \end{array} \quad (8.10)$$

The sum of two matrices of same dimensions is a matrix of same dimensions:

$$\mathbf{C} = \mathbf{A} + \mathbf{B} \quad c_{ij} = a_{ij} + b_{ij} \quad \begin{array}{l} \text{for } i = 1, \dots, n \\ \text{and } j = 1, \dots, m \end{array} \quad (8.11)$$

The product of a matrix by a number α is a matrix of same dimensions:

$$\mathbf{B} = \alpha \mathbf{A} \quad b_{ij} = \alpha a_{ij} \quad \begin{array}{l} \text{for } i = 1, \dots, n \\ \text{and } j = 1, \dots, m \end{array} \quad (8.12)$$

The transpose of a n by m matrix is a m by n matrix:

$$\mathbf{B} = \mathbf{A}^T \quad b_{ij} = a_{ji} \quad \begin{array}{l} \text{for } i = 1, \dots, n \\ \text{and } j = 1, \dots, m \end{array} \quad (8.13)$$

The product of a n by m matrix with a vector of dimension n is a vector of dimension m :

$$\mathbf{u} = \mathbf{A} \cdot \mathbf{v} \quad u_i = \sum_{j=1}^m a_{ij} v_j \quad \text{for } i = 1, \dots, m \quad (8.14)$$

The transposed product of a vector of dimension m by a n by m matrix is a vector of dimension n :

$$\mathbf{u} = \mathbf{v} \cdot \mathbf{A} \quad u_i = \sum_{j=1}^m a_{ji} v_j \quad \text{for } i = 1, \dots, n \quad (8.15)$$

The product of a n by p matrix with a p by m matrix is a n by m matrix:

$$\mathbf{C} = \mathbf{A} \cdot \mathbf{B} \quad c_{ij} = \sum_{k=1}^p a_{ik} a_{kj} \quad \begin{array}{l} \text{for } i = 1, \dots, n \\ j = 1, \dots, m \end{array} \quad (8.16)$$

There are of course other operations (the outer product for example) but they will not be used in this book.

To conclude this quick introduction, let us mention matrices with special properties.

A square matrix is a matrix which has the same number of rows and columns. To shorten sentences, we shall speak of a square matrix of dimension n instead of a n by n matrix.

An identity matrix \mathbf{I} is a matrix such that

$$\mathbf{I} \cdot \mathbf{v} = \mathbf{v} \quad (8.17)$$

for any vector \mathbf{v} . This implies that the identity matrix is a square matrix. the representation of the identity matrix contains 1 in the diagonal and 0 off the diagonal in any system of reference. For any square matrix \mathbf{A} we have:

$$\mathbf{I} \cdot \mathbf{A} = \mathbf{A} \cdot \mathbf{I} = \mathbf{A} \quad (8.18)$$

One important property for the algorithms discussed in this book is symmetry. A symmetrical matrix is a matrix such that $\mathbf{A}^T = \mathbf{A}$. In any system of reference the components of a symmetric matrix have the following property:

$$a_{ij} = a_{ji}, \quad \text{for all } i \text{ and } j. \quad (8.19)$$

The sum and product of two symmetric matrices is a symmetric matrix. The matrix $\mathbf{A}^T \cdot \mathbf{A}$ is a symmetric matrix for any matrix \mathbf{A} . If the matrix \mathbf{A} represented in equation 8.4 is symmetric, we have $\mathbf{D} = \mathbf{C}^T$.

8.1.1 Vector and matrix — Smalltalk implementation

Listings 8.1 and 8.2 show respectively the implementation of vectors and matrices as Smalltalk classes. A special implementation for symmetric matrices is shown in listing 8.3.

Figure 8.1 with the box Vector and Matrix grayed.

The public interface is designed as to map itself as close as possible to the mathematical definitions. Here are some example of code using operations between vectors and matrices:

Code example 8.1

```
l u v w a b c l
u := #(1 2 3) asVector.
v := #(3 4 5) asVector.
a := DhhMatrix rows: #( ( 1 0 1 ) (-1 -2 3)).
b := DhhMatrix rows: #( ( 1 2 3 ) (-2 1 7) (5 6 7)).
w := 4 * u + (3 * v).
c := a * b.
v := a * u.
w := c transpose * v.
w := v * c.
```

In the first two lines after the declarative statement, the vectors \mathbf{u} and \mathbf{v} are defined from their component array using the creator method `asVector`. They are 3-dimensional vectors. The matrices \mathbf{a} and \mathbf{b} are created by supplying the components to the class creation method `rows:`. The matrix \mathbf{a} is a 2 by 3 matrix, whereas the matrix \mathbf{b} is a

square matrix of dimension 3. In all cases the variable w is assigned to a vector and the variable c is assigned to a matrix. First, the vector w is assigned to a linear combination of the vectors u and v . Apart from the parentheses required for the second product, the expression is identical to what one would write in mathematics (compare this expression with equation 8.1).

Next the matrix c is defined as the product of the matrices a and b in this order. It is a direct transcription of the left part of equation 8.16 up to the case of the operands.

The next assignment redefines the vector v as the product of the matrix A with the vector u . It is now a 2-dimensional vector. Here again the correspondence between the Smalltalk and the mathematical expression is direct.

The last two lines compute the vector w as the transpose product with the matrix a . The result of both line is the same³. The first line makes the transposition of the matrix a explicit, whereas the second line used the implicit definition of the transpose product. The second line is faster than the previous one since no memory assignment is required for the temporary storage of the transpose matrix.

The use of other methods corresponding to the operations defined in equations 8.6 to 8.16 are left as an exercise to the reader.

Implementation

A vector is akin to an instance of the Smalltalk class `Array`, for which mathematical operations have been defined. Thus, a vector in Smalltalk can be implemented directly as a subclass of the class `Array`. A matrix is an object whose instance variable is an array of vectors.

The operations described in the preceding section can be assigned to the corresponding natural operators. The multiplication, however, can involve several types of operands. It can be applied between

1. a vector and a number,
2. a matrix and a number or
3. a vector and a matrix.

Thus, the multiplication will be implemented using double dispatching as explained in section 2.2.1 for operations between polynomials. Double dispatching is described in appendix (c.f. section B.4).

The method `asVector` is defined for compatibility with a similar method defined in the class `Array` to construct a vector out of an array object. A method `asVector` could be defined also for instances of the class `Collection`. This is left as an exercise for the reader.

The method `tensorProduct` returns an instance of a symmetric matrix. This class is defined in listing 8.3.

The method `accumulate` is meant to be used when there is a need to add several vectors. Indeed the following code

³There is a subtle difference between regular vectors and transposed vectors, which is overlooked by our choice of implementation, however. Transposed vectors or covariant vectors as they are called in differential geometry should be implemented in a proper class. This extension is left as an exercise to the reader.

Code example 8.2

```
la b c d el
a := #(1 2 3 4 5) asVector.
b := #(2 3 4 5 6) asVector.
c := #(3 4 5 6 7) asVector.
d := #(4 5 6 7 8) asVector.
e := a+b+c+d.
```

creates a lots of short-lived vectors, namely one for each addition. Using the method `accumulate` reduces the memory allocation:

Code example 8.3

```
la b c d el
a := #(1 2 3 4 5) asVector.
b := #(2 3 4 5 6) asVector.
c := #(3 4 5 6 7) asVector.
d := #(4 5 6 7 8) asVector.
e := a copy.
e accumulate: b; accumulate: c; accumulate: d.
```

If vectors of large dimension are used, using accumulation instead of addition can make a big difference in performance since many large short-lived objects put a heavy toll of the garbage collector.

Listing 8.1 Vector class in Smalltalk

Class **DhbVector**
Subclass of `Array`

Instance methods

* `aNumberOrMatrixOrVector`

`^aNumberOrMatrixOrVector productWithVector: self`

+ `aVector`

```
| answer n |
answer := self class new: self size.
n := 0.
self with: aVector do:
  [ :a :b |
    n := n + 1.
    answer at: n put: ( a + b).
  ].
^answer
```

- aVector

```

| answer n |
answer := self class new: self size.
n := 0.
self with: aVector do:
    [ :a :b |
        n := n + 1.
        answer at: n put: ( a - b).
    ].
^answer

```

accumulate: aVectorOrAnArray

```

1 to: self size do: [ :n | self at: n put: ( ( self at: n) + (
    aVectorOrAnArray at: n))].

```

accumulateNegated: aVectorOrAnArray

```

1 to: self size do: [ :n | self at: n put: ( ( self at: n) - (
    aVectorOrAnArray at: n))].

```

asVector

```

^self

```

dimension

```

^self size

```

negate

```

1 to: self size do: [ :n | self at: n put: (self at: n) negated].

```

norm

```

^(self * self) sqrt

```

normalized

```

^(1 / self norm) * self

```

productWithMatrix: aMatrix

```

^aMatrix rowsCollect: [ :each | each * self]

productWithVector: aVector

| n |
n := 0.
^self inject: 0
    into: [ :sum :each | n := n + 1. (aVector at: n) * each +
                                     sum]

scaleBy: aNumber

1 to: self size do: [ :n | self at: n put: ( ( self at: n) *
                                             aNumber)].

tensorProduct: aVector

self dimension = aVector dimension
    ifFalse:[ ^self error: 'Vector dimensions mismatch to build
                           tensor product' ].
^DhbSymmetricMatrix rows: ( self collect: [ :a | aVector collect:
                                             [ :b | a * b]])

```

The class `DhbMatrix` has two instance variables:

`rows` an array of vectors, each representing a row of the matrix and

`lupDecomposition` a pointer to an object of the class `DhbLUPDecomposition` containing the LUP decomposition of the matrix if already computed. LUP decomposition is discussed in section 8.3.

This implementation reuses the vector implementation of the vector scalar product to make the code as compact as possible. the iterator methods `columnsCollect:`, `columnsDo:`, `rowsCollect:` and `rowsDo:` are designed to limit the need for index management to these methods only.

An attentive reader will have noticed that the iterator methods `rowsDo:` and `rowsCollect:` present a potential breach of encapsulation. Indeed, the following expression

```
aMatrix rowsDo: [ :each | each at: 1 put: 0]
```

changes the matrix representation outside of the normal way. Similarly, the expression

```
aMatrix rowsCollect: [ :each | each]
```

gives direct access to the matrix's internal representation.

The method `square` implements the product of the transpose of a matrix with itself. This construct is used in several algorithms presented in this book.

The reader should compare the Smalltalk code with the Java code. The Java implementation makes the index management explicit. Since there are no iterator methods in Java, there is no other choice.

Note: The presented matrix implementation is straightforward. Depending on the problem to solve, however, it is not the most efficient one. Each multiplication allocates a lot of memory. If the problem is such that one can allocate memory once for all, more efficient methods can be designed.

The implementation of matrix operations — addition, subtraction, product — uses double or multiple dispatching to determine whether or not the result is a symmetric matrix. Double and multiple dispatching are explained in sections B.4 and B.5. The reader who is not familiar with multiple dispatching should trace down a few examples between simple matrices using the debugger.

Listing 8.2 Matrix classes in Smalltalk

Class	DhbMatrix
Subclass of	Object

Instance variable names: `rows` `lupDecomposition`

Class methods

`new: anInteger`

`^self new initialize: anInteger`

`rows: anArrayOfVector`

`^self new initializeRows: anArrayOfVector`

Instance methods

`* aNumberOrMatrixOrVector`

`^aNumberOrMatrixOrVector productWithMatrix: self`

`+ aMatrix`

`^aMatrix addWithRegularMatrix: self`

`- aMatrix`

```
^aMatrix subtractWithRegularMatrix: self
```

```
accumulate: aMatrix
```

```
| n |
n := 0.
self rowsCollect: [ :each | n := n + 1. each accumulate: (
    aMatrix rowAt: n)]
```

```
accumulateNegated: aMatrix
```

```
| n |
n := 0.
self rowsCollect: [ :each | n := n + 1. each accumulateNegated: (
    aMatrix rowAt: n)]
```

```
addWithMatrix: aMatrix class: aMatrixClass
```

```
| n |
n := 0.
^aMatrixClass rows: ( self rowsCollect: [ :each | n := n + 1.
    each + ( aMatrix rowAt: n)])
```

```
addWithRegularMatrix: aMatrix
```

```
^aMatrix addWithMatrix: self class: aMatrix class
```

```
addWithSymmetricMatrix: aMatrix
```

```
^aMatrix addWithMatrix: self class: self class
```

```
asSymmetricMatrix
```

```
^DhbSymmetricMatrix rows: rows
```

```
columnAt: anInteger
```

```
^rows collect: [ :each | each at: anInteger]
```

```
columnsCollect: aBlock
```

```
| n |
n := 0.
^rows last collect: [ :each | n := n + 1. aBlock value: ( self
    columnAt: n)]
```


columnsDo: aBlock

```
| n |
n := 0.
^rows last do: [ :each | n := n + 1. aBlock value: ( self
                                                    columnAt: n)]
```

initialize: anInteger

```
rows := ( 1 to: anInteger) asVector collect: [ :each | DhbVector
                                                    new: anInteger].
```

initializeRows: anArrayOfVector

```
rows := anArrayOfVector asVector collect: [ :each | each
                                                    asVector].
```

isSquare

```
^rows size = rows last size
```

isSymmetric

```
^false
```

lupDecomposition

```
lupDecomposition isNil
  ifTrue: [ lupDecomposition :=DhbLUPDecomposition equations:
                                                    rows].
^lupDecomposition
```

negate

```
rows do: [ :each |each negate].
```

numberOfColumns

```
^rows last size
```

numberOfRows

```
^rows size
```

printOn: aStream

```
| first |  
first := true.  
rows do:  
  [ :each |  
    first ifTrue: [ first := false]  
    ifFalse:[ aStream cr].  
    each printOn: aStream.  
  ].
```

productWithMatrix: aMatrix

```
^self productWithMatrixFinal: aMatrix
```

productWithMatrixFinal: aMatrix

```
^self class rows: ( aMatrix rowsCollect: [ :row | self  
  columnsCollect: [ :col | row * col]])
```

productWithSymmetricMatrix: aSymmetricMatrix

```
^self class rows: ( self rowsCollect: [ :row | aSymmetricMatrix  
  columnsCollect: [ :col | row * col]])
```

productWithTransposeMatrix: aMatrix

```
^self class rows: ( self rowsCollect: [ :row | aMatrix  
  rowsCollect: [ :col | row * col]])
```

productWithVector: aVector

```
^self columnsCollect: [ :each | each * aVector]
```

rowAt: anInteger

```
^rows at: anInteger
```

rowsCollect: aBlock

```
^rows collect: aBlock
```

rowsDo: aBlock

```

^rows do: aBlock

scaleBy: aNumber
    rows do: [ :each | each scaleBy: aNumber].

squared
    ^DhbSymmetricMatrix rows: ( self columnsCollect: [ :col | self
        columnsCollect: [ :colT | col * colT]])

subtractWithMatrix: aMatrix class: aMatrixClass
    | n |
    n := 0.
    ^aMatrixClass rows: ( self rowsCollect: [ :each | n := n + 1.
        each - ( aMatrix rowAt: n)])

subtractWithRegularMatrix: aMatrix
    ^aMatrix subtractWithMatrix: self class: aMatrix class

subtractWithSymmetricMatrix: aMatrix
    ^aMatrix subtractWithMatrix: self class: self class

transpose
    ^self class rows: ( self columnsCollect: [ :each | each])

transposeProductWithMatrix: aMatrix
    ^self class rows: ( self columnsCollect: [ :row | aMatrix
        columnsCollect: [ :col | row * col]])

```

Listing 8.3 shows the implementation of the class `DhbSymmetricMatrix` representing symmetric matrices. A few algorithms are implemented differently for symmetric matrices.

The reader should pay attention to the methods implementing addition, subtraction and products. Triple dispatching is used to ensure that the addition or subtraction of two symmetric matrices yields a symmetric matrix whereas the same operations between a symmetric matrix and a normal matrix yield a normal matrix. Product requires quadruple dispatching.

Listing 8.3 Symmetric matrix classes in Smalltalk

Class **DhbSymmetricMatrix**
Subclass of DhbMatrix

Class methods

identity: anInteger

 ^self new initializeIdentity: anInteger

Instance methods

+ aMatrix

 ^aMatrix addWithSymmetricMatrix: self

- aMatrix

 ^aMatrix subtractWithSymmetricMatrix: self

addWithSymmetricMatrix: aMatrix

 ^aMatrix addWithMatrix: self class: self class

clear

 rows do: [:each | each atAllPut: 0].

initializeIdentity: anInteger

 rows := (1 to: anInteger) asVector collect: [:n | (DhbVector
 new: anInteger) atAllPut: 0; at: n put: 1; yourself].

isSquare

 ^true

isSymmetric

 ^true

productWithMatrix: aMatrix

 ^aMatrix productWithSymmetricMatrix: self

```

productWithSymmetricMatrix: aSymmetricMatrix
    ^aSymmetricMatrix productWithMatrixFinal: self

subtractWithSymmetricMatrix: aMatrix
    ^aMatrix subtractWithMatrix: self class: self class

```

Figure 8.1 with the boxes Vector and Matrix grayed.

8.1.2 Vector and matrix — Java implementation

Listings 8.4 and 8.5 show respectively the implementation of vectors and matrices as Smalltalk classes.

The public interface is designed as to map itself as close as possible to the mathematical definitions. In Java, however, one cannot overload primitive operators such as `+` or `*`. Thus, named methods are used instead. The resulting code is admittedly less readable than its Smalltalk equivalent.

Some operations — a scalar product between two vectors of different dimension e.g.— are not possible. In our implementation, we have introduced a few specific exceptions — `DhbIllegalDimension`, `DhbNonSymmetricComponents` — which must be trapped on each operation.

Here is an example of code using operations between vectors and matrices:

Code example 8.4

```

double[] comp_u = { 1, 2, 3 };
double[] comp_v = { 3, 4, 5 };
double[][] comp_a = { { 1,0,1 }, { -1, -2, 3 } };
double[][] comp_b = { { 1,2,3 }, { -2, -1, 7 }, { 5,6,7 } };
try {
    DhbVector u = new DhbVector(components);
    DhbVector v = new DhbVector(components);
    Matrix a = new Matrix(comp_a);
    Matrix b = new Matrix(comp_b);
    DhbVector w = u.product(4).add(v.product(3));
    Matrix c = a.product(b);
    v = a.product(u);
    w = c.transpose().product(v);
    w = v.product(c);
} catch ( DhbIllegalDimension e ) { };

```

The first four lines define the component arrays used to define the vectors and matrices. Within the `try...catch`, the vectors `u` and `v` are created from their component arrays. They are 3-dimensional vectors. The matrices `a` and `b` are also created from their component arrays. The matrix `a` is a 2 by 3 matrix, whereas the matrix `b` is a square matrix of dimension 3. In all cases the variable `w` is assigned to a vector and the variable

c is assigned to a matrix. First, the vector w is assigned to a linear combination of the vectors u and v .

Next the matrix c is defined as the product of the matrices a and b in this order.

The next assignment redefines the vector v as the product of the matrix A with the vector u . It is now a 2-dimensional vector.

The last two lines compute the vector w as the transpose product with the matrix a . The result of both line is the same⁴. The first line builds the transpose of the matrix a explicitly, whereas the second line used the definition of the transpose product. The second line is faster than the previous one since no memory assignment is required for the temporary storage of the transpose matrix.

The use of named methods — as opposed to using primitive operators — obscures somewhat the readability of the code. For an objective comparison, Java programmers are advised to take a look at the code example in Smalltalk at the beginning of section 8.1.1.

The use of other methods corresponding to the operations defined in equations 8.6 to 8.16 are left as an exercise to the reader.

Since a Java array is not an object, a vector is created as a subclass of object. The name Vector could not be used since it is already part of the base classes of Java. An instance of the Java class Vector, however, has little to do with a mathematical vector. The Java class Vector has been totally misnamed. In fact it corresponds to the Smalltalk class OrderedCollection. An instance of the Java class Vector has a variable size whereas the dimension of a mathematical vector is fixed. In the rest of this book, whenever the word vector is used, it means a mathematical vector and not an instance of the Java class Vector.

Operations such as addition and products cannot be implemented with primitive operators (+, − e.g.) since Java does not allow the definition of primitive operators on objects. Thus, we are forced to use method names such as add or product.

In Java a method is not only defined by its name, but also by the type of its arguments. Thus, the same method name can be used to define implementations of an operator acting on different objects without the need to test the arguments. The Java compiler does that for us; double or triple dispatching is not needed.

Listing 8.4 defines the method tensorProduct using the class SymmetricMatrix that will be defined in section 8.5 (c.f. listing 8.12), as most of the code for symmetric matrices is devoted to matrix inversion.

An exception specific for this package has been created: DhbIllegalDimension. This exception is thrown when an operation is attempted between vectors or matrices with incompatible dimensions. The code for this exception is elementary and is not shown here. It is left as an exercise for the reader.

Vector components are kept in a protected instance variable. Protected means that it is available only to classes of the same package. This allows the optimization of many methods from other classes of the same package needing direct access to the components (Matrix e.g.). Other users of the class have no access to the vector's components. This choice is deliberate: it provides a strict hiding of the implementation. A vector as mathematical object should be manipulated without any access to its components.

⁴c.f. footnote 3 on page 200

The method `toString` is provided to allow printing of the components. The method `toComponents` provides a way to access the coordinates if needed. Note that this method is making a copy of the components, thus preserving the object from any indirect change.

Listing 8.4 Mathematical vector class in Java

```
package DhbMatrixAlgebra;

/**
 * Vector implementation
 *
 * @author Didier H. Besset
 */
public class DhbVector
{
    protected double[] components;

    /**
     * Create a vector of given dimension.
     * NOTE: The supplied array of components must not be changed.
     * @param comp double[]
     */
    public DhbVector( double comp[]) throws NegativeArraySizeException
    {
        int n = comp.length;
        if ( n <= 0 )
            throw new NegativeArraySizeException(
                "Vector components cannot be empty");
        components = new double[n];
        System.arraycopy( comp, 0, components, 0, n);
    }

    /**
     * Create a vector of given dimension.
     * @param dimension int dimension of the vector; must be positive.
     */
    public DhbVector ( int dimension) throws NegativeArraySizeException
    {
        if ( dimension <= 0 )
            throw new NegativeArraySizeException(
                "Requested vector size: "+dimension);
        components = new double[dimension];
        clear();
    }

    /**
     * @param v DHBmatrixAlgebra.DhbVector
     * @exception DHBmatrixAlgebra.DhbIllegalDimension if the vector
```

```

* and supplied vector do not have the same dimension.
*/
public void accumulate ( double[] x) throws DhbIllegalDimension
{
    if ( this.dimension() != x.length )
        throw new DhbIllegalDimension("Attempt to add a "
            +this.dimension()+"-dimension vector to a "
            +x.length+"-dimension array");
    for ( int i = 0; i < this.dimension(); i++)
        components[i] += x[i];
}
/**
 * @param v DHBmatrixAlgebra.DhbVector
 * @exception DHBmatrixAlgebra.DhbIllegalDimension if the vector
 * and supplied vector do not have the same dimension.
 */
public void accumulate ( DhbVector v) throws DhbIllegalDimension
{
    if ( this.dimension() != v.dimension() )
        throw new DhbIllegalDimension("Attempt to add a "
            +this.dimension()+"-dimension vector to a "
            +v.dimension()+"-dimension vector");
    for ( int i = 0; i < this.dimension(); i++)
        components[i] += v.components[i];
}
/**
 * @param v DHBmatrixAlgebra.DhbVector
 * @exception DHBmatrixAlgebra.DhbIllegalDimension if the vector
 * and supplied vector do not have the same dimension.
 */
public void accumulateNegated( double[] x) throws DhbIllegalDimension
{
    if ( this.dimension() != x.length )
        throw new DhbIllegalDimension("Attempt to add a "
            +this.dimension()+"-dimension vector to a "
            +x.length+"-dimension array");
    for ( int i = 0; i < this.dimension(); i++)
        components[i] -= x[i];
}
/**
 * @param v DHBmatrixAlgebra.DhbVector
 * @exception DHBmatrixAlgebra.DhbIllegalDimension if the vector
 * and supplied vector do not have the same dimension.
 */
public void accumulateNegated( DhbVector v) throws DhbIllegalDimension
{

```



```

        if ( this.dimension() != v.dimension() )
            throw new DhbIllegalDimension("Attempt to add a "
                +this.dimension()+"-dimension vector to a "
                +v.dimension()+"-dimension vector");
        for ( int i = 0; i < this.dimension(); i++)
            components[i] -= v.components[i];
    }
    /**
     * @return DHBmatrixAlgebra.DhbVector sum of the vector with
     *         the supplied vector
     * @param v DHBmatrixAlgebra.DhbVector
     * @exception DHBmatrixAlgebra.DhbIllegalDimension if the vector
     *         and supplied vector do not have the same dimension.
     */
    public DhbVector add ( DhbVector v) throws DhbIllegalDimension
    {
        if ( this.dimension() != v.dimension() )
            throw new DhbIllegalDimension("Attempt to add a "
                +this.dimension()+"-dimension vector to a "
                +v.dimension()+"-dimension vector");
        double[] newComponents = new double[ this.dimension()];
        for ( int i = 0; i < this.dimension(); i++)
            newComponents[i] = components[i] + v.components[i];
        return new DhbVector( newComponents);
    }
    /**
     * Sets all components of the receiver to 0.
     */
    public void clear()
    {
        for ( int i = 0; i < components.length; i++) components[i] = 0;
    }
    /**
     * @return double
     * @param n int
     */
    public double component( int n)
    {
        return components[n];
    }
    /**
     * Returns the dimension of the vector.
     * @return int
     */
    public int dimension ( )
    {

```

```

    return components.length;
}
/**
 * @return true if the supplied vector is equal to the receiver
 * @param v DHBmatrixAlgebra.DhbVector
 */
public boolean equals( DhbVector v)
{
    int n = this.dimension();
    if ( v.dimension() != n )
        return false;
    for ( int i = 0; i < n; i++)
    {
        if ( v.components[i] != components[i] )
            return false;
    }
    return true;
}
/**
 * Computes the norm of a vector.
 */
public double norm ( )
{
    double sum = 0;
    for ( int i = 0; i < components.length; i++)
        sum += components[i]*components[i];
    return Math.sqrt( sum);
}
/**
 * @param x double
 */
public DhbVector normalizedBy ( double x )
{
    for ( int i = 0; i < this.dimension(); i++)
        components[i] /= x;
    return this;
}
/**
 * Computes the product of the vector by a number.
 * @return DHBmatrixAlgebra.DhbVector
 * @param d double
 */
public DhbVector product( double d)
{
    double newComponents[] = new double[components.length];
    for ( int i = 0; i < components.length; i++)

```

```

        newComponents[i] = d * components[i];
    return new DhbVector(newComponents);
}
/**
 * Compute the scalar product (or dot product) of two vectors.
 * @return double the scalar product of the receiver with the argument
 * @param v DHBmatrixAlgebra.DhbVector
 * @exception DHBmatrixAlgebra.DhbIllegalDimension if the dimension
 *         of v is not the same.
 */
public double product ( DhbVector v) throws DhbIllegalDimension
{
    int n = v.dimension();
    if ( components.length != n )
        throw new DhbIllegalDimension(
            "Dot product with mismatched dimensions: "
            +components.length+", "+n);
    return secureProduct( v);
}
/**
 * Computes the product of the transposed vector with a matrix
 * @return MatrixAlgebra.DhbVector
 * @param a MatrixAlgebra.Matrix
 */
public DhbVector product ( Matrix a) throws DhbIllegalDimension
{
    int n = a.rows();
    int m = a.columns();
    if ( this.dimension() != n )
        throw new DhbIllegalDimension(
            "Product error: transposed of a "+this.dimension()
            +"-dimension vector cannot be multiplied with a "
            +n +" by "+m+" matrix");
    return secureProduct( a);
}
/**
 * @param x double
 */
public DhbVector scaledBy ( double x )
{
    for ( int i = 0; i < this.dimension(); i++)
        components[i] *= x;
    return this;
}
/**
 * Compute the scalar product (or dot product) of two vectors.

```

```

* No dimension checking is made.
* @return double the scalar product of the receiver with the argument
* @param v DHBmatrixAlgebra.DhbVector
*/
protected double secureProduct ( DhbVector v)
{
    double sum = 0;
    for ( int i = 0; i < v.dimension(); i++)
        sum += components[i]*v.components[i];
    return sum;
}
/**
* Computes the product of the transposed vector with a matrix
* @return MatrixAlgebra.DhbVector
* @param a MatrixAlgebra.Matrix
*/
protected DhbVector secureProduct ( Matrix a)
{
    int n = a.rows();
    int m = a.columns();
    double[] vectorComponents = new double[m];
    for ( int j = 0; j < m; j++ )
    {
        vectorComponents[j] = 0;
        for ( int i = 0; i < n; i++)
            vectorComponents[j] += components[i] * a.components[i][j];
    }
    return new DhbVector( vectorComponents);
}
/**
* @return DHBmatrixAlgebra.DhbVector subtract the supplied vector
* to the receiver
* @param v DHBmatrixAlgebra.DhbVector
* @exception DHBmatrixAlgebra.DhbIllegalDimension if the vector
* and supplied vector do not have the same dimension.
*/
public DhbVector subtract ( DhbVector v) throws DhbIllegalDimension
{
    if ( this.dimension() != v.dimension() )
        throw new DhbIllegalDimension("Attempt to add a "
            +this.dimension()+"-dimension vector to a "
            +v.dimension()+"-dimension vector");
    double[] newComponents = new double[ this.dimension()];
    for ( int i = 0; i < this.dimension(); i++)
        newComponents[i] = components[i] - v.components[i];
    return new DhbVector( newComponents);
}

```

```

    }
    /**
     * @return MatrixAlgebra.Matrix    tensor product with the specified
     *                                vector
     * @param v MatrixAlgebra.DhbVector second vector to build tensor
     *                                product with.
     */
    public Matrix tensorProduct ( DhbVector v)
    {
        int n = dimension();
        int m = v.dimension();
        double [][] newComponents = new double[n][m];
        for ( int i = 0; i < n; i++)
        {
            for ( int j = 0; j < m; j++)
                newComponents[i][j] = components[i] * v.components[j];
        }
        return n == m ? new SymmetricMatrix( newComponents)
            : new Matrix( newComponents);
    }
    /**
     * @return double[]    a copy of the components of the receiver.
     */
    public double[] toComponents ( )
    {
        int n = dimension();
        double[] answer = new double[ n];
        System.arraycopy( components, 0, answer, 0, n);
        return answer;
    }
    /**
     * Returns a string representation of the vector.
     * @return java.lang.String
     */
    public String toString()
    {
        StringBuffer sb = new StringBuffer();
        char[] separator = { '[', ' ' };
        for ( int i = 0; i < components.length; i++)
        {
            sb.append( separator);
            sb.append( components[i]);
            separator[0] = ',';
        }
        sb.append(']');
        return sb.toString();
    }

```

```
}
}
```

Listing 8.5 shows the implementation of matrices as Java classes.

Most of the remarks made for vectors are also valid for matrices. In particular, the components of the matrix are strictly encapsulated within the package using the protected declaration keyword.

The CRL inversion algorithm discussed in section 8.5, uses multiplication with a transpose matrix. Thus, a special method has been introduced to optimize this operation. This avoids allocating a memory block needed to hold the transposed matrix. Since a matrix can be very large, allocating a block of memory can seriously slow down execution, especially when it causes paging.

This implementation of the class `DhbMatrix` uses the fact that a class `SymmetricMatrix` will be introduced to implement the inversion algorithm discussed in section 8.5. Thus, the computation of addition, subtraction and product is made in two steps. First the components are created, then the resulting matrix is created with the appropriate class. In the class `SymmetricMatrix` only the creation of the resulting matrix needs to be implemented.

For further optimization, each matrix operation comes in two flavors: secure and normal. The secure version assumes that the dimensions of the supplied parameters are already checked. These methods are declared as protected since they are reserved for use by other classes of the same package. The normal operations available to any classes first check the arguments for the proper dimensions before carrying the operation using the secure version of the same method.

Unlike the vector implementation, however, strict encapsulation of matrix components is not guaranteed since creation of a matrix from a 2-dimensional array does not copy the components into a new array. This was done to reduce memory usage. Memory consumption does more to slow down program execution than anything else with modern day operating systems. The size of a matrix can become quite significant when working with large dimensions. Thus, a small breach to encapsulation was made in the constructor method. A fully encapsulated version would copy the components as it is done for the constructor method of vectors. The user of the package must be aware of this potential problem.

Listing 8.5 Matrix class in Java

```
package DhbMatrixAlgebra;

/**
 * Class representing matrix
 *
 * @author Didier H. Besset
 */
public class Matrix
{
```

```

    protected double[][] components;
    protected LUPDecomposition lupDecomposition = null;

    /**
     * Creates a matrix with given components.
     * NOTE: the components must not be altered after the definition.
     * @param a double[][]
     */
    public Matrix ( double[][] a)
    {
        components = a;
    }

    /**
     * Creates a null matrix of given dimensions.
     * @param n int    number of rows
     * @param m int    number of columns
     * @exception NegativeArraySizeException
     */
    public Matrix ( int n, int m) throws NegativeArraySizeException
    {
        if ( n <= 0 || m <= 0)
            throw new NegativeArraySizeException(
                "Requested matrix size: "+n+" by "+m);
        components = new double[n][m];
        clear();
    }

    /**
     * @param a MatrixAlgebra.Matrix
     * @exception MatrixAlgebra.DhbIllegalDimension if the supplied matrix
     *         does not have the same dimensions.
     */
    public void accumulate ( Matrix a) throws DhbIllegalDimension
    {
        if ( a.rows() != rows() || a.columns() != columns() )
            throw new DhbIllegalDimension("Operation error: cannot add a"
                +a.rows()+" by "+a.columns()
                +" matrix to a "+rows()+" by "
                +columns()+" matrix");
        int m = components[0].length;
        for ( int i = 0; i < components.length; i++)
        {
            for ( int j = 0; j < m; j++)
                components[i][j] += a.component(i,j);
        }
    }
    /**

```

```

    * @return MatrixAlgebra.Matrix    sum of the receiver with the
    *                                supplied matrix.
    * @param a MatrixAlgebra.Matrix
    * @exception MatrixAlgebra.DhbIllegalDimension if the supplied matrix
    *                                does not have the same dimensions.
    */
public Matrix add ( Matrix a) throws DhbIllegalDimension
{
    if ( a.rows() != rows() || a.columns() != columns() )
        throw new DhbIllegalDimension("Operation error: cannot add a"
            +a.rows()+" by "+a.columns()
            +" matrix to a "+rows()+" by "
            +columns()+" matrix");
    return new Matrix( addComponents( a));
}
/**
 * Computes the components of the sum of the receiver and
 * a supplied matrix.
 * @return double[][]
 * @param a MatrixAlgebra.Matrix
 */
protected double[][] addComponents ( Matrix a)
{
    int n = this.rows();
    int m = this.columns();
    double[][] newComponents = new double[n][m];
    for ( int i = 0; i < n; i++)
    {
        for ( int j = 0; j < m; j++)
            newComponents[i][j] = components[i][j] + a.components[i][j];
    }
    return newComponents;
}
public void clear()
{
    int m = components[0].length;
    for ( int i = 0; i < components.length; i++)
    {
        for ( int j = 0; j < m; j++) components[i][j] = 0;
    }
}
/**
 * @return int    the number of columns of the matrix
 */
public int columns ( )
{

```



```

        return components[0].length;
    }
    /**
     * @return double
     * @param n int
     * @param m int
     */
    public double component( int n, int m)
    {
        return components[n][m];
    }
    /**
     * @return double
     * @exception MatrixAlgebra.DhbIllegalDimension if the supplied
     * matrix is not square.
     */
    public double determinant () throws DhbIllegalDimension
    {
        return lupDecomposition().determinant();
    }
    /**
     * @return true if the supplied matrix is equal to the receiver.
     * @param a MatrixAlgebra.Matrix
     */
    public boolean equals( Matrix a)
    {
        int n = this.rows();
        if ( a.rows() != n )
            return false;
        int m = this.columns();
        if ( a.columns() != m )
            return false;
        for ( int i = 0; i < n; i++)
        {
            for ( int j = 0; j < n; j++)
            {
                if ( a.components[i][j] != components[i][j] )
                    return false;
            }
        }
        return true;
    }
    /**
     * @return DhbMatrixAlgebra.DhbMatrix      inverse of the receiver
     * or pseudoinverse if the receiver is not a square matrix.
     * @exception java.lang.ArithmeticException if the receiver is

```

```

*                               a singular matrix.
*/
public Matrix inverse ( ) throws ArithmeticException
{
    try { return new Matrix(
        lupDecomposition().inverseMatrixComponents());}
    catch ( DhbIllegalDimension e )
    { return new Matrix(
        transposedProduct().inverse()
        .productWithTransposedComponents( this));}
}
/**
 * @return boolean
 */
public boolean isSquare ( )
{
    return rows() == columns();
}
/**
 * @return LUPDecomposition  the LUP decomposition of the receiver.
 * @exception DhbIllegalDimension if the receiver is not
 *                               a square matrix.
 */
protected LUPDecomposition lupDecomposition()
    throws DhbIllegalDimension
{
    if ( lupDecomposition == null )
        lupDecomposition = new LUPDecomposition( this);
    return lupDecomposition;
}
/**
 * @return MatrixAlgebra.Matrix  product of the matrix with
 *                                a supplied number
 * @param a double  multiplicand.
 */
public Matrix product ( double a)
{
    return new Matrix( productComponents( a));
}
/**
 * Computes the product of the matrix with a vector.
 * @return DHBmatrixAlgebra.DhbVector
 * @param v DHBmatrixAlgebra.DhbVector
 */
public DhbVector product ( DhbVector v) throws DhbIllegalDimension
{

```

```

    int n = this.rows();
    int m = this.columns();
    if ( v.dimension() != m )
        throw new DhbIllegalDimension("Product error: "+n+" by "+m
            +" matrix cannot be multiplied with vector of dimension "
            +v.dimension());

    return secureProduct( v);
}
/**
 * @return MatrixAlgebra.Matrix    product of the receiver with the
 *                                 supplied matrix
 * @param a MatrixAlgebra.Matrix
 * @exception MatrixAlgebra.DhbIllegalDimension If the number of
 *        columns of the receiver are not equal to the
 *        number of rows of the supplied matrix.
 */
public Matrix product ( Matrix a) throws DhbIllegalDimension
{
    if ( a.rows() != columns() )
        throw new DhbIllegalDimension(
            "Operation error: cannot multiply a"
            +rows()+" by "+columns()
            +" matrix with a "+a.rows()
            +" by "+a.columns()
            +" matrix");

    return new Matrix( productComponents( a));
}
/**
 * @return double[][]
 * @param a double
 */
protected double[][] productComponents ( double a)
{
    int n = this.rows();
    int m = this.columns();
    double[][] newComponents = new double[n][m];
    for ( int i = 0; i < n; i++)
    {
        for ( int j = 0; j < m; j++)
            newComponents[i][j] = a * components[i][j];
    }
    return newComponents;
}
/**
 * @return double[][]    the components of the product of the receiver
 *                        with the supplied matrix

```

```

    * @param a MatrixAlgebra.Matrix
    */
protected double[][] productComponents ( Matrix a)
{
    int p = this.columns();
    int n = this.rows();
    int m = a.columns();
    double[][] newComponents = new double[n][m];
    for ( int i = 0; i < n; i++)
    {
        for ( int j = 0; j < m; j++)
        {
            double sum = 0;
            for ( int k = 0; k < p; k++)
                sum += components[i][k] * a.components[k][j];
            newComponents[i][j] = sum;
        }
    }
    return newComponents;
}
/**
 * @return MatrixAlgebra.Matrix   product of the receiver with the
 *                                transpose of the supplied matrix
 * @param a MatrixAlgebra.Matrix
 * @exception MatrixAlgebra.DhbIllegalDimension If the number of
 *                                columns of the receiver are not equal to
 *                                the number of columns of the supplied matrix.
 */
public Matrix productWithTransposed ( Matrix a)
                                throws DhbIllegalDimension
{
    if ( a.columns() != columns() )
        throw new DhbIllegalDimension(
            "Operation error: cannot multiply a "+rows()
            +" by "+columns()
            +" matrix with the transpose of a "
            +"a.rows()+" by "+a.columns()
            +" matrix");
    return new Matrix( productWithTransposedComponents( a));
}
/**
 * @return double[][]   the components of the product of the receiver
 *                                with the transpose of the supplied matrix
 * @param a MatrixAlgebra.Matrix
 */
protected double[][] productWithTransposedComponents ( Matrix a)

```

```

{
    int p = this.columns();
    int n = this.rows();
    int m = a.rows();
    double[][] newComponents = new double[n][m];
    for ( int i = 0; i < n; i++)
    {
        for ( int j = 0; j < m; j++)
        {
            double sum = 0;
            for ( int k = 0; k < p; k++)
                sum += components[i][k] * a.components[j][k];
            newComponents[i][j] = sum;
        }
    }
    return newComponents;
}
/**
 * @return int    the number of rows of the matrix
 */
public int rows ( )
{
    return components.length;
}
/**
 * Computes the product of the matrix with a vector.
 * @return DHBmatrixAlgebra.DhbVector
 * @param v DHBmatrixAlgebra.DhbVector
 */
protected DhbVector secureProduct ( DhbVector v)
{
    int n = this.rows();
    int m = this.columns();
    double[] vectorComponents = new double[n];
    for ( int i = 0; i < n; i++ )
    {
        vectorComponents[i] = 0;
        for ( int j = 0; j < m; j++)
            vectorComponents[i] += components[i][j] * v.components[j];
    }
    return new DhbVector( vectorComponents);
}
/**
 * Same as product(Matrix a), but without dimension checking.
 * @return MatrixAlgebra.Matrix    product of the receiver with the
 *                                  supplied matrix

```

```

    * @param a MatrixAlgebra.Matrix
    */
protected Matrix secureProduct ( Matrix a)
{
    return new Matrix( productComponents( a));
}
/**
 * Same as subtract ( DhhMatrix a), but without dimension checking.
 * @return MatrixAlgebra.Matrix
 * @param a MatrixAlgebra.Matrix
 */
protected Matrix secureSubtract ( Matrix a)
{
    return new Matrix( subtractComponents( a));
}
/**
 * @return MatrixAlgebra.Matrix    subtract the supplied matrix to
 *                                  the receiver.
 * @param a MatrixAlgebra.Matrix
 * @exception MatrixAlgebra.DhhIllegalDimension if the supplied matrix
 *                                  does not have the same dimensions.
 */
public Matrix subtract ( Matrix a) throws DhhIllegalDimension
{
    if ( a.rows() != rows() || a.columns() != columns() )
        throw new DhhIllegalDimension(
            "Product error: cannot subtract a"+a.rows()
            +" by "+a.columns()+" matrix to a "
            +rows()+" by "+columns()+" matrix");
    return new Matrix( subtractComponents( a));
}
/**
 * @return double[][]
 * @param a MatrixAlgebra.Matrix
 */
protected double[][] subtractComponents ( Matrix a)
{
    int n = this.rows();
    int m = this.columns();
    double[][] newComponents = new double[n][m];
    for ( int i = 0; i < n; i++)
    {
        for ( int j = 0; j < m; j++)
            newComponents[i][j] = components[i][j] - a.components[i][j];
    }
    return newComponents;
}

```

```

    }
    /**
     * @return double[][] a copy of the components of the receiver.
     */
    public double[][] toComponents ( )
    {
        int n = rows();
        int m = columns();
        double[][] answer = new double[ n][m];
        for ( int i = 0; i < n; i++ )
        {
            for ( int j = 0; j < m; j++ )
                answer[i][j] = components[i][j];
        }
        return answer;
    }
    /**
     * Returns a string representation of the system.
     * @return java.lang.String
     */
    public String toString()
    {
        StringBuffer sb = new StringBuffer();
        char[] separator = { '[', ' ' };
        int n = rows();
        int m = columns();
        for ( int i = 0; i < n; i++ )
        {
            separator[0] = '{';
            for ( int j = 0; j < m; j++ )
            {
                sb.append( separator );
                sb.append( components[i][j] );
                separator[0] = ' ';
            }
            sb.append('}');
            sb.append('\n');
        }
        return sb.toString();
    }
    /**
     * @return MatrixAlgebra.Matrix transpose of the receiver
     */
    public Matrix transpose ( )
    {
        int n = rows();

```

```

    int m = columns();
    double[][] newComponents = new double[m][n];
    for ( int i = 0; i < n; i++)
    {
        for( int j = 0; j < m; j++)
            newComponents[j][i] = components[i][j];
    }
    return new Matrix( newComponents);
}
/**
 * @return DhbMatrixAlgebra.SymmetricMatrix the transposed product
 * of the receiver with itself.
 */
public SymmetricMatrix transposedProduct()
{
    return new SymmetricMatrix( transposedProductComponents( this));
}
/**
 * @return MatrixAlgebra.Matrix product of the tranpose of the
 * receiver with the supplied matrix
 * @param a MatrixAlgebra.Matrix
 * @exception MatrixAlgebra.DhbIllegalDimension If the number of rows
 * of the receiver are not equal to
 * the number of rows of the supplied matrix.
 */
public Matrix transposedProduct ( Matrix a) throws DhbIllegalDimension
{
    if ( a.rows() != rows() )
        throw new DhbIllegalDimension(
            "Operation error: cannot multiply a tranposed "
            +rows()+" by "+columns()
            +" matrix with a "+a.rows()+" by "
            +a.columns()+" matrix");
    return new Matrix( transposedProductComponents( a));
}
/**
 * @return double[][] the components of the product of the
 * transpose of the receiver
 * with the supplied matrix.
 * @param a MatrixAlgebra.Matrix
 */
protected double[][] transposedProductComponents ( Matrix a)
{
    int p = this.rows();
    int n = this.columns();
    int m = a.columns();

```



```

double[][] newComponents = new double[n][m];
for ( int i = 0; i < n; i++)
{
    for ( int j = 0; j < m; j++)
    {
        double sum = 0;
        for ( int k = 0; k < p; k++)
            sum += components[k][i] * a.components[k][j];
        newComponents[i][j] = sum;
    }
}
return newComponents;
}
}

```

8.2 Linear equations

A linear equation is an equation in which the unknowns appear to the first order and are combined with the other unknowns only with addition or subtraction. For example, the following equation:

$$3x_1 - 2x_2 + 4x_3 = 0, \quad (8.20)$$

is a linear equation for the unknowns x_1 , x_2 and x_3 . The following equation

$$3x_1^2 - 2x_2 + 4x_3 - 2x_2x_3 = 0, \quad (8.21)$$

is not linear because x_1 appears as a second order term and a term containing the product of the unknowns x_2 and x_3 is present. However, equation 8.21 is linear for the unknown x_2 (or x_3) alone. A system of linear equation has the same number of equations as there are unknowns. For example

$$\begin{cases} 3x_1 + 2y_2 + 4z_3 &= 16 \\ 2x_1 - 5y_2 - z_3 &= 6 \\ x_1 - 2y_2 - 2z_3 &= 10 \end{cases} \quad (8.22)$$

is a system of linear equation which can be solved for the 3 unknowns x_1 , x_2 and x_3 . Its solution is $x_1 = 2$, $x_2 = -1$ and $x_3 = 3$.

A system of linear equations can be written in vector notation as follows:

$$\mathbf{A} \cdot \mathbf{x} = \mathbf{y}. \quad (8.23)$$

The matrix \mathbf{A} and the vector \mathbf{z} are given. Solving the system of equations is looking for a vector \mathbf{x} such that equation 8.23 holds. The vector \mathbf{x} is the solution of the system. A necessary condition for a unique solution to exist is that the matrix \mathbf{A} be a square

matrix. Thus, we shall only concentrate on square matrices⁵. A sufficient condition for the existence of a unique solution is that the rank of the matrix — that is the number of linearly independent rows — is equal to the dimension of the matrix. If the conditions are all fulfilled, the solution to equation 8.23 can be written in vector notation:

$$\mathbf{x} = \mathbf{A}^{-1}\mathbf{y}. \quad (8.24)$$

where \mathbf{A}^{-1} denotes the inverse of the matrix \mathbf{A} (c.f. section 8.5).

Computing the inverse of a matrix in the general case is numerically quite demanding (c.f. section 8.5). Fortunately, there is no need to explicitly compute the inverse of a matrix to solve a system of equations. Let us first assume that the matrix of the system is a triangular matrix, that is we have:

$$\mathbf{T}\mathbf{x} = \mathbf{y}'. \quad (8.25)$$

where \mathbf{T} is a matrix such that:

$$T_{ij} = 0 \quad \text{for } i > j. \quad (8.26)$$

Backward substitution

The solution of the system of equation 8.25 can be obtained using backward substitution. The name backward comes from the fact that the solution begins by calculating the component with the highest index; then it works its way backward on the index calculating each components using all components with higher index.

$$\begin{cases} x_n = \frac{y'_n}{t_{nn}}, \\ x_i = \frac{y'_i - \sum_{j=i+1}^n t_{ij}x_j}{a'_{nn}} \quad \text{for } i = n-1, \dots, 1. \end{cases} \quad (8.27) \quad \Leftarrow \text{Main equation}$$

Gaussian elimination

Any system as described by equation 8.23 can be transformed into a system based on a triangular matrix. This can be achieved through a series of transformations leaving the solution of the system of linear equations invariant. Let us first rewrite the system under the form of a single matrix \mathbf{S} defined as follows:

$$\mathbf{S} = \begin{pmatrix} a_{11} & a_{12} & \dots & a_{1n} & y_1 \\ a_{21} & a_{22} & \dots & a_{2n} & y_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nn} & y_n \end{pmatrix}. \quad (8.28)$$

⁵It is possible to solve system of linear equations defined with a non-square matrix using technique known as singular value decomposition (SVD). In this case, however, the solution of the system is not a unique vector, but a subspace of n -dimensional space where n is the number of columns of the system's matrix. The SVD technique is similar to the techniques used to find eigenvalues and eigenvectors.

Among all transformations leaving the solution of the system represented by the matrix \mathbf{S} invariant, there are two transformations, which help to obtain a system corresponding to triangular matrix. First, any row of the matrix \mathbf{S} can be exchanged. Second, a row of the matrix \mathbf{S} can be replaced by a linear combination⁶ of that row with another one. The trick is to replace all rows of the matrix \mathbf{S} , except for the first one, with rows having a vanishing first coefficient.

If $a_{11} = 0$, we permute the first row with row i such that $a_{i1} \neq 0$. Then, we replace each row j , where $j > 1$, by itself subtracted with the first row multiplied by a_{i1}/a_{11} . This yields a new system matrix \mathbf{S}' of the form:

$$\mathbf{S}' = \begin{pmatrix} a'_{11} & a'_{12} & \cdots & a'_{1n} & y'_1 \\ 0 & a'_{22} & \cdots & a'_{2n} & y'_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & a'_{n2} & \cdots & a'_{nn} & y'_n \end{pmatrix}. \quad (8.29)$$

This step is called pivoting the system on row 1 and the element a_{11} after the permutation is called the pivot. By pivoting the system on the subsequent rows, we shall obtain a system built on a triangular matrix as follows:

$$\mathbf{S}^{(n)} = \begin{pmatrix} a_{11}^{(n)} & a_{12}^{(n)} & a_{13}^{(n)} & \cdots & a_{1n}^{(n)} & y_1^{(n)} \\ 0 & a_{22}^{(n)} & a_{23}^{(n)} & \cdots & a_{2n}^{(n)} & y_2^{(n)} \\ 0 & 0 & a_{33}^{(n)} & \cdots & a_{3n}^{(n)} & y_3^{(n)} \\ \vdots & \vdots & & \ddots & \vdots & \vdots \\ 0 & 0 & \cdots & 0 & a_{nn}^{(n)} & y_n^{(n)} \end{pmatrix}. \quad (8.30)$$

This algorithm is called Gaussian elimination. Gaussian elimination will fail if we are unable to find a row with a non-zero pivot at one of the steps. In that case the system does not have a unique solution.

The first n columns of the matrix $\mathbf{S}^{(n)}$ can be identified to the matrix \mathbf{T} of equation 8.25 and the last column of the matrix $\mathbf{S}^{(n)}$ corresponds to the vector \mathbf{y}' of equation 8.25. Thus, the final system can be solved with backward substitution.

Note: The reader can easily understand that one could have made a transformation to obtain a triangular matrix with the elements above the diagonal all zero. In this case, the final step is called forward substitution since the first component of the solution vector is computed first. The two approaches are fully equivalent.

Fine points

A efficient way to avoid a null pivot is to systematically look for the row having the largest pivot at each step. To be precise, before pivoting row i , it is first exchanged with row j such that $|a_{ij}^{(i-1)}| > |a_{ik}^{(i-1)}|$ for all $k = i, \dots, n$ if such row exists. The systematic search for the largest pivot ensures optimal numerical precision [Phillips & Taylor].

⁶In such linear combination, the coefficient of the replaced row must not be zero.

The reader will notice that all operations can be performed in place since the original matrix \mathbf{S} is not needed to compute the final result.

Finally it should be noted that the pivoting step can be performed on several vectors \mathbf{y} at the same time. If one must solve the same system of equations for several sets of constants, pivoting can be applied to all constant vectors by extending the matrix \mathbf{S} with as many columns as there are additional constant vectors as follows:

$$\mathbf{S} = \begin{pmatrix} a_{11} & a_{12} & \dots & a_{1n} & y_{11} & \dots & y_{m1} \\ a_{21} & a_{22} & \dots & a_{2n} & y_{12} & \dots & y_{m2} \\ \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nn} & y_{1n} & \dots & y_{mn} \end{pmatrix}. \quad (8.31)$$

backward substitution must of course be evaluated separately for each constant vector.

Gaussian elimination is solely dedicated to solving systems of linear equations. The algorithm is somewhat slower than LUP decomposition described in section 8.3. When applied to systems with several constant vectors, however, Gaussian elimination is faster since the elimination steps are made only once. In the case of LUP decomposition, obtaining the solution for each vector requires more operations than those needed by backward substitution.

8.2.1 Linear equations — General implementation

Although using matrix and vector notation greatly simplifies the discussion of Gaussian elimination, there is little gain in making an implementation using matrices and vectors explicitly.

The class creation methods or constructors will take as arguments either a matrix and a vector or an array of arrays and an array. The class implementing Gaussian elimination has the following instance variables:

`rows` an array or a vector whose elements contain the rows of the matrix \mathbf{S} .

`solutions` an array whose elements contain the solutions of the system corresponding to each constant vector.

Solving the system is entirely triggered by retrieving the solutions. The instance variable `solutions` is used to keep whether or not Gaussian elimination has already taken place. If this variable is `nil` Gaussian elimination has not yet been performed. Gaussian elimination is performed by the method `solve`. At the end of the algorithm, the vector of solutions is allocated into the instance variable `solutions`. Similarly, backward substitution is triggered by the retrieving of a solution vector. If the solution for the specified index has not yet been computed, backward substitution is performed and the result is stored in the solution array.

8.2.2 Linear equations — Smalltalk implementation

Listing 8.6 shows the class `DhbLinearEquationSystem` implementing Gaussian elimination in Smalltalk.

Figure 8.1 with the box `LinearEquations` grayed.

To solve the system of equations 8.22 using Gaussian elimination, one needs to write to evaluate the following expression:

Code example 8.5

```
( DhbLinearEquationSystem equations: #( (3 2 4)
                                         (2 -5 -1)
                                         (1 -2 2))
  constant: #(16 6 10)
) solution.
```

This expression has been written on three lines to delineate the various steps. The first two lines create an instance of the class `DhbLinearEquationSystem` by feeding the coefficients of the equations, rows by rows, on the first line and giving the constant vector on the second line. The last line is a call to the method `solution` retrieving the solution of the system.

Solving the same system with an additional constant vector requires a little more code, but not much:

Code example 8.6

```
| s sol1 sol2|
s := DhbLinearEquationSystem equations: #( (3 2 4) (2 -5 -1) (1 -2 2))
      constants: #( (16 6 10)
                    (7 10 9)).

sol1 := s solutionAt: 1.
sol2 := s solutionAt: 2.
```

In this case, the creation method differs in that the two constant vectors are supplied in an array columns by columns. Similarly, the two solutions must be fetched one after the other.

The class `DhbLinearEquationSystem` The class method `equations:constants:` allows to create a new instance for a given matrix and a series of constant vectors.

The method `solutionAt:` returns the solution for a given constant vector. The index specified as argument to that method corresponds to that of the desired constant vector.

The method `solve` performs all required pivoting steps using a `do:` iterator. The method `pivotStepAt:` first swaps rows to bring a pivot having the largest absolute value in place and then invokes the method `pivotAt:` to perform the actual pivoting.

Convenience methods `equations:constant:` and `solution` are supplied to treat the most frequent case where there is only one single constant vector. However, the reader should be reminded that LUP decomposition is more effective in this case.

If the system does not have a solution — that is, if the system's matrix is singular — an arithmetic error occurs in the method `pivotAt:` when the division with the zero pivot is performed. The method `solutionAt:` traps this error within an exception handling structure and sets the solution vector to a special value — the integer 0 — as a flag to prevent attempting Gaussian elimination a second time. Then, the value `nil` is returned to represent the non-existent solution.

Listing 8.6 Smalltalk implementation of a system of linear equations

```
Class                DhbLinearEquationSystem
Subclass of         Object
Instance variable names: rows solutions
Pool dictionaries:   SystemExceptions
```

Class methods

```
equations: anArrayOfArrays constant: anArray
```

```
OfArrays constant: anArray
```

```
^self new initialize: anArrayOfArrays constants: (Array with:
                                     anArray)
```

```
equations: anArrayOfArrays constants: anArrayOfConstantArrays
```

```
^self new initialize: anArrayOfArrays constants:
                                     anArrayOfConstantArrays
```

Instance methods

```
backSubstitutionAt: anInteger
```

```
| size answer accumulator |
size := rows size.
answer := Array new: size.
size to: 1 by: -1 do:
    [ :n |
        accumulator := (rows at: n) at: (anInteger + size).
        ( n + 1) to: size
        do: [ :m | accumulator := accumulator - ( ( answer at: m)
                                                    * ( ( rows at: n) at: m))].
        answer at: n put: ( accumulator / ( ( rows at: n) at: n)).
    ].
solutions at: anInteger put: answer.
```

```
initialize: anArrayOfArrays constants: anArrayOfConstantArrays
```

```
| n |
n := 0.
rows := anArrayOfArrays collect: [ :each | n := n + 1. each, (
    anArrayOfConstantArrays collect: [ :c | c at: n])].
^self
```

```
largestPivotFrom: anInteger
```

```

| valueOfMaximum indexOfMaximum |
valueOfMaximum := ( rows at: anInteger) at: anInteger.
indexOfMaximum := anInteger.
(anInteger + 2) to: rows size do:
    [ :n |
        (( rows at: n) at: anInteger) > valueOfMaximum
            ifTrue: [ valueOfMaximum := ( rows at: n) at:
                    anInteger.
                    indexOfMaximum := n.
                ].
    ].
^indexOfMaximum

```

pivotAt: anInteger

```

| inversePivot rowPivotValue row pivotRow |
pivotRow := rows at: anInteger.
inversePivot := 1 / ( pivotRow at: anInteger).
(anInteger + 1) to: rows size do:
    [ :n |
        row := rows at: n.
        rowPivotValue := ( row at: anInteger) * inversePivot.
        anInteger to: row size do:
            [ :m |
                row at: m put: (( row at: m) - (( pivotRow at: m) *
                    rowPivotValue)).
            ].
    ].

```

pivotStepAt: anInteger

```

self swapRow: anInteger withRow: ( self largestPivotFrom:
    anInteger);
    pivotAt: anInteger.

```

printOn: aStream

```

| first delimitingString n k |
n := rows size.
first := true.
rows do:
    [ :row |
        first ifTrue: [ first := false]
            ifFalse: [ aStream cr].
        delimitingString := '('.

```

```

k := 0.
row do:
  [ :each |
    aStream nextPutAll: delimitingString.
    each printOn: aStream.
    k := k + 1.
    delimitingString := k < n ifTrue: [ ' ' ] ifFalse: [ '
      : ' ].
  ].
aStream nextPut: $).
].

```

solution

```

^self solutionAt: 1

```

solutionAt: anInteger

```

solutions isNil
  ifTrue: [ [self solve] when: ExError do: [ :signal | solutions
    := 0. signal exitWith: nil. ] ].
solutions = 0
  ifTrue: [ ^nil ].
( solutions at: anInteger ) isNil
  ifTrue: [ self backSubstitutionAt: anInteger ].
^solutions at: anInteger

```

solve

```

1 to: rows size do: [ :n | self pivotStepAt: n ].
solutions := Array new: ( (rows at: 1) size - rows size ).

```

swapRow: anInteger1 withRow: anInteger2

```

| swappedRow |
anInteger1 = anInteger2
  ifFalse: [ swappedRow := rows at: anInteger1.
    rows at: anInteger1 put: ( rows at:
      anInteger2 ).
    rows at: anInteger2 put: swappedRow.
  ].

```

8.2.3 Linear equations — Java implementation

Listing 8.7 shows the class `LinearEquations` implementing Gaussian elimination in Java.

To solve the system of equations 8.22, one needs to write to evaluate the following expression:

Code example 8.7

```
double[][] s = {{3,2,4}, {2, -5, -1},{1,-2,2}};
double[] c = { 16,6,10};
try{
    LinearEquations system = new LinearEquations( s, c);
    double[] solution = system.solution();
        <Processing of the solution>

    } catch(DhbIllegalDimension e){ }
```

The first two lines define the components of the components of the system's matrix (**A**) and the components of the constant vector (**y**). The `try...catch` block delineates the solution of the system proper. The first line of the block creates an instance of the class `LinearEquations` by feeding the coefficients of the equations, rows by rows, and the constant vector. The second line is a call to the method `solution` retrieving the solution of the system.

Solving the same system with an additional constant vector requires an additional line of code:

Code example 8.8

```
double[][] s = {{3,2,4}, {2, -5, -1},{1,-2,2}};
double[][] c = {{16,6,10}, {7,10,9}};
try{
    LinearEquations system = new LinearEquations( s, c);
    DhbVector solution1 = system.solution(1);
    DhbVector solution2 = system.solution(2);
        <Processing of the solutions>

    } catch(DhbIllegalDimension e){ }
```

In this case, the creation method differs in that the two constant vectors are supplied in an array columns by columns. Similarly, the two solutions must be fetched one after the other.

The class `LinearEquationSystem` has the following instance variables:

components an array whose elements contain the rows of the matrix **S**.

solutions a vector (an instance of the class `DhbVector`) whose elements contain the solutions of the system corresponding to each constant vector.

The class `LinearEquations` has 3 constructor methods. One uses the components of the matrix and the components of single constant vectors. However, the reader should

be reminded that LUP decomposition is more effective in this case. The next uses an arbitrary number of constant vectors. Finally a constructor with matrix and vector is supplied since the system's matrix and constant vector are often obtained after matrix and vector computations.

The solution to the system is always returned as a mathematical vector, that is, an object of class `DhbVector`. If the solutions were kept as array, an additional instance variable would be required to flag whether a solution was computed. Individual components, if needed, can be extracted from the vector using the method `toComponents`.

There are two methods named `solution`. The one with an integer argument returns the solution for a given constant vector. The index specified as argument to that method corresponds to that of the desired constant vector. The one without argument is a convenience method to retrieve the solution of a system with a single constant vector.

The method `pivotingStep` first swaps rows to bring a pivot having the largest absolute value in place and then invokes the method `pivot` to perform the actual pivoting.

If the system does not have a solution — that is, if the system's matrix is singular — an `ArithmeticException` is thrown. This exception must be caught by the calling object.

Listing 8.7 Java implementation of a system of linear equations

```
package DhbMatrixAlgebra;

/**
 * Class representing a system of linear equations.
 *
 * @author Didier H. Besset
 */
public class LinearEquations
{
    /**
     * components is a matrix build from the system's matrix and
     * the constant vector
     */
    private double[][] rows;
    /**
     * Array containing the solution vectors.
     */
    private DhbVector[] solutions;

    /**
     * Construct a system of linear equation  $Ax = y_1, y_2, \dots$ 
     * @param m double[][]
     * @param c double[][]
     * @exception DhbMatrixAlgebra.DhbIllegalDimension
     *             if the system's matrix is not square
     *             if constant dimension does not match
     *             that of the matrix
     */
}
```

```

*/
public LinearEquations ( double[][] m, double[][] c)
    throws DhbIllegalDimension
{
    int n = m.length;
    if ( m[0].length != n )
        throw new DhbIllegalDimension("Illegal system: a"+n+" by "
            +m[0].length+" matrix is not a square matrix");
    if ( c[0].length != n )
        throw new DhbIllegalDimension("Illegal system: a "+n+" by "+n
            +" matrix cannot build a system with a "
            +c[0].length+"-dimensional vector");
    rows = new double[n][n+c.length];
    for ( int i = 0; i < n; i++)
    {
        for ( int j = 0; j < n; j++)
            rows[i][j] = m[i][j];
        for ( int j = 0; j < c.length; j++)
            rows[i][n+j] = c[j][i];
    }
}
/**
 * Construct a system of linear equation  $Ax = y$ .
 * @param m double[][] components of the system's matrix
 * @param c double[] components of the constant vector
 * @exception DhbMatrixAlgebra.DhbIllegalDimension
 *             if the system's matrix is not square
 *             if constant dimension does not match
 *             that of the matrix
 */
public LinearEquations ( double[][] m, double[] c)
    throws DhbIllegalDimension
{
    int n = m.length;
    if ( m[0].length != n )
        throw new DhbIllegalDimension("Illegal system: a"+n+" by "
            +m[0].length+" matrix is not a square matrix");
    if ( c.length != n )
        throw new DhbIllegalDimension("Illegal system: a "+n+" by "+n
            +" matrix cannot build a system with a "
            +c.length+"-dimensional vector");
    rows = new double[n][n+1];
    for ( int i = 0; i < n; i++)
    {
        for ( int j = 0; j < n; j++)
            rows[i][j] = m[i][j];
    }
}

```

```

        rows[i][n] = c[i];
    }
}
/**
 * Construct a system of linear equation  $Ax = y$ .
 * @param a MatrixAlgebra.Matrix    matrix A
 * @param y MatrixAlgebra.DhbVector  vector y
 * @exception MatrixAlgebra.DhbIllegalDimension
 *           if the system's matrix is not square
 *           if vector dimension does not match
 *           that of the matrix
 */
public LinearEquations ( Matrix a, DhbVector y)
                        throws DhbIllegalDimension
{
    this( a.components, y.components);
}
/**
 * Computes the solution for constant vector p applying
 * backsubstitution.
 * @param p int
 * @exception java.lang.ArithmeticException if one diagonal element
 *           of the triangle matrix is zero.
 */
private void backSubstitution ( int p) throws ArithmeticException
{
    int n = rows.length;
    double [] answer = new double[n];
    double x;
    for ( int i = n - 1; i >= 0; i--)
    {
        x = rows[i][n+p];
        for ( int j = i + 1; j < n; j++)
            x -= answer[j] * rows[i][j];
        answer[i] = x / rows[i][i];
    }
    solutions[p] = new DhbVector( answer);
    return;
}
/**
 * Finds the position of the largest pivot at step p.
 * @return int
 * @param p int    step of pivoting.
 */
private int largestPivot ( int p)
{

```

```

    double pivot = Math.abs( rows[p][p]);
    int answer = p;
    double x;
    for ( int i = p + 1; i < rows.length; i++)
    {
        x = Math.abs(rows[i][p]);
        if ( x > pivot )
        {
            answer = i;
            pivot = x;
        }
    }
    return answer;
}
/**
 * Perform pivot operation at location p.
 * @param p int
 * @exception java.lang.ArithmeticException if the pivot element
 *                                     is zero.
 */
private void pivot ( int p) throws ArithmeticException
{
    double inversePivot = 1 / rows[p][p];
    double r;
    int n = rows.length;
    int m = rows[0].length;
    for ( int i = p + 1; i < n; i++)
    {
        r = inversePivot * rows[i][p];
        for ( int j = p; j < m; j++)
            rows[i][j] -= rows[p][j] * r;
    }
    return;
}
/**
 * Perform optimum pivot operation at location p.
 * @param p int
 */
private void pivotingStep ( int p)
{
    swapRows( p, largestPivot( p));
    pivot(p);
    return;
}
/**
 * @return DhbVector    solution for the 1st constant vector

```

```

*/
public DhbVector solution ( ) throws ArithmeticException
{
    return solution( 0);
}
/**
 * Return the vector solution of constants indexed by p.
 * @return DHBmatrixAlgebra.DhbVector
 * @param p int    index of the constant vector fed into the system.
 * @exception java.lang.ArithmeticException
 *                if the system cannot be solved.
 */
public DhbVector solution ( int p) throws ArithmeticException
{
    if ( solutions == null )
        solve();
    if ( solutions[p] == null )
        backSubstitution( p);
    return solutions[p];
}
/**
 * @exception java.lang.ArithmeticException
 *                if the system cannot be solved.
 */
private void solve ( ) throws ArithmeticException
{
    int n = rows.length;
    for ( int i = 0; i < n; i++)
        pivotingStep( i);
    solutions = new DhbVector[rows[0].length-n];
}
/**
 * Swaps rows p and q.
 * @param p int
 * @param q int
 */
private void swapRows ( int p, int q)
{
    if ( p != q)
    {
        double temp;
        int m = rows[p].length;
        for (int j = 0; j < m; j++)
        {
            temp = rows[p][j];
            rows[p][j] = rows[q][j];

```

```

        rows[q][j] = temp;
    }
}
return;
}
/**
 * Returns a string representation of the system.
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    char[] separator = { '[', ' ' };
    int n = rows.length;
    int m = rows[0].length;
    for ( int i = 0; i < n; i++)
    {
        separator[0] = '(';
        for ( int j = 0; j < n; j++)
        {
            sb.append( separator);
            sb.append( rows[i][j]);
            separator[0] = ',';
        }
        separator[0] = ':';
        for ( int j = n; j < m; j++)
        {
            sb.append( separator);
            sb.append( rows[i][j]);
            separator[0] = ',';
        }
        sb.append(')');
        sb.append('\n');
    }
    return sb.toString();
}
}

```

8.3 LUP decomposition

LUP decomposition is another technique to solve a system of linear equations. It is an alternative to the Gaussian elimination [Cormen et al.]. Gaussian elimination can solve a system with several constant vectors, but all constant vectors must be known before starting the algorithm.

LUP decomposition is done once for the matrix of a given system. Thus, the system can be solved for any constant vector obtained after the LUP decomposition. In addition, LUP decomposition gives a way to calculate the determinant of a matrix and it can be used to compute the inverse of a matrix.

LUP stands for Lower, Upper and Permutation. It comes from the observation that any non-singular square matrix \mathbf{A} can be decomposed into a product of 3 square matrices of the same dimension as follows:

$$\mathbf{A} = \mathbf{L} \cdot \mathbf{U} \cdot \mathbf{P}, \quad (8.32)$$

where \mathbf{L} is a matrix whose components located above the diagonal are zero (lower triangular matrix), \mathbf{U} is a matrix whose components located below the diagonal are zero (upper triangular matrix) and \mathbf{P} is a permutation matrix. The decomposition of equation 8.32 is non-unique. One can select a unique decomposition by requiring that all diagonal elements of the matrix \mathbf{L} be equal to 1.

The proof that such decomposition exists is the algorithm itself. It is a proof by recursion. We shall first start to construct an LU decomposition, that is an LUP decomposition with an identity permutation matrix. Let us write the matrix as follows:

$$\mathbf{A} = \begin{pmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} \end{pmatrix} = \begin{pmatrix} a_{11} & \mathbf{w}^T \\ \mathbf{v} & \mathbf{A}' \end{pmatrix}, \quad (8.33)$$

where \mathbf{v} and \mathbf{w} are two vectors of dimension $n-1$ and \mathbf{A}' is a square matrix of dimension $n-1$. Written in this form, one can write an LU decomposition of the matrix \mathbf{A} as follows:

$$\mathbf{A} = \begin{pmatrix} a_{11} & \mathbf{w}^T \\ \mathbf{v} & \mathbf{A}' \end{pmatrix} = \begin{pmatrix} 1 & 0 \\ \frac{\mathbf{v}}{a_{11}} & \mathbf{I}_{n-1} \end{pmatrix} \cdot \begin{pmatrix} a_{11} & \mathbf{w}^T \\ 0 & \mathbf{A}' - \frac{\mathbf{v} \otimes \mathbf{w}}{a_{11}} \end{pmatrix}, \quad (8.34) \quad \Leftarrow \text{Main equation}$$

where \mathbf{I}_{n-1} is an identity matrix of dimension $n-1$. The validity of equation 8.34 can be verified by carrying the product of the two matrices of the right-hand side using matrix components as discussed in section 8.1. We now are left with the problem of finding an LU decomposition for the matrix $\mathbf{A}' - \frac{\mathbf{v} \otimes \mathbf{w}}{a_{11}}$. This matrix is called the Shur's complement of the matrix with respect to the pivot element a_{11} . Let us assume that we have found such a decomposition for that matrix, that is, that we have:

$$\mathbf{A}' - \frac{\mathbf{v} \otimes \mathbf{w}}{a_{11}} = \mathbf{L}' \cdot \mathbf{U}'. \quad (8.35) \quad \Leftarrow \text{Main equation}$$

The we have:

$$\begin{aligned} \begin{pmatrix} 1 & 0 \\ \frac{\mathbf{v}}{a_{11}} & \mathbf{I}_{n-1} \end{pmatrix} \cdot \begin{pmatrix} a_{11} & \mathbf{w}^T \\ 0 & \mathbf{L}' \cdot \mathbf{U}' \end{pmatrix} &= \begin{pmatrix} 1 & 0 \\ \frac{\mathbf{v}}{a_{11}} & \mathbf{I}_{n-1} \end{pmatrix} \cdot \begin{pmatrix} 1 & 0 \\ 0 & \mathbf{L}' \end{pmatrix} \cdot \begin{pmatrix} a_{11} & \mathbf{w}^T \\ 0 & \mathbf{U}' \end{pmatrix} \\ &= \begin{pmatrix} 1 & 0 \\ \frac{\mathbf{v}}{a_{11}} & \mathbf{L}' \end{pmatrix} \cdot \begin{pmatrix} a_{11} & \mathbf{w}^T \\ 0 & \mathbf{U}' \end{pmatrix}. \end{aligned} \quad (8.36)$$

The above equality can be verified by carrying the multiplication with matrix components. The second line of equation 8.36 is the LU decomposition of the matrix \mathbf{A} , which we were looking for. The algorithm is constructed recursively on the successive Shur's complements. For practical implementation, however, it is best to use a loop.

Building the Shur's complement involves a division by the pivot element. As for Gaussian elimination, the algorithm runs into trouble if this element is zero. The expression for the Shur's complement (equations 8.35) shows that, if the pivot element is small, rounding errors occur when computing the elements of the Shur's complement. The strategy avoiding this problem is the same as for Gaussian elimination. One must find the row having the component with the largest absolute value and use this component as the pivot. This is where the permutation matrix comes into play. It will keep track of each row permutation required to bring the row selected for pivoting into the first position. If a non-zero pivot element cannot be found at one step, then the matrix \mathbf{A} is singular and no solution to the system of equation can be found (c.f. similar discussion in section 8.2).

Now that we have proved that an LUP decomposition exists for any non-singular matrix, let us see how it is used to solve a system of linear equations. Consider the system described by equation 8.23; using the LUP decomposition of the matrix \mathbf{A} , it can be rewritten as:

$$\mathbf{LU} \cdot \mathbf{x} = \mathbf{P} \cdot \mathbf{y}, \quad (8.37)$$

where we have used the fact that $\mathbf{P}^{-1} = \mathbf{P}$ for any permutation matrix. Equation 8.37 can be solved in two steps. First one solves the system

$$\mathbf{L} \cdot \tilde{\mathbf{x}} = \mathbf{P} \cdot \mathbf{y} \quad (8.38)$$

using forward substitution. Then, one solves the system

$$\mathbf{U} \cdot \mathbf{x} = \tilde{\mathbf{x}} \quad (8.39)$$

using backward substitution. One can see that the LUP decomposition can be used several times to solve a linear system of equations with the same matrix and different constant vectors.

Performing LUP decomposition is faster than performing Gaussian elimination because Gaussian elimination must also transform the constant vector. To compute the solution vector, however, Gaussian elimination only needs backward substitution whereas LUP requires both forward and backward substitution. The end result is that solving a system of linear equation for a single constant vector is slightly faster using LUP decomposition. If one needs to solve the same system of equations for several constant vectors known in advance, Gaussian elimination is faster. If the constant vectors are not known in advance — or cannot be all stored with the original matrix — LUP decomposition is the algorithm of choice.

8.3.1 LUP decomposition — General implementation

To implement the LUP algorithm, let us first note that we do not need much storage for the three matrices \mathbf{L} , \mathbf{U} and \mathbf{P} . Indeed, the permutation matrix \mathbf{P} can be represented

Figure 8.1 with the box LUPDecomposition grayed.

with a vector of integers of the same dimension as the matrix rows. Since the diagonal elements of the matrix **L** are set in advance, we only need to store elements located below the diagonal. These elements can be stored in the lower part of the matrix **U**. Looking at the definition of the Shur's complement (equation 8.35) and at equation 8.36 we can see that all operations can be performed within a matrix of the same size as the original matrix⁷.

The implementation of the LUP algorithm must create storage to place the components of the matrix whose LUP decomposition is to be computed. A method implements the solving of the system of equations for a given constant vector. Within the method the LUP decomposition itself is performed if it has not yet been made using lazy initialization. During the computation of the LUP decomposition the parity of the permutation is tracked. This information is used to compute the determinant of the matrix (c.f. section 8.4). Thus, the class implementing LUP decomposition has the following instance variables.

`rows` contains a copy of the rows of the matrix representing the system of linear equations, i.e. the matrix **A**; copying the matrix is necessary since LUP decomposition destroys the components; at the end of the LUP decomposition, it will contain the components of the matrices **L** and **U**,

`permutation` contains an array of integers describing the permutation, i.e. the matrix **P**,

`parity` contains parity of the permutation for efficiency purpose⁸.

The instance variable `permutation` is set to undefined (`nil` in Smalltalk, `null` in Java) at initialization time by default. It is used to check whether the decomposition has already been made or not.

The method `solve` implements the solving of the equation system for a given constant vector. It first checks whether the LUP decomposition has been performed. If not, LUP decomposition is attempted. Then, the methods implementing the forward and backward substitution algorithms are called in succession.

8.3.2 LUP decomposition — Smalltalk implementation

Listing 8.8 shows the methods of the class `DhbLUPDecomposition` implementing LUP decomposition in Smalltalk.

To solve the system of equations 8.22 using LUP decomposition, one needs to write to evaluate the following expression:

Code example 8.9

```
( DhbLUPDecomposition equations: #( (3 2 4) (2 -5 -1) (1 -2 2)) )
  solve: #(16 6 10).
```

⁷If the matrix **A** is no longer needed after solving the system of equation, the LUP decomposition can even be performed inside the matrix **A** itself.

⁸The parity is needed to compute the determinant. It could be computed from the parity matrix. However, the overhead of keeping track of the parity is negligible compared to the LUP steps and it is much faster than computing the parity.

This expression has been written on two lines to delineate the various steps. The first line creates an instance of the class `DhbLUPDecomposition` by giving the coefficients of the equations, rows by rows. The second line is a call to the method `solve`: retrieving the solution of the system for the supplied constant vector.

Solving the same system for several constant vectors requires storing the LUP decomposition in a variable:

Code example 8.10

```
l s sol1 sol2l
s := DhbLUPDecomposition equations: #( (3 2 4) (2 -5 -1) (1 -2 2)).
sol1 := s solve: #(16 6 10).
sol2 := s solve: #(7 10 9).
```

When the first solution is fetched, the LUP decomposition is performed; then the solution is computed using forward and backward substitutions. When the second solution is fetched, only forward and backward substitutions are performed.

The default creation class method `new` has been overloaded to prevent creating an object without initialized instance variables. The proper creation class method, `equations:`, takes an array of arrays, the components of the matrix **A**. When a new instance is initialized the supplied coefficients are copied into the instance variable `rows` and the parity of the permutation is set to one. Copying the coefficients is necessary since the storage is reused during the decomposition steps. In addition, some Smalltalk protect constants such as the one used in the code examples above. In this later case, copying is necessary to prevent a read-only exception.

A second creation method `direct:` allows the creation of an instance using the supplied system's coefficients directly. The user of this creation method must keep in mind that the coefficients are destroyed. This creation method can be used when the coefficients have been computed to the sole purpose of solving the system of equations (c.f. sections 10.9.2 and 10.10.2 for an example of use).

The method `protectedDecomposition` handles the case when the matrix is singular by trapping the exception occurring in the method `decompose` performing the actual decomposition. When this occurs, the instance variable `permutation` is set to the integer 0 to flag the singular case. Then, any subsequent calls to the method `solve:` returns `nil`.

Listing 8.8 Smalltalk implementation of the LUP decomposition

Class	DhbLUPDecomposition
Subclass of	Object
Instance variable names:	<code>rows</code> <code>permutation</code> <code>parity</code>

Class methods

`direct:` `anArrayOfArrays`

`^self new initialize:` `anArrayOfArrays`.

equations: anArrayOfArrays

```
^self new initialize: ( anArrayOfArrays collect: [ :each |
                        each deepCopy]).
```

Instance methods

backwardSubstitution: anArray

```
| n sum answer|
n := rows size.
answer := DhbVector new: n.
n to: 1 by: -1 do:
    [ :i |
        sum := anArray at: i.
        ( i + 1) to: n do: [ :j | sum := sum - ( ( ( rows at: i)
                                                    at: j) * ( answer at: j))].
        answer at: i put: sum / ( ( rows at: i) at: i).
    ].
^answer
```

decompose

```
| n |
n := rows size.
permutation := (1 to: n) asArray.
1 to: ( n - 1) do:
    [ :k |
        self swapRow: k withRow: ( self largestPivotFrom: k);
        pivotAt: k.
    ].
```

forwardSubstitution: anArray

```
| n sum answer|
answer := permutation collect: [ :each | anArray at: each].
n := rows size.
2 to: n do:
    [ :i |
        sum := answer at: i.
        1 to: ( i - 1) do: [ :j | sum := sum - ( ( ( rows at: i)
                                                    at: j) * ( answer at: j))].
        answer at: i put: sum.
    ].
^answer
```

initialize: anArrayOfArrays

```
rows := anArrayOfArrays.
parity := 1.
^self
```

largestPivotFrom: anInteger

```
| valueOfMaximum indexOfMaximum value |
valueOfMaximum := ( ( rows at: anInteger) at: anInteger) abs.
indexOfMaximum := anInteger.
( anInteger + 1) to: rows size do:
    [ :n |
        value := ( ( rows at: n) at: anInteger) abs.
        value > valueOfMaximum
            ifTrue: [ valueOfMaximum := value.
                    indexOfMaximum := n.
                    ].
    ].
^indexOfMaximum
```

pivotAt: anInteger

```
| inversePivot size k |
inversePivot := 1 / ( ( rows at: anInteger) at: anInteger).
size := rows size.
k := anInteger + 1.
k to: size
do: [ :i |
    ( rows at: i) at: anInteger put: (( rows at: i) at:
                                     anInteger) * inversePivot.
    k to: size
    do: [ :j |
        ( rows at: i) at: j put: ( ( rows at: i) at: j)
        - ( (( rows at: i) at: anInteger) * (( rows at: anInteger) at: j)).
    ]
].
```

printOn: aStream

```
| first delimitingString n k |
n := rows size.
first := true.
rows do:
    [ :row |
```

```

first ifTrue: [ first := false]
      ifFalse:[ aStream cr].
delimitingString := '('.
row do:
  [ :each |
    aStream nextPutAll: delimitingString.
    each printOn: aStream.
    delimitingString := ' '.
  ].
aStream nextPut: $).
].

```

protectedDecomposition

```

[ self decompose] when: ExAll do: [ :signal | permutation := 0.
      signal exitWith: nil].

```

solve: anArrayOfVector

```

permutation isNil
  ifTrue: [ self protectedDecomposition].
^permutation = 0
  ifTrue: [ nil]
  ifFalse:[ self backwardSubstitution: ( self
      forwardSubstitution: anArrayOfVector)]

```

swapRow: anInteger1 withRow: anInteger2

```

anInteger1 = anInteger2
  ifFalse:[ | swappedRow |
    swappedRow := rows at: anInteger1.
    rows at: anInteger1 put: ( rows at: anInteger2).
    rows at: anInteger2 put: swappedRow.
    swappedRow := permutation at: anInteger1.
    permutation at: anInteger1 put: ( permutation at:
      anInteger2).
    permutation at: anInteger2 put: swappedRow.
    parity := parity negated.
  ].

```

8.3.3 LUP decomposition — Java implementation

Listing 8.9 shows the methods of the class LUPDecomposition implementing LUP decomposition in Java.

To solve the system of equations 8.22, one needs to write to evaluate the following expression:

Code example 8.11

```
double[][] s = {{3,2,4}, {2, -5, -1},{1,-2,2}}};
double[] c = { 16,6,10};
try{
    LUPDecomposition system = new LUPDecomposition( s);
    double[] solution = system.solve(c);
    <Processing of the solution>

} catch(DhbIllegalDimension e){}
```

The first two lines define the components of the components of the system's matrix (**A**) and the components of the constant vector (**y**). The try...catch block delineates the solution of the system proper. The first line of the block creates an instance of the class LUPDecomposition using the constructor method taking a double dimensional array as argument. The second line is a call to the method solve retrieving the solution of the system for the supplied constant vector.

Solving the same system with an additional constant vector requires a couple additional lines of code:

Code example 8.12

```
double[][] s = {{3,2,4}, {2, -5, -1},{1,-2,2}}};
double[] c1 = {16,6,10};
double[] c2 = {7,10,9};
try{
    LUPDecomposition system = new LUPDecomposition( s, c);
    double[] solution1 = system.solve(c1);
    double[] solution2 = system.solve(c2);
    <Processing of the solutions>

} catch(DhbIllegalDimension e){}
```

In this case, the second solution is fetched with a second call to the method solve. Of course, LUP decomposition is only performed once, namely when the first solution is fetched. When the second solution is retrieved, only forward and backward substitution are performed.

The class has three constructor methods. The first constructor takes a two dimensional array of doubles as argument. The second constructor takes a matrix as argument. The third constructor takes a symmetric matrix as argument. When a matrix is supplied, the components of the matrix are copied into the instance variable rows. Copying the matrix is necessary since the storage is reused during the decomposition steps. The first two constructor methods throw the exception DhbIllegalDimension since LUP decomposition is only defined for square matrix. A symmetric matrix being always square, the exception is not needed for the third constructor method.

In case of a singular matrix the method `pivot` throws an `ArithmeticException`. When this happens, the instance variable `permutation` is set to an array of zero length to flag the singular case. This is to prevent the failing decomposition from being attempted more than one time. Then, any subsequent calls to the method `solve` throws an `ArithmeticException`.

Listing 8.9 Java implementation of the LUP decomposition

```
package DhbMatrixAlgebra;

/**
 * Lower Upper Permutation (LUP) decomposition
 *
 * @author Didier H. Besset
 */
public class LUPDecomposition
{
    /**
     * Rows of the system
     */
    private double[][] rows;
    /**
     * Permutation
     */
    private int[] permutation = null;
    /**
     * Permutation's parity
     */
    private int parity = 1;
    /**
     * Constructor method
     * @param components double[][]
     * @exception DhbMatrixAlgebra.DhbIllegalDimension
     *           the supplied matrix is not square
     */
    public LUPDecomposition ( double[][]components)
        throws DhbIllegalDimension
    {
        int n = components.length;
        if ( components[0].length != n )
            throw new DhbIllegalDimension("Illegal system: a"+n+" by "
                +components[0].length+" matrix is not a square matrix");
        rows = components;
        initialize();
    }
    /**
```



```

* Constructor method.
* @param m DhbMatrixAlgebra.Matrix
* @exception DhbMatrixAlgebra.DhbIllegalDimension
*           the supplied matrix is not square
*/
public LUPDecomposition ( Matrix m) throws DhbIllegalDimension
{
    if ( !m.isSquare() )
        throw new DhbIllegalDimension(
            "Supplied matrix is not a square matrix");
    initialize( m.components);
}
/**
* Constructor method.
* @param m DhbMatrixAlgebra.DhbSymmetricMatrix
*/
public LUPDecomposition ( SymmetricMatrix m)
{
    initialize( m.components);
}
/**
* @return double[]
* @param xTilde double[]
*/
private double[] backwardSubstitution( double[] xTilde)
{
    int n = rows.length;
    double[] answer = new double[n];
    for ( int i = n - 1; i >= 0; i--)
    {
        answer[i] = xTilde[i];
        for ( int j = i + 1; j < n; j++)
            answer[i] -= rows[i][j] * answer[j];
        answer[i] /= rows[i][i];
    }
    return answer;
}
private void decompose()
{
    int n = rows.length;
    permutation = new int[n];
    for ( int i = 0; i < n; i++ )
        permutation[i] = i;
    parity = 1;
    try {
        for ( int i = 0; i < n; i++)

```

```

        {
            swapRows( i, largestPivot( i));
            pivot( i);
        }
    } catch ( ArithmeticException e) { parity = 0;};
}
/**
 * @return boolean    true if decomposition was done already
 */
private boolean decomposed()
{
    if ( parity == 1 && permutation == null )
        decompose();
    return parity != 0;
}
/**
 * @return double[]
 * @param c double[]
 */
public double determinant()
{
    if ( !decomposed() )
        return Double.NaN;
    double determinant = parity;
    for ( int i = 0; i < rows.length; i++ )
        determinant *= rows[i][i];
    return determinant;
}
/**
 * @return double[]
 * @param c double[]
 */
private double[] forwardSubstitution( double[] c)
{
    int n = rows.length;
    double[] answer = new double[n];
    for ( int i = 0; i < n; i++)
    {
        answer[i] = c[permutation[i]];
        for ( int j = 0; j <= i - 1; j++)
            answer[i] -= rows[i][j] * answer[j];
    }
    return answer;
}
private void initialize ()
{

```

```

    permutation= null;
    parity = 1;
}
/**
 * @param components double[][] components obtained from constructor methods.
 */
private void initialize ( double[][] components)
{
    int n = components.length;
    rows = new double[n][n];
    for ( int i = 0; i < n; i++)
    {
        for ( int j = 0; j < n; j++)
            rows[i][j] = components[i][j];
    }
    initialize();
}
/**
 * @return double[]
 * @param c double[]
 */
public double[][] inverseMatrixComponents()
{
    if ( !decomposed() )
        return null;
    int n = rows.length;
    double[][] inverseRows = new double[n][n];
    double[] column = new double[n];
    for ( int i = 0; i < n; i ++ )
    {
        for ( int j = 0; j < n; j++ )
            column[j] = 0;
        column[i] = 1;
        column = solve( column);
        for ( int j = 0; j < n; j++ )
            inverseRows[i][j] = column[j];
    }
    return inverseRows;
}
/**
 * @return int
 * @param k int
 */
private int largestPivot(int k)
{
    double maximum = Math.abs( rows[k][k]);

```

```

    double abs;
    int index = k;
    for ( int i = k + 1; i < rows.length; i++)
    {
        abs = Math.abs( rows[i][k]);
        if ( abs > maximum )
        {
            maximum = abs;
            index = i;
        }
    }
    return index;
}
/**
 * @param k int
 */
private void pivot( int k)
{
    double inversePivot = 1 / rows[k][k];
    int k1 = k + 1;
    int n = rows.length;
    for ( int i = k1; i < n; i++)
    {
        rows[i][k] *= inversePivot;
        for ( int j = k1; j < n; j++)
            rows[i][j] -= rows[i][k] * rows[k][j];
    }
}
/**
 * @return double[]
 * @param c double[]
 */
public double[] solve( double[] c)
{
    return decomposed()
        ? backwardSubstitution( forwardSubstitution( c))
        : null;
}
/**
 * @return double[]
 * @param c double[]
 */
public DhbVector solve( DhbVector c)
{
    double[] components = solve( c.components);
    if ( components == null )

```

```

        return null;
    return components == null ? null : new DhbVector( components);
}
/**
 * @param i int
 * @param k int
 */
private void swapRows( int i, int k)
{
    if ( i != k )
    {
        double temp;
        for ( int j = 0; j < rows.length; j++ )
        {
            temp = rows[i][j];
            rows[i][j] = rows[k][j];
            rows[k][j] = temp;
        }
        int nTemp;
        nTemp = permutation[i];
        permutation[i] = permutation[k];
        permutation[k] = nTemp;
        parity = -parity;
    }
}
/**
 * Make sure the supplied matrix components are those of
 * a symmetric matrix
 * @param components double
 */
public static void symmetrizeComponents( double[][] components)
{
    for ( int i = 0; i < components.length; i++ )
    {
        for ( int j = i + 1; j < components.length; j++ )
        {
            components[i][j] += components[j][i];
            components[i][j] *= 0.5;
            components[j][i] = components[i][j];
        }
    }
}
/**
 * Returns a String that represents the value of this object.
 * @return a string representation of the receiver
 */

```

```

public String toString()
{
    StringBuffer sb = new StringBuffer();
    char[] separator = { '[', ' ' };
    int n = rows.length;
    for ( int i = 0; i < n; i++)
    {
        separator[0] = '{';
        for ( int j = 0; j < n; j++)
        {
            sb.append( separator);
            sb.append( rows[i][j]);
            separator[0] = ' ';
        }
        sb.append('}');
        sb.append('\n');
    }
    if ( permutation != null )
    {
        sb.append( parity == 1 ? '+' : '-' );
        sb.append(" ( " + permutation[0]);
        for ( int i = 1; i < n; i++)
            sb.append(", " + permutation[i]);
        sb.append(')');
        sb.append('\n');
    }
    return sb.toString();
}
}

```

8.4 Computing the determinant of a matrix

The determinant of a matrix is defined as

$$\det \mathbf{A} = \begin{vmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{vmatrix} = \sum_{\pi} \text{sign}(\pi) a_{1\pi(1)} a_{2\pi(1)} \cdots a_{n\pi(1)}, \quad (8.40)$$

where π represents a permutation of the indices 1 to n . The sum of equation 8.40 is made over all possible permutations. Thus, the sum contains $n!$ terms. Needless to say, the direct implementation of equation 8.40 to compute a determinant is highly inefficient.

Fortunately, the determinant of a matrix can be computed directly from its LUP decomposition. This comes from the following three properties of the determinants:

1. the determinant of a product of two matrices is the product of the determinants of the two matrices;
2. the determinant of a permutation matrix is 1 or -1 depending on the parity of the permutation;
3. the determinant of a triangular matrix is the product of its diagonal elements.

Applying the three properties above to equation 8.32 shows us that the determinant of the matrix **A** is simply the product of the diagonal elements of the matrix **U** times the sign of the permutation matrix **P**.

The parity of the permutation can be tracked while performing the LUP decomposition itself. The cost of this is negligible compared to the rest of the algorithm so that it can be done whether or not the determinant is needed. The initial parity is 1. Each additional permutation of the rows multiplies the parity by -1.

8.4.1 Computing the determinant of matrix — General implementation

Our implementation uses the fact that objects of the class **Matrix** have an instance variable in which the LUP decomposition is kept. This variable is initialized using lazy initialization: if no LUP decomposition exists, it is calculated. Then the computation of the determinant is delegated to the LUP decomposition object.

Since the LUP decomposition matrices are kept within a single storage unit, a matrix, the LUP decomposition object calculates the determinant by taking the product of the diagonal elements of the matrix of the LUP decomposition object and multiplies the product by the parity of the permutation to obtain the final result.

8.4.2 Computing the determinant of matrix — Smalltalk implementation

Listing 8.10 shows the methods of classes **DhbMatrix** and **DhbLUPDecomposition** needed to compute a matrix determinant.

Listing 8.10 Smalltalk methods to compute a matrix determinant

```

Class                DhbMatrix
Subclass of          Object
Instance variable names:  rows lupDecomposition

Instance methods
determinant
^self lupDecomposition determinant

```

Class **DhbLUPDecomposition**
 Subclass of Object
 Instance variable names: rows permutation parity

Instance methods

determinant

```

| n |
permutation isNil
  ifTrue: [ self protectedDecomposition].
permutation = 0
  ifTrue: [ ^0].

```

8.4.3 Computing the determinant of matrix — Java implementation

The code computing the determinant of a matrix consists of the method `determinant` of the class `Matrix` (c.f. listing 8.2) and the method `determinant` of the class `LUPDecomposition` (c.f. listing 8.9).

8.5 Matrix inversion

The inverse of a square matrix \mathbf{A} is denoted \mathbf{A}^{-1} . It is defined by the following equation:

$$\mathbf{A} \cdot \mathbf{A}^{-1} = \mathbf{I}, \quad (8.41)$$

where \mathbf{I} is the identity matrix.

To determine the coefficients of the inverse matrix, one could use equation 8.41 as a system of n^2 linear equations if n is the dimension of the matrix \mathbf{A} . this system could be solved using either Gaussian elimination (c.f. section 8.2) or LUP decomposition (c.f. section 8.3).

Using Gaussian elimination for such a system requires solving a system with n constant vectors. This could be done, but it is not very practical in terms of storage space except for matrices of small dimension. If we already have the LUP decomposition of the matrix \mathbf{A} , however, this is indeed a solution. One can solve equation 8.41 for each columns of the matrix \mathbf{A}^{-1} . Specifically, \mathbf{c}_i , the i^{th} column of the matrix \mathbf{A} , is the solution of the following system:

$$\mathbf{A} \cdot \mathbf{c}_i = \mathbf{e}_i \quad \text{for } i = 1, \dots, n, \quad (8.42)$$

where \mathbf{e}_i is the i^{th} column of the identity matrix, that is a vector with zero components except for the i^{th} component whose value is 1.

For large matrices, however, using LUP decomposition becomes quite slow. A cleverer algorithm for symmetric matrices is given in [Cormen et al.] with no name. In this book, we shall refer to this algorithm as the CLR algorithm (acronym of the authors of [Cormen et al.]).

Let \mathbf{A} be a symmetric matrix. In section 8.1 we have seen that it can be written in the form:

$$\mathbf{A} = \begin{pmatrix} \mathbf{B} & \mathbf{C}^T \\ \mathbf{C} & \mathbf{D} \end{pmatrix}, \quad (8.43)$$

where \mathbf{B} and \mathbf{D} are two symmetric matrices and \mathbf{C} is in general not a square matrix. Then the inverse can be written as:

Main equation⇒

$$\mathbf{A}^{-1} = \begin{pmatrix} \mathbf{B}^{-1} + \mathbf{B}^{-1} \cdot \mathbf{C}^T \cdot \mathbf{S}^{-1} \cdot \mathbf{C} \cdot \mathbf{B}^{-1} & \mathbf{B}^{-1} \cdot \mathbf{C}^T \cdot \mathbf{S}^{-1} \\ -\mathbf{S}^{-1} \cdot \mathbf{C} \cdot \mathbf{B}^{-1} & \mathbf{S}^{-1} \end{pmatrix}, \quad (8.44)$$

where the matrix \mathbf{S} is called the Schur's complement of the matrix \mathbf{A} with respect to the partition of equation 8.43. The matrix \mathbf{S} is also a symmetric matrix, given by the expression:

Main equation⇒

$$\mathbf{S} = \mathbf{D} - \mathbf{C} \cdot \mathbf{B}^{-1} \cdot \mathbf{C}^T. \quad (8.45)$$

The reader can verify that equation 8.44 gives indeed the inverse of \mathbf{A} by plugging 8.45 into 8.44 and carrying the multiplication with 8.43 in the conventional way. The result of the multiplication is an identity matrix. In particular, the result is independent of the type of partition described in equation 8.43.

The CRL algorithm consists of computing the inverse of a symmetric matrix using equations 8.43, 8.44 and 8.45 recursively. It is a divide-and-conquer algorithm in which the partition of equation 8.43 is further applied to the matrices \mathbf{B} and \mathbf{S} . First the initial matrix is divided into four parts of approximately the same size. At each step the two inverses, \mathbf{B}^{-1} and \mathbf{S}^{-1} are computed using a new partition until the matrices to be inverted are either 2 by 2 or 1 by 1 matrices.

In the book of Cormen et al. [Cormen et al.] the divide and conquer algorithm supposes that the dimension of the matrix is a power of two to be able to use Strassen's algorithm for matrix multiplication. We have investigated an implementation of Strassen's algorithm, unfortunately its performance is still inferior to that of regular multiplication for matrices of dimension up to 512, probably because of the impact of garbage collection⁹. Indeed, the increase in memory requirement can be significant for matrices of moderate size. A 600 by 600 matrix requires 2.7 megabytes of storage. Extending it to a 1024 by 1024 would require 8 megabytes of storage.

Implementation strategy

In our implementation, we have generalized the divide and conquer strategy to any dimension. The dimension of the upper partition is selected at each partition to be the largest power of two smaller than the dimension of the matrix to be partitioned. Although the dimension of the matrices is not an integral power of two the number of necessary partitions remains a logarithmic function of the dimension of the original matrix, thus preserving the original idea of the CRL algorithm. It turns out that the

⁹The author thanks Thomas Cormen for enlightening E-mails on this subject.




Figure 8.2: Comparison of inversion time for non-symmetrical matrices

number of necessary partitions is, in most cases, smaller than the number of partitions needed if the dimension of the original matrix is extended to the nearest largest power of two.

Both LUP and CRL algorithm perform within a time $O(n^2)$ where n is the dimension of the matrix. Figure 8.2 shows the time needed to inverse a non-symmetrical matrix using CRL algorithm (solid line) and LUP decomposition (broken line), as well as the ratio between the two times (dotted line). The CRL algorithm has a large overhead but a smaller factor for the dependency on dimension. Thus, computing the inverse of a matrix using LUP decomposition is faster than the CLR algorithm for small matrices and slower for large matrices. As a consequence, our implementation of matrix inversion uses a different algorithm depending on the dimension of the matrix: if the dimension of the matrix is below a critical dimension, LUP decomposition is used; otherwise the CRL algorithm is used. In addition, LUP decomposition is always used if it has already been computed for another purpose.

On figure 8.2 we can determine that the critical dimension, below which the LUP decomposition works faster than the CRL algorithm, is about 36. These data were collected on a Pentium II running Windows NT 4.0. As this value is depending on the performance of the operating system, the reader is advised to determine the critical dimension again when installing the classes on another system.

In practice, the CLR algorithm described in equations 8.43 to 8.45 can only be applied to symmetric matrices. In [Cormen et al.] Cormen et al. propose to generalize it to matrices of any size by observing the following identity:

$$\mathbf{A} \cdot [(\mathbf{A}^T \cdot \mathbf{A})^{-1} \cdot \mathbf{A}^T] = \mathbf{I} \quad (8.46)$$

which can be verified for any matrix \mathbf{A} . Thus, the expression in bracket can be

considered as the inverse of the matrix \mathbf{A} . In mathematics, it is called the pseudo-inverse or the Moore-Penrose inverse. Since the product $\mathbf{A}^T \cdot \mathbf{A}$ is always a symmetric matrix, its inverse can be computed with the CRL algorithm. In practice, however, this technique is plagued with rounding errors and should be used with caution (c.f. section 8.5.3).

8.5.1 Matrix inversion — Smalltalk implementation

Listing 8.11 shows the complete implementation in Smalltalk. It contains additional methods for the classes `DhbMatrix` and `DhbSymmetricMatrix`.

For symmetric matrices the method `inverse` first tests whether the dimension of the matrix is below a given threshold — defined by the class method `lupCRLCriticalDimension` — or whether the LUP decomposition of the matrix was already performed. In that case, the inverse is computed from the LUP decomposition using the method described at the beginning of section 8.5. Otherwise the CRL algorithm is used. The implementation of the CRL algorithm is straightforward thanks to the matrix operators defined in section 8.1.1.

For non-symmetric matrices the method `inverse` first tests whether the matrix is square or not. If the matrix is square, LUP decomposition is used. If the matrix is not square the pseudo inverse is computed using equation 8.46.

In both cases there is no error handling. Inverting a singular matrix produces an arithmetic error which must be handled by the calling method.

Listing 8.11 Smalltalk implementation of matrix inversion

Class **DhbSymmetricMatrix**
 Subclass of `DhbMatrix`

Class methods

`join: anArrayOfMatrices`

```

| rows n |
rows := OrderedCollection new.
n := 0.
( anArrayOfMatrices at: 1) rowsDo:
  [ :each |
    n := n + 1.
    rows add: each, ( ( anArrayOfMatrices at: 3) columnAt: n).
  ].
n := 0.
( anArrayOfMatrices at: 2) rowsDo:
  [ :each |
    n := n + 1.
    rows add: ( ( anArrayOfMatrices at: 3) rowAt: n), each.
  ].
^self rows: rows
```

lupCRLCriticalDimension

^36

Instance methods

crlInverse

```

| matrices b1 cb1ct cb1 |
matrices := self split.
b1 := (matrices at: 1) inverse.
cb1 := (matrices at: 3) * b1.
cb1ct := (cb1 productWithTransposeMatrix: (matrices at: 3))
        asSymmetricMatrix.
matrices at: 3 put: (matrices at: 2) * cb1.
matrices at: 2 put: ((matrices at: 2) accumulateNegated: cb1ct)
                  inverse.
matrices at: 1 put: (b1 accumulate: (cb1
                                     transposeProductWithMatrix: (matrices at: 3))).
(matrices at: 3) negate.
^self class join: matrices

```

inverse

```

^(rows size < self class lupCRLCriticalDimension or:
   [lupDecomposition notNil])
  ifTrue: [self lupInverse]
  ifFalse: [self crlInverse]

```

split

```

| n b d c |
n := self largestPowerOf2SmallerThan: rows size.
^Array with: ( self class rows: ( ( 1 to: n) asVector collect: [
    :k | ( rows at: k) copyFrom: 1 to: n]))
  with: ( self class rows: ( ( (n+1) to: rows size)
    asVector collect: [ :k | ( rows at: k) copyFrom: (n+1) to: rows
    size]))
  with: ( self class superclass rows: ( ( (n+1) to: rows
    size) asVector collect: [ :k | ( rows at: k) copyFrom: 1 to: n]))

```

Class
Subclass of

DhbMatrix
Object

Instance variable names: rows lupDecomposition

Class methods

lupCRLCriticalDimension

⁴⁰

Instance methods

inverse

[^]self isSquare

ifTrue: [self lupInverse]

ifFalse: [self squared inverse * self transpose]

largestPowerOf2SmallerThan: anInteger

| m m2|

m := 2.

[m2 := m * 2.

m2 < anInteger] whileTrue:[m := m2].

[^]m

lupInverse

[^]self class rows: self lupDecomposition inverseMatrixComponents

8.5.2 Matrix inversion — Java implementation

Listing 8.12 shows the implementation of the class `SymmetricMatrix` in Java.

Addition and product of two symmetric matrices yield a symmetric matrix. Thus, the corresponding methods have been redefined for that class.

Because the method to create a symmetric matrix from components is used frequently in the rest of the package, no check is made for the validity of the components. Of course, this method is declared as protected to prevent classes from outside the package from using it. Such classes must use the static method `fromComponents` to create a symmetric matrix with supplied components, which throws the exception `DhbNonSymmetricComponents` if the supplied components are not those of a symmetric matrix. The code for this exception is elementary and is not shown here. It is left as an exercise for the reader.

The method `inverse` for symmetric matrices first tests whether the dimension of the matrix is below a given threshold — defined by the static variable `lupCRLCriticalDimension` — or whether the LUP decomposition of the matrix was already performed. In that case,

the inverse is computed from the LUP decomposition using the method described at the beginning of section 8.5. Otherwise the CRL algorithm is used. The implementation of the CRL algorithm is a direct transcription of equations 8.43 to 8.45 using the methods defined in section 8.1.2.

Listing 8.12 Java implementation of the class SymmetricMatrix

```
package DhbMatrixAlgebra;

/**
 * Symmetric matrix
 *
 * @author Didier H. Besset
 */
public class SymmetricMatrix extends Matrix {

    private static int lupCRLCriticalDimension = 36;
    /**
     * Creates a symmetric matrix with given components.
     * @param a double[][]
     */
    protected SymmetricMatrix(double[][] a)
    {
        super(a);
    }
    /**
     * @param n int
     * @exception java.lang.NegativeArraySizeException if n <= 0
     */
    public SymmetricMatrix (int n ) throws NegativeArraySizeException
    {
        super( n, n);
    }
    /**
     * Constructor method.
     * @param n int
     * @param m int
     * @exception java.lang.NegativeArraySizeException if n,m <= 0
     */
    public SymmetricMatrix(int n, int m) throws NegativeArraySizeException {
        super(n, m);
    }
    /**
     * @return SymmetricMatrix    sum of the matrix with the supplied matrix.
     * @param a DhbMatrix
     * @exception DhbIllegalDimension if the supplied matrix does not
```

```

*                               have the same dimensions.
*/
public SymmetricMatrix add ( SymmetricMatrix a)
                               throws DhbIllegalDimension
{
    return new SymmetricMatrix( addComponents( a));
}
/**
 * Answers the inverse of the receiver computed via the CRL algorithm.
 * @return DhbMatrixAlgebra.SymmetricMatrix
 * @exception java.lang.ArithmeticException if the matrix is singular.
 */
private SymmetricMatrix crlInverse ( ) throws ArithmeticException
{
    if ( rows() == 1)
        return inverse1By1();
    else if ( rows() == 2)
        return inverse2By2();
    Matrix[] splitMatrices = split();
    SymmetricMatrix b1 = (SymmetricMatrix) splitMatrices[0].inverse();
    Matrix cb1 = splitMatrices[2].secureProduct( b1);
    SymmetricMatrix cb1cT = new SymmetricMatrix(
        cb1.productWithTransposedComponents(splitMatrices[2]));
    splitMatrices[1] = ( (SymmetricMatrix)
        splitMatrices[1]).secureSubtract( cb1cT).inverse();
    splitMatrices[2] = splitMatrices[1].secureProduct( cb1);
    splitMatrices[0] = b1.secureAdd( new SymmetricMatrix(
        cb1.transposedProductComponents( splitMatrices[2])));
    return SymmetricMatrix.join( splitMatrices);
}
/**
 * @return DhbMatrixAlgebra.SymmetricMatrix
 * @param comp double[][] components of the matrix
 * @exception DhbMatrixAlgebra.DhbIllegalDimension
 *           The supplied components are not those of a square matrix.
 * @exception DhbMatrixAlgebra.DhbNonSymmetricComponents
 *           The supplied components are not symmetric.
 */
public static SymmetricMatrix fromComponents ( double[][] comp)
    throws DhbIllegalDimension, DhbNonSymmetricComponents
{
    if ( comp.length != comp[0].length)
        throw new DhbIllegalDimension( "Non symmetric components: a "
            +comp.length+" by "+comp[0].length
            +" matrix cannot be symmetric");
    for ( int i = 0; i < comp.length; i++)

```

```

    {
        for ( int j = 0; j < i; j++)
        {
            if ( comp[i][j] != comp[j][i])
                throw new DhbNonSymmetricComponents(
                    "Non symmetric components: a["+i+"]["+j
                        +"]= "+comp[i][j]+", a["+j+"]["+
                            +i+"]= "+comp[j][i]);
        }
    }
    return new SymmetricMatrix( comp);
}
/**
 * @return SymmetricMatrix  an identity matrix of size n
 * @param n int
 */
public static SymmetricMatrix identityMatrix ( int n)
{
    double[][] a = new double[n][n];
    for ( int i = 0; i < n; i++)
    {
        for ( int j = 0; j < n; j++) a[i][j] = 0;
        a[i][i] = 1;
    }
    return new SymmetricMatrix( a);
}
/**
 * @return DhbMatrix      inverse of the receiver.
 * @exception java.lang.ArithmeticException if the receiver is
 *                                     a singular matrix.
 */
public Matrix inverse ( ) throws ArithmeticException
{
    return rows() < lupCRLCriticalDimension
        ? new SymmetricMatrix(
            (new LUPDecomposition( this)).inverseMatrixComponents())
        : crlInverse( );
}
/**
 * Compute the inverse of the receiver in the case of a 1 by 1 matrix.
 * Internal use only: no check is made.
 * @return DhbMatrixAlgebra.SymmetricMatrix
 */
private SymmetricMatrix inverse1By1 ( )
{
    double[][] newComponents = new double[1][1];

```



```

        newComponents[0][0] = 1 / components[0][0];
        return new SymmetricMatrix( newComponents);
    }
    /**
     * Compute the inverse of the receiver in the case of a 2 by 2 matrix.
     * Internal use only: no check is made.
     * @return DhbMatrixAlgebra.SymmetricMatrix
     */
    private SymmetricMatrix inverse2By2 ( )
    {
        double[][] newComponents = new double[2][2];
        double inverseDeterminant = 1 / ( components[0][0] * components[1][1]
            - components[0][1] * components[1][0]);
        newComponents[0][0] = inverseDeterminant * components[1][1];
        newComponents[1][1] = inverseDeterminant * components[0][0];
        newComponents[0][1] = newComponents[1][0] =
            -inverseDeterminant * components[1][0];
        return new SymmetricMatrix( newComponents);
    }
    /**
     * Build a matrix from 3 parts (inverse of split).
     * @return DhbMatrixAlgebra.SymmetricMatrix
     * @param a DhbMatrixAlgebra.Matrix[]
     */
    private static SymmetricMatrix join ( Matrix[] a)
    {
        int p = a[0].rows();
        int n = p + a[1].rows();
        double[][] newComponents = new double[n][n];
        for ( int i = 0; i < p; i++)
        {
            for ( int j = 0; j < p; j++)
                newComponents[i][j] = a[0].components[i][j];
            for ( int j = p; j < n; j++)
                newComponents[i][j] = newComponents[j][i] =
                    -a[2].components[j-p][i];
        }
        for ( int i = p; i < n; i++)
        {
            for ( int j = p; j < n; j++)
                newComponents[i][j] = a[1].components[i-p][j-p];
        }
        return new SymmetricMatrix( newComponents);
    }
    /**
     * @return int

```

```

    * @param n int
    */
private int largestPowerOf2SmallerThan ( int n)
{
    int m = 2;
    int m2;
    while (true)
    {
        m2 = 2 * m;
        if ( m2 >= n )
            return m;
        m = m2;
    }
}
/**
 * @return DhbMatrixAlgebra.SymmetricMatrix
 * @param a double
 */
public Matrix product ( double a)
{
    return new SymmetricMatrix( productComponents( a));
}
/**
 * @return Matrix    product of the receiver with the supplied matrix
 * @param a Matrix
 * @exception DhbIllegalDimension If the number of columns of
 * the receivers are not equal to the number of rows
 * of the supplied matrix.
 */
public SymmetricMatrix product ( SymmetricMatrix a) throws DhbIllegalDimension
{
    return new SymmetricMatrix( productComponents( a));
}
/**
 * @return DhbMatrixAlgebra.Matrix    product of the receiver with
 * the transpose of the supplied matrix
 * @param a DhbMatrixAlgebra.Matrix
 * @exception DhbMatrixAlgebra.DhbIllegalDimension If the number of
 * columns of the receiver are not equal to the number of
 * columns of the supplied matrix.
 */
public SymmetricMatrix productWithTransposed ( SymmetricMatrix a)
    throws DhbIllegalDimension
{
    if ( a.columns() != columns() )
        throw new DhbIllegalDimension(

```

```

        "Operation error: cannot multiply a "
        +rows()+" by "+columns()
        +" matrix with the transpose of a "
        +a.rows()+" by "+a.columns()+" matrix");
    return new SymmetricMatrix( productWithTransposedComponents( a));
}
/**
 * Same as add ( SymmetricMatrix a), but without dimension checking.
 * @return DhbMatrixAlgebra.SymmetricMatrix
 * @param a DhbMatrixAlgebra.SymmetricMatrix
 */
protected SymmetricMatrix secureAdd ( SymmetricMatrix a)
{
    return new SymmetricMatrix( addComponents( a));
}
/**
 * Same as product(DhbSymmetricMatrix a), but without dimension checking.
 * @return DhbMatrixAlgebra.SymmetricMatrix
 * @param a DhbMatrixAlgebra.SymmetricMatrix
 */
protected SymmetricMatrix secureProduct ( SymmetricMatrix a)
{
    return new SymmetricMatrix( productComponents( a));
}
/**
 * Same as subtract ( SymmetricMatrix a), but without dimension checking.
 * @return DhbMatrixAlgebra.SymmetricMatrix
 * @param a DhbMatrixAlgebra.SymmetricMatrix
 */
protected SymmetricMatrix secureSubtract ( SymmetricMatrix a)
{
    return new SymmetricMatrix( subtractComponents( a));
}
/**
 * Divide the receiver into 3 matrices or approximately equal dimension.
 * @return DhbMatrixAlgebra.Matrix[] Array of splitted matrices
 */
private Matrix[] split ( )
{
    int n = rows();
    int p = largestPowerOf2SmallerThan( n);
    int q = n - p;
    double[][] a = new double[p][p];
    double[][] b = new double[q][q];
    double[][] c = new double[q][p];
    for ( int i = 0; i < p; i++)

```

```

    {
        for ( int j = 0; j < p; j++)
            a[i][j] = components[i][j];
        for ( int j = p; j < n; j++)
            c[j-p][i] = components[i][j];
    }
    for ( int i = p; i < n; i++)
    {
        for ( int j = p; j < n; j++)
            b[i-p][j-p] = components[i][j];
    }
    Matrix[] answer = new Matrix[3];
    answer[0] = new SymmetricMatrix( a);
    answer[1] = new SymmetricMatrix( b);
    answer[2] = new Matrix( c);
    return answer;
}
/**
 * @return DHBmatrixAlgebra.SymmetricMatrix
 * @param a DHBmatrixAlgebra.SymmetricMatrix
 * @exception DHBmatrixAlgebra.DhbIllegalDimension (from constructor).
 */
public SymmetricMatrix subtract ( SymmetricMatrix a)
    throws DhbIllegalDimension
{
    return new SymmetricMatrix( subtractComponents( a));
}
/**
 * @return DHBmatrixAlgebra.Matrix    the same matrix
 */
public Matrix transpose()
{
    return this;
}
/**
 * @return DhbMatrixAlgebra.SymmetricMatrix    product of the tranpose
 *          of the receiver with the supplied matrix
 * @param a DhbMatrixAlgebra.SymmetricMatrix
 * @exception DhbMatrixAlgebra.DhbIllegalDimension If the number of
 *          rows of the receiver are not equal to
 *          the number of rows of the supplied matrix.
 */
public SymmetricMatrix transposedProduct ( SymmetricMatrix a)
    throws DhbIllegalDimension
{
    if ( a.rows() != rows() )

```

```

        throw new DhbIllegalDimension(
            "Operation error: cannot multiply a tranposed "
            +rows()+" by "+columns()+" matrix with a "+
            a.rows()+" by "+a.columns()+" matrix");
    return new SymmetricMatrix( transposedProductComponents( a));
}
}

```

For non-symmetric matrices the method `inverse` (c.f. listing 8.5) attempts to compute the inverse using LUP decomposition within a `try...catch` block trapping the exception `DhbIllegalDimension`. This kind of exception will occur for any non-square matrix. If the exception occurs, the pseudo inverse is computed using equation 8.46.

8.5.3 Matrix inversion — Rounding problems

Operations with large matrices are well known to exhibit serious rounding problems. The reason is that the computation of the vector product of each row and column is a sum: the higher the dimension and the longer the sum. For large matrix dimensions the magnitude of the sum can mask small contributions from single products. Successive multiplications thus amplify initial small deviations. This is especially the case when computing the inverse of a general matrix using the CRL algorithm combined with the pseudo-inverse (8.46).

Now is the time to unveil the mystery example of section 1.3.3 about rounding errors propagation. The problem solved in this example is matrix inversion. The parameter describing the complexity of the problem is the dimension of the matrix. This is the quantity plotted along the x -axis of figure 1.1. Let \mathbf{A} the matrix to be inverted. The matrix \mathbf{M} defined by

$$\mathbf{M} = \mathbf{A}^{-1} \cdot \mathbf{A} - \mathbf{I}, \quad (8.47)$$

should have all its components equal to zero. The precision of the result is defined as the largest absolute value over all components of the matrix \mathbf{M} . That quantity is plotted along the y -axis of figure 1.1.

Method A computes the inverse of the matrix using LUP decomposition, method B using the CRL algorithm. The general data correspond to a matrix whose components were generated by a random number generator (c.f. section 9.4). They were all comprised between 0 and 1. For the special data the matrix \mathbf{A} is a covariance matrix (c.f. section 12.2) obtained by generating 1000 vectors with random components comprised between 0 and 1. For method B general data, the inverse of a non-symmetrical matrix is computed using the CRL algorithm combined with equation 8.46. In this general form, the CRL algorithm is faster the LUP for matrices of dimensions larger than about 165. The precision, however, is totally unreliable as can be seen on Figure 1.1.

8.6 Matrix eigenvalues and eigenvectors of a non-symmetric matrix

A non-zero vector \mathbf{u} is called an eigenvector of the matrix \mathbf{M} if there exists a complex number λ such that:

$$\mathbf{M} \cdot \mathbf{u} = \lambda \mathbf{u}. \quad (8.48)$$

the number λ is called an eigenvalue of the matrix \mathbf{M} . Equation 8.48 implies that the matrix \mathbf{M} must be a square matrix. In general a non-singular matrix of dimension n has n eigenvalues and eigenvectors. Some eigenvalues, however, may be double¹⁰. Discussing the existence of eigenvalues in the general case goes beyond the scope of this book. Equation 8.48 shows that an eigenvector is defined up to a constant¹¹. One can prove that two eigenvectors of the same matrix, but corresponding to two different eigenvalues, are orthogonal to each other [Bass]. Thus, the eigenvectors of a matrix form a complete set of reference in a n dimensional space.

Computing the eigenvalues and eigenvectors of an arbitrary matrix is a difficult task. Solving this problem in the general case is quite demanding numerically. In the rest of this section we give an algorithm which works well when the absolute value of one of the eigenvalues is much larger than that of the others. The next section discusses Jacobi's algorithm finding all eigenvalues of a symmetrical matrix.

For an arbitrary square matrix the eigenvalue with the largest absolute value can be found with an iterative process. Let \mathbf{u} be an arbitrary vector and let λ_{\max} be the eigenvalue with the largest absolute value. Let us define the following series of vectors:

$$\begin{cases} \mathbf{u}_0 &= \mathbf{u}, \\ \mathbf{u}_k &= \frac{1}{\lambda_{\max}} \mathbf{M} \cdot \mathbf{u}_{k-1} \quad \text{for } k > 0. \end{cases} \quad (8.49) \quad \Leftarrow \text{Main equation}$$

It is easy to prove¹² that:

$$\lim_{k \rightarrow \infty} \mathbf{u}_k = \mathbf{u}_{\max}, \quad (8.50)$$

where \mathbf{u}_{\max} is the eigenvector corresponding to λ_{\max} . Using this property, the following algorithm can be applied.

1. Set $\mathbf{u} = (1, 1, \dots, 1)$.
2. Set $\mathbf{u}' = \mathbf{M}\mathbf{u}$.
3. Set $\lambda = u'_1$, that is the first component of the vector \mathbf{u}' .
4. Set $\mathbf{u} = \frac{1}{\lambda} \mathbf{u}'$.
5. Check for convergence of λ . Go to step 2 if convergence is not yet attained.

¹⁰Eigenvalues are the roots of a polynomial of degree n . A double eigenvalue has two different eigenvectors.

¹¹If the vector \mathbf{u} is an eigenvector of the matrix \mathbf{M} with eigenvalue λ , so are all vectors $\alpha \mathbf{u}$ for any $\alpha \neq 0$.

¹²Hint: one must write the vector \mathbf{u} as a linear combination of the eigenvectors of the matrix \mathbf{M} . Such linear combination exists because the eigenvectors of a matrix form a complete system of reference.

The algorithm will converge toward λ_{\max} if the initial vector \mathbf{u} is not an eigenvector corresponding to a null eigenvalue of the matrix \mathbf{M} . If that is the case, one can chose another initial vector.

Once the eigenvalue with the largest absolute value has been found, the remaining eigenvalues can be found by replacing the matrix \mathbf{M} with the matrix:

$$\mathbf{M}' = \mathbf{M} \cdot (\mathbf{I} - \mathbf{u}_{\max} \otimes \mathbf{v}_{\max}), \quad (8.51)$$

where \mathbf{I} is the identity matrix of same dimension as the matrix \mathbf{M} and \mathbf{v}_{\max} is the eigenvector of the matrix \mathbf{M}^T corresponding to λ_{\max} ¹³. Using the fact that eigenvectors are orthogonal to each other, one can prove that the matrix \mathbf{M}' of equation 8.51 has the same eigenvalues as the matrix except for λ_{\max} which is replaced by 0. A complete proof of the above can be found in [Bass].

All eigenvalues and eigenvectors of the matrix \mathbf{M} can be found by repeating the process above n times. However, this works well only if the absolute values of the eigenvalues differ from each consecutive ones by at least an order of magnitude. Otherwise, the convergence of the algorithm is not very good. In practice, this algorithm can only be used to find the first couple of eigenvalues.

8.6.1 Finding the largest eigenvalue — General implementation

The object in charge of finding the largest eigenvalue is of course an instance of a subclass of the iterative process class described in 4.1. As the reader can see very few methods are required because most of the work is already implemented in the framework for iterative processes. The implementation is identical in both languages and will be discussed here. The largest eigenvalue finder has the following instance variables:

`matrix` the matrix whose largest eigenvalue is sought,

`eigenValue` the sought eigenvalue,

`eigenVector` the sought eigenvector and

`transposedEigenVector` the eigenvector of the transposed matrix.

The creation method takes the matrix as argument. Two accessor methods are supplied to retrieve the results, the eigenvalue and the eigenvector.

The method `initializeIterations` creates a vector to the matrix dimension and sets all its components equal to 1. As the algorithm progresses this vector will contain the eigenvector of the matrix. Similarly, a vector, which will contain the eigenvector of the transposed matrix, is created in the same way. In principle one should add a part to verify that this vector does not correspond to a null eigenvalue of the matrix. This small improvement is left as an exercise to the reader.

The algorithm is implemented within the single method `evaluateIteration` as described in section 4.1. The relative precision of the sought eigenvalue is the precision used to break out of the iterative process.

¹³The transpose of a matrix has the same eigenvalues, but not necessarily the same eigenvectors.

Figure 8.1 with the box LargestEigenValueFinder grayed.

8.6. MATRIX EIGENVALUES AND EIGENVECTORS OF A NON-SYMMETRIC MATRIX 277

Since the algorithm determines both the eigenvalue and the eigenvector the object in charge of the algorithm keeps both of them and must give access to both of them. Two accessor methods are supplied to retrieve the results, the eigenvalue and the eigenvector.

The largest eigenvalue finder is responsible to create the object responsible for finding the next eigenvalue when needed. Thus, the eigenvector of the transposed matrix is also computed along with the regular eigenvector. The method `nextLargestEigenValueFinder` returns a new instance of the class, which can be used to compute the next largest eigenvalue, by computing a new matrix as described in equation 8.51.

8.6.2 Finding the largest eigenvalue — Smalltalk implementation

Listing 8.13 shows the Smalltalk implementation of the class `DhbLargestEigenValueFinder`, subclass of the class `DhbIterativeProcess`.

The following code example shows how to use the class to find the first two largest eigenvalues of a matrix.

Code example 8.13

```
l m finder eigenvalue eigenvector nextFinder nextEigenvalue nextEigenvector l
m := DhbMatrix rows: #((84 -79 58 55)
                        (-79 84 -55 -58)
                        (58 -55 84 79)
                        (55 -58 79 84)).

finder := DhbLargestEigenValueFinder matrix: m.
eigenvalue := finder evaluate.
eigenvector := finder eigenvector.
nextFinder := finder nextLargestEigenValueFinder.
nextEigenvalue := nextFinder evaluate.
nextEigenvector := nextFinder eigenvector.
```

First the matrix `m` is defined from its components. Then, an instance of the class `DhbLargestEigenValueFinder` is created for this matrix. The iterative process is started as described in section 4.1.1. Its result is the eigenvalue. The eigenvector is retrieved using an accessor method. Then, a new instance of `DhbLargestEigenValueFinder` is obtained from the first one. The next largest eigenvalue and its eigenvector are retrieved from this new instance exactly as before.

Listing 8.13 Smalltalk implementation of the search for the largest eigenvalue

Class	DhbLargestEigenValueFinder
Subclass of	DhbIterativeProcess
Instance variable names:	matrix eigenvector transposeEigenvector

Class methods

`defaultMaximumIterations`

¹⁰⁰

matrix: aMatrix

^{(self new) initialize: aMatrix; yourself}

matrix: aMatrix precision: aNumber

<sup>(self new) initialize: aMatrix; desiredPrecision: aNumber;
yourself</sup>

Instance methods

eigenvalue

^{result}

eigenvector

^{eigenvector * (1 / eigenvector norm)}

evaluateIteration

```

| oldEigenvalue |
oldEigenvalue := result.
transposeEigenvector := transposeEigenvector * matrix.
transposeEigenvector := transposeEigenvector
    * (1 / (transposeEigenvector at: 1)).
eigenvector := matrix * eigenvector.
result := eigenvector at: 1.
eigenvector := eigenvector * (1 / result).
oldEigenvalue isNil
    ifTrue: [2 * desiredPrecision]
    ifFalse: [(result - oldEigenvalue) abs]

```

initialize: aMatrix

matrix := aMatrix.

initializeIterations

```

eigenvector := DhbVector new: matrix numberOfRows.
eigenvector atAllPut: 1.0.
transposeEigenvector := DhbVector new: eigenvector size.
transposeEigenvector atAllPut: 1.0

```

8.6. MATRIX EIGENVALUES AND EIGENVECTORS OF A NON-SYMMETRIC MATRIX 279

nextLargestEigenValueFinder

```
| norm |  
norm := 1 / (eigenvector * transposeEigenvector).  
^self class  
  new: matrix * ((DhbSymmetricMatrix identity: eigenvector  
                  size)  
                - (eigenvector * norm tensorProduct:  
                   transposeEigenvector))  
precision: desiredPrecision
```

8.6.3 Finding the largest eigenvalue — Java implementation

Listing 8.13 shows the Java implementation of the class `LargestEigenValueFinder`, subclass of class `IterativeProcess`.

The following code example shows how to use the class to find the first two largest eigenvalues of a matrix.

Code example 8.14

```
double[][] c = {{84,-79,58,55},{-79,84,-55,-58},  
               {58,-55,84,79},{55,-58,79,84}};  
Matrix m = new Matrix(c);  
LargestEigenValueFinder finder = new LargestEigenValueFinder( m);  
finder.evaluate();  
double eigenvalue = finder.eigenvalue();  
DhbVector eigenvector = finder.eigenvector();  
LargestEigenValueFinder nextFinder = finder.nextLargestEigenValueFinder();  
nextFinder.evaluate();  
double nextEigenvalue = nextFinder.eigenvalue();  
nextEigenvector = nextFinder.eigenvector();
```

First the matrix `m` is defined from its components. Then, an instance of the class `LargestEigenValueFinder` is created for this matrix. The iterative process is started as described in section 4.1.2. The eigenvalue and the corresponding eigenvector are retrieved using accessor methods. Then, a new instance of `LargestEigenValueFinder` is obtained from the first one. The new iterative process is started. The next largest eigenvalue and its eigenvector are retrieved from this new instance exactly as before.

Listing 8.14 Java implementation of the search for the largest eigenvalue

```
package DhbMatrixAlgebra;
```

```
import DhbIterations.*;  
/**
```

```
 * Object used to find the largest eigen value and the corresponding
```

```

    * eigen vector of a matrix by successive approximations.
    *
    * @author Didier H. Besset
    */
public class LargestEigenvalueFinder extends IterativeProcess
{
    /**
     * Eigenvalue
     */
    private double eigenvalue;
    /**
     * Eigenvector
     */
    private DhbVector eigenvector;
    /**
     * Eigenvector of transposed matrix
     */
    private DhbVector transposedEigenvector;
    /**
     * Matrix.
     */
    private Matrix matrix;

    /**
     * Constructor method.
     * @param prec double
     * @param a DhbMatrixAlgebra.Matrix
     */
    public LargestEigenvalueFinder ( double prec, Matrix a)
    {
        this(a);
        this.setDesiredPrecision ( prec);
    }
    /**
     * Constructor method.
     * @param a DhbMatrixAlgebra.Matrix
     */
    public LargestEigenvalueFinder ( Matrix a)
    {
        matrix = a;
        eigenvalue = Double.NaN;
    }
    /**
     * Returns the eigen value found by the receiver.
     * @return double
     */

```

8.6. MATRIX EIGENVALUES AND EIGENVECTORS OF A NON-SYMMETRIC MATRIX 281

```

public double eigenvalue ( )
{
    return eigenvalue;
}
/**
 * Returns the normalized eigen vector found by the receiver.
 * @return DhbMatrixAlgebra.DhbVector
 */
public DhbVector eigenvector ( )
{
    return eigenvector.product( 1.0 / eigenvector.norm());
}
/**
 * Iterate matrix product in eigenvalue information.
 */
public double evaluateIteration()
{
    double oldEigenvalue = eigenvalue;
    transposedEigenvector =
        transposedEigenvector.secureProduct( matrix);
    transposedEigenvector = transposedEigenvector.product( 1.0
        / transposedEigenvector.components[0]);
    eigenvector = matrix.secureProduct( eigenvector);
    eigenvalue = eigenvector.components[0];
    eigenvector = eigenvector.product( 1.0 / eigenvalue);
    return Double.isNaN( oldEigenvalue)
        ? 10 * getDesiredPrecision()
        : Math.abs( eigenvalue - oldEigenvalue);
}
/**
 * Set result to undefined.
 */
public void initializeIterations()
{
    eigenvalue = Double.NaN;
    int n = matrix.columns();
    double [] eigenvectorComponents = new double[ n];
    for ( int i = 0; i < n; i++) { eigenvectorComponents [i] = 1.0;}
    eigenvector = new DhbVector( eigenvectorComponents);
    n = matrix.rows();
    eigenvectorComponents = new double[ n];
    for ( int i = 0; i < n; i++) { eigenvectorComponents [i] = 1.0;}
    transposedEigenvector = new DhbVector( eigenvectorComponents);
}
/**
 * Returns a finder to find the next largest eigen value of the receiver's matrix.

```

```

    * @return DhbMatrixAlgebra.LargestEigenvalueFinder
    */
    public LargestEigenvalueFinder nextLargestEigenvalueFinder ( )
    {
        double norm = 1.0 / eigenvector.secureProduct(
            transposedEigenvector);
        DhbVector v1 = eigenvector.product( norm);
        return new LargestEigenvalueFinder( getDesiredPrecision(),
            matrix.secureProduct(SymmetricMatrix.identityMatrix(
                v1.dimension()).secureSubtract(v1.tensorProduct(
                    transposedEigenvector))));
    }
    /**
    * Returns a string representation of the receiver.
    * @return java.lang.String
    */
    public String toString()
    {
        StringBuffer sb = new StringBuffer();
        sb.append( eigenvalue);
        sb.append(" (");
        sb.append( eigenvector.toString());
        sb.append(')');
        return sb.toString();
    }
}

```

8.7 Matrix eigenvalues and eigenvectors of a symmetric matrix

In the nineteenth century Carl Jacobi discovered an efficient algorithm to find the eigenvalues of a symmetric matrix. Finding the eigenvalues of a symmetric matrix is easier since all eigenvalues are real.

In the section 8.6 we have mentioned that the eigenvectors of a matrix are orthogonal. Let $\mathbf{u}^{(1)}, \dots, \mathbf{u}^{(n)}$ the set of eigenvectors of the matrix \mathbf{M} such that $\mathbf{u}^{(i)} \cdot \mathbf{u}^{(i)} = 1$ for all i . Then, the matrix

$$\mathbf{O} = \begin{pmatrix} u_1^{(1)} & u_1^{(2)} & \dots & u_1^{(n)} \\ u_2^{(1)} & u_2^{(2)} & \dots & u_2^{(n)} \\ \vdots & \vdots & \ddots & \vdots \\ u_n^{(1)} & u_n^{(2)} & \dots & u_n^{(n)} \end{pmatrix}, \quad (8.52)$$

where $u_i^{(k)}$ is the i^{th} component of the k^{th} eigenvector, is an orthogonal¹⁴ matrix. That

¹⁴An orthogonal matrix of dimension n is a rotation in the n -dimensional space.

is, we have:

$$\mathbf{O}^T \cdot \mathbf{O} = \mathbf{I}. \quad (8.53)$$

Equation 8.53 is just another way of stating that the vectors $\mathbf{u}^{(1)}, \dots, \mathbf{u}^{(n)}$ are orthogonal to each other and are all normalized to 1. Combining this property with the definition of an eigenvector (equation 8.48) yields:

$$\mathbf{O}^T \cdot \mathbf{M} \cdot \mathbf{O} = \begin{pmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \lambda_n \end{pmatrix}, \quad (8.54)$$

where $\lambda_1, \dots, \lambda_n$ are the eigenvalues of the matrix \mathbf{M} .

The gist of Jacobi's algorithm is to apply a series of orthogonal transformations such that the resulting matrix is a diagonal matrix. It uses the fact that, for any orthogonal matrix \mathbf{R} , the matrix $\mathbf{R}^T \mathbf{M} \mathbf{R}$ has the same eigenvalues as the matrix \mathbf{M} . This follows from the definition of an eigenvector (equation 8.48) and the property of an orthogonal matrix (equation 8.53).

An orthogonal matrix corresponds to a rotation of the system of reference axes. Each step of Jacobi's algorithm is to find an rotation, which annihilates one of the off-diagonal elements of the matrix resulting from that orthogonal transformation. Let \mathbf{R}_1 be such matrix and let us define

$$\mathbf{M}_1 = \mathbf{R}_1^T \cdot \mathbf{M} \cdot \mathbf{R}_1. \quad (8.55)$$

Now, let us define the orthogonal transformation \mathbf{R}_2 , which annihilates one of the off-diagonal elements of the matrix \mathbf{M}_1 . The hope is that, after a certain number of steps m , the matrix

$$\begin{aligned} \mathbf{M}_m &= \mathbf{R}_m^T \cdot \mathbf{M}_{m-1} \cdot \mathbf{R}_m \\ &= \mathbf{R}_m^T \cdots \mathbf{R}_1^T \cdot \mathbf{M} \cdot \mathbf{R}_1 \cdots \mathbf{R}_m \end{aligned} \quad (8.56)$$

becomes a diagonal matrix. Then the diagonal elements of the matrix \mathbf{M}_m are the eigenvalues and the matrix

$$\mathbf{O}_m = \mathbf{R}_1 \cdots \mathbf{R}_m \quad (8.57)$$

is the matrix containing the eigenvectors.

Instead of annihilating just any diagonal element, one tries to annihilate the element with the largest absolute value. This ensures the fastest possible convergence of the algorithm. Let m_{kl} be the off-diagonal element of the matrix \mathbf{M} with the largest

absolute value. We define a matrix \mathbf{R}_1 with components:

$$\begin{cases} r_{kk}^{(1)} = \cos \vartheta, \\ r_{ll}^{(1)} = \cos \vartheta, \\ r_{kl}^{(1)} = -\sin \vartheta, \\ r_{lk}^{(1)} = \sin \vartheta, \\ r_{ii}^{(1)} = 1 \text{ for } i \neq k, l, \\ r_{ij}^{(1)} = 0 \text{ for } i \neq j, i \text{ and } j \neq k, l. \end{cases} \quad (8.58)$$

The reader can verify that the matrix \mathbf{R}_1 is an orthogonal matrix. The new matrix $\mathbf{M}_1 = \mathbf{R}_1^T \cdot \mathbf{M} \cdot \mathbf{R}_1$ has the same components as the matrix \mathbf{M} except for the rows and columns k and l . That is, we have

$$\begin{cases} m_{kk}^{(1)} = \cos^2 \vartheta m_{kk} + \sin^2 \vartheta m_{ll} - 2 \sin \vartheta \cos \vartheta m_{kl}, \\ m_{ll}^{(1)} = \sin^2 \vartheta m_{kk} + \cos^2 \vartheta m_{ll} + 2 \sin \vartheta \cos \vartheta m_{kl}, \\ m_{kl}^{(1)} = (\cos^2 \vartheta - \sin^2 \vartheta) m_{kl} + \sin \vartheta \cos \vartheta (m_{kk} - m_{ll}), \\ m_{ik}^{(1)} = \cos \vartheta m_{ik} - \sin \vartheta m_{il} \text{ for } i \neq k, l, \\ m_{il}^{(1)} = \cos \vartheta m_{il} + \sin \vartheta m_{ik} \text{ for } i \neq k, l, \\ m_{ij}^{(1)} = m_{ij} \text{ for } i \neq k, l \text{ and } j \neq k, l. \end{cases} \quad (8.59)$$

In particular, the angle of rotation can be selected such that $m_{kl}^{(1)} = 0$. That condition yields the following equation for the angle of the rotation:

$$\frac{\cos^2 \vartheta - \sin^2 \vartheta}{\sin \vartheta \cos \vartheta} = \frac{m_{ll} - m_{kk}}{m_{kl}} = \alpha, \quad (8.60)$$

where the constant α is defined by equation 8.60. Introducing the variable $t = \tan \vartheta$, equation 8.60 can be rewritten as:

$$t^2 + 2\alpha t - 1 = 0. \quad (8.61)$$

Since equation 8.61 is a second order equation, there are two solutions. To minimize rounding errors, it is preferable to select the solution corresponding to the smallest rotation [Press et al.]. The solution of equation 8.61 has already been discussed in section 1.3.4 for the case where α is positive. For any α , it can be written as:

$$t = \frac{\text{sign}(\alpha)}{|\alpha| + \sqrt{\alpha^2 + 1}}. \quad (8.62)$$

8.7. MATRIX EIGENVALUES AND EIGENVECTORS OF A SYMMETRIC MATRIX 285

In fact, the value of the angle ϑ does not need to be determined. We have:

$$\begin{cases} \cos \vartheta &= \frac{1}{\sqrt{t^2 + 1}}, \\ \sin \vartheta &= t \cos \vartheta. \end{cases} \quad (8.63)$$

Let us now introduce the quantities σ and τ defined as

$$\begin{cases} \sigma &= \sin \vartheta, \\ \tau &= \frac{\sin \vartheta}{1 + \cos \vartheta}. \end{cases} \quad (8.64)$$

Then equations 8.59 can be rewritten as

$$\begin{cases} m_{kk}^{(1)} &= m_{kk} - tm_{kl}, \\ m_{ll}^{(1)} &= m_{ll} + tm_{kl}, \\ m_{kl}^{(1)} &= 0, \\ m_{ik}^{(1)} &= m_{ik} - \sigma(m_{il} + \tau m_{ik}) \quad \text{for } i \neq k, l, \\ m_{il}^{(1)} &= m_{il} + \sigma(m_{ik} - \tau m_{il}) \quad \text{for } i \neq k, l, \\ m_{ij}^{(1)} &= m_{ij} \quad \text{for } i \neq k, l \text{ and } j \neq k, l. \end{cases} \quad (8.65) \quad \Leftarrow \text{Main equation}$$

Finally, we must prove that the transformation above did not increase the absolute values of the remaining off-diagonal elements of the matrix \mathbf{M}_1 . Using equations 8.59 the sum of the off-diagonal elements of the matrix \mathbf{M}_1 is:

$$\sum_{i \neq j} \left(m_{ij}^{(1)} \right)^2 = \sum_{i \neq j} m_{ij}^2 - 2m_{kl}^2. \quad (8.66)$$

Thus, this sum is always less than the sum of the squared off-diagonal elements of the matrix \mathbf{M} . In other words the algorithm will always converge.

Jacobi's algorithm

Now we have all the elements to implement Jacobi's algorithm. The steps are described hereafter:

1. Set the matrix \mathbf{M} to the matrix whose eigenvalues are sought.
2. Set the matrix \mathbf{O} to an identity matrix of the same dimension as the matrix \mathbf{M} .
3. Find the largest off-diagonal element, m_{kl} , of the matrix \mathbf{M} .
4. Build the orthogonal transformation \mathbf{R}_1 annihilating the element m_{kl} .
5. Build the matrix $\mathbf{M}_1 = \mathbf{R}_1^T \cdot \mathbf{M} \cdot \mathbf{R}_1$.

6. If $|m_{kl}|$ is less than the desired precision go to step 8.
7. Let $\mathbf{M} = \mathbf{M}_1$ and $\mathbf{O} = \mathbf{O} \cdot \mathbf{R}_1$; go to step 3.
8. The eigenvalues are the diagonal elements of the matrix \mathbf{M} and the eigenvectors are the rows of the matrix \mathbf{O} .

Strictly speaking, Jacobi's algorithm should be stopped if the largest off-diagonal element of matrix \mathbf{M}_1 is less than the desired precision. However, equation 8.66 guaranties that the largest off-diagonal element of the matrix after each step of Jacobi's algorithm is always smaller than the largest off-diagonal element of the matrix before the step. Thus, the stopping criteria proposed above can safely be used. This slight overkill prevents us from scanning the off-diagonal elements twice per step.

As the algorithm converges, α becomes very large. As discussed in section 1.3.4, the solution of equation 8.61 can be approximated with

$$t \approx \frac{1}{2\alpha}. \quad (8.67)$$

This expression is used when the computation of α^2 causes an overflow while evaluating equation 8.62.

8.7.1 Jacobi's algorithm — General implementation

Figure 8.1 with the box JacobiTransformation grayed.

Jacobi's algorithm is an iterative algorithm. The object implementing Jacobi's algorithm is a instance of the class JacobiTransform; it is a subclass of the iterative process discussed in section 4.1. The instance variables of this class are different in the two language implementations.

When an instance of the class JacobiTransform is created, the matrix whose eigenvalues are sought is copied into the matrix \mathbf{M} . This permits to use the same storage over the duration of the algorithm since equations 8.65 can be evaluated in place. Actually, only the upper half of the components needs to be stored since the matrix is a symmetric matrix.

The method evaluateIteration finds the largest off-diagonal element and performs the Jacobi step (equations 8.65) for that element. During the search for the largest off-diagonal element, the precision of the iterative process is set to the absolute value of the largest off-diagonal element. This is one example where it does not make sense to compute a relative precision. Actually, the precision returned by the method evaluateIteration is that of the previous iteration, but it does not really matter to make one iteration too much.

The method finalizeIterations performs a bubble sort to place the eigenvalues in decreasing order of absolute value. Bubble sorting is used instead of using a SortedCollection because one must also exchange the corresponding eigenvectors.

The result of the iterative process is an array containing the sorted eigenvalues plus the transformation matrix \mathbf{O} containing the eigenvectors. Extracting these results is language dependent.

8.7.2 Jacobi's algorithm — Smalltalk implementation

Listing 8.15 shows the Smalltalk implementation of Jacobi's algorithm.

The following code example shows how to use the class to find the eigenvalues and eigenvectors of a symmetric matrix.

Code example 8.15

```

| m jacobi eigenvalues eigenvectors |
m := DhbSymmetricMatrix rows: #((84 -79 58 55)
                                (-79 84 -55 -58)
                                (58 -55 84 79)
                                (55 -58 79 84)).

jacobi := DhbJacobiTransformation matrix: m.
eigenvalues := jacobi evaluate.
eigenvectors := jacobi transform columnsCollect: [ :each | each].

```

First the matrix *m* is defined from its components. Then, an instance of the class *DhbJacobiTransformation* is created for this matrix. The iterative process is started as described in section 4.1.1. Its result is an array containing the eigenvalues sorted in decreasing order. The corresponding eigenvectors are retrieved from the columns of the matrix *O* obtained from the method *transform*.

The class *DhbJacobiTransformation* has two instance variables

lowerRows an array of array containing the lower part of the matrix and

transform the components of the matrix *O*.

Since the matrix *M* is symmetric there is no need to keep all of its components. This not only reduces storage but also speeds up somewhat the algorithm because one only need to transform the lower part of the matrix.

The instance variable *result* contains the sorted eigenvalues at the end of the iterative process. The method *transform* returns the symmetric matrix *O* whose columns contain the eigenvectors in the same order. The code example shown at the beginning of this section shows how to obtain the eigenvectors from the matrix.

Listing 8.15 Smalltalk implementation of Jacobi's algorithm

Class	DhbJacobiTransformation
Subclass of	DhbIterativeProcess
Instance variable names:	lowerRows transform
Class methods	
matrix:	aSymmetricMatrix
^super new initialize:	aSymmetricMatrix
new	

```
^self error: 'Illegal creation message for this class'
```

Instance methods

evaluateIteration

```
| indices |
indices := self largestOffDiagonalIndices.
self transformAt: ( indices at: 1) and: ( indices at: 2).
^precision
```

exchangeAt: anInteger

```
| temp n |
n := anInteger + 1.
temp := result at: n.
result at: n put: ( result at: anInteger).
result at: anInteger put: temp.
transform do:
[ :each |
temp := each at: n.
each at: n put: ( each at: anInteger).
each at: anInteger put: temp.
].
```

finalizeIterations

```
| n |
n := 0.
result := lowerRows collect:
[ :each |
n := n + 1.
each at: n].
self sortEigenValues
```

initialize: aSymmetricMatrix

```
| n m |
n := aSymmetricMatrix numberOfRows.
lowerRows := Array new: n.
transform := Array new: n.
1 to: n do:
[ :k |
lowerRows at: k put: ( ( aSymmetricMatrix rowAt: k)
copyFrom: 1 to: k).
```

8.7. MATRIX EIGENVALUES AND EIGENVECTORS OF A SYMMETRIC MATRIX 289

```

transform at: k put: ( ( Array new: n) atAllPut: 0; at: k
                        put: 1; yourself).
].
^self

```

largestOffDiagonalIndices

```

| n m abs |
n := 2.
m := 1.
precision := ( ( lowerRows at: n) at: m) abs.
1 to: lowerRows size do:
    [ :i |
        1 to: ( i - 1) do:
            [ :j |
                abs := ( ( lowerRows at: i) at: j) abs.
                abs > precision
                ifTrue: [ n := i.
                        m := j.
                        precision := abs.
                        ].
            ].
    ].
^Array with: m with: n

```

printOn: aStream

```

| first |
first := true.
lowerRows do:
    [ :each |
        first ifTrue: [ first := false]
        ifFalse:[ aStream cr].
        each printOn: aStream.
    ].

```

sortEigenValues

```

| n bound m |
n := lowerRows size.
bound := n.
[ bound = 0 ]
whileFalse: [ m := 0.
              1 to: bound - 1 do:
                  [ :j |

```

```

      ( result at: j) abs > ( result at: j + 1)
      abs
      ifFalse:[ self exchangeAt: j.
        m := j.
      ].
    ].
    bound := m.
  ].

```

transform

[^]DhbMatrix rows: transform

transformAt: anInteger1 and: anInteger2

```

  | d t s c tau apq app aqq arp arq |
  apq := ( lowerRows at: anInteger2) at: anInteger1.
  apq = 0
    ifTrue: [ ^nil].
  app := ( lowerRows at: anInteger1) at: anInteger1.
  aqq := ( lowerRows at: anInteger2) at: anInteger2.
  d := aqq - app.
  arp := d * 0.5 / apq.
  t := arp > 0 ifTrue: [ 1 / ( ( arp squared + 1) sqrt + arp)]
    ifFalse:[ 1 / ( arp - ( arp squared + 1) sqrt)].
  c := 1 / ( t squared + 1) sqrt.
  s := t * c.
  tau := s / ( 1 + c).
  1 to: ( anInteger1 - 1)
    do: [ :r |
      arp := ( lowerRows at: anInteger1) at: r.
      arq := ( lowerRows at: anInteger2) at: r.
      ( lowerRows at: anInteger1) at: r put: ( arp - ( s *
        (tau * arp + arq))).
      ( lowerRows at: anInteger2) at: r put: ( arq + ( s *
        (arp - (tau * arq)))).
    ].
  ( anInteger1 + 1) to: ( anInteger2 - 1)
    do: [ :r |
      arp := ( lowerRows at: r) at: anInteger1.
      arq := ( lowerRows at: r) at: anInteger2.
      ( lowerRows at: r) at: anInteger1 put: ( arp - ( s *
        (tau * arp + arq))).
      ( lowerRows at: r) at: anInteger2 put: ( arq + ( s *
        (arp - (tau * arq)))).
    ].

```

8.7. MATRIX EIGENVALUES AND EIGENVECTORS OF A SYMMETRIC MATRIX 291

```

    ].
  ( anInteger2 + 1) to: lowerRows size
  do: [:r |
    arp := ( lowerRows at: r) at: anInteger1.
    arq := ( lowerRows at: r) at: anInteger2.
    ( lowerRows at: r) at: anInteger1 put: ( arp - ( s *
      (tau * arp + arq))).
    ( lowerRows at: r) at: anInteger2 put: ( arq + ( s *
      (arp - (tau * arq)))).
  ].
  1 to: lowerRows size
  do: [:r |
    arp := ( transform at: r) at: anInteger1.
    arq := ( transform at: r) at: anInteger2.
    ( transform at: r) at: anInteger1 put: ( arp - ( s *
      (tau * arp + arq))).
    ( transform at: r) at: anInteger2 put: ( arq + ( s *
      (arp - (tau * arq)))).
  ].
  ( lowerRows at: anInteger1) at: anInteger1 put: ( app - (t *
    apq)).
  ( lowerRows at: anInteger2) at: anInteger2 put: ( aqq + (t *
    apq)).
  ( lowerRows at: anInteger2) at: anInteger1 put: 0.

```

8.7.3 Jacobi's algorithm — Java implementation

Listing 8.16 shows the Java implementation of Jacobi's algorithm.

The following code example shows how to use the class to find the eigenvalues and eigenvectors of a symmetric matrix.

Code example 8.16

```

double[][] c = {{84,-79,58,55},{-79,84,-55,-58},
                {58,-55,84,79},{55,-58,79,84}};
SymmetricMatrix m;
try { m = SymmetricMatrix.fromComponents(c);
    } catch( DhbNonSymmetricComponents e) { return;}
    catch( DhbIllegalDimension e) { return;};
JacobiTransformation jacobi = new JacobiTransformation( m);
jacobi.evaluate().
double[] eigenvalues = jacobi.eigenvalues();
DhbVector[] eigenvectors = jacobi.eigenvectors();

```

First the matrix *m* is defined from its components. Then, an instance of the class *JacobiTransformation* is created for this matrix. The iterative process is started as

described in section 4.1.2. The eigenvalues and the corresponding eigenvectors are retrieved using accessor methods.

The class `JacobiTransformation` has four instance variables

`rows` an double dimensional array containing the components of the matrix,

`transform` the components of the matrix O ,

`p` the row index of the largest off-diagonal element and.

`q` the column index of the largest off-diagonal element.

Unlike Smalltalk, Java does not permit to have double dimensional arrays whose components are different sizes. Thus, one is forced to keep the components of the entire matrix.

The method `eigenvalues` returns an array containing the sorted eigenvalues and the method `eigenvectors` returns an array of vectors containing the corresponding eigenvectors.

Listing 8.16 Java implementation of Jacobi's algorithm

```
package DhbMatrixAlgebra;

import DhbIterations.IterativeProcess;
/**
 * JacobiTransformation
 *
 * @author Didier H. Besset
 */
public class JacobiTransformation extends IterativeProcess
{
    double[][] rows;
    double[][] transform;
    int p,q; //Indices of the largest off-diagonal element
/**
 * Create a new instance for a given symmetric matrix.
 * @param m DhbMatrixAlgebra.SymmetricMatrix
 */
    public JacobiTransformation (SymmetricMatrix m)
    {
        int n = m.rows();
        rows = new double[n][n];
        for ( int i = 0; i < n; i++)
        {
            for ( int j = 0; j < n; j++)
                rows[i][j] = m.components[i][j];
        }
    }
}
```

8.7. MATRIX EIGENVALUES AND EIGENVECTORS OF A SYMMETRIC MATRIX 293

```

/**
 * @return double[]
 */
public double[] eigenvalues ( )
{
    int n = rows.length;
    double[] eigenvalues = new double[n];
    for ( int i = 0; i < n; i++ )
        eigenvalues[i] = rows[i][i];
    return eigenvalues;
}
/**
 * @return DhbMatrixAlgebra.SymmetricMatrix
 */
public DhbVector[] eigenvectors ( )
{
    int n = rows.length;
    DhbVector[] eigenvectors = new DhbVector[n];
    double[] temp = new double[n];
    for ( int i = 0; i < n; i++ )
    {
        for ( int j = 0; j < n; j++ )
            temp[j] = transform[j][i];
        eigenvectors[i] = new DhbVector( temp );
    }
    return eigenvectors;
}
public double evaluateIteration()
{
    double offDiagonal = largestOffDiagonal();
    transform();
    return offDiagonal;
}
/**
 * @param m int
 */
private void exchange ( int m )
{
    int m1 = m + 1;
    double temp = rows[m][m];
    rows[m][m] = rows[m1][m1];
    rows[m1][m1] = temp;
    int n = rows.length;
    for ( int i = 0; i < n; i++ )
    {
        temp = transform[i][m];

```



```

        transform[i][m] = transform[i][m1];
        transform[i][m1] = temp;
    }
}
public void finalizeIterations ( )
{
    int n = rows.length;
    int bound = n - 1;
    int i, m;
    while ( bound >= 0 )
    {
        m = -1;
        for ( i = 0; i < bound; i++ )
        {
            if ( Math.abs( rows[i][i]) < Math.abs( rows[i+1][i+1]) )
            {
                exchange( i);
                m = i;
            }
        }
        bound = m;
    }
    return;
}
public void initializeIterations()
{
    transform = SymmetricMatrix.identityMatrix( rows.length).components;
}
/**
 * @return double   absolute value of the largest off diagonal element
 */
private double largestOffDiagonal( )
{
    double value = 0;
    double r;
    int n = rows.length;
    for (int i = 0; i < n; i++)
    {
        for ( int j = 0; j < i; j++)
        {
            r = Math.abs( rows[i][j]);
            if ( r > value )
            {
                value = r;
                p = i;
                q = j;
            }
        }
    }
}

```

8.7. MATRIX EIGENVALUES AND EIGENVECTORS OF A SYMMETRIC MATRIX 295

```

        }
    }
}
return value;
}
/**
 * Returns a string representation of the system.
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    char[] separator = { '[', ' ' };
    int n = rows.length;
    for ( int i = 0; i < n; i++)
    {
        separator[0] = '{';
        for ( int j = 0; j <= i; j++)
        {
            sb.append( separator);
            sb.append( rows[i][j]);
            separator[0] = ' ';
        }
        sb.append('}');
        sb.append('\n');
    }
    return sb.toString();
}
/**
 * @return DhbMatrixAlgebra.SymmetricMatrix
 */
private void transform ( )
{
    double apq = rows[p][q];
    if ( apq == 0 )
        return;
    double app = rows[p][p];
    double aqq = rows[q][q];
    double arp = ( aqq - app ) * 0.5 / apq;
    double t = arp > 0 ? 1 / ( Math.sqrt( arp * arp + 1 ) + arp )
        : 1 / ( arp - Math.sqrt( arp * arp + 1 ));
    double c = 1 / Math.sqrt( t * t + 1 );
    double s = t * c;
    double tau = s / ( 1 + c );
    rows[p][p] = app - t * apq;
    rows[q][q] = aqq + t * apq;

```

```
rows[p][q] = 0;
rows[q][p] = 0;
int n = rows.length;
for ( int i = 0; i < n; i++ )
{
    if ( i != p && i != q )
    {
        rows[p][i] = rows[i][p] - s * ( rows[i][q]
                                         + tau * rows[i][p] );
        rows[q][i] = rows[i][q] + s * ( rows[i][p]
                                         - tau * rows[i][q] );

        rows[i][p] = rows[p][i];
        rows[i][q] = rows[q][i];
    }
    arp = transform[i][p];
    aqq = transform[i][q];
    transform[i][p] = arp - s * ( aqq + tau * arp );
    transform[i][q] = aqq + s * ( arp - tau * aqq );
}
}
```

Chapter 9

Elements of statistics

La statistique est la première des sciences inexactes.¹
Edmond et Jules de Goncourt

Statistical analysis comes into play when dealing with a large amount of data. Obtaining information from the statistical analysis of data is the subject of chapter 10. Some sections of chapter 12 are also using statistics. Concepts needed by statistics are based on probability theory.

This chapter makes a quick overview of the concepts of probability theory. It is the third (and last) chapter of this book where most of the material is not useful per se. Figure 9.1 shows the classes described in this chapter. All these classes, however, are used extensively in the remaining chapters of this book. The example on how to use the code are kept to a minimum since real examples of use can be found in the next chapters.

An in-depth description of probability theory is beyond the scope of this book. The reader in the need for additional should consult the numerous textbooks on the subject, [Phillips & Taylor] or [Law & Kelton] for example.

9.1 Statistical moments

When one measures the values of an observable random variable, each measurement gives a different magnitude. Assuming measurement errors are negligible, the fluctuation of the measured values correspond to the distribution of the random variable. The problem to be solved by the experimenter is to determine the parameters of the distribution from the observed values. Statistical moments can contribute to the characterization of the distribution².

¹Statistics is the first of the inexact sciences.

²Central moments are related to the coefficients of the Taylor expansion of the Fourier transform of the distribution function.

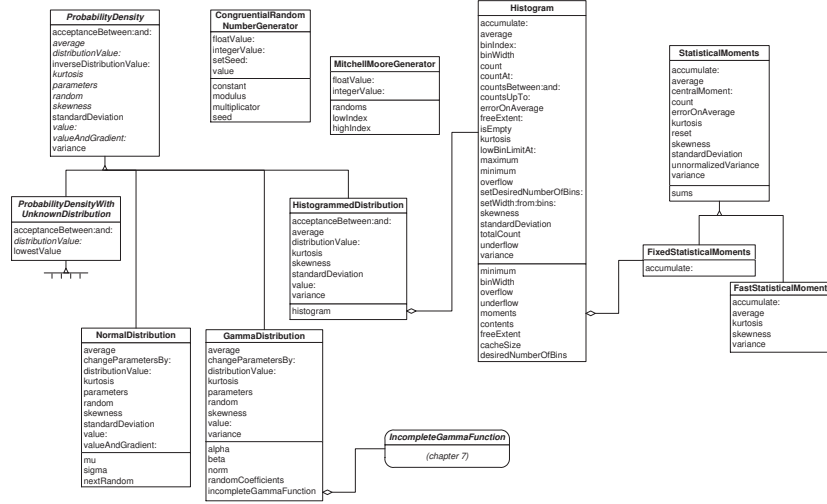


Figure 9.1: Classes related to statistics

Given a set of measurements, x_1, \dots, x_n , of the values measured for a random variable one defines the moment of k^{th} order by:

$$M_k = \frac{1}{n} \sum_{i=1}^n x_i^k. \quad (9.1)$$

In particular, the moment of first order is the mean or average of the set of data:

$$\bar{x} = M_1 = \frac{1}{n} \sum_{i=1}^n x_i. \quad (9.2)$$

The central moments of k^{th} order is defined by:

$$m_k = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^k. \quad (9.3)$$

where k is larger than 1. The central moments are easily expressed in terms of the moments. We have:

$$m_k = \sum_{j=0}^k \binom{k}{j} (-\bar{x})^{k-j} M_j, \quad (9.4)$$

where $\binom{k}{j}$ are the binomial coefficients.

Some statistical parameters are defined on the central moments. The variance of a set of measurement is the central moment of second order. The standard deviation, s , is the square root of the variance given by the following formula:

$$s^2 = \frac{n}{n-1} m_2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2. \quad (9.5)$$

The factor in front the central moment of second order is called Bessel's correction factor. This factor removes the bias of the estimation when the standard deviation is evaluated over a finite sample. The standard deviation measures the spread of the data around the average.

Many people believe that the standard deviation is the error of the average. This is not true: the standard deviation describes how much the data are spread around the average. It thus represents the error of a single measurement. An estimation of the standard deviation of the average value is given by the following formula:

$$s_{\bar{x}}^2 = \frac{s^2}{n} \quad \text{or} \quad s_{\bar{x}} = \frac{s}{\sqrt{n}}. \quad (9.6)$$

This expression must be taken as the error on the average when performing a least square fit on averaged data, for example.

Two quantities are related to the central moments of 3rd and 4th order. Each of these quantities are normalized by the adequate power of the standard deviation needed to yield a quantity without dimension.

The skewness is defined by:

$$a = \frac{n}{(n-1)(n-2)s^3} m_3 = \frac{1}{(n-1)(n-2)} \sum_{i=1}^n \left(\frac{x_i - \bar{x}}{s} \right)^3. \quad (9.7)$$

The skewness is a measure of the asymmetry of a distribution. If the skewness is positive, the observed distribution is asymmetric toward large values and vice-versa.

The kurtosis is defined by

$$\begin{aligned} k &= \frac{n(n+1)}{(n-1)(n-2)(n-3)s^4} m_4 - \frac{3(n-1)^2}{(n-2)(n-3)} \\ &= \frac{(n+1)}{(n-1)(n-2)(n-3)} \sum_{i=1}^n \left(\frac{x_i - \bar{x}}{s} \right)^4 - \frac{3(n-1)^2}{(n-2)(n-3)} \end{aligned} \quad (9.8)$$

The kurtosis is a measure of the peakedness or flatness of a distribution in the region of the average. The subtracted term in equation 9.8 is a convention defining the kurtosis of the normal distribution as 0³.

As we have seen, the average, standard deviation, skewness and kurtosis are parameters, which helps characterizing a distribution of observed values. To keep track of these parameters, it is handy to define an object whose responsibility is to accumulate the moments up to order 4. One can then easily compute the parameters of the distribution. It can be used in all cases where distribution parameters are needed.

9.1.1 Statistical moments — General implementation

To describe this implementation we must anticipated on the next section: the class `FastStatisticalMoments` implementing statistical moments as described in Section 9.1

Figure 9.1 with the box `FastStatisticalMoments` grayed.

³One talks about a platykurtic distribution when the kurtosis is negative, that is the peak of the distribution is flatter than that of the normal distribution. Student (c.f. section 10.2) and Cauchy (c.f. section D.2) distributions are platykurtic. The opposite is called leptokurtic. The Laplace (c.f. section D.5) distribution is leptokurtic.

is a subclass of the class defined in section 9.2.

Space allocation is handled by the superclass. The class `FastStatisticalMoments` uses this allocation to store the moments (instead of the central moments). The method `accumulate`: perform the accumulation of the moments. The methods `average`, `variance`, `skewness` and `kurtosis` compute the respective quantities using explicit expansion of the central moments as a function of the moments.

The computation of the standard deviation and of the error on the average are handled by the superclass (c.f. listings 9.3 and 9.5).

9.1.2 Statistical moments — Smalltalk implementation

Listing 9.1 shows the Smalltalk implementation. The class `DhbFastStatisticalMoments` is a subclass of class `DhbStatisticalMoments` presented in listing 9.3 of section 9.2.2. The reason for the split into two classes will become clear in section 9.2.

The following code shows how to use the class `DhbFastStatisticalMoments` to accumulate measurements of a random variable and to extract the various distribution parameters discussed in section 9.1.

Code example 9.1

```
| accumulator valueStream average stdev skewness kurtosis |
accumulator := DhbFastStatisticalMoments new.
[ valueStream atEnd]
  whileFalse:[ accumulator accumulate: valueStream next].
average := accumulator average.
stdev := accumulator standardDeviation.
skewness := accumulator skewness.
kurtosis := accumulator kurtosis.
```

This example assumes that the measurement of the random variable are obtained from a stream. The exact implementation of the stream is not shown here.

After the declarative statements, the first executable statement creates a new instance of the class `DhbFastStatisticalMoments` with the default dimension. This default allocates enough storage to accumulate up to the moment of 4th order. The next two lines are the accumulation proper using a `whileFalse:` construct and the general behavior of a stream. The last four lines extract the main parameters of the distribution.

If any of the distribution's parameters — average, variance, skewness or kurtosis — cannot be computed, the returned value is `nil`.

Listing 9.1 Smalltalk fast implementation of statistical moments

```
Class      DhbFastStatisticalMoments
Subclass of DhbStatisticalMoments
```

```
Instance methods
```

```
accumulate: aNumber
```

```

| var |
var := 1.
1 to: moments size
do:
  [:n |
    moments at: n put: (moments at: n) + var.
    var := var * aNumber]

```

average

```

self count = 0 ifTrue: [^nil].
^(moments at: 2) / self count

```

kurtosis

```

| var x1 x2 x3 x4 kFact kConst n m4 xSquared |
n := self count.
n < 4 ifTrue: [^nil].
var := self variance.
var = 0 ifTrue: [^nil].
x1 := (moments at: 2) / n.
x2 := (moments at: 3) / n.
x3 := (moments at: 4) / n.
x4 := (moments at: 5) / n.
xSquared := x1 squared.
m4 := x4 - (4 * x1 * x3) + (6 * x2 * xSquared) - (xSquared
    squared * 3).
kFact := n * (n + 1) / (n - 1) / (n - 2) / (n - 3).
kConst := 3 * (n - 1) * (n - 1) / (n - 2) / (n - 3).
^kFact * (m4 * n / var squared) - kConst

```

skewness

```

| x1 x2 x3 n stdev |
n := self count.
n < 3 ifTrue: [^nil].
stdev := self standardDeviation.
stdev = 0 ifTrue: [^nil].
x1 := (moments at: 2) / n.
x2 := (moments at: 3) / n.
x3 := (moments at: 4) / n.
^(x3 - (3 * x1 * x2) + (2 * x1 * x1 * x1)) * n * n
    / (stdev squared * stdev * (n - 1) * (n - 2))

```

variance


```

| n |
n := self count.
n < 2 ifTrue: [^nil].
^((moments at: 3) - ((moments at: 2) squared / n)) / (n - 1)

```

9.1.3 Statistical moments — Java implementation

Listing 9.2 shows the complete implementation in Java.

The following code shows how to use the class `FastStatisticalMoments` to accumulate measurements of a random variable and to extract the various distribution parameters discussed in section 9.1.

Code example 9.2

```

double[] values;

<Collecting measurements into the array values>

FastStatisticalMoments accumulator = new FastStatisticalMoments();
for ( int i = 0; i < values.length; i++)
    accumulator.accumulate( values[i]);
double average = accumulator.average();
double stdev = accumulator.standardDeviation();
double skewness = accumulator.skewness();
double kurtosis = accumulator.kurtosis();

```

This example assumes that the measurements of the random variable are collected into an array of `double`.

After collecting the measurements, an instance of the class `FastStatisticalMoments` is created with the default dimension. This default allocates enough storage to accumulate up to the moment of 4th order. The next two lines are the loop accumulating the values into the moments. The last four lines extract the main parameters of the distribution.

If any of the distribution's parameters — average, variance, skewness or kurtosis — cannot be computed, the returned value is `NaN`.

The computation of the standard deviation and of the error on the average are handled by the superclass (c.f. listing 9.5).

Listing 9.2 Java implementation of statistical moments

```

package DhbStatistics;

/**
 * Fast StatisticalMonents (at the cost of accuracy)
 *
 * @author Didier H. Besset
 */

```

```
public class FastStatisticalMoments extends StatisticalMoments
{
    /**
     * Default constructor method.
     */
    public FastStatisticalMoments()
    {
        super();
    }
    /**
     * Constructor method.
     * @param n int
     */
    public FastStatisticalMoments(int n)
    {
        super(n);
    }
    /**
     * Accumulate a random variable.
     * @param x value of the random variable.
     */
    public void accumulate( double x)
    {
        double value = 1.0;
        for ( int n=0; n < moments.length; n++)
        {
            moments[n] += value;
            value *= x;
        }
    }
    /**
     * @return double average.
     */
    public double average()
    {
        return moments[1] / moments[0];
    }
    /**
     * The kurtosis measures the sharpness of the distribution near the maximum.
     * Note: The kurtosis of the Normal distribution is 0 by definition.
     * @return double kurtosis.
     */
    public double kurtosis() throws ArithmeticException
    {
        if ( moments[0] < 4 )
            return Double.NaN;
    }
}
```

```

double x1 = average();
double x2 = moments[2] / moments[0];
double x3 = moments[3] / moments[0];
double x4 = moments[4] / moments[0];
double xSquared = x1 * x1;
double m4 = x4 - (4 * x1 * x3) + 3 * xSquared
               * (2 * x2 - xSquared);
double kFact = moments[0] * (moments[0] + 1)
               / ((moments[0] - 1) * (moments[0] - 2)
                  * (moments[0] - 3));
double kConst = 3 * (moments[0] - 1) * (moments[0] - 1)
                / ((moments[0] - 2) * (moments[0] - 3));
x4 = variance();
x4 *= x4;
return kFact * (m4 * moments[0] / x4) - kConst;
}
/**
 * @return double skewness.
 */
public double skewness() throws ArithmeticException
{
    if (moments[0] < 3)
        return Double.NaN;
    double x1 = average();
    double x2 = moments[2] / moments[0];
    double x3 = moments[3] / moments[0];
    double m3 = x3 + x1 * (2 * x1 * x1 - 3 * x2);
    x1 = standardDeviation();
    x2 = x1 * x1;
    x2 *= x1;
    return m3 * moments[0] * moments[0] / (x2 * (moments[0] - 1)
                                           * (moments[0] - 2));
}
/**
 * Unnormalized central moment of 2nd order
 * (needed to compute the t-test).
 * @return double
 */
public double unnormalizedVariance()
{
    double average = average();
    return moments[2] - average * average * moments[0];
}
/**
 * Note: the variance includes the Bessel correction factor.
 * @return double variance.

```

```

*/
public double variance() throws ArithmeticException
{
    if ( moments[0] < 2 )
        return Double.NaN;
    double average = average();
    return ( moments[0] / ( moments[0] - 1 ) )
        * ( moments[2] / moments[0] - average * average );
}
}

```

Note: The method `unnormalizedVariance` computes the expression $\sum_{i=1}^n (x_i - \bar{x})^2$ needed to compute the confidence level for the t -test (c.f. section 10.1).

9.2 Robust implementation of statistical moments

The methods used to implement the computation of the central moments in the previous section is prone to rounding errors. Indeed, contribution from values distant from the average can totally offset a result, however infrequent they are. Such an effect is worse when the central moments are derived from the moments. This section gives an algorithm ensuring minimal rounding errors.

The definition of statistical moments is based on the concept of expectation value. The expectation value is a linear operator over all functions of the random variable. If one measures the values of the random variable n times, the expectation value of a function $f(x)$ of a random variable x is estimated by the following expression:

$$\langle f(x) \rangle_n = \frac{1}{n} \sum_{i=1}^n f(x_i), \quad (9.9)$$

where the values x_1, \dots, x_n are the measurements of the random variable. A comparison of equation 9.9 with 9.2 shows that the average is simply the expectation value of the function $f(x) = x$. The central moment of order k is the expectation value of the function $(x - \bar{x})^k$:

$$\langle (x - \bar{x})^k \rangle_n = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^k. \quad (9.10)$$

To minimize rounding errors, one computes the changes occurring to the central moments when a new value is taken into account. In other words, one computes the value of a central moment over $n + 1$ values as a function of the central moment over

n values and the $(n + 1)^{\text{th}}$ value. For the average, we have

$$\begin{aligned}
 \langle x \rangle_{n+1} &= \frac{1}{n+1} \sum_{i=1}^{n+1} x_i \\
 &= \frac{x_{n+1}}{n+1} + \frac{1}{n+1} \sum_{i=1}^n x_i \\
 &= \frac{x_{n+1}}{n+1} + \frac{n}{n+1} \langle x \rangle_n \\
 &= \frac{x_{n+1}}{n+1} + \left(1 - \frac{1}{n+1}\right) \langle x \rangle_n \\
 &= \langle x \rangle_n - \frac{\langle x \rangle_n - x_{n+1}}{n+1}.
 \end{aligned} \tag{9.11}$$

Thus, the estimator of the average over $n + 1$ measurements can be computed from the estimator of the average over n measurements by subtracting a small correction, Δ_{n+1} , given by:

Main equation \Rightarrow

$$\begin{aligned}
 \Delta_{n+1} &= \langle x \rangle_n - \langle x \rangle_{n+1} \\
 &= \frac{\langle x \rangle_n - x_{n+1}}{n+1}.
 \end{aligned} \tag{9.12}$$

The expression in the numerator of equation 9.12 subtracts two quantities of comparable magnitude. This ensures a minimization of the rounding errors.

A similar derivation can be made for the central moments of higher orders. A complete derivation is given in appendix E.1. The final expression is

Main equation \Rightarrow

$$\left\langle (x - \bar{x})^k \right\rangle_{n+1} = \frac{n}{n+1} \left\{ \left[1 - (-n)^{k-1} \right] \Delta_{n+1}^k + \sum_{l=2}^k \binom{l}{k} \left\langle (x - \mu)^l \right\rangle_n \Delta_{n+1}^{k-l} \right\}. \tag{9.13}$$

The reader can verify the validity of equation 9.13 by verifying that it gives 1 for $k = 0$ and 0 for $k = 1$. Put in this form, the computation of the central moment estimators minimizes indeed rounding errors. For the central moment of order 2 we have:

Main equation \Rightarrow

$$\left\langle (x - \bar{x})^2 \right\rangle_{n+1} = \frac{n}{n+1} \left\{ (1+n) \Delta_{n+1}^2 + \left\langle (x - \bar{x})^2 \right\rangle_n \right\}. \tag{9.14}$$

For the central moment of order 3 we have:

Main equation \Rightarrow

$$\left\langle (x - \bar{x})^3 \right\rangle_{n+1} = \frac{n}{n+1} \left\{ (1-n^2) \Delta_{n+1}^3 + 3 \left\langle (x - \bar{x})^2 \right\rangle_n \Delta_{n+1} + \left\langle (x - \bar{x})^3 \right\rangle_n \right\}. \tag{9.15}$$

For the central moment of order 4 we have:

Main equation \Rightarrow

$$\begin{aligned}
 \left\langle (x - \bar{x})^4 \right\rangle_{n+1} &= \frac{n}{n+1} \left\{ (1+n^3) \Delta_{n+1}^4 + 6 \left\langle (x - \bar{x})^2 \right\rangle_n \Delta_{n+1}^2 \right. \\
 &\quad \left. + 4 \left\langle (x - \bar{x})^3 \right\rangle_n \Delta_{n+1} + \left\langle (x - \bar{x})^4 \right\rangle_n \right\}.
 \end{aligned} \tag{9.16}$$

9.2.1 Robust central moments — General implementation

The class `StatisticalMoments` has a single instance variable `moments` used to store the accumulated central moments.

The evaluation of equation 9.13 is not as hard as it seems from a programming point of view. One must remember that the binomial coefficients can be obtained by recursion (Pascal triangle). Furthermore, the terms of the sum can be computed recursively from those of the previous order so that raising the correction Δ_{n+1} to an integer power is never made explicitly. Equation 9.13 is implemented in method `accumulate`. The reader will notice that the binomial coefficients are computed inside the loop computing the sum.

Accumulating the central moments using equation 9.13 has the advantage that the estimated value of the central moment is always available. Nevertheless, accumulation is about 2 times slower than with the brute force method exposed in section 9.1. The reader must decide between speed and accuracy to choose between the two implementations.

The class `FixedStatisticalMoments` is a subclass of class `StatisticalMoments` specialized in the accumulation of central moments up to order 4. Instead of implementing the general equation 9.13, the central moments are accumulated using equations 9.14, 9.15 and 9.16. The only instance method redefined by this class is the method `accumulate`. All other computations are performed using the methods of the superclass.

9.2.2 Robust central moments — Smalltalk implementation

Listing 9.3 shows the implementation of the robust statistical moments. Listing 9.4 shows a specialization to optimize the speed of accumulation for the most frequently used case (accumulation up to the 4th order).

Using the class is identical for all classes of the hierarchy. Thus, the code example presented in section 9.1.2 is also valid for these two classes.

The creation method `new:` takes as argument the highest order of the accumulated moments. The corresponding initialization method allocates the required storage. The creation method `new` corresponds to the most frequent usage: the highest order is 4.

The methods computing the distribution parameters — average, variance, skewness and kurtosis — are using the method `centralMoment:` retrieving the central moment of a given order. They will return `nil` if not enough data as been accumulated in the moments.

Listing 9.3 Smalltalk implementation of accurate statistical moments

```

Class                DhbStatisticalMoments
Subclass of          Object
Instance variable names:  moments

Class methods
new
```

^self new: 4

new: anInteger

^super new initialize: anInteger + 1

Instance methods

accumulate: aNumber

```

| correction n n1 oldSums pascal nTerm cTerm term |
n := moments at: 1.
n1 := n + 1.
correction := ((moments at: 2) - aNumber) / n1.
oldSums := moments copyFrom: 1 to: moments size.
moments
    at: 1 put: n1;
    at: 2 put: (moments at: 2) - correction.
pascal := Array new: moments size.
pascal atAllPut: 0.
pascal
    at: 1 put: 1;
    at: 2 put: 1.
nTerm := -1.
cTerm := correction.
n1 := n / n1.
n := n negated.
3 to: moments size
do:
    [:k |
        cTerm := cTerm * correction.
        nTerm := n * nTerm.
        term := cTerm * (1 + nTerm).
        k to: 3
        by: -1
        do:
            [:l |
                pascal at: l put: (pascal at: l - 1) + (pascal
                    at: l).
                term := (pascal at: l) * (oldSums at: l) + term.
                oldSums at: l put: (oldSums at: l) * correction].
            pascal at: 2 put: (pascal at: 1) + (pascal at: 2).
            moments at: k put: term * n1]

```

average

```
self count = 0 ifTrue: [^nil].
^moments at: 2
```

```
centralMoment: anInteger
    ^moments at: anInteger + 1
```

```
count
    ^moments at: 1
```

```
errorOnAverage
    ^( self variance / self count) sqrt
```

```
initialize: anInteger
    moments := Array new: anInteger.
    self reset.
    ^self
```

```
kurtosis
    | n n1 n23 |
    n := self count.
    n < 4 ifTrue: [^nil].
    n23 := (n - 2) * (n - 3).
    n1 := n - 1.
    ^((moments at: 5) * n squared * (n + 1) / (self variance squared
                                                * n1)
       - (n1 squared * 3)) / n23
```

```
reset
    moments atAllPut: 0
```

```
skewness
    | n v |
    n := self count.
    n < 3 ifTrue: [^nil].
    v := self variance.
    ^((moments at: 4) * n squared / ((n - 1) * (n - 2) * (v sqrt * v))
```



```
standardDeviation
```

```
    ^self variance sqrt
```

```
unnormalizedVariance
```

```
    ^ ( self centralMoment: 2) * self count
```

```
variance
```

```
    | n |
    n := self count.
    n < 2
        ifTrue: [ ^nil].
    ^self unnormalizedVariance / ( n - 1)
```

The class `DhbFixedStatisticalMoments` is a specialization of the class `DhbStatisticalMoments` for a fixed number of central moments going up to the 4th order.

The class creation method `new:` is barred from usage as the class can only be used for a fixed number of moment orders. As a consequence the default creation method must be redefined to delegate the parametric creation to the method of the superclass.

Listing 9.4 Smalltalk implementation of accurate statistical moments with fixed orders

```
Class      DhbFixedStatisticalMoments
Subclass of DhbStatisticalMoments
```

```
    Class methods
```

```
new
```

```
    ^super new: 4
```

```
new: anInteger
```

```
    ^self error: 'Illegal creation message for this class'
```

```
    Instance methods
```

```
accumulate: aNumber
```

```

| correction n n1 c2 c3 |
n := moments at: 1.
n1 := n + 1.
correction := ((moments at: 2) - aNumber) / n1.
c2 := correction squared.
c3 := c2 * correction.
moments
  at: 5
    put: (((moments at: 5) + ((moments at: 4) * correction *
                                     4)
          + ((moments at: 3) * c2 * 6) + (c2 squared * (n
          squared * n + 1))))
        * n / n1;
  at: 4
    put: (((moments at: 4) + ((moments at: 3) * correction *
                                     3)
          + (c3 * (1 - n squared)))) * n
        / n1;
  at: 3 put: ((moments at: 3) + (c2 * (1 + n))) * n / n1;
  at: 2 put: (moments at: 2) - correction;
  at: 1 put: n1

```

9.2.3 Robust central moments — Java implementation

Listing 9.5 shows the implementation of the robust statistical moments. Listing 9.6 shows a specialization to optimize the speed of accumulation for the most frequently used case (accumulation up to the 4th order).

Using the class is identical for all classes of the hierarchy. Thus, the code example presented in section 9.1.3 is also valid for these two classes.

The class constructor method of class `StatisticalMoments` takes as argument the highest order of the accumulated moments and allocates the required storage. A class constructor method without argument creates an instance corresponding to the most frequent usage: the highest order is 4.

The methods computing the distribution parameters — average, variance, skewness and kurtosis — return `NaN` if the result is undefined for lack of measurements.

Listing 9.5 Java implementation of accurate statistical moments

```

package DhbStatistics;

/**
 * A StatisticalMoments accumulates statistical moments of a random variable.
 *
 * @author Didier H. Besset

```

```

*/
public class StatisticalMoments
{
    /**
     * Vector containing the points.
     */
    protected double[] moments;
    /**
     * Default constructor methods: declare space for 5 moments.
     */
    public StatisticalMoments()
    {
        this( 5);
    }
    /**
     * General constructor methods.
     * @param n number of moments to accumulate.
     */
    public StatisticalMoments( int n)
    {
        moments = new double[n];
        reset();
    }
    /**
     * @param x double value to accumulate
     */
    public void accumulate( double x)
    {
        double n = moments[0];
        double n1 = n + 1;
        double delta = ( moments[1] - x ) / n1;
        double[] sums = new double[ moments.length];
        sums[0] = moments[0];
        moments[0] = n1;
        sums[1] = moments[1];
        moments[1] -= delta;
        int[] pascal = new int[moments.length];
        pascal[0] = 1;
        pascal[1] = 1;
        double r1 = (double) n / (double) n1;
        double nk = -1;
        n = -n;
        double cterm = delta;
        double term;
        for ( int k = 2 ; k < moments.length; k++)
        {

```

```

        sums[k] = moments[k];
        nk = nk * n;
        cterm *= delta;
        term = (1 + nk) * cterm;
        for ( int l = k; l >= 2; l--)
        {
            pascal[l] += pascal[l-1];
            term += pascal[l] * sums[l];
            sums[l] *= delta;
        }
        pascal[1] += pascal[0];
        moments[k] = term * r1;
    }
}
/**
 * @return double average.
 */
public double average()
{
    return moments[1];
}
/**
 * Returns the number of accumulated counts.
 * @return number of counts.
 */
public long count()
{
    return (long) moments[0];
}
/**
 * Returns the error on average. May throw divide by zero exception.
 * @return error on average.
 */
public double errorOnAverage()
{
    return Math.sqrt( variance() / moments[0]);
}
/**
 * @return double    F-test confidence level with data accumulated
 *                  in the supplied moments.
 * @param m DhbStatistics.StatisticalMoments
 */
public double fConfidenceLevel( StatisticalMoments m)
{
    FisherSnedecorDistribution fDistr = new FisherSnedecorDistribution( (int) count(), (int) m.count());
    return fDistr.confidenceLevel( variance() / m.variance());
}

```

```

    }
    /**
     * The kurtosis measures the sharpness of the distribution near
     * the maximum.
     * Note: The kurtosis of the Normal distribution is 0 by definition.
     * @return double kurtosis or NaN.
     */
    public double kurtosis() throws ArithmeticException
    {
        if ( moments[0] < 4 )
            return Double.NaN;
        double kFact = ( moments[0] - 2 ) * ( moments[0] - 3 );
        double n1 = moments[0] - 1;
        double v = variance();
        return ( moments[4] * moments[0] * moments[0] * ( moments[0] + 1 )
                / ( v * v * n1 ) - n1 * n1 * 3 ) / kFact;
    }
    /**
     * Reset all counters.
     */
    public void reset( )
    {
        for ( int n=0; n < moments.length; n++)
            moments[n] = 0;
    }
    /**
     * @return double skewness.
     */
    public double skewness() throws ArithmeticException
    {
        if ( moments[0] < 3 )
            return Double.NaN;
        double v = variance();
        return moments[3] * moments[0] * moments[0]
            / ( Math.sqrt(v) * v * ( moments[0] - 1 )
              * ( moments[0] - 2 ));
    }
    /**
     * Returns the standard deviation. May throw divide by zero exception.
     * @return double standard deviation.
     */
    public double standardDeviation()
    {
        return Math.sqrt( variance());
    }
    /**

```

```

* @return double    t-test confidence level with data accumulated
*                  in the supplied moments.
* Approximation for the case where the variance of both sets may
*                  differ.
* @param m DhbStatistics.StatisticalMoments
*/
public double tApproximateConfidenceLevel( StatisticalMoments m)
{
    StudentDistribution tDistr = new StudentDistribution(
        (int) ( count()+m.count()-2));
    return tDistr.confidenceLevel( ( average() / standardDeviation()
        - m.average()
        / m.standardDeviation())
        / Math.sqrt(1/count()
        +1/m.count()));
}
/**
* @return double    t-test confidence level with data accumulated
*                  in the supplied moments.
* The variance of both sets is assumed to be the same.
* @param m DhbStatistics.StatisticalMoments
*/
public double tConfidenceLevel( StatisticalMoments m)
{
    int dof = (int) ( count()+m.count()-2);
    double sbar = Math.sqrt( ( unnormalizedVariance()
        + m.unnormalizedVariance()) / dof);
    StudentDistribution tDistr = new StudentDistribution( dof);
    return tDistr.confidenceLevel( ( average() - m.average())
        / ( sbar * Math.sqrt(1/count()
        +1/m.count())));
}
/**
* @return double
*/
public double unnormalizedVariance()
{
    return moments[2] * moments[0];
}
/**
* Note: the variance includes the Bessel correction factor.
* @return double variance.
*/
public double variance() throws ArithmeticException
{
    if ( moments[0] < 2 )

```

```

        return Double.NaN;
    return unnormalizedVariance() / ( moments[0] - 1);
}
}

```

Note: The method `tConfidenceLevel` implements the *t*-Test discussed in section 10.2. The method `fConfidenceLevel` implements the *F*-Test discussed in section 10.1.

The class `FixedStatisticalMoments` is a specialization of the class `StatisticalMoments` for a fixed number of central moments going up to the 4th order.

There is only one class constructor method, namely the one corresponding to the default number of orders.

Listing 9.6 Java implementation of accurate statistical moments with fixed order

package DhbStatistics;

```

/**
 * Statistical moments for a fixed set (1-4th order)
 *
 * @author Didier H. Besset
 */
public class FixedStatisticalMoments extends StatisticalMoments
{
    /**
     * Constructor method.
     */
    public FixedStatisticalMoments()
    {
        super();
    }
    /**
     * Quick implementation of statistical moment accumulation up to order 4.
     * @param x double
     */
    public void accumulate ( double x)
    {
        double n = moments[0];
        double n1 = n + 1;
        double n2 = n * n;
        double delta = ( moments[1] - x) / n1;
        double d2 = delta * delta;
        double d3 = delta * d2;
        double r1 = (double) n / (double) n1;
    }
}

```




Figure 9.2: A typical histogram

```

moments[4] += 4 * delta * moments[3] + 6 * d2 * moments[2]
              + (1 + n * n2) * d2 * d2;
moments[4] *= r1;
moments[3] += 3 * delta * moments[2] + (1 - n2) * d3;
moments[3] *= r1;
moments[2] += (1 + n) * d2;
moments[2] *= r1;
moments[1] -= delta;
moments[0] = n1;
return;
}
}

```

9.3 Histograms

Whereas statistical moments provides a quick way of obtaining information about the distribution of a measured random variable, the information thus provided is rather terse and quite difficult to interpret by humans. Histograms provide a more complete way of analyzing an experimental distribution. A histogram has a big advantage over statistical moments: it can easily be represented graphically. Figure 9.2 shows a typical histogram.

A histogram is defined by three main parameters: x_{\min} , the minimum of all values accumulated into the histogram, w , the bin width and n , the number of bins. A bin is defined as an interval. The i^{th} bin of a histogram is the interval $[x_{\min} + (i - 1)w, x_{\min} + iw[$.

The customary convention is that the lower limit is included in the interval and the higher limit excluded from the interval. The bin contents of a histogram — or histogram contents for short — is the number of times a value falls within each bin interval. Sometimes, a histogram is defined by the minimum and maximum of the accumulated values and the number of bins. The bin width is then computed as:

$$w = \frac{x_{\max} - x_{\min}}{n}, \quad (9.17)$$

where x_{\max} is the maximum of the accumulated values.

In section 10.10 we shall need the error on the contents of a histogram. In absence of any systematic effects⁴ the contents of each bin are distributed according to a Poisson distribution. The standard deviation of a Poisson distribution is the square root of the average. The standard deviation is used as an estimator of the error on the bin contents⁵. If n_i is the content of the i^{th} bin of the histogram, the estimated error on the contents is $\sqrt{n_i}$.

To obtain more information about the measured distribution, one can also keep track of the number of values falling outside of the histogram limits. The underflow of a histogram is defined as the number of values falling below the minimum of the accumulated values. Similarly, the overflow of a histogram is defined as the number of values falling on⁶ or above the maximum of the accumulated values.

9.3.1 Histograms — General implementation

Our implementation of histogram also accumulates the values into statistical moments. One can in principle compute the statistical moments of the measured distribution from the histogram contents. This determination, however, depends strongly on the bin width, especially if the bin width is large compared to the standard deviation. Thus, it is preferable to use the original data when accumulating the statistical moments. The net result is that a histogram has the same polymorphic behavior as a statistical moment.

When defining a histogram, the histogram limits — x_{\min} and x_{\max} — must be known in advance. This is not always practical since it implies a first scan of the measurements to determine the limits and a second scan to perform the accumulation into the defined histogram. Thus, our implementation offers the possibility of defining a histogram without predefined limits. In this mode, the first values are cached into an array until a sufficient number of data is available. When this happens, the histogram limits are determined from the data and the cached values are accumulated.

There are some cases when one would like to accumulate all the values within the histogram limits. The proposed implementation allows this by changing the histogram limits accordingly when a new value falls outside of the current histogram limits. When a histogram is accumulated in this mode the underflow and overflow counts are always zero.

⁴A good example of systematic effect is when values are computed from measurements made with an ADC. In this case, the integer rounding of the ADC may interfere with the bin sorting of the histogram.

⁵This is not a contradiction to what was said in section 9.1: the bin content is not an average, but a counting

⁶This is different from the definition of the underflow to be consistent with the fact that the definition of a bin interval is open ended at the upper limit.

Figure 9.1 with the box
Histogram grayed.

When the histogram limits are computed automatically, it can happen that these limits have odd values. For example, if the minimum value is 2.13456 and the maximum value is 5.1245, selecting a number of bins of 50 would yield a bin width of 0.0597988. Of course such value for the bin width is quite undesirable in practice. A similar thing can happen if the application creating the histogram obtains the minimum and maximum values from a computation or an automatic measuring device. To avoid such silly parameters, our implementation computes a reasonable limit and bin width by rounding the bin width to the nearest reasonable scale at the order of magnitude⁷ of the bin width. The possible scales are chosen to be easily computed by a human. In our example, the order of magnitude is -2 . The bin width is then selected to be 0.075 and the minimum and maximum are adjusted to be integral multiples of the bin width enclosing the given limits. In our example, there are 2.1 and 5.175 and the number of bins becomes 41 instead of 50.

9.3.2 Histograms — Smalltalk implementation

Listing 9.7 shows the implementation of a histogram in Smalltalk. The following code shows how to use the class `DhbHistogram` to accumulate measurements into a histogram.

Code example 9.3

```
| histogram valueStream |
  histogram := DhbHistogram new.
  [ valueStream atEnd ]
    whileFalse: [ histogram accumulate: valueStream next ].
    <printing or display of the histogram>
```

This example assumes that the measurement of the random variable are obtained from a stream. The exact implementation of the stream is not shown here.

After the declarative statements, the first executable statement creates a new instance of the class `DhbHistogram` with the default settings: automatic determination of the limits for 50 desired bins. The next two lines are the accumulation proper using a `whileFalse:` construct and the general behavior of a stream. This code is very similar to the code example presented in section 9.1.2. Extracting the parameters of the distribution can also be performed from the histogram.

The next example shows how to declare a histogram with given limits (2 and 7) and a desired number of bins of 50:

Code example 9.4

```
| histogram valueStream |
  histogram := DhbHistogram new.
  histogram setRangeFrom: 2.0 to: 7.0 bins: 100.
  <the rest is identical to the previous example>
```

The class `DhbHistogram` has the following instance variables:

`minimum` the minimum of the accumulated values, that is x_{\min} ,

⁷Let us recall that the order of magnitude is the power of ten of a number.

`binWidth` the bin width, that is w ,

`overflow` a counter to accumulate the overflow of the histogram,

`underflow` a counter to accumulate the underflow of the histogram,

`moments` an instance of the class `DhbFixedStatisticalMoments` to accumulate statistical moments up to the 4th order (c.f. section 9.2.2) with minimal rounding errors.

`contents` the contents of the histogram, that is an array of integers,

`freeExtent` a Boolean flag denoting whether the limits of the histogram can be adjusted to include all possible values,

`cacheSize` the size of the cache allocated to collect values for an automatic determination of the histogram limits,

`desiredNumberOfBins` the number of bins desired by the calling application.

Since there are many ways to declare a histogram, there is a single creation method `new`, which calls in turn a single standard initialization method `initialize`. In this mode, the histogram is created with undefined limits — that is, the first accumulated values are cached until a sufficient number is available for an automatic determination of the limits — and a default number of bins. The default number of bins is defined by the class method `defaultNumberOfBins`.

Four methods allow to change the default initialization.

The method `setRangeFrom:to:bins:` allows the definition of the parameters x_{\min} , x_{\max} and n , respectively. The method `setWidth:from:bins:` allows the definition of the parameters w , x_{\min} and n , respectively. In both cases, the histogram limits and number of bins are adjusted to reasonable values as explained at the end of section 9.3. These methods generate an error if the histogram is not cached, as limits cannot be redefined while the histogram is accumulating. The method `setDesiredNumberOfBins:` allows to overrule the default number of bins. Finally, the method `freeExtent:` takes a Boolean argument to define whether or not the limits must be adjusted when an accumulated value falls outside of the histogram limits. This method generates an error if any count has been accumulated in the underflow or overflow.

The method `accumulate` is used to accumulate the values. If the histogram is still cached — that is when values are directly accumulated into the cache for later determination of the limits — accumulation is delegated to the method `accumulateInCache:`. In this mode, the instance variable `contents` is an `OrderedCollection` collecting the values. When the size of the collection is reaching the maximum size allocated to the cache, limits are computed and the cache is flushed. In direct accumulation mode, the bin index corresponding to the value is computed. If the index is within the range, the value is accumulated. Otherwise it is treated like an overflow or an underflow. The method `processOverflows:for:` handles the case where the accumulated values falls outside of the histogram limits. If histogram limits cannot be adjusted it simply counts the overflow or the underflow. Otherwise processing is delegated to the methods `growsContents:`, `growsPositiveContents` and `growsNegativeContents`, which adjust the histogram limits according to the accumulated value.

The adjustment of the histogram limits to reasonable values is performed by the method `adjustDimensionUpTo:`. This is made when the limits are determined automatically. This method is also used when the limits are specified using one of the initialization methods.

There are many methods used to retrieve information from a histogram. Enumerating them here would be too tedious. Method names are explicit enough to get a rough idea of what each method is doing; looking at the code should suffice for a detailed understanding. The reader should just note that all methods retrieving the parameters of the distribution, as discussed in section 9.1.2, are implemented by delegating the method to the instance variable moments.

The iterator method `pointsAndErrorsDo:` is used for maximum likelihood fit of a probability distribution. Smalltalk implementation of maximum likelihood fit is discussed in section 10.10.2.

Listing 9.7 Smalltalk implementation of histograms

```

Class                DhbHistogram
Subclass of          Object
Instance variable names: minimum binWidth overflow underflow moments contents freeExtent
                        cacheSize desiredNumberOfBins

Class methods
defaultCacheSize
    ^100

defaultNumberOfBins
    ^50

integerScales

new
    ^super new initialize

scales

semiIntegerScales

```

Instance methods

accumulate: aNumber

```

| bin |
self isCached
  ifTrue: [ ^self accumulateInCache: aNumber].
bin := self binIndex: aNumber.
( bin between: 1 and: contents size)
  ifTrue: [ contents at: bin put: ( contents at: bin) + 1.
           moments accumulate: aNumber.
          ]
  ifFalse:[ self processOverflows: bin for: aNumber].

```

accumulateInCache: aNumber

```

contents add: aNumber.
contents size > cacheSize
  ifTrue: [ self flushCache].

```

adjustDimensionUpTo: aNumber

```

| maximum |
binWidth := self roundToScale: ( aNumber - minimum) /
                               desiredNumberOfBins.
minimum := ( minimum / binWidth) floor * binWidth.
maximum := ( aNumber / binWidth) ceiling * binWidth.
contents := Array new: ( ( maximum - minimum) / binWidth)
                      ceiling.
contents atAllPut: 0.

```

average

```

^moments average

```

binIndex: aNumber

```

^(( aNumber - minimum) / binWidth) floor + 1

```

binWidth

```

self isCached
  ifTrue: [ self flushCache].
^binWidth

```

collectIntegralPoints: aBlock

```

| answer bin lastContents integral norm x |
self isCached
  ifTrue: [ self flushCache].
answer := OrderedCollection new: ( contents size * 2 + 1).
bin := self minimum.
answer add: ( aBlock value: bin @ 0).
integral := self underflow.
norm := self totalCount.
contents do:
  [ :each |
    integral := integral + each.
    x := integral / norm.
    answer add: ( aBlock value: bin @ x).
    bin := bin + binWidth.
    answer add: ( aBlock value: bin @ x).
  ].
answer add: ( aBlock value: bin @ 0).
^answer asArray

```

collectPoints: aBlock

```

| answer bin lastContents |
self isCached
  ifTrue: [ self flushCache].
answer := OrderedCollection new: ( contents size * 2 + 1).
bin := self minimum.
answer add: ( aBlock value: bin @ 0).
contents do:
  [ :each |
    answer add: ( aBlock value: bin @ each).
    bin := bin + binWidth.
    answer add: ( aBlock value: bin @ each).
  ].
answer add: ( aBlock value: bin @ 0).
^answer asArray

```

count

^moments count

countAt: aNumber

| n |

```

n := self binIndex: aNumber.
^( n between: 1 and: contents size)
  ifTrue: [ contents at: n]
  ifFalse:[ 0]

```

countOverflows: anInteger

```

anInteger > 0
  ifTrue: [ overflow := overflow + 1]
  ifFalse:[ underflow := underflow + 1].

```

countsBetween: aNumber1 and: aNumber2

```

| n1 n2 answer |
n1 := self binIndex: aNumber1.
n2 := self binIndex: aNumber2.
answer := ( contents at: n1) * ( ( binWidth * n1 + minimum) -
                                aNumber1) / binWidth.

n2 > contents size
  ifTrue: [ n2 := contents size + 1]
  ifFalse:[ answer := answer + ( ( contents at: n2) * (
    aNumber2 - ( binWidth * ( n2 - 1) + self maximum)) / binWidth)].
(n1 + 1) to: ( n2 - 1) do: [ :n | answer := answer + ( contents
    at: n)].

^answer

```

countsUpTo: aNumber

```

| n answer |
n := self binIndex: aNumber.
n > contents size
  ifTrue: [ ^self count].
answer := ( contents at: n) * ( aNumber - ( binWidth * ( n - 1) +
    self minimum)) / binWidth.
1 to: ( n - 1) do: [ :m | answer := answer + ( contents at: m)].
^answer + underflow

```

errorOnAverage

```

^moments errorOnAverage

```

flushCache

```

| maximum values |
minimum isNil
  ifTrue: [ minimum := contents isEmpty ifTrue: [ 0]

              ifFalse:[ contents first].

            ].
maximum := minimum.
contents do:
  [ :each |
    each < minimum
      ifTrue: [ minimum := each]
      ifFalse:[ each > maximum
        ifTrue: [ maximum := each].
      ].
  ].
maximum = minimum
  ifTrue: [ maximum := minimum + desiredNumberOfBins].
values := contents.
self adjustDimensionUpTo: maximum.
values do: [ :each | self accumulate: each].

```

freeExtent: aBoolean

```

( underflow = 0 and: [ overflow = 0])
  ifFalse: [ self error: 'Histogram extent cannot be
                      redefined'].
freeExtent := aBoolean.

```

growContents: anInteger

```

anInteger > 0
  ifTrue: [ self growPositiveContents: anInteger]
  ifFalse:[ self growNegativeContents: anInteger].

```

growNegativeContents: anInteger

```

| n newSize newContents |
n := 1 - anInteger.
newSize := contents size + n.
newContents := Array new: newSize.
newContents
  at: 1 put: 1;
  replaceFrom: 2 to: n withObject: 0;
  replaceFrom: ( n + 1) to: newSize with: contents.
contents := newContents.

```



```
minimum := ( anInteger - 1 ) * binWidth + minimum.
```

```
growPositiveContents: anInteger
```

```
| n newContents |
n := contents size.
newContents := Array new: anInteger.
newContents
    replaceFrom: 1 to: n with: contents;
    replaceFrom: ( n + 1 ) to: ( anInteger - 1 ) withObject: 0;
    at: anInteger put: 1.
contents := newContents.
```

```
initialize
```

```
freeExtent := false.
cacheSize := self class defaultCacheSize.
desiredNumberOfBins := self class defaultNumberOfBins.
contents := OrderedCollection new: cacheSize.
moments := DhbFixedStatisticalMoments new.
overflow := 0.
underflow := 0.
^self
```

```
inverseDistributionValue: aNumber
```

```
| count x integral |
count := self count * aNumber.
x := self minimum.
integral := 0.
contents do:
    [ :each || delta |
        delta := count - integral.
        each > delta
            ifTrue: [ ^self binWidth * delta / each + x ].
        integral := integral + each.
        x := self binWidth + x.
    ].
^self maximum
```

```
isCached
```

```
^binWidth isNil
```

isEmpty

^false

kurtosis

^moments kurtosis

lowBinLimitAt: anInteger

^(anInteger - 1) * binWidth + minimum

maximum

self isCached
 ifTrue: [self flushCache].
 ^contents size * binWidth + minimum

maximumCount

self isCached
 ifTrue: [self flushCache].
 ^contents inject: (contents isEmpty ifTrue: [1] ifFalse:[
 contents at: 1])
 into: [:max :each | max max: each]

minimum

self isCached
 ifTrue: [self flushCache].
 ^minimum

overflow

^overflow

processOverflows: anInteger for: aNumber

freeExtent
 ifTrue: [self growContents: anInteger.
 moments accumulate: aNumber
]
 ifFalse:[self countOverflows: anInteger].

roundToScale: aNumber

setDesiredNumberOfBins: anInteger

anInteger > 0

ifFalse: [self error: 'Desired number of bins must be
positive'].

desiredNumberOfBins := anInteger.

setRangeFrom: aNumber1 to: aNumber2 bins: anInteger

self isCached

ifFalse: [self error: 'Histogram limits cannot be
redefined'].

minimum := aNumber1.

self setDesiredNumberOfBins: anInteger;

adjustDimensionUpTo: aNumber2.

setWidth: aNumber1 from: aNumber2 bins: anInteger

self isCached

ifFalse: [self error: 'Histogram limits cannot be
redefined'].

minimum := aNumber2.

self setDesiredNumberOfBins: anInteger;

adjustDimensionUpTo: (aNumber1 * anInteger + aNumber2).

skewness

^moments skewness

standardDeviation

^moments standardDeviation

totalCount

^moments count + underflow + overflow

underflow

^underflow

variance

^moments variance

9.3.3 Histograms — Java implementation

Listing 9.8 shows the implementation of a histogram in Java. The following code shows how to use the class `Histogram` to accumulate measurements of a random variable and to extract the various distribution parameters discussed in section 9.1.

Code example 9.5

```
double[] values;

    <Collecting measurements into the array values>

Histogram histogram = new Histogram();
for ( int i = 0; i < values.length; i++)
    histogram.accumulate( values[i]);

    <printing or display of the histogram>
```

This example assumes that the measurements of the random variable are collected into an array of double. As you can see, there is very little difference with the example of section 9.1.3. One can also extract the parameters of the distribution from the histogram in the exact same way as it was done for statistical moments.

A histogram can also be declared with given limits and a desired number of bins. If the limits are 2.0 and 7.0 and the desired number of bins is 50, the histogram declaration of the previous example must be replaced by the line:

```
Histogram histogram = new Histogram(2.0,7.0,50);
```

The class `Histogram` has the following instance variables:

`minimum` the minimum of the accumulated values, that is x_{\min} ,

`binWidth` the bin width, that is w ,

`contents` the contents of the histogram, that is an array of integers,

`growthAllowed` a Boolean flag denoting whether the limits of the histogram can be adjusted to include all possible values,

`integerBinWidth` a Boolean flag denoting whether the limits of the histogram must be adjusted to integer scales only,

`underflow` a counter to accumulate the underflow of the histogram,

`overflow` a counter to accumulate the overflow of the histogram,

`moments` an instance of the class `DhbFixedStatisticalMoments` to accumulate statistical moments up to the 4th order (c.f. section 9.2.2).

`cached` a Boolean flag denoting whether values are currently accumulated in a cache for later determination of the limits,

`cache` a cache to accumulate values.

There are six constructor methods for a histogram. We have used the fact that the type of the parameters uniquely defines the method. Thus, each constructor methods depends on a different combination of the parameters defining the limits:

`Histogram(double,double,int)` defines a histogram by specifying the parameters x_{\min} , x_{\max} and n in this order,

`Histogram(int,double,double)` defines a histogram by specifying the parameters n , x_{\min} and w in this order,

`Histogram(double,double)` defines a histogram by specifying the parameters x_{\min} and x_{\max} in this order; the number of bin is set to 50,

`Histogram(int,int)` defines a histogram with automatic determination of the limits by specifying the size of the cache and the number of bins n in this order,

`Histogram(int)` defines a histogram with automatic determination of the limits by specifying the size of the cache; the number of bin is set to 50,

`Histogram()` defines a histogram with automatic determination of the limits with a cache size of 100 and a number of bin of 50,

The method `setGrowthAllowed` sets the flag to allow the limits of the histogram to be adjusted to the accumulated values. A run time exception is thrown by this method if the histogram has already accumulated some counts in the underflow or the overflow.

The method `accumulate` is used to accumulate the values. If the histogram is still cached — that is when values are directly accumulated into the cache for later determination of the limits — the value is added to the cache. In this mode, the instance variable `underflow` is used to index the cache position. When the cache is full, limits are computed and the cache is flushed. The corresponding instance variable is reset to null. In direct accumulation mode, the value is compared to the histogram limits. If the value lies within the histogram limits, the value is accumulated. Otherwise it is treated like an overflow or an underflow. The methods `expandDown` and `expandUp` handle the case where the limits must be adjusted.

The adjustment of the histogram limits to reasonable values is performed by the methods `getBinParameters` and `defineParameters`. The former is used when the limits are determined automatically. The latter is used when the limits are specified using one of the initialization methods. The Java implementation gives the choice to force the limits to integer values. The scale values and final algorithm can be found in the class `DhbMath` (c.f. listing 1.2).

There are many methods used to retrieve information from a histogram. Enumerating them here would be too tedious. Method names are explicit enough to get a rough idea of what each method is doing; looking at the code should suffice for a detailed understanding. The reader should just note that all methods retrieving the parameters of the distribution, as discussed in section 9.1.3, are implemented by delegating the method to the instance variable `moments`.

```

package DhbScientificCurves;

import DhbInterfaces.PointSeries;
import DhbStatistics.FixedStatisticalMoments;
import DhbStatistics.ScaledProbabilityDensityFunction;
import DhbStatistics.ChiSquareDistribution;
import DhbFunctionEvaluation.DhbMath;
import DhbEstimation.WeightedPoint;
/**
 * A Histogram stores the frequency of hits into bins of equal width.
 *
 *
 * @author Didier H. Besset
 */
public class Histogram implements PointSeries
{
    /**
     * Lower limit of first histogram bin.
     */
    private double minimum;
    /**
     * Width of a bin.
     */
    private double binWidth;
    /**
     * Histogram contents.
     */
    private int[] contents;
    /**
     * Flag to allow automatical growth.
     */
    private boolean growthAllowed = false;
    /**
     * Flag to enforce integer bin width.
     */
    private boolean integerBinWidth = false;
    /**
     * Counts of values located below first bin.
     * Note: also used to count cached values when caching is in effect.
     */
    private int underflow;
    /**
     * Counts of values located above last bin.
     * Note: also used to hold desired number of bins when caching is in effect.
     */
    private int overflow;

```

```

/**
 * Statistical moments of values accumulated within the histogram limits.
 */
private FixedStatisticalMoments moments;
/**
 * Flag indicating the histogram is caching values to compute adequate range.
 */
private boolean cached = false;
/**
 * Cache for accumulated values.
 */
private double cache[];

/**
 * Constructor method with unknown limits and a desired number
 * of 50 bins. The first 100 accumulated values are cached.
 * Then, a suitable range is computed.
 */
public Histogram( )
{
    this( 100);
}
/**
 * Constructor method for approximate range for a desired number
 * of 50 bins.
 * All parameters are adjusted so that the bin width is a round number.
 * @param from approximate lower limit of first histogram bin.
 * @param to approximate upper limit of last histogram bin.
 */
public Histogram( double from, double to)
{
    this( from, to, 50);
}
/**
 * Constructor method for approximate range and desired number of bins.
 * All parameters are adjusted so that the bin width is a round number.
 * @param from approximate lower limit of first histogram bin.
 * @param to approximate upper limit of last histogram bin.
 * @param bins desired number of bins.
 */
public Histogram( double from, double to, int bins)
{
    defineParameters( from, to, bins);
}
/**

```

```

* Constructor method with unknown limits and a desired number of
* 50 bins.
* Accumulated values are first cached. When the cache is full,
* a suitable range is computed.
* @param n size of cache.
*/
public Histogram( int n)
{
    this( n, 50);
}
/**
* General constructor method.
* @param n number of bins.
* @param min lower limit of first histogram bin.
* @param width bin width (must be positive).
* @exception java.lang.IllegalArgumentException
*             if the number of bins is non-positive,
*             if the limits are inversed.
*/
public Histogram( int n, double min, double width)
    throws IllegalArgumentException
{
    if ( width <= 0 )
        throw new IllegalArgumentException(
            "Non-positive bin width: "+width);
    contents = new int[n];
    minimum = min;
    binWidth = width;
    reset();
}
/**
* Constructor method with unknown limits.
* Accumulated values are first cached. When the cache is full,
* a suitable range is computed.
* @param n size of cache.
* @param m desired number of bins
*/
public Histogram( int n, int m)
{
    cached = true;
    cache = new double[n];
    underflow = 0;
    overflow = m;
}
/**
* Accumulate a random variable.

```



```

    * @param x value of the random variable.
    */
    public void accumulate( double x)
    {
        if ( cached )
        {
            cache[ underflow++] = x;
            if ( underflow == cache.length)
                flushCache();
        }
        else if ( x < minimum )
        {
            if ( growthAllowed )
            {
                expandDown( x);
                moments.accumulate( x);
            }
            else
                underflow++;
        }
        else
        {
            int index = binIndex(x);
            if ( index < contents.length)
            {
                contents[ index]++;
                moments.accumulate( x);
            }
            else if ( growthAllowed )
            {
                expandUp( x);
                moments.accumulate( x);
            }
            else
                overflow++;
        }
    }
}
/**
 * Returns the average of the values accumulated in the histogram bins.
 * @return average.
 */
public double average()
{
    if ( cached )
        flushCache();
    return moments.average();
}

```

```

    }
    /**
     * @return int    index of the bin where x is located
     * @param x double
     */
    public int binIndex( double x)
    {
        return (int) Math.floor( (x - minimum) / binWidth);
    }
    /**
     * @param pdf DhbStatistics.ScaledProbabilityDensityFunction
     * @return double    chi2 of histogram compared to supplied
     *                   probability distribution.
     *
     */
    public double chi2Against( ScaledProbabilityDensityFunction pdf)
    {
        double chi2 = 0;
        for ( int i = 0; i < contents.length; i++ )
            chi2 += ( new WeightedPoint( 1,
                                         contents[i])).chi2Contribution( pdf);
        return chi2;
    }
    /**
     * @param pdf DhbStatistics.ScaledProbabilityDensityFunction
     * @return double    chi2 of histogram compared to supplied
     *                   probability distribution.
     *
     */
    public double chi2ConfidenceLevel(
        ScaledProbabilityDensityFunction pdf)
    {
        return (new ChiSquareDistribution(contents.length -
                                         pdf.parameters().length))
            .confidenceLevel(chi2Against(pdf));
    }
    /**
     * Returns the number of accumulated counts.
     * @return number of counts.
     */
    public long count()
    {
        return cached ? underflow : moments.count();
    }
    /**
     * Compute suitable limits and bin width.

```

```

* @param from approximate lower limit of first histogram bin.
* @param to approximate upper limit of last histogram bin.
* @param bins desired number of bins.
* @exception java.lang.IllegalArgumentException
*         if the number of bins is non-positive,
*         if the limits are inversed.
*/
private void defineParameters( double from, double to, int bins)
    throws IllegalArgumentException
{
    if ( from >= to )
        throw new IllegalArgumentException(
            "Inverted range: minimum = "+from+", maximum = "+to);
    if ( bins < 1)
        throw new IllegalArgumentException(
            "Non-positive number of bins: "+bins);
    binWidth = DhbMath.roundToScale( (to - from) / bins,
        integerBinWidth);
    minimum = binWidth * Math.floor( from / binWidth);
    int numberOfBins = (int) Math.ceil( ( to - minimum) / binWidth);
    if ( minimum + numberOfBins * binWidth <= to )
        numberOfBins++;
    contents = new int[numberOfBins];
    cached = false;
    cache = null;
    reset();
}
/**
* Returns the error on average. May throw divide by zero exception.
* @return error on average.
*/
public double errorOnAverage()
{
    if ( cached )
        flushCache();
    return moments.errorOnAverage();
}
/**
* Expand the contents so that the lowest bin include the specified
*         value.
* @param x value to be included.
*/
private void expandDown( double x)
{
    int addSize = (int) Math.ceil(( minimum - x) / binWidth);
    int newContents[] = new int[ addSize + contents.length];

```

```

    minimum -= addSize * binWidth;
    int n;
    newContents[0] = 1;
    for ( n = 1; n < addSize; n++ )
        newContents[n] = 0;
    for ( n = 0; n < contents.length; n++)
        newContents[n+addSize] = contents[n];
    contents = newContents;
}
/**
 * Expand the contents so that the highest bin include the specified
 *                                     value.
 * @param x value to be included.
 */
private void expandUp( double x)
{
    int newSize = (int) Math.ceil((x - minimum) / binWidth);
    int newContents[] = new int[ newSize];
    int n;
    for ( n = 0; n < contents.length; n++)
        newContents[n] = contents[n];
    for ( n = contents.length; n < newSize - 1; n++)
        newContents[n] = 0;
    newContents[n] = 1;
    contents = newContents;
}
/**
 * @return double    F-test confidence level with data accumulated
 *                  in the supplied histogram.
 * @param h DhbScientificCurves.Histogram
 */
public double fConfidenceLevel( Histogram h)
{
    return moments.fConfidenceLevel( h.moments);
}
/**
 * Flush the values from the cache.
 */
private void flushCache()
{
    double min = cache[0];
    double max = min;
    int cacheSize = underflow;
    double[] cachedValues = cache;
    int n;
    for ( n = 1; n < cacheSize; n++)

```

```

    {
        if ( cache[n] < min )
            min = cache[n];
        else if ( cache[n] > max )
            max = cache[n];
    }
    defineParameters( min, max, overflow);
    for ( n = 0; n < cacheSize; n++)
        accumulate( cachedValues[n]);
}
/**
 * @return int
 * @param x double
 */
public double getBinContent( double x)
{
    if ( x < minimum)
        return Double.NaN;
    int n = binIndex(x);
    return n < contents.length ? yValueAt(n) : Double.NaN;
}
/**
 * Returns the low and high limits and the content of the bin
 * containing the specified number or nul if the specified number
 * lies outside of the histogram limits.
 * @return a 3-dimensional array containing the bin limits and
 *         the bin content.
 */
public double[] getBinParameters( double x)
{
    if ( x >= minimum )
    {
        int index = (int) Math.floor( (x - minimum) / binWidth);
        if ( index < contents.length)
        {
            double[] answer = new double[3];
            answer[0] = minimum + index * binWidth;
            answer[1] = answer[0] + binWidth;
            answer[2] = contents[index];
            return answer;
        }
    }
    return null;
}
/**
 * Returns the bin width.

```

```

    * @return bin width.
    */
    public double getBinWidth()
    {
        return binWidth;
    }
    /**
    * @return double
    * @param x double
    * @param y double
    */
    public double getCountsBetween( double x, double y)
    {
        int n = binIndex(x);
        int m = binIndex(y);
        double sum = contents[n] * ( (minimum - x) / binWidth - (n+1))
            + contents[m] * ( (y - minimum) / binWidth - m);
        while ( ++n < m)
            sum += contents[n];
        return sum;
    }
    /**
    * @return double integrated count up to x
    * @param x double
    */
    public double getCountsUpTo( double x)
    {
        int n = binIndex( x);
        double sum = contents[n] * ( (x - minimum) / binWidth - n)
            + underflow;

        for ( int i = 0; i < n; i++)
            sum += contents[i];
        return sum;
    }
    /**
    * Returns the number of bins of the histogram.
    * @return number of bins.
    */
    public double getDimension()
    {
        if ( cached )
            flushCache();
        return contents.length;
    }
    /**
    * @return double

```

```

*/
public double getMaximum()
{
    return minimum + ( contents.length - 1 ) * binWidth;
}
/**
 * Returns the lower bin limit of the first bin.
 * @return minimum histogram range.
 */
public double getMinimum()
{
    return minimum;
}
/**
 * Returns the range of values to be plotted.
 * @return An array of 4 double values as follows
 * index 0: minimum of X range
 *      1: maximum of X range
 *      2: minimum of Y range
 *      3: maximum of Y range
 */
public double[] getRange( )
{
    if ( cached )
        flushCache();
    double[] range = new double[4];
    range[0] = minimum;
    range[1] = getMaximum();
    range[2] = 0;
    range[3] = 0;
    for ( int n = 0; n < contents.length; n++ )
        range[3] = Math.max( range[3], contents[n] );
    return range;
}
/**
 * Returns the kurtosis of the values accumulated in the histogram bins.
 * The kurtosis measures the sharpness of the distribution near the maximum.
 * Note: The kurtosis of the Normal distribution is 0 by definition.
 * @return double kurtosis.
 */
public double kurtosis()
{
    if ( cached )
        flushCache();
    return moments.kurtosis();
}

```

```

/**
 * @return FixedStatisticalMoments
 */
protected FixedStatisticalMoments moments()
{
    return moments;
}
/**
 * Returns the number of counts accumulated below the lowest bin.
 * @return overflow.
 */
public long overflow()
{
    return cached ? 0 : overflow;
}
/**
 * Reset histogram.
 */
public void reset()
{
    if ( moments == null )
        moments = new FixedStatisticalMoments();
    else
        moments.reset();
    underflow = 0;
    overflow = 0;
    for ( int n = 0; n < contents.length; n++ )
        contents[n] = 0;
}
/**
 * Allows histogram contents to grow in order to contain all
 * accumulated values.
 * Note: Should not be called after counts have been accumulated in
 * the underflow and/or overflow of the histogram.
 * @exception java.lang.RuntimeException
 * if the histogram has some contents.
 */
public void setGrowthAllowed() throws RuntimeException
{
    if ( underflow != 0 || overflow != 0 )
    {
        if ( !cached )
            throw new RuntimeException(
                "Cannot allow growth to a non-empty histogram");
    }
    growthAllowed = true;
}

```



```

}
/**
 * Forces the bin width of the histogram to be integer.
 * Note: Can only be called when the histogram is cached.
 * @exception java.lang.RuntimeException
 *         if the histogram has some contents.
 */
public void setIntegerBinWidth() throws RuntimeException
{
    if ( !cached)
        throw new RuntimeException(
            "Cannot change bin width of a non-empty histogram");
    integerBinWidth = true;
}
/**
 * Returns the number of points in the series.
 */
public int size()
{
    if ( cached )
        flushCache();
    return contents.length;
}
/**
 * Returns the skewness of the values accumulated in the histogram bins.
 * @return double skewness.
 */
public double skewness()
{
    if ( cached )
        flushCache();
    return moments.skewness();
}
/**
 * Returns the standard deviation of the values accumulated in the histogram bins.
 * @return double standard deviation.
 */
public double standardDeviation()
{
    if ( cached )
        flushCache();
    return moments.standardDeviation();
}
/**
 * @return double    t-test confidence level with data accumulated
 *                  in the supplied histogram.

```

```

    * @param h DhbScientificCurves.Histogram
    */
    public double tConfidenceLevel( Histogram h)
    {
        return moments.tConfidenceLevel( h.moments);
    }
    /**
    * @return long
    */
    public long totalCount()
    {
        return cached ? underflow
            : moments.count() + overflow + underflow;
    }
    /**
    * Returns the number of counts accumulated below the lowest bin.
    * @return underflow.
    */
    public long underflow()
    {
        return cached ? 0 : underflow;
    }
    /**
    * Returns the variance of the values accumulated in the histogram bins.
    * @return double variance.
    */
    public double variance()
    {
        if ( cached )
            flushCache();
        return moments.variance();
    }
    /**
    * @return DhbEstimation.WeightedPoint    corresponding to bin n.
    * @param n int
    */
    public WeightedPoint weightedPointAt( int n)
    {
        return new WeightedPoint( xValueAt( n), contents[n]);
    }
    /**
    * Returns the middle of the bin at the specified index.
    * @param index the index of the bin.
    * @return middle of bin
    */
    public double xValueAt( int index)

```

```

{
    return ( index + 0.5 ) * binWidth + minimum;
}
/**
 * Returns the content of the bin at the given index.
 * @param index the index of the bin.
 * @return bin content
 */
public double yValueAt( int index)
{
    if ( cached )
        flushCache();
    return ( index >= 0 && index < contents.length ) ? (double) contents[index] : 0;
}
}

```

Note: The method `tConfidenceLevel` implements the t -Test discussed in section 10.2. The method `fConfidenceLevel` implements the F -Test discussed in section 10.1.

9.4 Random number generator

When studying statistical processes on a computer one often has to simulate the behavior of a random variable⁸. As we shall see in section 9.5 it suffice to implement a random generator for a uniform distribution, that is a random variable whose probability density function is constant over a given interval. Once such an implementation is available, any probability distribution can be simulated.

Linear congruential random generators

The most widely used random number generators are linear congruential random generators. Random numbers are obtained from the following series [Knudth 2]:

$$X_{n+1} = (aX_n + c) \bmod m, \quad (9.18)$$

where m is called the modulus, a the multiplier and c the increment. By definition, we have $0 \leq X_n < m$ for all n . The numbers X_n are actually pseudo-random numbers since, for a given modulus, multiplier and increment, the sequence of numbers X_1, \dots, X_n is fully determined by the value X_0 . The value X_0 is called the seed of the series. In spite of its reproducibility the generated series behaves very close to that of random variable uniformly distributed between 0 and $m - 1$. Then the following variable:

$$x_n = \frac{X_n}{m}, \quad (9.19)$$

⁸Another wide use for random number generators are games.

is a random rational number uniformly distributed between 0 and 1, 1 excluded.

In practice, the modulus, multiplier and increment must be chosen quite carefully to achieve a good randomness of the series. Don Knuth [Knudth 2] gives a series of criteria for choosing the parameters of the random number generator. If the parameters are correctly chosen, the seed X_0 can be assigned to any value.

Additive sequence generators

Another class of random generators are additive sequence generators [Knudth 2]. The series of pseudo-random numbers is generated as follows:

$$X_n = (X_{n-l} + X_{n-k}) \bmod m, \quad (9.20)$$

where m is the modulus as before and l and k are two indices such that $l < k$. These indices must be selected carefully. [Knudth 2] contains a table of suitable indices. The initial series of numbers X_1, \dots, X_k can be any integers not all even.

Generators based on additive sequences are ideal to generate floating point numbers. If this case, the modulo operation on the modulus is not needed. Instead, one simply checks whether or not the newly generated number is larger than 1. Thus, the series becomes:

$$\begin{aligned} y_n &= x_{n-l} + x_{n-k}, \\ x_n &= \begin{cases} y_n & \text{if } y_n < 1, \\ y_n - 1 & \text{if } y_n \geq 1, \end{cases} \end{aligned} \quad (9.21)$$

It is clear that the evaluation above is much faster than that of equation 9.18. In practice, the additive sequence generator is about 4 times faster. In addition, the length of the sequence is larger than that of a congruential random generator with the same modulus.

In our implementation we have selected the pair of numbers (24, 55) corresponding to the generator initially discovered by G.J. Mitchell and D.P. Moore [Knudth 2]. The corresponding sequence has a length of $2^{55} - 1$. In our tests (c.f. below) we have found that the randomness of this generator is at least as good as that of the congruential random generator. The initial series x_1, \dots, x_{55} is obtained from the congruential random generator.

In [Knudth 2] Don Knuth describes a wealth of test to investigate the randomness of random number generators. Some of these tests are also discussed in [Law & Kelton]. To study the goodness of our proposed random generators, we have performed two types of tests: a χ^2 test and a correlation test.

The χ^2 test is performed on a histogram, in which values generated according to a probability distributions have been accumulated. Then, a χ^2 confidence level (c.f. section 10.3) is calculated against the theoretical curve computed using the histogram bin width, the number of generated values and the parameters of the distribution. A confidence level should be larger than 60% indicates that the probability distribution is correctly generated. When using distributions requiring the generation of several random numbers to obtain one value — Normal distribution (2 values), gamma distribution (2 values) and beta distribution (4 values) — one can get a good confidence that short

term correlations⁹ are not creating problems. The code for this test is given in the code examples 10.3 for Smalltalk and 10.4 for Java. In this test the Mitchell-Moore generator gives results similar to that of the congruential random generator.

The correlation test is performed by computing the covariance matrix (c.f. section 12.2) of vectors of given dimension (between 5 and 10). The covariance matrix should be a diagonal matrix with all diagonal elements equal to $1/12$, the variance of a uniform distribution. Deviation from this theoretical case should be small. Here longer correlations can be investigated by varying the dimension of the generated vectors. In this test too, the Mitchell-Moore generator gives results similar to that of the congruential random generator.

Bit-pattern generators

The generators described above are suitable to the generation of random values, but not for the generation of random bit patterns [Knudth 2], [Press et al.]. The generation of random bit patterns can be achieved with generator polynomials. Such polynomials are used in error correction codes for their abilities to produce sequences of numbers with a maximum number of different bits. For example the following polynomial

$$G(x) = x^{16} + x^{12} + x^5 + 1, \quad (9.22)$$

is a good generator for random patterns of 16 bits¹⁰. Of course, the evaluation of equation 9.22 does not require the computation of the polynomial. The following algorithm can be used:

1. Set X_{n+1} to X_n shifted by one position to the left and truncated to 16 bits ($X_{n+1} = 2X_n \bmod 2^{16}$),
2. If bit 15 (least significant bit being 0) of X_n is set, set X_{n+1} to the bit wise exclusive OR of X_{n+1} with 0x1021.

Other polynomials are given in [Press et al.].

Random bit patterns are usually used to simulate hardware behavior. They are rarely used in statistical analysis. A concrete implementation of a random bit pattern generator is left as an exercise to the reader.

9.4.1 Random number generator — Smalltalk implementation

Figure 9.1 with the boxes `CongruentialRandomNumberGenerator` and `MitchellMooreGenerator` grayed.

Listing 9.9 shows the implementation of a congruential random generator in Smalltalk. Listing 9.10 shows the implementation of an additive sequence random generator in Smalltalk. Listing 9.11 shows usage of the generator for standard use.

The class `DhbCongruentialRandomNumberGenerator` has three public methods:

⁹Pseudo random number generators have a tendency to exhibit correlations in the series. That is, the number X_n can be correlated to the number X_{n-k} for each n and a given k .

¹⁰c.f. O. Yamada, K. Yamazaki and D.H.Besset, An Error-Correction Scheme for an Optical Memory Card System, 1990 International Symposium on Information Theory and its Applications (ISITA'90), Hawaii, USA, November 27-30, 1990.

value returns the next random number of the series, that is X_n , a number between 0 and m ,

floatValue returns the next random floating number, that is the value X_n/m ,

integerValue: returns a random integer, whose values are between 0 and the specified argument.

When calling any of the above methods, the next random number of the series is obtained using equation 9.18.

There are several ways of using a random number generator. If there is no specific requirement the easiest approach is to use the instance provided by default creation method (new) returning a singleton. The next example shows how to proceed assuming the application uses the values generated by the random generator directly:

Code example 9.6

```
| generator x |
generator := DhbCongruentialRandomNumberGenerator new.

<Here is where the generator is used>

x := generator value.
```

If one needs several series which must be separately reproducible, one can assign several generators, one for each series. The application can use predefined numbers for the seed of each series. Here is a complete example assuming that the generated numbers must be floating numbers.

Code example 9.7 | generators seeds index x |

```
seeds := <an array containing the desired seeds>

generators := seeds collect:
    [ :each | DhbCongruentialRandomNumberGenerator seed: each].

<Here is where the various generators are used>
<index is the index of the desired series>

x := ( generators at: index) floatValue.
```

In game applications, it is of course not desirable to have a reproducible series. In this case, the easiest way is to use the time in milliseconds as the seed of the series. This initial value is sufficiently hard to reproduce to give the illusion of randomness. Furthermore the randomness of the series guarantees that two series generated at almost the same time are quite different. Here is how to do it.

Code example 9.8

```
| generator x |
generator := DhbCongruentialRandomNumberGenerator
    seed: Time millisecondClockValue.

<Here is where the generator is used>

x := (generator integerValue: 20) + 1.
```

In this last example, the generated numbers are integers between 1 and 20.

Implementation

The class `DhbCongruentialRandomNumberGenerator` has the following instance variables:

constant the increment c ,

modulus the modulus m ,

multiplicator the multiplier a and

seed the last generated number X_{n-1} .

There are three instance creation methods. The method `new` returns a singleton instance containing parameters from [Knuth 2]. The method `seed:` allows one to create an instance to generate a series of random number starting at a given seed. Finally the method `constant:multiplicator:modulus:` creates a congruential random number generator based on supplied parameters. Readers tempted to use this method are strongly advised to read [Knuth 2] and the references therein thoroughly. Then, they should perform tests to verify that their parameters are indeed producing a series with acceptable randomness.

The modulus of the standard parameters has 32 bits. In the Smalltalk implementation, however, the evaluation of equation 9.18 generates integers larger than 32 bits. As a result, the generation of the random numbers is somewhat slow, as it is using multiple precision integers. Using floating number¹¹ does not disturb the evaluation of equation 9.18 and is significantly faster since floating point evaluation is performed on the hardware. The generation of random numbers using floating point parameters is about 3 times faster than with integer parameters. This can easily be verified by the reader.

Listing 9.9 Smalltalk implementation of congruential random number generators

Class	DhbCongruentialRandomNumberGenerator
Subclass of	Object
Instance variable names:	constant modulus multiplicator seed
Class variable names:	UniqueInstance

Class methods

constant: aNumber1 multiplicator: aNumber2 modulus: aNumber3

^super new

initialize: aNumber1

multiplicator: aNumber2

modulus: aNumber3

new

¹¹The author is grateful to Dave N. Smith of IBM for this useful tip.

```

UniqueInstance isNil
  ifTrue: [ UniqueInstance := super new initialize.
           UniqueInstance setSeed: 1.
         ].
^UniqueInstance

seed: aNumber

  ^( super new) initialize; setSeed: aNumber; yourself

Instance methods
floatValue
  ^self value asFloat / modulus

initialize

  self initialize: 2718281829.0 multiplier: 3141592653.0
              modulus: 4294967296.0.

initialize: aNumber1 multiplier: aNumber2 modulus: aNumber3

  constant := aNumber1.
  modulus := aNumber2.
  multiplier := aNumber3.
  self setSeed: 1.

integerValue: anInteger

  ^( self value \ ( anInteger * 1000)) // 1000

setSeed: aNumber

  seed := aNumber.

value

  seed := ( seed * multiplier + constant) \ modulus.
^seed

```

The class `DhbMitchellMooreGenerator` implements a random number generator with additive sequence. It has two public methods:

`floatValue` returns the next random floating number, that is the value x_n ,

`integerValue:` returns a random integer, whose values are between 0 and the specified argument.

When calling any of the above methods, the next random number of the series is obtained using equation 9.21. The series of generated numbers are all floating points.

The creation methods `new` and `seed:` are used exactly as the corresponding methods of the class `DhbCongruentialRandomNumberGenerator`. Please refer to the code examples 9.6 and 9.7. Both methods use the congruential random number generator to generate the initial series of numbers x_1, \dots, x_{55} . The class method `constants:lowIndex:` offers a way to define the numbers k and l as well as the initial series of numbers. The reader wishing to use this method should consult the table of good indices k and l in [Knudth 2].

Listing 9.10 Smalltalk implementation of an additive sequence random number generator

```

Class                DhbMitchellMooreGenerator
Subclass of          Object
Instance variable names:  randoms lowIndex highIndex
Class variable names:    UniqueInstance

    Class methods
constants: anArray lowIndex: anInteger
^super new initialize: anArray lowIndex: anInteger

default
| congruentialGenerator |
congruentialGenerator := DhbCongruentialRandomNumberGenerator
                        new.
^self generateSeeds: congruentialGenerator

generateSeeds: congruentialGenerator

new
    UniqueInstance isNil
    ifTrue: [ UniqueInstance := self default].
^UniqueInstance

seed: anInteger

```

```

| congruentialGenerator |
congruentialGenerator := DhbCongruentialRandomNumberGenerator
                        seed: anInteger.
^self generateSeeds: congruentialGenerator

```

Instance methods

floatValue

```

| x |
x := (randoms at: lowIndex) + (randoms at: highIndex).
x < 1.0 ifFalse: [x := x - 1.0].
randoms at: highIndex put: x.
highIndex := highIndex + 1.
highIndex > randoms size ifTrue: [highIndex := 1].
lowIndex := lowIndex + 1.
lowIndex > randoms size ifTrue: [lowIndex := 1].
^x

```

initialize: anArray lowIndex: anInteger

```

randoms := anArray.
lowIndex := anInteger.
highIndex := randoms size.
^self

```

integerValue: anInteger

```

^( self floatValue * anInteger) truncated

```

For simple simulation, one wishes to generate a random number — floating or integer — within certain limits. Here are convenience methods implemented for the class `Number` and `Integer`. These methods frees the user from keeping track of the instance of the random number generator. For example, the following Smalltalk expression

```
50 random
```

generates an integer random number between 0 and 49 included. Similarly the following Smalltalk expression

```
2.45 random
```

generates a floating random number between 0 and 2.45 excluded. Finally the following Smalltalk expression

```
Number random
```

generates a floating random number between 0 and 1 excluded.

Listing 9.11 Smalltalk implementation of random number generators

Class **Integer**
 Subclass of Number

Instance methods

random

`^DhbMitchellMooreGenerator new integerValue: self`

Class **Number**
 Subclass of Magnitude

Class methods

random

`^DhbMitchellMooreGenerator new floatValue`

Instance methods

random

`^self class random * self`

Figure 9.1 with the box
 MitchellMooreGenerator
 grayed.

9.4.2 Random number generator — Java implementation

Java provides its own congruential random number generator, which is sufficient for most purposes. If one needs a better generator the Smalltalk code presented in listing 9.9 can be directly transposed into Java. Listing 9.12 shows the Java implementation of an additive sequence random generator.

The class MitchellMooreGenerator has three constructor methods.

`MitchellMooreGenerator()` returns an instance of the generator with standard parameters — $k = 55$ and $l = 24$ — with the initial series filled using the random number generator proposed by Java.

`MitchellMooreGenerator(int, int)` returns an instance of the generator with given parameters — k and l respectively — with the initial series filled using the random number generator proposed by Java.

`MitchellMooreGenerator(double[] ,int)` returns an instance of the generator by specifying the initial series given by the first argument — the size of the argument being the parameter k — and the parameter l given by the second argument.

The class `MitchellMooreGenerator` has two public methods. The method `nextDouble` returns a random floating number between 0 and 1 excluded. The method `nextIntInteger` returns a random integer between 0 and $n - 1$ where n is the specified argument. When calling any of these methods, the next random number of the series is obtained using equation 9.18.

Listing 9.12 Java implementation of an additive sequence random number generator

package DhbStatistics;

```
import java.util.Random;
/**
 * MitchellMoore random number generator
 *
 * @author Didier H. Besset
 */
public class MitchellMooreGenerator
{
    /**
     * List of previously generated numbers
     */
    private double[] randoms;
    /**
     * Index of last generated number
     */
    int highIndex;
    /**
     * Index of number to add to last number
     */
    int lowIndex;
    /**
     * Default constructor.
     */
    public MitchellMooreGenerator()
    {
        this(55,24);
    }
    /**
     * Constructor method.
     * @param seeds double[]
     * @param index int
     */
}
```

```

public MitchellMooreGenerator( double[] seeds, int index)
{
    highIndex = seeds.length;
    randoms = new double[ highIndex];
    System.arraycopy( seeds, 0, randoms, 0, --highIndex);
    lowIndex = index - 1;
}
/**
 * Constructor method.
 * @param indexH int    high index
 * @param indexL int    low index
 */
public MitchellMooreGenerator( int indexH, int indexL)
{
    Random generator = new Random();
    randoms = new double[indexH];
    for ( int i = 0; i < indexH; i++)
        randoms[i] = generator.nextDouble();
    highIndex = indexH - 1;
    lowIndex = indexL - 1;
}
/**
 * @return double    the next random number
 */
public double nextDouble()
{
    double x = randoms[ highIndex--] + randoms[ lowIndex--];
    if( highIndex < 0)
        highIndex = randoms.length - 1;
    if( lowIndex < 0)
        lowIndex = randoms.length - 1;
    return ( randoms[highIndex] = x < 1.0 ? x : x - 1);
}
/**
 * @return long    returns a long integer between 0 and n-1
 * @param n long
 */
public long nextInteger( long n)
{
    return (long) (n * nextDouble());
}
}

```

9.5 Probability distributions

A probability density function defines the probability of finding a continuous random variable within an infinitesimal interval. More precisely, the probability density function $P(x)$ gives the probability for a random variable to take a value lying in the interval $[x, x + dx[$. A probability density function is a special case of a one variable function described in section 2.1.

The moment of k^{th} order for a probability density function $P(x)$ is defined by:

$$M_k = \int x^k P(x) dx, \quad (9.23)$$

where the range of the integral is taken over the range where the function $P(x)$ is defined. By definition probability density functions are normalized, that is, M_0 is equal to 1.

As for statistical moments, one defines the mean or average of the distribution as:

$$\mu = M_1 = \int x P(x) dx. \quad (9.24)$$

Then the central moment of k^{th} order is defined by:

$$m_k = \int (x - \mu)^k P(x) dx. \quad (9.25)$$

In particular the variance is defined as m_2 , the central moment of second order. The standard deviation, σ , is the square root of m_2 . The skewness and kurtosis¹² of a probability density function are respectively defined as:

$$\omega = \frac{m_3}{\sqrt[3]{m_2}} = \frac{m_3}{\sigma^3} \quad \text{and} \quad (9.26)$$

$$\kappa = \frac{m_4}{m_2^2} - 3 = \frac{m_4}{\sigma^4} - 3. \quad (9.27)$$

The distribution function, also called acceptance function or repartition function, is defined as the probability for the random variable to have a value smaller or equal to a given value. For a probability density function defined over all real numbers we have:

$$F(t) = \text{Prob}(x < t) = \int_{-\infty}^t P(x) dx. \quad (9.28)$$

If the probability density function $P(x)$ is defined over a restricted range, the lower limit of the integral in equation 9.28 must be changed accordingly. For example, if the probability density function is defined for $x \geq x_{\min}$, the distribution function is given by:

$$F(t) = \text{Prob}(x < t) = \int_{x_{\min}}^t P(x) dx. \quad (9.29)$$

¹²In old references the kurtosis, as defined here, is called excess; then, the kurtosis is defined as the square of the excess; [Abramovitz & Stegun] e.g..

Instead of the distribution function, the name centile is sometimes used to refer to the value of the distribution function expressed in percent. This kind of terminology is frequently used in medical and natural science publications. For example, a value x is said to be at the 10th centile if $F(t) = 1/10$; in other words, there is a ten-percent chance of observing a value less than or equal to t ¹³.

The interval acceptance function measures the probability for the random variable to have a value between two given values. That is

$$F(x_1, x_2) = \text{Prob}(x_1 \leq x < x_2) = \int_{x_1}^{x_2} P(x) dx, \quad (9.30)$$

$$F(x_1, x_2) = F(x_2) - F(x_1). \quad (9.31)$$

If the integral of equation 9.28 can be resolved into a closed expression or if it has a numerical approximation, the evaluation of the interval acceptance function is made using equation 9.31. Otherwise the interval acceptance function must be evaluated using Romberg integration (c.f. section 6.4) applied to equation 9.30.

The inverse of the repartition function is frequently used. For example, in order to determine an acceptance threshold in a decision process, one needs to determine the variable t such that the repartition function is equal to a given value p . In other words, $t = F^{-1}(p)$. For example, the threshold of a coin detection mechanism to only reject 99.9% of the good coins is $F^{-1}(0.999)$. If the distribution function is not invertible, one can solve this equation using the Newton's zero-finding method exposed in section 5.3.2. Newton's zero-finding method is especially handy since the derivative of the function is known: it is the probability density function. Since $F(x)$ is strictly monotonous between 0 and 1 a unique solution is guaranteed for any p within the open interval $]0, 1[$. The initial value for the search can be obtained from Markov's inequality [Cormen et al.], which can be written in the form:

$$t \leq \frac{\mu}{1 - F(t)} \quad (9.32)$$

If no closed expression exists for the distribution function (it is determined using numerical integration e.g.) the computation of the inverse value is best obtained by interpolating the inverse function over a set of by tabulated values (c.f. section 3).

The inverse of the distribution function is also used to generate a random variable distributed according to the distribution. Namely, if r is a random variable uniformly distributed between 0 and 1, then the variable $x = F^{-1}(r)$ is a random variable distributed according to the distribution whose distribution function is $F(x)$. In practice this method can only be used if a closed expression exists for the distribution function, otherwise the function must be tabulated and Newton interpolation can be used on the inverse function (c.f. section 3.3). For most known distributions, however, special algorithms exist to generate random values distributed according to a given distribution. Such algorithms are described in [Law & Kelton]. They are not discussed here, but the code is included in each implementation of the specific probability distribution.

In the rest of this chapter we shall present the distributions used in this book. Other important distributions are presented in appendix D.

¹³A centile can also be quoted for a value relative to a set of observed values.

9.5.1 Probability distributions — General implementation

Table 9.1 shows the description of the public methods of the implementations of both languages.

Figure 9.1 with the boxes ProbabilityDensity and ProbabilityDensityWithUnknownDistr grayed.

Table 9.1: Public methods for probability density functions

Description	Smalltalk	Java	
$P(x)$	value:	value(double)	† x
$F(x)$	distributionValue:	distributionValue(double)	
$F(x_1, x_2)$	acceptanceBetween:and:	distributionValue(double,double)	
$F^{-1}(x)$	inverseDistributionValue:	inverseDistributionValue(double)	
x^\dagger	random	random()	
\bar{x}	average	average()	
σ^2	variance	variance()	
σ	standardDeviation	standardDeviation()	
skewness	skewness	skewness()	
kurtosis	kurtosis	kurtosis()	

represents the random variable itself. In other words, the method random returns a random value distributed according to the distribution.

Depending on the distribution, closed expressions for the variance or the standard deviation exist. Here general methods are supplied to compute one from the other. Subclasses must implement at least one of them; otherwise a stack overflow will result.

Methods to compute skewness and kurtosis are supplied, but return nil in Smalltalk and NaN in Java. A very general implementation could have used explicit integration. The integration interval, however, maybe infinite and a general integration strategy cannot easily be supplied. Subclasses are expected to implement both methods.

As we have quoted earlier a probability density function is a function, as defined in section 2.1. Since the distribution function is also a function, a Adapter must be provided to create a function (as defined in section 2.1) for the distribution function.

9.5.2 Probability distributions — Smalltalk implementation

Listing 9.13 shows the implementation of a general probability density distribution in Smalltalk. The class DhbProbabilityDensity is an abstract implementation. Concrete implementation of probability density distributions are subclass of it.

The method distributionValue: returning the value of the distribution function must be implemented by the subclass. The method to compute the interval acceptance functions is using equation 9.31 .

The inverse acceptance function is defined with two methods, one public and one private. The public method verifies the range of the argument, which must lie between 0 and 1. The private method uses the class DhbNewtonZeroFinder discussed in section 5.3.1. The derivative needed by the Newton zero finder is the probability density function itself since, by definition, it is the derivative of the acceptance function (c.f. equation 9.28).

Finally the class creation method `fromHistogram:` creates a new instance of a probability distribution with parameters derived using a quick approximation from the data accumulated into the supplied histogram; the derivation assumes that the accumulated data are distributed according to the distribution. This method is used to compute suitable starting values for least square or maximum likelihood fits (c.f. chapter 10). The convention is that this methods returns `nil` if the parameters cannot be obtained. Thus, returning `nil` is the default behavior for the superclass since this method is specific to each distribution. The estimation of the parameters is usually made using the statistical moments of the histogram and comparing them to the analytical expression of the distribution's parameter.

Listing 9.13 Smalltalk implementation of a probability distribution

```

Class      DhbProbabilityDensity
Subclass of Object

    Class methods
distributionName
    ^'Unknown distribution'

fromHistogram: aHistogram
    ^nil

    Instance methods
acceptanceBetween: aNumber1 and: aNumber2
    ^( self distributionValue: aNumber2) - ( self distributionValue:
                                           aNumber1)

approximatedValueAndGradient: aNumber
    | delta parameters dp gradient n |
    parameters := self parameters.
    n := parameters size.
    dp := self value: aNumber.
    delta := Array new: n.
    delta atAllPut: 0.
    gradient := DhbVector new: n.
    1 to: n do:
        [ :k |
            delta at: k put: ( parameters at: k) * 0.0001.
            self changeParametersBy: delta.

```

```

    gradient at: k put: ( ( ( self value: aNumber) - dp) / (
                        delta at: k)).
    delta at: k put: ( delta at: k ) negated.
    k > 1
    ifTrue: [ delta at: ( k - 1) put: 0].
  ].
  self changeParametersBy: delta.
  ^Array with: dp with: gradient

```

average

```

    self subclassResponsibility.

```

distributionFunction

```

    ^DhbProbabilityDistributionFunction density: self

```

distributionValue: aNumber

```

    ^self subclassResponsibility

```

inverseDistributionValue: aNumber

```

    ^( aNumber between: 0 and: 1)
    ifTrue: [ self privateInverseDistributionValue: aNumber]
    ifFalse:[ self error: 'Illegal argument for inverse
                        distribution value']

```

kurtosis

```

    ^nil

```

parameters

```

    ^self subclassResponsibility

```

printOn: aStream

```

    aStream nextPutAll: self class distributionName.
    self parameters ifNotNil: [ :params | | first |
    first := true.
    aStream nextPut: $(.
    params do:
    [ :each |

```

```

    first ifTrue: [ first := false]
      ifFalse:[ aStream nextPut: $,].
    aStream space.
    each printOn: aStream.
  ].
  aStream nextPut: $).
].

```

```
privateInverseDistributionValue: aNumber
```

```

^( DhbNewtonZeroFinder function: [ :x | ( self distributionValue:
      x) - aNumber] derivative: self)
  initialValue: self average / (1 - aNumber); evaluate

```

```
random
```

```

^self privateInverseDistributionValue: DhbMitchellMooreGenerator
  new floatValue

```

```
skewness
```

```
^nil
```

```
standardDeviation
```

```
^self variance sqrt
```

```
value: aNumber
```

```
self subclassResponsibility.
```

```
valueAndGradient: aNumber
```

```
^self approximatedValueAndGradient: aNumber
```

```
variance
```

```
^self standardDeviation squared
```

The class `DhbProbabilityDensityWithUnknownDistribution` is the abstract class for probability distribution having neither an analytical expression nor a numerical approximation for the distribution function.

Therefore, methods computing the acceptance function (`distributionValue:`) and interval acceptance (`acceptanceBetween:and:`) are using equations 9.28 and 9.30 respectively, using the class `DhbRombergIntegrator` discussed in section 6.4.2. The lower limit of the integral for the distribution function — x_{\min} of equation 9.29 — is defined by the method `lowestValue`. Since the majority of the probability density distributions are defined for non-negative numbers, this method returns 0. If the supplied default is not appropriate, the method `lowestValue` must be redefined by the subclass.

Listing 9.14 Smalltalk implementation of a probability distribution with unknown distribution function

Class **DhbProbabilityDensityWithUnknownDistribution**
 Subclass of `DhbProbabilityDensity`

Instance methods

`acceptanceBetween: aNumber1 and: aNumber2`

```
^(DhbRombergIntegrator new: self from: aNumber1 to: aNumber2)
   evaluate
```

`distributionValue: aNumber`

```
^(DhbRombergIntegrator new: self from: self lowestValue to:
   aNumber) evaluate
```

`lowestValue`

```
^0
```

Listing 9.15 shows the implementation of the Adapter for the distribution function. The class `DhbProbabilityDistributionFunction` has a single instance variable containing the corresponding probability density function. The creation method `density:` takes an instance of class `DhbProbabilityDensity` as argument.

Listing 9.15 Smalltalk implementation of a probability distribution function

Class **DhbProbabilityDistributionFunction**
 Subclass of `Object`

Instance variable names: `probabilityDensity`

Class methods

`density: aProbabilityDensity`

```
^self new initialize: aProbabilityDensity
```

```
Instance methods
```

```
initialize: aProbabilityDensity
```

```
probabilityDensity := aProbabilityDensity.  
^self
```

```
value: aNumber
```

```
^probabilityDensity distributionValue: aNumber
```

9.5.3 Probability distributions — Java implementation

Listings 9.13 and 9.17 show the implementation of two abstract classes implementing the general behavior of a probability density distribution in Java. Concrete probability density distributions are implemented as subclass of one of them.

The class `ProbabilityDensityFunction` assumes that the distribution function can be computed by the subclass, either using an analytical expression (Weibull distribution in section D.9.1 e.g.) or using a numerical approximation (Normal distribution in section 9.6 e.g.). Therefore, the method `distributionValue` with one argument, in charge of computing the distribution function, has been declared abstract. The interval acceptance function — implemented by the method `distributionValue` with two arguments — is implemented using equation 9.31.

The inverse acceptance function is defined with two methods, one public and one private. The public method verifies the range of the argument, which must lie between 0 and 1. It throws the exception `IllegalArgumentException` if this is not the case. The private method uses the class `NewtonZeroFinder` discussed in section 5.3.1. The derivative for the Newton zero finder is the instance itself since, by definition, the probability density distribution is the derivative of the acceptance function (c.f. equation 9.28).

Strangely enough, Java does not allow abstract constructor methods. For each subclass a constructor method using an histogram as argument must be provided to be consistent with the rest of the statistical analysis code. This method is the equivalent of the Smalltalk class creation method `fromHistogram:`, creating a new instance of a probability distribution with parameters derived using a quick approximation from the data accumulated into the supplied histogram; the derivation assumes that the accumulated data are distributed according to the distribution. This constructor method is used to compute suitable starting values for least square or maximum likelihood fits (c.f. chapter 10).

Listing 9.16 Java implementation of a probability distribution

```

package DhbStatistics;

import DhbFunctionEvaluation.DhbMath;
import DhbIterations.NewtonZeroFinder;
import DhbInterfaces.ParametrizedOneVariableFunction;
import java.util.Random;
/**
 * Subclasses of this class represent probability density function.
 * The value of the function  $f(x)$  represents the probability that a
 * continuous random variable takes values in the interval  $[x, x+dx[$ .
 * A norm is defined for the case where the function is overlaid over
 * a set of experimental points or a histogram.
 *
 * @author Didier H. Besset
 */
public abstract class ProbabilityDensityFunction
    implements ParametrizedOneVariableFunction
{
    /**
     * Random generator needed if random numbers are needed
     * (lazy initialization used).
     */
    private Random generator = null;

    /**
     * Compute an approximation of the gradient.
     * @return double[]
     * @param x double
     */
    public double[] approximateValueAndGradient( double x)
    {
        double temp, delta;
        double[] params = parameters();
        double[] answer = new double[ params.length + 1];
        answer[0] = value( x);
        for ( int i = 0; i < params.length; i++ )
        {
            temp = params[i];
            delta = Math.abs( temp) > DhbMath.defaultNumericalPrecision()
                ? 0.0001 * temp : 0.0001;
            params[i] += delta;
            setParameters( params);
            answer[i+1] = ( value(x) - answer[0]) / delta;
            params[i] = temp;
        }
        setParameters( params);
    }
}

```

```

        return answer;
    }
    /**
     * @return double average of the distribution.
     */
    public abstract double average ( );
    /**
     * Returns the probability of finding a random variable smaller than
     * or equal to x.
     * This method assumes that the probability density is 0 for  $x < 0$ .
     * If this is not the case, the subclass must implement this method.
     * @return integral of the probability density function from 0 to x.
     * @param x double upper limit of integral.
     */
    public abstract double distributionValue ( double x);
    /**
     * Returns the probability of finding a random variable between x1 and x2.
     * Computing is made using the method distributionValue(x).
     * This method should be used by distributions whose distributionValue
     * is computed using a method overriding the default one.
     * @return double integral of the probability density function from x1 to x2.
     * @param x1 double lower limit of integral.
     * @param x2 double upper limit of integral.
     */
    public double distributionValue ( double x1, double x2)
    {
        return distributionValue(x2) - distributionValue(x1);
    }
    /**
     * @return java.util.Random a random number generator.
     */
    protected Random generator( )
    {
        if ( generator == null)
            generator = new Random();
        return generator;
    }
    /**
     * @return double the value for which the distribution function is
     * equal to x.
     * @param x double value of the distribution function.
     * @exception java.lang.IllegalArgumentException
     *         if the argument is not between 0 and 1.
     */
    public double inverseDistributionValue ( double x) throws IllegalArgumentException
    {

```

```

    if ( x < 0 || x > 1)
        throw new IllegalArgumentException( "argument must be between 0 and 1");
    return privateInverseDistributionValue( x);
}
/**
 * @return double kurtosis of the distribution.
 */
public double kurtosis( )
{
    return Double.NaN;
}
/**
 * @return java.lang.String the name of the distribution.
 */
public abstract String name ( );
/**
 * This method assumes that the range of the argument has been checked.
 * Computation is made using the Newton zero finder.
 * @return double the value for which the distribution function
 *         is equal to x.
 * @param x double value of the distribution function.
 */
protected double privateInverseDistributionValue ( double x)
{
    OffsetDistributionFunction distribution = new OffsetDistributionFunction( this, x);
    NewtonZeroFinder zeroFinder = new NewtonZeroFinder( distribution, this, average());
    zeroFinder.setDesiredPrecision( DhbmMath.defaultNumericalPrecision());
    zeroFinder.evaluate();
    return zeroFinder.getResult();
}
/**
 * @return double a random number distributed according to the receiver.
 */
public double random ( )
{
    return privateInverseDistributionValue( generator().nextDouble());
}
/**
 * Set the seed of the random generator used by thr receiver.
 * @param seed long
 */
public void setSeed(long seed)
{
    generator().setSeed( seed);
    return;
}

```



```

/**
 * @return double skewness of the distribution.
 */
public double skewness( )
{
    return Double.NaN;
}
/**
 * NOTE: subclass MUST implement one of the two method variance
 *       or standardDeviation.
 * @return double standard deviation of the distribution from the variance.
 */
public double standardDeviation( )
{
    return Math.sqrt( variance());
}
/**
 * @return java.lang.String name and parameters of the distribution.
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    sb.append( name());
    char[] separator = { '(', ' ' };
    double[] parameters = parameters();
    for ( int i = 0; i < parameters.length; i++)
    {
        sb.append( separator);
        sb.append( parameters[i]);
        separator[0] = ',';
    }
    sb.append(')');
    return sb.toString();
}
/**
 * Evaluate the distribution and the gradient of the distribution with respect
 * to the parameters.
 * @return double[] 0: distribution's value, 1,2,...,n distribution's gradient
 * @param x double
 */
public double[] valueAndGradient( double x)
{
    return approximateValueAndGradient( x);
}
/**
 * NOTE: subclass MUST implement one of the two method variance

```

```

*      or standardDeviation.
* @return double variance of the distribution from the standard deviation.
*/
public double variance ( )
{
    double v = standardDeviation();
    return v*v;
}
}

```

Note: the interface `ParametrizedOneVariableFunction` implemented by the probability density function is needed for non-linear least square fits and maximum likelihood fits. It is described in section 10.9.3.

The class `ProbabilityDensityFunctionWithUnknownDistribution` is the superclass of probability distributions whose distribution function cannot be computed neither analytically nor numerically. The methods `distributionValue` with one or two arguments are computing the integral of equation 9.29 using the class `RombergIntegrator` discussed in section 6.4.3. The lower limit of the integral for the distribution function — x_{\min} of equation 9.29 — is defined by the abstract method `lowestValue`. This method must be implemented by the subclass.

Listing 9.17 Java implementation of a probability distribution with unknown distribution

```

package DhbStatistics;

import DhbFunctionEvaluation.DhbMath;
import DhbIterations.RombergIntegrator;
/**
 * Subclasses oare distributions whose distribution cannot be expressed
 * as a closed expression or have no numerical approximation.
 *
 * @author Didier H. Besset
 */
public abstract class ProbabilityDensityWithUnknownDistribution
    extends ProbabilityDensityFunction
{
    /**
     * Returns the probability of finding a random variable smaller than
     * or equal to x.
     * Computation is done using Romberg integration.
     * @return integral of the probability density function from lowValue() to x.
     * @param x double upper limit of intergral.

```

```

*/
public double distributionValue ( double x)
{
    RombergIntegrator integrator = new RombergIntegrator( this, lowValue(), x);
    integrator.setDesiredPrecision( DhbMath.defaultNumericalPrecision());
    integrator.evaluate();
    return integrator.getResult();
}
/**
 * Returns the probability of finding a random variable between x1 and x2.
 * @return double integral of the probability density function from x1 to x2.
 * @param x1 double lower limit of integral.
 * @param x2 double upper limit of integral.
 */
public double distributionValue ( double x1, double x2)
{
    RombergIntegrator integrator = new RombergIntegrator( this, x1, x2);
    integrator.setDesiredPrecision(
        DhbMath.defaultNumericalPrecision());
    integrator.evaluate();
    return integrator.getResult();
}
/**
 * @return double    lower limit of the integral used to compute
 *                    the distribution function
 */
protected abstract double lowValue();
}

```

Listing 9.15 shows the implementation of the Adapter for the distribution function. The class `OffsetDistributionFunction` has the following instance variables:

`probabilityDensity` an instance of class `ProbabilityDensityFunction`,

`offset` a value subtracted to the distribution function.

The constructor method takes as arguments an instance of class `ProbabilityDensityFunction` and a double. Instances of this class are used to compute the inverse value of the distribution function. It can also be used to implement the distribution function as described in section 2.1.2 using an offset of 0. If parameters cannot be estimated, the constructor method throws the exception `IllegalArgumentException`. The estimation of the parameters is usually made using the statistical moments of the histogram and comparing them to the analytical expression of the distribution's parameter.

Listing 9.18 Java implementation of a probability distribution function

```
package DhbStatistics;

import DhbInterfaces.OneVariableFunction;
/**
 * This class is used to find the inverse distribution function of
 * a probability density function.
 *
 * @author Didier H. Besset
 */
public final class OffsetDistributionFunction
    implements OneVariableFunction
{
    /**
     * Probability density function.
     */
    private ProbabilityDensityFunction probabilityDensity;
    /**
     * Value for which the inverse value is desired.
     */
    private double offset;

    /**
     * Create a new instance with given parameters.
     * @param p statistics.ProbabilityDensityFunction
     * @param x double
     */
    protected OffsetDistributionFunction ( ProbabilityDensityFunction p,
                                           double x)
    {
        probabilityDensity = p;
        offset = x;
    }
    /**
     * @return distribution function minus the offset.
     */
    public double value(double x)
    {
        return probabilityDensity.distributionValue( x ) - offset;
    }
}
```

9.6 Normal distribution

The normal distribution is the most important probability distribution. Most other distributions tend toward it when some of their parameters become large. Experimental data subjected only¹⁴ to measurement fluctuation usually follow a normal distribution. Table 9.2 shows the properties of the normal distribution. Figure 9.3 shows the well-known

Table 9.2: Properties of the Normal distribution

Range of random variable	$]-\infty, +\infty[$
Probability density function	$P(x) = \frac{1}{\sqrt{2\pi}\sigma^2} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$
Parameters	$-\infty < \mu < +\infty$ $0 < \sigma < +\infty$
Distribution function	$F(x) = \text{erf}\left(\frac{x-\mu}{\sigma}\right)$ (c.f. section 2.3)
Average	μ
Variance	σ^2
Skewness	0
Kurtosis	0

bell shape of the normal distribution for various values of the parameters. The reader can see that the peak of the distribution is always located at μ and that the width of the bell curve is proportional to σ .

Figure 9.1 with the box NormalDistribution grayed.

9.6.1 Normal distribution — Smalltalk implementation

Listing 9.19 shows the implementation of the normal distribution in Smalltalk.

The distribution function of the normal distribution can be computed with the error function (c.f. section 2.3). Therefore the class DhbNormalDistribution is implemented as a subclass of DhbProbabilityDensity.

Listing 9.19 Smalltalk implementation of the normal distribution

```

Class          DhbNormalDistribution
Subclass of    DhbProbabilityDensity

```

¹⁴The presence of systematic errors is a notable exception to this rule.




Figure 9.3: Normal distribution for various values of the parameters

Instance variable names: `mu` `sigma` `nextRandom`

Class variable names: `NextRandom`

Class methods

`distributionName`

`^^Normal distribution'`

`fromHistogram: aHistogram`

`^self new: aHistogram average sigma: aHistogram standardDeviation`

`new`

`^self new: 0 sigma: 1`

`new: aNumber1 sigma: aNumber2`

`^super new initialize: aNumber1 sigma: aNumber2`

`random`

```

| v1 v2 w y |
NextRandom isNil
  ifTrue: [ [ v1 := Number random * 2 - 1.
             v2 := Number random * 2 - 1.
             w := v1 squared + v2 squared.
             w > 1 ] whileTrue: [].
            y := ( ( w ln * 2 negated) / w) sqrt.
            v1 := y * v1.
            NextRandom := y * v2.
          ]
  ifFalse:[ v1 :=NextRandom.
           NextRandom := nil.
          ].
^v1

```

Instance methods

average

^mu

changeParametersBy: aVector

```

mu := mu + ( aVector at: 1).
sigma := sigma + ( aVector at: 2).

```

distributionValue: aNumber

```

^DhbErfApproximation new value: ( ( aNumber - mu) / sigma)

```

initialize: aNumber1 sigma: aNumber2

```

mu := aNumber1.
sigma := aNumber2.
^self

```

kurtosis

^0

parameters

```

^Array with: mu with: sigma

```

random

```

    ^self class random * sigma + mu

skewness
    ^0

standardDeviation
    ^sigma

value: aNumber

    ^( DhbErfApproximation new normal: (aNumber - mu) / sigma) /
      sigma

valueAndGradient: aNumber

    | dp y |
    y := ( aNumber - mu) / sigma.
    dp := ( DhbErfApproximation new normal: y) / sigma.
    ^Array with: dp
      with: ( DhbVector with: dp * y / sigma
        with: dp * ( y squared - 1) / sigma)

```

9.6.2 Normal distribution — Java implementation

Listing 9.20 shows the implementation of the normal distribution in Java.

The distribution function of the normal distribution can be computed with the error function (c.f. section 2.3). Therefore the class `NormalDistribution` is implemented as a subclass of `ProbabilityDensityFunction`.

Figure 9.1 with the box `NormalDistribution` grayed.

Listing 9.20 Java implementation of the normal distribution

```

package DhbStatistics;

import DhbFunctionEvaluation.PolynomialFunction;
import DhbScientificCurves.Histogram;
import DhbInterfaces.ParametrizedOneVariableFunction;
/**
 * Normal distribution, a.k.a. Gaussian distribution.
 *
 * @author Didier H. Besset
 */

```



```

public final class NormalDistribution
    extends ProbabilityDensityFunction
{
    /**
     * Average of the distribution.
     */
    private double mu;
    /**
     * Standard deviation of the distribution.
     */
    private double sigma;
    /**
     * Constant needed to compute the norm.
     */
    private static double baseNorm = Math.sqrt( 2 * Math.PI);
    /**
     * Series to compute the error function.
     */
    private static PolynomialFunction errorFunctionSeries;
    static {
        double[] coeffs = { 0.31938153, -0.356563782, 1.781477937,
                           -1.821255978, 1.330274429};
        errorFunctionSeries = new PolynomialFunction( coeffs);
    };
    /**
     * Constant needed to compute the argument to the error function series.
     */
    private static double errorFunctionConstant = 0.2316419;

    /**
     * Defines a normalized Normal distribution with average 0
     * and standard deviation 1.
     */
    public NormalDistribution ( ) throws IllegalArgumentException
    {
        this( 0, 1);
    }
    /**
     * Defines a Normal distribution with known average
     * and standard deviation.
     * @param average of the distribution
     * @param standard deviation of the distribution
     * @exception java.lang.IllegalArgumentException
     * if the standard deviation is non-positive
     */
    public NormalDistribution ( double average, double standardDeviation)

```

```

        throws IllegalArgumentException
    {
        if ( standardDeviation <= 0 )
            throw new IllegalArgumentException(
                "Standard deviation must be positive");
        mu = average;
        sigma = standardDeviation;
    }
    /**
     * Create an instance of the receiver with parameters estimated from
     * the given histogram using best guesses. This method can be used to
     * find the initial values for a fit.
     * @param h DhbScientificCurves.Histogram
     */
    public NormalDistribution( Histogram h)
    {
        this( h.average(), h.standardDeviation());
    }
    /**
     * @return double average of the distribution.
     */
    public double average ( )
    {
        return mu;
    }
    /**
     * Returns the probability of finding a random variable smaller
     * than or equal to x.
     * @return integral of the probability density function from -infinity to x.
     * @param x double upper limit of integral.
     */
    public double distributionValue( double x)
    {
        return errorFunction( ( x - mu) / sigma);
    }
    /**
     * @return error function for the argument.
     * @param x double
     */
    public static double errorFunction ( double x)
    {
        if ( x == 0 )
            return 0.5;
        else if ( x > 0 )
            return 1 - errorFunction( -x);
        double t = 1 / (1 - errorFunctionConstant * x);
    }

```

```

        return t * errorFunctionSeries.value( t ) * normal( x );
    }
    /**
     * @return double kurtosis of the distribution.
     */
    public double kurtosis( )
    {
        return 0;
    }
    /**
     * @return java.lang.String    name of the distribution
     */
    public String name ( )
    {
        return "Normal distribution";
    }
    /**
     * @return the density probability function for a (0,1) normal distribution.
     * @param x double    value for which the probability is evaluated.
     */
    static public double normal( double x )
    {
        return Math.exp( -0.5 * x * x ) / baseNorm;
    }
    /**
     * @return double[]    array containing mu and sigma
     */
    public double[] parameters ( )
    {
        double[] answer = new double[2];
        answer[0] = mu;
        answer[1] = sigma;
        return answer;
    }
    /**
     * @return double a random number distributed according to the receiver.
     */
    public double random( )
    {
        return generator().nextGaussian() * sigma + mu;
    }
    /**
     * @param average double
     */
    public void setAverage( double average )
    {

```

```
    mu = average;
}
/**
 * @param p double[]  assigns the parameters
 */
public void setParameters( double[] params)
{
    setAverage( params[0]);
    setStandardDeviation( params[1]);
}
/**
 * @param average double
 */
public void setStandardDeviation( double standardDeviation)
{
    sigma = standardDeviation;
}
/**
 * @return double skewness of the distribution.
 */
public double skewness( )
{
    return 0;
}
/**
 * @return double standard deviation of the distribution
 */
public double standardDeviation( )
{
    return sigma;
}
/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    java.text.DecimalFormat fmt = new java.text.DecimalFormat(
        "0.00000");
    sb.append("Normal distribution (");
    sb.append(fmt.format(mu));
    sb.append(',');
    sb.append(fmt.format(sigma));
    sb.append(')');
    return sb.toString();
}
```

```

/**
 * @return double probability density function
 * @param x double random variable
 */
public double value( double x)
{
    return normal( ( x - mu) / sigma) / sigma;
}
/**
 * Evaluate the distribution and the gradient of the distribution with respect
 * to the parameters.
 * @return double[] 0: distribution's value, 1,2,...,n distribution's gradient
 * @param x double
 */
public double[] valueAndGradient( double x)
{
    double[] answer = new double[3];
    double y = ( x - mu) / sigma;
    answer[0] = normal( y) / sigma;
    answer[1] = answer[0] * y / sigma;
    answer[2] = answer[0] * ( y * y - 1) / sigma;
    return answer;
}
}

```

9.7 Gamma distribution

The gamma distribution is used to describe the time between some task, for example the time between repairs.

The generalization of the gamma distribution to a range of the random variable of the type $[x_{\min}, +\infty[$ is called a Pearson type III distribution. The average of a Pearson type III distribution is $x_{\min} + \alpha\beta$. The central moments are the same as those of the gamma distribution.

Table 9.3 shows the properties of the gamma distribution. Figure 9.4 shows the shape of the gamma distribution for several values of the parameter α with $\beta = 1$. The shape of the distribution for values of the parameter β can be obtained by modifying the scale of the x -axis since β is just a scale factor of the random variable.

9.7.1 Gamma distribution — Smalltalk implementation

Listing 9.21 shows the implementation of the gamma distribution in Smalltalk.

The distribution function of the gamma distribution can be computed with the incomplete gamma function (c.f. section 2.4.1). Therefore the class `DhbGammaDistribution` is implemented as a subclass of `DhbProbabilityDensity`.

Figure 9.1 with the box `GammaDistribution` grayed.

Table 9.3: Properties of the gamma distribution

Range of random variable	$[0, +\infty[$
Probability density function	$P(x) = \frac{x^{\alpha-1}}{\beta^\alpha \Gamma(\alpha)} e^{-\frac{x}{\beta}}$
Parameters	$0 < \alpha < +\infty$ $0 < \beta < +\infty$
Distribution function	$F(x) = \Gamma\left(\frac{x}{\beta}, \alpha\right)$ (c.f. section 7.4.1)
Average	$\alpha\beta$
Variance	$\alpha\beta^2$
Skewness	$\frac{2}{\sqrt{\alpha}}$
Kurtosis	$\frac{6}{\alpha}$

Listing 9.21 Smalltalk implementation of the gamma distribution


```

Class                DhbGammaDistribution
Subclass of          DhbProbabilityDensity
Instance variable names:  alpha beta norm randomCoefficients incompleteGammaFunction

    Class methods
distributionName
    ^'Gamma distribution'

fromHistogram: aHistogram
    | alpha beta |
    aHistogram minimum < 0
        ifTrue: [ ^nil].
    alpha := aHistogram average.
    beta := aHistogram variance / alpha.
    ^[ self shape: alpha / beta scale: beta] when: ExAll do: [
        :signal | signal exitWith: nil]

```

Figure 9.4: Gamma distribution for various values of α

new

^self error: 'Illegal creation message for this class'

shape: aNumber1 scale: aNumber2

^super new initialize: aNumber1 scale: aNumber2

Instance methods

average

^alpha * beta

changeParametersBy: aVector

alpha := alpha + (aVector at: 1).

beta := beta + (aVector at: 2).

self computeNorm.

incompleteGammaFunction := nil.

randomCoefficients := nil.

computeNorm

norm := beta ln * alpha + alpha logGamma.

distributionValue: aNumber

^self incompleteGammaFunction value: aNumber / beta

incompleteGammaFunction

incompleteGammaFunction isNil

ifTrue:

[incompleteGammaFunction := DhbIncompleteGammaFunction
shape: alpha].

^incompleteGammaFunction

initialize: aNumber1 scale: aNumber2

(aNumber1 > 0 and: [aNumber2 > 0])

ifFalse: [self error: 'Illegal distribution parameters'].

alpha := aNumber1.

beta := aNumber2.

self computeNorm.

^self

initializeRandomCoefficientsForLargeAlpha

| a b q d |

a := 1 / (2 * alpha - 1) sqrt.

b := alpha - (4 ln).

q := 1 / a + alpha.

d := 4.5 ln + 1.

^Array with: a with: b with: q with: d

initializeRandomCoefficientsForSmallAlpha

| e |

e := 1 exp.

^(e + alpha) / e

kurtosis

^6 / alpha

parameters

^Array with: alpha with: beta

random

```
^( alpha > 1 ifTrue: [ self randomForLargeAlpha]
  ifFalse:[ self randomForSmallAlpha]) * beta
```

randomCoefficientsForLargeAlpha

```
randomCoefficients isNil
  ifTrue: [ randomCoefficients := self
            initializeRandomCoefficientsForLargeAlpha].
^randomCoefficients
```

randomCoefficientsForSmallAlpha

```
randomCoefficients isNil
  ifTrue: [ randomCoefficients := self
            initializeRandomCoefficientsForSmallAlpha].
^randomCoefficients
```

randomForLargeAlpha

```
[ true] whileTrue: [
| u1 u2 c v y z w|
u1 := DhbMitchellMooreGenerator new floatValue.
u2 := DhbMitchellMooreGenerator new floatValue.
c := self randomCoefficientsForLargeAlpha.
v := ( u1 / ( 1 - u1)) ln * (c at: 1).
y := v exp * alpha.
z := u1 squared * u2.
w := ( c at: 3) * v + ( c at: 2) - y.
( c at: 4) + w >= ( 4.5 * z) ifTrue: [ ^y].
z ln <= w ifTrue: [ ^y].
].
```

randomForSmallAlpha

```
[ true] whileTrue: [
| p |
p := DhbMitchellMooreGenerator new floatValue * self
    randomCoefficientsForSmallAlpha.
p > 1
  ifTrue: [ | y |
            y := ( ( self randomCoefficientsForSmallAlpha -
                    p) / alpha) ln negated.
            DhbMitchellMooreGenerator new floatValue <= ( y
```

```

        raisedTo: ( alpha - 1))
      ifTrue: [ ^y].
    ]
    ifFalse: [ | y |
      y := p raisedTo: ( 1 / alpha).
      DhbMitchellMooreGenerator new floatValue <= ( y
        negated exp)
      ifTrue: [ ^y].
    ].
  ].

skewness
  ^2 / alpha sqrt

value: aNumber
  ^aNumber > 0
    ifTrue: [ ( aNumber ln * (alpha - 1) - (aNumber / beta) -
      norm) exp]
    ifFalse:[ 0].

variance
  ^beta squared * alpha

```

9.7.2 Gamma distribution — Java implementation

Listing 9.22 shows the implementation of the gamma distribution in Java.

The distribution function of the gamma distribution can be computed with the incomplete gamma function (c.f. section 2.4.1). Therefore the class `GammaDistribution` is implemented as a subclass of `ProbabilityDensityFunction`.

Figure 9.1 with the box `GammaDistribution` grayed.

Listing 9.22 Java implementation of the gamma distribution

```

package DhbStatistics;

import DhbIterations.IncompleteGammaFunction;
import DhbFunctionEvaluation.GammaFunction;
import DhbScientificCurves.Histogram;
import DhbInterfaces.ParametrizedOneVariableFunction;
/**
 * Gamma distribution.

```

```

*
* @author Didier H. Besset
*/
public class GammaDistribution extends ProbabilityDensityFunction
{
    /**
     * Shape parameter of the distribution.
     */
    protected double alpha;
    /**
     * Scale parameter of the distribution.
     */
    private double beta;
    /**
     * Norm of the distribution (cached for efficiency).
     */
    private double norm;
    /**
     * Constants used in random number generator (cached for efficiency).
     */
    private double a;
    private double b;
    private double q;
    private double d;
    /**
     * Incomplete gamma function for the distribution (cached for efficiency).
     */
    private IncompleteGammaFunction incompleteGammaFunction;
    /**
     * Constructor method (for internal use only).
     */
    protected GammaDistribution()
    {
    }
    /**
     * Create a new instance of the Gamma distribution with given shape and scale.
     * @param shape double shape parameter of the distribution (alpha).
     * @param scale double scale parameter of the distribution (beta).
     * @exception java.lang.IllegalArgumentException The exception description.
     */
    public GammaDistribution ( double shape, double scale) throws IllegalArgumentException
    {
        if ( shape <= 0 )
            throw new IllegalArgumentException( "Shape parameter must be positive");
        if ( scale <= 0 )
            throw new IllegalArgumentException( "Scale parameter must be positive");
    }

```

```

    defineParameters( shape, scale);
}
/**
 * Create an instance of the receiver with parameters estimated from the
 * given histogram using best guesses. This method can be used to
 * find the initial values for a fit.
 * @param h DhbScientificCurves.Histogram
 * @exception java.lang.IllegalArgumentException when no suitable parameter can be found.
 */
public GammaDistribution( Histogram h) throws IllegalArgumentException
{
    if ( h.getMinimum() < 0 )
        throw new IllegalArgumentException("Gamma distribution is only defined for non-negative values");
    double shape = h.average();
    if ( shape <= 0 )
        throw new IllegalArgumentException("Gamma distribution must have a non-negative shape parameter");
    double scale = h.variance() / shape;
    if ( scale <= 0 )
        throw new IllegalArgumentException("Gamma distribution must have a non-negative scale parameter");
    defineParameters( shape / scale, scale);
}
/**
 * @return double average of the distribution.
 */
public double average()
{
    return alpha * beta;
}
/**
 * Assigns new values to the parameters.
 * This method assumes that the parameters have been already checked.
 */
public void defineParameters ( double shape, double scale)
{
    alpha = shape;
    beta = scale;
    norm = Math.log( beta ) * alpha + GammaFunction.logGamma( alpha);
    if ( alpha < 1)
        b = (Math.E + alpha) / Math.E;
    else if ( alpha > 1)
    {
        a = Math.sqrt( 2 * alpha - 1);
        b = alpha - Math.log( 4.0);
        q = alpha + 1 / a;
        d = 1 + Math.log( 4.5);
    }
}

```

```
        incompleteGammaFunction = null;
        return;
    }
    /**
     * Returns the probability of finding a random variable smaller
     * than or equal to x.
     * @return integral of the probability density function from 0 to x.
     * @param x double upper limit of integral.
     */
    public double distributionValue ( double x)
    {
        return incompleteGammaFunction().value( x / beta);
    }
    /**
     * @return DhbIterations.IncompleteGammaFunction
     */
    private IncompleteGammaFunction incompleteGammaFunction()
    {
        if ( incompleteGammaFunction == null )
            incompleteGammaFunction = new IncompleteGammaFunction( alpha);
        return incompleteGammaFunction;
    }
    /**
     * @return double kurtosis of the distribution.
     */
    public double kurtosis( )
    {
        return 6 / alpha;
    }
    /**
     * @return java.lang.String name of the distribution.
     */
    public String name()
    {
        return "Gamma distribution";
    }
    /**
     * @return double[] an array containing the parameters of
     * the distribution.
     */
    public double[] parameters()
    {
        double[] answer = new double[2];
        answer[0] = alpha;
        answer[1] = beta;
        return answer;
    }
}
```

```

    }
    /**
     * @return double a random number distributed according to the receiver.
     */
    public double random( )
    {
        double r;

        if ( alpha > 1)
            r = randomForAlphaGreaterThan1();
        else if (alpha < 1)
            r = randomForAlphaLessThan1();
        else
            r = randomForAlphaEqual1();

        return r * beta;
    }
    /**
     * @return double
     */
    private double randomForAlphaEqual1( )
    {
        return -Math.log( 1 - generator().nextDouble());
    }
    /**
     * @return double
     */
    private double randomForAlphaGreaterThan1( )
    {
        double u1, u2, v, y, z, w;
        while ( true)
        {
            u1 = generator().nextDouble();
            u2 = generator().nextDouble();
            v = a * Math.log( u1 / ( 1 - u1));
            y = alpha * Math.exp( v);
            z = u1 * u1 * u2;
            w = b + q * v - y;
            if ( w + d - 4.5 * z >= 0 || w >= Math.log( z) )
                return y;
        }
    }
    /**
     * @return double
     */
    private double randomForAlphaLessThan1( )

```

```

{
    double p, y;
    while ( true)
    {
        p = generator().nextDouble() * b;
        if ( p > 1)
        {
            y = -Math.log( ( b - p) / alpha);
            if ( generator().nextDouble() <= Math.pow( y, alpha - 1) )
                return y;
        }
        y = Math.pow( p, 1 / alpha);
        if ( generator().nextDouble() <= Math.exp( -y) )
            return y;
    }
}
/**
 * @param p double[]  assigns the parameters
 */
public void setParameters( double[] params)
{
    defineParameters( params[0], params[1]);
}
/**
 * @return double skewness of the distribution.
 */
public double skewness( )
{
    return 2 / Math.sqrt( alpha);
}
/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    java.text.DecimalFormat fmt = new java.text.DecimalFormat("####0.00000");
    sb.append("Gamma distribution (");
    sb.append(fmt.format(alpha));
    sb.append(',');
    sb.append(fmt.format(beta));
    sb.append(')');
    return sb.toString();
}
/**
 * @return double probability density function

```

```

* @param x double random variable
*/
public double value( double x)
{
    return x > 0 ? Math.exp( Math.log( x) * ( alpha - 1) - x / beta - norm) : 0;
}
/**
* @return double variance of the distribution.
*/
public double variance( )
{
    return alpha * beta * beta;
}
}

```

9.8 Experimental distribution

A histogram described in section 9.3 can be used as a probability distribution. After all, a histogram can be considered as the representation of a distribution, which has been measured experimentally.

If N is the total count of the histogram and n_i the count in bin number i the probability of finding a measurement within the bin number i is simply given by:

$$P_i = \frac{n_i}{N}. \quad (9.33)$$

If w is the width of each bin, the probability density function of the distribution measured by the histogram can be estimated by:

$$P(x) = \frac{P_i}{w} = \frac{n_i}{wN} \quad \text{where } i = \left\lfloor \frac{x - x_{\min}}{w} \right\rfloor. \quad (9.34)$$

Equation 9.34 is only valid for $x_{\min} \leq x < x_{\max}$. Outside of the histogram's limits there is no information concerning the shape of the probability density function.

The distribution function is computed by evaluating the sum of all bins located below the argument and by adding a correction computed by linear interpolation over the bin in which the value is located. Thus, we have:

$$F(x) = \frac{1}{N} \left(\sum_{j=1}^{i-1} n_j + \frac{x - x_i}{w} n_i \right) \quad \text{where } i = \left\lfloor \frac{x - x_{\min}}{w} \right\rfloor. \quad (9.35)$$

If $x < x_{\min}$, $F(x) = 0$ and if $x \geq x_{\max}$, $F(x) = 1$. A similar equation can be derived for the acceptance interval function.

9.8.1 Experimental distribution — General implementation

Adding the responsibility of behaving like a probability density function to a histogram is not desirable. In a good object oriented design, objects should have only one responsibility or type of behavior.

Thus, a good object oriented implementation implements an Adapter pattern. One creates an object, having the behavior of a probability density function. A single instance variable inside this object refers to the histogram, over which the experimental distribution is defined. The Adapter object is a subclass of the abstract class describing all probability density functions.

The parameters of the distribution — average, variance, skewness and kurtosis — are obtained from the histogram itself.

The computation of the distribution and the interval acceptance function is delegated to the histogram. The floor operation of equation 9.35 is evaluated by the method `binIndex` of the histogram.

Note: In both implementation, there is no protection against illegal arguments. Illegal arguments can occur when computing the distribution value when the histogram underflow and overflow counts are non zero. Below the minimum and above the maximum, no information can be obtained for the distribution. Within the histogram limits, there is no need for protection. Therefore the implementation assumes that the histogram was collected using automatic adjustment of the limits (c.f. section 9.3).

Figure 9.1 with the HistogrammedDistribution grayed.

9.8.2 Experimental distribution — Smalltalk implementation

Listing 9.23 shows the implementation of an experimental distribution in Smalltalk.

The class `DhbHistogrammedDistribution` is a subclass of the class `DhbProbabilityDensity`. The class creation method `histogram:` takes as argument the histogram over which the instance is defined. To prevent creating a instance with undefined instance variable, the default class creation method `new` returns an error.

Listing 9.23 Smalltalk implementation of an experimental distribution

```

Class          DhbHistogrammedDistribution
Subclass of    DhbProbabilityDensity
Instance variable names: histogram

Class methods
distributionName
    ^'Experimental distribution'

histogram: aHistogram
```

```
^super new initialize: aHistogram
```

```
new
```

```
^self error: 'Illegal creation message for this class'
```

```
Instance methods
```

```
acceptanceBetween: aNumber1 and: aNumber2
```

```
^( histogram countsBetween: ( aNumber1 max: histogram minimum)
    and: ( aNumber2 min: histogram maximum) ) /
    histogram totalCount
```

```
average
```

```
^histogram average
```

```
distributionValue: aNumber
```

```
^aNumber < histogram minimum
    ifTrue: [ 0]
    ifFalse:[ aNumber < histogram maximum
        ifTrue: [ ( histogram countsUpTo:
            aNumber) / histogram totalCount]
        ifFalse:[ 1]
    ]
```

```
initialize: aHistogram
```

```
aHistogram count = 0
    ifTrue: [ self error: 'Cannot define probability density on
        an empty histogram'].
    histogram := aHistogram.
    ^self
```

```
kurtosis
```

```
^histogram kurtosis
```

```
privateInverseDistributionValue: aNumber
```

```
^histogram inverseDistributionValue: aNumber
```

skewness

^histogram skewness

standardDeviation

^histogram standardDeviation

value: aNumber

```
^( aNumber >= histogram minimum and: [ aNumber < histogram
                                     maximum])
  ifTrue: [ ( histogram countAt: aNumber) / ( histogram
                                     totalCount * histogram binWidth)]
  ifFalse:[ 0]
```

variance

^histogram variance

9.8.3 Experimental distribution — Java implementation

Listing 9.24 shows the implementation of an experimental distribution in Java.

The class `HistogrammedDistribution` is a subclass of the class `ProbabilityDensityFunction`. The class constructor method takes as argument the histogram over which the instance is defined.

Listing 9.24 Java implementation of an experimental distribution

```
package DhbStatistics;

import DhbScientificCurves.Histogram;
/**
 * Distribution constructed on a histogram.
 *
 * @author Didier H. Besset
 */
public class HistogrammedDistribution extends ProbabilityDensityFunction
{
    Histogram histogram;
    /**
     * @return double average of the histogram.
     */
}
```

```

public double average()
{
    return histogram.average();
}
/**
 * Returns the probability of finding a random variable smaller
 * than or equal to x.
 * @return integral of the probability density function from -infinity to x.
 * @param x double upper limit of integral.
 */
public double distributionValue(double x)
{
    if ( x < histogram.getMinimum())
        return 0;
    else if ( x < histogram.getMaximum())
        return histogram.getCountsUpTo( x ) / histogram.totalCount();
    else
        return 1;
}
/**
 * @return double
 * @param x1 double
 * @param x2 double
 */
public double distributionValue ( double x1, double x2)
{
    return histogram.getCountsBetween(Math.max(x1,
        histogram.getMinimum()),
        Math.min(x2,
        histogram.getMaximum()))
        / histogram.totalCount();
}
/**
 * @return double kurtosis of the histogram.
 */
public double kurtosis()
{
    return histogram.kurtosis();
}
/**
 * @return java.lang.String name of the distribution.
 */
public String name()
{
    return "Experimental distribution";
}

```

```

/**
 * NOTE: this method is a dummy because the distribution
 * cannot be fitted.
 * @return double[] an array containing the parameters of
 *                 the distribution.
 */
public double[] parameters()
{
    return new double[0];
}
/**
 * This method is a dummy method, needed for the compiler because
 * the superclass requires implementation of the
 * interface ParametrizedOneVariableFunction.
 * Histogrammed distributions cannot be fitted.
 */
public void setParameters( double[] params)
{
}
/**
 * @return double skewness of the histogram.
 */
public double skewness()
{
    return histogram.skewness();
}
/**
 * @return double probability density function
 * @param x double random variable
 */
public double value( double x)
{
    return ( x >= histogram.getMinimum()
            || x < histogram.getMaximum() )
        ? histogram.getBinContent(x)
          / (histogram.totalCount()
            + histogram.getBinWidth())
        : 0;
}
/**
 * @return double variance of the histogram.
 */
public double variance()
{
    return histogram.variance();
}

```

}
}

Chapter 10

Statistical analysis

L'expérience instruit plus sûrement que le conseil.¹
André Gide

This chapter is dedicated on how to extract information from large amount of data using statistical analysis. One of the best book I have read on this subject is titled *How to lies with statistics*². This admittedly slanted title seems a little pessimistic. The truth, however, is that most people in their daily job ignore the little statistics they have learned in high school. As a result, statistical argumentation is often used wrongly to produce the wrong conclusions.

The problems addressed in this section pertain to the interpretation of experimental measurements. For a long time such a discipline was reserved to physicists only. Recently natural science disciplines discovered that statistics could be used effectively to verify hypotheses or to determine parameters based on experimental observations. Today, the best papers on statistics and estimations are found primarily in natural science publications (*Biometrika* e.g.). Recently, the use of statistics has been extended to financial analysis.

Statistics can be applied to experimental data in two ways. Firstly, one can test the consistency and/or accuracy of the data. These tests are the subject of the first 3 sections. Secondly, the values of unknown parameters can be derived from experimental data. This very important aspect of statistical analysis is treated in the remaining of the chapter.

Figure 10.1 shows the classes described in this chapter. The reader should be aware that the techniques used for non-linear least square fits (section 10.9) can also be also applied to solve systems of non-linear equations.

¹Experience teaches more surely than counseling.

²D. Huff, *How to lies with statistics*, Norton and Co., New York 1954.

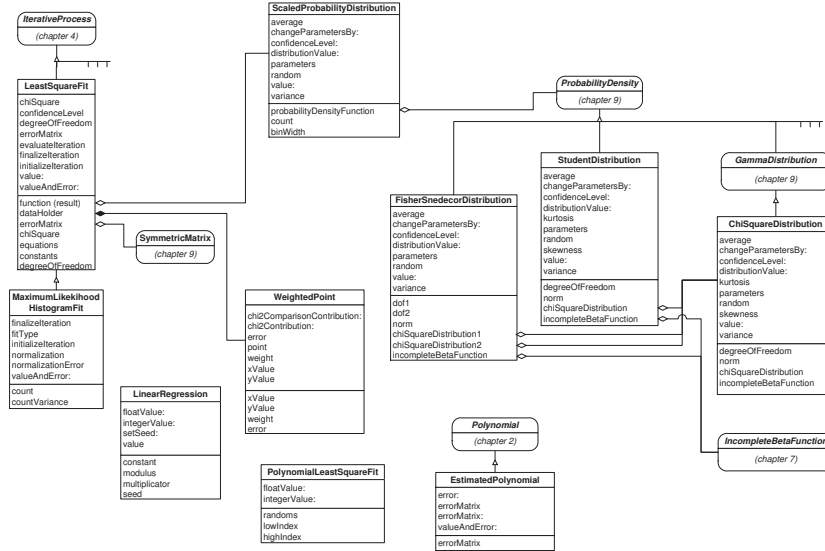


Figure 10.1: Classes related to estimation

10.1 F -test and the Fisher-Snedecor distribution

The F -test tries to answer the following question: given two series of measurements, x_1, \dots, x_n and y_1, \dots, y_m , what is the probability that the two measurements have the same standard deviation? The F -test is used when there are not enough statistics to perform a detailed analysis of the data.

Let us assume that the distribution of the two random variables, x and y , are normal distributions with respective averages μ_x and μ_y , and respective standard deviations σ_x and σ_y . Then, \bar{s}_x , the standard deviation of x_1, \dots, x_n is an estimator of σ_x ; \bar{s}_y , the standard deviation of y_1, \dots, y_m is an estimator of σ_y . The following statistics

$$F = \frac{\bar{s}_x^2}{\sigma_x^2} \cdot \frac{\sigma_y^2}{\bar{s}_y^2} \quad (10.1)$$

can be shown to be distributed according to a Fisher-Snedecor distribution with degrees of freedom n and m . In particular, if one wants to test for the equality of the two standard deviations, one constructs the following statistics:

$$F = \frac{\bar{s}_x^2}{\bar{s}_y^2}. \quad (10.2)$$

Traditionally one chooses $\bar{s}_x > \bar{s}_y$ so that the variable F is always greater than one.

It is important to recall that the expression above is distributed according to a Fisher-Snedecor distribution if and only if the two sets of data are distributed according to a normal distribution. For experimental measurements this is often the case unless

systematic errors are present. Nevertheless this assumption must be verified before making an *F*-test.

Table 10.1 shows the properties of the Fisher-Snedecor distribution. The Fisher-Snedecor distribution is itself rarely used as a probability density function, however.

Table 10.1: Properties of the Fisher-Snedecor distribution

Range of random variable	$[0, +\infty[$
Probability density function	$P(x) = \frac{n_1^{\frac{n_1}{2}} n_2^{\frac{n_2}{2}} x^{\frac{n_1-1}{2}}}{B\left(\frac{n_1}{2}, \frac{n_2}{2}\right) (n_1 + n_2 x)^{\frac{n_1+n_2}{2}}}$
Parameters	n_1, n_2 two positive integers
Distribution function	$F(x) = B\left(\frac{n_1}{n_1 + n_2 x}, \frac{n_1}{2}, \frac{n_2}{2}\right)$
Average	$\frac{n_2}{n_2 - 2}$ for $n > 2$ undefined otherwise
Variance	$\frac{2n_2^2 (n_1 + n_2 - 2)}{n_1 (n_2 - 2)^2 (n_2 - 4)}$ for $n > 4$ undefined otherwise
Skewness	
Kurtosis	

The main part of figure 10.2 shows the shape of the Fisher-Snedecor distribution for some values of the degrees of freedom. For large n and m , the Fisher-Snedecor distribution tends toward a normal distribution.

The confidence level of a Fisher-Snedecor distribution is defined as the probability expressed in percent of finding a value larger than F defined in equation 10.2 or lower than $1/F$. The confidence level is thus related to the distribution function. We have:

$$CL_F(x, n) = 100 \left\{ 1 - \left[F(x) - F\left(\frac{1}{x}\right) \right] \right\}, \quad (10.3)$$

where x is the value F defined in equation 10.2. The change of notation was made to avoid confusion between the acceptance function and the random variable. The confidence level corresponds to the surface of the shaded areas in the insert in the upper right corner of figure 10.2.

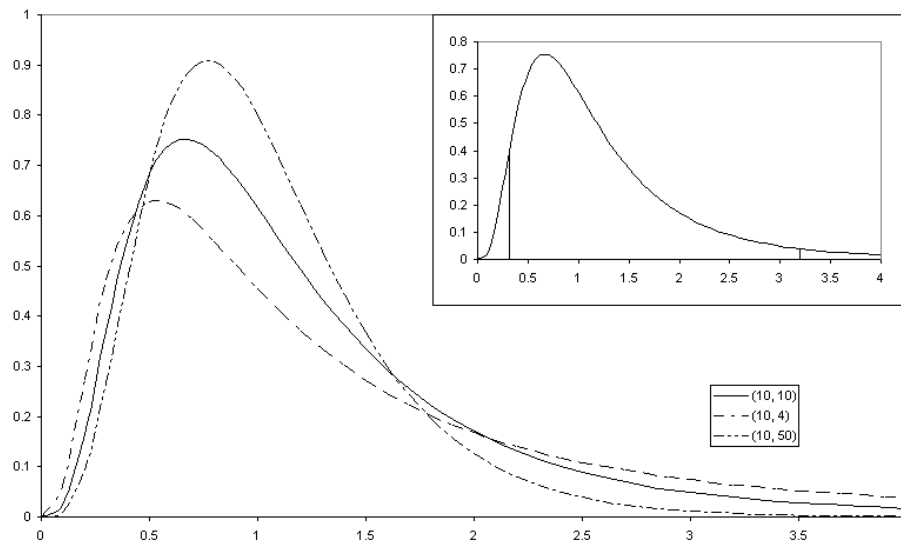


Figure 10.2: Fisher-Snedecor distribution for a few parameters

Example

. Here is an example used to investigate the goodness of the two random number generators discussed in section 9.4. The collected data are the error of the covariance test described in section 9.4. The dimension of the covariance matrix is 7 and the number of trials on each measurement was 1000. The table 10.2 presents the obtained results, expressed in $\frac{1}{1000}$. the ratio of the two variances is 1.18 and the degrees of freedom are both 10. The F -test applied to the two variances gives a confidence level of 20%. This means that there is only a 20% probability that the variances of the two series of measurements are different.

10.1.1 Fisher-Snedecor distribution — Smalltalk implementation

Listing 10.1 shows the implementation of the Fisher-Snedecor distribution in Smalltalk. Listing 10.2 shows the implementation of the F -test in Smalltalk. The following code example shows how to perform a F -test between two sets of experimental measurements.

Code example 10.1

```

| mom1 mom2 confidenceLevel |
mom1 := DhbFixedStatisticalMoments new.
    <Collecting measurements of set 1 into mom1>

mom2 := DhbFixedStatisticalMoments new.
    <Collecting measurements of set 2 into mom1>

confidenceLevel := mom1 fConfidenceLevel: mom2.

```

Figure 10.1 with the box FisherSnedecorDistribution grayed.

Table 10.2: Covariance test of random number generator

	Congruential	Mitchell-Moore
1	5.56	7.48
2	5.89	6.75
3	4.66	3.77
4	5.69	5.71
5	5.34	7.25
6	4.79	4.73
7	4.80	6.23
8	7.86	5.60
9	3.64	5.94
10	5.70	4.58
Average	5.40	5.80
Std. dev.	1.10	1.19

Two instances of statistical moments (c.f. section 9.2.2) are created. Experimental data are accumulated into each set separately (c.f. code example 9.1). The last line returns the probability in percent that the two sets of data have the same standard deviation.

The class `DhbFisherSnedecorDistribution` is implemented as a subclass of `DhbProbabilityDensity` because its distribution function can be computed numerically using the incomplete beta function (c.f. section 7.5).

Listing 10.1 Smalltalk implementation of the Fisher-Snedecor distribution

```

Class                DhbFisherSnedecorDistribution
Subclass of          DhbProbabilityDensity

Instance variable names:  dof1  dof2  norm  chiSquareDistribution1  chiSquareDistribution2
                           incompleteBetaFunction

Class methods

degreeOfFreedom: anInteger1 degreeOfFreedom: anInteger2

^super new initialize: anInteger1 and: anInteger2

distributionName

^^Fisher-Snedecor distribution'

fromHistogram: aHistogram

| n1 n2 a |
aHistogram minimum < 0 ifTrue: [^nil].

```

```

n2 := (2 / (1 - (1 / aHistogram average))) rounded.
n2 > 0 ifFalse: [^nil].
a := (n2 - 2) * (n2 - 4) * aHistogram variance / (n2 squared *
2).
n1 := (0.7 * (n2 - 2) / (1 - a)) rounded.
^n1 > 0
  ifTrue: [self degreeOfFreedom: n1 degreeOfFreedom: n2]
  ifFalse: [nil]

```

```

new

```

```

^self error: 'Illegal creation message for this class'

```

```

test: aStatisticalMoment1 with: aStatisticalMoment2

```

```

^(self class degreeOfFreedom: aStatisticalMoment1 count
degreeOfFreedom: aStatisticalMoment2 count)
distributionValue: aStatisticalMoment1 variance
/ aStatisticalMoment2 variance

```

```

Instance methods

```

```

average

```

```

^dof2 > 2
  ifTrue: [ dof2 / ( dof2 - 2)]
  ifFalse:[ nil]

```

```

changeParametersBy: aVector

```

```

dof1 := ( dof1 + ( aVector at: 1)) max: 1.
dof2 := ( dof2 + ( aVector at: 2)) max: 1.
self computeNorm.
chiSquareDistribution1 := nil.
chiSquareDistribution2 := nil.
incompleteBetaFunction := nil.

```

```

computeNorm

```

```

norm := ( dof1 ln * ( dof1 / 2) ) + ( dof2 ln * ( dof2 / 2) )
- ( ( dof1 / 2) logBeta: ( dof2 / 2) ).

```

```

confidenceLevel: aNumber

```

```

aNumber < 0
  ifTrue: [ self error: 'Confidence level argument must be
                      positive' ].
^( ( self distributionValue: aNumber) - ( self distributionValue:
    aNumber reciprocal) ) * 100

distributionValue: aNumber

^1 - ( self incompleteBetaFunction value: ( dof2 / ( aNumber *
    dof1 + dof2)))

incompleteBetaFunction

incompleteBetaFunction isNil
  ifTrue:
    [incompleteBetaFunction := DhbIncompleteBetaFunction
      shape: dof2 / 2
      shape: dof1 / 2].
^incompleteBetaFunction

initialize: anInteger1 and: anInteger2

dof1 := anInteger1.
dof2 := anInteger2.
self computeNorm.
^self

parameters

^Array with: dof1 with: dof2

random

chiSquareDistribution1 isNil
  ifTrue: [ chiSquareDistribution1 := DhbChiSquareDistribution
    degreeOfFreedom: dof1.
    chiSquareDistribution2 := DhbChiSquareDistribution
    degreeOfFreedom: dof2.
  ].
^chiSquareDistribution1 random * dof2 / ( chiSquareDistribution2
    random * dof1)

value: aNumber

```

```

^aNumber > 0
  ifTrue: [ ( norm + ( aNumber ln * ( dof1 / 2 - 1 ) ) - (
    (aNumber * dof1 + dof2) ln * ( ( dof1 + dof2) / 2))) exp]
  ifFalse:[ 0]

```

variance

```

^dof2 > 4 ifTrue: [ dof2 squared * 2 * ( dof1 + dof2 - 2) / ( (
  dof2 - 2) squared * dof1 * ( dof2 - 4))]
  ifFalse:[ nil]

```

The computation of the confidence level for the F -test is implemented in the method `fConfidenceLevel:` of the class `DhbStatisticalMoments`. It calculates the statistics F according to equation 10.2, creates an instance of a Fisher-Snedecor distribution and passes the value of F to the method `confidenceLevel:` of the distribution. The method `fConfidenceLevel:` is also implemented by the class `Histogram` where it is simply delegated to the statistical moments accumulated by the histogram. The argument of the method can be a statistical moment or a histogram since the messages sent by the method are polymorphic to both classes.

Listing 10.2 Smalltalk implementation of the F -test

Class	DhbStatisticalMoments
Subclass of	Object
Instance variable names:	moments

Instance methods

`fConfidenceLevel: aStatisticalMomentsOrHistogram`

| fValue |

fValue := self variance/ aStatisticalMomentsOrHistogram variance.

^fValue < 1

```

  ifTrue: [ (DhbFisherSnedecorDistribution degreeOfFreedom:
    aStatisticalMomentsOrHistogram count
    degreeOfFreedom: self count)
    confidenceLevel: fValue
    reciprocal]

```

```

  ifFalse:[ (DhbFisherSnedecorDistribution degreeOfFreedom:
    self count
    degreeOfFreedom:
    aStatisticalMomentsOrHistogram count)
    confidenceLevel: fValue]

```

Class **DhbHistogram**
 Subclass of Object
 Instance variable names: minimum binWidth overflow underflow moments contents freeExtent
 cacheSize desiredNumberOfBins

Instance methods

fConfidenceLevel: aStatisticalMomentsOrHistogram

^moments fConfidenceLevel: aStatisticalMomentsOrHistogram

10.1.2 Fisher-Snedecor distribution — Java implementation

Listing 10.3 shows the implementation of the Fisher-Snedecor distribution in Java. The following code example shows how to perform a *F*-test between two sets of experimental measurements.

Figure 10.1 with the box FisherSnedecorDistribution grayed.

Code example 10.2

```
StatisticalMoments mom1 = new StatisticalMoments();
    <Collecting measurements of set 1 into mom1>

StatisticalMoments mom2 = new StatisticalMoments();

    <Collecting measurements of set 2 into mom1>

double confidenceLevel = mom1.fConfidenceLevel( mom2);
```

Two instances of statistical moments (c.f. section 9.2.3) are created. Experimental data are accumulated into each set separately (c.f. code example 9.2). The last line returns the probability in percent that the two sets of data have the same standard deviation.

Listing 10.3 Java implementation of the Fisher-Snedecor distribution

```
package DhbStatistics;

import DhbFunctionEvaluation.GammaFunction;
import DhbIterations.IncompleteBetaFunction;
import DhbScientificCurves.Histogram;
/**
 * Fisher-Snedecor distribution
 * (distribution used to perform the F-test).
 *
 * @author Didier H. Besset
```



```

*/
public final class FisherSnedecorDistribution
    extends ProbabilityDensityFunction
{
    /**
     * First degree of freedom.
     */
    protected int dof1;
    /**
     * Second degree of freedom.
     */
    private int dof2;
    /**
     * Norm (stored for efficiency).
     */
    private double norm;
    /**
     * Function used to compute the distribution.
     */
    private IncompleteBetaFunction incompleteBetaFunction = null;
    /**
     * Auxiliary distributions for random number generation.
     */
    private ChiSquareDistribution chiSquareDistribution1 = null;
    private ChiSquareDistribution chiSquareDistribution2 = null;
    /**
     * Create a new instance of the Fisher-Snedecor distribution with
     * given degrees of freedom.
     * @param n1 int first degree of freedom
     * @param n2 int second degree of freedom
     * @exception java.lang.IllegalArgumentException
     * one of the specified degrees of freedom is non-positive.
     */
    public FisherSnedecorDistribution( int n1, int n2)
        throws IllegalArgumentException
    {
        if ( n1 <= 0 )
            throw new IllegalArgumentException(
                "First degree of freedom must be positive");
        if ( n2 <= 0 )
            throw new IllegalArgumentException(
                "Second degree of freedom must be positive");
        defineParameters( n1, n2);
    }
    /**
     * Create an instance of the receiver with parameters estimated from

```

```

* the given histogram using best guesses. This method can be used to
* find the initial values for a fit.
* @param h Histogram
* @exception java.lang.IllegalArgumentException
*         when no suitable parameter can be found.
*/
public FisherSnedecorDistribution( Histogram h) throws IllegalArgumentException
{
    if ( h.getMinimum() < 0 )
        throw new IllegalArgumentException(
"Fisher-Snedecor distribution is only defined for non-negative values");
    int n2 = (int) Math.round(2 / (1 - 1 / h.average()));
    if ( n2 <= 0 )
        throw new IllegalArgumentException(
"Fisher-Snedecor distribution has positive degrees of freedom");
    double a = 1 - ( n2 - 2 ) * ( n2 - 4 ) * h.variance() / ( 2 * 2 * n2);
    int n1 = (int) Math.round( 0.7 * ( n2 - 2 ) / a);
    if ( n1 <= 0 )
        throw new IllegalArgumentException(
"Fisher-Snedecor distribution has positive degrees of freedom");
    defineParameters( n1, n2);
}
/**
* @return double average of the distribution.
*/
public double average()
{
    return dof2 > 2
        ? dof2 / ( dof2 - 2)
        : Double.NaN;
}
/**
* @return double
* @param x double
* @exception java.lang.IllegalArgumentException
*         if the argument is outside the expected range.
*/
public double confidenceLevel( double x)
        throws IllegalArgumentException
{
    return x < 0 ? Double.NaN
        : distributionValue( x ) * 100;
}
/**
* Assigns new degrees of freedom to the receiver.
* Compute the norm of the distribution after a change of parameters.

```

```

    * @param n1 int    first degree of freedom
    * @param n2 int    second degree of freedom
    */
    public void defineParameters ( int n1, int n2)
    {
        dof1 = n1;
        dof2 = n2;
        double nn1 = 0.5 * n1;
        double nn2 = 0.5 * n2;
        norm = nn1 * Math.log( n1) + nn2 * Math.log( n2)
              - GammaFunction.logBeta( nn1, nn2);
        incompleteBetaFunction = null;
        chiSquareDistribution1 = null;
        chiSquareDistribution2 = null;
    }
    /**
    * Returns the probability of finding a random variable smaller
    * than or equal to x.
    * @return integral of the probability density function from 0 to x.
    * @param x double upper limit of integral.
    */
    public double distributionValue ( double x)
    {
        return incompleteBetaFunction().value( dof2 / ( x * dof1 + dof2));
    }
    private IncompleteBetaFunction incompleteBetaFunction()
    {
        if ( incompleteBetaFunction == null )
            incompleteBetaFunction = new IncompleteBetaFunction(
                0.5 * dof1, 0.5 * dof2);
        return incompleteBetaFunction;
    }
    /**
    * @return java.lang.String    name of the distribution.
    */
    public String name ( )
    {
        return "Fisher-Snedecor distribution";
    }
    /**
    * @return double[] an array containing the parameters of
    *                  the distribution.
    */
    public double[] parameters ( )
    {
        double[] answer = new double[2];

```

```

    answer[0] = dof1;
    answer[1] = dof2;
    return answer;
}
/**
 * @return double a random number distributed according to the receiver.
 */
public double random( )
{
    if ( chiSquareDistribution1 == null )
    {
        chiSquareDistribution1 = new ChiSquareDistribution( dof1);
        chiSquareDistribution2 = new ChiSquareDistribution( dof2);
    }
    return chiSquareDistribution1.random() * dof2
        / ( chiSquareDistribution2.random() * dof1);
}
/**
 * This function cannot be fitted because the parameters are integers.
 * @param p double[] assigns the parameters
 */
public void setParameters( double[] params)
{
    defineParameters( Math.round( (float) params[0]),
        Math.round( (float) params[1]));
}
/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    sb.append("Fisher-Snedecor distribution (");
    sb.append(dof1);
    sb.append(',');
    sb.append(dof2);
    sb.append(')');
    return sb.toString();
}
/**
 * @return double probability density function
 * @param x double random variable
 */
public double value( double x)
{
    return x > 0

```

```

        ? Math.exp( norm + Math.log( x ) * ( dof1 / 2 - 1 )
          - Math.log( x * dof1 + dof2 )
            * ( ( dof1 + dof2 ) / 2 ))
      : 0;
    }
  /**
   * @return double variance of the distribution.
   */
  public double variance ( )
  {
    return dof2 > 4
      ? dof2 * dof2 * 2 * ( dof1 + dof2 + 2 )
        / ( dof1 * ( dof2 - 2 ) * ( dof2 - 4 ))
      : Double.NaN;
  }
}

```

10.2 *t*-test and the Student distribution

The *t*-test tries to answer the following question: given two series of measurements, x_1, \dots, x_n and y_1, \dots, y_m , what is the probability that the two measurements have the same average? The *t*-test is used when there are not enough statistics to perform a detailed analysis of the data.

Let us assume that the distribution of the two random variables, x and y , are normal distributions with respective averages μ_x and μ_y , and the same standard deviation σ . Then \bar{x} , the average of x_1, \dots, x_n is an estimator of μ_x ; \bar{y} , the average of y_1, \dots, y_m is an estimator of μ_y . An estimation \bar{s} of the standard deviation σ can be made using both measurement samples. We have:

$$\bar{s}^2 = \frac{\sum_{i=1}^n (x_i - \bar{x})^2 + \sum_{i=1}^m (y_i - \bar{y})^2}{n + m - 2}. \quad (10.4)$$

One can prove that the following statistics:

$$t = \frac{(\bar{x} - \bar{y}) - (\mu_x - \mu_y)}{\bar{s} \sqrt{\frac{1}{n} + \frac{1}{m}}} \quad (10.5)$$

is distributed according to a Student distribution with $n + m - 2$ degrees of freedom. In particular, to test for the probability that the two series of measurements have the same average, one uses the following statistics:

$$t = \frac{\bar{x} - \bar{y}}{\bar{s} \sqrt{\frac{1}{n} + \frac{1}{m}}}. \quad (10.6)$$

It is important to recall the two fundamental hypotheses that have been made so far.

1. The two sets of data must be distributed according to a normal distribution.
2. The two sets of data must have the same standard deviation.

Too many people use the *t*-test without first checking the assumptions. Assumption 1 is usually fulfilled with experimental measurements in the absence of systematic errors. Assumption 2, however, must be checked, for example using the *F*-test discussed in section 10.1.

Because the random variable of the distribution is traditionally labeled *t*, this distribution is often called the *t*-distribution. Table 10.3 shows the properties of the Student distribution. The Student distribution is itself rarely used as a probability density function, however.

Table 10.3: Properties of the Student distribution

Range of random variable	$]-\infty, +\infty[$
Probability density function	$P(x) = \frac{1}{\sqrt{n}B\left(\frac{n}{2}, \frac{1}{2}\right)} \left(1 + \frac{t^2}{n}\right)^{-\frac{n+1}{2}}$
Parameters	n a positive integer
Distribution function	$F(x) = \begin{cases} \frac{1+B\left(\frac{n}{n+x^2}, \frac{n}{2}, \frac{1}{2}\right)}{2} & \text{for } x \geq 0 \\ \frac{1-B\left(\frac{n}{n+x^2}, \frac{n}{2}, \frac{1}{2}\right)}{2} & \text{for } x < 0 \end{cases}$
Average	0
Variance	$\frac{n}{n-2}$ for $n > 2$ undefined otherwise
Skewness	0
Kurtosis	$\frac{6}{n-4}$ for $n > 4$ undefined otherwise

For $n = 1$, the Student distribution is identical to a Cauchy distribution with $\mu = 0$ and $\beta = 1$. For large n , the Student distribution tends toward a normal distribution with average 0 and variance 1. The main part of figure 10.3 shows the shapes of the Student distribution for a few values of the degrees of freedom. The normal distribution is also given for comparison.

The confidence level of a Student distribution is defined as the probability to find a value whose absolute value is larger than a given value. Thus, it estimates the level




Figure 10.3: Student distribution for a few degrees of freedom

of confidence that the hypothesis — namely, that the two sets of measurements have the same average — cannot be accepted. Traditionally the confidence level is given in percent. The confidence level corresponds to the surface of shaded area in the insert in the upper left corner of figure 10.3. By definition, the confidence level is related to the interval acceptance function:

$$CL_t(t, n) = 100 [1 - F(-|t|, |t|)], \quad (10.7)$$

using the definition of the interval acceptance function (equation 9.31). The value of t in equation 10.7 is obtained from equation 10.6.

The distribution function of the Student distribution is calculated with the incomplete beta function (c.f. section 7.5). Using the fact that the distribution is symmetric, one can derive the following expression

$$F(-|t|, |t|) = B\left(\frac{n}{n+t^2}; \frac{n}{2}, \frac{1}{2}\right), \quad (10.8)$$

from the properties of the distribution (c.f. table 10.3) and using equations 10.7 and 7.12.

Example

Now, we shall continue the analysis of the results of table 10.2. The t value computed from the two sets of measurements is 0.112 for a degree of freedom of 18. the corresponding confidence level is 8.76%. That is, there is only a 8.76% probability that the two generators have a different behavior. Thus, we can conclude that the Mitchell-Moore random generator is as good as the congruential random generator.

10.2.1 Student distribution — Smalltalk implementation

Listing 10.4 shows the implementation of the Student distribution in Smalltalk. Listing 10.5 shows the implementation of the *t*-test in Smalltalk. Performing a *t*-test between two sets of experimental measurements is very similar to performing a *F*-test. In code example 10.1 it suffices to replace the last line with the following:

```
confidenceLevel := mom1 fConfidenceLevel: mom2.
```

This last line returns the probability in percent that the two sets of data have the same average provided that the two sets have the same standard deviation.

The class `DhbStudentDistribution` is implemented as a subclass of `DhbProbabilityDensity` because its distribution function can be computed numerically using the incomplete beta function (c.f. section 7.5).

The method `symmetricAcceptance:` computes the symmetric acceptance function defined by equation 10.8. This method is used to compute the distribution function and the confidence level. The method `confidenceLevel:` gives the confidence level in percent.

Figure 10.1 with the box `StudentDistribution` grayed.

Listing 10.4 Smalltalk implementation of the Student distribution

```

Class                DhbStudentDistribution
Subclass of          DhbProbabilityDensity
Instance variable names: degreeOfFreedom norm chiSquareDistribution incompleteBetaFunction

Class methods
asymptoticLimit
^30

degreeOfFreedom: anInteger
^anInteger > self asymptoticLimit
  ifTrue: [DhbNormalDistribution new]
  ifFalse:
    [anInteger = 1
     ifTrue: [DhbCauchyDistribution shape: 0 scale: 1]
     ifFalse: [super new initialize: anInteger]]

distributionName
^^'Student distribution'

fromHistogram: aHistogram
```



```

| dof var |
var := aHistogram variance.
var = 0
  ifTrue: [ ^nil].
dof := ( 2 / (1 - (1 / aHistogram variance))) rounded max: 1.
^dof > self asymptoticLimit ifTrue: [ nil]
      ifFalse:[ self degreeOfFreedom: dof]

new
  ^self error: 'Illegal creation message for this class'

test: aStatisticalMoment1 with: aStatisticalMoment2

| t |
t := ( aStatisticalMoment1 average - aStatisticalMoment2 average)
      abs.
^1 - ( ( self class degreeOfFreedom: ( aStatisticalMoment1 count
+ aStatisticalMoment2 count - 2)) acceptanceBetween: t negated and:
t)

Instance methods
average
  ^0

changeParametersBy: aVector
  degreeOfFreedom := degreeOfFreedom + ( aVector at: 1).
  self computeNorm.

chiSquareDistribution
  chiSquareDistribution isNil
    ifTrue: [ chiSquareDistribution := DhbChiSquareDistribution
      degreeOfFreedom: (degreeOfFreedom - 1)].
  ^chiSquareDistribution

computeNorm
  norm := ( ( degreeOfFreedom / 2 logBeta: ( 1 / 2) ) + (
    degreeOfFreedom ln / 2)) negated.

confidenceLevel: aNumber

```

```
^( 1 - ( self symmetricAcceptance: aNumber abs)) * 100
```

```
distributionValue: aNumber
```

```
aNumber = 0
  ifTrue: [ ^0.5].
^( aNumber > 0
  ifTrue: [ 2 - ( self symmetricAcceptance: aNumber abs)]
  ifFalse:[ self symmetricAcceptance: aNumber abs]) / 2
```

```
incompleteBetaFunction
```

```
incompleteBetaFunction isNil
  ifTrue:
    [incompleteBetaFunction := DhbIncompleteBetaFunction
      shape: degreeOfFreedom / 2
      shape: 0.5].
^incompleteBetaFunction
```

```
initialize: anInteger
```

```
anInteger > 0
  ifFalse: [ self error: 'Degree of freedom must be positive'].
degreeOfFreedom := anInteger.
self computeNorm.
^self
```

```
kurtosis
```

```
^degreeOfFreedom > 4 ifTrue: [ 6 / ( degreeOfFreedom - 4)]
  ifFalse:[ nil]
```

```
parameters
```

```
^Array with: degreeOfFreedom
```

```
random
```

```
^DhbNormalDistribution random * ( ( (degreeOfFreedom - 1) / self
  chiSquareDistribution random ) sqrt)
```

```
skewness
```

```
^0
```

symmetricAcceptance: aNumber

```
^ self incompleteBetaFunction value: ( degreeOfFreedom / (
    aNumber squared + degreeOfFreedom))
```

value: aNumber

```
^( norm - ( ( aNumber squared / degreeOfFreedom + 1) ln * ( (
    degreeOfFreedom + 1) / 2))) exp
```

variance

```
^degreeOfFreedom > 2 ifTrue: [ degreeOfFreedom / (
    degreeOfFreedom - 2)]
ifFalse:[ nil]
```

The computation of the confidence level for the t -test is implemented in the method `tConfidenceLevel:` of the class `DhbStatisticalMoments`. It calculates the statistics t according to equation 10.6, creates an instance of a Student distribution and passes the value of t to the method `confidenceLevel:` of the distribution. The method `tConfidenceLevel:` is also implemented by the class `Histogram` where it is simply delegated to the statistical moments accumulated by the histogram. The argument of the method can be a statistical moment or a histogram since the messages sent by the method are polymorphic to both classes.

The method `unnormalizedVariance` of class `DhbStatisticalMoments` corresponds to each sums in the numerator of equation 10.4. To allow performing a t -test also with instances of class `DhbFastStatisticalMoments`, it was necessary to define this for that class.

Listing 10.5 Smalltalk implementation of the t -test

```
Class          DhbStatisticalMoments
Subclass of    Object
Instance variable names:  moments
```

Instance methods

```
tConfidenceLevel: aStatisticalMomentsOrHistogram
```

```
| sbar dof |
dof := self count + aStatisticalMomentsOrHistogram count - 2.
sbar := ( ( self unnormalizedVariance +
    aStatisticalMomentsOrHistogram unnormalizedVariance) / dof) sqrt.
^( DhbStudentDistribution degreeOfFreedom: dof)
```

```
confidenceLevel: ( self average -
    (aStatisticalMomentsOrHistogram average))
    / ( ( 1 / self count + (
    aStatisticalMomentsOrHistogram count)) sqrt * sbar)
```

Class **DhbFastStatisticalMoments**

Subclass of DhbStatisticalMoments

Instance methods

unnormalizedVariance

```
^(moments at: 3) - ((moments at: 2) squared * self count)
```

Class **DhbHistogram**

Subclass of Object

Instance variable names: minimum binWidth overflow underflow moments contents freeExtent
cacheSize desiredNumberOfBins

Instance methods

tConfidenceLevel: aStatisticalMomentsOrHistogram

```
^moments tConfidenceLevel: aStatisticalMomentsOrHistogram
```

unnormalizedVariance

```
^moments unnormalizedVariance
```

10.2.2 Student distribution — Java implementation

Listing 10.6 shows the implementation of the Student distribution in Java. Performing a *t*-test between two sets of experimental measurements is very similar to performing a *F*-test. In code example 10.2 it suffices to replace the last line with the following:

```
confidenceLevel = mom1.tConfidenceLevel( mom2);
```

This last line returns the probability in percent that the two sets of data have the same average provided that the two sets have the same standard deviation.

The class StudentDistribution is implemented as a subclass of ProbabilityDensityFunction because its distribution function can be computed numerically using the incomplete beta function (c.f. section 7.5).

Figure 10.1 with the box StudentDistribution grayed.

The private method `symmetricAcceptance` computes the symmetric acceptance function defined by equation 10.8. This method is used to compute the distribution function and the confidence level. The method `confidenceLevel` gives the confidence level in percent.

Listing 10.6 Java implementation of the Student distribution

```
package DhbStatistics;

import DhbFunctionEvaluation.GammaFunction;
import DhbIterations.IncompleteBetaFunction;
import DhbScientificCurves.Histogram;
/**
 * Student distribution
 * used in computing the t-test.
 *
 * @author Didier H. Besset
 */
public final class StudentDistribution
    extends ProbabilityDensityFunction
{
    /**
     * Degree of freedom.
     */
    protected int dof;
    /**
     * Norm (stored for efficiency).
     */
    private double norm;
    /**
     * Function used to compute the distribution.
     */
    private IncompleteBetaFunction incompleteBetaFunction = null;
    /**
     * Auxiliary distribution for random number generation.
     */
    private ChiSquareDistribution chiSquareDistribution = null;
    /**
     * Constructor method.
     * @param n int    degree of freedom
     * @exception java.lang.IllegalArgumentException\
     *         when the specified degree of freedom is non-positive.
     */
    public StudentDistribution( int n) throws IllegalArgumentException
    {
        if ( n <= 0 )
            throw new IllegalArgumentException(
```

```

        "Degree of freedom must be positive");
    defineParameters( n);
}
/**
 * Create an instance of the receiver with parameters estimated from
 * the given histogram using best guesses. This method can be used to
 * find the initial values for a fit.
 * @param h DhbScientificCurves.Histogram
 * @exception java.lang.IllegalArgumentException
 *         when no suitable parameter can be found.
 */
public StudentDistribution( Histogram h)
{
    double variance = h.variance();
    if ( variance <= 0 )
        throw new IllegalArgumentException(
            "Student distribution is only defined for positive variance");
    defineParameters( (int) Math.max( 1,
        Math.round( 2 / (1 - 1 / variance))));
}
/**
 * @return double average of the distribution.
 */
public double average()
{
    return 0;
}
/**
 * @return double
 * @param x double
 * @exception java.lang.IllegalArgumentException
 *         if the argument is illegal.
 */
public double confidenceLevel( double x)
    throws IllegalArgumentException
{
    return x < 0
        ? Double.NaN
        : symmetricAcceptance( x) * 100;
}
/**
 * @param n int    degree of freedom
 */
public void defineParameters ( int n)
{
    dof = n;
}

```

```

        norm = -( Math.log( dof ) * 0.5
                  + GammaFunction.logBeta( dof * 0.5, 0.5));
    }
    /**
     * Returns the probability of finding a random variable smaller
     * than or equal to x.
     * @return integral of the probability density function from -infinity to x.
     * @param x double upper limit of integral.
     */
    public double distributionValue(double x)
    {
        if ( x == 0 )
            return 0.5;
        double acc = symmetricAcceptance( Math.abs( x));
        return x > 0 ? 1 + acc : 1 - acc;
    }
    /**
     * @return DhbIterations.IncompleteBetaFunction
     */
    private IncompleteBetaFunction incompleteBetaFunction( )
    {
        if ( incompleteBetaFunction == null)
            incompleteBetaFunction = new IncompleteBetaFunction( dof / 2,
                                                                    0.5);
        return incompleteBetaFunction;
    }
    /**
     * @return double kurtosis of the distribution.
     */
    public double kurtosis( )
    {
        return dof > 4 ? 6 / ( dof - 4 ) : Double.NaN;
    }
    /**
     * @return java.lang.String the name of the distribution.
     */
    public String name()
    {
        return "Student distribution";
    }
    /**
     * @return double[] an array containing the parameters of
     * the distribution.
     */
    public double[] parameters()
    {

```

```

    double[] answer = new double[1];
    answer[0] = dof;
    return answer;
}
/**
 * @return double a random number distributed according to the receiver.
 */
public double random( )
{
    if ( chiSquareDistribution == null )
    {
        chiSquareDistribution = new ChiSquareDistribution( dof - 1);
    }
    return generator().nextGaussian() * Math.sqrt(( dof - 1)/chiSquareDistribution.random());
}
/**
 * This distribution cannot be fitted because the parameter is an integer.
 * @param p double[] assigns the parameters
 */
public void setParameters( double[] params)
{
    defineParameters( Math.round( (float) params[0]));
}
/**
 * @return double skewness of the distribution.
 */
public double skewness( )
{
    return 0;
}
/**
 * @return double integral from -x to x
 * @param x double
 */
private double symmetricAcceptance( double x)
{
    return incompleteBetaFunction().value( dof / ( x * x + dof));
}
/**
 * @return double probability density function
 * @param x double random variable
 */
public double value( double x)
{
    return Math.exp( norm - Math.log( x * x / dof + 1) * ( dof + 1) / 2);
}

```



```

/**
 * @return double variance of the distribution.
 */
public double variance ( )
{
    return dof > 2 ? dof / ( dof - 2) : Double.NaN;
}
}

```

Note: The method `tConfidenceLevel` for classes `StatisticalMoments` and `Histogram` are part of listing 9.5 and 9.8 respectively.

10.3 χ^2 -test and χ^2 distribution

The χ^2 -test tries to answer the following question: how well a theory is able to predict observed results? Alternatively a χ^2 -test can also tell whether two independent sets of observed results are compatible. This latter formulation is less frequently used than the former. Admittedly these two questions are somewhat vague. We shall now put them in mathematical terms for a more precise definition.

Let us assume that the measurement of an observable quantity depends on some parameters. These parameters cannot be adjusted by the experimenter but can be measured exactly³. Let x_p be the measured values of the observed quantity where p is a label for the parameters; let σ_p be the standard deviation of x_p .

The first question assumes that one can predict the values of the observed quantity: let μ_p be the predicted value of x_p . Then the quantity:

$$y_p = \frac{x_p - \mu_p}{\sigma_p} \quad (10.9)$$

is distributed according to a normal distribution with average 0 and standard deviation 1 if and only if the quantities x_p are distributed according to a normal distribution with average μ_p and standard deviation σ_p .

A χ^2 distribution with n degrees of freedom describes the distribution of the sum of the squares of n random variables distributed according to a normal distribution with mean 0 and standard deviation 1. Thus, the following quantity

$$S = \sum_p \frac{(x_p - \mu_p)^2}{\sigma_p^2} \quad (10.10)$$

is distributed according to a χ^2 distribution with n degrees of freedom where n is the number of available measurements (that is the number of terms in the sum of equation 10.10).

³Of course, there is no such thing as an exact measurement. The measurement of the parameters must be far more precise than that of the observed quantity.

To formulate the second question one must introduce a second set of measurement of the same quantity and at the same values of the parameters. Let x'_p be the second set of measured values and σ'_p the corresponding standard deviations. The estimated standard deviation for the difference $x_p - x'_p$ is $\sqrt{\sigma_p^2 + \sigma_p'^2}$. If the two sets of measurements are compatible, the quantity

$$y'_p = \frac{x_p - x'_p}{\sqrt{\sigma_p^2 + \sigma_p'^2}} \quad (10.11)$$

is distributed according to a normal distribution with average 0 and standard deviation 1. Then the following quantity

$$S = \sum_p \frac{(x_p - x'_p)^2}{\sigma_p^2 + \sigma_p'^2} \quad (10.12)$$

is distributed according to a χ^2 distribution with n degrees of freedom.


Table 10.4 shows the properties of the χ^2 distribution. It is a special case of the

Table 10.4: Properties of the χ^2 distribution

Range of random variable	$[0, +\infty[$
Probability density function	$P(x) = \frac{x^{\frac{n}{2}-1} e^{-\frac{x}{2}}}{2^{\frac{n}{2}} \Gamma\left(\frac{n}{2}\right)}$
Parameters	n a positive integer
Distribution function	$F(x) = \Gamma\left(\frac{x}{2}; \frac{n}{2}\right)$
Average	n
Variance	$2n$
Skewness	$2\sqrt{\frac{2}{n}}$
Kurtosis	$\frac{12}{n}$

gamma distribution with $\alpha = \frac{n}{2}$ and $\beta = 2$ (c.f. section 9.7). For $n > 30$ one can prove that the variable $y = \sqrt{2x} - \sqrt{2n-1}$ is approximately distributed according to a normal distribution with average 0 and standard deviation 1. If n is very large the χ^2 distribution tends toward a normal distribution with average n and standard deviation $\sqrt{2n}$.

Figure 10.4 shows the shape of the χ^2 distribution for a few values of the degree of freedom.

Figure 10.4: χ^2 distribution for a few degrees of freedom

To perform a χ^2 -test, it is customary to evaluate the probability of finding a value larger than the value obtained in equations 10.10 or 10.12. In this form, the result of a χ^2 -test gives the probability that the set of measurements is not compatible with the prediction or with another set of measurements. The confidence level of a χ^2 value is defined as the probability of finding a value larger than χ^2 expressed in percent. It is thus related to the distribution function as follows:

$$CL_S = 100 [1 - F(S)], \quad (10.13)$$

where S is the quantity defined in equation 10.10 or 10.12. The confidence level corresponds to the surface of the shaded area of the insert in the upper right corner of figure 10.4.

Since the χ^2 distribution is a special case of the gamma distribution the confidence level can be expressed with the incomplete gamma function (c.f. section 7.4):

$$CL_S = 100 \left[1 - \Gamma \left(\frac{S}{2}, \frac{n}{2} \right) \right]. \quad (10.14)$$

For large n the χ^2 confidence level can be computed from the error function (c.f. section 2.3.1):

$$CL_S = 100 \left[1 - \operatorname{erf} \left(\sqrt{2S} - \sqrt{2n-1} \right) \right]. \quad (10.15)$$

10.3.1 χ^2 distribution — Smalltalk implementation

Listing 10.7 shows the implementation of the χ^2 distribution in Smalltalk. The asymptotic limit is implemented directly in the class creation method.

Figure 10.1 with the box ChiSquaredDistribution grayed.

Figure 10.1 with the box
ChiSquaredDistribution
grayed.

10.3.2 χ^2 distribution — Java implementation

Listing 10.8 shows the implementation of the chi-square distribution in Java. The asymptotic form of the distribution has not been implemented. The reason is that a constructor method can only return an instance of the class it belongs to. The alternative requires to keep the asymptotic form in an instance variable and to overload all methods of the super class to add a test whether or not the computation must be delegated to the asymptotic form. This is somewhat tedious. In practice, however, one seldom needs to use the χ^2 distribution as a probability density function. Only the distribution function is necessary to perform the χ^2 -test. Therefore, we did not implement an asymptotic form.

Listing 10.8 Java implementation of the χ^2 distribution

```
package DhbStatistics;

import DhbScientificCurves.Histogram;
/**
 * Chi square distribution.
 * (as special case of the gamma distribution)
 *
 * @author Didier H. Besset
 */
public final class ChiSquareDistribution extends GammaDistribution
{
    /**
     * Create a new instance as Gamma( n/2, 2).
     * @param n int degrees of freedom of the receiver.
     */
    public ChiSquareDistribution ( int n)
    {
        super( 0.5 * n, 2.0);
    }
    /**
     * Create an instance of the receiver with parameters estimated from
     * the given histogram using best guesses. This method can be used to
     * find the initial values for a fit.
     * @param h Histogram
     * @exception java.lang.IllegalArgumentException
     *         when no suitable parameter can be found.
     */
    public ChiSquareDistribution( Histogram h) throws IllegalArgumentException
    {
        if ( h.getMinimum() < 0 )
            throw new IllegalArgumentException(
```

```

        "Chi square distribution is only defined for non-negative values");
    int dof = (int) Math.round( h.average());
    if ( dof <= 0 )
        throw new IllegalArgumentException(
            "Chi square distribution is only defined for positive degrees of freedom");
    setDegreesOfFreedom( dof);
}
/**
 * @return double
 * @param x double
 * @exception java.lang.IllegalArgumentException
 *         if the argument is outside the expected range.
 */
public double confidenceLevel( double x)
        throws IllegalArgumentException
{
    return x < 0
        ? Double.NaN
        : ( 1 - distributionValue( x)) * 100;
}
/**
 * @return java.lang.String name of the distribution.
 */
public String name()
{
    return "Chi square distribution";
}
/**
 * @return double[] an array containing the parameters of
 *         the distribution.
 */
public double[] parameters()
{
    double[] answer = new double[1];
    answer[0] = alpha * 2;
    return answer;
}
/**
 * @param n int
 */
public void setDegreesOfFreedom( int n)
{
    super.defineParameters( 0.5 * n, 2.0);
}
/**
 * Note: for fitting, non-integer degree of freedom is allowed

```

```

* @param params double[]  assigns the parameters
*/
public void setParameters( double[] params)
{
    defineParameters( params[0] * 0.5, 2);
}
/**
* @return java.lang.String
*/
public String toString()
{
    StringBuffer sb = new StringBuffer();
    java.text.DecimalFormat fmt = new java.text.DecimalFormat("0.00");
    sb.append("Chi square distribution (");
    sb.append(fmt.format(alpha * 2));
    sb.append(')');
    return sb.toString();
}
}

```

Figure 10.1 with the box
WeightedPoint grayed.

10.3.3 Weighted point implementation

As we shall see in the rest of this chapter, the evaluation of equation 10.10 is performed at many places. Thus, it is convenient to create a new class handling this type of calculation. The new class is called `DhbWeightedPoint` in Smalltalk and `WeightedPoint` in Java. As the implementation in both languages are fully identical the two classes are described in a common section. An open-minded reader will see that, in this case, each language has one shortcoming over the other.

A weighted point has the following instance variables:

`xValue` the x value of the data point, that is x_i ,

`yValue` the y value of the data point, that is y_i ,

`weight` the weight of the point, that is $1/\sigma_i^2$ and

`error` the error of the y value, that is σ_i .

Accessor methods for each of these instance variables are provided. The accessor method for the error is using lazy initialization to compute the error from the weight in case the error has not yet been defined.

The method `chi2Contribution` — with an added semicolon at the end of the name for Smalltalk — implements the computation of one term of the sum in equation 10.10. The argument of the method is any object implementing the behavior of a one-variable function defined in section 2.1. In Java one can use the same method name to define a similar method to compute the terms of the sum of equation 10.12: in this case, the argument of the method is another weighted point. This is not possible in Smalltalk,

which cannot distinguish the types of the arguments. Thus, for Smalltalk the second method must have a different name: `chi2ComparisonContribution:`. Here Java marks a point over Smalltalk.

Creating instances of the classes can be done in many ways. The fundamental method takes as arguments x_i , y_i and the weight $1/\sigma_i^2$. However convenience methods are provided for frequent cases:

1. x_i , y_i and the error on y_i , σ_i ;
2. x_i and the content of a histogram bin; the weight is derived from the bin contents as explained in section 10.4;
3. x_i , y_i without known error; the weight of the point is set to 1; points without error should not be used together with points with errors;
4. x_i and a statistical moment; in this case, the value y_i is an average over a set of measurements; the weight is determined from the error on the average (c.f. section 9.1);

Examples of use of weighted points appear in many sections of this chapter (10.4, 10.8, 10.9).

In the Smalltalk class `DhbWeightedPoint` the values x_i and y_i are always supplied as an instance of the class `Point`. The class `DhbWeightedPoint` has the following class creation methods:

`point:weight:` fundamental method;

`point:error:` convenience method 1;

`point:count:` convenience method 2;

`point:` convenience method 3

The convenience method 4 is implemented by the method `asWeightedPoint` of the class `DhbStatisticalMoments`. This kind of technique is quite common in Smalltalk instead of making a class creation method with an explicit name (`fromMoment:` e.g.).

Listing 10.9 Smalltalk implementation of the weighted point class

```

Class                DhbWeightedPoint
Subclass of          Object
Instance variable names: xValue yValue weight error

```

Class methods

```
point: aPoint
```

```
^self new initialize: aPoint weight: 1
```


point: aNumber count: anInteger

```

^self point: aNumber @ anInteger
  weight: ( anInteger > 0 ifTrue: [ 1 / anInteger]
           ifFalse:[ 1])

```

point: aPoint error: aNumber

```

^self new initialize: aPoint error: aNumber

```

point: aPoint weight: aNumber

```

^self new initialize: aPoint weight: aNumber

```

Instance methods

chi2ComparisonContribution: aWeightedPoint

```

^(aWeightedPoint yValue - yValue) squared / ( 1 / aWeightedPoint
weight + ( 1 / weight))

```

chi2Contribution: aFunction

```

^(yValue - ( aFunction value: xValue)) squared * weight

```

error

```

error isNil
  ifTrue: [ error := 1 / weight sqrt].
^error

```

initialize: aPoint error: aNumber

```

error := aNumber.
^self initialize: aPoint weight: 1 / aNumber squared

```

initialize: aPoint weight: aNumber

```

xValue := aPoint x.
yValue := aPoint y.
weight := aNumber.
^self

```

point

`^xValue @ yValue`

`weight`

`^weight`

`xValue`

`^xValue`

`yValue`

`^yValue`

Class	DhbStatisticalMoments
Subclass of	Object

Instance variable names: moments

Instance methods

`asWeightedPoint: aNumber`

`^DhbWeightedPoint point: aNumber @ self average error: self
errorOnAverage`

The constructor methods for the Java class `WeightedPoint` are:

`(double,double,double)` fundamental method;

`(double,int)` convenience method 2;

`(double,double)` convenience method 3;

`(double,StatisticalMoments)` convenience method 4.

Convenience method 1 could not be implemented with a constructor method as the type of the arguments is the same as those in the fundamental method. So, it was necessary to create a method `setError` which must be used in combination with convenience method 3 to implement a definition with values and error. A point for Smalltalk which makes it even with Java!

Listing 10.10 Java implementation of the weighted point class

```

package DhbEstimation;

import DhbInterfaces.OneVariableFunction;
import DhbStatistics.StatisticalMoments;
/**
 * Point with error used in chi-square test and least square fits
 *
 * @author Didier H. Besset
 */
public class WeightedPoint
{
    private double xValue;
    private double yValue;
    private double weight;
    private double error = Double.NaN;
/**
 * Constructor method.
 * @param x double
 * @param y double
 */
    public WeightedPoint( double x, double y)
    {
        this( x, y, 1);
    }
/**
 * Constructor method.
 * @param x double
 * @param y double
 * @param w double
 */
    public WeightedPoint( double x, double y, double w)
    {
        xValue = x;
        yValue = y;
        weight = w;
    }
/**
 * Constructor method.
 * @param x double
 * @param n int    a Histogram bin content
 */
    public WeightedPoint( double x, int n)
    {
        this( x, n, 1.0 / Math.max(n,1));
    }
/**

```

```

* Constructor method.
* @param x double
* @param m DhbStatistics.StatisticalMoments
*/
public WeightedPoint( double x, StatisticalMoments m)
{
    this( x, m.average());
    setError( m.errorOnAverage());
}
/**
* @return double    contribution to chi^2 sum against
*                   a theoretical function
* @param wp WeightedPoint
*/
public double chi2Contribution( WeightedPoint wp)
{
    double residue = yValue - wp.yValue();
    return residue * residue / ( 1 / wp.weight() + 1 / weight);
}
/**
* @return double    contribution to chi^2 sum against
*                   a theoretical function
* @param f DhbInterfaces.OneVariableFunction
*/
public double chi2Contribution( OneVariableFunction f)
{
    double residue = yValue - f.value( xValue);
    return residue * residue * weight;
}
/**
* @return double    error of the receiver
*/
public double error()
{
    if( Double.isNaN( error) )
        error = 1 / Math.sqrt( weight);
    return error;
}
/**
* @param e double error on the point
*/
public void setError( double e)
{
    error = e;
    weight = 1 / ( e * e);
}

```

```

/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    sb.append( '(');
    sb.append( xValue);
    sb.append( ',');
    sb.append( yValue);
    sb.append( "+-");
    sb.append( error());
    sb.append( ')');
    return sb.toString();
}
/**
 * @return double    weight of the receiver
 */
public double weight() {
    return weight;
}
/**
 * @return double    x value of the receiver
 */
public double xValue() {
    return xValue;
}
/**
 * @return double    y value of the receiver
 */
public double yValue() {
    return yValue;
}
}

```

10.4 χ^2 -test on histograms

As we have seen in section 9.3 histograms are often used to collect experimental data. Performing a χ^2 -test of data accumulated into a histogram against a function is a frequent task of data analysis.

The χ^2 statistics defined by equation 10.10 requires an estimate of the standard deviation of the content of each bin. One can show that the contents of a histogram bin is distributed according to a Poisson distribution. The Poisson distribution is a discrete

distribution⁴ whose average is equal to the variance. The probability of observing the integer k is defined by:

$$P_{\mu}(k) = \frac{\mu^k}{k!} e^{-\mu}, \quad (10.16)$$

where μ is the average of the distribution. In the case of a histogram, the estimated variance of the bin content is then the bin content itself. Therefore equation 10.10 becomes:

$$S = \sum_{i=1}^n \frac{\{n_i - \mu [x_{\min} + (i + \frac{1}{2})w]\}^2}{n_i}, \quad (10.17)$$

where n is the number of bins of the histogram, x_{\min} its minimum and w its bin width. The estimation of the bin content against which the χ^2 statistics is computed, μ , is now a function evaluated at the middle of each bin to average out variations of the function over the bin interval.

In fact, the function μ is often related to a probability density function since histograms are measuring probability distributions. In this case the evaluation is somewhat different. Let $P(x)$ be the probability density function against which the χ^2 statistics is computed. Then, the predicted bin content for bin i is given by:

$$\mu_i = wNP \left[x_{\min} + \left(i + \frac{1}{2} \right) w \right], \quad (10.18)$$

where N is the total number of values accumulated in the histogram. This is a symmetric version of the definition of a probability density function: w plays the role of dx in equation 9.23. Plugging equation 10.18 into equation 10.17 yields the expression of the χ^2 statistics for a histogram computed against a probability density function $P(x)$:

$$S = \sum_{i=1}^n \frac{\{n_i - wNP [x_{\min} + (i + \frac{1}{2})w]\}^2}{n_i}, \quad (10.19)$$

This equation cannot be applied for empty bins. If the bin is empty one can set the weight to 1. This corresponds to a 63% probability of observing no counts if the expected number of measurement is larger than 0.

In both implementations a single class is in charge of evaluating the predicted bin contents. This class is called a scaled probability density function. It is defined by a probability distribution and a histogram.

10.4.1 χ^2 -test on histograms — Smalltalk implementation

Listing 10.11 shows the implementation of a scaled probability density function in Smalltalk. Listing 10.12 shows the additional methods for the class `DhbHistogram` needed to perform a χ^2 -test. Examples of use are given in sections 10.9.2 and 10.10.2. Here is a simple example showing how to compute a χ^2 -confidence level to estimate the goodness of a random number generator.

Figure 10.1 with the box `ScaledProbabilityDistribution` grayed.

⁴A discrete distribution is a probability distribution whose random variable is an integer.

Code example 10.3

```

| trials probDistr histogram |
trials := 5000.
probDistr := DhbNormalDistribution new.
histogram := DhbHistogram new.
histogram freeExtent: true; setDesiredNumberOfBins: 100.
trials timesRepeat: [ histogram accumulate: probDistr random].
histogram chi2ConfidenceLevelAgainst:
    ( DhbScaledProbabilityDensityFunction histogram: histogram
      against: probDistr)

```

The first line after the declaration defines the number of data to be generated to 5000. After, an new instance of a probability distribution — in this case a normal distribution with average 0 and variance 1 — is created. Then, a new instance of a histogram is created and the next line defines it with a rough number of bins of 100 and the ability to automatically adjust its limits. After all instances have been created, random data generated by the probability distribution are generated. The last statement — extending itself over the last three lines — calculates the confidence level. The argument of the method `chi2ConfidenceLevelAgainst:` is a scaled probability distribution constructed over the histogram and the probability distribution used to generate the accumulated data.

The class `DhbScaledProbabilityDensityFunction` has two class creation methods. The class method `histogram:against:` takes two arguments, a histogram and a probability distribution. This method is used to perform a χ^2 -test of the specified histogram against the given probability distribution. The class method `histogram:distributionClass:` first create a probability distribution of the given class using parameters estimated from the histogram. This method is used to create a scaled probability density function whose parameters will be determined with least square or maximum likelihood fits.

Listing 10.11 Smalltalk implementation of a scaled probability density function

Class	DhbScaledProbabilityDensityFunction
Subclass of	Object

Instance variable names: probabilityDensityFunction count binWidth

Class methods

histogram: aHistogram against: aProbabilityDensityFunction

^self new

initialize: aProbabilityDensityFunction

binWidth: aHistogram binWidth

count: aHistogram totalCount

histogram: aHistogram distributionClass: aProbabilityDensityFunctionClass

```
^(aProbabilityDensityFunctionClass fromHistogram: aHistogram)
  ifNotNil: [:dp | self histogram: aHistogram against: dp]
```

Instance methods

changeParametersBy: aVector

```
count := count + aVector last.
probabilityDensityFunction changeParametersBy: aVector.
```

distributionFunction

```
^probabilityDensityFunction distributionFunction
```

initialize: aProbabilityDensityFunction binWidth: aNumber count: anInteger

```
probabilityDensityFunction := aProbabilityDensityFunction.
binWidth := aNumber.
count := anInteger.
^self
```

parameters

```
^probabilityDensityFunction parameters copyWith: count
```

printOn: aStream

```
super printOn: aStream.
aStream nextPut: $[;
  nextPutAll: probabilityDensityFunction class
              distributionName;
  nextPut: $].
```

setCount: aNumber

```
count := aNumber.
```

value: aNumber

```
^(probabilityDensityFunction value: aNumber) * binWidth * count
```

valueAndGradient: aNumber


```

| g temp |
g := probabilityDensityFunction valueAndGradient: aNumber.
temp := binWidth * count.
^Array with: g first * temp
  with: ( (g last collect: [:each | each * temp]) copyWith:
    g first * binWidth)

```

The evaluation of equation 10.19 is performed by the method `chi2Against:` of the class `DhbHistogram`. This method uses the iterator method `pointsAndErrorsDo:`. This method iterates on all bins and performs on each of them a block using as argument a weighted point as described in section 10.3.3. This iterator method is also used for least square and maximum likelihood fits (c.f. sections 10.9.2 and 10.10.2).

Listing 10.12 Smalltalk implementation of χ^2 -test on histograms

Class	DhbHistogram
Subclass of	Object
Instance variable names:	minimum binWidth overflow underflow moments contents freeExtent cacheSize desiredNumberOfBins

Instance methods

`chi2Against: aScaledDistribution`

```

| chi2 |
chi2 := 0.
self pointsAndErrorsDo:
  [ :each | chi2 := ( each chi2Contribution:
    aScaledDistribution) + chi2].
^chi2

```

`chi2ConfidenceLevelAgainst: aScaledDistribution`

```

^( DhbChiSquareDistribution degreeOfFreedom: ( contents size -
  aScaledDistribution parameters size))
  confidenceLevel: ( self chi2Against: aScaledDistribution)

```

`pointsAndErrorsDo: aBlock`

```

| x |
x := self minimum - ( self binWidth / 2).
contents do:
  [ :each |
    x := x + self binWidth.

```

```
aBlock value: ( DhbWeightedPoint point: x count: each).
].
```

10.4.2 χ^2 -test on histograms — Java implementation

Listing 10.13 shows the implementation of a scaled probability density function in Java. Examples of use are given in sections 10.9.3 and 10.10.3. Here is a simple example showing how to compute a χ^2 -confidence level to estimate the goodness of a random number generator.

Figure 10.1 with the box ScaledProbabilityDistribution grayed.

Code example 10.4

```
NormalDistribution probDistr = new NormalDistribution();
Histogram histogram = new Histogram( 50, 100);
histogram.setGrowthAllowed();
for( int i = 0; i < 5000; i++ )
    histogram.accumulate( probDistr.random());
double chiCL = histogram.chi2ConfidenceLevelAgainst(
    new DhbScaledProbabilityDensityFunction
        (histogram,probDistr));
```

The first line creates an instance of a normal distribution with average 0 and variance 1. The second line creates an instance of a histogram with 100 desired bins and a cache of 50 elements. The next line defines that the limits of the histogram can be adjusted automatically. The for loop accumulate random numbers generated according to the normal distribution into the histogram. The final statement — spread over the last three lines — compute the χ^2 -confidence level of the histogram's content against the normal distribution.

The method `chi2Against` of class `Histogram` evaluates the sum of equation 10.17. The method `chi2ConfidenceLevelAgainst` in the same class calculates the confidence level of the χ^2 -test. The code of these two methods can be found in listing 9.7.

The class `ScaledProbabilityDensityFunction` has two constructor methods. The first one takes as arguments the probability density function, the total count and the bin width of the histogram. The second only takes the probability density function and the histogram as arguments.

Listing 10.13 Java implementation of a scaled probability density function

```
package DhbStatistics;

import DhbInterfaces.ParametrizedOneVariableFunction;
import DhbScientificCurves.Histogram;
/**
 * Construct a function from a probability density function
 * for a given norm.
```

```

*
* @author Didier H. Besset
*/
public class ScaledProbabilityDensityFunction
    implements ParametrizedOneVariableFunction
{
    /**
     * Total count of the histogram.
     */
    private double count;
    /**
     * Bin width of the histogram.
     */
    private double binWidth;
    /**
     * Probability density function
     */
    private ProbabilityDensityFunction density;

    /**
     * @param pdf DhbStatistics.ProbabilityDensityFunction
     * @param n long
     * @param w double
     */
    public ScaledProbabilityDensityFunction(
        ProbabilityDensityFunction pdf, long n, double w)
    {
        density = pdf;
        setCount( n);
        binWidth = w;
    }
    /**
     * @param f statistics.ProbabilityDensity
     * @param hist curves.Histogram
     */
    public ScaledProbabilityDensityFunction (
        ProbabilityDensityFunction f, Histogram hist)
    {
        this( f, hist.count(), hist.getBinWidth());
    }
    /**
     * The array contains the parameters of the distribution
     * and the estimated number of events.
     * @return double[] an array containing the parameters of
     * the distribution.
     */
}

```

```

public double[] parameters()
{
    double[] parameters = density.parameters();
    double[] answer = new double[ parameters.length + 1];
    for ( int i = 0; i < parameters.length; i++ )
        answer[i] = parameters[i];
    answer[parameters.length] = count;
    return answer;
}
/**
 * @param x double    total count in the receiver
 */
public void setCount(double x)
{
    count = x;
}
/**
 * @param n int    total count in the receiver
 */
public void setCount(int n)
{
    count = n;
}
/**
 * @param n int    total count in the receiver
 */
public void setCount(long n)
{
    count = n;
}
/**
 * @param p double[]  assigns the parameters
 */
public void setParameters( double[] params)
{
    count = params[params.length-1];
    density.setParameters( params);
}
/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    sb.append("Scaled ");
    sb.append(density);
}

```

```

        return sb.toString();
    }
    /**
     * @return double the value of the function.
     * @param x double
     */
    public double value( double x)
    {
        return count * binWidth * density.value( x);
    }
    /**
     * Evaluate the function and the gradient of the function with respect
     * to the parameters.
     * @return double[] 0: function's value, 1,2,...,n function's gradient
     * @param x double
     */
    public double[] valueAndGradient( double x)
    {
        double[] dpG = density.valueAndGradient(x);
        double[] answer = new double[dpG.length+1];
        double r = binWidth * count;
        for ( int i = 0; i < dpG.length; i++ )
            answer[i] = dpG[i] * r;
        answer[dpG.length] = dpG[0] * binWidth;
        return answer;
    }
}

```

10.5 Definition of estimation

Let us assume that an observable quantity y is following a probability distribution described by a set of observable quantities $x_1, x_2 \dots$ - called the experimental conditions - and a set of parameters $p_1, p_2 \dots$. In other words, the probability density function of the random variable⁵ corresponding to the observable quantity y can be written as

$$P(y) = P(y; \mathbf{x}, \mathbf{p}), \quad (10.20)$$

where \mathbf{x} is the vector $(x_1, x_2 \dots)$ and \mathbf{p} the vector $(p_1, p_2 \dots)$.

The estimation of the values of the parameters $p_1, p_2 \dots$ is the determination of the parameters $p_1, p_2 \dots$ by performing several measurements of the observable y for different experimental conditions $x_1, x_2 \dots$.

Let N be the number of measurements; let y_i be the i^{th} measured value of the observable y under the experimental conditions \mathbf{x}_i .

⁵For simplicity we shall use the same notation for the random variable and the observable quantity.

10.5.1 Maximum likelihood estimation

The maximum likelihood estimation of the parameters \mathbf{p} is the set of values $\bar{\mathbf{p}}$ maximizing the following function:

$$L(\mathbf{p}) = \prod_{i=1}^N P(y_i; \mathbf{x}_i, \mathbf{p}) \quad (10.21)$$

By definition the likelihood function $L(\mathbf{p})$ is the probability of making the N measurements. The maximum likelihood estimation determines the estimation $\bar{\mathbf{p}}$ of the parameters \mathbf{p} such that the series of measurements performed is the most probable, hence the name maximum likelihood.

One can show that the maximum likelihood estimation is robust and unbiased. The robustness and the bias of an estimation are defined mathematically. For short, robust means that the estimation converges toward the true value of the parameters for an infinite number of measurements; unbiased means that the deviations between the estimated parameters and their true value are symmetrically distributed around 0 for any finite number of measurements.

Equation 10.21 is often rewritten in logarithmic form to ease the computation of the likelihood function.

$$I(\mathbf{p}) = \ln L(\mathbf{p}) = \sum_{i=1}^N \ln P(y_i; \mathbf{x}_i, \mathbf{p}) \quad (10.22)$$

The function $I(\mathbf{p})$ is related to information and is used in information theory.

10.5.2 Least square estimation

Let us assume that the random variable y is distributed according to a normal distribution of given standard deviation σ and that the average of the normal distribution is given by a function $F(\mathbf{x}, \mathbf{p})$ of the experimental conditions and the parameters. In this case equation 10.20 becomes:

$$P(y) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{[y - F(\mathbf{x}, \mathbf{p})]^2}{2\sigma^2}} \quad (10.23)$$

Plugging equation 10.23 into equation 10.22 yields:

$$I(\mathbf{p}) = -N\sqrt{2\pi\sigma^2} - \sum_{i=1}^N \frac{[y - F(\mathbf{x}, \mathbf{p})]^2}{2\sigma^2} \quad (10.24)$$

The problem of finding the maximum of $I(\mathbf{p})$ is now equivalent to the problem of finding the minimum of the function:

$$S_{ML}(\mathbf{p}) = \sum_{i=1}^N \frac{[y - F(\mathbf{x}, \mathbf{p})]^2}{\sigma^2}, \quad (10.25)$$

where a redundant factor 2 has been removed. This kind of estimation is called least square estimation. Written as in equation 10.25 least square estimation is fully

equivalent to maximum likelihood estimation. By definition, the quantity $S_{ML}(\mathbf{p})$ is distributed as a χ^2 random variable with $N - m$ degrees of freedom where m is the number of parameters, that is the dimension of the vector \mathbf{p} .

In practice, however, the standard deviation σ is not known and frequently depends on the parameters \mathbf{p} . In that case, one uses instead an estimation for the standard deviation. Either the standard deviation of each measurement is determined experimentally by making several measurements under the same experimental conditions or it is estimated from the measurement error. Then, equation 10.25 can be rewritten as:

$$S(\mathbf{p}) = \sum_{i=1}^N \frac{[y - F(\mathbf{x}, \mathbf{p})]^2}{\sigma_i^2}. \quad (10.26)$$

The least square estimation is obtained by minimizing the quantity $S(\mathbf{p})$ with respect to the parameters \mathbf{p} . This kind of estimation can be used to determine the parameters of a functional dependence of the variable y from the observable quantities \mathbf{x} . For this reason it is also called a least square fit when it is used to fit the parameter of a functional dependence to the measurements.

In general the distribution of the random variable y may not be a normal distribution. One can nevertheless show that the least square estimation is robust. However, it is biased. Depending on the nature of the distribution of the random variable y the parameters may be over- or underestimated. This is especially the case when working with histograms.

We have said that all measurements of the observable quantities y must be distributed according to a normal distribution so that the quantity $S(\mathbf{p})$ of equation 10.26 is distributed as a χ^2 random variable. In general this is often the case⁶ when dealing with a large quantity of measurements. Thus, a least square fit is also called a χ^2 fit. In this case one can apply the χ^2 -test described in section 10.3 to assess the goodness of the fitted function.

If $S(\mathbf{p})$ has a minimum respective to \mathbf{p} then all partial derivatives of the function $S(\mathbf{p})$ respective to each of the components of the vector \mathbf{p} are zero. Since the function is positive and quadratic in \mathbf{p} , it is clear that the function must have at least one minimum. Under this circumstances the minimum can be obtained by solving the following set of equations:

$$\frac{\partial}{\partial p_j} F(\mathbf{x}_i; p_1, \dots, p_m) = 0 \quad \text{for } j = 1, \dots, m \quad (10.27)$$

where m is the number of parameters, that is the dimension of the vector \mathbf{p} . When a solution is found, one should in principle verify that it is really a minimum. Solving equation 10.27 gives the following system of equations:

$$\sum_{i=1}^N \frac{y - F(\mathbf{x}, \mathbf{p})}{\sigma_i^2} \cdot \frac{\partial}{\partial p_j} F(\mathbf{x}_i; p_1, \dots, p_m) = 0 \quad \text{for } j = 1, \dots, m \quad (10.28)$$

Once the system above has been solved, one can compute the value of $S(\mathbf{p})$ using equations 10.26 or, better, the value $S_{ML}(\mathbf{p})$ using equation 10.25. Computing the χ^2

⁶This is a consequence of a theorem known as the law of large numbers.

confidence level of that value (c.f. section 10.3) using a χ^2 distribution with $N - m$ degrees of freedom gives the probability that the fit is acceptable.

10.6 Least square fit with linear dependence

If the function $F(\mathbf{x}, \mathbf{p})$ is a linear function of the vector \mathbf{p} , it can be written in the following form:

$$F(\mathbf{x}, \mathbf{p}) = \sum_{j=1}^m f_j(\mathbf{x}) \cdot p_j. \quad (10.29)$$

In that case, equation 10.28 become a system of linear equations of the form:

$$\mathbf{M} \cdot \mathbf{p} = \mathbf{c}, \quad (10.30)$$

where the coefficients of the matrix \mathbf{M} are given by:

$$M_{jk} = \sum_{i=1}^N \frac{f_j(\mathbf{x}_i) f_k(\mathbf{x}_i)}{\sigma_i^2} \quad \text{for } j, k = 1, \dots, m, \quad (10.31)$$

and the components of the constant vector \mathbf{c} are given by:

$$c_j = \sum_{i=1}^N \frac{y_i f_j(\mathbf{x}_i)}{\sigma_i^2} \quad \text{for } j = 1, \dots, m. \quad (10.32)$$

Equation 10.30 is a system of linear equation which can be solved according to the algorithms exposed in sections 8.2 and 8.3. If one is interested only in the solution this is all there is to do.

A proper fit, however, should give an estimation of the error in estimating the parameters. The inverse of the matrix \mathbf{M} is the error matrix for the fit. The error matrix is used to compute the estimation of variance on the function $F(\mathbf{x}, \mathbf{p})$ as follows:

$$\text{var}[F(\mathbf{x}, \mathbf{p})] = \sum_{j=1}^m \sum_{k=1}^m M_{jk}^{-1} f_j(\mathbf{x}) f_k(\mathbf{x}). \quad (10.33)$$

The estimated error on the function $F(\mathbf{x}, \mathbf{p})$ is the square root of the estimated variance.

The diagonal elements of the error matrix are the variance of the corresponding parameter. That is:

$$\text{var}(p_j) = M_{jj}^{-1} \quad \text{for } j = 1, \dots, m. \quad (10.34)$$

The off diagonal elements describe the correlation between the errors on the parameters. One defines the correlation coefficient of parameter p_j and p_k by:

$$\text{cor}(p_j, p_k) = \frac{M_{jk}^{-1}}{\sqrt{M_{jj}^{-1} M_{kk}^{-1}}} \quad \text{for } j, k = 1, \dots, m \text{ and } j \neq k. \quad (10.35)$$

All correlation coefficients are comprised between -1 and 1. If the absolute value of a correlation coefficient is close to 1, it means that one of the two corresponding two parameters is redundant for the fit. In other word, one parameter can be expressed as a function of the other.

10.7 Linear regression

A linear regression is a least square fit with a linear function of a single variable. The dimension of the vector \mathbf{x} is one and the dimension of the vector \mathbf{p} is two. The function to fit has only two parameters. The following convention is standard:

$$\begin{cases} p_1 &= a, \\ p_2 &= b, \\ F(\mathbf{x}, \mathbf{p}) &= ax + b. \end{cases} \quad (10.36)$$

With these definitions, the system of equations 10.30 becomes:

$$\begin{cases} \sum_{i=1}^N \frac{x_i^2}{\sigma^2} a + \sum_{i=1}^N \frac{x_i}{\sigma^2} b &= \sum_{i=1}^N \frac{x_i y_i}{\sigma^2} \\ \sum_{i=1}^N \frac{x_i}{\sigma^2} a + \sum_{i=1}^N \frac{1}{\sigma^2} b &= \sum_{i=1}^N \frac{y_i}{\sigma^2}. \end{cases} \quad (10.37)$$

This system can easily be solved. Before giving the solution, let us introduce a short hand notation for the weighted sums:

$$\langle Q \rangle = \sum_{i=1}^N \frac{Q_i}{\sigma_i^2}. \quad (10.38)$$

Using this notation the solution of the system of equations 10.37 can be written as:

$$\begin{cases} a &= \frac{\langle xy \rangle \cdot \langle 1 \rangle - \langle x \rangle \cdot \langle y \rangle}{\langle xx \rangle \cdot \langle 1 \rangle - \langle x \rangle \cdot \langle x \rangle} \\ b &= \frac{\langle xx \rangle \cdot \langle y \rangle - \langle xy \rangle \cdot \langle x \rangle}{\langle xx \rangle \cdot \langle 1 \rangle - \langle x \rangle \cdot \langle x \rangle} \end{cases} \quad (10.39)$$

where the symmetry of the expression is quite obvious. It is interesting to note that if we had fitted x as a linear function of y , that is $x = \tilde{a}y + \tilde{b}$, we would have the following expression for the slope:

$$\tilde{a} = \frac{\langle xy \rangle \cdot \langle 1 \rangle - \langle x \rangle \cdot \langle y \rangle}{\langle yy \rangle \cdot \langle 1 \rangle - \langle y \rangle \cdot \langle y \rangle}. \quad (10.40)$$

If the dependence between x and y is truly a linear function, the product $a\tilde{a}$ ought to be 1. The square root of the product $a\tilde{a}$ is defined as the correlation coefficient of the linear regression, the sign of the square root being the sign of the slope. The correlation coefficient r is thus given by:

$$r = \frac{\langle xy \rangle \cdot \langle 1 \rangle - \langle x \rangle \cdot \langle y \rangle}{\sqrt{(\langle xx \rangle \cdot \langle 1 \rangle - \langle x \rangle \cdot \langle x \rangle)(\langle yy \rangle \cdot \langle 1 \rangle - \langle y \rangle \cdot \langle y \rangle)}}. \quad (10.41)$$

Since the least square fit is a biased estimator for the parameters, the square of the correlation coefficient is less than 1 in practice. The value of the correlation coefficient lies between -1 and 1. A good linear fit ought to have the absolute value of r close to 1.

Finally the error matrix of a linear regression is given by:

$$\mathbf{M}^{-1} = \frac{1}{\langle xx \rangle \cdot \langle 1 \rangle - \langle x \rangle \cdot \langle x \rangle} \begin{pmatrix} \langle xx \rangle & -\langle x \rangle \\ -\langle x \rangle & \langle 1 \rangle \end{pmatrix} \quad (10.42)$$

when the vector representing the parameters of the fit is defined as (b, a) in this order.

When fitting a functional dependence with one variable, x and many parameters, one can use a linear regression to reduce rounding errors when the observed values y_1, \dots, y_N cover a wide numerical range. Let a and b be the result of the linear regression of the values y_i as a function of x_i . One defines the new quantity $y'_i = y_i - (ax_i + b)$ for all i . The standard deviation of y'_i is the same as that of y_i since the subtracted expression is just a change of variable. In fact, the linear regression does not need to be a good fit at all. Then, the functional dependence can be fitted on the quantities y'_1, \dots, y'_N . We shall give a detailed example on this method in section 10.8

10.7.1 Linear regression — General implementation

Linear regression is implemented within a single class using a similar implementation as that of the statistical moments. This means that individual measurements are accumulated and not stored. The drawback is that the object cannot compute the confidence level of the fit. This is not so much a problem since the correlation coefficient is usually sufficient to estimate the goodness of the fit.

The class has the following instance variables:

- `sum1` is used to accumulate the sum of weights, that is, $\langle 1 \rangle$,
- `sumX` is used to accumulate the weighted sum of x_i , that is, $\langle x \rangle$,
- `sumY` is used to accumulate the weighted sum of y_i , that is, $\langle y \rangle$,
- `sumXY` is used to accumulate the weighted sum of $x_i \times y_i$, that is, $\langle xy \rangle$,
- `sumXX` is used to accumulate the weighted sum of x_i^2 , that is, $\langle xx \rangle$,
- `sumYY` is used to accumulate the weighted sum of y_i^2 , that is, $\langle yy \rangle$,
- `slope` the slope of the linear regression, that is, a ,
- `intercept` the value of the linear regression at $x = 0$, that is, b ,
- `correlationCoefficient` the correlation coefficient, that is, r in equation 10.41.

When either one of the instance variables `slope`, `intercept` or `correlationCoefficient` is needed, the method `computeResults` calculating the values of the three instance variables is called using lazy initialization. When new data is added to the object, these variables are reset. It is thus possible to investigate the effect of adding new measurements on the results.

The methods `asPolynomial` and `asEstimatedPolynomial` return an object used to compute the predicted value for any x . The estimated polynomial is using the error matrix of the least square fit to compute the error on the predicted value. Estimated polynomials are explained in section 10.8

Figure 10.1 with the box `LinearRegression` grayed.

10.7.2 Linear regression — Smalltalk implementation

Listing 10.14 shows the complete implementation in Smalltalk. The following code shows how to use the class `DhbLinearRegression` to perform a linear regression over a series of measurements .

Code example 10.5

```

| linReg valueStream measurement slope intercept
  correlationCoefficient estimation value error|
linReg := DhbLinearRegression new.
[ valueStream atEnd]
  whileFalse:[ measurement := valueStream next.
              linReg addPoint: measurement point
                weight: measurement weight
              ].
slope := linReg slope.
intercept := linReg intercept.
correlationCoefficient := linReg correlationCoefficient.
estimation := linReg asEstimatedPolynomial.
value := estimation value: 0.5.
error := estimation error: 0.5.

```

This example assumes that the measurement of the random variable are obtained from a stream. The exact implementation of the stream is not shown here. The first line after the declaration creates a new instance of class `DhbLinearRegression`. Next comes the loop over all values found in the stream. This examples assumes that the values are stored on the stream as a single object implementing the following methods:

`point` returns a point containing the measurement, that is, the pair (x_i, y_i) for all i ,

`weight` returns the weight of the measurement, that is, $1/\sigma_i^2$.

Each point is accumulated into the linear regression object with the method `addPoint:weight:`.

After all measurements have been read, the results of the linear regression are fetched. The last three lines show how to obtain a polynomial object used to compute the value predicted by the linear regression at $x = 0.5$ and the error on that prediction.

The mechanism of lazy initialization is implemented by setting the three instance variables `slope`, `intercept` and `correlationCoefficient` to `nil` in the method `reset`.

Listing 10.14 Smalltalk implementation of linear regression

Class	DhbLinearRegression
Subclass of	Object
Instance variable names:	sum1 sumX sumY sumXX sumYY sumXY slope intercept correlationCoefficient
Class methods	
new	

```
^( super new) reset; yourself
```

Instance methods

```
add: aPoint
```

```
self add: aPoint weight: 1.
```

```
add: aPoint weight: aNumber
```

```
sum1 := sum1 + aNumber.
sumX := sumX + (aPoint x * aNumber).
sumY := sumY + (aPoint y * aNumber).
sumXX := sumXX + (aPoint x squared * aNumber).
sumYY := sumYY + (aPoint y squared * aNumber).
sumXY := sumXY + (aPoint x * aPoint y * aNumber).
self resetResults
```

```
asEstimatedPolynomial
```

```
^( DhbEstimatedPolynomial coefficients: self coefficients)
  errorMatrix: self errorMatrix;
  yourself
```

```
asPolynomial
```

```
^DhbPolynomial coefficients: self coefficients
```

```
coefficients
```

```
^Array with: self intercept with: self slope
```

```
computeResults
```

```
| xNorm xyNorm |
xNorm := sumXX * sum1 - (sumX * sumX).
xyNorm := sumXY * sum1 - (sumX * sumY).
slope := xyNorm / xNorm.
intercept := (sumXX * sumY - (sumXY * sumX)) / xNorm.
correlationCoefficient := xyNorm
  / (xNorm * (sumYY * sum1 - (sumY * sumY))) sqrt
```

```
correlationCoefficient
```

```
correlationCoefficient isNil
  ifTrue: [ self computeResults].
^correlationCoefficient
```

errorMatrix

```
| c1 cx cxx |
c1 := 1.0 / (sumXX * sum1 - sumX squared).
cx := sumX negated * c1.
cxx := sumXX * c1.
c1 := sum1 * c1.
^DhbSymmetricMatrix rows: (Array with: (Array with: cxx with: cx)
  with: (Array with: cx with: c1))
```

errorOnIntercept

```
^(sumXX / (sumXX * sum1 - sumX squared)) sqrt
```

errorOnSlope

```
^(sum1 / (sumXX * sum1 - sumX squared)) sqrt
```

intercept

```
intercept isNil
  ifTrue: [ self computeResults].
^intercept
```

remove: aPoint

```
sum1 := sum1 - 1.
sumX := sumX - aPoint x.
sumY := sumY - aPoint y.
sumXX := sumXX - aPoint x squared.
sumYY := sumYY - aPoint y squared.
sumXY := sumXY - (aPoint x * aPoint y).
self resetResults
```

reset

```
sum1 := 0.
sumX := 0.
sumY := 0.
sumXX := 0.
```

```

sumYY := 0.
sumXY := 0.
self resetResults

```

```

resetResults

```

```

slope := nil.
intercept := nil.
correlationCoefficient := nil.

```

```

slope

```

```

slope isNil
  ifTrue: [ self computeResults].
^slope

```

```

value: aNumber

```

```

^aNumber * self slope + self intercept

```

10.7.3 Linear regression — Java implementation

Listing 10.15 shows the complete implementation in Java. The following code shows how to use the class `LinearRegression` to perform a linear regression over a series of measurements .

Code example 10.6

```

double[] x, y, w;

<Gathering measurements into the arrays x ( $x_i$ ), y ( $y_i$ )
and w ( $1/\sigma_i^2$ )>

LinearRegression linReg = new LinearRegression();
for ( int i = 1, i < x.length, i++ )
    linReg.add( x[i], y[i], w[i]);
double slope = linReg.slope();
double intercept = linReg.intercept();
double correlationCoefficient = linReg.correlationCoefficient();
EstimatedPolynomial estimation = linReg.asEstimatedPolynomial();
double value = estimation.value(0.5);
double error = estimation.error(0.5);

```

This example assumes that the measurements and the weights are gathered into 3 arrays. The exact implementation of the gathering is not shown here. The first line after the gathering creates a new instance of class `LinearRegression`. Next comes the loop accumulating the measurements into the linear regression object with the method `add`.

After all measurements have been accumulated, the results of the linear regression are fetched. The last three lines show how to obtain a polynomial object used to compute the value predicted by the linear regression at $x = 0.5$ and the error on that prediction.

The mechanism of lazy initialization is implemented by setting the three instance variables `slope`, `intercept` and `correlationCoefficient` to `NaN` in the method `reset`.

Listing 10.15 Java implementation of linear regression

```
package DhbEstimation;

import DhbFunctionEvaluation.PolynomialFunction;
import DhbInterfaces.PointSeries;
import DhbMatrixAlgebra.SymmetricMatrix;
import DhbMatrixAlgebra.DhbIllegalDimension;
import DhbMatrixAlgebra.DhbNonSymmetricComponents;
/**
 * linear regression
 *
 * @author Didier H. Besset
 */
public class LinearRegression
{
/**
 * Number of accumulated points
 */
    private int sumI;
/**
 * Sum of X
 */
    private double sumX;
/**
 * Sum of Y
 */
    private double sumY;
/**
 * Sum of XX
 */
    private double sumXX;
/**
 * Sum of XY
 */
```

```

    private double sumXY;
/**
 * Sum of YY
 */
    private double sumYY;
/**
 * Slope
 */
    private double slope;
/**
 * Intercept
 */
    private double intercept;
/**
 * Correlation coefficient
 */
    private double correlationCoefficient;
/**
 * Cnstructor method.
 */
public LinearRegression()
{
    super();
    reset();
}
/**
 * @param x double
 * @param y double
 */
public void add( double x, double y)
{
    add( x, y, 1);
}
/**
 * @param x double
 * @param y double
 * @param w double
 */
public void add( double x, double y, double w)
{
    double wx = w * x;
    double wy = w * y;
    sum1 += w;
    sumX += wx;
    sumY += wy;
    sumXX += wx * x;

```



```

        sumYY += wy * y;
        sumXY += wx * y;
        resetResults();
    }
    /**
     * @return DhbFunctionEvaluation.PolynomialFunction
     */
    public EstimatedPolynomial asEstimatedPolynomial()
    {
        return new EstimatedPolynomial( coefficients(), errorMatrix());
    }
    /**
     * @return DhbFunctionEvaluation.PolynomialFunction
     */
    public PolynomialFunction asPolynomial()
    {
        return new PolynomialFunction( coefficients());
    }
    /**
     * @return double[]
     */
    private double[] coefficients()
    {
        double[] answer = new double[2];
        answer[0] = getIntercept();
        answer[1] = getSlope();
        return answer;
    }
    private void computeResults()
    {
        double xNorm = sumXX * sum1 - sumX * sumX;
        double xyNorm = sumXY * sum1 - sumX * sumY;
        slope = xyNorm / xNorm;
        intercept = ( sumXX * sumY - sumXY * sumX ) / xNorm;
        correlationCoefficient = xyNorm / Math.sqrt( xNorm *
            ( sumYY * sum1 - sumY * sumY ));
    }
    /**
     * @return DhbMatrixAlgebra.SymmetricMatrix
     */
    public SymmetricMatrix errorMatrix()
    {
        double[][] rows = new double[2][2];
        rows[1][1] = 1./ ( sumXX * sum1 - sumX * sumX);
        rows[0][1] = sumXX * rows[1][1];
        rows[1][0] = rows[0][1];
    }

```

```

    rows[0][0] = sumXX * rows[1][1];
    SymmetricMatrix answer = null;
    try { try { answer = SymmetricMatrix.fromComponents( rows);
        } catch( DhbIllegalDimension e){ };
        } catch( DhbNonSymmetricComponents e){ };
    return answer;
}
/**
 * @return double
 */
public double getCorrelationCoefficient()
{
    if( Double.isNaN( correlationCoefficient ) )
        computeResults();
    return correlationCoefficient;
}
/**
 * @return double
 */
public double getIntercept()
{
    if( Double.isNaN( intercept ) )
        computeResults();
    return intercept;
}
/**
 * @return double
 */
public double getSlope()
{
    if( Double.isNaN( slope ) )
        computeResults();
    return slope;
}
/**
 * @param x double
 * @param y double
 */
public void remove( double x, double y)
{
    sumI -= 1;
    sumX -= x;
    sumY -= y;
    sumXX -= x * x;
    sumYY -= y * y;
    sumXY -= x * y;
}

```

```

    resetResults();
}
public void reset()
{
    sum1 = 0;
    sumX = 0;
    sumY = 0;
    sumXX = 0;
    sumYY = 0;
    sumXY = 0;
    resetResults();
}
private void resetResults()
{
    slope = Double.NaN;
    intercept = Double.NaN;
    correlationCoefficient = Double.NaN;
}
/**
 * @return double
 * @param x double
 */
public double value( double x)
{
    return x * getSlope() + getIntercept();
}
}

```

10.8 Least square fit with polynomials

In a polynomial fit the fit function is a polynomial of degree m . In this case, the parameters are usually numbered starting from 0; the number of free parameters is $m + 1$ and the number of degrees of freedom is $N - m - 1$. We have:

$$F(x; p_0, p_1, \dots, p_m) = \sum_{k=0}^m p_k x^k. \quad (10.43)$$

The partial derivative of equation 10.27 is easily computed since a polynomial is a linear function of its coefficients:

$$\frac{\partial}{\partial p_j} F(\mathbf{x}_i; p_1, \dots, p_m) = x_i^j \quad \text{for } j = 1, \dots, m. \quad (10.44)$$

Such a matrix is called a Van Der Monde matrix. The system of equations 10.28 then becomes:

$$\sum_{k=0}^m p_k \cdot \sum_{i=1}^N \frac{x_i^{j+k}}{\sigma_i^2} = \sum_{i=1}^N \frac{x_i^j y_i}{\sigma_i^2}. \quad (10.45)$$




Figure 10.5: Example of polynomial fit

Equation 10.45 is of the same form as equation 10.30 where the coefficients of the matrix M are given by:

$$M_{jk} = \sum_{i=1}^N \frac{x_i^{j+k}}{\sigma_i^2}, \quad (10.46)$$

and the vector c has for components:

$$c_j = \sum_{i=1}^N \frac{x_i^j y_i}{\sigma_i^2}. \quad (10.47)$$

Polynomial least square fit provides a way to construct an ad-hoc representation of a functional dependence defined by a set of point. Depending on the type of data it can be more efficient than the interpolation methods discussed in chapter 3. In general the degree of the polynomial should be kept small to prevent large fluctuations between the data points.

Let us now shows a concrete example of polynomial fit.

In order to determine whether or not a fetus is developing itself normally within the womb, the dimension of the fetus' bones are measured during an ultrasound examination of the mother-to-be. The dimensions are compared against a set of standard data measured on a control population. Such data⁷ are plotted in figure 10.5: the y -axis is the length of the femur expressed in mm; the x -axis represents the duration in weeks of the pregnancy based on the estimated date of conception. Each measurement has been

⁷These numbers are reproduced with permission of Prof. P.J. Steer. from the department of obstetrics and gynecology of the Chelsea & Westminster Hospital of London.

obtained by measuring the length of different fetuses at the same gestational age. The measurement are averaged and the error on the average is also calculated (c.f. section 9.1).

The obtained data do not follow a smooth curve since the data have been determined experimentally. The exact date of conception cannot be exactly determined so some fluctuation is expected, but the major limitation of the data is to obtain an sufficient number of measurements to smooth out the natural variations between individuals. The data from figure 10.5 have been fitted with a second order polynomial: the result is shown with a black thin curve. As one can see, the fit is excellent in spite of the fluctuation of the measurements. Figure 10.6 shows the fit results. This χ^2 confidence level is rather good. However, the correlation coefficients are quite high. This is usually

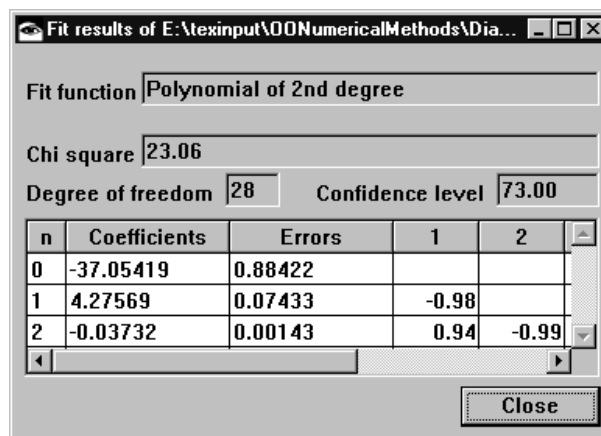


Figure 10.6: Fit results for the fit of figure 10.5

the case with polynomial fits as each coefficient strongly depends on the other.

The thick gray line on figure 10.5 shows the interpolation polynomial⁸ for comparison (interpolation is discussed in chapter 3). As the reader can see the interpolation polynomial gives unrealistic results because of the fluctuations of the experimental data.

Polynomial fits have something in common with interpolation: a fitted polynomial can seldom be used to extrapolate the data outside of the interval defined by the reference points. This is illustrated on figure 10.7 showing a second order polynomial fit made on a series of points. This series is discussed in section 3.1 (figure 3.1). The reader can see that the fitted polynomial does not reproduce the behavior of the data in the lower part of the region. Nevertheless, the data are well within the estimated error. Thus, the fit results are consistent. Unfortunately too many fit results are presented without their estimated error. This kind of information, however, is an essential part of a fit and should always be deliver along with the fitted function.

This is the idea behind what we call estimated polynomials. An estimated polynomial

⁸An attentive reader will notice that the interpolation curve dos not go through some data points. This is an artifact of the plotting over a finite sample of points which do not coincide with the measured data.




Figure 10.7: Limitation of polynomial fits

is a polynomial whose coefficients have been determined by a least square fit. An estimated polynomial keeps the error matrix of the fit is along with the coefficients of the fitted polynomial.

Note: The Smalltalk and Java implementations are quite different from each other. The reader is advised to read both sections. As an exercise, the reader may attempt to write the alternate implementation in the desired language.

10.8.1 Polynomial least square fits — Smalltalk implementation

Listing 10.16 shows the complete implementation in Smalltalk. The following code example show how to perform the fit made in figure 10.5.

Figure 10.1 with the boxes PolynomialLeastSquareFit and EstimatedPolynomial grayed.

Code example 10.7

```
| fit valueStream dataHolder estimation value error|
    <Accumulation of data into dataHolder>

    fit := DhhPolynomialLeastSquareFit new: 2.
    dataHolder pointsAndErrorsDo: [ :each | fit add: each].
    estimation := fit evaluate.
    value := estimation value: 20.5.
    error := estimation error: 20.5.
```

The data are accumulated into a object called dataHolder implementing the iterator method pointsAndErrorsDo. The argument of the block used by this method is an

instance of class `DhbWeightedPoint` described in section 10.3.3. The iterator method acts on all experimental data stored in `dataHolder`. Next, an instance of class `DhbPolynomialLeastSquareFit` is created. The argument of the method is the degree of the polynomial; here a second order polynomial is used. After data have been accumulated into this object, the fit is performed by sending the method `evaluate` to the fit object. This method returns a polynomial including the error matrix. The last three lines compute the predicted femur length and its error in the middle of the 20th week of pregnancy.

The Smalltalk implementation assumes the points are stored in an object implementing the iterator method `do:`. Any instance of `Collection` or its subclasses will work. Each element of the collection must be an array containing the values x_i , y_i and $1/\sigma_i$. The class `DhbPolynomialLeastSquareFit` keeps this collection in the instance variable `pointCollection`. A second instance variable, `degreePlusOne`, keeps the number of coefficients to be estimated by the fit.

The class creation method `new:` is used to create an instance by supplying the degree of the fit polynomial as argument. `pointCollection` is set to a new instance of an `OrderedCollection`. Then, new values can be added to the fit instance with the method `add:`.

The other class creation method, `new: on:` takes two arguments: the degree of the fit polynomial and the collection of points. The fit result can be fetched directly after the creation.

The method `evaluate` solves equation 10.28 by first computing the inverse of the matrix **M** to get the error matrix. The coefficients are then obtained from the multiplication of the constant vector by the error matrix.

Listing 10.16 Smalltalk implementation of a polynomial least square fit

```

Class                DhbPolynomialLeastSquareFit
Subclass of          Object
Instance variable names:  pointCollection degreePlusOne

Class methods
new: anInteger
    ^super new initialize: anInteger

new: anInteger on: aCollectionOfPoints
    ^super new initialize: anInteger on: aCollectionOfPoints

Instance methods
accumulate: aWeightedPoint into: aVectorOfVectors and: aVector
OfVectors and: aVector
    | t p powers |

```

```

p := 1.0.
powers := aVector collect: [ :each | t := p. p := p *
                           aWeightedPoint xValue. t].
aVector accumulate: powers * ( aWeightedPoint yValue *
                              aWeightedPoint weight).
1 to: aVector size do:
  [ :k |
    ( aVectorOfVectors at: k) accumulate: powers * ( ( powers
      at: k) * aWeightedPoint weight).
  ].

```

```
add: aWeightedPoint
```

```
  ^pointCollection add: aWeightedPoint
```

```
computeEquations
```

```

| rows vector |
vector := ( DhbVector new: degreePlusOne) atAllPut: 0 ; yourself.
rows := ( 1 to: degreePlusOne) collect: [ :k | ( DhbVector new:
  degreePlusOne) atAllPut: 0 ; yourself].
pointCollection do:
  [ :each | self accumulate: each into: rows and: vector].
^Array with: ( DhbSymmetricMatrix rows: rows) with: vector

```

```
evaluate
```

```

| system errorMatrix |
system := self computeEquations.
errorMatrix := ( system at: 1) inverse.
^( DhbEstimatedPolynomial coefficients: errorMatrix * (system at:
  2))
  errorMatrix: errorMatrix;
  yourself

```

```
initialize: anInteger
```

```
  ^self initialize: anInteger on: OrderedCollection new
```

```
initialize: anInteger on: aCollectionOfPoints
```

```

  pointCollection := aCollectionOfPoints.
  degreePlusOne := anInteger + 1.
  ^self

```

Listing 10.17 show the implementation of the class `DhbEstimatedPolynomial` which is a subclass of the class `DhbPolynomial` containing the error matrix of the fit performed to make a estimation of the polynomial's coefficients. The method `error`: returns the estimated error of its value based on the error matrix of the fit using equation 10.33. The convenience method `valueAndError`: returns an array containing the estimated value and its error in single method call. This is suitable for plotting the resulting curve.

Listing 10.17 Smalltalk implementation of a polynomial with error

```

Class                DhbEstimatedPolynomial
Subclass of          DhbPolynomial
Instance variable names:  errorMatrix

    Instance methods

error: aNumber
| errorVector term nextTerm |
nextTerm := 1.
errorVector := ( coefficients collect: [ :each | term :=
    nextTerm. nextTerm := aNumber * nextTerm. term ]) asVector.
^( errorVector * errorMatrix * errorVector ) sqrt

errorMatrix
    ^errorMatrix

errorMatrix: aMatrix
    errorMatrix := aMatrix.

valueAndError: aNumber
    ^Array with: ( self value: aNumber ) with: ( self error: aNumber )

```

Figure 10.1 with the boxes `PolynomialLeastSquareFit` and `EstimatedPolynomial` grayed.

10.8.2 Polynomial least square fits — Java implementation

Listing 10.18 shows the complete implementation in Java. The following code example show how to perform the fit made in figure 10.5.

Code example 10.8

```

double[] x, y, dy;

    <Gathering measurements into the arrays x ( $x_i$ ), y ( $y_i$ )
    and dy ( $\sigma_i$ )>

```

```

PolynomialLeastSquareFit fit = new PolynomialLeastSquareFit(2);
for( int i = 0; i < x.length; i++ )
    fit.accumulatePoint( x, y, w);
EstimatedPolynomial estimation = fit.evaluate();
double value = estimation.value(20.5);
double error = estimation.error(20.5);

```

The class `PolynomialLeastSquareFit` is organized like the linear regression class: data points are accumulated directly on each call as each value is not needed. One drawback is that the value of the sum of equation 10.26 must be computed by another object when the χ^2 -test of the fit is required. The advantage is that the implementation of the class is very simple. The class `PolynomialLeastSquareFit` has one constructor method whose argument is the degree of the polynomial used in the fit.

The main accumulation method, `accumulateWeightedPoint` takes the arguments x_i , y_i and $1/\sigma_i^2$ in this order. Other constructor methods are convenience methods calling the method `accumulateWeightedPoint` to do the real work:

`accumulatePoint(double,double,double)` to accumulate points with error, that is, the arguments are x_i , y_i and σ_i in this order

`accumulateBin(double,int)` to accumulate the bin contents of a histogram; the first argument is the middle of the bin and the second the bin contents; the weight is computed as described in section 10.4.

`accumulateAverage(double,StatisticalMoments)` to accumulate points whose y value is determined from an average over several measurements; the first argument is x_i ; y_i and σ_i are obtained respectively from the average and error on average of the statistical moments.

`accumulatePoint(double,double)` to accumulate a point with a weight of 1; the arguments are x_i and y_i .

The method `evaluate` returns an estimated polynomial corresponding to the fit. The system of equation 10.45 is solved using LUP decomposition since the dimension of the matrix is usually small⁹. The calculation of the error matrix is obtained from the LUP decomposition. To avoid possible rounding errors, the off-diagonal elements are made exactly symmetric before being passed to the constructor method of a symmetric matrix.

Listing 10.18 Java implementation of a polynomial least square fit

```
package DhbEstimation;
```

```
import DhbFunctionEvaluation.PolynomialFunction;
import DhbMatrixAlgebra.LUPDecomposition;
import DhbMatrixAlgebra.DhbVector;
```

⁹Fitted polynomials of very high degree have the same bad behavior as the interpolation polynomials

```

import DhbMatrixAlgebra.SymmetricMatrix;
import DhbMatrixAlgebra.DhbIllegalDimension;
import DhbMatrixAlgebra.DhbNonSymmetricComponents;
import DhbInterfaces.PointSeriesWithErrors;
import DhbStatistics.StatisticalMoments;
/**
 * Polynomial least square fit
 *
 * @author Didier H. Besset
 */
public class PolynomialLeastSquareFit
{
    double[][] systemMatrix;
    double[] systemConstants;
    /**
     * Constructor method.
     */
    public PolynomialLeastSquareFit( int n)
    {
        int n1 = n + 1;
        systemMatrix = new double[n1][n1];
        systemConstants = new double[n1];
        reset();
    }
    /**
     * @param x double
     * @param m StatisticalMoments
     */
    public void accumulateAverage( double x, StatisticalMoments m)
    {
        accumulatePoint( x, m.average(), m.errorOnAverage());
    }
    /**
     * @param x double
     * @param n int    bin content
     */
    public void accumulateBin( double x, int n)
    {
        accumulateWeightedPoint( x, n, 1.0 / Math.max( 1, n));
    }
    /**
     * @param x double
     * @param y double
     */
    public void accumulatePoint( double x, double y)
    {

```

```

    accumulateWeightedPoint( x, y, 1);
}
/**
 * @param x double
 * @param y double
 * @param error double    standard deviation on y
 */
public void accumulatePoint( double x, double y, double error)
{
    accumulateWeightedPoint( x, y, 1.0 / (error * error));
}
/**
 * @param x double
 * @param y double
 * @param w double    weight of point
 */
public void accumulateWeightedPoint( double x, double y, double w)
{
    double xp1 = w;
    double xp2;
    for ( int i = 0; i < systemConstants.length; i++ )
    {
        systemConstants[i] += xp1 * y;
        xp2 = xp1;
        for ( int j = 0; j <= i; j++ )
        {
            systemMatrix[i][j] += xp2;
            xp2 *= x;
        }
        xp1 *= x;
    }
}
/**
 * @return DhbEstimation.EstimatedPolynomial
 */
public EstimatedPolynomial evaluate()
{
    for ( int i = 0; i < systemConstants.length; i++ )
    {
        for ( int j = i + 1; j < systemConstants.length; j++ )
            systemMatrix[i][j] = systemMatrix[j][i];
    }
    try {
        LUPDecomposition lupSystem = new LUPDecomposition(
            systemMatrix);
        double [][] components = lupSystem.inverseMatrixComponents();
    }
}

```

```

        LUPDecomposition.symmetrizeComponents( components);
    return new EstimatedPolynomial(
        lupSystem.solve( systemConstants),
        SymmetricMatrix.fromComponents( components));
    } catch ( DhbIllegalDimension e) {}
    catch ( DhbNonSymmetricComponents ex) {};
    return null;
}
public void reset()
{
    for ( int i = 0; i < systemConstants.length; i++ )
    {
        systemConstants[i] = 0;
        for ( int j = 0; j < systemConstants.length; j++ )
            systemMatrix[i][j] = 0;
    }
}
}

```

Listing 10.19 show the implementation of the class `EstimatedPolynomial` which is a subclass of the class `Polynomial` containing the error matrix of the fit performed to make a estimation of the polynomial's coefficients. The class `EstimatedPolynomial` returns the estimated error of its value based on the error matrix of the fit using equation 10.33.

Listing 10.19 Java implementation of a polynomial with error

```

package DhbEstimation;

import DhbFunctionEvaluation.PolynomialFunction;
import DhbMatrixAlgebra.DhbVector;
import DhbMatrixAlgebra.SymmetricMatrix;
import DhbMatrixAlgebra.DhbIllegalDimension;
/**
 * Polynomial with error estimation
 *
 * @author Didier H. Besset
 */
public class EstimatedPolynomial extends PolynomialFunction
{
    /**
     * Error matrix.
     */
    SymmetricMatrix errorMatrix;
}

```

```

* Constructor method.
* @param coeffs double[]
* @param e double[] error matrix
*/
public EstimatedPolynomial(double[] coeffs, SymmetricMatrix e)
{
    super(coeffs);
    errorMatrix = e;
}
/**
* @return double
* @param x double
*/
public double error( double x)
{
    int n = degree() + 1;
    double[] errors = new double[n];
    errors[0] = 1;
    for ( int i = 1; i < n; i++)
        errors[i] = errors[i-1] * x;
    DhbVector errorVector = new DhbVector( errors);
    double answer;
    try { answer = errorVector.product(
        errorMatrix.product( errorVector));
    } catch (DhbIllegalDimension e) { answer = Double.NaN;};
    return Math.sqrt( answer);
}
}

```

10.9 Least square fit with non-linear dependence

In the case of a non-linear function, the fit can be reduced to a linear fit and a search by successive approximations.

Let us assume that we have an approximate estimation \mathbf{p}_0 of the parameters \mathbf{p} . Let us define the vector $\Delta\mathbf{p} = \mathbf{p} - \mathbf{p}_0$. One redefines the function $F(\mathbf{x}, \mathbf{p})$ as:

$$F(\mathbf{x}, \mathbf{p}) = F(\mathbf{x}, \mathbf{p}_0) + \left. \frac{\partial F(\mathbf{x}, \mathbf{p})}{\partial \mathbf{p}} \right|_{\mathbf{p}=\mathbf{p}_0} \cdot \Delta\mathbf{p}. \quad (10.48)$$

Equation 10.48 is a linear expansion¹⁰ of the function $F(\mathbf{x}, \mathbf{p})$ around the vector \mathbf{p}_0 respective to the vector \mathbf{p} . In equation 10.48 $\left. \frac{\partial F(\mathbf{x}, \mathbf{p})}{\partial \mathbf{p}} \right|_{\mathbf{p}=\mathbf{p}_0}$ is the gradient of the

¹⁰That is, the first couples of terms of a Taylor expansion of the function $F(\mathbf{x}, \mathbf{p})$ around the vector \mathbf{p}_0 in an m dimensional space.

function $F(\mathbf{x}, \mathbf{p})$ relative to the vector \mathbf{p} evaluated for $\mathbf{p} = \mathbf{p}_0$; this is a vector with the same dimension as the vector \mathbf{p} . Then, one minimizes the expression in equation 10.26 respective to the vector $\Delta\mathbf{p}$. This is of course a linear problem as described in section 10.6. Equation 10.30 becomes:

$$\mathbf{M} \cdot \Delta\mathbf{p} = \mathbf{c}, \quad (10.49)$$

where the components of the matrix \mathbf{M} are now defined by:

$$M_{jk} = \sum_{i=1}^N \frac{1}{\sigma_i^2} \cdot \left. \frac{\partial F(\mathbf{x}_i, \mathbf{p})}{\partial p_j} \right|_{\mathbf{p}=\mathbf{p}_0} \cdot \left. \frac{\partial F(\mathbf{x}_i, \mathbf{p})}{\partial p_k} \right|_{\mathbf{p}=\mathbf{p}_0} \quad (10.50)$$

and the components of the vector \mathbf{c} are defined by:

$$c_j = \sum_{i=1}^N \frac{y_i - F(\mathbf{x}_i, \mathbf{p}_0)}{\sigma_i^2} \cdot \left. \frac{\partial F(\mathbf{x}_i, \mathbf{p})}{\partial p_j} \right|_{\mathbf{p}=\mathbf{p}_0}. \quad (10.51)$$

The vector $\Delta\mathbf{p}$ is obtained by solving equation 10.49 using the algorithms described in sections 8.2 and 8.3. Then, we can use the vector $\mathbf{p}_0 + \Delta\mathbf{p}$ as the new estimate and repeat the whole process. One can show¹¹ that iterating this process converges toward the vector $\bar{\mathbf{p}}$ minimizing the function $S(\mathbf{p})$ introduced in equation 10.26.

As explained in section 10.6, the inverse of the matrix \mathbf{M} is the error matrix containing the variance of each parameter and their correlation. The expression for the estimated variance on the function $F(\mathbf{x}, \mathbf{p})$ becomes:

$$\text{var}[F(\mathbf{x}, \mathbf{p})] = \sum_{j=1}^m \sum_{k=1}^m M_{jk}^{-1} \cdot \frac{\partial F(\mathbf{x}_i, \bar{\mathbf{p}})}{\partial p_j} \cdot \frac{\partial F(\mathbf{x}_i, \bar{\mathbf{p}})}{\partial p_k}. \quad (10.52)$$

A careful examination of the error matrix can tell whether or not the fit is meaningful.

Figure 10.8 shows an example of a least square fit performed on a histogram with a probability density function. The histogram of figure 10.8 was generated using a random generator distributed according to a Fisher-Tippett distribution (c.f. D.4) with parameters $\alpha = 0$ and $\beta = 1$. Only 1000 events have been accumulated into the histogram. The inset window in the upper right corner shows the fit results. The order of the parameter are α, β and the number of generated events. The solid curve laid onto the histogram is the prediction of the fitted function; the two dotted lines indicate the error on the prediction. The reader can verify that the fit is excellent. The number of needed iterations is quite low: the convergence of the algorithm is quite good in general.

10.9.1 Non-linear fit — General implementation

As we have seen the solution of a non-linear fit can be approximated by successive approximations. Thus, non-linear fits are implemented with a subclass of the iterative process class described in section 4.1. Data points must be kept in a structure maintained

Figure 10.1 with the box
LeastSquareFit grayed.

¹¹A mathematically oriented reader can see that this is a generalization of the Newton zero-finding algorithm (c.f. section 5.3) to m dimensions.

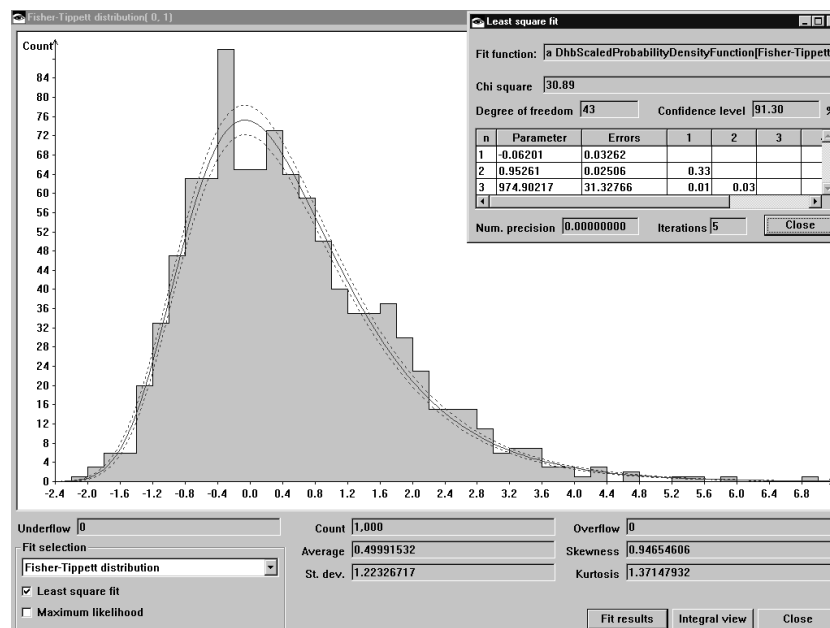


Figure 10.8: Example of a least square fit

by the object implementing linear least square fit to be readily available at each iteration. Thus, the data point are kept in an instance variable.

The result of the iterative process are the parameters. Our implementation assumes that the supplied function contains and maintains its parameter. Thus, the instance variable corresponding to the result of the iterative process is the fit function itself. The parameters determined by the fit — the result proper — are kept within the object implementing the supplied function. In particular the determination of the initial values for the iterative process are the responsibility of the fit function. Thus, the method `initializeIterations` does not do anything.

In most cases, the number of parameters in a least square fits is relatively small. Thus, LUP decomposition — described in section 8.3 — is sufficient to solve equation 10.49 at each iteration. Except for the last iteration, there is no need to compute the error matrix (the inverse of the matrix M). The components of the error matrix can be obtained from the LUP decomposition when the algorithm converges.

Convergence is attained when the largest of the relative variation of the components of the vector becomes smaller than a given value. In addition to the fact that we are dealing with floating point numbers, the reason for using relative precision is that the components of the vector p usually have different ranges.

When the fit has been obtained, convenience methods allows to retrieve the sum of equation 10.25 (`chiSquare`) and the confidence level of the fit (`confidenceLevel`). Another convenience method, `valueAndError` computes the prediction of the fit and its estimated error using equation 10.52.

Note: The details of implementation are quite different between Smalltalk and Java. I therefore advise the reader to read both. Some are admittedly arbitrary, some are caused by the specific features of each language. A good exercise is to explain the reasons for the differences.

10.9.2 Non-linear fit — Smalltalk implementation

Listing 10.20 shows the complete implementation in Smalltalk. The following code example shows how the data of figure 10.8 were generated¹².

Code example 10.9

```
| genDistr hist fit |
hist := DhbHistogram new.
hist freeExtent: true.
genDistr := DhbFisherTippettDistribution shape: 0 scale: 1.
1000 timesRepeat: [ hist accumulate: genDistr random].
fit := DhbLeastSquareFit histogram: hist
      distributionClass: DhbFisherTippettDistribution.
fit evaluate.
```

The first two lines after the declaration define an instance of class `DhbHistogram` with automatic adjustment of the limits (c.f. section 9.3.2). The next line defines an instance of a Fisher-Tippett distribution. Then, 1000 random numbers generated according to this distribution are accumulated into the histogram. Next, an instance of the class `DhbLeastSquareFit` is defined with the histogram for the data points and the desired class of the probability density function. The corresponding scaled probability is created within the method (c.f. listing 10.20). the final line performs the fit proper. After this line, the calling application can either use the fit object or extract the fit result to make predictions with the fitted distribution.

The class `DhbLeastSquareFit` is a subclass of `DhbIterativeProcess` described in section 4.1.1. It has the following instance variables:

`dataHolder` is the object containing the experimental data; this object must implement the iterator method `pointsAndErrorsDo;`; the block supplied to the iterator method takes as argument an instance of the class `DhbWeightedPoint` described in section 10.3.3;

`equations` contains the components of the matrix M ;

`constants` contains the components of the vector c ;

`errorMatrix` contains the LUP decomposition of the matrix M ;

`chiSquare` contains the sum of equation 10.25;

`degreeOfFreedom` contains the degree of freedom of the fit.

¹². . . up to the plotting facilities. This could be the topic of a future book.

accumulateEquationSystem

```
dataHolder pointsAndErrorsDo: [ :each | self accumulate: each].
```

chiSquare

```
chiSquare isNil
  ifTrue: [ self computeChiSquare].
^chiSquare
```

computeChanges

```
errorMatrix := DhbLUPDecomposition direct: equations.
^errorMatrix solve: constants
```

computeChiSquare

```
chiSquare := 0.
degreeOfFreedom := self numberOfFreeParameters negated.
dataHolder pointsAndErrorsDo:
  [ :each |
    chiSquare := ( each chi2Contribution: result) + chiSquare.
    degreeOfFreedom := degreeOfFreedom + 1.
  ].
```

computeEquationSystem

```
constants atAllPut: 0.
equations do: [ :each | each atAllPut: 0].
self accumulateEquationSystem.
```

confidenceLevel

```
^( DhbChiSquareDistribution degreeOfFreedom: self
    degreeOfFreedom) confidenceLevel: self chiSquare
```

degreeOfFreedom

```
degreeOfFreedom isNil
  ifTrue: [ self computeChiSquare].
^degreeOfFreedom
```

errorMatrix

```
^DhbSymmetricMatrix rows: errorMatrix inverseMatrixComponents
```

```
evaluateIteration
```

```
  | changes maxChange|
  self computeEquationSystem.
  changes := self computeChanges.
  result changeParametersBy: changes.
  maxChange := 0.
  result parameters with: changes do:
    [ :r :d | maxChange := ( d / r) abs max: maxChange].
  ^maxChange
```

```
finalizeIterations
```

```
  equations := nil.
  constants := nil.
  degreeOfFreedom := nil.
  chiSquare := nil
```

```
fitType
```

```
  ^'Least square fit'
```

```
initialize: aDataHolder data: aParametricFunction
```

```
  dataHolder := aDataHolder.
  result := aParametricFunction.
  ^self
```

```
initializeIterations
```

```
  | n |
  n := self numberOfParameters.
  constants := (DhbVector new: n)
    atAllPut: 0;
    yourself.
  equations := (1 to: n) collect:
    [:k |
      (DhbVector new: n)
        atAllPut: 0;
        yourself]
```

```
numberOfFreeParameters
```

```

^self numberOfParameters

numberOfParameters

^result parameters size

value: aNumber

^result value: aNumber

valueAndError: aNumber

| valueGradient |
valueGradient := result valueAndGradient: aNumber.
^Array with: valueGradient first
with: ( valueGradient last * ( self errorMatrix *
valueGradient last)) sqrt

```

10.9.3 Non-linear fit — Java implementation

In Java, the fit function must be declared as an interface. Here we not only need to compute the functions value, but we need to value of the gradient of the function respective to the vector of parameters; in addition, the parameters of the fit function must be changed. Thus, the interface `OneVariableFunction` is not sufficient as it only covers the retrieving of the function's value. A new interface `ParametrizedOneVariableFunction` was derived from the interface `OneVariableFunction` to handle the management of the function's parameters. It is shown in listing 10.21. This interface has the following instance methods:

`valueAndGradient` returns an array containing the value of the function for the supplied argument in the first position; the rest of the array contains the components of the function's gradient or the same argument; in other word, it returns the array $F(x, \mathbf{p}_0), \frac{\partial F(x, \mathbf{p})}{\partial p_1} \Big|_{\mathbf{p}=\mathbf{p}_0}, \dots, \frac{\partial F(x, \mathbf{p})}{\partial p_m} \Big|_{\mathbf{p}=\mathbf{p}_0}$ for the supplied argument x (let us recall that m is the number of parameters);

`setParameters` defines the value of the function's parameters; in other words, this method is used to set the value of \mathbf{p} in the fit function.

Listing 10.21 Java parametrized function interface

```

package DhbInterfaces;

/**
 * ParametrizedOneVariableFunction is an interface for mathematical
 * functions of one variable depending on several parameters,
 * that is functions of the form  $f(x;p)$ , where  $p$  is a vector.
 *
 * @author Didier H. Besset
 */
public interface ParametrizedOneVariableFunction
    extends OneVariableFunction
{
    /**
     * @return double[] array containing the parameters
     */
    double[] parameters();
    /**
     * @param p double[] assigns the parameters
     */
    void setParameters( double[] p);
    /**
     * Evaluate the function and the gradient of the function with respect
     * to the parameters.
     * @return double[] 0: function's value, 1,2,...,n function's gradient
     * @param x double
     */
    double[] valueAndGradient( double x);
}

```

Note: This implementation of the least square fit is only¹³ for functions with a real argument (a double). However, the mathematical derivation we gave is valid for an argument x of any type. As an exercise, the reader can generalize this implementation for a function of a vector (or array) argument.

Listing 10.22 shows the implementation of a least square fit in Java. The following code example shows how figure 10.8 was generated up to the plotting facilities.

Code example 10.10

```

FisherTippettDistribution genDistr = new FisherTippettDistribution( 0, 1);
Histogram histogram = new Histogram();
histogram.setGrowthAllowed();
for ( int i = 0; i < 1000; i++ )
    histogram.accumulate( genDistr.random());

```

¹³This is not the case in Smalltalk since one does not need to define explicit types.

```

WeightedPoint[] wps = new WeightedPoint[histogram.size()];
for ( int i = 0; i < wps.length; i++)
    wps[i] = histogram.weightedPointAt(i);
FisherTippettDistribution fitDistr =
    new FisherTippettDistribution( histogram);
LeastSquareFit fit = new LeastSquareFit( wps,
    new ScaledProbabilityDensityFunction( fitDistr, histogram));
fit.evaluate();
ScaledProbabilityDensityFunction fittedScaledDistr =
    (ScaledProbabilityDensityFunction) fit.getResult();

```

The first line creates an instance of a Fisher-Tippett distribution with parameters $\alpha = 0$ and $\beta = 1$. The next two lines create an instance of a histogram with automatic adjustment of the limits. Then a loop accumulates random numbers generated according to the distribution into the histogram. After data have been accumulated, an array of weighted points is created and the bins of the histogram are collected into the array of weighted points. Next, the variable `fitDistr` is created on the histogram. This constructor method calculate rough estimates of the distribution's parameters based on the histogram's content, this providing initial values for the fit. The last two statements respectively perform the fit and retrieve the fitted function as an instance of `ScaledProbabilityDensityFunction`. A cast is necessary since the type of the object returned by the method `getResult` is defined only by an interface.

It is not necessary to retrieve the result as a scaled probability density function: after the fit has been performed the parameters of the distribution stored in variable `fitDistr` are the fitted parameters because our implementation keeps the fit parameters in the supplied function object. In other words, the variable `fitDistr` contains the fitted probability density function.

The class `LeastSquareFit` has the following instance variables:

`result` the fit function; this object must implement the `ParametrizedOneVariableFunction` interface described at the beginning of this section;

`points` an array of weighted points containing the data, over which the fit is made;

`systemMatrix` the components of the matrix **M** calculated at each iterations;

`systemConstants` the components of the vector **c** calculated at each iterations;

`systemLUP` an LUP decomposition object (the class `LUPDecomposition`) is described in section 8.3) needed to solve equation 10.49 at each iterations;

`errorMatrix` the error matrix of the fit, that is the inverse of the matrix **M** obtained at the last iteration;

`chiSquare` the sum of equation 10.26;

`degreeOfFreedom` the degree of freedom of the fit.

The class `LeastSquareFit` has two constructor methods, each with two arguments. The arguments of the first method are respectively an array of weighted points (the class `WeightedPoint` described in section 10.3.3) representing the data, over which the fit is performed, and the fit function. The second method is a convenience method used to define the fit of a probability density function (second argument) over a histogram (first argument). In code example 10.10 one could have replaced the generation of the weighted points and the definition of the fit function by the following line:

```
LeastSquareFit fit = new LeastSquareFit( histogram, fitDistr);
```

Many methods may seem redundant to the reader. Their presence will become clear once the reader discovers the implementation of the maximum likelihood fit described in section 10.24. In the maximum likelihood fit the weighted points are directly taken from a histogram. Thus, initializing the system (method `initializeSystem`), reset the system at each iteration (method `resetSystem`), the number of data points (method `getDataSetSize`, access to each data point `weightedPointAt` have been implemented with the up-coming subclassing¹⁴ in mind.

The method `evaluateIteration` first calculates the matrix M and the vector Δp of equation 10.49. Then, this system of linear equations is solved using LUP decomposition described in section 8.3. Then, the obtained vector Δp is added to the vector p to get a new estimation of the parameters. This estimation is fed back into the fit function. The precision returned by the method `evaluateIteration` is the largest of the relative change of the parameter's components.

The method `accumulate` calculates each term of the series of equations 10.58 and 10.59. To optimize speed only the lower half of the symmetric matrix M is accumulated. At end of the summation, the method `symmetrizeMatrix` is used to fill up the upper half of the matrix.

A series of methods allows to retrieve information after the fit has been performed. They are all working with lazy initialization. When the fit is terminated, the method `finalizeIterations` sets the instance variables `systemMatrix` and `systemConstants` to release no longer needed storage. It also sets the values of instance variables `errorMatrix`, `chiSquare` and `degreeOfFreedom` to undefined values to force the computation of these when their respective accessor methods are called. The reader will note that space used by the LUP decomposition is reclaimed as soon as the error matrix is computed since it is no longer needed.

Listing 10.22 Java implementation of a non-linear least square fit

```
package DhbEstimation;

import DhbIterations.IterativeProcess;
import DhbMatrixAlgebra.DhbIllegalDimension;
import DhbMatrixAlgebra.DhbNonSymmetricComponents;
import DhbMatrixAlgebra.LUPDecomposition;
```

¹⁴For `Smalltalkers`: these methods must be declared `protected` since private methods are not inherited in Java.


```

import DhbMatrixAlgebra.SymmetricMatrix;
import DhbStatistics.ChiSquareDistribution;
import DhbInterfaces.ParametrizedOneVariableFunction;
import DhbScientificCurves.Histogram;
import DhbStatistics.ScaledProbabilityDensityFunction;
import DhbStatistics.ProbabilityDensityFunction;
/**
 * Non-linear least square fit
 *
 * @author Didier H. Besset
 */
public class LeastSquareFit extends IterativeProcess
{
    protected ParametrizedOneVariableFunction result;
    private WeightedPoint[] points;
    protected double[][] systemMatrix;
    protected double[] systemConstants;
    private LUPDecomposition systemLUP;
    private SymmetricMatrix errorMatrix;
    private double chiSquare;
    private int degreeOfFreedom;
    /**
     * Default constructor method (internal use only)
     */
    protected LeastSquareFit()
    {
    }
    /**
     * Constructor method
     * @param n int
     */
    public LeastSquareFit(WeightedPoint[] pts,
                          ParametrizedOneVariableFunction f)
    {
        points = pts;
        result = f;
        initializeSystem( result.parameters().length);
    }
    /**
     * Constructor method
     * @param histogram Histogram
     * @param distr ProbabilityDensityFunction
     */
    public LeastSquareFit(Histogram histogram,
                          ProbabilityDensityFunction distr)
    {

```

```

    points = new WeightedPoint[histogram.size()];
    for ( int i = 0; i < points.length; i++)
        points[i] = histogram.weightedPointAt(i);
    result = new ScaledProbabilityDensityFunction( distr, histogram);
    initializeSystem( result.parameters().length);
}
/**
 * @param wp DhbEstimation.WeightedPoint
 */
protected void accumulate( WeightedPoint wp)
{
    double[] fg = result.valueAndGradient( wp.xValue());
    for( int i = 0; i < systemConstants.length; i++ )
    {
        systemConstants[i] += ( wp.yValue() - fg[0])
                               * fg[i+1] * wp.weight();
        for( int j = 0; j <= i; j++ )
            systemMatrix[i][j] += fg[i+1] * fg[j+1] * wp.weight();
    }
}
/**
 * Append the name of the fit to the supplied string buffer
 * @param sb java.lang.StringBuffer
 */
protected void appendFitName( StringBuffer sb)
{
    sb.append("Least square fit with ");
}
/**
 * Append the results of the fit to the supplied string buffer
 * @param sb java.lang.StringBuffer
 */
private void appendFitResults( StringBuffer sb)
{
    java.text.DecimalFormat fmt =
        new java.text.DecimalFormat("###0.00000");
    java.text.DecimalFormat corFmt =
        new java.text.DecimalFormat("0.000");
    sb.append("\n");
    sb.append("\tcompleted in ");
    sb.append( getIterations());
    sb.append(" iterations\n");
    sb.append("\tParams\tErrors\tCorrelation");
    double[][] comp = errorMatrix().toComponents();
    double[] params = result.parameters();
    double[] errors = new double[comp.length];

```

```

char separator;
for ( int i = 0; i < comp.length; i++ )
{
    sb.append("\n\t");
    sb.append( fmt.format(params[i]));
    errors[i] = Math.sqrt(comp[i][i]);
    sb.append("\t+-");
    sb.append( fmt.format(errors[i]));
    separator = '\t';
    for ( int j = 0; j < i; j++ )
    {
        sb.append(separator);
        sb.append(' ');
        sb.append(corFmt.format(comp[i][j] /
                                (errors[i] * errors[j])));
    }
}
appendNormalization( sb);
sb.append("\n\tChi square =");
sb.append( fmt.format(chiSquare()));
sb.append("\tDegree of freedom =");
sb.append( degreeOfFreedom());
sb.append("\tConfidence level =");
sb.append( corFmt.format(confidenceLevel()));
}
/**
 * This method does nothing (compatibility with maximum likelihood fit)
 * @param sb java.lang.StringBuffer
 */
protected void appendNormalization( StringBuffer sb) { }
/**
 * @return double
 */
public double chiSquare()
{
    if ( Double.isNaN( chiSquare) )
        computeChiSquare();
    return chiSquare;
}
/**
 * @return double[] changes on parameters
 */
protected double[] computeChanges()
{
    return systemLUP.solve( systemConstants);
}

```

```

/**
 * Compute the chi^2 of the fit function.
 */
private void computeChiSquare()
{
    chiSquare = 0;
    for( int i = 0; i < getDataSetSize(); i++ )
        chiSquare += weightedPointAt(i).chi2Contribution( result);
}
/**
 * @return DhbMatrixAlgebra.SymmetricMatrix
 */
private void computeErrorMatrix()
{
    double [][] components = systemLUP.inverseMatrixComponents();
    LUPDecomposition.symmetrizeComponents( components);
    try { errorMatrix = SymmetricMatrix.fromComponents( components);
        systemLUP = null;
    }
    catch ( DhbNonSymmetricComponents e) {}
    catch ( DhbIllegalDimension ex) {};
}
private void computeSystem()
{
    resetSystem();
    for ( int i = 0; i < getDataSetSize(); i++ )
        accumulate( weightedPointAt(i));
    symmetrizeMatrix();
}
/**
 * @return double    confidence level of the fit.
 */
public double confidenceLevel()
{
    return (new ChiSquareDistribution(
        degreeOfFreedom())).confidenceLevel( chiSquare());
}
/**
 * @return long    the degree of freedom of the fit.
 */
public int degreeOfFreedom()
{
    if ( degreeOfFreedom < 0 )
        degreeOfFreedom = getDataSetSize() -
            result.parameters().length;
    return degreeOfFreedom;
}

```

```

    }
    /**
     * @return DhbMatrixAlgebra.SymmetricMatrix the error matrix of the fit.
     */
    public SymmetricMatrix errorMatrix()
    {
        if ( errorMatrix == null )
            computeErrorMatrix();
        return errorMatrix;
    }
    /**
     * @return double
     */
    public double evaluateIteration()
    {
        double[] parameters = result.parameters();
        computeSystem();
        try { systemLUP = new LUPDecomposition( systemMatrix);}
        catch( DhbIllegalDimension e) { };
        double[] changes = computeChanges();
        double eps = 0;
        for ( int i = 0; i < parameters.length; i++ )
        {
            parameters[i] += changes[i];
            eps = Math.max( eps, Math.abs(
                relativePrecision( changes[i], parameters[i])));
        }
        result.setParameters( parameters);
        return eps;
    }
    public void finalizeIterations()
    {
        systemMatrix = null;
        systemConstants = null;
        errorMatrix = null;
        chiSquare = Double.NaN;
        degreeOfFreedom = -1;
    }
    /**
     * @return int number of data points.
     */
    protected int getDataSetSize()
    {
        return points.length;
    }
    /**

```

```

    * @return ParametrizedOneVariableFunction the fitted function
    */
    public ParametrizedOneVariableFunction getResult()
    {
        return result;
    }
    /**
    * @param n int
    */
    protected void initializeSystem( int n)
    {
        systemConstants = new double[ n];
        systemMatrix = new double[n][n];
    }
    protected void resetSystem( )
    {
        for( int i = 0; i < systemConstants.length; i++ )
        {
            systemConstants[i] = 0;
            for( int j = 0; j <= i; j++ )
                systemMatrix[i][j] = 0;
        }
    }
    private void symmetrizeMatrix( )
    {
        for( int i = 0; i < systemConstants.length; i++ )
        {
            for( int j = 0; j < i; j++ )
                systemMatrix[j][i] = systemMatrix[i][j];
        }
    }
    /**
    * @return java.lang.String
    */
    public String toString()
    {
        StringBuffer sb = new StringBuffer();
        appendFitName(sb);
        sb.append( getResult());
        if ( hasConverged() )
            appendFitResults(sb);
        return sb.toString();
    }
    /**
    * @return DhbEstimation.WeightedPoint n-th weighted data point
    * @param n int

```

```

*/
protected WeightedPoint weightedPointAt( int n)
{
    return points[n];
}
}

```

10.10 Maximum likelihood fit of a probability density function

In section 9.3 histograms have been discussed as a way to represent a probability density function directly from experimental data. In this section we shall show that the maximum likelihood estimation can easily be applied to the data gathered in a histogram in order to determine the parameters of a hypothesized probability density function.

In general the maximum likelihood fit of a probability density function to a histogram is much faster than the corresponding least square fit because the number of free parameters is lower, as we shall see in this section. In addition, the maximum likelihood estimation is unbiased and is therefore a better estimation than the least square fit estimation, especially when the histogram is sparsely populated. Thus, a maximum likelihood fit is the preferred way of finding the parameters of a probability density function from experimental data collected in a histogram.

Let m be the number of bins in the histogram and let n_i be the content of the i^{th} bin. Let $P_i(\mathbf{p})$ the probability of observing a value in the i^{th} bin. The likelihood function $L(\mathbf{p})$ is the probability of observing the particular histogram. Since the hypothesis of a probability density function does not constrain the total number of values collected into the histogram, the total number of collected values can be considered as constant. As a consequence, a maximum likelihood fit has one parameter less than a least square fit using the same function. Since the total number is unconstrained, the probability of observing the particular histogram is given by a multinomial probability. Thus, the likelihood function can be written as:

$$L(\mathbf{p}) = N! \prod_{i=1}^m \frac{P_i(\mathbf{p})^{n_i}}{n_i!}, \quad (10.53)$$

where $N = \sum_{i=1}^m n_i$ is the total number of values collected into the histogram. As we have seen in section 10.5.1, finding the maximum of $L(\mathbf{p})$ is equivalent of finding the maximum of the function $I(\mathbf{p})$. Since N is a constant, we use a renormalized function:

$$I(\mathbf{p}) = \ln \frac{L(\mathbf{p})}{N!} = \sum_{i=1}^m n_i \ln P_i(\mathbf{p}). \quad (10.54)$$

Finding the maximum of the function $I(\mathbf{p})$ is equivalent to solving the following system of non-linear equations:

$$\frac{\partial I(\mathbf{p})}{\partial \mathbf{p}} = \sum_{i=1}^m \frac{n_i}{P_i(\mathbf{p})} \cdot \frac{\partial P_i(\mathbf{p})}{\partial \mathbf{p}} = 0. \quad (10.55)$$

This system can be solved with a search by successive approximations, where a system of linear equations must be solved at each step. The technique used is similar to the one described in section 10.9. In this case, however, it is more convenient to expand the inverse of the probability density function around a previous approximation as follows:

$$\frac{1}{P_i(\mathbf{p})} = \frac{1}{P_i(\mathbf{p}_0)} - \frac{1}{P_i(\mathbf{p}_0)^2} \cdot \left. \frac{\partial P_i(\mathbf{p})}{\partial \mathbf{p}} \right|_{\mathbf{p}=\mathbf{p}_0} \cdot \Delta \mathbf{p}. \quad (10.56)$$

This expansion can only be defined over a range where the probability density function is not equal to zero. Therefore, this expansion of the maximum likelihood estimation cannot be used on a histogram where bins with non-zero count are located on a range where the probability density function is equal to zero¹⁵. Contrary to a least square fit, bins with zero count do not participate to the estimation.

Now equation 10.55 becomes a system of linear equations of the type:

$$\mathbf{M} \cdot \Delta \mathbf{p} = \mathbf{c}, \quad (10.57)$$

where the components of the matrix \mathbf{M} are now defined by:

$$M_{jk} = \sum_{i=1}^m \frac{n_i}{P_i(\mathbf{p}_0)^2} \cdot \left. \frac{\partial P_i(\mathbf{p}_0)}{\partial p_j} \right|_{\mathbf{p}=\mathbf{p}_0} \cdot \left. \frac{\partial P_i(\mathbf{p}_0)}{\partial p_k} \right|_{\mathbf{p}=\mathbf{p}_0}, \quad (10.58)$$

and those of the vector \mathbf{c} by:

$$c_j = \sum_{i=1}^m \frac{n_i}{P_i(\mathbf{p}_0)} \cdot \left. \frac{\partial P_i(\mathbf{p}_0)}{\partial p_j} \right|_{\mathbf{p}=\mathbf{p}_0}. \quad (10.59)$$

As discussed at the beginning of this section, the maximum likelihood estimation for a histogram cannot determine the total count in the histogram. The estimated total count, \bar{N} , is estimated with the following hypothesis:

$$n_i = \bar{N} P(\bar{\mathbf{p}}), \quad (10.60)$$

where $\bar{\mathbf{p}}$ is the maximum likelihood estimation of the distribution parameters. The estimation is performed using \bar{N} as the only variable. The maximum likelihood estimation cannot be solved analytically, however, the least square estimation can.

As we have seen in section 10.4 the variance of the bin count is the estimated bin content. Thus, the function to minimize becomes:

$$S(\bar{N}) = \sum_{i=1}^m \frac{[n_i - \bar{N} P_i(\bar{\mathbf{p}})]^2}{\bar{N} P_i(\bar{\mathbf{p}})} \quad (10.61)$$

The value of \bar{N} minimizing the expression of equation 10.61 is:

$$\bar{N} = \sqrt{\frac{\sum_{i=1}^m n_i^2 / P_i(\bar{\mathbf{p}})}{\sum_{i=1}^m P_i(\bar{\mathbf{p}})}}. \quad (10.62)$$

¹⁵Equation 10.54 shows that the bins over which the probability density function is zero give no information.

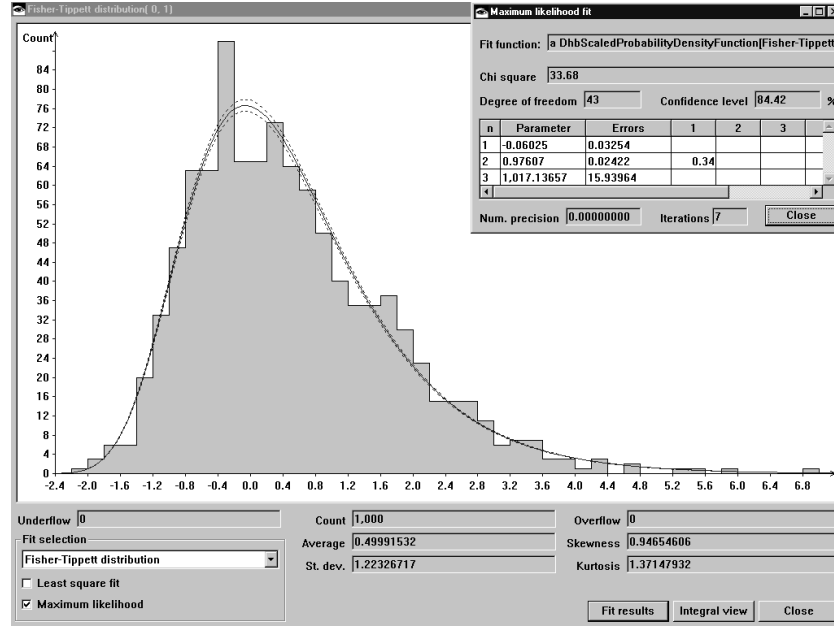


Figure 10.9: Example of a maximum likelihood fit

and the estimated error on \bar{N} is given by

$$\sigma_{\bar{N}} = \sqrt{\frac{\sum_{i=1}^m n_i^2 / P_i(\bar{\mathbf{p}})}{2\bar{N}}}. \quad (10.63)$$

After computing \bar{N} using equation 10.62, the goodness of the maximum likelihood fit can be estimated by calculating the χ^2 confidence level of $S(\bar{N})$ given by equation 10.61.

Figure 10.9 shows an example of a maximum likelihood fit performed on the same histogram as in figure 10.8. The inset window in the upper right corner shows the fit results in the same order as figure 10.8. The correlation coefficients, however, are not shown for the normalization since it is not determined as part of the fit. The solid curve laid onto the histogram is the prediction of the fitted function; the two dotted lines indicate the error on the prediction. The reader can see that the fit is as good as the least square fit. Of course, the χ^2 test is significantly higher with a correspondingly lower confidence level. This mostly comes from the fact that a maximum likelihood fit does not use the bins with zero count. In fact, the reader can see that the count in the histogram (normalization) estimated by the maximum likelihood fit is higher than in the case of the least square fit.

10.10.1 Maximum likelihood fit — General implementation

A maximum likelihood fit of a probability density function on a histogram is very similar to a least square fit of a histogram with a scaled probability distribution. There are two major differences: first the number of parameters is lower; second the computation of the matrix and vectors is not the same. Otherwise, most of the structure of a least square fit can be reused.

Instead of creating special methods to compute the gradient of the fitted function using a new set of parameters, our implementation uses the same gradient calculation than the one used by the least square fit. This is possible if the component of the gradient relative to the normalization is placed at the end. Since the computation of this component does not require additional calculation, the additional time required by the re-using of the gradient's computation is negligible. Since the fit function is a scaled probability distribution the current normalization is kept in an instance variable and the normalization of the fitted function is set to 1 for the duration of the iterations. When the algorithm is completed, the estimated normalization is put back into the fit function.

The computation of the normalization (equation 10.62) and that of its error (equation 10.63) is performed in the method `finalizeIterations`.

Figure 10.1 with the box `MaximumLikelihoodHistogramFit` grayed.

10.10.2 Maximum likelihood fit — Smalltalk implementation

Listing 10.23 shows the complete implementation in Smalltalk. The following code example shows how figure 10.9 was generated up to the plotting facilities.

Code example 10.11

```
| genDistr hist fit |
hist := DhbHistogram new.
hist freeExtent: true.
genDistr := DhbFisherTippettDistribution shape: 0 scale: 1.
1000 timesRepeat: [ hist accumulate: genDistr random ].
fit := DhbMaximumLikelihoodHistogramFit histogram: hist
      distributionClass: DhbFisherTippettDistribution.
fit evaluate.
```

As the reader can see the only difference with code example 10.9 is the name of the class in the statement where the instance of the fit is created.

The class `DhbMaximumLikelihoodHistogramFit` is a subclass of the class `DhbLeastSquareFit`. It has the following additional instance variables:

`count` the estimated normalization, that is \bar{N} ;

`countVariance` the estimated variance of \bar{N} .

The variance is kept instead of the error because the most frequent use of this quantity is in computing the estimated error on the predicted value. In the method `valueAndError`: this computation requires the combination of the error of the fit — that is, equation 10.33 — with the error on the normalization. An accessor method is provided for the variable `count`. The method `normalizationError` calculates the error on the normalization.

The method `accumulate`: uses the vector operations to calculate the terms of the sums in equations 10.58 and 10.59. Because of the lower number of parameters, the routine `computeChanges`: places in the vector Δ_p an additional zero element corresponding to the normalization in the case of the least square fit.

The method `finalizeIterations` calculates the estimated value of the normalization (equation 10.61) and its variance (square of equation 10.62). After this, it sets the obtained normalization into the scaled probability distribution.

Listing 10.23 Smalltalk implementation of a maximum likelihood fit

```

Class                DhbMaximumLikekihoodHistogramFit
Subclass of          DhbLeastSquareFit
Instance variable names:  count countVariance

```

Instance methods

`accumulate: aWeightedPoint`

```

| f g temp inverseProbability |
f := result valueAndGradient: aWeightedPoint xValue.
g := f last copyFrom: 1 to: ( f last size - 1).
f := f first.
f = 0 ifTrue: [ ^nil].
inverseProbability := 1 / f.
temp := aWeightedPoint yValue * inverseProbability.
constants accumulate: g * temp.
temp := temp * inverseProbability.
1 to: g size do:
    [ :k |
        ( equations at: k) accumulate: g * ( ( g at: k) * temp).
    ].

```

`computeChanges`

```

^super computeChanges copyWith: 0

```

`computeNormalization`

```

| numerator denominator temp |
numerator := 0.
denominator := 0.
dataHolder pointsAndErrorsDo:
    [ :each |
        temp := result value: each xValue.
        temp = 0
    ]

```

10.10. MAXIMUM LIKELIHOOD FIT OF A PROBABILITY DENSITY FUNCTION 489

```
ifFalse:
    [numerator := numerator + (each yValue squared /
                               temp).
      denominator := denominator + temp]].
count := ( numerator / denominator) sqrt.
countVariance := numerator / ( 4 * count).
```

finalizeIterations

```
self computeNormalization.
result setCount: count.
super finalizeIterations
```

fitType

```
^^Maximum likelihood fit'
```

initializeIterations

```
result setCount: 1.
count := dataHolder totalCount.
super initializeIterations
```

normalization

```
^count
```

normalizationError

```
^countVariance sqrt
```

numberOfFreeParameters

```
^super numberOfParameters
```

numberOfParameters

```
^super numberOfParameters - 1
```

valueAndError: aNumber

```

| valueGradient gradient gVar |
valueGradient := result valueAndGradient: aNumber.
gradient := valueGradient last copyFrom: 1 to: valueGradient last
              size - 1.
gVar := gradient * (self errorMatrix * gradient) / count.
^Array with: valueGradient first
  with: ((valueGradient first / count) squared * countVariance
        + gVar) sqrt

```

10.10.3 Maximum likelihood fit — Java implementation

Listing 10.24 shows the complete implementation in Java. The following code example shows how figure 10.9 was generated up to the plotting facilities.

Code example 10.12

```

FisherTippettDistribution genDistr = new FisherTippettDistribution( 0, 1);
Histogram histogram = new Histogram();
histogram.setGrowthAllowed();
for ( int i = 0; i < 1000; i++ )
    histogram.accumulate( genDistr.random());
FisherTippettDistribution fitDistr = new FisherTippettDistribution( histogram);
MaximumLikelihoodHistogramFit fit = new MaximumLikelihoodHistogramFit( histogram, fitDistr);
fit.evaluate();

```

In this example the constructor method using a histogram and a probability density function is used. Otherwise, it is very similar to code example 10.10.

The class `MaximumLikelihoodHistogramFit` is a subclass of class `LeastSquareFit`. It has the following additional instance variables:

`histogram` the histogram over which the fit is performed; because of strong typing, it was not possible to re-use the variable points;

`count` the estimated normalization, that is \bar{N} ;

`countError` the estimated error on \bar{N} .

The class `MaximumLikelihoodHistogramFit` has only one constructor method corresponding to the convenience constructor method of the class `LeastSquareFit`. Assignment of the instance variable is different since data is accessed directly in the histogram.

The method `accumulate` calculates the terms of the sums in equations 10.58 and 10.59. The methods `getDataSetSize`, `computeChanges` and `weightedPointAt` overloads the corresponding methods of class `LeastSquareFit` to reflect the fact that the data points are obtained directly from the histogram.

The method `finalizeComputation` first calls the method `computeNormalization`, which calculates the estimation of the normalization and its error using equations 10.62 and

10.10. MAXIMUM LIKELIHOOD FIT OF A PROBABILITY DENSITY FUNCTION 491

10.63. Then the obtained normalization is set into the fitted function. Unused space is release in the method of the superclass.

Listing 10.24 Java implementation of a maximum likelihood fit

```
package DhbEstimation;

import DhbInterfaces.ParametrizedOneVariableFunction;
import DhbScientificCurves.Histogram;
import DhbStatistics.ScaledProbabilityDensityFunction;
import DhbStatistics.ProbabilityDensityFunction;
/**
 * Maximum likelihood fit
 *
 * @author Didier H. Besset
 */
public class MaximumLikelihoodHistogramFit extends LeastSquareFit
{
    /**
     * Histogram containing the data
     */
    private Histogram histogram;
    /**
     * Estimated total count in histogram
     */
    private double count;
    /**
     * Estimated error on total count in histogram
     */
    private double countError;
    /**
     * Constructor method.
     * @param pts Histogram
     * @param f DhbInterfaces.ParametrizedOneVariableFunction
     */
    public MaximumLikelihoodHistogramFit(Histogram hist,
                                         ProbabilityDensityFunction f)
    {
        histogram = hist;
        result = new ScaledProbabilityDensityFunction ( f, hist);
        initializeSystem( f.parameters().length);
    }
    /**
     * @param wp DhbEstimation.WeightedPoint
     */
    protected void accumulate( WeightedPoint wp)
```

```

{
    double[] fg = result.valueAndGradient( wp.xValue());
    if ( fg[0] == 0 )
        return;
    double invProb = 1 / fg[0];
    double temp = wp.yValue() * invProb;
    for( int i = 0; i < systemConstants.length; i++ )
    {
        systemConstants[i] += fg[i+1] * temp;
        for( int j = 0; j <= i; j++ )
            systemMatrix[i][j] += fg[i+1] * fg[j+1] * temp * invProb;
    }
}
/**
 * Append the name of the fit to the supplied string buffer
 * @param sb java.lang.StringBuffer
 */
protected void appendFitName( StringBuffer sb)
{
    sb.append("Maximum likelihood fit with ");
}
/**
 * Append the normalization and its error to the fit results
 * @param sb java.lang.StringBuffer
 */
protected void appendNormalization( StringBuffer sb)
{
    java.text.DecimalFormat fmt = new java.text.DecimalFormat("###0.0");
    sb.append("\n\t");
    sb.append( fmt.format( count));
    sb.append("\t+-");
    sb.append( fmt.format( countError));
}
/**
 * Computes the changes in the parameters:
 * since the normalization is not fitted, the change to the
 * normalization (last parameter) is set to zero.
 */
protected double[] computeChanges()
{
    double[] changes = super.computeChanges();
    double[] answer = new double[changes.length+1];
    for ( int i = 0; i < changes.length; i++ )
        answer[i] = changes[i];
    answer[changes.length] = 0;
    return answer;
}

```

```

    }
    /**
     * Computes the estimated normalization and variance on it.
     */
    private void computeNormalization()
    {
        double numerator = 0;
        double denominator = 0;
        double temp;
        WeightedPoint wp;
        for ( int i = 0; i < getDataSetSize(); i++ )
        {
            wp = weightedPointAt(i);
            temp = result.value( wp.xValue());
            if ( temp != 0 )
            {
                numerator += wp.yValue() * wp.yValue() / temp;
                denominator += temp;
            }
        }
        count = Math.sqrt( numerator / denominator);
        countError = Math.sqrt( 0.25 * numerator / count);
    }
    public void finalizeIterations()
    {
        computeNormalization();
        getDistribution().setCount( count);
        super.finalizeIterations();
    }
    /**
     * @return int    number of data points.
     */
    protected int getDataSetSize()
    {
        return histogram.size();
    }
    /**
     * @return ScaledProbabilityDensityFunction    the fitted function
     */
    public ScaledProbabilityDensityFunction getDistribution()
    {
        return (ScaledProbabilityDensityFunction) getResult();
    }
    public void initializeIterations()
    {
        getDistribution().setCount( 1);
    }

```



```
        count = histogram.totalCount();
    }
    /**
     * @return DhbEstimation.WeightedPoint n-th weighted data point
     * @param n int
     */
    protected WeightedPoint weightedPointAt( int n)
    {
        return histogram.weightedPointAt(n);
    }
}
```

Chapter 11

Optimization

Cours vite au but, mais gare á la chute.¹
Alexandre Soljenitsyne

An optimization problem is a numerical problem where the solution is characterized by the largest or smallest value of a numerical function depending on several parameters. Such function is often called the goal function. Many kinds of problems can be expressed into optimization, that is, finding the maximum or the minimum of a goal function. This technique has been applied to a wide variety of fields going from operation research to game playing or artificial intelligence. In chapter 10 for example, the solution of maximum likelihood or least square fits was obtained by finding the maximum, respectively the minimum of a function.

In fact generations of high energy physicists have used the general purpose minimization program MINUIT² written by Fred James³ of CERN to perform least square fits and maximum likelihood fits. To achieve generality, MINUIT uses several strategies to reach a minimum. In this chapter we shall discuss a few techniques and conclude with a program quite similar in spirit to MINUIT. Our version, however, will not have all the features offered by MINUIT.

If the goal function can be expressed with an analytical form, the problem of optimization can be reduced into calculating the derivatives of the goal function respective to all parameters, a tedious but manageable job. In most cases, however, the goal function cannot always be expressed analytically.

The classes described in this chapter are different in Smalltalk and in Java. Therefore we present two class diagrams: figure 11.1 shows the Smalltalk class diagram and figure 11.2 shows the Java class diagram. The main reason for the difference is the strong typing imposed in Java preventing the reuse of instance variables.

¹Run fast to the goal, but beware of the fall.

²F.James, M. Roos, MINUIT — a system for function minimization and analysis of the parameter errors and corrections, *Comput. Phys. Commun.*, 10 (1975) 343-367.

³I take this opportunity to thank Fred for the many useful discussions we have had on the subject of minimization.

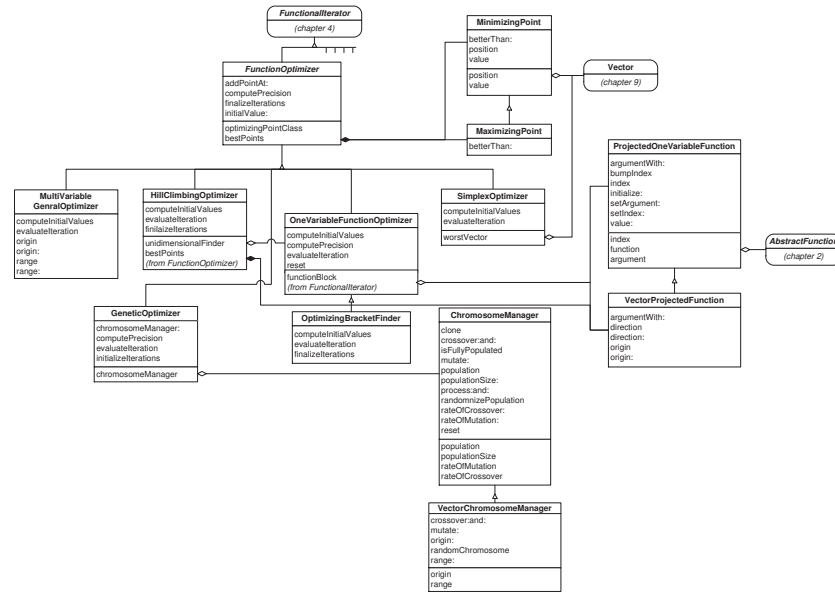


Figure 11.1: Smalltak classes used in optimization

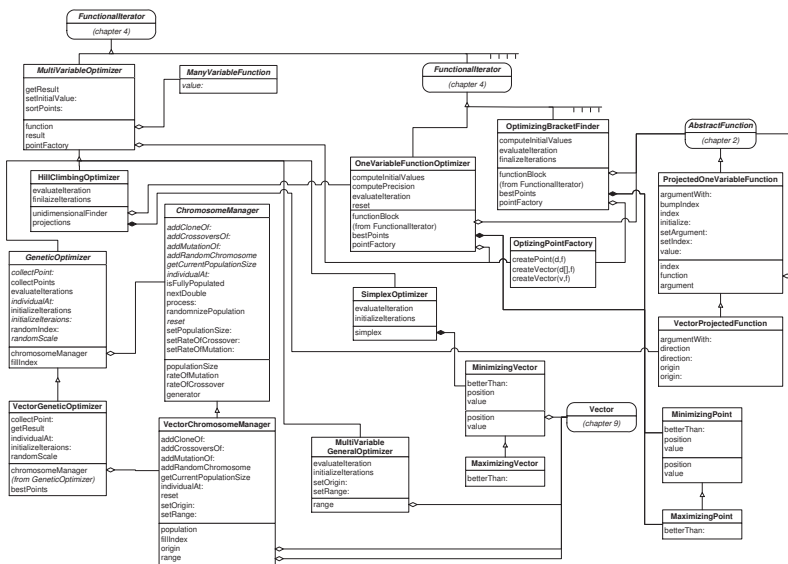


Figure 11.2: Java classes used in optimization




Figure 11.3: Local and absolute optima

11.1 Introduction

Let us state the problem in general terms. Let $f(\mathbf{x})$ be a function of a vector \mathbf{x} of dimension n . The n -dimensional space is called the search space of the problem. Depending on the problem the space can be continuous or not. In this section we shall assume that the space is continuous.

If the function is derivable, the gradient of the function respective to the vector \mathbf{x} must vanish at the optimum. Finding the optimum of the function can be replaced by the problem of finding the vector \mathbf{x} such that:

$$\frac{df(\mathbf{x})}{d\mathbf{x}} = 0. \quad (11.1)$$

Unfortunately, the above equation is not a necessary condition for an optimum. It can be either a maximum, a minimum or a saddle point, that is a point where the function has a minimum in one projection and a maximum in another projection. Furthermore, the function may have several optima. Figure 11.3 shows an example of a function having two minima. Some problems require to find the absolute optimum of the function. Thus, one must verify that the solution of 11.1 corresponds indeed to an optimum with the expected properties. The reader can already see at this point that searching for an optimum in the general case is a very difficult task.

All optimization algorithms can be classified in two broad categories:

- Greedy algorithms: these algorithms are characterized by a local search in the most promising direction. They are usually efficient and quite good at finding

local optima. Among greedy algorithms, one must distinguish those requiring the evaluation of the function's derivatives.

- Random based algorithms: these algorithms are using a random approach. They are not efficient; however, they are good at finding absolute optima. Simulated annealing [Press et al.] and genetic algorithms[Berry & Linoff] belong to this class.

The table 11.1 lists the properties of the algorithms presented in this chapter.

Table 11.1: Optimizing algorithms presented in this book

Name	Category	Derivatives
Extended Newton	greedy	yes
Powell's hill climbing	greedy	no
Simplex	greedy	no
Genetic algorithm	random based	no

11.2 Extended Newton algorithms

Extended Newton algorithms are using a generalized version of Newton's zero finding algorithm. These algorithms assume that the function is continuous and has only one optimum in the region where the search is initiated.

Let us expand the function $f(\mathbf{x})$ around a point $\mathbf{x}^{(0)}$ near the solution. We have in components:

$$f(\mathbf{x}) = f[\mathbf{x}^{(0)}] + \sum_j \left. \frac{\partial f(\mathbf{x})}{\partial x_j} \right|_{\mathbf{x}=\mathbf{x}^{(0)}} [x_j - x_j^{(0)}]. \quad (11.2)$$

Using the expansion above into equation 11.1 yields:

$$\sum_j \left. \frac{\partial^2 f(\mathbf{x})}{\partial x_i \partial x_j} \right|_{\mathbf{x}=\mathbf{x}^{(0)}} [x_j - x_j^{(0)}] + \left. \frac{\partial f(\mathbf{x})}{\partial x_i} \right|_{\mathbf{x}=\mathbf{x}^{(0)}} = 0. \quad (11.3)$$

This equation can be written as a system of linear equations of the form

$$\mathbf{M}\Delta = \mathbf{c}, \quad (11.4)$$

where $\Delta_j = x_j - x_j^{(0)}$. The components of the matrix \mathbf{M} — called the Hessian matrix — are given by:

$$m_{ij} = \left. \frac{\partial^2 f(\mathbf{x})}{\partial x_i \partial x_j} \right|_{\mathbf{x}=\mathbf{x}^{(0)}}, \quad (11.5)$$

and the components of the vector \mathbf{c} are given by:

$$c_i = - \left. \frac{\partial f(\mathbf{x})}{\partial x_i} \right|_{\mathbf{x}=\mathbf{x}^{(0)}}. \quad (11.6)$$

Like in section 10.9 one can iterate this process by replacing $\mathbf{x}^{(0)}$ with $\mathbf{x}^{(0)} + \Delta$. This process is actually equivalent to the Newton zero finding method (c.f. section 5.3). The final solution is a minimum if the matrix \mathbf{M} is positive definite; else it is a maximum.

This technique is used by MINUIT in the vicinity of the goal function's optimum. It is the region where the algorithm described above works well. Far from the optimum, the risk of reaching a point where the matrix \mathbf{M} cannot be inverted is quite high in general. In addition, the extended Newton algorithm requires that the second order derivatives of the function can be computed analytically; at least the first order derivatives must be provided, otherwise the cost of computation at each step becomes prohibitive. A concrete implementation of the technique is not given here. The reader can find in this book all the necessary tools to make such an implementation. It is left as an exercise for the reader. In the rest of this chapter, we shall present other methods which work without an analytical knowledge of the function.

11.3 Hill climbing algorithms

Hill climbing is a generic term covering many algorithms trying to reach an optimum by determining the optimum along successive directions. The general algorithm is outlined below.

1. select an initial point \mathbf{x}_0 and a direction \mathbf{v} ;
2. find \mathbf{x}_1 , the optimum of the function along the selected direction;
3. if convergence is attained, terminate the algorithm;
4. set $\mathbf{x}_0 = \mathbf{x}_1$, select a different direction and go back to step 2.

The simplest of these algorithms simply follows each axis in turn until a convergence is reached. More elaborate algorithms exist [Press et al.]. One of them is described in section 11.6.

Hill climbing algorithms can be applied to any continuous function, especially when the function's derivatives are not easily calculated. The core of the hill climbing algorithm is finding the optimum along one direction. Let \mathbf{v} be the direction, then finding the optimum of the vector function $f(\mathbf{x})$ along the direction \mathbf{v} starting from point \mathbf{x}_0 is equivalent to finding the optimum of the one-variable function $g(\lambda) = f(\mathbf{x}_0 + \lambda\mathbf{v})$.

Therefore, in order to implement a hill climbing algorithm, we first need to implement an algorithm able to find the optimum of a one-variable function. This is the topic of the sections 11.4 and 11.5. Before this, we need to discuss the implementation details providing a common framework to all classes discussed in the rest of this chapter.

11.3.1 Optimizing — General implementation

At this point the reader may be a little puzzled by the use of optimum instead of speaking of minimum or maximum. We shall now disclose a general implementation which works both for finding a minimum or a maximum. This should not come to a surprise since, in mathematics, a minimum or a maximum are both very similar — position where the

derivative of a function vanishes — and can be easily turned into each other — e.g. by negating the function.

To implement a general purpose optimizing framework, we introduce two new classes: `MinimizingPoint` and `MaximizingPoint`, a subclass of `MinimizingPoint`. These two classes are used as Strategy by the optimizing algorithms. The class `MinimizingPoint` has two instance variables

value the value of the goal function, that is $g(\lambda)$ or $f(\mathbf{x})$;

position the position at which the function has been evaluated, that is λ or \mathbf{x} .

The class `MinimizingPoint` contains most of the methods. The only method overloaded by the class `MaximizingPoint` is the method `betterThan`, which tells whether an optimizing point is better than another. The method `betterThan` can be used in all parts of the optimizing algorithms to find out which point is the optimum so far. In algorithms working in multiple dimensions, the method `betterThan` is also used to sort the points from the best to the worst. In Java, the architecture is a little more complex because of typing requirements, but the basic design concept is the same.

A convenience instance creation method allows to create instances for a given function with a given argument. The instance is then initialized with the function's value evaluated at the argument. Thus, all optimizing algorithms described here do not call the goal function explicitly.

Otherwise the implementation of the one dimensional optimum search uses the general framework of the iterative process. More specifically it uses the class `FunctionalIterator` described in section 4.2.

A final remark concerns the method `initializeIteration`. The golden search algorithm assume that the 3 points λ_0 , λ_1 and λ_2 have been determined. What if they have not been? In this case, the method `initializeIteration` uses the optimum bracket finder described in section 11.5

11.3.2 Common optimizing classes — Smalltalk implementation

Figure 11.1 with the boxes `FunctionOptimizer`, `MinimizingPoint`, `MaximizingPoint`, `ProjectedOneVariableFunction` and `VectorProjectedFunction` grayed.

In Smalltalk we have two classes of optimizing points: `DhbMinimizingPoint` and its subclass `DhbMaximizingPoint`. These classes are shown in listing 11.1. The class `DhbFunctionOptimizer` is in charge of handling the management of the optimizing points. This class is shown in listing 11.2.

The class `DhbMinimizingPoint` has the following instance variables:

position contains the position at which the function is evaluated; this instance variable is a number if the function to optimize is a one variable function and an array or a vector if the function to evaluate is a function of many variables;

value contains the value of the function evaluated at the point's position;

Accessor methods corresponding to these variables are supplied. As we noted in section 11.3.1, the only method redefined by the subclass `DhbMaximizingPoint` is the method `betterThan`: used to decide whether a point is better than another.

Optimizing points are created with the convenience method `vector:function:` which evaluates the function supplied as second argument at the position supplied as the first argument.

Listing 11.1 Smalltalk classes common to all optimizing classes

Class	DhbMinimizingPoint
Subclass of	Object

Instance variable names: value position

Class methods

`new: aVector value: aNumber`

`^self new vector: aVector; value: aNumber; yourself`

`vector: aVector function: aFunction`

`^self new: aVector value: (aFunction value: aVector)`

Instance methods

`betterThan: anOptimizingPoint`

`^value < anOptimizingPoint value`

`position`

`^position`

`printOn: aStream`

`position printOn: aStream.`

`aStream`

`nextPut: $;;`

`space.`

`value printOn: aStream`

`value`

`^value`

`value: aNumber`

`value := aNumber.`


```
vector: aVector
    position := aVector
```

```
Class      DhbMaximizingPoint
Subclass of DhbMinimizingPoint
```

```
    Instance methods
```

```
    betterThan: anOptimizingPoint
    ^value > anOptimizingPoint value
```

The class `DhbFunctionOptimizer` is in charge of handling the optimizing points. it has the following instance variables:

`optimizingPointClass` is the class of the optimizing points used as Strategy by the optimizer; it is used to create instances of points at a given position for a given function;

`bestPoints` contains a sorted collection of optimizing points; the best point is the first and the worst point is the last; all optimizers rely on the fact that sorting is done by this sorted collection.

The method `addPointAt:` creates an optimizing point at the position supplied as argument and adds this point to the collection of best points. Since that collection is sorted, one is always certain to find the best result in the first position. This fact is used by the method `finalizeIterations`, which retrieves the result from the collection of best points.

Instances of the function optimizer are created with the two convenience methods `minimizingFuntion:` and `maximizingFuntion:` helping to define the type of optimum. An additional convenience method, `forOptimizer:` allows to create a new optimizer with the same strategy — that is, the same class of optimizing points — and the same function as the optimizer supplied as argument. This method is used to create optimizers used in intermediate steps.

Because finding an optimum cannot be determined numerically with high precision [Press et al.] the class `DhbFunctionOptimizer` redefines the method `defaultPrecision` to be 100 times the default numerical precision.

Listing 11.2 Smalltalk abstract class for all optimizing classes

```
Class      DhbFunctionOptimizer
Subclass of DhbFunctionalIterator
Instance variable names:  optimizingPointClass bestPoints
```

```
    Class methods
```

```
    defaultPrecision
```

```

^super defaultPrecision * 100

forOptimizer: aFunctionOptimizer
    ^self new initializeForOptimizer: aFunctionOptimizer

maximizingFunction: aFunction
    ^super new initializeAsMaximizer; setFunction: aFunction

minimizingFunction: aFunction
    ^super new initializeAsMinimizer; setFunction: aFunction

    Instance methods
addPointAt: aNumber
    bestPoints add: (optimizingPointClass vector: aNumber
                    function: functionBlock)

bestPoints
    ^bestPoints

finalizeIterations
    result := bestPoints first position.

functionBlock
    ^functionBlock

initialize
    bestPoints := SortedCollection sortBlock:
        [ :a :b | a betterThan: b].
    ^super initialize

initializeAsMaximizer
    optimizingPointClass := DhbMaximizingPoint.
    ^self initialize

```

```
initializeAsMinimizer
```

```
    optimizingPointClass := DhbMinimizingPoint.  
    ^self
```

```
initializeForOptimizer: aFunctionOptimizer
```

```
    optimizingPointClass := aFunctionOptimizer pointClass.  
    functionBlock := aFunctionOptimizer functionBlock.  
    ^self initialize
```

```
initialValue: aVector
```

```
    result := aVector copy.
```

```
pointClass
```

```
    ^optimizingPointClass
```

```
printOn: aStream
```

```
    super printOn: aStream.  
    bestPoints do: [ :each | aStream cr. each printOn: aStream].
```

In order to find an optimum along a given direction, one needs to construct an object able to transform a vector function into a one variable function. The class `DhbProjectedOneVariableFunction` and its subclass `DhbVectorProjectedFunction` provide this functionality. They are shown in listing 11.3. The class `DhbProjectedOneVariableFunction` has the following instance variables:

function the goal function $f(\mathbf{x})$;

argument the vector argument of the goal function, that is the vector \mathbf{x} ;

index the index of the axis along which the function is projected.

The instance variables `argument` and `index` can be read and modified using direct accessor methods. The goal function is set only at creation time: the instance creation method `function:` take the goal function as argument. A convenience method `bumpIndex` allows to alter the index in circular fashion⁴.

The class `DhbVectorProjectedFunction` has no additional variables. Instead it is reusing the instance variable `index` as the direction along which the function is evaluated.

⁴We do not give the implementation of the simplest of the hill climbing algorithms using alternatively each axes of the reference system. This implementation, which uses the method `bumpIndex`, is left as an exercise for the reader.

For clarity, the accessor methods have been renamed `direction`, `direction:`, `origin` and `origin:`.

For both classes, the method `argumentAt:` returns the argument vector for the goal function, that is the vector \mathbf{x} . The method `value:` returns the value of the function $g(\lambda)$ for the supplied argument λ .

Listing 11.3 Smalltalk projected function classes

```

Class                DhbProjectedOneVariableFunction
Subclass of          Object
Instance variable names:  index function argument

```

Class methods

`function:` aVectorFunction

`^super new initialize: aVectorFunction`

Instance methods

`argumentWith:` aNumber

`^argument at: index put: aNumber; yourself`

`bumpIndex`

```

index isNil
  ifTrue: [ index := 1]
  ifFalse:[ index := index + 1.
    index > argument size
      ifTrue: [ index := 1].
    ].

```

`index`

```

index isNil
  ifTrue: [ index := 1].
^index

```

`initialize:` aFunction

```

function := aFunction.
^self

```

`setArgument:` anArrayOfVector

argument := anArrayOfVector copy.

setIndex: anInteger

index := anInteger.

value: aNumber

^function value: (self argumentWith: aNumber)

Class DhbVectorProjectedFunction

Subclass of DhbProjectedOneVariableFunction

Instance methods

argumentWith: aNumber

^aNumber * self direction + self origin

direction

^index

direction: aVector

index := aVector.

origin

^argument

origin: aVector

argument := aVector.

printOn: aStream

self origin printOn: aStream.
aStream nextPutAll: ' ('.
self direction printOn: aStream.
aStream nextPut: \$).

position a double containing the position at which the function is evaluated;

Both instance variables have a corresponding getter accessor method, but no setter method. Otherwise, the functional relation between position and value could not be ensured. As we noted in section 11.3.1, the only method redefined by the subclass `MaximizingPoint` is the method `betterThan` used to decide whether a point is better than another. Instances of optimizing points are created with a single constructor method taking as arguments the position at which the function must be evaluated and the function itself. The supplied function must implement the `OneVariableFunction` interface described in section 2.1.2.

```
package DhbOptimizing;
```

```
import DhhInterfaces.OneVariableFunction;
/**
 * Point & function holder used in optimizing one-variable functions.
 *
 * @author Didier H. Besset
 */
public abstract class OptimizingPoint
{
    /**
     * Value of the function to optimize.
     */
    private double value;
    /**
     * Position at which the value was evaluated.
     */
    private double position;
    /**
     * Constructor method
     * @param x double    position at which the goal function is evaluated.
     * @param f OneVariableFunction    function to optimize.
     */
    public OptimizingPoint( double x, OneVariableFunction f)
    {
```

```

        position = x;
        value = f.value(x);
    }
    /**
     * @return boolean true if the receiver is "better" than
     *                  the supplied point
     * @param point OptimizingPoint
     */
    public abstract boolean betterThan( OptimizingPoint entity);
    /**
     * @return double the receiver's position
     */
    public double getPosition()
    {
        return position;
    }
    /**
     * @return double the value of the function at the receiver's
     *                  position
     */
    public double getValue()
    {
        return value;
    }
    /**
     * (used by method toString).
     * @return java.lang.String
     */
    protected abstract String printedKey();
    /**
     * @return java.lang.String
     */
    public String toString()
    {
        StringBuffer sb = new StringBuffer();
        sb.append( value);
        sb.append(printedKey());
        sb.append( position);
        return sb.toString();
    }
}

package DhbOptimizing;

import DhbInterfaces.OneVariableFunction;
/**

```

```

* Point & function holder used in minimizing one-variable functions.
*
* @author Didier H. Besset
*/
public class MinimizingPoint extends OptimizingPoint {
/**
* Constructor method.
* @param x double
* @param f DhbInterfaces.OneVariableFunction
*/
public MinimizingPoint(double x, OneVariableFunction f)
{
    super(x, f);
}
/**
* @return boolean true if the receiver is "better" than
*                  the supplied point
* @param point OptimizingPoint
*/
public boolean betterThan(OptimizingPoint point)
{
    return getValue() < point.getValue();
}
/**
* (used by method toString).
* @return java.lang.String
*/
protected final String printedKey()
{
    return " min@";
}
}

package DhbOptimizing;

import DhbInterfaces.OneVariableFunction;
/**
* Point & function holder used in maximizing one-variable functions.
*
* @author Didier H. Besset
*/
public class MaximizingPoint extends OptimizingPoint {
/**
* Constructor method.
* @param x double

```



```

    * @param f DhbInterfaces.OneVariableFunction
    */
    public MaximizingPoint(double x, OneVariableFunction f)
    {
        super(x, f);
    }
    /**
     * @return boolean true if the receiver is "better" than
     *                  the supplied point
     * @param point OptimizingPoint
     */
    public boolean betterThan(OptimizingPoint point)
    {
        return getValue() > point.getValue();
    }
    /**
     * (used by method toString).
     * @return java.lang.String
     */
    protected final String printedKey()
    {
        return " max@";
    }
}

```

The abstract class `OptimizingVector` has two subclasses `MinimizingVector` and `Maximizingvector`. The only differences with the corresponding optimizing point classes are that the instance variable position is an array of double and that the function supplied as the second argument of the constructor method must implement the `ManyVariableFunction` interface. These classes and the interface `ManyVariableFunction` are shown in listing 11.5. The need for the additional hierarchy comes primarily from the fact that we cannot handle double and array of double as a common object since these entities are primitive types and not objects.

Listing 11.5 Java optimizing vector classes

```

package DhbInterfaces;

/**
 * ManyVariableFunction is an interface for mathematical functions
 * of many variables, that is functions of the form:
 *      f(X) where X is a vector.
 *
 * @author Didier H. Besset
 */

```

```

public interface ManyVariableFunction
{

    /**
     * Returns the value of the function for the specified vector.
     */
    public double value ( double[] x);
}

package DhbOptimizing;

import DhbInterfaces.ManyVariableFunction;
/**
 * Vector & function holder used in optimizing many-variable functions.
 *
 * @author Didier H. Besset
 */
public abstract class OptimizingVector
{
    /**
     * Value of the function to optimize.
     */
    private double value;
    /**
     * Position at which the value was evaluated.
     */
    private double[] position;
    /**
     * Value of the function to optimize.
     */
    protected ManyVariableFunction f;
    /**
     * Constructor method.
     * @param v double[]
     * @param f DhbInterfaces.OneVariableFunction
     */
    public OptimizingVector(double[] v, ManyVariableFunction func)
    {
        position = v;
        f = func;
        value = f.value( position);
    }
    /**
     * @return boolean true if the receiver is "better" than
     * the supplied point
     * @param point OptimizingVector

```

```

*/
public abstract boolean betterThan( OptimizingVector entity);
/**
 * (used by the Simplex algorithm).
 * @param v double[]
 */
public void contractFrom( double[] v)
{
    for ( int i = 0; i < position.length; i++ )
    {
        position[i] += v[i];
        position[i] *= 0.5;
    }
    value = f.value( position);
}
/**
 * @return double    the receiver's position
 */
public double[] getPosition()
{
    return position;
}
/**
 * @return double    the value of the function
 *                  at the receiver's position
 */
public double getValue()
{
    return value;
}
/**
 * (used by method toString)..
 * @return java.lang.String
 */
protected abstract String printedKey();
/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    sb.append( value);
    sb.append(printedKey());
    for ( int i = 0; i < position.length; i++ )
    {
        sb.append( ' ');
    }
}

```

```

        sb.append( position[i]);
    }
    return sb.toString();
}
}

```

```
package DhbOptimizing;
```

```

import DhbInterfaces.ManyVariableFunction;
/**
 * Vector & function holder used in minimizing many-variable functions.
 *
 * @author Didier H. Besset
 */
public class MinimizingVector extends OptimizingVector {
/**
 * Constructor method.
 * @param v double[]
 * @param f DhbInterfaces.ManyVariableFunction
 */
    public MinimizingVector(double[] v, DhbInterfaces.ManyVariableFunction f) {
        super(v, f);
    }
/**
 * @return boolean true if the receiver is "better" than
 *                  the supplied point
 * @param point OptimizingVector
 */
    public boolean betterThan(OptimizingVector point)
    {
        return getValue() < point.getValue();
    }
/**
 * (used by method toString).
 * @return java.lang.String
 */
    protected final String printedKey()
    {
        return " min@";
    }
}

```

```
package DhbOptimizing;
```

```
import DhbInterfaces.ManyVariableFunction;
```

```

/**
 * Vector & function holder used in maximizing many-variable functions.
 *
 * @author Didier H. Besset
 */
public class MaximizingVector extends OptimizingVector {
/**
 * Constructor method.
 * @param v double[]
 * @param f DhbInterfaces.ManyVariableFunction
 */
    public MaximizingVector(double[] v,
                           DhbInterfaces.ManyVariableFunction f)
    {
        super(v, f);
    }
/**
 * @return boolean true if the receiver is "better" than
 *                  the supplied point
 * @param point OptimizingVector
 */
    public boolean betterThan(OptimizingVector point)
    {
        return getValue() > point.getValue();
    }
/**
 * (used by method toString).
 * @return java.lang.String
 */
    protected final String printedKey()
    {
        return " max@";
    }
}

```

The abstract class `OptimizingPointFactory` and its two concrete subclasses, `MinimizingPointFactory` and `MaximizingPointFactory`, are in charge of creating the optimizing points. These factory classes⁵ play the role of the optimizing Strategy described in section 11.3.1. The concrete classes implement the two methods `createPoint` and `createVector`. The method `createPoint` creates an optimizing point of the desired type for a supplied position (a double) and function (a `OneVariableFunction`). The method `createVector` creates an

⁵For people reading code in both languages: Smalltalk does not need any factory: objects of type `Class` are instance factories built into the language. Reflection in Java could have been used here to build an architecture similar to that of Smalltalk. The use of reflection, however, implies the use of casting, which, in my humble opinion, spoils the purpose of strong typing.

Listing 11.6 Java optimizing point factory classes

[illegible]

```

{
    return createVector( v.toComponents(), f);
}
}

```

```
package DhbOptimizing;
```

```

/**
 * Factory of point/vector & function holders for minimizing functions.
 *
 * @author Didier H. Besset
 */
public class MinimizingPointFactory extends OptimizingPointFactory
{
    /**
     * Constructor method.
     */
    public MinimizingPointFactory() {
        super();
    }
    /**
     * @return OptimizingPoint an minimizing point strategy.
     */
    public OptimizingPoint createPoint(double x, DhbInterfaces.OneVariableFunction f)
    {
        return new MinimizingPoint( x, f);
    }
    /**
     * @return OptimizingVector an minimizing vector strategy.
     */
    public OptimizingVector createVector(double[] v, DhbInterfaces.ManyVariableFunction f)
    {
        return new MinimizingVector( v, f);
    }
}

```

```
package DhbOptimizing;
```

```

/**
 * Factory of point/vector & function holders for maximizing functions.
 *
 * @author Didier H. Besset
 */
public class MaximizingPointFactory extends OptimizingPointFactory {
    /**

```

```

* Constructor method.
*/
public MaximizingPointFactory() {
    super();
}
/**
 * @return OptimizingPoint  an maximizing point strategy.
 */
public OptimizingPoint createPoint(double x,
                                   DhbInterfaces.OneVariableFunction f)
{
    return new MaximizingPoint( x, f);
}
/**
 * @return OptimizingVector  an maximizing vector strategy.
 */
public OptimizingVector createVector(double[] v,
                                     DhbInterfaces.ManyVariableFunction f)
{
    return new MaximizingVector( v, f);
}
}

```

In the Smalltalk implementation we have been able to bring all optimizing classes under a single abstract class. The Java architecture is similar. In the case of Java, however, the abstract class only handles optimization in n -dimensional space.

Listing 11.7 shows the code of class `MultiVariableOptimizer`. This class has the following instance variables:

`f` the goal function; this object must implement the `ManyVariableFunction` interface shown in listing 11.5.

`pointFactory` the factory used to create optimizing points; this must be one concrete instance of the classes shown in listing 11.6.

`result` an array of double containing the initial value where to start the algorithm; at the end of the algorithm, this variable contains the position of the optimum if convergence was attained;

The default constructor method provided by the abstract class take three arguments: the goal function, the optimizing point factory and the initial value. These three arguments correspond to the three instance variables of the class.

The method `setInitialValue` allows to change the initial value in case another search is made with the same instance. The accessor method `getResult` allows one to retrieve the result.

The method `sortPoints` sorts a supplied array of optimizing points according to their `Strategy`. Since sorting is made in situ, a bubble sort algorithm is used.

Listing 11.7 Java abstract class for all optimizing classes

```

package DhbOptimizing;

import DhbInterfaces.ManyVariableFunction;
import DhbIterations.IterativeProcess;
/**
 * Abstract optimizer of many-variable functions.
 *
 * @author Didier H. Besset
 */
public abstract class MultiVariableOptimizer extends IterativeProcess
{
    /**
     * Value of the function to optimize.
     */
    protected ManyVariableFunction f;
    /**
     * Best value found so far: must be set to determine the dimension
     * of the argument of the function.
     */
    protected double[] result;
    /**
     * Optimizing strategy (minimum or maximum).
     */
    protected OptimizingPointFactory pointFactory;
    /**
     * Constructor method.
     */
    public MultiVariableOptimizer(ManyVariableFunction func,
        OptimizingPointFactory pointCreator, double[] initialValue)
    {
        f = func;
        pointFactory = pointCreator;
        setInitialValue( initialValue);
    }
    /**
     * @return double[] result of the receiver
     */
    public double[] getResult()
    {
        return result;
    }
    /**
     * @param v double[] educated guess for the optimum's location

```

```

*/
public void setInitialValue( double[] v)
{
    result = v;
}
/**
 * Use bubble sort to sort the best points
 */
protected void sortPoints( OptimizingVector[] bestPoints)
{
    OptimizingVector temp;
    int n = bestPoints.length;
    int bound = n - 1;
    int i, m;
    while ( bound >= 0 )
    {
        m = -1;
        for ( i = 0; i < bound; i++ )
        {
            if ( bestPoints[i+1].betterThan( bestPoints[i]) )
            {
                temp = bestPoints[i];
                bestPoints[i] = bestPoints[i+1];
                bestPoints[i+1] = temp;
                m = i;
            }
        }
        bound = m;
    }
}
}

```

In order to find an optimum along a given direction, one needs to construct an object able to transform a vector function into a one variable function. The class `VectorProjectedFunction` providing this functionality has the following instance variables:

`f` the goal function $f(\mathbf{x})$;

`origin` the vector argument of the goal function, that is the vector \mathbf{x}_0 ;

`direction` the direction along which the function is projected, that is the vector \mathbf{v} .

The constructor method takes three argument corresponding to the three instance variables. In addition, the instance variables argument and direction can be read and modified using direct accessor methods. For convenience the set methods exist in two versions of the argument: an array of double or a `DhbVector`. This is also true for the constructor method.

The method `argumentAt`: returns the argument vector for the goal function, that is the vector \mathbf{x} . The method `value`: returns the value of the function $g(\lambda)$ for the supplied argument λ .

Listing 11.8 Java projected function class

```
package DhbOptimizing;

import DhbInterfaces.OneVariableFunction;
import DhbInterfaces.ManyVariableFunction;
import DhbMatrixAlgebra.DhbVector;
import DhbMatrixAlgebra.DhbIllegalDimension;
/**
 * Projection of a many-variable function
 *          onto a one-dimensional direction.
 *
 * @author Didier H. Besset
 */
public class VectorProjectedFunction implements OneVariableFunction
{
    /**
     * Value of the function to optimize.
     */
    private ManyVariableFunction f;
    /**
     * Origin for function evaluation.
     */
    private DhbVector origin;
    /**
     * Direction along which the function is evaluated.
     */
    private DhbVector direction;
    /**
     * Constructor method.
     * @param func ManyVariableFunction  function to project
     * @param x double[]  origin of projected function
     * @param d double[]  direction of projection
     * @exception NegativeArraySizeException if dimension of x or d is 0.
     */
    public VectorProjectedFunction( ManyVariableFunction func,
                                   double[] x, double[] d)
        throws NegativeArraySizeException
    {
        f = func;
        setOrigin( x);
        setDirection( d);
    }
}
```

```

    }
    /**
     * Constructor method.
     * @param func ManyVariableFunction  function to project
     * @param x DhbVector  origin of projected function
     * @param d DhbVector  direction of projection
     */
    public VectorProjectedFunction( ManyVariableFunction func,
                                   DhbVector x, DhbVector d)
    {
        f = func;
        setOrigin( x);
        setDirection( d);
    }
    /**
     * @param x double[]  origin of projected function
     * @exception DhbIllegalDimension
     *             if dimension of x is not that of the origin.
     */
    public DhbVector argumentAt(double x) throws DhbIllegalDimension
    {
        DhbVector v = direction.product( x);
        v.accumulate( origin);
        return v;
    }
    /**
     * @return DhbMatrixAlgebra.DhbVector  direction of the receiver
     */
    public DhbVector getDirection()
    {
        return direction;
    }
    /**
     * @return DhbMatrixAlgebra.DhbVector  origin of the receiver
     */
    public DhbVector getOrigin()
    {
        return origin;
    }
    /**
     * @param v DhbMatrixAlgebra.DhbVector
     * @exception NegativeArraySizeException if dimension of v is 0.
     */
    public void setDirection( double[] v) throws NegativeArraySizeException
    {
        direction = new DhbVector( v);
    }

```

```

    }
    /**
     * @param v DhbMatrixAlgebra.DhbVector
     */
    public void setDirection( DhbVector v)
    {
        direction = v;
    }
    /**
     * @param v DhbMatrixAlgebra.DhbVector
     * @exception NegativeArraySizeException if dimension of v is 0.
     */
    public void setOrigin( double[] v) throws NegativeArraySizeException
    {
        origin = new DhbVector( v);
    }
    /**
     * @param v DhbMatrixAlgebra.DhbVector
     */
    public void setOrigin( DhbVector v)
    {
        origin = v;
    }
    /**
     * Returns a String that represents the value of this object.
     * @return a string representation of the receiver
     */
    public String toString()
    {
        StringBuffer sb = new StringBuffer();
        sb.append( origin);
        sb.append(" -> ");
        sb.append( direction);
        return sb.toString();
    }
    /**
     * @return double    value of the function
     * @param x double    distance from the origin in unit of direction.
     */
    public double value(double x)
    {
        try{ return f.value( argumentAt(x).toComponents());
        } catch ( DhbIllegalDimension e) { return Double.NaN;}
    }
}

```

11.4 Optimizing in one dimension

To find the optimum of a one-variable function, $g(\lambda)$, whose derivative is unknown, the most robust algorithm is an algorithm similar to the bisection algorithm described in section 5.2.

Let us assume that we have found three points λ_0, λ_1 and λ_2 such that $\lambda_0 < \lambda_1 < \lambda_2$ and such that $g(\lambda_1)$ is better than both $g(\lambda_0)$ and $g(\lambda_2)$. If the function g is continuous over the interval $[\lambda_0, \lambda_2]$, then we are certain that an optimum of the function is located in the interval $[\lambda_0, \lambda_2]$. As for the bisection algorithm, we shall try to find a new triplet of values with similar properties while reducing the size of the interval. A point is picked in the largest of the two intervals $[\lambda_0, \lambda_1]$ or $[\lambda_1, \lambda_2]$ and is used to reduce the initial interval.

If $\lambda_1 - \lambda_0 \leq \lambda_2 - \lambda_1$ we compute $\lambda_4 = \lambda_1 + \omega(\lambda_2 - \lambda_1)$ where ω is the golden section⁶ from Pythagorean lore. Choosing ω instead of $1/2$ ensures that successive intervals have the same relative scale. A complete derivation of this argument can be found in [Press et al.]. If λ_4 yields a better function value than λ_1 , the new triplet of point becomes λ_1, λ_4 and λ_2 ; otherwise, the triplet becomes λ_0, λ_1 and λ_4 .

If we have $\lambda_1 - \lambda_0 > \lambda_2 - \lambda_1$ we compute $\lambda_4 = \lambda_1 + \omega(\lambda_0 - \lambda_1)$. Then the new triplets can be either λ_0, λ_4 and λ_1 , or λ_4, λ_1 and λ_2 .

The reader can verify that the interval decreases steadily although not as fast as in the case of bisection where the interval is halved at each iteration. Since the algorithm is using the golden section it is called golden section search.

By construction the golden section search algorithm makes sure that the optimum is always located between the points λ_0 and λ_2 . Thus, at each iteration, the quantity $\lambda_2 - \lambda_0$ give an estimate of the error on the position of the optimum.

11.4.1 Optimizing in one dimension — Smalltalk implementation

Listing 11.9 shows the class `DhbOneVariableFunctionOptimizer` implementing the search for an optimum of a one-variable function using the golden section search. The following code example shows how to use this class to find the maximum of the gamma distribution discussed in section 9.7.

Figure 11.1 with the box `OneVariableFunctionOptimizer` grayed.

Code example 11.1

```
| distr finder maximum |
distr := DhbGammaDistribution shape: 2 scale: 5.
finder := DhbOneVariableFunctionOptimizer maximizingFunction: distr.
maximum := finder evaluate.
```

The first line after the declarations creates a new instance of a gamma distribution with parameters $\alpha = 2$ and $\beta = 5$. The next line creates an instance of the optimum finder. The selector used to create the instance selects a search for a maximum. The last line is the familiar statement to evaluate the iterations — that is, performing the search for the maximum — and to retrieve the result. Since no initial value was supplied the search begins at a random location.

⁶ $\omega = \frac{3-\sqrt{5}}{2} \approx 0.38197$

The class `DhbOneVariableFunctionOptimizer` is a subclass of the class `FunctionOptimizer`. It does not need any additional instance variables. The golden section is kept as a class variable and is calculated at the first time it is needed.

At each iteration the method `nextXValue` is used to compute the next position at which the function is evaluated. This corresponding new optimizing point is added to the collection of best points. Then, the method `indexOfOuterPoint` is used to determine which point must be discarded: it is always the second point on either side of the best point. The precision of the result is estimated from the bracketing interval in the method `computePrecision`, using relative precision (of course!).

The method `addPoint:` of the superclass can be used to supply an initial bracketing interval. The method `computeInitialValues` first checks whether a valid bracketing interval has been supplied into the collection of best points. If this is not the case, a search for a bracketing interval is conducted using the class `DhbOptimizingBracketFinder` described in section 11.5.1. The instance of the bracket finder is created with the method `forOptimizer:` so that its strategy and goal function are taken over from the golden section optimum finder.

Listing 11.9 Smalltalk golden section optimum finder

```

Class                DhbOneVariableFunctionOptimizer
Subclass of          DhbFunctionOptimizer
Class variable names: GoldenSection

    Class methods
defaultPrecision
    ^DhbFloatingPointMachine new defaultNumericalPrecision * 10

goldenSection
    GoldenSection isNil ifTrue: [GoldenSection := (3 - 5 sqrt) / 2].
    ^GoldenSection

    Instance methods
computeInitialValues
    [ bestPoints size > 3 ] whileTrue: [ bestPoints removeLast].
    bestPoints size = 3
        ifTrue: [ self hasBracketingPoints
            ifFalse:[ bestPoints removeLast].
        ].
    bestPoints size < 3
        ifTrue: [ ( DhbOptimizingBracketFinder forOptimizer: self)
            evaluate].

```

computePrecision

```
^self precisionOf: ( ( bestPoints at: 2) position - ( bestPoints
                    at: 3) position) abs
relativeTo: ( bestPoints at: 1) position abs
```

evaluateIteration

```
self addPointAt: self nextXValue.
bestPoints removeAtIndex: self indexOfOuterPoint.
^self computePrecision
```

hasBracketingPoints

```
| x1 |
x1 := ( bestPoints at: 1) position.
^( ( bestPoints at: 2) position - x1) * (( bestPoints at: 3)
    position - x1) < 0
```

indexOfOuterPoint

```
| inferior superior x |
inferior := false.
superior := false.
x := bestPoints first position.
2 to: 4 do:
    [ :n |
        ( bestPoints at: n) position < x
        ifTrue: [ inferior
                  ifTrue: [ ^n].
                  inferior := true.
                ]
        ifFalse:[ superior
                  ifTrue: [ ^n].
                  superior := true.
                ].
    ].
```

nextXValue

```
| d3 d2 x1 |
x1 := ( bestPoints at: 1) position.
d2 := ( bestPoints at: 2) position - x1.
d3 := ( bestPoints at: 3) position - x1.
^( d3 abs > d2 abs ifTrue: [ d3]
   ifFalse:[ d2]) * self class goldenSection + x1
```


reset

[bestPoints isEmpty] whileFalse: [bestPoints removeLast].

Figure 11.2 with the box OneVariableFunctionOptimizer grayed.

11.4.2 Optimizing in one dimension — Java implementation

Listing 11.10 shows the class OneVariableFunctionOptimizer implementing the search for an optimum of a one-variable function using the golden section search in Java. The following code example shows how to use this class to find the maximum of the gamma distribution discussed in section 9.7.

Code example 11.2

```
CauchyDistribution distr = new CauchyDistribution( 10, 5);
MaximizingPointFactory strategy = new MaximizingPointFactory();
OneVariableFunctionOptimizer finder =
    new OneVariableFunctionOptimizer( distr, strategy);
finder.setDesiredPrecision( 1.0e-5);
finder.evaluate();
double result = finder.getResult();
```

The first line creates an instance of a gamma distribution with parameters $\alpha = 2$ and $\beta = 5$. The next line creates an instance of a maximizing point factory. This factory is used as the strategy for the optimum finder created on the next line, the first argument of the constructor method being the function to maximize. A medium precision is given explicitly since the default provided by the class IterativeProcess is likely to be meaningless (c.f. [Press et al.]). Since no initial value is supplied, the search begins at a random point. After a call to the method evaluate — common to all iterative processes — the result is retrieved on the last line.

The class OneVariableFunctionOptimizer is a subclass of the class FunctionalIterator described in section 4.2.2. It has the following instance variables

bestPoints an array of 3 optimizing points corresponding to the values λ_0 , λ_1 and λ_2 in this order; thus, the best is always the middle one;

pointFactory the factory used to create optimizing points (c.f. section 11.3.3).

The private static variable **goldenSection** holds the value of the golden section.

The method **evaluateIteration** determines which side of the bracketing interval the bisection step will be performed. The bisection step of the algorithm is performed within the method **reducePoints**. The arguments of the method is either 0 or 2 depending on the selected side. The precision returned by the method **evaluateIteration** is of course calculated using relative precision.

The method **initializeIterations** uses the optimum bracket finder described in section 11.5.2 to obtain the first bracketing interval. The search for the interval is conducted from the supplied initial value. The instance of the bracket finder is created by supplying the same goal function and same strategy as the initial golden section finder.

Listing 11.10 Java implementation of the golden section optimum search

```

package DhbOptimizing;

import DhbInterfaces.OneVariableFunction;
import DhbIterations.FunctionalIterator;
/**
 * Optimizer of one-variable functions
 * (uses golden search algorithm).
 *
 * @author Didier H. Besset
 */
public class OneVariableFunctionOptimizer extends FunctionalIterator
{
    private static double goldenSection = ( 3 - Math.sqrt(5)) / 2;
    /**
     * Best points found so far.
     */
    private OptimizingPoint[] bestPoints = null;
    /**
     * Optimizing strategy (minimum or maximum).
     */
    private OptimizingPointFactory pointFactory;
    /**
     * Constructor method
     * @param func OneVariableFunction
     * @param pointCreator OptimizingPointFactory a factory to create
     *                    strategy points
     */
    public OneVariableFunctionOptimizer(OneVariableFunction func,
                                       OptimizingPointFactory pointCreator)
    {
        super(func);
        pointFactory = pointCreator;
    }
    /**
     * @return double the relative precision on the result
     */
    private double computePrecision()
    {
        return relativePrecision( Math.abs( bestPoints[2].getPosition()
                                           - bestPoints[1].getPosition()),
                                Math.abs( bestPoints[0].getPosition()));
    }
    /**
     * @return double current precision of result

```

```

*/
public double evaluateIteration()
{
    if ( bestPoints[2].getPosition() - bestPoints[1].getPosition()
        > bestPoints[1].getPosition() - bestPoints[0].getPosition() )
        reducePoints(2);
    else
        reducePoints(0);
    result = bestPoints[1].getPosition();
    return computePrecision();
}
public void initializeIterations()
{
    OptimizingBracketFinder bracketFinder =
        new OptimizingBracketFinder( f, pointFactory);
    bracketFinder.setInitialValue( result);
    bracketFinder.evaluate();
    bestPoints = bracketFinder.getBestPoints();
}
/**
 * Apply bisection on points 1 and n
 * @param n int    index of worst point of bisected interval
 */
private void reducePoints( int n)
{
    double x = bestPoints[1].getPosition();
    x += goldenSection * ( bestPoints[n].getPosition() - x);
    OptimizingPoint newPoint = pointFactory.createPoint( x, f);
    if ( newPoint.betterThan( bestPoints[1]) )
    {
        bestPoints[2-n] = bestPoints[1];
        bestPoints[1] = newPoint;
    }
    else
        bestPoints[n] = newPoint;
}
/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    sb.append( getIterations());
    sb.append( " iterations, precision = ");
    sb.append( getPrecision());
    for ( int i = 0 ; i < bestPoints.length; i++ )

```

```

{
    sb.append( '\n');
    sb.append( bestPoints[i]);
}
return sb.toString();
}
}

```

11.5 Bracketing the optimum in one dimension

As we have seen in section 11.4 the golden section algorithm requires the knowledge of a bracketing interval. This section describes a very simple algorithm to obtain a bracketing interval with certainty if the function is continuous and does indeed have an optimum of the sought type.

The algorithm goes as follows. Take two points λ_0 and λ_1 . If they do not exist, pick up some random values (random generators are described in section 9.4). Let us assume that $g(\lambda_1)$ is better than $g(\lambda_0)$.

1. Let $\lambda_2 = 3\lambda_1 - 2\lambda_0$, that is, λ_2 is twice as far from λ_1 than λ_0 and is located on the other side, toward the optimum.
2. If $g(\lambda_1)$ is better than $g(\lambda_2)$ we have a bracketing interval; the algorithm is stopped.
3. Otherwise, set $\lambda_0 = \lambda_1$ and $\lambda_1 = \lambda_2$ and go back to step 1.

The reader can see that the interval $[\lambda_0, \lambda_1]$ is increasing at each step. Thus, if the function has no optimum of the sought type, the algorithm will cause a floating overflow exception quite rapidly.

The implementation in each language have too little in common. The common section is therefore omitted.

11.5.1 Bracketing the optimum — Smalltalk implementation

Listing 11.11 shows the Smalltalk code of the class implementing the search algorithm for an optimizing bracket. The class `DhbOptimizingBracketFinder` is a subclass of class `DhbOneVariableFunctionOptimizer` from section 11.9. This was a convenient, but not necessary, choice to be able to reuse most of the management and accessor methods. The methods pertaining to the algorithm are of course quite different.

Example of use of the optimizing bracket finder can be found in method `computeInitialValues` of class `DhbOneVariableFunctionOptimizer` (c.f. listing 11.9).

The method `setInitialPoints:` allows to use the collection of best points of another optimizer inside the class. This breach to the rule of hiding the implementation can be tolerated here because the class `DhbOptimizingBracketFinder` is used exclusively with the class `DhbOneVariableFunctionOptimizer`. It allows the two class to use the

Figure 11.1 with the box `OptimizingBracketFinder` grayed.

same sorted collection of optimizing points. If no initial point has been supplied, it is obtained from a random generator.

The precision calculated in the method `evaluateIteration` is a large number, which becomes negative as soon as the condition to terminate the algorithm is met. Having a negative precision causes an iterative process as defined in chapter 4 to stop.

Listing 11.11 Smalltalk optimum bracket finder

Class **DhbOptimizingBracketFinder**

Subclass of `DhbOneVariableFunctionOptimizer`

Class methods

`initialPoints: aSortedCollection function: aFunction`

`^super new setInitialPoints: aSortedCollection; setFunction:
aFunction`

Instance methods

`computeInitialValues`

`[bestPoints size < 2] whileTrue: [self addPointAt: Number random]`

`evaluateIteration`

`| x1 x2 |`

`x1 := (bestPoints at: 1) position.`

`x2 := (bestPoints at: 2) position.`

`self addPointAt: (x1 * 3 - (x2 * 2)).`

`precision := (x2 - x1) * ((bestPoints at: 3) position - x1).`

`self hasConverged`

`ifFalse:[bestPoints removeLast].`

`^precision`

`finalizeIterations`

`result := bestPoints.`

`initializeForOptimizer: aFunctionOptimizer`

`super initializeForOptimizer: aFunctionOptimizer.`

`bestPoints := aFunctionOptimizer bestPoints.`

`^self`

`setInitialPoints: aSortedCollection`

`bestPoints := aSortedCollection.`

11.5.2 Bracketing the optimum — Java implementation

Listing 11.12 shows the Smalltalk code of the class implementing the search algorithm for an optimizing bracket. The class `OptimizingBracketFinder` is a subclass of the class `FunctionalIterator` defined in section 4.2.2. It has the same instance variables as the class `OneVariableFunctionOptimizer`.

The method `initializeIterations` gets an initial value using a random generator if no initial value has been supplied. The other points of the initial interval are obtained from a random generator. At the end of the method, the array of best points contains 3 optimizing points sorted in ascending order of their position.

The method `evaluateIteration` first tests whether the optimum seems to be located toward negative or positive values. After the test the interval is expanded toward to corresponding direction by one of the methods `moveTowardNegative` or `moveTowardPositive`.

The method `evaluateIteration` calculate a pseudo precision: it is a large number, which becomes negative as soon as the condition to terminate the algorithm is met. Having a negative precision causes an iterative process as defined in chapter 4 to stop.

The object using the bracket finder can retrieve the array of best points using the method `getBestPoints`. This method constitutes a breach to the rule of hiding the implementation. In this case, however, this is an acceptable breach: the class `OneVariableFunctionOptimizer` can retrieve the initial bracketing interval and use it immediately without having to evaluate again the function at these positions. After all, the bracketing interval is the result of the optimizing bracket search.

Figure 11.2 with the box `OptimizingBracketFinder` grayed.

Listing 11.12 Java optimum bracket finder

```
package DhbOptimizing;

import java.util.Random;
import DhbInterfaces.OneVariableFunction;
import DhbIterations.FunctionalIterator;
/**
 * Finds a bracket for the optimum of a one-variable function.
 *
 * @author Didier H. Besset
 */
public class OptimizingBracketFinder extends FunctionalIterator
{
    private OptimizingPoint[] bestPoints = null;
    private OptimizingPointFactory pointFactory;
    /**
     * Constructor method
     * @param func OneVariableFunction
     * @param pointCreator OptimizingPointFactory a factory to create
     *                    strategy points
     */
    public OptimizingBracketFinder(OneVariableFunction func,
```

```

        OptimizingPointFactory pointCreator)
    {
        super(func);
        pointFactory = pointCreator;
    }
    /**
     * @return double    1 as long as no bracket has been found
     */
    private double computePrecision()
    {
        return bestPoints[1].betterThan( bestPoints[0]) &&
            bestPoints[1].betterThan( bestPoints[2])
            ? 0 : 1;
    }
    /**
     * @return double    pseudo-precision of the current search
     */
    public double evaluateIteration()
    {
        if ( bestPoints[0].betterThan( bestPoints[1]) )
            moveTowardNegative();
        else if ( bestPoints[2].betterThan( bestPoints[1]) )
            moveTowardPositive();
        return computePrecision();
    }
    /**
     * @return OptimizingPoint[]    a triplet bracketing the optimum
     */
    public OptimizingPoint[] getBestPoints()
    {
        return bestPoints;
    }
    /**
     * Use random locations (drunkard's walk algorithm).
     */
    public void initializeIterations()
    {
        Random generator = new Random();
        bestPoints = new OptimizingPoint[3];
        if ( Double.isNaN( result ) )
            result = generator.nextDouble();
        bestPoints[0] = pointFactory.createPoint( result, f);
        bestPoints[1] = pointFactory.createPoint( generator.nextDouble()
            + bestPoints[0].getPosition(), f);
        bestPoints[2] = pointFactory.createPoint( generator.nextDouble()
            + bestPoints[1].getPosition(), f);
    }

```

```

    }
    /**
     * Shift the best points toward negative positions.
     */
    private void moveTowardNegative()
    {
        OptimizingPoint newPoint = pointFactory.createPoint(
            3 * bestPoints[0].getPosition()
            - 2 * bestPoints[1].getPosition(), f);
        bestPoints[2] = bestPoints[1];
        bestPoints[1] = bestPoints[0];
        bestPoints[0] = newPoint;
    }
    /**
     * Shift the best points toward positive positions.
     */
    private void moveTowardPositive()
    {
        OptimizingPoint newPoint = pointFactory.createPoint(
            3 * bestPoints[2].getPosition()
            - 2 * bestPoints[1].getPosition(), f);
        bestPoints[0] = bestPoints[1];
        bestPoints[1] = bestPoints[2];
        bestPoints[2] = newPoint;
    }
}

```

11.6 Powell's algorithm

Powell's algorithm is one of many hill climbing algorithms [Press et al.]. The idea underlying Powell's algorithm is that once an optimum has been found in one direction, the chance for the biggest improvement lies in the direction perpendicular to that direction. A mathematical formulation of this sentence can be found in [Press et al.] and references therein. Powell's algorithm provides a way to keep track of the next best direction at each iteration step.

The original steps of Powell's algorithm are as follow:

1. Let \mathbf{x}_0 the best point so far and initialize a series of vectors $\mathbf{v}_1, \dots, \mathbf{v}_n$ forming the system of reference (n is the dimension of the vector \mathbf{x}_0); in other words the components of the vector \mathbf{v}_k are all zero except for the k^{th} components, which is one.
2. Set $k = 1$.
3. Find the optimum of the goal function along the direction \mathbf{v}_k starting from point \mathbf{x}_{k-1} . Let \mathbf{x}_k be the position of that optimum.

4. Set $k = k + 1$. If $k \leq n$, go back to step 3.
5. For $k = 1, \dots, n - 1$, set $\mathbf{v}_k = \mathbf{v}_{k-1}$.
6. Set $\mathbf{v}_n = \frac{\mathbf{x}_n - \mathbf{x}_0}{|\mathbf{x}_n - \mathbf{x}_0|}$.
7. Find the optimum of the goal function along the direction \mathbf{v}_n . Let \mathbf{x}_{n+1} be the position of that optimum.
8. If $|\mathbf{x}_n - \mathbf{x}_0|$ is less than the desired precision, terminate.
9. Otherwise, set $\mathbf{x}_0 = \mathbf{x}_{n+1}$ and go back to step 1.

There is actually two hill climbing algorithms within each other. The progression obtained by the inner loop is taken as the direction in which to continue the search.

Powell recommends to use this algorithm on goal functions having a quadratic behaviour near the optimum. It is clear that this algorithm cannot be used safely on any function. If the goal function has narrow valleys, all directions $\mathbf{v}_1, \dots, \mathbf{v}_n$ will become colinear when the algorithm ends up in such a valley. Thus, the search is likely to end up in a position where no optimum is located. Press et al. [Press et al.] mention two methods avoiding such problems: one method is quite complex and the other slows down the convergence of the algorithm.

In spite of this caveat, we implement Powell's algorithm in its original form. However, we recommend its use only in the vicinity of the minimum. In section 11.9 we show how other techniques can be utilized to read the vicinity of the optimum, where Powell's algorithm can safely be used to make the final determination of the optimum's position.

11.6.1 Powell's algorithm — General implementation

Since the class implementing the vector projected function $g(\lambda)$ described in sections 11.3.2 and 11.3.3 keep the vector \mathbf{x}_0 and \mathbf{v} in instance variables, there is no need to allocate explicit storage for the vectors $\mathbf{x}_1, \dots, \mathbf{x}_n$ and $\mathbf{v}_1, \dots, \mathbf{v}_n$. Instead, the class implementing Powell's algorithm keep an array of vector projected functions with the corresponding parameters. Then, the manipulation of the vector $\mathbf{x}_1, \dots, \mathbf{x}_n$ and $\mathbf{v}_1, \dots, \mathbf{v}_n$ is made directly on the projected function.

Since the origin of the projected function is always the starting value, \mathbf{x}_k , the initial value for the search of the optimum of the function $g(\lambda)$ is always 0.

The method `initializeIterations` allocated a series of vector projected functions starting with the axes of the reference system. This method also creates an instance of a one dimensional optimum finder kept in the instance variable, `unidimensionalFinder`. The goal function of the finder is alternatively assigned to each of the projected functions.

We made a slight modification to Powell's algorithm. If the norm of the vector $\mathbf{x}_n - \mathbf{x}_0$ at step 6 is smaller than the desired precision, the directions are only rotated, the assignment of step 6 is not done and the search of step 7 is omitted.

The precision computed at the end of each iterations is the maximum of the relative change on all components between the vectors \mathbf{x}_n and \mathbf{x}_0 .

11.6.2 Powell's algorithm — Smalltalk implementation

Listing 11.13 shows the implementation of Powell's algorithm in Smalltalk. The following code example shows how to find the maximum of a vector function

Figure 11.1 with the box HillClimbingOptimizer grayed.

Code example 11.3

```
| fBlock educatedGuess hillClimber result |
fBlock := <the goal function>
educatedGuess := <a vector not too far from the optimum>

hillClimber := DhbHillClimbingOptimizer maximizingFunction: fBlock.
hillClimber initialValue: educatedGuess.
result := hillClimber evaluate.
```

The class `DhbHillClimbingOptimizer` is a subclass of class `DhbFunctionOptimizer`. It has only one additional instance variable, `unidimensionalFinder`, to hold the one-dimensional optimizer used to find an optimum of the goal function along a given direction.

The method `evaluateIteration` uses the method `inject:into:` to perform steps 2-4 of the algorithm. Similarly step 5 of the algorithm is performed with a method `inject:into:` within the method `shiftDirection`. This mode of using the iterator method `inject:into:` — performing an action involving two consecutive elements of an indexed collection — is somewhat unusual, but highly convenient[Beck]. The method `minimizeDirection:` implements step 3 of the algorithm.

Listing 11.13 Smalltalk implementation of Powell's algorithm

Class	DhbHillClimbingOptimizer
Subclass of	DhbFunctionOptimizer

Instance variable names: `unidimensionalFinder`

Instance methods

`computeInitialValues`

```
unidimensionalFinder := DhbOneVariableFunctionOptimizer
                        forOptimizer: self.
unidimensionalFinder desiredPrecision: desiredPrecision.
bestPoints := ( 1 to: result size)
               collect: [ :n | ( DhbVectorProjectedFunction
                               function: functionBlock)
                           direction: ( ( DhbVector
                                           new: result size)
                                       atAllPut: 0;
                                       at: n put: 1;
                                       yourself);
                           yourself
               ].
```

evaluateIteration

```

| oldResult |
precision := 1.
bestPoints inject: result
    into: [ :prev :each | ( self minimizeDirection: each
                           from: prev)].

self shiftDirections.
self minimizeDirection: bestPoints last.
oldResult := result.
result := bestPoints last origin.
precision := 0.
result with: oldResult do:
    [ :x0 :x1 |
        precision := ( self precisionOf: (x0 - x1) abs relativeTo:
                                   x0 abs) max: precision.
    ].
^precision

```

finalizeIterations

minimizeDirection: aVectorFunction

```

^unidimensionalFinder
    reset;
    setFunction: aVectorFunction;
    addPointAt: 0;
    addPointAt: precision;
    addPointAt: precision negated;
    evaluate

```

minimizeDirection: aVectorFunction from: aVector

Function from: aVector

```

^aVectorFunction
    origin: aVector;
    argumentWith: ( self minimizeDirection: aVectorFunction)

```

shiftDirections

```

| position delta firstDirection |
firstDirection := bestPoints first direction.
bestPoints inject: nil
    into: [ :prev :each |

```

```

        position isNil
        ifTrue: [ position := each origin]
        ifFalse:[ prev direction: each
                    direction].
        each
    ].
    position := bestPoints last origin - position.
    delta := position norm.
    delta > desiredPrecision
        ifTrue: [ bestPoints last direction: (position scaleBy: (1 /
                                                                delta))]
        ifFalse:[ bestPoints last direction: firstDirection].

```

11.6.3 Powell's algorithm — Java implementation

Listing 11.14 shows the implementation of Powell's algorithm in Java. The following code example shows how to use this class to find the maximum of a vector function.

Figure 11.2 with the box HillClimbingOptimizer grayed.

Code example 11.4

```

ManyVariableFunction func = <the goal function>
double[] educatedGuess =<an array of double not too far from the optimum>

    MaximizingPointFactory strategy = new MaximizingPointFactory();
    HillClimbingOptimizer hillClimber =
        new HillClimbingOptimizer( func, strategy);
    hillClimber.setInitialValue( educatedGuess);
    hillClimber.evaluate();
    double[] result = hillClimber.getResult();

```

The class HillClimbingOptimizer is a subclass of class MultiVariableOptimizer. It has the following instance variables

unidimensionalFinder an instance of the class OneVariableFunctionOptimizer; this instance is used to find the optimum of the function in steps 3 and 7 of the algorithm;

projections an array of projected functions.

Listing 11.14 Java implementation of Powell's algorithm

```
package DhbOptimizing;
```

```
import DhbMatrixAlgebra.DhbVector;
import DhbMatrixAlgebra.DhbIllegalDimension;
```

```

import DhbInterfaces.ManyVariableFunction;
/**
 * Hill climbing optimizer using Powell's algorithm.
 *
 * @author Didier H. Besset
 */
public class HillClimbingOptimizer extends MultiVariableOptimizer
{
    /**
     * One dimensional optimizer used in each direction.
     */
    private OneVariableFunctionOptimizer unidimensionalFinder;
    /**
     * Projected goal function on independent directions.
     */
    private VectorProjectedFunction[] projections;
    /**
     * Constructor method.
     * @param func DhbInterfaces.ManyVariableFunction
     * @param pointCreator DhbOptimizing.OptimizingPointFactory
     */
    public HillClimbingOptimizer( ManyVariableFunction func,
                                OptimizingPointFactory pointCreator, double[] v)
    {
        super(func, pointCreator, v);
    }
    private void adjustLastDirection( DhbVector start)
    {
        try {
            int n = projections.length - 1;
            projections[n].setOrigin( result);
            DhbVector newDirection = projections[n].getOrigin()
                                .subtract( start);
            double norm = newDirection.norm();
            if ( norm > getDesiredPrecision() )
            {
                newDirection.scaledBy( 1 / norm);
                projections[n].setDirection( newDirection);
                unidimensionalFinder.setFunction( projections[n]);
                unidimensionalFinder.setInitialValue( 0);
                unidimensionalFinder.evaluate();
                result = projections[n].argumentAt(
                    unidimensionalFinder.getResult()).toComponents();
            }
        } catch ( DhbIllegalDimension e){};
    }
}

```

```

/**
 * @return double   relative precision of current result
 * @param x double[] result at previous iteration
 */
private double computePrecision( double[] x)
{
    double eps = 0;
    for ( int i = 0; i < result.length; i++)
        eps = Math.max( eps, relativePrecision(
            Math.abs( result[i] - x[i]), result[i]));
    return eps;
}
public double evaluateIteration()
{
    try {
        DhbVector start;
        start = new DhbVector( result);
        int n = projections.length;
        for( int i = 0; i < n; i++ )
        {
            projections[i].setOrigin( result);
            unidimensionalFinder.setFunction( projections[i]);
            unidimensionalFinder.setInitialValue( 0);
            unidimensionalFinder.evaluate();
            result = projections[i].argumentAt(
                unidimensionalFinder.getResult()).toComponents();
        }
        rotateDirections();
        adjustLastDirection( start);
        return computePrecision( start.toComponents());
    } catch (NegativeArraySizeException e){ return Double.NaN;}
    catch ( DhbIllegalDimension e){ return Double.NaN;};
}

public void initializeIterations()
{
    projections = new VectorProjectedFunction[ result.length];
    double [] v = new double[ result.length];
    for ( int i = 0; i < projections.length; i++ )
        v[i] = 0;
    for ( int i = 0; i < projections.length; i++ )
    {
        v[i] = 1;
        projections[i] = new VectorProjectedFunction( f, result, v);
        v[i] = 0;
    }
}

```

```

        unidimensionalFinder = new OneVariableFunctionOptimizer(
            projections[0], pointFactory);
        unidimensionalFinder.setDesiredPrecision( getDesiredPrecision());
    }
    private void rotateDirections()
    {
        DhbVector firstDirection = projections[0].getDirection();
        int n = projections.length;
        for ( int i = 1; i < n; i++ )
            projections[i-1].setDirection(projections[i].getDirection());
        projections[n-1].setDirection( firstDirection);
    }
    /**
     * Returns a String that represents the value of this object.
     * @return a string representation of the receiver
     */
    public String toString()
    {
        StringBuffer sb = new StringBuffer();
        sb.append( getIterations());
        sb.append( " iterations, precision = ");
        sb.append( getPrecision());
        sb.append( "\nResult:");
        for ( int i = 0; i < result.length; i++ )
        {
            sb.append( ' ');
            sb.append( result[i]);
        }
        for ( int i = 0 ; i < projections.length; i++ )
        {
            sb.append( '\n');
            sb.append( projections[i]);
        }
        return sb.toString();
    }
}

```

11.7 Simplex algorithm

The simplex algorithm, invented by Nelder and Mead, provides an efficient way to find a good approximation of the optimum of a function starting from any place [Press et al.]. The only trap into which the simplex algorithm can run into is a local optimum. On the other hand, this algorithm does not converge very well in the vicinity of the optimum. Thus, it must not be used with the desired precision set to a very low value. Once the

optimum has been found with the simplex algorithm, other more precise algorithms can then be used, such as the ones describes in section 11.1 or 11.6. MINUIT uses a combination of simplex and Newton algorithms. Our implementation of general purpose optimizer uses a combination of simplex and Powell algorithms.

A simplex in a n -dimensional space is a figure formed with $n + 1$ summits. For example, a simplex in a 2-dimensional space is a triangle; a simplex in a 3- dimensional space is a tetrahedron. Let us now discuss the algorithm for finding the optimum of a function with a simplex.

1. Pick up $n + 1$ points in the search space and evaluate the goal function at each of them. Let \mathbf{A} be the summit yielding the worst function's value.
2. if the size of the simplex is smaller than the desired precision, terminate the algorithm.
3. Calculate \mathbf{G} , the center of gravity of the n best points, that is all points except \mathbf{A} .
4. Calculate the location of the symmetric point of \mathbf{A} relative to \mathbf{G} : $\mathbf{A}' = 2\mathbf{G} - \mathbf{A}$.
5. If $f(\mathbf{A}')$ is not the best value found so far go to step 9.
6. Calculate the point $\mathbf{A}'' = 2\mathbf{A}' - \mathbf{G}$, that is a point twice as far from \mathbf{G} as \mathbf{A}' .
7. If $f(\mathbf{A}'')$ is a better value than $f(\mathbf{A}')$ build a new simplex with the point \mathbf{A} replaced by the point \mathbf{A}'' and go to step 2.
8. Otherwise, build a new simplex with the point \mathbf{A} replaced by the point \mathbf{A}' and go to step 2.
9. Calculate the point $\mathbf{B} = \frac{(\mathbf{G} + \mathbf{A})}{2}$.
10. If $f(\mathbf{B})$ yields the best value found so far build a new simplex with the point \mathbf{A} replaced by the point \mathbf{A}'' and go to step 2.
11. Otherwise build a new simplex obtained by dividing all edges leading to the point yielding the best value by 2 and go back to step 2.

Figure 11.4 shows the meaning of the operations involved in the algorithm in the 3 dimensional case. Step 6 makes the simplex grow into the direction where the function is the best so far. Thus, the simplex becomes elongated in the expected direction of the optimum. Because of its geometrical shape, the next step is necessarily taken along another direction, causing an exploration of the regions surrounding the growth obtained at the preceding step. Over the iterations, the shape of the simplex adapts itself to narrow valleys where the hill climbing algorithms notoriously get into trouble. Steps 9 and 11 ensures the convergence of the algorithm when the optimum lies inside the simplex. In this mode the simplex works very much like the golden section search or the bisection algorithms.

Finding the initial points can be done in several ways. If a good approximation of the region where the maximum might be located can be obtained one uses that




Figure 11.4: Operations of the simplex algorithm

approximation as a start and generate n other points by finding the optimum of the function along each axis. Otherwise, one can generate random points and select $n + 1$ points yielding the best values to build the initial simplex. In all cases, one must make sure that the initial simplex has a non-vanishing size in all dimensions of the space. Otherwise the algorithm will not reach the optimum.

11.7.1 Simplex algorithm — General implementation

The class implementing the simplex algorithm belong to the hierarchy of the iterative processes discussed in chapter 4. The method `evaluateIteration` directly implements the steps of the algorithm as described above. The points G , A' , A'' and B are calculated using the vector operations described in section 8.1.

The routine `initializeIterations` assumes that an initial value has been provided. It then finds the location of an optimum of the goal function along each axis of the reference system starting each time from the supplied initial value, unlike hill climbing algorithms. Restarting from the initial value is necessary to avoid creating a simplex with a zero volume. Such mishaps can arise when the initial value is located on an axis of symmetry of the goal function. This can happen quite frequently with educated guesses.

11.7.2 Simplex algorithm — Smalltalk implementation

Listing 11.15 shows the Smalltalk implementation of the simplex algorithm. The following code example shows how to invoke the class to find the minimum of a vector

Figure 11.1 with the box `SimplexOptimizer` grayed.

function.

Code example 11.5

```
| fBlock educatedGuess simplex result |
fBlock :=<the goal function>
educatedGuess :=<a vector in the search space>

simplex := DhbSimplexOptimizer minimizingFunction: fBlock.
simplex initialValue: educatedGuess.
result := simplex evaluate.
```

Except for the line creating the instance of the simplex optimizer, this code example is identical to the example of Powell's hill climbing algorithm (code example 11.3).

The class `DhbSimplexOptimizer` is a subclass of class `DhbFunctionOptimizer`. In order to be able to use the iterator methods efficiently, the worst point of the simplex, **A**, is held in a separate instance variable `worstPoint`. As we do not need to know the function's value $f(\mathbf{A})$, it is kept as a vector. The remaining points of the simplex are kept in the instance variable `bestPoints` of the superclass. Since this collection is sorted automatically when points are inserted to it, there is no explicit sorting step.

Listing 11.15 Smalltalk implementation of simplex algorithm

```
Class                DhbSimplexOptimizer
Subclass of         DhbFunctionOptimizer

Instance variable names:  worstVector

Class methods
defaultPrecision
    ^DhbFloatingPointMachine new defaultNumericalPrecision * 1000

Instance methods
buildInitialSimplex
    | projectedFunction finder partialResult |
    projectedFunction := DhbProjectedOneVariableFunction
        function: functionBlock.
    finder := DhbOneVariableFunctionOptimizer forOptimizer: self.
    finder setFunction: projectedFunction.
    [bestPoints size < (result size + 1)] whileTrue:
        [projectedFunction
            setArgument: result;
            bumpIndex.
            partialResult := finder
                reset;
```

```

        evaluate.
    bestPoints add: (optimizingPointClass
        vector: (projectedFunction argumentWith:
            partialResult)
        function: functionBlock)]

```

computeInitialValues

```

    bestPoints
        add: (optimizingPointClass vector: result function:
            functionBlock).
    self buildInitialSimplex.
    worstVector := bestPoints removeLast position

```

computePrecision

```

    | functionValues bestFunctionValue |
    functionValues := bestPoints collect: [ :each | each value].
    bestFunctionValue := functionValues removeFirst.
    ^functionValues inject: 0
        into: [ :max :each | ( self precisionOf: ( each -
            bestFunctionValue) abs relativeTo: bestFunctionValue abs) max: max]

```

contract

```

    | bestVector oldVectors |
    bestVector := bestPoints first position.
    oldVectors := OrderedCollection with: worstVector.
    [bestPoints size > 1] whileTrue: [oldVectors add: bestPoints
        removeLast position].
    oldVectors do: [:each | self contract: each around: bestVector].
    worstVector := bestPoints removeLast position.
    ^self computePrecision

```

contract: aVector around: bestVector

```

    bestPoints
        add: (optimizingPointClass vector: bestVector * 0.5 +
            (aVector * 0.5)
            function: functionBlock)

```

evaluateIteration

```

| centerOfGravity newPoint nextPoint |
centerOfGravity := (bestPoints inject: ((worstVector copy)
    atAllPut: 0;
    yourself)
    into: [:sum :each | each position + sum]) * (1 /
    bestPoints size).
newPoint := optimizingPointClass vector: 2 * centerOfGravity -
    worstVector
    function: functionBlock.
(newPoint betterThan: bestPoints first)
ifTrue:
    [nextPoint := optimizingPointClass
        vector: newPoint position * 2 -
            centerOfGravity
        function: functionBlock.
    (nextPoint betterThan: newPoint) ifTrue: [newPoint :=
        nextPoint]]
ifFalse:
    [newPoint := optimizingPointClass
        vector: centerOfGravity * 0.666667 +
            (worstVector * 0.333333)
        function: functionBlock.
    (newPoint betterThan: bestPoints first) ifFalse: [^self
        contract]].
worstVector := bestPoints removeLast position.
bestPoints add: newPoint.
result := bestPoints first position.
^self computePrecision

printOn: aStream

super printOn: aStream.
aStream cr.
worstVector printOn: aStream.

```

11.7.3 Simplex algorithm — Java implementation

Listing 11.15 shows the Java implementation of the simplex algorithm. The following code example shows how to invoke the class to find the minimum of a vector function.

Code example 11.6

```

ManyVariableFunction func = <the goal function>
double[] educatedGuess =<an array of double representing one point in the
search space>

```

Figure 11.2 with the box SimplexOptimizer grayed.

```

MinimizingPointFactory strategy = new MinimizingPointFactory();
SimplexOptimizer simplex = new SimplexOptimizer( func, strategy);
simplex.setInitialValue( educatedGuess);
simplex.evaluate();
double[] result = simplex.getResult();

```

Except for the lines creating the strategy and the optimizer, this code example is identical to that of the Powell's algorithm (code example 11.4).

The class `SimplexOptimizer` is a subclass of class `MultiVariableOptimizer`. The additional instance variable `simplex` contains an array of optimizing vectors whose positions are the summit of the simplex.

At the end of the contraction operation, the best points must be sorted — using the method `sortPoints` of the superclass — because this transformation can alter the order of the points considerably. In all other case, one has found the best point so far; the method `addBestPoint` is then used to discard to worst point, shift the remaining points and add the best point in the first position of the array `simplex`.

Listing 11.16 Java implementation of simplex algorithm

```

package DhbOptimizing;

import DhbMatrixAlgebra.DhbIllegalDimension;
import DhbMatrixAlgebra.DhbVector;
/**
 * Simplex optimizer of many-variable functions.
 *
 * @author Didier H. Besset
 */
public class SimplexOptimizer extends MultiVariableOptimizer
{
    /**
     * Best value found so far.
     */
    private OptimizingVector[] simplex;
    /**
     * Constructor method.
     * @param func DhbInterfaces.ManyVariableFunction
     * @param pointCreator DhbOptimizing.OptimizingPointFactory
     * @param initialValue double[]
     */
    public SimplexOptimizer(DhbInterfaces.ManyVariableFunction func,
        OptimizingPointFactory pointCreator, double[] initialValue)
    {
        super(func, pointCreator, initialValue);
    }
    /**

```

```

* Add a new best point to the simplex
* @param v DhbOptimizing.OptimizingVector
*/
private void addBestPoint( OptimizingVector v)
{
    int n = simplex.length;
    while ( --n > 0 )
        simplex[n] = simplex[n-1];
    simplex[0] = v;
}
/**
* @return boolean    true if a better point was found
* @param g DhbVector    summit whose median is contracted
* @exception DhbIllegalDimension if dimension of initial value is 0.
*/
private boolean addContraction(DhbVector g)
    throws DhbIllegalDimension
{
    g.accumulate( simplex[result.length].getPosition());
    g.scaledBy( 0.5);
    OptimizingVector contractedPoint =
        pointFactory.createVector( g, f);
    if ( contractedPoint.betterThan( simplex[0]) )
    {
        addBestPoint( contractedPoint);
        return true;
    }
    else
        return false;
}
/**
* @return boolean    true if a better point was found
* @exception DhbIllegalDimension if dimension of initial value is 0.
*/
private boolean addReflection( DhbVector centerOfgravity)
    throws DhbIllegalDimension
{
    DhbVector reflectedVector = centerOfgravity.product( 2);
    reflectedVector.accumulateNegated(
        simplex[result.length].getPosition());
    OptimizingVector reflectedPoint =
        pointFactory.createVector( reflectedVector, f);
    if ( reflectedPoint.betterThan( simplex[0]) )
    {
        reflectedVector.scaledBy( 2);
        reflectedVector.accumulateNegated( centerOfgravity);
    }
}

```

```

        OptimizingVector expandedPoint =
            pointFactory.createVector( reflectedVector, f);
        if ( expandedPoint.betterThan( reflectedPoint ) )
            addBestPoint( expandedPoint);
        else
            addBestPoint( reflectedPoint);
        return true;
    }
    else
        return false;
}
/**
 * @return DhbVector    center of gravity of best points of simplex,
 *                      except worst one
 */
private DhbVector centerOfGravity() throws DhbIllegalDimension
{
    DhbVector g = new DhbVector( result.length);
    for ( int i = 0; i < result.length; i++ )
        g.accumulate( simplex[i].getPosition());
    g.scaledBy( 1.0 / result.length);
    return g;
}
/**
 * @return double    maximum simplex extent in each direction
 */
private double computePrecision()
{
    int i, j;
    double[] position = simplex[0].getPosition();
    double[] min = new double[ position.length];
    double[] max = new double[ position.length];
    for ( i = 0; i < position.length; i++ )
    {
        min[i] = position[i];
        max[i] = position[i];
    }
    for ( j = 1; j < simplex.length; j++ )
    {
        position = simplex[j].getPosition();
        for ( i = 0; i < position.length; i++ )
        {
            min[i] = Math.min( min[i], position[i]);
            max[i] = Math.max( max[i], position[i]);
        }
    }
}

```

```

    double eps = 0;
    for ( i = 1; i < position.length; i++ )
        eps = Math.max( eps, relativePrecision( max[i]-min[i], result[i]));
    return eps;
}
/**
 * Reduce the simplex from the best point.
 */
private void contractSimplex()
{
    double[] bestPoint = simplex[0].getPosition();
    for ( int i = 1; i < simplex.length; i++)
        simplex[i].contractFrom( bestPoint);
    sortPoints( simplex);
}
/**
 * Here precision is the largest extent of the simplex.
 */
public double evaluateIteration()
{
    try {
        double bestValue = simplex[0].getValue();
        DhbVector g = centerOfGravity();
        if ( !addReflection( g ) )
        {
            if ( !addContraction( g ) )
                contractSimplex();
        }
        result = simplex[0].getPosition();
        return computePrecision();
    } catch ( DhbIllegalDimension e ) { return 1;};
}
/**
 * Create a Simplex by finding the optimum in each direction
 * starting from the initial value..
 */
public void initializeIterations()
{
    double [] v = new double[ result.length];
    for ( int i = 0; i < result.length; i++ )
        v[i] = 0;
    VectorProjectedFunction projection =
        new VectorProjectedFunction( f, result, v);
    OneVariableFunctionOptimizer unidimensionalFinder =
        new OneVariableFunctionOptimizer( projection, pointFactory);
    unidimensionalFinder.setDesiredPrecision( getDesiredPrecision());
}

```



```

simplex = new OptimizingVector[result.length+1];
try {
    for ( int i = 0; i < result.length; i++ )
    {
        v[i] = 1;
        projection.setDirection( v);
        v[i] = 0;
        unidimensionalFinder.setInitialValue( 0);
        unidimensionalFinder.evaluate();
        simplex[i] = pointFactory.createVector(
            projection.argumentAt(
                unidimensionalFinder.getResult()), f);
    }
} catch ( DhbIllegalDimension e) { };
simplex[result.length] = pointFactory.createVector( result, f);
sortPoints( simplex);
}
/**
 * Returns a String that represents the value of this object.
 * @return a string representation of the receiver
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    sb.append( simplex[0]);
    for ( int i = 1; i < simplex.length; i++ )
    {
        sb.append( '\n');
        sb.append( simplex[i]);
    }
    return sb.toString();
}
}

```

11.8 Genetic algorithm

All optimizing algorithm discussed so far have one common flaw: they all terminate when a local optimum is encountered. In most problems, however, one wants to find the absolute optimum of the function. This is especially true if the goal function represents some economical merit.

One academic example is the maximization of the function

$$f(\mathbf{x}) = \frac{\sin^2 |\mathbf{x}|}{|\mathbf{x}|^2}. \quad (11.7)$$




Figure 11.5: Mutation and crossover reproduction of chromosomes

This function has an absolute maximum at $x = 0$, but all algorithms discussed so far will end up inside a ring corresponding to $|x| = n\pi/2$ where n is any positive odd integer.

In 1975 John Holland introduced a new type of algorithm — dubbed genetic algorithm — because it tries to mimic the evolutionary process identified as the cause for the diversity of living species by Charles Darwin. In a genetic algorithm the elements of the search space are considered as the chromosomes of individuals; the goal function is considered as the measure of the fitness of the individual to adapt itself to its environment[Berry & Linoff][Koza et al.]. The iterations are aping (pun intended) the Darwinian principle of survival and reproduction. At each iteration, the fittest individuals survive and reproduce themselves. To bring some variability to the algorithm mutation and crossover of chromosomes are taken into account.

Mutation occurs when one gene of a chromosome is altered at reproduction time. Crossover occurs when two chromosomes break themselves and recombine with the piece coming from the other chromosome. These processes are illustrated on figure 11.5. The point where the chromosomes are breaking is called the crossover point. Which individual survives and reproduces itself, when and where mutation occurs and when and where a crossover happens is determined randomly. This is precisely the random nature of the algorithm which gives it the ability to jump out of a local optimum to look further for the absolute optimum.

Mapping the search space on chromosomes

To be able to implement a genetic algorithm one must establish how to represent the genes of a chromosome. At the smallest level the genes could be the bits of the structure representing the chromosome. If the search space of the goal function do cover the domain generated by all possible permutations of the bits, this is a good approach. However, this is not always a practical solution since some bit combinations may be forbidden by the structure. For example, some of the combinations of a 64 bit word do not correspond to a valid floating point number.

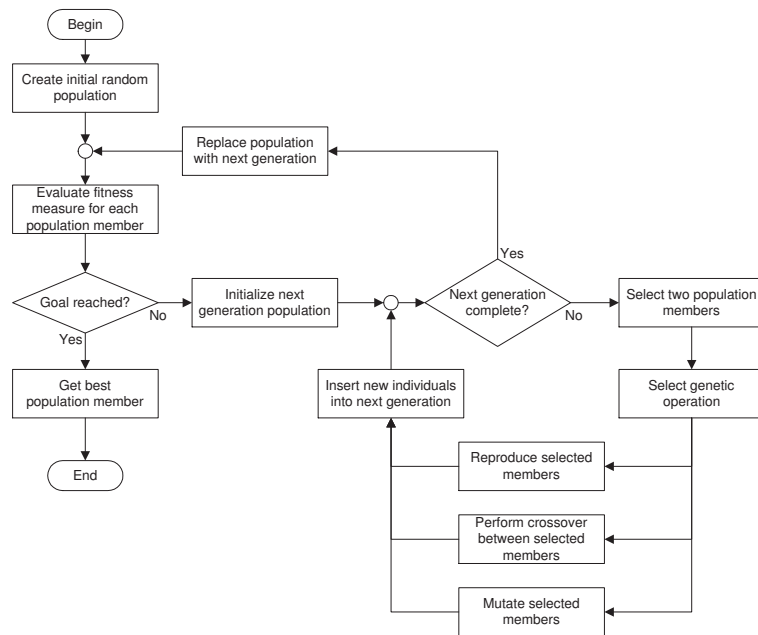


Figure 11.6: General purpose genetic algorithm

In the case of the optimization of a vector function, the simplest choice is to take the components of the vector as the genes. Genetic algorithms are used quite often to adjust the parameters of a neural network [Berry & Linoff]. In this case, the chromosomes are the coefficients of each neuron. Chromosomes can even be computer subprograms in the case of genetic programming [Koza et al.]. In this latter case, each individual is a computer program trying to solve a given problem.

Figure 11.6 shows a flow diagram of a general genetic algorithm. The reproduction of the individual is taken literally: a copy of the reproducing individual is copied into the next generation. The important feature of a generic algorithm is that the building of the next generation is a random process. To ensure the survival of the fittest, the selection of the parents of the individuals of the next generation is performed at random with uneven probability: the fittest individuals have a larger probability of being selected than the others. Mutation enables the algorithm to create individuals having genes corresponding to unexplored regions of the search space. Most of the times such mutants will be discarded at the next iteration; but, in some cases, a mutation may uncover a better candidate. In the case of the function of equation 11.7, this would correspond to jumping from one ring to another ring closer to the function's maximum. Finally the crossover operation mixes good genes in the hope of building a better individual out of the properties of good individuals. Like mutation the crossover operation gives a stochastic behavior to the algorithm enabling it to explore uncharted regions of the search space.

Note: Because of its stochastic nature a genetic algorithm is the algorithm

of choice when the goal function is expressed on integers.

11.8.1 Genetic algorithm — General implementation

The left hand side of the diagram of figure 11.6 is quite similar to the flow diagram of an iterative process (c.f. figure 4.2 in chapter 4). Thus, the class implementing the genetic algorithm is a subclass of the iterative process class discussed in chapter 4.

The genetic nature of the algorithm is located in the right hand side of the diagram of figure 11.6. As we have mentioned before the implementation of the chromosomes is highly problem dependent. All operations located in the top portion of the mentioned area can be expressed in generic terms without any knowledge of the chromosomic implementation. To handle the lower part of the right hand side of the diagram of figure 11.6, we shall implement a new object, the chromosome manager.

One should also notice that the value of the function is not needed when the next generation is build. Thus, the chromosome manager does not need to have any knowledge of the goal function. The goal function comes into play when transferring the next generation to the mature population, that is, the population used for reproduction at the next iteration. At the maturity stage, the value of the goal function is needed to identify the fittest individuals. In our implementation, the next generation is maintained by the chromosome manager whereas the population of mature individuals is maintained by the object in charge of the genetic algorithm which has the knowledge of the goal function.

The chromosome manager has the following instance variables:

`populationSize` contains the size of the population; one should pick up a large enough number to be able to cover the search space efficiently: the larger the dimension of the space search space, the larger must be the population size;

`rateOfMutation` contains the probability of having a mutation while reproducing;

`rateOfCrossover` contains the probability of having a crossover while reproducing.

All of these variables have getter and setter accessor methods. In addition a convenience instance creation method is supplied to create a chromosome manager with given values for all three instance variables. The chromosome manager implements the following methods:

`isFullyPopulated` to signal that a sufficient number of individuals has been generated into the population;

`process` to process a pair of individuals; this method does the selection of the genetic operation and applies it; individuals are processed by pair to always have a possibility of crossover;

`randomizePopulation` to generate a random population;

`reset` to create an empty population for the next generation.

Finally the chromosome manager must also implement methods performing each of the genetic operations: reproduction, mutation and crossover. The Smalltalk implementation supplies methods that returns a new individual; the Java implementation supplies methods that add a new individual to the population. The reason for this difference come from the static typing requirements of Java.

The genetic optimizer is the object implementing the genetic algorithm proper. It is a subclass of the iterative process class described in chapter 4.2. In addition to the handling of the iterations the genetic optimizer implements the steps of the algorithm drawn on the top part of the right hand side of the diagram of figure 11.6. It has one instance variable containing the chromosome manager with which it will interact. The instance creation method take three arguments: the function to optimize, the optimizing strategy and the chromosome manager.

The method `initializeIteration` asks the chromosome manager to supply a random population. The method `evaluateIteration` performs the loop of the right hand side of the diagram of figure 11.6. It selects a pair of parents on which the chromosome manager performs the genetic operation.

Selecting the genetic operation is performed with a random generator. The values of the goal function are used as weights. Let $f(p_i)$ be the value of the goal function for individual p_i and let p_b and p_w be respectively the fittest and the lest fit individual found so far (b stands for best and w stands for worst). One first computes the unnormalized probability:

$$\tilde{P}_i = \frac{f(p_i) - f(p_w)}{f(p_b) - f(p_w)}. \quad (11.8)$$

This definition ensures that \tilde{P}_i is always comprised between 0 and 1 for any goal function. Then we can use the discrete probability

$$P_i = \frac{1}{\sum \tilde{P}_i} \tilde{P}_i. \quad (11.9)$$

The sum in equation 11.9 is taken over the entire population. An attentive reader will notice that this definition assigns a zero probability of selecting the worst individuals. This gives a slight bias to our implementation compared to the original algorithm. This is can be easily compensated by taking a sufficiently large population. The method `randomScale` calculates the P_i of equation 11.9 and returns an array containing the integrated sums:

$$R_i = \sum_{k=0}^i P_k. \quad (11.10)$$

The array R_i is used to generate a random index to select individuals for reproduction. The transfer between the next generation and the mature population is performed by the method `collectPoints`.

In the general case, there is no possibility to decide when to terminate the algorithm. In practice, it is possible that the population stays stable for quite a while until suddenly a new individual is found to be better than the rest. Therefore a criteria based on the stability of the first best points is likely to be beat the purpose of the algorithm, namely

to jump out of a local optimum. Some problems can define a threshold at which the goal function is considered sufficiently good. In this case, the algorithm can be stopped as soon as the value of the goal function for the fittest individual becomes better than that threshold. In the general case, however, the implementation of the genetic algorithm simply returns a constant pseudo precision — set to one — and runs until the maximum number of iterations becomes exhausted.

11.8.2 Genetic algorithm — Smalltalk implementation

Listing 11.17 shows the code of an abstract chromosome manager in Smalltalk and of a concrete implementation for vector chromosomes. The class `DhbChromosomeManager` has one instance variable in addition to the variables listed in section 11.8.1: `population`. This variable is an instance of an `OrderedCollection` containing the individuals of the next generation being prepared.

The class `DhbVectorChromosomeManager` is a subclass of class `DhbChromosomeManager` implementing vector chromosomes. It has two instance variables

`origin` a vector containing the minimum possible values of the generated vectors;

`range` a vector containing the range of the generated vectors.

In other words `origin` and `range` are delimiting an hypercube defining the search space.

Figure 11.1 with the boxes `GeneticOptimizer`, `ChromosomeManager` and `VectorChromosomeManager` grayed.

Listing 11.17 Smalltalk chromosome: abstract and concrete

```

Class                                DhbChromosomeManager
Subclass of                          Object

Instance variable names:  population populationSize rateOfMutation rateOfCrossover

Class methods

new: anInteger mutation: aNumber1 crossover: aNumber2
    ^self new populationSize: anInteger; rateOfMutation: aNumber1;
        rateOfCrossover: aNumber2; yourself

Instance methods

clone: aChromosome
    ^aChromosome copy

crossover: aChromosome1 and: aChromosome2
    ^self subclassResponsibility

isFullyPopulated
```

```

    ^population size >= populationSize

mutate: aChromosome
    ^self subclassResponsibility

population
    ^population

populationSize: anInteger
    populationSize := anInteger.

process: aChromosome1 and: aChromosome2
    | roll |
    roll := Number random.
    roll < rateOfCrossover
        ifTrue: [population addAll: (self crossover: aChromosome1
                                   and: aChromosome2)]
        ifFalse:
            [roll < (rateOfCrossover + rateOfMutation)
             ifTrue:
                 [population
                  add: (self mutate: aChromosome1);
                  add: (self mutate: aChromosome2)]
             ifFalse:
                 [population
                  add: (self clone: aChromosome1);
                  add: (self clone: aChromosome2)]]

randomizePopulation
    self reset.
    [ self isFullyPopulated] whileFalse: [ population add: self
                                           randomChromosome].

rateOfCrossover: aNumber
    (aNumber between: 0 and: 1)
    ifFalse: [self error: 'Illegal rate of cross-over'].
    rateOfCrossover := aNumber

```

rateOfMutation: aNumber

```
(aNumber between: 0 and: 1)
  ifFalse: [self error: 'Illegal rate of mutation'].
rateOfMutation := aNumber
```

reset

```
population := OrderedCollection new: populationSize.
```

Class **DhbVectorChromosomeManager**
 Subclass of DhbChromosomeManager

Instance variable names: origin range

Instance methods

crossover: aChromosome1 and: aChromosome2

```
| index new1 new2 |
index := ( aChromosome1 size - 1) random + 2.
new1 := self clone: aChromosome1.
new1 replaceFrom: index to: new1 size with: aChromosome2
      startingAt: index.
new2 := self clone: aChromosome2.
new2 replaceFrom: index to: new2 size with: aChromosome1
      startingAt: index.
^Array with: new1 with: new2
```

mutate: aVector

```
| index |
index := aVector size random + 1.
^( aVector copy)
  at: index put: ( self randomComponent: index);
  yourself
```

origin: aVector

```
origin := aVector.
```

randomChromosome

```
^( ( 1 to: origin size) collect: [ :n | self randomComponent: n])
  asVector
```


randomComponent: anInteger

\wedge (range at: anInteger) random + (origin at: anInteger)

range: aVector

range := aVector.

Listing 11.18 shows how the genetic optimizer is implemented in Smalltalk. The following code example shows how to use a genetic optimizer to find the maximum of a vector function.

Code example 11.7

```

| fBlock optimizer manager origin range result |
fBlock := <the goal function>
origin := <a vector containing the minimum expected value of the component>

range := <a vector containing the expected range of the component>

optimizer := DhbGeneticOptimizer maximizingFunction: fBlock.
manager := DhbVectorChromosomeManager new: 100 mutation: 0.1 crossover: 0.1.
manager origin: origin; range: range.
optimizer chromosomeManager: manager.
result := optimizer evaluate.

```

After establishing the goal function and the search space, an instance of the genetic optimizer is created. The next line creates an instance of a vector chromosome manager for a population of 100 individuals (sufficient for a 2-3 dimensional space) and rates of mutation and crossover equal to 10%. The next line defines the search space into the chromosome manager. The final line performs the genetic search and returns the result.

In Smalltalk the population of the next generation is maintained in the instance variable population. Each time a next generation has been established, it is transferred into a collection of best points by the method collectPoints. Each element of the collection bestPoints is an instance of an subclass of OptimizingPoint. The exact type of the class is determined by the search strategy. Since best points are sorted automatically, the result is always the position of the first element of bestPoints.

Listing 11.18 Smalltalk implementation of genetic algorithm

Class	DhbGeneticOptimizer
Subclass of	DhbFunctionOptimizer
Instance variable names:	chromosomeManager

Class methods

defaultMaximumIterations

^500

defaultPrecision

^0

Instance methods

chromosomeManager: aChromosomeManager

chromosomeManager := aChromosomeManager.
^self

collectPoints

| bestPoint |
bestPoints notEmpty
 ifTrue: [bestPoint := bestPoints removeFirst].
bestPoints removeAll: bestPoints asArray.
chromosomeManager population do: [:each | self addPointAt: each].
bestPoint notNil
 ifTrue: [bestPoints add: bestPoint].
result := bestPoints first position.

computePrecision

^1

evaluateIteration

| randomScale |
randomScale := self randomScale.
chromosomeManager reset.
[chromosomeManager isFullyPopulated]
 whileFalse: [self processRandomParents: randomScale].
self collectPoints.
^self computePrecision

initializeIterations

chromosomeManager randomizePopulation.
self collectPoints

processRandomParents: aNumberArray

```

chromosomeManager process: ( bestPoints at: ( self randomIndex:
                                aNumberArray)) position
and: ( bestPoints at: ( self randomIndex:
                                aNumberArray)) position.

```

```

randomIndex: aNumberArray

```

```

| x n |
x := Number random.
n := 1.
aNumberArray do:
    [ :each |
        x < each
            ifTrue: [ ^n].
        n := n + 1.
    ].
^aNumberArray size

```

```

randomScale

```

```

| norm fBest fWorst answer|
fBest := bestPoints first value.
fWorst := bestPoints last value.
norm := 1 / ( fBest - fWorst).
answer := bestPoints collect: [ :each | (each value - fWorst) *
                                         norm].
norm := 1 / ( answer inject: 0 into: [ :sum :each | each + sum]).
fBest := 0.
^answer collect: [ :each | fBest := each * norm + fBest. fBest]

```

Figure 11.2 with the boxes ChromosomeManager and GeneticOptimizer grayed.

11.8.3 Genetic algorithm — Java abstract implementation

The Java implementation of the genetic algorithm is done in two parts because the type of the population is not known until one uses a concrete chromosome manager. One could use casting in order to preserve anonymity of the chromosome type at the level of the genetic algorithm. I tend to avoid cast operator like the plague and use them only when all other alternatives have failed. In this case, casting is avoided by maintaining a concrete implementation of the chromosome manager in pair with a class implementing the genetic algorithm. This section describes the features of the abstract classes. The code for these classes is shown in listings 11.19 and 11.20. The next section describes a concrete implementation for vector functions.

The class ChromosomeManager has an additional instance variable — called generator — used to keep an instance of a random generator. Our implementation uses the

default random generator provided by Java. The method `reset` allows the creation of an empty next generation. The current size of the next generation is obtained by calling method `getPopulationSize`. Other methods are provided to provide interaction with the genetic algorithm class. The methods `individualAt` returns the k^{th} individual of the next generation, where k is the supplied integer argument. The methods `addCloneOf`, `addMutationOf` add a new individual to the next generation by applying respectively reproduction or mutation. The method `addCrossoversOf` add the two crossover offsprings of two individuals. Here, the reader will see that casting becomes unavoidable by the subclasses. However, since the objects taken as arguments by the methods `addCloneOf`, `addMutationOf` and `addCrossoversOf` are coming from the method `individualAt` and only that method, the risk of error is minimal.

Listing 11.19 Java abstract implementation of a chromosome

```
package DhbOptimizing;

import java.util.Random;
/**
 * Abstract chromosome manager.
 * (genetic algorithm)
 *
 * @author Didier H. Besset
 */
public abstract class ChromosomeManager
{
    /**
     * Population size.
     */
    private int populationSize = 100;
    /**
     * Rate of mutation.
     */
    private double rateOfMutation = 0.1;
    /**
     * Rate of crossover.
     */
    private double rateOfCrossover = 0.1;
    /**
     * Random generator.
     */
    private Random generator = new Random();
    /**
     * Constructor method.
     */
    public ChromosomeManager()
    {
```

```

    super();
}
/**
 * Constructor method.
 * @param n int
 * @param mRate double
 * @param cRate double
 */
public ChromosomeManager(int n, double mRate, double cRate)
{
    populationSize = n;
    rateOfMutation = mRate;
    rateOfCrossover = cRate;
}
/**
 * @param x java.lang.Object
 */
public abstract void addCloneOf( Object x);
/**
 * @param x java.lang.Object
 */
public abstract void addCrossoversOf( Object x, Object y);
/**
 * @param x java.lang.Object
 */
public abstract void addMutationOf( Object x);
public abstract void addRandomChromosome();
/**
 * @return int    the current size of the population
 */
public abstract int getCurrentPopulationSize();
/**
 * @return int    desired population size.
 */
public int getPopulationSize( )
{
    return populationSize;
}
/**
 * @return java.lang.Object (must be casted into the proper type
 *                          of chromosome)
 * @param n int
 */
public abstract Object individualAt( int n);
/**
 * @return boolean    true if the new generation is complete

```

```

    */
    public boolean isFullyPopulated()
    {
        return getCurrentPopulationSize() >= populationSize;
    }
    /**
     * @return double    a random number (delegated to the generator)
     */
    public double nextDouble()
    {
        return generator.nextDouble();
    }
    /**
     * @param x java.lang.Object
     * @param y java.lang.Object
     */
    public void process( Object x, Object y)
    {
        double roll = generator.nextDouble();
        if ( roll < rateOfCrossover )
            addCrossoversOf( x, y);
        else if ( roll < rateOfCrossover + rateOfMutation )
        {
            addMutationOf( x);
            addMutationOf( y);
        }
        else
        {
            addCloneOf( x);
            addCloneOf( y);
        }
    }
    /**
     * Create a population of random chromosomes.
     */
    public void randomizePopulation()
    {
        reset();
        while ( !isFullyPopulated() )
            addRandomChromosome();
    }
    /**
     * Reset the population of the receiver.
     */
    public abstract void reset();
    /**

```

```

    * @param n int    desired population size.
    */
    public void setPopulationSize( int n)
    {
        populationSize = n;
    }
    /**
    * @param n int    desired rate of crossover
    */
    public void setRateOfCrossover( int cRate)
    {
        rateOfCrossover = cRate;
    }
    /**
    * @param n int    desired rate of mutation
    */
    public void setRateOfMutation( int mRate)
    {
        rateOfMutation = mRate;
    }
}

```

The class `GeneticOptimizer` is an abstract class. The methods providing the functionality to fill up the new population from the next generation established by the chromosome manager are all abstract methods. The method `collectPoint` transfers a single individual from the next generation to the population of mature individuals.

A method `initializeIterations` with one integer argument is needed to allow the concrete class to initialize its memory for the desired population size. This allows the instance variable `populationSize` to remain private.

Listing 11.20 Java implementation of genetic algorithm

```

package DhbOptimizing;

/**
 * Abstract genetic algorithm.
 *
 * @author Didier H. Besset
 */
public abstract class GeneticOptimizer extends MultiVariableOptimizer
{
    /**
     * Chromosome manager.
     */
    private ChromosomeManager chromosomeManager;
}

```

```

/**
 * Constructor method.
 * @param func DhbInterfaces.ManyVariableFunction
 * @param pointCreator DhbOptimizing.OptimizingPointFactory
 * @param chrManager ChromosomeManager
 */
public GeneticOptimizer(DhbInterfaces.ManyVariableFunction func,
    OptimizingPointFactory pointCreator, ChromosomeManager chrManager)
{
    super(func, pointCreator, null);
    chromosomeManager = chrManager;
}
/**
 * @param x java.lang.Object
 */
public abstract void collectPoint(Object x);
/**
 * Collect points for the entire population.
 */
public void collectPoints()
{
    reset();
    for ( int i = 0; i < chromosomeManager.getPopulationSize(); i++ )
        collectPoint( chromosomeManager.individualAt(i));
}
/**
 * This method causes the receiver to exhaust the maximum number of
 * iterations. It may be overloaded by a subclass (hence "protected")
 * if a convergence criteria can be defined.
 * @return double
 */
protected double computePrecision()
{
    return 1;
}
/**
 * @return double
 */
public double evaluateIteration()
{
    double[] randomScale = randomScale();
    chromosomeManager.reset();
    while ( !chromosomeManager.isFullyPopulated() )
    {
        chromosomeManager.process(
            individualAt( randomIndex(randomScale)),

```



```

        individualAt( randomIndex(randomScale)));
    }
    collectPoints();
    return computePrecision();
}
/**
 * @return java.lang.Object (must be casted into the proper type)
 * @param n int
 */
public abstract Object individualAt(int n);
/**
 * Create a random population.
 */
public void initializeIterations()
{
    initializeIterations( chromosomeManager.getPopulationSize());
    chromosomeManager.randomizePopulation();
    collectPoints();
}
/**
 * @param n int size of the initial population
 */
public abstract void initializeIterations( int n);
/**
 * @return int an index generated randomly
 * @param randomScale double[] fitness scale (integral)
 */
protected int randomIndex( double[] randomScale)
{
    double roll = chromosomeManager.nextDouble();
    if ( roll < randomScale[0] )
        return 0;
    int n = 0;
    int m = randomScale.length;
    int k;
    while ( n < m - 1 )
    {
        k = ( n + m ) / 2;
        if ( roll < randomScale[k] )
            m = k;
        else
            n = k;
    }
    return m;
}
/**

```

```

* @return double[]    integral fitness scale.
*/
public abstract double[] randomScale();
public abstract void reset();
}

```

11.8.4 Genetic algorithm — Java implementation with vectors

Listings 11.21 and 11.22 show the concrete Java classes for the chromosome manager and the genetic optimizer respectively. They provide a concrete implementation of a genetic algorithm for finding the optimum of a vector function. The following code example shows how to find the maximum of a vector function using these classes.

Figure 11.2 with the boxes `VectorChromosomeManager` and `VectorGeneticOptimizer` grayed.

Code example 11.8

```

ManyVariableFunction func = <the goal function>
DhbVector origin =<a vector containing the minimum expected value of
the component>
DhbVector range =<a vector containing the expected range of the component>

```

```

MaximizingPointFactory strategy = new MaximizingPointFactory();
VectorChromosomeManager manager =
    new VectorChromosomeManager( 100, 0.1, 0.1);
manager.setOrigin( origin);
manager.setRange( range);
VectorGeneticOptimizer finder =
    new VectorGeneticOptimizer( func, strategy, manager);
finder.evaluate();
double[] result = finder.getResult();

```

The line after the definition of the goal function and the hypercube defining the search space creates an instance of a maximizing vector. This will be the Strategy of the genetic optimizer. The next statement creates an instance of a vector chromosome manager. Right after, two statements define the search space. Then, the instance of the genetic optimizer is created. The next statement performs the genetic algorithm and the last statement retrieves the result.

The class `VectorChromosomeManager` is a concrete subclass of the class `ChromosomeManager`. It maintains the next generation in the instance variable `population`. The instance variable `fillIndex` is used as an index when filling up the next generation.

Listing 11.21 Java implementation of a vector chromosome

```

package DhbOptimizing;

```

```

import DhbMatrixAlgebra.DhbVector;
/**
 * Chromosome manager for vector chromosomes.
 * (genetic algorithm)
 *
 * @author Didier H. Besset
 */
public class VectorChromosomeManager extends ChromosomeManager
{
    /**
     * Population.
     */
    private DhbVector[] population;
    /**
     * Current population size.
     */
    private int fillIndex;
    /**
     * Origin of values.
     */
    private DhbVector origin;
    /**
     * Range of values.
     */
    private DhbVector range;
    /**
     * Default constructor method.
     */
    public VectorChromosomeManager() {
        super();
    }
    /**
     * Constructor method.
     * @param n int
     * @param mRate double
     * @param cRate double
     */
    public VectorChromosomeManager(int n, double mRate, double cRate) {
        super(n, mRate, cRate);
    }
    /**
     * @param x DhbVector
     */
    public void addCloneOf(Object x)
    {
        double[] v = ((DhbVector) x).toComponents();

```

```

    try { population[fillIndex++] = new DhbVector( v);}
        catch( NegativeArraySizeException e) { };
    }
    /**
    * @param x DhbVector
    * @param y DhbVector
    */
    public void addCrossoversOf(Object x, Object y)
    {
        double[] v = ((DhbVector) x).toComponents();
        double[] w = ((DhbVector) x).toComponents();
        int n = (int) ( nextDouble() * ( origin.dimension() - 1));
        double temp;
        for ( int i = 0; i < n; i++ )
        {
            temp = v[i];
            v[i] = w[i];
            w[i] = temp;
        }
        try { population[fillIndex++] = new DhbVector( v);
            population[fillIndex++] = new DhbVector( w);
        } catch( NegativeArraySizeException e) { };
    }
    /**
    * @param x DhbVector
    */
    public void addMutationOf(Object x)
    {
        double[] v = ((DhbVector) x).toComponents();
        int i = (int) ( nextDouble() * origin.dimension());
        v[i] = randomComponent(i);
        try { population[fillIndex++] = new DhbVector( v);}
            catch( NegativeArraySizeException e) { };
    }
    public void addRandomChromosome()
    {
        double[] v = new double[origin.dimension()];
        for ( int i = 0; i < origin.dimension(); i++ )
            v[i] = randomComponent(i);
        try { population[fillIndex++] = new DhbVector( v);}
            catch( NegativeArraySizeException e) { };
    }
    /**
    * @return int    the current size of the population
    */
    public int getCurrentPopulationSize()

```

```
{
    return fillIndex;
}
/**
 * @return Vector    vector at given index
 * @param n int
 */
public Object individualAt( int n)
{
    return population[n];
}
/**
 * @return double
 * @param n int
 */
private double randomComponent(int n)
{
    return origin.component(n) + nextDouble() * range.component(n);
}
/**
 * Allocated memory for a new generation.
 */
public void reset()
{
    population = new DhbVector[getPopulationSize()];
    fillIndex = 0;
}
/**
 * @param x double    component of the origin of the hypercube
 *    constraining the domain of definition of the function to optimize
 * @exception java.lang.NegativeArraySizeException
 *    when the size of the array is 0
 */
public void setOrigin( double[] x) throws NegativeArraySizeException
{
    setOrigin( new DhbVector( x));
}
/**
 * @param v DhbVector    origin of the hypercube
 *    constraining the domain of definition of the function to optimize
 */
public void setOrigin( DhbVector v)
{
    origin = v;
}
/**
```

```

* @param x double    components of the lengths of the hypercube
*    constraining the domain of definition of the function to optimize
* @exception java.lang.NegativeArraySizeException
*    when the size of the array is 0
*/
public void setRange( double[] x) throws NegativeArraySizeException
{
    setRange( new DhbVector( x));
}
/**
* @param v DhbVector    lengths of the hypercube
*    constraining the domain of definition of the function to optimize
*/
public void setRange( DhbVector v)
{
    range = v;
}
}

```

The class `VectorGeneticOptimizer` is a concrete subclass of the class `GeneticOptimizer`. Our implementation chooses to sort the points when they are collected. This is not really needed as the algorithm for selecting the individual does not assumes that the individuals are sorted by fitness. This is quite practical, however, when following the behavior of the algorithm with the debugger. For heavy duty usage the sorting ought to be removed. If that is the case, the method `getResult` must be rewritten to fetch the fittest individual of the mature population.

Listing 11.22 Java implementation of genetic algorithm for vectors

```

package DhbOptimizing;

import DhbMatrixAlgebra.DhbVector;
import DhbInterfaces.ManyVariableFunction;
/**
* Genetic optimizer of many-variable functions.
*
* @author Didier H. Besset
*/
public class VectorGeneticOptimizer extends GeneticOptimizer
{
    /**
    * Best values found so far.
    */
    private OptimizingVector[] bestPoints;
    /**

```

```

        * Number of points filled so far.
        */
    private int fillIndex;
/**
 * Constructor method.
 * @param func DhbInterfaces.ManyVariableFunction
 * @param pointCreator DhbOptimizing.OptimizingPointFactory
 * @param chrManager DhbOptimizing.ChromosomeManager
 */
    public VectorGeneticOptimizer( ManyVariableFunction func,
                                   OptimizingPointFactory pointCreator,
                                   ChromosomeManager chrManager)
    {
        super(func, pointCreator, chrManager);
    }
/**
 * @param x DhbVector
 */
    public void collectPoint(Object x)
    {
        OptimizingVector v = pointFactory.createVector( (DhbVector) x, f);
        if ( fillIndex == 0 || bestPoints[fillIndex-1].betterThan(v) )
        {
            bestPoints[fillIndex++] = v;
            return;
        }
        int n = 0;
        int m = fillIndex - 1;
        if ( bestPoints[0].betterThan(v) )
        {
            int k;
            while ( m - n > 1 )
            {
                k = ( n + m ) / 2;
                if ( v.betterThan(bestPoints[k]) )
                    m = k;
                else
                    n = k;
            }
            n = m;
        }
        for ( m = fillIndex; m > n; m-- )
            bestPoints[m] = bestPoints[m-1];
        bestPoints[n] = v;
        fillIndex += 1;
    }

```

```

/**
 * @return double[]    best point found so far
 */
public double[] getResult()
{
    return bestPoints[0].getPosition();
}
/**
 * @return DhbVector    vector at given index
 * @param n int
 */
public Object individualAt( int n)
{
    try { return new DhbVector( bestPoints[n].getPosition());}
    catch( NegativeArraySizeException e) { return null;};
}
/**
 * @param n int    size of the initial population
 */
public void initializeIterations( int n)
{
    bestPoints = new OptimizingVector[n];
}
/**
 * @return double[]    fitness scale for random generation
 */
public double[] randomScale()
{
    double[] f = new double[ bestPoints.length];
    double sum = 0;
    for ( int i = 0; i < bestPoints.length; i++ )
    {
        f[i] = bestPoints[i].getValue() + sum;
        sum += bestPoints[i].getValue();
    }
    sum = 1 / sum;
    for ( int i = 0; i < bestPoints.length; i++ )
        f[i] *= sum;
    return f;
}
public void reset()
{
    fillIndex = 0;
}
/**
 * Returns a String that represents the value of this object.

```



```

    * @return a string representation of the receiver
    */
    public String toString()
    {
        StringBuffer sb = new StringBuffer();
        sb.append( bestPoints[0]);
        for ( int i = 1; i < Math.min( bestPoints.length, 30); i++ )
        {
            sb.append( '\n');
            sb.append( bestPoints[i]);
        }
        return sb.toString();
    }
}

```

11.9 Multiple strategy approach

As we have seen most of the optimizing algorithms described so far have some limitation:

- Hill climbing algorithms may get into trouble far from the optimum and may get caught into a local optimum. This is exemplified in figure 11.7.
- The simplex algorithm may get caught into a local optimum and does not converge well near the optimum.
- Genetic algorithms do not have a clear convergence criteria.

After reading the above summary of the pro and cons of each algorithm, the reader may have already come to the conclusion that mixing the three algorithms together can make a very efficient strategy to find the optimum of a wide variety of functions.

One can start with a genetic optimizer for a sufficient number of iterations. This should ensure that the best points found at the end of the search does not lie too far from the absolute optimum. Then, one can use the simplex algorithm to get rapidly near the optimum. The final location of the optimum is obtained using a hill climbing optimizer.

11.9.1 Multiple strategy approach — General implementation

This multiple strategy approach, inspired from the program MINUIT, has been adapted to the use of the algorithms discussed here. The class `MultiVariableGeneralOptimizer` combines the three algorithms: genetic, simplex and hill climbing, in this order. We could have made it a subclass of `Object`, but we decided to reuse all the management provided by the abstract optimizer class discussed in section 11.3.1. Therefore, our general purpose optimizer is a subclass of the abstract optimizer class although it does not really use the framework of an iterative process. We only need one additional instance variable: the range used to construct the hypercube search space for the vector genetic chromosome manager. A corresponding setting method is provided: `setRange`.




Figure 11.7: Compared behavior of hill climbing and random based algorithms.

The method `initializeIterations` performs search using the genetic algorithm as an option and, then, the simplex algorithm. Since the genetic algorithm require a great deal of function evaluate — due to its stochastic nature — it is a good idea to give the user the choice of by-passing the use of the genetic algorithm. If no range has been defined, only the simplex algorithm is used from the supplied initial value. Otherwise a search is made with the genetic algorithm using the initial value and the range to define the search space. Then the simplex algorithm is started from the best point found by the genetic algorithm. The precision for the simplex search is set to the square root of the precision for the final search. Less precision is required for this step because the final search will give a better precision.

The method `evaluateIteration` performs the hill climbing algorithm and returns its precision. As the desired precision of the hill climbing algorithm is set to that of the general purpose optimizer. As a consequence, there will only be a single iteration.

Listing 11.23 shows the implementation in Smalltalk. Listing 11.24 gives the code for the Java implementation. At this point we shall abstain from commenting the code as the reader should have no more need for such thing. . . Hopefully!

Figure 11.1 with the box `MultiVariableGeneralOptimizer` grayed.

Listing 11.23 Smalltalk implementation of a general optimizer

```
Class      DhbMultiVariableGeneralOptimizer
Subclass of DhbFunctionOptimizer
```

Instance methods

computeInitialValues

```
self range notNil
  ifTrue: [ self performGeneticOptimization].
self performSimplexOptimization.
```

evaluateIteration

```
| optimizer |
optimizer := DhbHillClimbingOptimizer forOptimizer: self.
optimizer desiredPrecision: desiredPrecision;
           maximumIterations: maximumIterations.
result := optimizer evaluate.
^optimizer precision
```

origin

```
^result
```

origin: anArrayOfVector

```
result := anArrayOfVector.
```

performGeneticOptimization

```
| optimizer manager |
optimizer := DhbGeneticOptimizer forOptimizer: functionBlock.
manager := DhbVectorChromosomeManager new: 100 mutation: 0.1
           crossover: 0.1.
manager origin: self origin asVector; range: self range asVector.
optimizer chromosomeManager: manager.
result := optimizer evaluate.
```

performSimplexOptimization

```
| optimizer manager |
optimizer := DhbSimplexOptimizer forOptimizer: self.
optimizer desiredPrecision: desiredPrecision sqrt;
           maximumIterations: maximumIterations;
           initialValue: result asVector.
result := optimizer evaluate.
```

range

```

    ^self bestPoints

range: anArrayOfVector
    bestPoints := anArrayOfVector.

```

Figure 11.2 with the box
MultiVariableGeneralOptimizer
grayed.

Listing 11.24 Java implementation of a general optimizer

```

package DhbOptimizing;

/**
 * Multi-strategy optimizer of many-variable functions.
 *
 * @author Didier H. Besset
 */
public class MultiVariableGeneralOptimizer extends MultiVariableOptimizer
{
    /**
     * Initial range for random search.
     */
    protected double[] range;
    /**
     * Constructor method.
     * @param func DhbInterfaces.ManyVariableFunction
     * @param pointCreator DhbOptimizing.OptimizingPointFactory
     * @param initialValue double[]
     */
    public MultiVariableGeneralOptimizer(DhbInterfaces.ManyVariableFunction func,
        OptimizingPointFactory pointCreator, double[] initialValue)
    {
        super(func, pointCreator, initialValue);
    }
    public double evaluateIteration()
    {
        HillClimbingOptimizer finder = new HillClimbingOptimizer( f, pointFactory,
                                                                    result);
        finder.setDesiredPrecision( getDesiredPrecision());
        finder.setMaximumIterations( getMaximumIterations());
        finder.evaluate();
        result = finder.getResult();
        return finder.getPrecision();
    }
    public void initializeIterations()

```

```
{
    if ( range != null )
        performGeneticOptimization();
        performSimplexOptimization();
}
private void performGeneticOptimization()
{
    VectorChromosomeManager manager = new VectorChromosomeManager();
    manager.setRange( range);
    manager.setOrigin( result);
    VectorGeneticOptimizer finder = new VectorGeneticOptimizer( f, pointFactory, manager);
    finder.evaluate();
    result = finder.getResult();
}
private void performSimplexOptimization()
{
    SimplexOptimizer finder = new SimplexOptimizer( f, pointFactory, result);
    finder.setDesiredPrecision( Math.sqrt( getDesiredPrecision()));
    finder.setMaximumIterations( getMaximumIterations());
    finder.evaluate();
    result = finder.getResult();
}
/**
 * @param x double    component of the origin of the hypercube
 *                  constraining the domain of definition of the function
 */
public void setOrigin( double[] x)
{
    result = x;
}
/**
 * @param x double    components of the lengths of the hypercube
 *                  constraining the domain of definition of the function
 */
public void setRange( double[] x)
{
    range = x;
}
}
```

Chapter 12

Data mining

Creusez, fouillez, bêchez, ne laissez nulle place
Où la main ne passe et repasse,¹
Jean de La Fontaine

Data mining is a catchy buzz-word of recent introduction covering activities formerly known as data analysis. The problem is akin to what we already have seen in chapter 10. In the case of data mining, however, the emphasis is put on large data sets. Large must be understood in two ways: first, each data point is actually made of a large number of measurements; second, the number of data points is large, even huge. The expression data mining was coined when large corporate databases became common place. The original reason for the presence of these databases was the day to day dealing with the business events. A few people started realizing that these databases are containing huge amount of information, mostly of statistical nature, about the type of business. That information was just waiting to be mined just like the Mother Lode waited for the coming of the 49ers. Hence the term data mining.

Figure 12.1 shows the classes described in this chapter. There are two aspects to the data mining activity: one is preparing the data in order to render them suitable for the processing; the second consists of extracting the information, that is, the processing of the data. Depending on the problems and on the technique used during the second step, the first step is not always needed. In any case, the preparation of the data is usually very specific to the type of data to be analyzed. Therefore one can only make very general statements about this problem: one must watch for rounding errors; one must adjust the scale of data having widely different ranges; one must check the statistical validity of the data sample; etc. We shall say no more about this first aspect of data mining, but we wanted to warn the reader about this important issue.

Finally, data mining differs from estimation in that, in many cases the type of information to be extracted is not really known in advance. Data mining tries to identify trends, common properties and correlation between the data. The goal of data mining is usually to reduce large sample to much smaller manageable sets, which can be efficiently

¹Dig, search, excavate, do not leave a place where your hands did not go once or more.

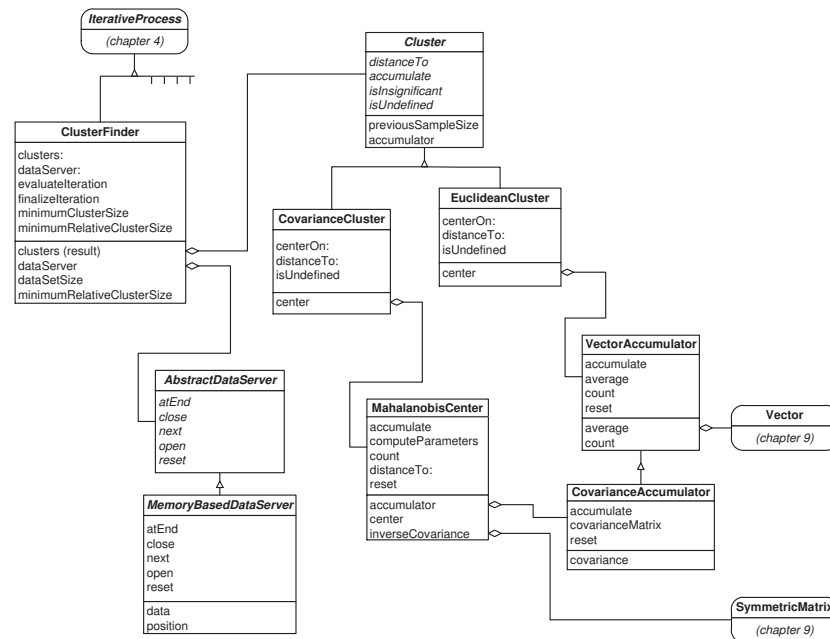


Figure 12.1: Classes used in data mining

targeted. One example, is the selection of a sample of customers suitable to respond to a new product offered by a bank or an insurance company. Mailing operations are costly; thus, any reduction of the mailing target set with a significant improvement of the probability of response can bring a significant saving. Another example in the medical domain is the scanning for certain type of cancer, which are expensive². If one can identify a population with high risk of cancer, the scanning only needs to be done for that population, thus keeping the cost low.

A good collection of articles about the techniques exposed in this chapter can be found in [Achtley & Bryant].

12.1 Data server

Figure 12.1 with the boxes `AbstractDataServer` and `MemoryBasedDataServer` grayed.

As we have said in the introduction, data mining means handling large amounts of data, most likely more than the computer memory can hold. Thus, we need an object to handle these data for all objects implementing data mining techniques. The data server object needs to implement five functionalities:

1. opening the physical support of the data,
2. getting the next data item,

²In medical domain, expensive is not necessarily a matter of money. It can mean high risk for the patient or the examination is regarded as too intrusive by the patient.

3. checking whether more data items are available,
4. repositioning the data stream at its beginning and
5. closing the physical support of the data.

Depending on the problem at hand, the responsibility of a data server can extend beyond the simple task of handling data. In particular, it could be charged of performing the data preparation step mentioned in the introduction. In our implementation, we give two classes. One is an abstract class from which all data servers used by the data mining objects described in this chapter must derive. The data item returned by the method returning the next item is a vector object whose components are the data corresponding to one measurement. The second class of data server is a concrete class implementing the data server on a collection or array kept in the computer's memory. Such server is handy for making tests.

Note: Examples of use of data server are given in the other sections; no code example are given here.

12.1.1 Data server — Smalltalk implementation

Listing 12.1 shows the implementation of the abstract data server in Smalltalk. The implementation of the concrete class is shown in listing 12.2.

Our implementation uses the same methods used by the hierarchy of the class Stream.

Listing 12.1 Smalltalk abstract data server

```
Class      DhbAbstractDataServer
Subclass of Object
```

```
    Instance methods
```

```
atEnd
```

```
    self subclassResponsibility
```

```
close
```

```
next
```

```
    self subclassResponsibility
```

```
open
```

```
    self subclassResponsibility
```


reset

self subclassResponsibility

Listing 12.2 Smalltalk memory based data server

Class **DhbMemoryBasedDataServer**
Subclass of DhbAbstractDataServer
Instance variable names: data position

Instance methods

atEnd

^data size < position

data: anOrderedCollection

data := anOrderedCollection.
self reset.

dimension

^data first size

next

| answer |
answer := data at: position.
position := position + 1.
^answer

open

self reset

reset

position := 1.

12.1.2 Data server — Java implementation

Listing 12.3 shows the implementation of the abstract data server in Java. The implementation of the concrete class is shown in listing 12.4.

Our implementation follows the naming used in the `java.io` package. There is no method checking for the availability of more data; instead, the exception `EOFException` is thrown by the `read` method when reading beyond the end of the stream occurs .

Listing 12.3 Java abstract data server

package DhbDataMining;

```
import DhbMatrixAlgebra.DhbVector;
/**
 * Abstract data server for data mining.
 *
 * @author Didier H. Besset
 */
public abstract class AbstractDataServer
{
    /**
     * Constructor method.
     */
    public AbstractDataServer() {
        super();
    }
    /**
     * Closes the stream of data.
     */
    public abstract void close();
    /**
     * Opens the stream of data.
     */
    public abstract void open();
    /**
     * @return DhbVector next data point found on the stream
     * @exception java.io.EOFException when no more data point can be found.
     */
    public abstract DhbVector read() throws java.io.EOFException;
    /**
     * Rewind the stream of data.
     */
    public abstract void reset();
}
```

Listing 12.4 Java memory based data server

```
package DhbDataMining;

import DhbMatrixAlgebra.DhbVector;
/**
 * Data server containing data in memory for simulation purposes.
 *
 * @author Didier H. Besset
 */
public class MemoryBasedDataServer extends AbstractDataServer
{
    private int index;
    private DhbVector[] dataPoints;
    /**
     * Constructor method (for internal use only)
     */
    protected MemoryBasedDataServer()
    {
        super();
    }
    /**
     * @param points DhbVector[] supplied data points
     * (must not be changed after creation)
     */
    public MemoryBasedDataServer( DhbVector[] points)
    {
        dataPoints = points;
    }
    /**
     * Nothing to do
     */
    public void close()
    {
    }
    /**
     * Nothing to do
     */
    public void open()
    {
    }
    /**
     * @return DhbMatrixAlgebra.DhbVector next data point
     * @exception java.io.EOFException no more data.
     */
    public DhbVector read() throws java.io.EOFException
```

```

{
    if( index >= dataPoints.length )
        throw new java.io.EOFException();
    return dataPoints[index++];
}
/**
 * Data index is reset
 */
public void reset()
{
    index = 0;
}
}

```

12.2 Covariance and covariance matrix

When one deals with two or more random variables an important question to ask is whether or not the two variables are dependent from each other.

For example, if one collects the prices of homes and the incomes of the home owners, one will find that inexpensive homes are mostly owned by low income families, mostly but not always. It is said that the price of a home is correlated with the income of the home owner. As soon as one deals with more than two variables things stop being clear cut. Correlations become hard to identify especially because of these mostly but not always cases. Therefore, one must find a way to express mathematically how much two random variables are correlated.

Let x_1, \dots, x_m be several random variable. They can be considered as the components of a m -dimensional (random) vector \mathbf{x} . The probability density function of the vector \mathbf{x} is denoted $P(\mathbf{x})$; it measures the probability of observing a vector within the differential volume element located at \mathbf{x} . The average of the vector \mathbf{x} is defined in a way similar to the case of a single random variable. The i^{th} component of the average is defined by

$$\mu_i = \int \dots \int x_i P(\mathbf{x}) dx_1 \dots dx_m. \quad (12.1)$$

The covariance matrix of the random vector \mathbf{x} gives a measure of the correlations between the components of the vector \mathbf{x} . The components of the covariance matrix are defined by

$$\varrho_{ij} = \int \dots \int (x_i - \mu_i)(x_j - \mu_j) P(\mathbf{x}) dx_1 \dots dx_m. \quad (12.2)$$

As one can see the covariance matrix is a symmetric matrix. It is also positive definite. Furthermore ϱ_{ii} is the variance of the i^{th} component of the random vector.

Note: The error matrix of a least square or maximum likelihood fit — discussed in chapter 10 — is the covariance matrix of the fit parameters.

If two components are independent, their covariance — that is, the corresponding element of the covariance matrix — is zero. The inverse is not true, however. For example, consider a 2-dimensional vector with components (zz^2) where z is a random variable. If z is distributed according to a symmetric distribution, the covariance between the two components of the vector is zero. Yet, the components are 100% dependent from each other by construction.

The correlation coefficient between components i and j of the vector \mathbf{x} is then defined by

$$\rho_{ij} = \frac{\varrho_{ij}}{\sigma_i \sigma_j}, \quad (12.3)$$

where $\sigma_i = \sqrt{\varrho_{ii}}$ is the standard deviation of the i^{th} component of the vector \mathbf{x} . By definition, the correlation coefficient is comprised between -1 and 1 . If the absolute value of a correlation coefficient is close to 1 then one can assert that the two corresponding components are indeed correlated.

If the random vector is determined experimentally, one calculates the estimated covariance with the following statistics

$$\text{cov}(x_i, x_j) = \frac{1}{n} \sum_{k=1}^n (x_{i,k} - \mu_i)(x_{j,k} - \mu_j), \quad (12.4)$$

where $x_{i,k}$ is the i^{th} component of the k^{th} measurement of the vector \mathbf{x} . Similarly the estimated correlation coefficient of the corresponding components is defined as

$$\text{cor}(x_i, x_j) = \frac{\text{cov}(x_i, x_j)}{s_i s_j}, \quad (12.5)$$

where s_i is the estimated standard deviation of the i^{th} component.

Like for the central moment, there is a way to compute the components of the covariance matrix while they are accumulated. If $\text{cov}_n(x_i, x_j)$ denotes the estimated covariance over n measurements, one has:

$$\text{cov}_{n+1}(x_i, x_j) = \frac{n}{n+1} \text{cov}_n(x_i, x_j) + n \Delta_{i,n+1} \Delta_{j,n+1}, \quad (12.6)$$

where $\Delta_{x,n+1}$ and $\Delta_{y,n+1}$ are the corrections to the averages of each variable defined in equation 9.12 of section 9.2. The derivation of equation 12.6 is given in appendix E.2.

Using covariance information

A covariance matrix contains statistical information about the set of measurements over which it has been determined. There are several ways of using this information.

The first approach uses the covariance matrix directly. The best example is the analysis known as the shopping cart analysis [Berry & Linoff]. For example, one can observe that consumers buying cereals are buying low fat milk. This can give useful information on how to target special sales efficiently. Application working in this mode can use the code of sections 12.2.2 or 12.2.3 as is.

Another approach is to use the statistical information contained in a covariance matrix to process data coming from measurements which were not used to determine the covariance matrix. In the rest of this chapter we shall call the set of measurements, which is used to determine the covariance matrix, the calibrating set. In this second mode of using a covariance matrix, measurements are compared or evaluated against those of the calibrating set. It is clear that the quality of the information contained in the covariance matrix depends on the quality of the calibrating set. We shall assume that this is always the case. Techniques working according to this second mode are described in sections 12.4, 12.5 and 12.7.

12.2.1 Covariance matrix — General implementation

The object in charge of computing the covariance matrix of a series of measurements is implemented as for central moments. Because we shall need to only compute the average of a vector, the implementation is spread over two classes, one being the subclass of the other for efficient reuse.

The class `VectorAccumulator` has two instance variables:

`count` counts the number of vectors accumulated in the object so far;

`average` keeps the average of the accumulated vector;

The class `CovarianceAccumulator` is a subclass of the class `VectorAccumulator`. It has one additional instance variable:

`covariance` accumulates the components of the covariance matrix; for efficiency reason, only the lower half of the matrix is computed since it is symmetric.

The topmost class implements equation 9.12 in the method `accumulate`. The subclass overloads this method to implement equation 12.6.

12.2.2 Covariance matrix — Smalltalk implementation

Listing 12.5 shows the implementation of the accumulation of a vector in Smalltalk. Listing 12.6 shows the implementation of the accumulation of the covariance matrix. The following code example shows how to accumulate the average of a series of vectors read from a data stream.

Code example 12.1

```
| accumulator valueStream average |
accumulator := DhbVectorAccumulator new.
valueStream open.
[ valueStream atEnd]
  whileFalse:[ accumulator accumulate: valueStream next].
valueStream close.
average := accumulator average.
```

Figure 12.1 with the boxes `VectorAccumulator` and `CovarianceAccumulator` grayed.

The reader can see that this example is totally equivalent to the code example 9.1. Here the method next of the data stream must return a vector instead of a number; all vectors must have the same dimension. The returned average is a vector of the same dimension.

The next code example shows how to accumulate the both average and covariance matrix. The little differences with the preceding example should be self explanatory to the reader.

Code example 12.2

```
| accumulator valueStream average covarianceMatrix |
accumulator := DhbCovarianceAccumulator new.
valueStream open.
[ valueStream atEnd]
  whileFalse:[ accumulator accumulate: valueStream next].
valueStream close.
average := accumulator average.
covarianceMatrix := accumulator covarianceMatrix.
```

The method accumulate of class DhbVectorAccumulator answers the corrections to each component of the average vector. This allows the class DhbCovarianceAccumulator to reuse the results of this method. In class DhbVectorAccumulator, vector operations are used. The method accumulate of class DhbCovarianceAccumulator works with indices because one only computes the lower half of the matrix.

Listing 12.5 Smalltalk implementation of vector average

```
Class                DhbVectorAccumulator
Subclass of          Object

Instance variable names:  count average

Class methods

new: anInteger

^self new initialize: anInteger

Instance methods

accumulate: aVectorOrArray

| delta |
count := count + 1.
delta := average - aVectorOrArray asVector scaleBy: 1 / count.
average accumulateNegated: delta.
^delta

average
```

```

^average

count

^count

initialize: anInteger
    average := DhbVector new: anInteger.
    self reset.
    ^self

printOn: aStream
    super printOn: aStream.
    aStream space.
    count printOn: aStream.
    aStream space.
    average printOn: aStream.

reset
    count := 0.
    average atAllPut: 0.

```

Listing 12.6 Smalltalk implementation of covariance matrix

Class	DhbCovarianceAccumulator
Subclass of	DhbVectorAccumulator
Instance variable names:	covariance

```

    Instance methods
accumulate: anArray

    | delta count1 r |
    count1 := count.
    delta := super accumulate: anArray.
    r := count1 / count.
    1 to: delta size
        do: [ :n |
            1 to: n do:

```



```

[ :m |
  ( covariance at: n) at: m put: ( count1 * ( delta
at: n) * ( delta at: m) + ( r * ( ( covariance at: n) at: m))).
].

```

covarianceMatrix

```

| rows n |
n := 0.
rows := covariance collect:
  [ :row | n := n + 1. row, ( ( ( n + 1) to: covariance
size) collect: [ :m | ( covariance at: m) at: n ])].
^DhbSymmetricMatrix rows: rows

```

initialize: anInteger

```

covariance := ( ( 1 to: anInteger) collect: [ :n | DhbVector new:
n]) asVector.
^super initialize: anInteger

```

reset

```

super reset.
covariance do: [ :each | each atAllPut: 0].

```

12.2.3 Covariance matrix — Java implementation

Listing 12.5 shows the implementation of the accumulation of a vector in Smalltalk. Listing 12.6 shows the implementation of the accumulation of the covariance matrix. The following code example shows how to accumulate the average of a series of vectors read from a data stream.

Code example 12.3

<Creating an instance of a concrete subclass of AbstractDataServer into the variable dataServer>

```

VectorAccumulator accumulator = new VectorAccumulator(dimension);
dataServer.open();
try {
  while (true) accumulator.accumulate( server.read());
} catch ( java.io.EOFException e) { };
DhbVector averages = accumulator.averageVector();

```

This example is very similar to code example 9.2. Here the data to accumulate are prepared as an array of vectors. The next code example shows how to accumulate the both average and covariance matrix. The little differences with the preceding example should be self explanatory to the reader.

Code example 12.4

```
VectorAccumulator accumulator = new VectorAccumulator(dimension);
dataServer.open();
try {
    while (true) accumulator.accumulate( server.read());
} catch ( java.io.EOFException e) { };
DhbVector averages = accumulator.averageVector();
SymmetricMatrix covariance = accumulator.covarianceMatrix();
```

In our implementation the method `accumulate` has been implemented without any exception. This departure from the recommended practice is justified by the fact that the handling of exception — even when they do not occur — increases the execution of a method. To optimize the loop, the computation of the average is duplicated in the method `accumulate` of each class.

Listing 12.7 Java implementation of vector average

```
package DhbDataMining;

import DhbMatrixAlgebra.DhbVector;
/**
 * Statistical average for vectors
 *
 * @author Didier H. Besset
 */
public class VectorAccumulator
{
    protected long count = 0;
    protected double[] average;
    /**
     * Default constructor method.
     */
    public VectorAccumulator(int n)
    {
        this(n,1);
        reset();
    }
    /**
     * Constructor method.
     * @param n int
     * @param dummy int
```

```

*/
protected VectorAccumulator(int n, int dummy)
{
    average = new double[n];
}
/**
 * @param v double[] values to accumulate in the receiver
 */
public void accumulate( double[] v)
{
    count += 1;
    for ( int i = 0; i < average.length; i++ )
        average[i] -= (average[i] - v[i]) / count;
}
/**
 * @param v DhbVector vector of values to accumulate in the receiver
 */
public void accumulate( DhbVector v)
{
    accumulate( v.toComponents());
}
/**
 * @return DhbVector vector containing the average
 */
public DhbVector averageVector()
{
    try { return new DhbVector( average);}
    catch( NegativeArraySizeException e){ return null;}
}
/**
 * @return long number of accumulated data points.
 */
public long getCount()
{
    return count;
}
/**
 * @param sb java.lang.StringBuffer
 */
protected void printOn( StringBuffer sb)
{
    sb.append("Counts: "+count);
    char separator = '\n';
    for ( int i = 0; i < average.length; i++)
    {
        sb.append( separator);
    }
}

```

```

        sb.append( average[i]);
        separator = ' ';
    }
}
public void reset()
{
    count = 0;
    for ( int i = 0; i < average.length; i++ )
        average[i] = 0;
}
/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    printOn( sb);
    return sb.toString();
}
}

```

Listing 12.8 Java implementation of covariance matrix

```

package DhbDataMining;

import DhbMatrixAlgebra.SymmetricMatrix;
import DhbMatrixAlgebra.DhbIllegalDimension;
import DhbMatrixAlgebra.DhbNonSymmetricComponents;
/**
 * Statistical average and covariance for vectors
 *
 * @author Didier H. Besset
 */
public class CovarianceAccumulator extends VectorAccumulator
{
    private double[][] covariance;
/**
 * Constructor method
 * @param n int
 */
public CovarianceAccumulator(int n)
{
    super(n,1);
    covariance = new double[n][n];
}
}

```

```

        reset();
    }
    /**
     * @param v DhbVector    vector to accumulate in the receiver
     */
    public void accumulate( double[] v)
    {
        long n = count;
        count += 1;
        double[] deltas = new double[average.length];
        int j;
        double r = (double) n / (double) count;
        for ( int i = 0; i < average.length; i++ )
        {
            deltas[i] = ( average[i] - v[i] ) / count;
            average[i] -= deltas[i];
            for ( j = 0; j <= i; j++ )
                covariance[i][j] = r * covariance[i][j] + n * deltas[i] * deltas[j];
        }
    }
    /**
     * @return double
     * @param n int
     * @param m int
     */
    public double correlationCoefficient( int n, int m)
    {
        return covariance[n][m] / Math.sqrt(covariance[n][n] * covariance[m][m]);
    }
    /**
     * @return SymmetricMatrix    covariance matrix
     */
    public SymmetricMatrix covarianceMatrix()
    {
        double[][] components = new double[average.length][average.length];
        int j;
        for ( int i = 0; i < average.length; i++ )
        {
            for ( j = 0; j <= i; j++ )
            {
                components[i][j] = covariance[i][j];
                components[j][i] = components[i][j];
            }
        }
        try { return SymmetricMatrix.fromComponents(components);}
        catch (DhbNonSymmetricComponents e) { return null;}
    }

```

```

        catch (DhbIllegalDimension e) { return null;}
    }
    /**
     * @param sb java.lang.StringBuffer
     */
    protected void printOn( StringBuffer sb)
    {
        super.printOn( sb);
        for ( int i = 0; i < average.length; i++)
        {
            char separator = '\n';
            for ( int j = 0; j <= i; j++)
            {
                sb.append( separator);
                sb.append( covariance[i][j]);
                separator = ' ';
            }
        }
    }
    public void reset()
    {
        super.reset();
        int j;
        for ( int i = 0; i < average.length; i++ )
        {
            for ( j = 0; j <= i; j++ )
                covariance[i][j] = 0;
        }
    }
    /**
     * @return double
     * @param n int
     * @param m int
     */
    public double standardDeviation( int n)
    {
        return Math.sqrt(variance(n));
    }
    /**
     * @return double
     * @param n int
     */
    public double variance( int n)
    {
        return covariance[n][n] * count / (count - 1);
    }

```

}

12.3 Multidimensional probability distribution

To get a feeling of what the covariance matrix represents, let us now consider a vector \mathbf{y} whose components are independent random variables distributed according to a normal distribution. The probability density function is given by

$$P(\mathbf{y}) = \prod_{i=1}^m \frac{1}{\sqrt{2\pi\sigma_i^2}} e^{-\frac{(y_i - \mu_i)^2}{2\sigma_i^2}}. \quad (12.7)$$

In this case, the covariance matrix of the vector \mathbf{y} is a diagonal matrix $\tilde{\mathbf{V}}$, whose diagonal elements are the variance of the vector's components. If $\tilde{\mathbf{C}}$ is the inverse of the matrix $\tilde{\mathbf{V}}$, equation 12.7 can be rewritten as

$$P(\mathbf{y}) = \sqrt{\frac{\det \tilde{\mathbf{C}}}{(2\pi)^m}} e^{-\frac{1}{2}(\mathbf{y} - \bar{\mathbf{y}})^T \tilde{\mathbf{C}}(\mathbf{y} - \bar{\mathbf{y}})}, \quad (12.8)$$

where $\bar{\mathbf{y}}$ is the vector whose components are μ_1, \dots, μ_m . Let us now consider a change of coordinates $\mathbf{x} = \mathbf{O}\mathbf{y}$, where \mathbf{O} is an orthogonal matrix. We have already met such transformations in section 8.7. Because the matrix is orthogonal, the differential volume element is invariant under the change of coordinates. Thus, the probability density function of the vector \mathbf{x} is

$$P(\mathbf{x}) = \sqrt{\frac{\det \mathbf{C}}{(2\pi)^m}} e^{-\frac{1}{2}(\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{C}(\mathbf{x} - \bar{\mathbf{x}})}, \quad (12.9)$$

where the matrix \mathbf{C} is equal to $\mathbf{O}^T \tilde{\mathbf{C}} \mathbf{O}$. The vector $\bar{\mathbf{x}}$ is simply³ equal to $\mathbf{O}\bar{\mathbf{y}}$. The covariance matrix, \mathbf{V} , of the vector \mathbf{x} is then equal to $\mathbf{O}^T \tilde{\mathbf{V}} \mathbf{O}$. It is also the inverse of the matrix \mathbf{C} . Thus, equation 12.9 can be rewritten as

$$P(\mathbf{x}) = \frac{1}{\sqrt{(2\pi)^m \det \mathbf{V}}} e^{-\frac{1}{2}(\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{V}^{-1}(\mathbf{x} - \bar{\mathbf{x}})}. \quad (12.10)$$

In the case of a normal distribution in a multi-dimensional space, the covariance matrix plays the same role as the variance in one dimension.

12.4 Covariance data reduction

Reversing the derivation of the preceding section one can see that the eigenvalues of the covariance matrix of the vector \mathbf{x} correspond to the variances of a series of independent⁴ random variables y_1, \dots, y_m . Since the covariance matrix is symmetric

³All this is a consequence of the linear property of the expectation value operator.

⁴Not necessarily normally distributed!

these eigenvalues as well as the matrix \mathbf{O} describing the change of coordinates can be obtained using Jacobi's algorithm described in section 8.7.

If some eigenvalues are much larger than the others, one can state that the information brought by the corresponding variables brings little information to the problem. Thus, one can omit the corresponding variable from the rest of the analysis.

Let $\sigma_1^2, \dots, \sigma_m^2$ be the eigenvalues of the covariance matrix such that $\sigma_i^2 < \sigma_j^2$ for $i < j$. Let us assume that there exist an index k such that $\sigma_k^2 \gg \sigma_{k-1}^2$. Then, The rest of the data analysis can be made with a vector with components y_k, \dots, y_m where the vector \mathbf{y} is defined by $\mathbf{y} = \mathbf{O}^T \mathbf{x}$.

This reduction technique has been used successfully in high energy physics⁵ under the name principal component analysis. The data reduction allows to extract the relevant parameters of a complex particle detector to facilitate the quick extraction of the physical data — momentum and energy of the particle — from the observed data.

This kind of data reduction can be implemented within the framework of the data server described in section 12.1. The concrete implementation of such a server is straight forward. All needed objects — covariance accumulation, eigenvalues of a symmetric matrix, vector manipulation — have been discussed in different chapters. The rest of the implementation is specific to the problem at hand and can therefore not be treated on a general basis.

12.5 Mahalanobis distance

Mahalanobis, an Indian statistician, introduced this distance in the 30's when working on anthropometric statistics. A paper by Mahalanobis himself can be found in [Achtley & Bryant]. Readers interested by the historical dimension of the Mahalanobis distance can consult the paper by Das Gupta⁶.

By definition, the exponent of equation 12.7 is distributed according to a χ^2 distribution with m degrees of freedom. This exponent remains invariant under the change of coordinates discussed in section 12.3. Thus, the exponent of equation 12.10 is also distributed according to a χ^2 distribution with m degrees of freedom, where m is the dimension of the vector x . The Mahalanobis distance, d_M , is defined as the square root of the exponent, up to a factor $\sqrt{2}$. We have

$$d_M^2 = (\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{V}^{-1} (\mathbf{x} - \bar{\mathbf{x}}). \quad (12.11)$$

The Mahalanobis distance is a distance using the inverse of the covariance matrix as the metric. It is a distance in the geometrical sense because the covariance matrix as well as its inverse are positive definite matrices. The metric defined by the covariance matrix provides a normalization of the data relative to their spread.

⁵H.Wind, Pattern recognition by principal component analysis of border regions, Proceedings 3th topical meeting on multi-dimensional analysis of high energy data, Nijmegen, 8-11 March 1978, W. Kittel, University of Nijmegen, 1978, pp. 99-106.

⁶Somesh Das Gupta, The evolution of the D^2 -statistics of Mahalanobis, Indian J. Pure Appl. Math., 26(1995), no. 6, 485-501.

The Mahalanobis distance — or its square — can be used to measure how close an object is from another when these objects can be characterized by a series of numerical measurements. Using the Mahalanobis distance is done as follows

1. the covariance matrix of the measured quantities, \mathbf{V} , is determined over a calibrating set;
2. one compute the inverse of the covariance matrix, \mathbf{V}^{-1} ;
3. the distance of a new object to the calibrating set is estimated using equation 12.11; if the distance is smaller than a given threshold value, the new object is considered as belonging to the same

One interesting property of the Mahalanobis distance is that it is normalized. Thus, it is not necessary to normalize the data provided rounding errors in inverting the covariance matrix are kept under control. If the data are roughly distributed according to a normal distribution the threshold for accepting whether or not an object belongs to the calibrating set can be determined from the χ^2 distribution.

Examples of use

The Mahalanobis distance can be applied in all problems where measurements must be classified.

A good example is the detection of coins in a vending machine. When a coin is inserted into the machine, a series of sensors gives several measurements, between a handful and a dozen. The detector can be calibrated using a set of good coins forming a calibration set. The coin detector can differentiate good coins from the fake coins using the Mahalanobis distance computed on the covariance matrix of the calibration set. Figure 12.2 shows an example of such data⁷. The light gray histogram is the distribution of the Mahalanobis distance of the good coins; the dark gray histogram that of the fake coins. The dotted line is the χ^2 -distribution corresponding to the degrees of freedom of the problem; in other words, the distribution was not fitted. The reader can see that the curve of the χ^2 -distribution reproduces the distribution of the experimental data. Figure 12.2 also shows the power of separation achieved between good and fake coins.

Another field of application is the determination of cancer cells from a biopsy. Parameters of cells — size and darkness of nucleus, granulation of the membrane — can be measured automatically and expressed in numbers. The covariance matrix can be determined either using measurements of healthy cells or using measurements of malignant cells. Identification of cancerous cells can be automatized using the Mahalanobis distance.

12.5.1 Mahalanobis distance — General implementation

The final goal of the object implementing the Mahalanobis distance is to compute the square Mahalanobis distance as defined in equation 12.11. This is achieved with the

Figure 12.1 with the box MahalanobisCenter grayed.

⁷Data are reproduced with permission; the owner of the data wants to remain anonymous, however.




Figure 12.2: Using the Mahalanobis distance to differentiate between good and fake coins.

method `distanceTo`. The inverse of the covariance matrix as well as the average vector \bar{x} are contained within the object. We have called the object a Mahalanobis center since it describes the center of the calibrating set.

To have a self-contained object, the Mahalanobis center is acting as a facade to the covariance accumulator of section 12.2 for the accumulation of measurements. The methods `accumulate` and `reset` are delegated to an instance variable holding a covariance accumulator.

The Mahalanobis center has the following variables

`accumulator` a covariance accumulator as described in section 12.2;

`center` the vector \bar{x} and

`inverseCovariance` the inverse of the covariance matrix, that is V^{-1} .

Our implementation is dictated by its future reuse in cluster analysis (c.f. section 12.7). There, we need to be able to accumulate measurements while using the result of a preceding accumulation. Thus, computation of the center and the inverse covariance matrix must be done explicitly with the method `computeParameters`.

There are two ways of creating a new instance. One is to specify the dimension of the vectors which will be accumulated into the object. The second supplies a vector as the tentative center. This mode is explained in section 12.7.

12.5.2 Mahalanobis distance — Smalltalk implementation

Listing 12.9 shows the implementation of a Mahalanobis center in Smalltalk. The following code example shows how to sort measurements using the Mahalanobis distance.

Code example 12.5

```

| center calibrationServer dataServer data threshold|
center := DhbMahalanobisCenter new: 5.

<The variable calibrationServer is setup to read measurements from the
calibrating set>

calibrationServer open.
[ calibrationServer atEnd]
  whileFalse: [ center accumulate: calibrationServer next].
calibrationServer close.
center computeParameters.

<The variable dataServer is setup to read the measurements to be sorted
between accepted and rejected; the variable threshold must also be determined
or given>

dataServer open.
[ dataServer atEnd]
  whileFalse: [ data := dataServer next.
    ( center distanceTo: data) > threshold
    ifTrue: [ self reject: data]
    ifFalse:[ self accept: data].
  ].
dataServer close.

```

The first line after the declaration creates a new instance of a Mahalanobis center for vectors of dimension 5. After setting up the server for the calibrating set data from the calibrating set are accumulated into the Mahalanobis center. At the end the parameters of the Mahalanobis center are computed. Then, the server for the other measurements is set up. The loop calculates the distance to the Mahalanobis center for each measurements. Our example supposes that the object executing this code has implemented two methods accept and reject to process accepted and rejected data.

Listing 12.9 Smalltalk Mahalanobis center

Class	DhbMahalanobisCenter
Subclass of	Object
Instance variable names:	center inverseCovariance accumulator
Class methods	
new: anInteger	
^self new initialize: anInteger	
onVector: aVector	

^self new center: aVector

Instance methods

accumulate: aVector

accumulator accumulate: aVector.

center: aVector

accumulator := DhbcovarianceAccumulator new: aVector size.

center := aVector.

inverseCovariance := DhbsymmetricMatrix identity: aVector size.

^self

computeParameters

center := accumulator average copy.

inverseCovariance := accumulator covarianceMatrix inverse.

count

^accumulator count

distanceTo: aVector

| delta |

delta := aVector - center.

^delta * inverseCovariance * delta

initialize: anInteger

accumulator := DhbcovarianceAccumulator new: anInteger.

^self

printOn: aStream

accumulator count printOn: aStream.

aStream nextPutAll: ' '.
center printOn: aStream.

reset

accumulator reset.

12.5.3 Mahalanobis distance — Java implementation

Listing 12.9 shows the implementation of a Mahalanobis center in Smalltalk. The following code example shows how to sort measurements using the Mahalanobis distance.

Code example 12.6

```
MahalanobisCenter center = new MahalanobisCenter( 5);
```

<The variable calibrationServer is setup to read measurements from the calibrating set>

```
calibrationServer.open();
try{ while( true)
    { center.accumulate( calibrationServer.next());}}
  catch( EOFException e){};
calibrationServer.close();
center.computeParameters();
```

<The variable dataServer is setup to read the measurements to be sorted between accepted and rejected; the double variable threshold must also be determined or given>

```
dataServer.open();
try {
  while (true)
  {
    DhbVector data = dataServer.next();
    if ( center.distanceTo( data) > threshold )
      reject(data);
    else
      accept(data);
  };}
  catch( EOFException e){};
dataServer.close();
```

The first line creates a new instance of a Mahalanobis center for vectors of dimension 5. After setting up the server for the calibrating set data from the calibrating set are accumulated into the Mahalanobis center. At the end the parameters of the Mahalanobis center are computed. Then, the server for the other measurements is set up. The loop calculates the distance to the Mahalanobis center for each measurements. Our example supposes that the object executing this code has implemented two methods accept and reject to process accepted and rejected data.

Listing 12.10 Java Mahalanobis center

```
package DhbDataMining;
```

```

import DhbMatrixAlgebra.DhbIllegalDimension;
import DhbMatrixAlgebra.DhbNonSymmetricComponents;
import DhbMatrixAlgebra.DhbVector;
import DhbMatrixAlgebra.SymmetricMatrix;
/**
 * This object is used to compute the Mahalanobis distance
 * to a set of data.
 *
 * @author Didier H. Besset
 */
public class MahalanobisCenter
{
    private DhbVector center = null;
    private SymmetricMatrix inverseCovariance = null;
    private CovarianceAccumulator accumulator;
    /**
     * Constructor method.
     * @param int dimension of the receiver
     */
    public MahalanobisCenter( int dimension)
    {
        accumulator = new CovarianceAccumulator( dimension);
    }
    /**
     * Constructor method.
     * @param DhbVector center of the receiver
     */
    public MahalanobisCenter( DhbVector v)
    {
        accumulator = new CovarianceAccumulator( v.dimension());
        center = v;
        inverseCovariance = SymmetricMatrix.identityMatrix(v.dimension());
    }
    /**
     * Accumulation is delegated to the covariance accumulator.
     * @param v DhbVector vector of values to accumulate in the receiver
     */
    public void accumulate( DhbVector v)
    {
        accumulator.accumulate(v);
    }
    /**
     * Computes the parameters of the receiver.
     */
    public void computeParameters()
    {

```

```

        center = accumulator.averageVector();
        inverseCovariance = (SymmetricMatrix)
            accumulator.covarianceMatrix().inverse();
    }
    /**
     * @return double    Mahalanobis distance of the data point from the
     *                    center of the receiver.
     * @param dataPoint DhbVector    data point
     */
    public double distanceTo(DhbVector dataPoint)
    {
        try {
            DhbVector v = dataPoint.subtract( center);
            return v.product( inverseCovariance.product(v));
        }
        catch (DhbIllegalDimension e) { return Double.NaN;}
    }
    /**
     * @return long    number of data points inside the receiver
     */
    public long getCount()
    {
        return accumulator.getCount();
    }
    /**
     * Keep the center and covariance matrix.
     */
    public void reset()
    {
        accumulator.reset();
    }
    /**
     * @return java.lang.String
     */
    public String toString() {
        return center.toString();
    }
}

```

12.6 Cluster analysis

Cluster analysis — also known as K -cluster — is a method to identify similarities between data. If the dimension of the data is less than or equal than 3, graphical data representation provides an easy way to identify data points having some similarity. For more than 3 measurements, the human brain is unable to clearly identify clustering.

Cluster analysis have been used successfully by the US army to define a new classification of uniform sizes and in banks [Berry & Linoff]. British Telecom has used cluster analysis to detect a phone fraud of large scale in the early 90's⁸. The K -cluster algorithm goes as follows [Berry & Linoff]:

1. Pick up a set of centers where possible clusters may exist;
2. place each data point into a cluster corresponding to the nearest center;
3. when all data points have been processed, compute the center of each cluster;
4. if the centers have changed, go back to point 2.

This algorithm nicely maps itself to the framework of successive approximations discussed in section 4.1. We now will investigate the steps of the algorithm in details.

Algorithm details

Picking up a set of centers corresponds to the box labeled Compute or choose initial object of figure 4.3. Since one is looking for unknown structure in the data there is little chance to make a good guess on the starting values. The most general approach is to pick up a few points at random from the existing data points to be the initial cluster's centers.

The next two steps correspond to the box labeled Compute next object of figure 4.3. Here lies the gist of the K -cluster algorithm. For each data point one first finds the cluster whose center is the nearest to the data point. What is the meaning of near? It depends on the problem. Let us just say at this point that one needs to have a way of expressing the distance between two data points. For the algorithm to converge the distance must be a distance in the geometrical sense of the term. In geometry, a distance is a numerical function of two vectors, $d(\mathbf{x}, \mathbf{y})$. For all vectors \mathbf{x} , \mathbf{y} and \mathbf{z} the following conditions must be met by the function

$$\begin{aligned} d(\mathbf{x}, \mathbf{y}) &\geq 0, \\ d(\mathbf{x}, \mathbf{y}) &= d(\mathbf{y}, \mathbf{x}), \\ d(\mathbf{x}, \mathbf{y}) &\leq d(\mathbf{x}, \mathbf{z}) + d(\mathbf{z}, \mathbf{y}). \end{aligned} \tag{12.12}$$

Furthermore, $d(\mathbf{x}, \mathbf{y}) = 0$ if and only if $\mathbf{x} = \mathbf{y}$. The simplest known distance function is the Euclidean distance expressed as

$$d_E(\mathbf{x}, \mathbf{y}) = \sqrt{(\mathbf{x} - \mathbf{y}) \cdot (\mathbf{x} - \mathbf{y})}. \tag{12.13}$$

The square root in equation 12.13 is required for the distance function to behave linearly on a one dimensional subspace. The Euclidean distance corresponds to the notion of distance in everyday life. In assessing the proximity of two points, the square root is not needed.

After all points have been processed and assigned to a cluster, the center of each cluster is obtained by taking the vector average of all data points belonging to that cluster.

⁸Private communication to the author.




Figure 12.3: Example of cluster algorithm

Then, one needs to determine whether the clusters have changed since the last iteration. In the case of the K -cluster algorithm it is sufficient to count the number of data points changing clusters at each iteration. When the same data point are assigned to the same clusters at the end of an iteration, the algorithm is completed. The precision used to stop the iterative process is an integer in this case.

Figure 12.3 shows how the algorithm works for a 2-dimensional case. Data points were generated randomly centered around 3 separated clusters. Three random points were selected as starting points: in this example the random choice was particularly unlucky since the three starting points all belong to the same cluster. Nevertheless, convergence was attained after 5 iterations⁹.

Fine points

This simple example is admittedly not representative. However, this is not because of the small dimension nor because of the well separated clusters. It is because we knew in advance the number of clusters, 3 in this case. If we had started with 5 initial clusters, we would have ended up with 5 clusters, according to the expression, garbage in, garbage out, well-known among programmers.

One can modify the original algorithm to prune insignificant clusters from the search. This additional step must be added between steps 2 and 3. How to characterize an insignificant cluster? This often depends on the problem at hand. One thing for sure is that clusters with 0 elements should be left out. Thus, one easy method to prune insignificant clusters is to place a limit of the number of data points contained in each cluster. That limit should not be too high, however, otherwise most clusters will never get a chance of accumulating a significant amount of data points during the first few

⁹The last iteration is not visible since the centers did not change.

iterations.

12.6.1 Cluster analysis — General implementation

Our implementation of the K -cluster algorithm uses two classes: `ClusterFinder` and `Cluster`.

Figure 12.1 with the boxes `ClusterFinder` and `Cluster` grayed.

The class `Cluster` describes the behavior of a cluster. It is an abstract class. Subclasses implements the various strategies needed by the algorithm: distance calculation and pruning of insignificant clusters. The abstract class has only one instance variable, `previousSampleSize`, keeping the count of data points accumulated in the previous iteration. A cluster must be able to return the number of data points contained in the cluster. The method `changes` gives the number of data points which have changed cluster since the last iteration. The method `isInsignificantIn` determines whether a cluster must be pruned from the process. The method `isUndefined` allows the identification of clusters whose centers have not yet been defined. This method is used by the cluster finder to initialize undefined clusters with the data points before starting the iterative process.

An instance of a subclass of cluster must implement the following methods to interact with the class `ClusterFinder`:

- `distanceTo` the argument of this method is a vector representing the data point; the returned value is the distance between the supplied vector and the center of the cluster; any subclass of `Cluster` can implement a suitable distance function as well as its own representation of the center;
- `accumulate` the argument of this method is a vector representing the data point; this method is called when the cluster finder has determined that the data point must be placed within this cluster;
- `changes` this method returns the number of data points which have been added to and removed from the cluster since the last iteration; the default implementation only calculates the difference between the number currently in the cluster and the number in the previous iteration; this simple approach works in practice¹⁰;
- `sampleSize` this method returns the number of data points actually accumulated into the cluster;
- `reset` this method calculates the position of the new center at the end of an iteration; the default implementation stores the size of the previous sample.

Note: A reader knowledgeable in patterns will think of using a Strategy pattern to implement the distance. I have tried this approach, but the resulting classes were much more complex for little gain. It is easier to implement subclasses of `Cluster`. Each subclass can not only implement the distance function, but can also choose how to accumulate data points: accumulation or storage.

¹⁰Of course, one can construct cases where this simple approach fails. Such cases, however, correspond to situations where data points are oscillating between clusters and, therefore, do not converge. I have never met such cases in practice.

The concrete class `EuclideanCluster` calculates the square of the Euclidean distance as defined in equation 12.13. Using Euclidean distance requires that the components of the data vector have comparable scales. Cluster analysis with Euclidean clusters requires data preparation in general.

The class `ClusterFinder` implements the algorithm itself. This class is a subclass of the class for iterative processes described in section 4.1. The result of the iterative process is an array of clusters. The class `ClusterFinder` needs the following instance variables:

`dataServer` a data server object as described in section 12.1; this object is used to iterate on the data points;

`dataSetSize` a counter to keep track of the number of data points; this instance variable is combined with the next to provide a first order pruning strategy;

`minimumRelativeClusterSize` the minimum relative size of a significant cluster; the minimum cluster size is computed at each iteration by taking the product of this variable with the variable `dataSetSize`.

The class `ClusterFinder` uses an instance of the data server to iterate over the data points. before starting the search the client application can assign the list of initial clusters (step 1 of the algorithm). By default, the minimum relative size of the cluster is set to 0. A convenience method allows creating instances for a given data server and an initial set of clusters.

The method `initializeIterations` scans all clusters and looks for undefined clusters. The center of each undefined cluster is assigned to a data point read from the data server. This assumes that the data points have been collected randomly.

The method `evaluateIteration` processes each data point: first, it finds the index of the cluster nearest to the data point; then, the data point is accumulated into that cluster. After processing the data points, the clusters are processed to compute the new position of their centers and insignificant clusters are removed from the search. These tasks are performed within a method named `collectChangesAndResetClusters`. This method returns the number of data points which have changed cluster. The determination whether or not a cluster is significant is delegated to each cluster with the method `isInsignificant`. Thus, any subclass of cluster can implement its own strategy. The argument of the method `isInsignificant` is the cluster finder to provide each cluster with global information if needed.

The method `finalizeIterations` just closes the data server.

12.6.2 Cluster analysis — Smalltalk implementation

Listing 12.11 shows the implementation of the K -cluster algorithm in Smalltalk. The following code example shows how to implement a search for clusters.

Code example 12.7

```
| dataServer finder clusters|
```

```
<The variable dataServer is setup to read measurements from the calibrating
set>
```

```
finder := DhbClusterFinder new: 5 server: dataServer

type: <a concrete subclass of Cluster>.

finder minimumRelativeClusterSize: 0.1.
clusters := finder evaluate.
```

After setting up a data server to read the data point, an instance of class `DhbClusterFinder` is created. the number of desired clusters is set to 5. The class of the clusters is specified. The next line sets the minimum relative size of each cluster to be kept during the iteration. Finally, the K -cluster algorithm is performed and clusters are retrieved from the finder object.

The abstract class `DhbCluster` is implemented with an instance variable accumulator. The cluster delegates the responsibility of accumulating the data points to this variable. It is assumed that the object in accumulator implements the interface defined by the vector accumulators described in section 12.2.2.

The class `DhbClusterFinder` can be created in two ways. An application can set the list of initial clusters and the data server using the methods `cluster` and `dataServer` respectively. The convenience class creation method `new:server:type:` allows to specify the initial number of clusters, the data server and the class of the clusters. When this method is used, the collection of clusters is created when the instance of `DhbClusterFinder` is initialized; each cluster is created in an undefined state.

Listing 12.11 Smalltalk K -cluster algorithm

Class	DhbCluster
Subclass of	Object

Instance variable names: accumulator previousSampleSize

Class methods

```
new
^super new initialize
```

Instance methods

```
accumulate: aVector
    accumulator accumulate: aVector.

centerOn: aVector
    self subclassResponsibility

changes
```

$\wedge(\text{self sampleSize} - \text{previousSampleSize}) \text{ abs}$

distanceTo: aVector

$\wedge\text{self subclassResponsibility}$

initialize

previousSampleSize := 0.
 $\wedge\text{self}$

isInsignificantIn: aClusterFinder

$\wedge\text{self sampleSize} \leq \text{aClusterFinder minimumClusterSize}$

isUndefined

$\wedge\text{self subclassResponsibility}$

reset

previousSampleSize := self sampleSize.
 self collectAccumulatorResults.
 accumulator reset

sampleSize

$\wedge\text{accumulator count}$

Class **DhbClusterFinder**

Subclass of DhbIterativeProcess

Instance variable names: dataServer dataSetSize minimumRelativeClusterSize

Class methods

new: anInteger server: aClusterDataServer type: aClusterClass

$\wedge\text{super new initialize: anInteger server: aClusterDataServer type:}$
 aClusterClass

Instance methods

accumulate: aVector

```
( result at: ( self indexOfNearestCluster: aVector)) accumulate:
    aVector.
```

```
clusters: aCollectionOfClusters
```

```
result := aCollectionOfClusters.
```

```
collectChangesAndResetClusters
```

```
| hasEmptyClusters changes |
changes := 0.
hasEmptyClusters := false.
result do:
    [:each |
        changes := each changes + changes.
        ( each isInsignificantIn: self)
            ifTrue:
                [each centerOn: nil.
                 hasEmptyClusters := true]
            ifFalse: [each reset].
    ].
hasEmptyClusters
    ifTrue: [result := result reject: [:each | each
                                       isUndefined]].
^changes / 2
```

```
dataServer: aClusterDataServer
```

```
dataServer := aClusterDataServer.
```

```
evaluateIteration
```

```
dataServer reset.
dataSetSize := 0.
[ dataServer atEnd]
    whileFalse:[ self accumulate: dataServer next.
                 dataSetSize := dataSetSize + 1.
    ].
^self collectChangesAndResetClusters
```

```
finalizeIterations
```

```
dataServer close
```

indexOfNearestCluster: aVector

```

| distance index |
index := 1.
distance := ( result at: 1) distanceTo: aVector.
2 to: result size do:
    [ :n | | x |
      x := ( result at: n) distanceTo: aVector.
      x < distance
      ifTrue: [ distance := x.
                index := n.
              ].
    ].
^index

```

initialize: anInteger server: aClusterDataServer type: aClusterClass

```

self dataServer: aClusterDataServer.
self clusters: ( (1 to: anInteger) collect: [ :n | aClusterClass
                                                new]).
minimumRelativeClusterSize := 0.
^self

```

initializeIterations

```

dataServer open.
result
do: [:each | each isUndefined ifTrue: [each centerOn:
                                       dataServer next]]

```

minimumClusterSize

```

^(minimumRelativeClusterSize * dataSetSize) rounded

```

minimumRelativeClusterSize: aNumber

```

minimumRelativeClusterSize := aNumber max: 0.

```

printOn: aStream

```

aStream nextPutAll: 'Iterations: '.
iterations printOn: aStream.
result do: [ :each | aStream cr. each printOn: aStream].

```

Listing 12.12 shows the implementation of the concrete cluster class `DhbEuclideanCluster`. The corresponding accumulator is an instance of class `DhbVectorAccumulator`. Data points are directly accumulated into the accumulator; individual data points are not kept.

Listing 12.12 Smalltalk implementation of an Euclidean cluster

```

Class                DhbEuclideanCluster
Subclass of          DhbCluster
Instance variable names: center

    Instance methods
centerOn: aVector

    center := aVector.
    accumulator := DhbVectorAccumulator new: aVector size.

collectAccumulatorResults

    center := accumulator average copy.

distanceTo: aVector

    ^ (aVector - center) norm

isUndefined

    ^center isNil

printOn: aStream

    accumulator count printOn: aStream.
    aStream nextPutAll: ' '.
    center printOn: aStream.
```

12.6.3 Cluster analysis — Java implementation

Listing 12.13 shows the abstract implementation of the K -cluster algorithm in Java. Listing 12.14 shows a concrete implementation using Euclidean distance. The following code example shows how to implement a search for Euclidean clusters.

Code example 12.8

<Creating an instance of a concrete subclass of AbstractDataServer into the variable dataServer>

```
EuclideanCluster[] clusters = new EuclideanCluster[6];
for ( int i = 0; i < clusters.length; i++ )
    clusters[i] = new EuclideanCluster();
ClusterFinder finder =
    new ClusterFinder( clusters, server);
try { finder.setMinimumRelativeClusterSize( 0.09);}
    catch( IllegalArgumentException e){ };
finder.evaluate();
EuclideanCluster[] clusters = finder.getClusters();
```

The first line after setting up the data server creates an array of 6 Euclidean clusters. The loop after this line populates this array with empty clusters. Next, an instance of a cluster finder is created with the array of clusters and the data server. The statement within the try...catch block defines the minimum relative size of significant clusters. The next statement performs the algorithm. The last statement retrieves the array of found clusters.

The constructor for the class ClusterFinder takes two arguments: an array of clusters and the data server. The array of clusters may contain undefined clusters. The undefined clusters are initialized in the method initializeIterations. The data server must be an instance of a concrete data server class as described in section 12.1.2.

Listing 12.13 Java *K*-cluster algorithm

```
package DhbDataMining;

import DhbMatrixAlgebra.DhbVector;
/**
 * Abstract cluster.
 *
 * @author Didier H. Besset
 */
public abstract class Cluster
{
    protected long previousSampleSize = 0;
/**
 * Default constructor method.
 */
public Cluster() {
}
/**
 * Constructor method.
 */
```

```

public Cluster(DhbVector v)
{
    initialize(v);
}
/**
 * @param Object    data point
 */
public abstract void accumulate(DhbVector dataPoint);
/**
 * @param Object    data point
 */
public abstract double distanceTo(DhbVector dataPoint);
/**
 * @return long    number of data points taken from or added to the receiver
 */
public long getChanges()
{
    return Math.abs( getSampleSize() - previousSampleSize);
}
/**
 * @return long    number of data points accumulated in the receiver
 */
public abstract long getSampleSize();
/**
 * @param v DhbMatrixAlgebra.DhbVector
 */
public abstract void initialize( DhbVector v);
/**
 * @return boolean    true if no data was accumulated in the receiver
 */
public boolean isEmpty()
{
    return getSampleSize() == 0;
}
/**
 * @return boolean    true if the receiver should be dropped from
 *                    the cluster finder
 * @param finder DhbDataMining.ClusterFinder
 */
public boolean isInsignificantIn( ClusterFinder finder)
{
    return getSampleSize() <= finder.minimumClusterSize();
}
/**
 * @return boolean    true if the cluster is in an undefined state.
 */

```

```

public abstract boolean isUndefined( );
public void reset()
{
    previousSampleSize = getSampleSize();
}
}

package DhbDataMining;

import DhbIterations.IterativeProcess;
import DhbMatrixAlgebra.DhbVector;
/**
 * Implements k-cluster algorithm
 *
 * @author Didier H. Besset
 */
public class ClusterFinder extends IterativeProcess
{
    private Cluster[] clusters;
    private AbstractDataServer server;
    private double minimumRelativeClusterSize = 0;
    private long dataSetSize;
    /**
     * Constructor method
     * @param Cluster[] clusterArray  initial clusters
     * @param server AbstractDataServer  server for the data points
     */
    public ClusterFinder(Cluster[] clusterArray,
                        AbstractDataServer clusterServer)
    {
        clusters = clusterArray;
        server = clusterServer;
    }
    /**
     * Constructor method
     * @param numberOfCluster int  maximum number of clusters foreseen
     * @param server AbstractDataServer  server for the data points
     */
    public ClusterFinder(int numberOfCluster,
                        AbstractDataServer clusterServer)
    {
        clusters = new Cluster[numberOfCluster];
        server = clusterServer;
    }
    /**
     * Accumulate all data points into the nearest cluster.

```

```

*/
private void accumulateData()
{
    dataSetSize = 0;
    try {
        while (true)
        {
            DhbVector dataPoint = server.read();
            nearestCluster( dataPoint).accumulate( dataPoint);
            dataSetSize += 1;
        }
    } catch ( java.io.EOFException e) { }
}
/**
 * @return int    number of data points which changed clusters since
 *                the last iteration
 */
private int collectChangesAndResetClusters()
{
    int n = 0;
    int emptyClusters = 0;
    for ( int i = 0; i < clusters.length; i++ )
    {
        n += clusters[i].getChanges();
        if ( clusters[i].isInsignificantIn( this ) )
        {
            emptyClusters += 1;
            clusters[i] = null;
        }
        else
            clusters[i].reset();
    }
    if ( emptyClusters > 0 )
        removeEmptyClusters( emptyClusters);
    return n / 2;
}
/**
 * Perform one iteration step.
 * @return double  number of data points which changed clusters
 *                since the last iteration
 */
public double evaluateIteration()
{
    server.reset();
    accumulateData();
    return collectChangesAndResetClusters();
}

```

```

    }
    /**
     * Closes the data point server
     */
    public void finalizeIterations()
    {
        server.close();
    }
    /**
     * @return DhbDataMining.Cluster[] clusters contained in the receiver
     */
    public Cluster[] getClusters()
    {
        return clusters;
    }
    /**
     * Opens the data stream and creates the initial clusters.
     */
    public void initializeIterations()
    {
        server.open();
        try {
            for ( int i = 0; i < clusters.length; i++ )
            {
                if ( clusters[i].isUndefined() )
                    clusters[i].initialize(server.read());
            }
        } catch (java.io.EOFException e) {};
    }
    /**
     * @return long    minimum cluster size to be considered in the next iteration
     */
    public long minimumClusterSize()
    {
        return Math.round( minimumRelativeClusterSize * dataSetSize);
    }
    /**
     * @param dataPoint ClusterData
     * @return Cluster    nearest to the data point
     */
    private Cluster nearestCluster( DhbVector dataPoint)
    {
        int index = 0;
        int nearestIndex = index;
        double closestDistance = clusters[index].distanceTo( dataPoint);
        double distance;

```

```

    while ( ++index < clusters.length )
    {
        distance = clusters[index].distanceTo( dataPoint);
        if ( distance < closestDistance )
        {
            closestDistance = distance;
            nearestIndex = index;
        }
    }
    return clusters[nearestIndex];
}
/**
 * Removes empty clusters. The array of clusters is reconstructed.
 * @param n int    number of empty clusters
 */
private void removeEmptyClusters( int n)
{
    Cluster[] newClusters = new Cluster[ clusters.length - n];
    int index = 0;
    for ( int i = 0; i < clusters.length; i++ )
    {
        if ( clusters[i] != null )
            newClusters[index++] = clusters[i];
    }
    clusters = newClusters;
}
/**
 * @return DhbDataMining.Cluster[] clusters contained in the receiver
 */
public void setClusters( Cluster[] clusterArray)
{
    clusters = clusterArray;
}
/**
 * @param r double    the minimum relative size of a cluster to be kept
 *                    in the next iteration
 * @exception java.lang.IllegalArgumentException
 *                    argument cannot be negative.
 */
public void setMinimumRelativeClusterSize( double r)
    throws IllegalArgumentException
{
    if ( r < 0 )
        throw new IllegalArgumentException(
            "Relative cluster size cannot be negative");
    minimumRelativeClusterSize = r;
}

```

```

    }
    /**
     * @return java.lang.String
     */
    public String toString()
    {
        StringBuffer sb = new StringBuffer();
        sb.append( "Iterations: "+getIterations());
        for ( int i = 0; i < clusters.length; i++)
        {
            sb.append( '\n');
            sb.append( clusters[i]);
        }
        return sb.toString();
    }
}

```

The concrete class `EuclideanCluster` implements clusters using the Euclidean distance defined by equation 12.13. It has the following instance variables

`center` the center of the cluster — that is, the position from which distances to the cluster are calculated;

`accumulator` an instance of the class `VectorAccumulator` described in section 12.2.3.

The accumulation of the data points is delegated the accumulator.

Listing 12.14 Java implementation of an Euclidean cluster

```

package DhbDataMining;

import DhbMatrixAlgebra.DhbVector;
import DhbMatrixAlgebra.DhbIllegalDimension;
/**
 * Cluster using Euclidean distance.
 *
 * @author Didier H. Besset
 */
public class EuclideanCluster extends Cluster
{
    private DhbVector center;
    private VectorAccumulator accumulator;
    /**
     * Default constructor method.
     */
    public EuclideanCluster()

```

```

{
    super();
}
/**
 * Constructor method.
 * @param DhbVector center of the receiver
 */
public EuclideanCluster(DhbVector dataPoint)
{
    super( dataPoint);
}
/**
 * @param dataPoint DhbVector    data point
 */
public void accumulate(DhbVector dataPoint)
{
    accumulator.accumulate( dataPoint);
}
/**
 * @param dataPoint DhbVector    data point
 * @return DhbVector    square of the Euclidian distance from the data
 *                        point to the center of the receiver.
 */
public double distanceTo( DhbVector dataPoint)
{
    try{ DhbVector v = dataPoint.subtract( center);
        return v.product( v);}
    catch(DhbIllegalDimension e) { return Double.NaN;}
}
/**
 * @return long    number of data points accumulated in the receiver
 */
public long getSampleSize()
{
    return accumulator.getCount();
}
/**
 * @param v DhbVector    center for the receiver
 */
public void initialize(DhbVector dataPoint)
{
    center = dataPoint;
    accumulator = new VectorAccumulator( dataPoint.dimension());
}
/**
 * @return boolean    true if the cluster is in an undefined state.

```



```

*/
public boolean isUndefined()
{
    return center == null;
}
public void reset()
{
    super.reset();
    center = accumulator.averageVector();
    accumulator.reset();
}
/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    sb.append( previousSampleSize);
    sb.append( ' ');
    sb.append( center);
    return sb.toString();
}
}

```

12.7 Covariance clusters

As we have seen in section 12.5 the Mahalanobis distance is a distance in the geometrical sense. Thus, this distance can be used by the K -cluster algorithm. We call clusters using the Mahalanobis distance covariance clusters since the metric for the distance is based on the covariance matrix.

The normalizing properties of the Mahalanobis distance makes it ideal for this task. When Euclidean distance is used, the metric remains the same in all directions. Thus, the extent of each cluster has more or less circular shapes. With the Mahalanobis distance the covariance metric is unique for each cluster. Thus, covariance clusters can have different shapes since the metric adapts itself to the shape of each cluster. As the algorithm progresses the metric changes dynamically.

12.7.1 Covariance clusters — General implementation

Covariance clusters need little implementation. All tasks are delegated to a Mahalanobis center described in section 12.5. Listing 12.15 shows the Smalltalk implementation and the Java implementation is shown in listing 12.16.

Figure 12.1 with the boxes
CovarianceCluster grayed.

Listing 12.15 Smalltalk covariance cluster

```

Class                DhbCovarianceCluster
Subclass of         DhbCluster
Instance variable names: center

    Instance methods
centerOn: aVector

    accumulator := aVector ifNotNil: [ :v | DhbMahalanobisCenter
                                           onVector: v].

collectAccumulatorResults

    accumulator computeParameters.

distanceTo: aVector

    ^accumulator distanceTo: aVector

isUndefined

    ^accumulator isNil

printOn: aStream

    accumulator printOn: aStream.

```

Listing 12.16 Java covariance cluster

```

package DhbDataMining;

import DhbMatrixAlgebra.DhbVector;
import DhbMatrixAlgebra.SymmetricMatrix;
import DhbMatrixAlgebra.DhbIllegalDimension;
/**
 * Cluster using Mahalanobis distance.
 *
 * @author Didier H. Besset
 */
public class CovarianceCluster extends Cluster
{

```

```
    private MahalanobisCenter center = null;
/**
 * Default constructor method.
 */
public CovarianceCluster()
{
    super();
}
/**
 * Constructor method.
 * @param DhbVector center of the receiver
 */
public CovarianceCluster(DhbVector v)
{
    super(v);
}
/**
 * Accumulation is delegated to the Mahalanobis center.
 */
public void accumulate(DhbMatrixAlgebra.DhbVector dataPoint)
{
    center.accumulate( dataPoint);
}
/**
 * Distance computation is delegated to the Mahalanobis center.
 */
public double distanceTo( DhbVector dataPoint)
{
    return center.distanceTo( dataPoint);
}
/**
 * @return long    number of data points accumulated in the receiver
 */
public long getSampleSize()
{
    return center.getCount();
}
/**
 * @param v DhbVector    center for the receiver
 */
public void initialize(DhbVector v)
{
    center = new MahalanobisCenter( v);
}
/**
 * @return boolean    true if the cluster is in an undefined state.
```

```
*/
public boolean isUndefined()
{
    return center == null;
}
public void reset()
{
    super.reset();
    center.computeParameters();
    center.reset();
}
/**
 * @return java.lang.String
 */
public String toString() {
    return center.toString();
}
}
```

Appendix A

Decimal floating-point simulation

The class `DhbDecimalFloatingNumber` is intended to demonstrate rounding problems with floating-point number representation. It models the floating-point number representation. The radix of the floating-point representation is decimal to ease the reading of the results. It also allows people to carry some of the operations by hand. This model is almost equivalent to what is done inside a computer. One notable difference is the absence of exponent offset. Another difference is that there is no limit on the size of the exponent. Simple as it is this model can be used to illustrate rounding problems to beginners. This class is only intended for didactical purposes.

Only the Smalltalk implementation is given here, as Java does not lend itself to operator overloading. Moreover, fraction arithmetic is not available in Java. Thus, making an equivalent class would require much more code.

Instances of the class are created with the method `new`: supplying any number as argument. For example,

```
DhbDecimalFloatingNumber new : 3.141592653589793238462643
```

Arithmetic operations are performed as usual. For example, the first expression of section 1.3.2 can be coded as:

```
Code example A.1      | a b |
a := DhbDecimalFloatingNumber new: (2 raisedToInteger: 64).
b := DhbDecimalFloatingNumber new: 300.
a + b
```

The class has two instance variables:

`mantissa` contains the mantissa of the number normalized to a finite number of digits,

`exponent` contains the negative of the decimal exponent of the number.

The class has one class variable:

`Digits` contains the maximum number of digits kept in the mantissa of all instances.

The method `normalize:` computes the mantissa and exponent of the argument and stores it into the corresponding instance variables. The method `value` allows retrieving the value of the number in numerical form. The method `printOn:` allows to visualize the representation of the number.

The number of decimal digits kept in the mantissa can be changed by the user using the method `setDigits:`. By default it is set to 15, which corresponds roughly to the precision of a 64-bit IEEE standard floating-point number.

The four arithmetic operations and the square root have been defined. The mechanism of each operation is the same. First each operand is converted to a fraction. Then, the operation is carried using normal numbers. In the case of arithmetic operation the result is exact since the use of Smalltalk fractions guarantees absolute precision. Then, a new instance of the class `DhbDecimalFloatingNumber` is created from the result and is returned. Rounding errors will occur during this step.

Listing A.1 Smalltalk code simulating decimal floating number arithmetic

```

Class                DhbDecimalFloatingNumber
Subclass of          Object
Instance variable names: mantissa exponent
Class variable names:  Digits

    Class methods
defaultDigits
    ^15

defaultDigits: anInteger
    Digits := anInteger.

digits
    Digits isNil
        ifTrue: [ Digits := self defaultDigits].
    ^Digits

new: aNumber
    ^self new normalize: aNumber

resetDigits
    Digits := nil.
```

Instance methods

* aNumber

^self class new: (self value * aNumber value)

+ aNumber

^self class new: (self value + aNumber value)

- aNumber

^self class new: (self value - aNumber value)

/ aNumber

^self class new: (self value / aNumber value)

normalize: aNumber

```
exponent := (self class digits - (aNumber log: 10)) floor.
mantissa := ( aNumber * ( 10 raisedToInteger: exponent))
              truncated.
^self
```

printOn: aStream

```
mantissa printOn: aStream.
aStream nextPutAll: 'xE'.
exponent negated printOn: aStream.
```

sqrt

^self class new: self value sqrt

value

^mantissa / (10 raisedToInteger: exponent)

Appendix B

Smalltalk primer for Java programmers

This appendix is meant as a quick introduction to the Smalltalk language for Java programmers. The goal of the explanations is not to expose all features of Smalltalk, but to give enough background so that the reader is able to read the Smalltalk code presented in this book.

There are a few Smalltalk books on the market. The most recent is the book by David Smith [Smith]. The book of Kent Beck [Beck] is a good choice to deepen Smalltalk knowledge and Smalltalk specific object-oriented approach.

B.1 Syntax in a nutshell

In Smalltalk there is no primitive type. Everything is an object. Objects are represented in the code by a symbol. Objects communicate together by sending messages to each other. As in Java, a symbol may contain any alphanumerical characters and must begin with a lower case letter.

B.1.1 Smalltalk expressions

A Smalltalk expression is composed of an identifier representing the object receiving the message followed by the message. The message is placed after the name of the object to which it is directed, separated by at least a blank separator (space, tabulation or carriage return).

Objects are represented by variables in the conventional way. There are three kinds of messages:

1. unary messages,
2. binary messages and
3. keyword messages.

Unary messages correspond to a Java method without argument. There are represented by symbols containing any alphanumerical characters and beginning with a lower case letter.

Binary messages correspond to conventional operators in other languages, Java included. Examples are the arithmetic operators (+, -, * and /) or the relational operators (=, >, <, >= and <=). In Smalltalk, the inequality operator is noted `~=`. Other non-alphabetical operator symbols can be used as binary messages.

Keyword messages correspond to a Java method with one or more arguments. In a Smalltalk keyword message each argument is placed after a keyword. Each keyword is written as a symbol containing any alphanumerical characters, beginning with a lower case letter and ending with a semi colon (:).

Table B.1 shows a few examples of Smalltalk messages together with their Java equivalent. To make the examples easy to follow, the objects used are either constants or objects from the single class String. These objects are denoted by the symbols `s1` and `s2`. The class String has the advantage of being the same in both Smalltalk and Java.

Table B.1: Sample Smalltalk messages with their Java equivalent

Message type	Smalltalk	Java
Unary	<code>s1 size</code>	<code>s1.length()</code>
	<code>s1 hash</code>	<code>s1.hashCode()</code>
Binary	<code>s1 = s2</code>	<code>s1.equals(s2)</code>
	<code>s1 < s2</code>	<code>s1.compare(s2)</code>
Keyword	<code>s1 at: 3</code>	<code>s1.charAt(3)</code>
	<code>s1 indexOfSubCollection: s2 startingAt: 5</code>	<code>s1.indexOf(s2,5)</code>
	<code>s1 copyFrom: 3 to: 5</code>	<code>s1.substring(3,5)</code>

B.1.2 Precedence

Arguments of binary and keyword messages may be other Smalltalk expressions. In case of combined expressions, the precedence goes as follows: unary messages are evaluated first, then binary messages and finally keyword messages. Messages of same type are evaluated from left to right. This gives a somewhat unconventional precedence rule for arithmetic expressions. As in any other computer language expressions enclosed within parentheses are always evaluated first, starting with the innermost pair of parentheses.

As a consequence, keyword messages used as arguments to other messages must always be enclosed within parentheses.

B.1.3 Assignment, equality and identity

In Smalltalk, the assignment operator is composed of an equal sign followed by a colon (`:=`). This corresponds to the equal sign in Java.

Like Java, Smalltalk differentiates between equality and identity. The equality operator is an equal sign, corresponding to the method `equals` of Java. The inequality

operator is written `~=` corresponding to the Java `!=` equals. The identity operator is a double equal sign (`==`) like in Java. The negation of the identity is written as `~~`. Table

Table B.2: Smalltalk assignment and equalities

Operator	Smalltalk	Java
Assignment	<code>:=</code>	<code>=</code>
Equality	<code>=</code>	<code>equals()</code>
Inequality	<code>~=</code>	<code>!=</code>
Identity	<code>==</code>	<code>==</code>
Non-identity	<code>~~</code>	<code>!=</code>

B.2 Class and methods

A Smalltalk class is quite similar to a Java class. The main difference is that Smalltalk is not file-oriented. Classes are not assigned to a file like in Java. Instead they reside in the Smalltalk image, that is a copy of the memory used by

As a consequence, any class can be extended by anyone. If an application designer is missing a method from the base classes, the method is simply added by the designer. This book contains numerous example of methods added to the base classes `Number` and `Integer`.

A class is declared by stating its superclass and the instance variables. There are other parameters defining a class, but we do not mention them as they are not used in this book. As in Java, the class `Object` is the topmost class.

Smalltalk instance variable are listed as symbols without types. Smalltalk is a dynamically typed language. In principle an instance variable can hold any object. At run time, however, the type of the instance is known. This is how the virtual machine knows how to retrieve the methods supported by the object.

B.2.1 Instance methods

An instance method is very similar to a Java instance method. Of course, the syntax is quite different. At best, we shall discuss an example, taken from listing 8.15. The lines are numbered for easier references.

```

1 evaluateIteration
2 | indices |
3 indices := self largestOffDiagonalIndices.
4 self transformAt:(indices at: 1) and:(indices at: 2).
5 ^precision

```

The first line is the method declaration, that is the declaration of the message sent when this method is executed. In this example, this is an unary message named `evaluateIteration`.

Line 2 is the declaration of the variables local to the method. Since Smalltalk is not typed, only the names of the variable are enumerated between two vertical bars. If a method does not have local variables, this line is omitted. Here the only declared variable is `indices`.

Line 3 is an assignment statement: the local variable `indices` is assigned to the result of sending the message `largestOffDiagonalIndices` to the variable `self`. `self` is the instance, which is executing the method. In other words, it is equivalent to the Java variable `this`. The statement is terminated with a dot (`.`) corresponding to the semicolon used in Java. Line 4 is a simple statement. The message `transformAt:and:` is sent to the instance executing the method. The arguments are the results of two keyword messages (`at:`) sent to the variable `indices`. In this case, the variable `indices` was set to an array with two elements. These elements are used as arguments for the message `transformAt:and:`. Here again the statement is terminated by a dot.

Line 5 is the last statement of the method. The wedge (`^`) indicates that the expression that follows is the result of the method. In this case, the result of the method is the value of the instance variable `precision`. A return statement is not terminated.

The next example is taken from listing 8.8. It is far from being simple, but it covers more advance features. Together with the first example we have covered the entire syntax of Smalltalk.

```

1  decompose
2  | n |
3  n := rows size.
4  permutation := (1 to: n) asArray.
5  1 to: (n - 1) do:
6    [ :k |
7      self swapRow: k withRow: (self largestPivotFrom: k);
8      pivotAt: k.
9    ].

```

The first line is the method declaration for the unary message named `decompose`.

Line 2 is the declaration of the local variable `n`.

Line 3 is an assignment statement: the local variable `n` is set to the number of rows. The variable `rows` is an instance variable of the class and is set to an array; the message `size` returns the number of elements located in the array. The statement is terminated with a dot (`.`).

Line 4 is another assignment statement. It assigns the instance variable `permutation` to an array containing the series of integers 1 to `n`. The message `to:` sent to an integer answers an interval. It must be converted to an array with the message `asArray`. Here again the statement is terminated by a dot.

Line 5 is the beginning of an iterator message ending at line 9. Iterator methods are described in section B.3. The object to which the iterator message is sent is an interval from 1 to `n - 1`. This line is equivalent to the Java statement `for (int k = 1; k < n; k++)`. The reader should notice that indices in Smalltalk begin at 1 instead of 0.

Line 6 is the beginning of the block, argument of the `do:` message. This line is declarative and states that the variable used to evaluate the block is named `k`.

Line 7 contains two messages sent to the variable `self`. The first message to be executed is a keyword message — `largestPivotFrom:` — with one argument `k`. The second message is a keyword message `swapRow:withRow:` with two arguments: the first argument is `k` and the second argument is the result of the message `largestPivotFrom:`.

Unlike the preceding statements, this statement is terminated by a semicolon (;). In Smalltalk a semicolon is not a statement terminator. The semicolon indicates that the message on line 8 — written without explicit target — is directed to the same target as the preceding message. In this case, the target of the message `pivotAt:` — a keyword message with one argument `k` — is `self`.

Line 9 is the end of the statement beginning on line 5. It is the terminating bracket of the block beginning on line 6. This statement is terminator with a dot. Because this method does not return an explicit result, the result of the method is `self`, that is the instance which was executing the method.

B.2.2 Class methods

The biggest difference between Smalltalk and Java lies in the class methods. As a first approximation, a Java programmer can think that class methods are equivalent to static methods. Class methods, however, are methods like any other methods. In particular class methods are fully inherited.

Here the Java programmer must be in mind that everything is an object in Smalltalk. This is also true for classes. A class is an object and, as such, has methods. Thus, class methods are exactly like instance methods, but they are defined on the class as opposed to the instance. In particular, the variable `self` in a class method now refers to the class itself.

Class methods are also used as class constructor methods. Unlike Java, Smalltalk allows class constructor methods with any name. The default creation method `new` is provided by the superclass of all classes. It creates a new instance of the class with all instance variables set to `nil`. An application designer can chose to redefine the method `new` for a given subclass. This book shows several example of this.

B.2.3 Block

In Smalltalk everything is an object. And yes! This is the third times this important statement is made. This is true for Smalltalk code itself. A Smalltalk code object is specified with a special syntax: the block. Here is an example of a block computing polar coordinates from supplied Cartesian coordinates:

```
[ :x :y |
  | radius angle |
  radius := ( x squared + y squared) sqrt.
  angle := y arctan: x.
  Array with: radius with: angle]
```

The block is delimited by a pair of brackets ([]). The first line contains the declaration of the block variables. These variables are supplied to the block at execution time. The next line contains the declaration of the local variables, that is variables used within the

block. The next two lines perform computation needed to obtain the final result. The last expression of the block is the value returned by the block at execution time.

Here is an example on how the block of the previous example is used.

```
| block |
block := [ :x :y |
  | radius angle |
  radius := ( x squared + y squared) sqrt.
  angle := y arctan: x.
  Array with: radius with: angle].
block value: 2 value: 3.
```

Of course, this is a contrived example. Usually block are much simpler. They are also seldom used to perform computation for the sake of computation, although functions described in chapter 2 are a good example of using blocks.

B.3 Iterator methods

The most important use of blocks is within so-called iterator methods. Quickly said, iterators method provide the functionality of the Enumeration interface. Smalltalk iterator methods, however, provide far most flexibility that the Enumeration interface.

Iterator methods can be applied to any instance of any subclass of the class Collection. A Smalltalk collection is simply a container of objects, not necessarily of the same class. Depending on the particular subclass, a collection has some specific structure. They cover many type including the Java arrays and the Java classes Vector or HashTable.

There are many iterator methods. This section only describes the one used in this book.

B.3.1 do:

The do: iterator corresponds to the Java for loop. In Smalltalk, however, there is no need to supply an index. The block supplied as argument of the do: message is evaluated for each element of the collection. For example, to perform the message sample on each element of a collection one simply writes:

```
aCollection do: [ :each | each sample].
```

It is customary to use the variable name each as the argument of the block used with an iterator method.

B.3.2 collect:

The iterator method collect: has no equivalent in Java. Its results is a collection of the same type¹ as the collection to which the message was sent. The elements of the new collection are the result of the block supplied as argument to the method. For example,

¹There are some special collections for which the type of the result is different.

here is how one can construct an array containing the sum of the squares of the integers 1 to 9:

```
#(1 2 3 4 5 6 7 8 9) collect: [ :each | each squared].
```

The result of the code above is the array #(1 4 9 16 25 36 49 64 81).

B.3.3 inject:into:

The iterator method `inject:into:` is a great tool for mathematical evaluation. As it is rather complex, we shall start with an example. Here is how to compute the sum of the elements of an array:

```
anArray inject: 0 into: [ :sum :each | sum + each].
```

The iterator method `inject:into:` is a keyword message. The first argument is the initial value used in the summation. The second argument is a block with two arguments: the first argument is the result of the evaluation of the block with the preceding element or the initial value if this is the first evaluation; the second is the element of the collection over which the iterator method is iterating. The result of the message is the value of evaluating the block on the last element of the collection.

B.4 Double dispatching

Since the arguments of Smalltalk methods are not typed, a method is able to accept arguments of different classes as long as all possible argument types have the same behavior. However, a problem arises when the method must behave differently depending on the type of the argument.

For example, the multiplication operator defined for polynomials (c.f. section 2.2.3) can accept a number or another polynomial. The implementation of the method strongly differ depending on the type of argument. One possible implementation could be to test for the type of the argument and switch to the proper method accordingly. In this case the code would look as follows:

Code example B.1

```
* aNumberOrPolynomial
^aNumberOrPolynomial class = DhbPolynomial
  ifTrue: [ self productWithNumber: aNumberOrPolynomial]
  ifFalse:[ self productWithPolynomial: aNumberOrPolynomial]
```

The code above is a bad example. This is usually what beginners do, especially those who still think with a strong legacy coming from other languages.

The elegant solution is called double dispatching and is illustrated on figure B.1. It merely uses the fact that testing which class an object belongs to is exactly what the virtual machine is doing when looking for the code of the method corresponding to a message received by an object. Thus, it is not necessary to repeat such test. Instead, one delegates the execution of the method to the argument object. as follows:




Figure B.1: Triple dispatching

Code example B.2

```

* aNumberOrPolynomial
^aNumberOrPolynomial timesPolynomial: self

```

In fact, the sending of the message `timesPolynomial: self` to the method's argument ensures that the types of both operands are now completely known. One must now implement two versions of the method `timesPolynomial:`, one for instances of the class `DhbPolynomial` and one for the class `Number`. Within both versions of the method, the programmer can be sure that the type of the argument is indeed an instance of class `DhbPolynomial`.

Figure B.1 shows this process schematically. Each box shows the class executing the method on the top and the method with the type of the argument explicited at the bottom. The arrows represent the sequence of execution.

One caveat with double dispatching is that the order of the arguments is inverted. Thus, implementing double dispatching for non commutative operators — such as subtract, divide, or matrix multiplication — requires some thinking.

It is easy to understand that double dispatching is much faster than testing the type of the argument. It only requires the invocation of the operation message, whereas testing requires evaluating the test itself plus a message invocation².

B.5 Multiple dispatching

The technique of double dispatching can be extended to multiple levels. This is required when an operation is implemented by a class having several subclasses, each subclasses requiring a different behavior.

A good example of multiple dispatching is given by the addition between matrices (c.f. section 8.1.1). The product of two symmetric matrices is a symmetric matrix. In all other cases, the result is a non-symmetric matrix. Thus, the addition operation is delegated a first time to the proper subclass of the argument and a second time to the class of the first operand. Here we have triple dispatching. This is shown schematically in figure B.2.

²`ifTrue:ifFalse:` is a message sent to the result of the testing.

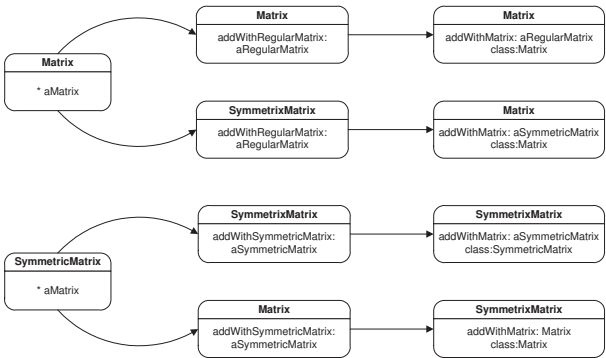


Figure B.2: Triple dispatching

In the case of matrix multiplication, the situation is more complex since the product is already being dispatched to distinguish between three possible arguments. Here quadruple dispatching is necessary.

Appendix C

Java primer for Smalltalk programmers

This appendix is meant as a quick introduction to the Java language for Smalltalkers. We shall give minimum explanations to give the Smalltalker enough background to be able to read the Java code presented in this book.

There are many books about Java on the market. My favorite book is *Java in a nutshell*[Flanagan]. This book teaches Java assuming that the reader is already familiar with C. A similar approach is used here. This makes sense since the syntax of Java is very close to that of C.

C.1 Remarks on the syntax

The designer of Java elected to use the existing syntax of C. Compared to Smalltalk, this means a lot of redundant notation: curly braces everywhere and a lot of reserved keywords. Statements are terminated by a semicolon (;). We shall not explain the syntax because it would be too long and we shall assume that most Smalltalk programmers have already been exposed to C. Under this assumption, we shall concentrate on the object-oriented features of Java.

Like C, Java is file oriented. An application is made of several files loaded together. One file may contain several class definitions. However, a class definition must be contained within a single file. As a consequence, base classes as well as any class obtained from a third party cannot be extended. The only possibility is to extend them when it is possible¹

C.1.1 Classes

A Java class is defined as follows:

¹Java has a qualifier `final` which is a compiler optimization. Classes with the `final` qualifier cannot have subclasses.

```
public class MyClass { }
```

This defines a public class named `MyClass` as a subclass of the class `Object` by default. Like in Smalltalk the superclass of all classes is the class `Object`. This definition of `MyClass` is not a very useful one since it has no instance variable and no method. Instance variables and methods are declared within the brackets following the class name. The keyword `public` states that this class is visible by any other class.

The class `MySubclass` is defined as a subclass of `MyClass` by the following declaration

```
public class MySubclass extends MyClass { }
```

Instances of class are created via a constructor method. The constructor method has the same name as the class. The constructor method is called by supplying the keyword `new` before the class name. For example, here is how to create a new instance of the class `MyClass`:

```
MyClass x = new MyClass();
```

After this statement, the local variable `x` points to the newly created instance.

Constructors methods may have parameters. Since Java differentiates methods by their names and by the type of its arguments, a class can have several constructor methods.

C.1.2 Instance variables

Instance variables must be declared with a type. An optional qualifier defines the scope (or visibility) of the variable. An instance variable is declared as follows

```
protected double var1;
```

This declares the variable `var1` to be of type `double`, a floating number with double precision. The keyword `protected` indicates that the variable is visible to all subclasses of this class. This is the default when no qualifier is supplied. `Protected` instance variables behaves like instance variables in Smalltalk.

C.1.3 Method declaration and method calling

A Java method is declared with a type. An optional qualifier defines the scope (or visibility) of the method. The name of the method is followed by the list of parameters enclosed in parentheses. The parentheses are always supplied, even if the list of parameters is empty. Each parameter is defined by a type and an identifier. This identifier can be used within the method code to refer to the parameters. Like in Smalltalk, all parameters are pointer to the objects. After the declaration the code of the method follows enclosed within curly brackets.

Here is an example of a method declaration:

```
public DhbVector gradient( double x) { ... }
```

This method is declared as returning an object of type `DhbVector`. The argument of the method is a double. The code of the method must be located within curly brackets following the method's declaration: our example does not show the code.

One interesting aspect of Java is that one can declare two methods with the same name provided the type of the arguments is different. For example, the following two methods can also be defined within the same class where the above gradient method is defined:

```
public DhbVector gradient( float x) { ... }
```

```
public DhbVector gradient( double x, double[] parameters) { ... }
```

At calling time, the Java virtual machine selects the proper method depending of the type of the supplied argument(s).

The code of a Java method is of procedural nature. This is mainly due to the use of the C syntax and constructs: for and while loops, if and if...else.

Calling a method is done by appending the name of the method to the object to which the method is sent separated with a dot. For example, the above methods can be called as follows:

```
MyClass x = new MyClass();  
DhbVector v = x.gradient( Math.PI);
```

C.1.4 Objects and non-objects

In Smalltalk everything is an object. This is not the case in Java. There are so-called primitive types which are handled as values like in conventional programming languages. The primitive types are: char, int, long, float, double and boolean.

In order to use values as objects, Java provides so-called wrappers for each of the primitive type. For example, if a double needs to be passed to a method expecting an argument of type `Object`, the double must be wrapped using the `Double` class wrapper as follows:

```
double pi = 3.14159;  
myMethod( new Double(pi));
```

C.1.5 packages

A Java package provide a solution to the name space, which has been the holy grail of Smalltalk vendors for a long time. A Java package provides a way to group cooperating classes together. Classes from one package are not visible to classes of other packages. Thus, two classes in an application may have the same name as long as they are in different packages. To refer to both of these classes in one common place, the designer must use the package name followed by a dot and the class name. For example, if you already have a class named `Matrix` and that you want to use the class `Matrix` of this books, you must refer to the class of this book as: `DhbMatrixAlgebra.Matrix`.

The import statement allows to make classes from another package visible for all methods of a class. Then, a class can be referred by its name only. Our Java code make ample use of import statements.

C.1.6 Scope qualifiers

Java supports 3 types of scope qualifiers:

- private
- protected
- public

The meaning of the qualifiers is slightly different when applied to classes, methods or instance variables.

A private instance variable is accessible only to the class where this instance variable is declared. This means that a private instance variable is not visible to the subclass of a class where the instance variable is declared. A protected instance variable is accessible to all subclasses of the class where this instance variable is declared. It is also visible to all classes located in the same package as the class defining the variable. A public instance variable is accessible by any object.

The scope qualifiers applied to a method have the same properties than for instance variables.

A private class is accessible by all classes within the same file. A protected class is visible to all classes located in the same package. A public class is visible to any class.

The scope qualifiers are also used to optimize code execution. This has an odd consequence for Smalltalkers: a private method of a subclass called in a method from a superclass will not be executed; the private method of the superclass is executed instead. Let us make an example. Consider the two classes shown in listing C.1

Listing C.1 Inheritance of private methods

```
package InheritanceCheck;

public class MyClass {
    public MyClass() {
        super();
    }
    public static void main(java.lang.String[] args)
    {
        MyClass objectA = new MyClass();
        System.out.println("Start running MyClass.privateCommon");
        objectA.privateCommon();
        MySubclass objectB = new MySubclass();
        System.out.println("Start running MySubclass.privateCommon");
        objectB.privateCommon();
    }
}
```

```

    System.out.println("Start running MyClass.protectedCommon");
    objectA.protectedCommon();
    System.out.println("Start running MySubclass.protectedCommon");
    objectB.protectedCommon();
}
private void privateCommon()
{
    System.out.println("\t=> Executing private method in superclass.");
}
private void privateSuper() {
    privateCommon();
}
protected void protectedCommon()
{
    System.out.println("\t=> Executing protected method in superclass.");
}
protected void protectedSuper() {
    protectedCommon();
}
}

```

```
package InheritanceCheck;
```

```

public class MySubclass extends MyClass {
    public MySubclass() {
        super();
    }
    public static void main(java.lang.String[] args)
    {
        MySubclass objectB = new MySubclass();
        System.out.println("Start running MySubclass.privateSub");
        objectB.privateSub();
    }
    private void privateCommon()
    {
        System.out.println("\t=> Executing private method in subclass.");
    }
    private void privateSub() {
        privateCommon();
    }
    protected void protectedCommon()
    {
        System.out.println("\t=> Executing protected method in subclass.");
    }
}

```

The method `privateCommon` executed within the method `privateSuper` is always that of the class `MyClass`. Smalltalkers should carefully study the code as it works against their logic. Executing the main method of class `MyClass` yields the following result:

```
Start running MyClass.privateCommon
=> Executing private method in superclass.
Start running MySubclass.privateCommon
=> Executing private method in superclass.
Start running MyClass.protectedCommon
=> Executing protected method in superclass.
Start running MySubclass.protectedCommon
=> Executing protected method in subclass.
```

Executing the main method of class `MySubclass` yields the following result:

```
Start running MySubclass.privateSub
=> Executing private method in subclass.
```

This surprising behavior is not a bug. It is a documented feature of the Java compiler and virtual machine optimization. This apparent breach of inheritance is a consequence of code optimization.

C.1.7 static qualifier

The static qualifier defines that a variable or a method is attached to the class and not to an instance.

Smalltalkers can think of static variables as class variables. Static methods can be considered as class variables. They are, however, some subtle differences with the Smalltalk equivalents.

C.2 Abstract class and interface

In Java a class or a method can be declared as abstract.

An abstract class is meant to be the start of a sub-hierarchy of classes. In Smalltalk the concept of abstract class is purely a matter of convention. In Java the abstract qualifier is enforced: an abstract class cannot have any instances.

An abstract method can only be defined within an abstract class or an interface (c.f. below). It is defined as follows:

```
protected abstract double computePrecision();
```

Declaring a method as abstract instructs the Java compiler to require that any concrete subclass must implement such a method: a concrete subclass omitting to define concrete implementations of the abstract methods of its superclasses is flagged as an error at compiling time.

An interface is a new concept introduced by Java (at least to my knowledge). One can consider an interface as an abstract class, which cannot have any instance

variable or concrete methods. Interfaces come handy when constructing polymorphic behavior between classes not part of the same hierarchy. Here is an example of interface definition:

```
public interface PlanarShape {
    protected abstract double getPerimeter();
    protected abstract double getArea();
}
```

Once an interface has been declared, it can be used as a type declaration. For example the following method can be written in a class having access to the interface defined just above:

```
public double totalArea( PlanarShape[] p) {
    double totalArea = 0;
    for ( int i = 0; i < p.length; i++)
        totalArea += p.getArea();
    return totalArea;
}
```

At this point, the designer of the interface has no idea how the method `getArea` can be implemented.

A class declares itself as adhering to an interface as follows:

```
public class Square implements PlanarShape {
    private double side;

    .....

    double getArea() {
        return side * side;
    }
}
```

A class can adhere to several interfaces at the same time.

C.3 Exception handling

Exception handling is built in the language. By this I mean that not only exception handling structures exist, like in Smalltalk, but also that these structures are enforced by the compiler.

Exceptions are subclasses of the class `Exception`. Like in Smalltalk, a designer can define its own type of exception. Exceptions have two constructor methods: one with no argument and one with a string argument. The string argument is used to pass additional information to the exception. This information is displayed by the debugger or on the console when the exception is not handled by the application.

In the Java jargon, an exception is thrown by a method and is caught (or better catch-ed) by a caller of that method, not necessarily the direct caller. Here is an example of a method throwing an exception:

```
public double inverse( double x) throws MyZeroDivideException {
    if ( Math.abs( x) < DhbMath.getMachinePrecision() )
        throw new MyZeroDivideException();
    return 1 / x;
}
```

The declaration of the method uses the modifier `throws MyZeroDivideException` to signal the compiler that this method may throw an exception of the specified class. The actual throwing of the exception is made by the statement `throw` followed by a new instance of the exception.

If a method is defined as throwing an exception, any method calling the former must catch the exception. Otherwise, the compiler issues an error. Catching the exception can be made explicitly or implicitly.

Explicit exception catching is made with the statement `try...catch`. Here is an example

```
try{ double x = pivot();
    double inversePivot = inverse(x);
} catch ( MyZeroDivideException e)

    {<code to handle the exception>}
```

To catch an exception, the statements where an exception might occurs must be enclosed inside curly brackets preceded by the key word `try` and followed by the keyword `catch`. The keyword `catch` is followed by a parentheses enclosing a declaration of the exception and a series of statements, enclosed within curly brackets, which are handling the exception. The variable corresponding to the exception — `e` in the example above — is visible inside the block handling the exception.

Implicit exception catching is made by declaring that the method is throwing the same exception. For example,

```
public double inversePivot() throws MyZeroDivideException {
    double pivot = pivot();
    return inverse( pivot);
}
```

In this case the method `inversePivot` delegates the handling of the exception to one of its callers.

C.4 Collections and related topics

Smalltalkers, when first exposed to Java, surely are appalled at the scarcity of the collection classes and at the absence of decent iterator methods. Well, the sad thing is they are right! The Java designer completely omitted this aspect of programming.

Arrays do exist. However, they are primitive types, not objects. An array is declared with the keyword `new` followed by the type of the array elements and the dimension of the array enclosed in square brackets. For example

```
Cluster[] clusterArray = new Cluster[10];
```

declares an array of 10 objects of the class `Cluster`. The reader should note that the type of an array is the type of the element followed by `[]`. In our example the type of the variable `clusterArray` is `Cluster[]`. All elements of an array must be of the same type. Iterating over the elements of an array must be done using index in a `for` or `while` loop.

To be complete, let us mention a few Java constructs which provide more flexible collection behavior. They are not used in this book, but I do not want to leave the Smalltalkers on a bad impression of the Java language.

The Java class `Vector` corresponds to the Smalltalk class `OrderedCollection`. The elements of a `Vector` are of type `Object`. Its elements do not need to be of the same type theoretically. In practice, however, handling a multi-type `Vector` is difficult as one needs to cast the element to the proper type when extracting them from the `Vector`.

The Smalltalk class `Dictionary` corresponds to the Java class `HashTable`.

An equivalent to the Smalltalk `do: iterator` method is provided by the `Enumeration` interface to which both `Vector` and `HashTable` are adhering. Each element of the collection is retrieved under the type `Object`. However, this forces the designer to cast each elements to its proper type before using it within the iterator method.

Missing file

Appendix D

Additional probability distributions

D.1 Beta distribution

Table D.1 shows the properties of the beta distribution. If the parameters α_1 and α_2 are equal to 1 the beta distribution is identical to the uniform distribution over the interval $[0, 1]$.

The beta distribution is an ad-hoc distribution, which can take many shapes. Figure D.1 shows a few characteristic shapes for several values of the parameters. The random variable is limited to the interval $[0, 1]$. However, any distribution of a random variable varying over a finite interval can be mapped to a beta distribution with a suitable variable transformation.

D.1.1 Beta distribution — Smalltalk implementation

Listing D.1 shows the implementation of the beta distribution in Smalltalk.

Listing D.1 Smalltalk implementation of the beta distribution

```
Class                DhbBetaDistribution
Subclass of         DhbProbabilityDensity
Instance variable names:  alpha1 alpha2 gamma1 gamma2 logNorm incompleteBetaFunction

Class methods
distributionName
    ^'Beta distribution'

fromHistogram: aHistogram
```

Table D.1: Properties of the beta distribution

Range of random variable	$[0, 1]$
Probability density function	$P(x) = \frac{1}{B(\alpha_1, \alpha_2)} x^{\alpha_1-1} (1-x)^{\alpha_2-1}$
Parameters	$0 < \alpha_1 < +\infty$ $0 < \alpha_2 < +\infty$
Distribution function	$F(x) = B(x; \alpha_1, \alpha_2)$ (c.f. section 7.5)
Average	$\frac{\alpha_1}{\alpha_1 + \alpha_2}$
Variance	$\frac{\alpha_1 \alpha_2}{(\alpha_1 + \alpha_2)^2 (\alpha_1 + \alpha_2 + 1)}$
Skewness	$2 \frac{\alpha_1 - \alpha_2}{(\alpha_1 + \alpha_2 + 2)}$
Kurtosis	$\sqrt[3]{\frac{\alpha_1 + \alpha_2 + 2}{\alpha_1 \alpha_2}} \left\{ \frac{(\alpha_1 + \alpha_2 + 1)}{\alpha_1 \alpha_2 (\alpha_1 + \alpha_2 + 2) (\alpha_1 + \alpha_2 + 3)} \right.$ $\left. \times \left[2 (\alpha_1 + \alpha_2)^2 + \alpha_1 \alpha_2 (\alpha_1 + \alpha_2 - 6) \right] - 1 \right\}$

```

| average variance a b c discr |
( aHistogram minimum < 0 or: [ aHistogram maximum > 1])
  ifTrue: [ ^nil].
average := aHistogram average.
variance := aHistogram variance.
a := ( ( 1 - average) / variance - 1) * average.
a > 0
  ifFalse:[ ^nil].
b := ( 1 / average - 1) * a.
b > 0
  ifFalse:[ ^nil].
^self shape: a shape: b

```

```
new
```

```
^self error: 'Illegal creation message for this class'
```

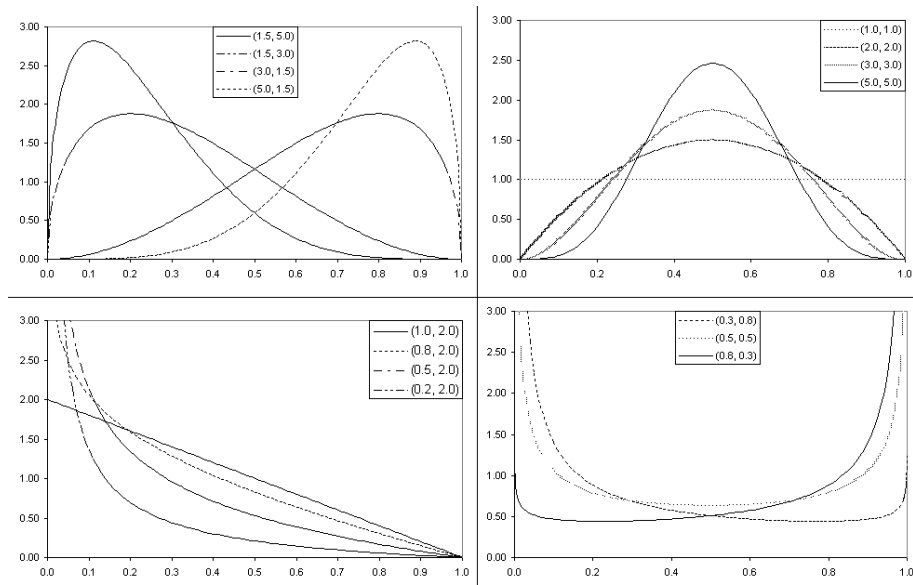


Figure D.1: Many shapes of the beta distribution

shape: aNumber1 shape: aNumber2

$\hat{\text{super new initialize: aNumber1 shape: aNumber2}}$

Instance methods

average

$\hat{\alpha}1 / (\alpha1 + \alpha2)$

changeParametersBy: aVector

$\alpha1 := \alpha1 + (\text{aVector at: } 1).$

$\alpha2 := \alpha2 + (\text{aVector at: } 2).$

self computeNorm.

gamma1 := nil.

gamma2 := nil.

incompleteBetaFunction := nil.

computeNorm

$\log\text{Norm} := (\alpha1 + \alpha2) \log\Gamma - \alpha1 \log\Gamma - \alpha2 \log\Gamma.$

distributionValue: aNumber

```
incompleteBetaFunction isNil
  ifTrue: [ incompleteBetaFunction := DhbIncompleteBetaFunction
    shape: alpha1 shape: alpha2].
^incompleteBetaFunction value: aNumber
```

firstGammaDistribution

```
gamma1 isNil
  ifTrue: [ gamma1 := DhbGammaDistribution shape: alpha1 scale:
    1].
^gamma1
```

initialize: aNumber1 shape: aNumber2

```
(aNumber1 > 0 and: [aNumber2 > 0])
  ifFalse: [self error: 'Illegal distribution parameters'].
alpha1 := aNumber1.
alpha2 := aNumber2.
self computeNorm.
^self
```

kurtosis

```
^3 * ( alpha1 + alpha2 + 1) * ((alpha1 + alpha2) squared * 2 + (
  (alpha1 + alpha2 - 6) * alpha1 * alpha2)
  / ((alpha1 + alpha2 + 2) * (alpha1 + alpha2 + 3) *
    alpha1 * alpha2)) - 3
```

parameters

```
^Array with: alpha1 with: alpha2
```

random

```
| r |
r := self firstGammaDistribution random.
^r / ( self secondGammaDistribution random + r)
```

secondGammaDistribution

```
gamma2 isNil
  ifTrue: [ gamma2 := DhbGammaDistribution shape: alpha2 scale:
    1].
^gamma2
```

skewness

$$\frac{(\alpha_1 + \alpha_2 + 1) \sqrt{2} (\alpha_2 - \alpha_1)}{(\alpha_1 \alpha_2) \sqrt{(\alpha_1 + \alpha_2 + 2)}}$$

value: aNumber

```
^(aNumber > 0 and: [ aNumber < 1])
  ifTrue:
    [( ( aNumber ln * (alpha1 - 1) ) + ( ( 1 - aNumber) ln *
      ( alpha2 - 1)) + logNorm) exp]
  ifFalse: [0]
```

variance

$$\frac{\alpha_1 * \alpha_2}{((\alpha_1 + \alpha_2) \text{ squared} * (\alpha_1 + \alpha_2 + 1))}$$

D.1.2 Beta distribution — Java implementation

Listing D.2 shows the implementation of the beta distribution in Java.

Listing D.2 Java implementation of the beta distribution

```
package DhbStatistics;

import DhbIterations.IncompleteBetaFunction;
import DhbFunctionEvaluation.GammaFunction;
import DhbScientificCurves.Histogram;
import DhbInterfaces.ParametrizedOneVariableFunction;
/**
 * Beta distribution
 *
 * @author Didier H. Besset
 */
public final class BetaDistribution
    extends ProbabilityDensityFunction
{
    /**
     * First shape parameter of the distribution.
     */
    protected double alpha1;
    /**
```

```

    * Second shape parameter of the distribution.
    */
    private double alpha2;
    /**
     * Norm of the distribution (cached for efficiency).
     */
    private double norm;
    /**
     * Gamma distribution for alpha1 used for random generation (cached for efficiency).
     */
    private GammaDistribution gamma1;
    /**
     * Gamma distribution for alpha2 used for random generation (cached for efficiency).
     */
    private GammaDistribution gamma2;
    /**
     * Incomplete beta function for the distribution (cached for efficiency).
     */
    private IncompleteBetaFunction incompleteBetaFunction;
    /**
     * Create a new instance of the Beta distribution with given shape and scale.
     * @param shape1 double first shape parameter of the distribution (alpha1).
     * @param shape2 double second shape parameter of the distribution (alpha2).
     * @exception java.lang.IllegalArgumentException
     *         if the parameters of the distribution are illegal.
     */
    public BetaDistribution ( double shape1, double shape2) throws IllegalArgumentException
    {
        if ( shape1 <= 0 )
            throw new IllegalArgumentException(
                "First shape parameter must be positive");
        if ( shape2 <= 0 )
            throw new IllegalArgumentException(
                "Second shape parameter must be positive");
        defineParameters( shape1, shape2);
    }
    /**
     * Create an instance of the receiver with parameters estimated from
     * the given histogram using best guesses. This method can be used to
     * find the initial values for a fit.
     * @param h DhbScientificCurves.Histogram
     * @exception java.lang.IllegalArgumentException
     *         when no suitable parameter can be found.
     */
    public BetaDistribution( Histogram h) throws IllegalArgumentException
    {

```

```

    if ( h.getMinimum() < 0 || h.getMaximum() > 1)
        throw new IllegalArgumentException(
            "Beta distribution is only defined over [0,1]");
    double average = h.average();
    double variance = h.variance();
    double a = ( ( 1 - average) / variance - 1) * average;
    if ( a <= 0 )
        throw new IllegalArgumentException("Negative shape parameter");
    double b = ( 1 /average - 1) * a;
    if ( b <= 0 )
        throw new IllegalArgumentException("Negative shape parameter");
    defineParameters( a, b);
}
/**
 * @return double average of the distribution.
 */
public double average()
{
    return alpha1 / ( alpha1 + alpha2);
}
/**
 * Assigns new values to the parameters.
 * This method assumes that the parameters have been already checked.
 */
private void defineParameters ( double shape1, double shape2)
{
    alpha1 = shape1;
    alpha2 = shape2;
    norm = GammaFunction.logBeta( alpha1, alpha2);
    gamma1 = null;
    gamma2 = null;
    incompleteBetaFunction = null;
    return;
}
private void defineRandomGenerator ( )
{
    gamma1 = new GammaDistribution( alpha1, 1.0);
    gamma2 = new GammaDistribution( alpha2, 1.0);
    return;
}
/**
 * Returns the probability of finding a random variable smaller
 * than or equal to x.
 * @return integral of the probability density function from 0 to x.
 * @param x double upper limit of integral.
 */

```

```

public double distributionValue ( double x)
{
    return incompleteBetaFunction().value(x);
}
/**
 * @return DhbIterations.IncompleteBetaFunction
 */
private IncompleteBetaFunction incompleteBetaFunction()
{
    if ( incompleteBetaFunction == null )
        incompleteBetaFunction = new IncompleteBetaFunction( alpha1,
                                                                alpha2);
    return incompleteBetaFunction;
}
/**
 * @return double kurtosis of the distribution.
 */
public double kurtosis( )
{
    double s = alpha1 + alpha2;
    return 3 * ( alpha1 + alpha2 + 1) * ( 2 * s * s +
        ( alpha1 + alpha2 - 6) * alpha1 * alpha2)
        / ( ( alpha1 + alpha2 + 2) *
            ( alpha1 + alpha2 + 3) * alpha1 * alpha2)
        - 3;
}
/**
 * @return java.lang.String name of the distribution.
 */
public String name()
{
    return "Beta distribution";
}
/**
 * @return double[] an array containing the parameters of
 * the distribution.
 */
public double[] parameters()
{
    double[] answer = new double[2];
    answer[0] = alpha1;
    answer[1] = alpha2;
    return answer;
}
/**
 * @return double a random number distributed according to the receiver.

```

```

*/
public double random ( )
{
    if ( gamma1 == null )
        defineRandomGenerator();
    double y1 = gamma1.random();
    return y1 / ( y1 + gamma2.random());
}
/**
 * @param a1 double
 */
public void setAlpha1( double a1)
{
    defineParameters( a1, alpha2);
}
/**
 * @param a2 double
 */
public void setAlpha2( double a2)
{
    defineParameters( alpha1, a2);
}
/**
 * @param p double[]  assigns the parameters
 */
public void setParameters( double[] params)
{
    defineParameters( params[0], params[1]);
}
/**
 * @return double skewness of the distribution.
 */
public double skewness( )
{
    return 2 * Math.sqrt( alpha1 + alpha2 + 1) * (alpha2 - alpha1)
           / ( Math.sqrt( alpha1 * alpha2)
              * ( alpha1 + alpha2 + 2));
}
/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    java.text.DecimalFormat fmt = new java.text.DecimalFormat("0.00000");
    sb.append("Beta distribution (");

```

```

        sb.append(fmt.format(alpha1));
        sb.append(', ');
        sb.append(fmt.format(alpha2));
        sb.append(' ');
        return sb.toString();
    }
    /**
     * @return double probability density function
     * @param x double random variable
     */
    public double value( double x)
    {
        return Math.exp( Math.log( x) * ( alpha1 - 1)
                        + Math.log( 1 - x) * ( alpha2 - 1) - norm);
    }
    /**
     * @return double variance of the distribution.
     */
    public double variance( )
    {
        double s = alpha1 + alpha2;
        return alpha1 * alpha2 / ( s * s * ( alpha1 + alpha2 + 1));
    }
}

```

D.2 Cauchy distribution

Table D.2 shows the properties of the Cauchy distribution. Physicists use the Cauchy distribution under the name Breit-Wigner or resonance curve. All moments of order greater than 0 are not defined as the corresponding integrals diverge.

Figure D.2 shows the shapes taken by the Cauchy distribution for a few values of the parameters. These parameter are identical to the parameters of the normal distributions shown in figure 9.3 so that the reader can compare them.

D.2.1 Cauchy distribution — Smalltalk implementation

Listing D.3 shows the implementation of the Cauchy distribution in Smalltalk.

This implementation returns μ for the average although the average is not defined mathematically. Other moment related quantities are returning nil.

Listing D.3 Smalltalk implementation of the Cauchy distribution

Class	DhbCauchyDistribution
Subclass of	DhbProbabilityDensity

Table D.2: Properties of the Cauchy distribution

Range of random variable	$]-\infty, +\infty[$
Probability density function	$P(x) = \frac{\beta}{\pi \left[(x - \mu)^2 + \beta^2 \right]}$
Parameters	$-\infty < \mu < +\infty$ $0 < \beta < +\infty$
Distribution function	$F(x) = \frac{1}{2} + \frac{1}{\pi} \arctan \left(\frac{x - \mu}{\beta} \right)$
Average	(undefined)
Variance	(undefined)
Skewness	(undefined)
Kurtosis	(undefined)

Instance variable names: mu beta

Class methods

distributionName

^^Cauchy distribution'

fromHistogram: aHistogram

^self shape: aHistogram average
scale: 4 * aHistogram variance
/ (Float pi * (aHistogram maximum squared +
aHistogram minimum squared))
sqrt

new

^self shape: 0 scale: 1

shape: aNumber1 scale: aNumber2

^super new initialize: aNumber1 scale: aNumber2

Instance methods

acceptanceBetween: aNumber1 and: aNumber2




Figure D.2: Cauchy distribution for a few parameters

`^self privateAcceptanceBetween: aNumber1 and: aNumber2`

`average`

`^mu`

`changeParametersBy: aVector`

`mu := mu + (aVector at: 1).`
`beta := beta + (aVector at: 2).`

`distributionValue: aNumber`

`^((aNumber - mu) / beta) arcTan / Float pi + (1 / 2)`

`initialize: aNumber1 scale: aNumber2`

`mu := aNumber1.`
`beta := aNumber2.`
`^self`

`parameters`

```

    ^Array with: mu with: beta

privateInverseDistributionValue: aNumber
    ^(( aNumber - (1 / 2)) * Float pi) tan * beta + mu

standardDeviation
    ^nil

value: aNumber
    ^beta / ( Float pi * ( beta squared + ( aNumber - mu) squared))

valueAndGradient: aNumber
    | dp denominator |
    dp := self value: aNumber.
    denominator := 1 / ( ( aNumber - mu) squared + beta squared).
    ^Array with: dp
        with: ( DhbVector with: 2 * dp * ( aNumber - mu) *
            denominator
            with: dp * ( 1 / beta - ( 2 * beta *
                denominator)))

variance
    ^nil

```

D.2.2 Cauchy distribution — Java implementation

Listing D.4 shows the implementation of the Cauchy distribution in Java.

This implementation returns μ for the average although the average is not defined mathematically. Other moment related quantities are returning the special value `Double.NaN`.

Listing D.4 Java implementation of the Cauchy distribution

```

package DhbStatistics;

import DhbScientificCurves.Histogram;
import DhbInterfaces.ParametrizedOneVariableFunction;
/**

```

```

* Cauchy distribution
*
* @author Didier H. Besset
*/
public final class CauchyDistribution
    extends ProbabilityDensityFunction
{
    /**
     * Center of the distribution.
     */
    private double mu;
    /**
     * Scale of the distribution.
     */
    private double beta;

    /**
     * Create an instance centered at 0 with width 1.
     */
    public CauchyDistribution ( )
    {
        this( 0, 1);
    }
    /**
     * @param middle double middle point of the distribution.
     * @param width double width of the distribution.
     */
    public CauchyDistribution ( double middle, double width)
    {
        mu = middle;
        beta = width;
    }
    /**
     * Create an instance of the receiver with parameters estimated from
     * the given histogram using best guesses. This method can be used to
     * find the initial values for a fit.
     * @param h Histogram
     */
    public CauchyDistribution( Histogram h)
    {
        this( h.average(),
            4 * h.variance() /Math.sqrt(Math.PI *( h.getMinimum()
                * h.getMinimum()+h.getMaximum() * h.getMaximum()))
            );
    }
    /**

```

```

    * @return double average of the distribution.
    */
    public double average()
    {
        return mu;
    }
    /**
    * Returns the probability of finding a random variable smaller
    * than or equal to x.
    * @return integral of the probability density function from -infinity to x.
    * @param x double upper limit of integral.
    */
    public double distributionValue ( double x)
    {
        return Math.atan( ( x - mu) / beta) / Math.PI + 0.5;
    }
    /**
    * @return java.lang.String    name of the distribution
    */
    public String name ( )
    {
        return "Cauchy distribution";
    }
    /**
    * @return double[] an array containing the parameters of
    *                  the distribution.
    */
    public double[] parameters ( )
    {
        double[] answer = new double[2];
        answer[0] = mu;
        answer[1] = beta;
        return answer;
    }
    /**
    * This method assumes that the range of the argument has been checked.
    * @return double the value for which the distribution function
    *                is equal to x.
    * @param x double value of the distribution function.
    */
    protected double privateInverseDistributionValue ( double x)
    {
        return Math.tan( (x - 0.5) * Math.PI) * beta + mu;
    }
    /**
    * @param center double

```

```
*/
public void setBeta( double width)
{
    beta = width;
}
/**
 * @param center double
 */
public void setMu( double center)
{
    mu = center;
}
/**
 * @param p double[]  assigns the parameters
 */
public void setParameters( double[] params)
{
    setMu( params[0]);
    setBeta( params[1]);
}
/**
 * @return NaN since the standard deviation of the distribution is
 *                               not defined.
 */
public double standardDeviation( )
{
    return Double.NaN;
}
/**
 * This method was created in VisualAge.
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    java.text.DecimalFormat fmt = new java.text.DecimalFormat(
        "####0.00000");
    sb.append("Cauchy distribution (");
    sb.append(fmt.format(mu));
    sb.append(',');
    sb.append(fmt.format(beta));
    sb.append(')');
    return sb.toString();
}
/**
 * @return double probability density function
```

```

* @param x double random variable
*/
public double value( double x)
{
    double dev = x - mu;
    return beta / ( Math.PI *( beta * beta + dev * dev));
}
/**
 * Evaluate the distribution and the gradient of the distribution with respect
 * to the parameters.
 * @return double[] 0: distribution's value, 1,2,...,n distribution's gradient
 * @param x double
 */
public double[] valueAndGradient( double x)
{
    double[] answer = new double[3];
    answer[0] = value( x);
    double y = x - mu;
    double r = 1 / ( y * y + beta * beta);
    answer[1] = 2 * answer[0] * y * r;
    answer[2] = answer[0] * ( 1 / beta - 2 * beta * r);
    return answer;
}
}

```

D.3 Exponential distribution

Table D.3 shows the properties of the exponential distribution.

The exponential distribution describes the distribution of the time of occurrence between independent random events with a constant probability of occurrence. It is used in queuing theory and in nuclear physics. Figure D.3 shows the shapes taken by the exponential distribution for a few values of the parameters.

D.3.1 Exponential distribution — Smalltalk implementation

Listing D.5 shows the implementation of the exponential distribution in Smalltalk.

Listing D.5 Smalltalk implementation of the exponential distribution

Class	DhbExponentialDistribution
Subclass of	DhbProbabilityDensity
Instance variable names:	beta

Class methods

Table D.3: Properties of the exponential distribution

Range of random variable	$[0, +\infty[$
Probability density function	$P(x) = \frac{1}{\beta} e^{-\frac{x}{\beta}}$
Parameters	$0 < \beta < +\infty$
Distribution function	$F(x) = 1 - e^{-\frac{x}{\beta}}$
Average	β
Variance	β^2
Skewness	2
Kurtosis	6

distributionName

^Exponential distribution'

fromHistogram: aHistogram

| mu |

aHistogram minimum < 0

ifTrue: [^nil].

mu := aHistogram average.

^mu > 0 ifTrue: [self scale: aHistogram average]

ifFalse:[nil]

new

^super new initialize: 1

scale: aNumber

^super new initialize: aNumber

Instance methods

acceptanceBetween: aNumber1 and: aNumber2

^self privateAcceptanceBetween: aNumber1 and: aNumber2




Figure D.3: Exponential distribution for a few parameters

average

$\hat{\text{beta}}$

changeParametersBy: aVector

beta := beta + (aVector at: 1).

distributionValue: aNumber

$[1 - ((\text{aNumber} / \text{beta negated}) \exp)]$
when: ExAll do: [:signal | signal exitWith: 0]

initialize: aNumber

aNumber > 0
ifFalse: [self error: 'Illegal distribution parameters'].
beta := aNumber.
 $\hat{\text{self}}$

kurtosis

$\hat{6}$

parameters

$\hat{\text{Array}}$ with: beta

privateInverseDistributionValue: aNumber

$\hat{(1 - \text{aNumber}) \ln \text{negated} * \text{beta}}$

random

$\hat{\text{DhbMitchellMooreGenerator new floatValue ln} * \text{beta negated}}$

skewness

$\hat{2}$

standardDeviation

$\hat{\text{beta}}$

value: aNumber

$\hat{[(\text{aNumber} / \text{beta}) \text{negated exp} / \text{beta}]}$
 when: ExAll do: [:signal | signal exitWith: 0]

valueAndGradient: aNumber

| dp |
 dp := self value: aNumber.
 $\hat{\text{Array}}$ with: dp
 with: (DhbVector with: (aNumber / beta - 1) * dp / beta)

D.3.2 Exponential distribution — Java implementation

Listing D.6 shows the implementation of the exponential distribution in Java.

Listing D.6 Java implementation of the exponential distribution

package DhbStatistics;

import DhbScientificCurves.Histogram;
 import DhbInterfaces.ParametrizedOneVariableFunction;

```
/**
 * Exponential distribution.
 *
 * @author Didier H. Besset
 */
public final class ExponentialDistribution
    extends ProbabilityDensityFunction
{
    /**
     * Exponential term.
     */
    private double beta;

    /**
     * General constructor method.
     * @param exponential fall-off
     * @exception java.lang.IllegalArgumentException
     *             if the fall-off is non-positive.
     */
    public ExponentialDistribution ( double fallOff)
        throws IllegalArgumentException
    {
        if ( fallOff <= 0 )
            throw new IllegalArgumentException(
                "Exponential fall-off must be positive");
        beta = fallOff;
    }

    /**
     * Create an instance of the receiver with parameters estimated from
     * the given histogram using best guesses. This method can be used to
     * find the initial values for a fit.
     * @param h DhbScientificCurves.Histogram
     * @exception java.lang.IllegalArgumentException
     *             when no suitable parameter can be found.
     */
    public ExponentialDistribution( Histogram h)
        throws IllegalArgumentException
    {
        if ( h.getMinimum() < 0 )
            throw new IllegalArgumentException(
                "Exponential distribution is only defined for non-negative values");
        double average = h.average();
        if ( h.average() < 0 )
            throw new IllegalArgumentException(
                "Exponential distribution is only defined for positive scale");
        setScale( average);
    }
}
```

```
}
/**
 * @return double average of the distribution.
 */
public double average ( )
{
    return beta;
}
/**
 * Returns the probability of finding a random variable smaller
 * than or equal to x.
 * @return integral of the probability density function from 0 to x.
 * @param x double upper limit of integral.
 */
public double distributionValue ( double x)
{
    return 1 - Math.exp( -x / beta);
}
/**
 * @return double kurtosis of the distribution.
 */
public double kurtosis ( )
{
    return 6;
}
/**
 * @return java.lang.String    name of the distribution
 */
public String name ( )
{
    return "Exponential distribution";
}
/**
 * @return double[] an array containing the parameters of
 *                  the distribution.
 */
public double[] parameters ( )
{
    double[] answer = new double[1];
    answer[0] = beta;
    return answer;
}
/**
 * This method assumes that the range of the argument has been checked.
 * @return double the value for which the distribution function
 *                is equal to x.
```

```
* @param x double value of the distribution function.
*/
protected double privateInverseDistributionValue ( double x)
{
    return -Math.log( 1 - x) * beta;
}
/**
 * @return double a random number distributed according to the receiver.
 */
public double random( )
{
    return -beta * Math.log(generator().nextDouble());
}
/**
 * @param p double[] assigns the parameters
 */
public void setParameters( double[] params)
{
    setScale( params[0]);
}
/**
 * @param falloff double
 */
public void setScale( double falloff)
{
    beta = falloff;
}
/**
 * @return double skewness of the distribution.
 */
public double skewness( )
{
    return 2;
}
/**
 * @return double standard deviation of the distribution
 */
public double standardDeviation( )
{
    return beta;
}
/**
 * @return java.lang.String
 */
public String toString()
{

```

```

    StringBuffer sb = new StringBuffer();
    java.text.DecimalFormat fmt = new java.text.DecimalFormat(
        "###0.00000");

    sb.append("Exponential distribution (");
    sb.append(fmt.format(beta));
    sb.append(')');
    return sb.toString();
}
/**
 * @return double probability density function
 * @param x double random variable
 */
public double value( double x)
{
    return Math.exp( -x / beta) / beta;
}
/**
 * Evaluate the distribution and the gradient of the distribution with respect
 * to the parameters.
 * @return double[] 0: distribution's value, 1,2,...,n distribution's gradient
 * @param x double
 */
public double[] valueAndGradient( double x)
{
    double[] answer = new double[2];
    answer[0] = value(x);
    answer[1] = ( x / beta - 1) * answer[0] / beta;
    return answer;
}
}

```

D.4 Fisher-Tippett distribution

Table D.4 shows the properties of the fishertippett distribution. In this table $\gamma = 0.5772156649 \dots$ is the Euler constant.

The Fisher-Tippett distribution describes the distribution of extreme values. Figure D.4 shows the shapes taken by the Fisher-Tippett distribution for a few values of the parameters. These parameter are identical to the parameters of the normal distributions shown in figure 9.3 so that the reader can compare them.

D.4.1 Fisher-Tippett distribution — Smalltalk implementation

Listing D.7 shows the implementation of the Fisher-Tippett distribution in Smalltalk.

Table D.4: Properties of the Fisher-Tippett distribution

Range of random variable	$]-\infty, +\infty[$
Probability density function	$P(x) = \frac{1}{\beta} e^{-\frac{x-\alpha}{\beta}} - e^{-\frac{x-\alpha}{\beta}}$
Parameters	$-\infty < \alpha < +\infty$ $0 < \beta < +\infty$
Distribution function	$F(x) = e^{-e^{-\frac{x-\alpha}{\beta}}}$
Average	$\alpha + \gamma\beta$
Variance	$\frac{\pi\beta}{\sqrt{6}}$
Skewness	1.3
Kurtosis	2.4

Listing D.7 Smalltalk implementation of the Fisher-Tippett distribution

```

Class                DhbFisherTippettDistribution
Subclass of          DhbProbabilityDensity
Instance variable names:  alpha beta

    Class methods
distributionName
    ^'Fisher-Tippett distribution'

fromHistogram: aHistogram
    | beta |
    beta := aHistogram standardDeviation.
    beta = 0 ifTrue: [^nil].
    beta := beta * (6 sqrt / Float pi).
    ^self shape: aHistogram average - (0.5772156649 * beta) scale:
        beta

new
    ^self shape: 0 scale: 1

```




Figure D.4: Fisher-Tippett distribution for a few parameters

shape: aNumber1 scale: aNumber2

^{super new initialize:} aNumber1 scale: aNumber2

Instance methods

average

^{0.577256649} * beta + alpha

changeParametersBy: aVector

alpha := alpha + (aVector at: 1).

beta := beta + (aVector at: 2).

distributionValue: aNumber

| arg |

arg := (aNumber - alpha) / beta.

arg := arg < DhhFloatingPointMachine new largestExponentArgument
negated

ifTrue: [^0]

ifFalse:[arg negated exp].

^{arg} > DhhFloatingPointMachine new largestExponentArgument
ifTrue: [1]

```
ifFalse:[ arg negated exp]
```

```
initialize: aNumber1 scale: aNumber2
```

```
aNumber2 > 0
  ifFalse: [ self error: 'Illegal distribution parameters'].
alpha := aNumber1.
beta := aNumber2.
^self
```

```
integralFrom: aNumber1 to: aNumber2
```

```
^( DhbRombergIntegrator new: self from: aNumber1 to: aNumber2)
  evaluate
```

```
integralUpTo: aNumber
```

```
^( DhbRombergIntegrator new:
  [ :x | x = 0 ifTrue: [ 0] ifFalse: [ ( self value: 1 / x)
    / x squared] ]
  from: 1 / aNumber to: 0) evaluate
```

```
kurtosis
```

```
^2.4
```

```
parameters
```

```
^Array with: alpha with: beta
```

```
random
```

```
| t |
[ t := DhbMitchellMooreGenerator new floatValue ln negated.
  t > 0] whileFalse: [].
^t ln negated * beta + alpha
```

```
skewness
```

```
^1.3
```

```
standardDeviation
```



```

^Float pi * beta / ( 6 sqrt)

value: aNumber

| arg |
arg := ( aNumber - alpha ) / beta.
arg := arg > DhbFloatingPointMachine new largestExponentArgument
      ifTrue: [ ^0]
      ifFalse:[arg negated exp + arg].
^arg > DhbFloatingPointMachine new largestExponentArgument
      ifTrue: [ 0]
      ifFalse:[ arg negated exp / beta]

valueAndGradient: aNumber

| dp dy y |
dp := self value: aNumber.
y := ( aNumber - alpha ) / beta.
dy := ( y negated exp - 1).
^Array with: dp
      with: ( DhbVector with: dy * dp / beta negated
                  with: dp * ( y * dy + 1 ) / beta negated)

```

D.4.2 Fisher-Tippett distribution — Java implementation

Listing D.8 shows the implementation of the Fisher-Tippett distribution in Java.

Listing D.8 Java implementation of the Fisher-Tippett distribution

```

package DhbStatistics;

import DhbFunctionEvaluation.DhbMath;
import DhbScientificCurves.Histogram;
import DhbInterfaces.ParametrizedOneVariableFunction;
/**
 * Fisher-Tippett distribution
 *
 * @author Didier H. Besset
 */
public final class FisherTippettDistribution
    extends ProbabilityDensityFunction
{
    /**

```

```

    * Center of the distribution.
    */
    protected double alpha;
    /**
     * Scale parameter of the distribution.
     */
    private double beta;
    /**
     * Constructor method
     * @param center double
     * @param scale double
     * @exception java.lang.IllegalArgumentException if the scale parameter is non-positive.
     */
    public FisherTippettDistribution( double center, double scale) throws IllegalArgumentException
    {
        if ( scale <= 0 )
            throw new IllegalArgumentException("Scale parameter must be positive");
        alpha = center;
        beta = scale;
    }
    /**
     * Create an instance of the receiver with parameters estimated from the
     * given histogram using best guesses. This method can be used to
     * find the initial values for a fit.
     * @param h DhbScientificCurves.Histogram
     * @exception java.lang.IllegalArgumentException when no suitable parameter can be found.
     */
    public FisherTippettDistribution( Histogram h) throws IllegalArgumentException
    {
        double beta = h.standardDeviation();
        if ( beta < 0 )
            throw new IllegalArgumentException("Histogram has vanishing standard deviation");
        beta *= Math.sqrt(6)/Math.PI;
        defineParameters( h.average() - 0.5772156649 * beta, beta);
    }
    /**
     * @return double average of the distribution.
     */
    public double average()
    {
        return 0.5772156649 * beta + alpha;
    }
    /**
     * @param center double
     * @param scale double
     */

```

```
public void defineParameters ( double center, double scale)
{
    alpha = center;
    beta = scale;
}
/**
 * Returns the probability of finding a random variable smaller
 * than or equal to x.
 * @return integral of the probability density function from -infinity to x.
 * @param x double upper limit of integral.
 */
public double distributionValue ( double x)
{
    double y = ( x - alpha ) / beta;
    if ( y < -DhbMath.getLargestExponentialArgument() )
        return 0;
    y = Math.exp( -y);
    if ( y > DhbMath.getLargestExponentialArgument() )
        return 1;
    return Math.exp( -y);
}
/**
 * @return double kurtosis of the distribution.
 */
public double kurtosis( )
{
    return 2.4;
}
/**
 * @return java.lang.String name of the distribution
 */
public String name ( )
{
    return "Fisher-Tippett distribution";
}
/**
 * @return double[] an array containing the parameters of
 * the distribution.
 */
public double[] parameters ( )
{
    double[] answer = new double[2];
    answer[0] = alpha;
    answer[1] = beta;
    return answer;
}
```

```

/**
 * @return double a random number distributed according to the receiver.
 */
public double random( )
{
    double t;
    while ( ( t = -Math.log( generator().nextDouble())) == 0 );
    return alpha - beta * Math.log( t);
}
/**
 * @param p double[]  assigns the parameters
 */
public void setParameters( double[] params)
{
    defineParameters ( params[0], params[1]);
}
/**
 * @return double skewness of the distribution.
 */
public double skewness( )
{
    return 1.3;
}
/**
 * @return double standard deviation of the distribution
 */
public double standardDeviation( )
{
    return Math.PI * beta / Math.sqrt( 6);
}
/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    java.text.DecimalFormat fmt = new java.text.DecimalFormat("####0.00000");
    sb.append("Fisher-Tippett distribution (");
    sb.append(fmt.format(alpha));
    sb.append(',');
    sb.append(fmt.format(beta));
    sb.append(')');
    return sb.toString();
}
/**
 * @return double probability density function

```

```

* @param x double random variable
*/
public double value( double x)
{
    double y = ( x - alpha) / beta;
    if ( y > DhbMath.getLargestExponentialArgument() )
        return 0;
    y += Math.exp( -y);
    if ( y > DhbMath.getLargestExponentialArgument() )
        return 0;
    return Math.exp( -y) / beta;
}
/**
 * Evaluate the distribution and the gradient of the distribution with respect
 * to the parameters.
 * @return double[] 0: distribution's value, 1,2,...,n distribution's gradient
 * @param x double
 */
public double[] valueAndGradient( double x)
{
    double[] answer = new double[3];
    answer[0] = value( x);
    double y = ( x - alpha) / beta;
    double dy = Math.exp( -y) - 1;
    double r = -1 / beta;
    answer[1] = dy * answer[0] * r;
    answer[2] = answer[0] * ( y * dy + 1) * r;
    return answer;
}
}

```

D.5 Laplace distribution

Table D.5 shows the properties of the Laplace distribution. The Laplace distribution is an ad-hoc distribution made of two exponential distributions, one on each side of the peak. Figure D.5 shows the shapes taken by the Laplace distribution for a few values of the parameters. These parameter are identical to the parameters of the normal distributions shown in figure 9.3 so that the reader can compare them.

D.5.1 Laplace distribution — Smalltalk implementation

Listing D.9 shows the implementation of the Laplace distribution in Smalltalk.

Table D.5: Properties of the Laplace distribution

Range of random variable	$] -\infty, +\infty[$
Probability density function	$P(x) = \frac{1}{2\beta} e^{-\frac{ x-\alpha }{\beta}}$
Parameters	$-\infty < \alpha < +\infty$ $0 < \beta < +\infty$
Distribution function	$F(x) = \begin{cases} \frac{1}{2} e^{-\frac{\alpha-x}{\beta}} & \text{for } x < \alpha \\ 1 - \frac{1}{2} e^{-\frac{x-\alpha}{\beta}} & \text{for } x \geq \alpha \end{cases}$
Average	$\alpha + \beta$
Variance	$2\beta^2$
Skewness	0
Kurtosis	3

Listing D.9 Smalltalk implementation of the Laplace distribution

```
Class                DhbLaplaceDistribution
Subclass of         DhbProbabilityDensity
Instance variable names: mu beta

Class methods
distributionName
    ^'Laplace distribution'

fromHistogram: aHistogram
    ^self shape: aHistogram average scale: (aHistogram variance / 2)
        sqrt

new
    ^self shape: 0 scale: 1

shape: aNumber1 scale: aNumber2
```




Figure D.5: Laplace distribution for a few parameters

`^super new initialize: aNumber1 scale: aNumber2`

Instance methods

`average`

`^mu`

`changeParametersBy: aVector`

`mu := mu + (aVector at: 1).`
`beta := beta + (aVector at: 2).`

`distributionValue: aNumber`

`^aNumber > mu`
`ifTrue: [1 - ((aNumber - mu) / beta) negated exp / 2)]`
`ifFalse:[((aNumber - mu) / beta) exp / 2)]`

`initialize: aNumber1 scale: aNumber2`

`mu := aNumber1.`
`beta := aNumber2.`
`^self`

kurtosis

3

parameters

$^{\wedge}$ Array with: mu with: beta

random

```
| r |
r := DhbMitchellMooreGenerator new floatValue ln * beta negated.
 $^{\wedge}$ DhbMitchellMooreGenerator new floatValue > 0.5
  ifTrue: [ mu + r]
  ifFalse:[ mu - r]
```

skewness

0

standardDeviation

$^{\wedge}$ beta * (2 sqrt)

value: aNumber

$^{\wedge}$ ((aNumber - mu) / beta) abs negated exp / (2 * beta)

valueAndGradient: aNumber

```
| dp |
dp := self value: aNumber.
 $^{\wedge}$ Array with: dp
  with: ( DhbVector with: ( aNumber - mu) sign * dp / beta
    with: ( ( ( aNumber - mu) abs / beta -
      1) * dp / beta))
```

D.5.2 Laplace distribution — Java implementation

Listing D.10 shows the implementation of the Laplace distribution in Java.

Listing D.10 Java implementation of the Laplace distribution


```

package DhbStatistics;

import DhbScientificCurves.Histogram;
import DhbInterfaces.ParametrizedOneVariableFunction;
/**
 * Laplace distribution.
 *
 * @author Didier H. Besset
 */
public final class LaplaceDistribution
    extends ProbabilityDensityFunction
{
    /**
     * Average of the distribution.
     */
    private double mu;
    /**
     * Scale of the distribution.
     */
    private double beta;
    /**
     * Constructor method.
     * @param center double
     * @param scale double
     * @exception java.lang.IllegalArgumentException
     *             when the scale parameter is non-positive
     */
    public LaplaceDistribution( double center, double scale)
        throws IllegalArgumentException
    {
        if ( scale <= 0 )
            throw new IllegalArgumentException(
                "Scale parameter must be positive");
        mu = center;
        beta = scale;
    }
    /**
     * Create an instance of the receiver with parameters estimated from
     * the given histogram using best guesses. This method can be used to
     * find the initial values for a fit.
     * @param h DhbScientificCurves.Histogram
     */
    public LaplaceDistribution( Histogram h)
    {
        this( h.average(), Math.sqrt( 0.5 * h.variance()));
    }
}

```

```

/**
 * @return double average of the distribution.
 */
public double average()
{
    return mu;
}
/**
 * @param center double
 * @param scale double
 */
public void defineParameters ( double center, double scale)
{
    mu = center;
    beta = scale;
}
/**
 * Returns the probability of finding a random variable smaller
 * than or equal to x.
 * @return integral of the probability density function from -infinity to x.
 * @param x double upper limit of integral.
 */
public double distributionValue(double x)
{
    return x > mu
        ? 1 - Math.exp( -( x - mu) / beta) / 2
        : Math.exp( -( x - mu) / beta) / 2;
}
/**
 * @return double kurtosis of the distribution.
 */
public double kurtosis ( )
{
    return 3;
}
/**
 * @return java.lang.String name of the distribution.
 */
public String name() {
    return null;
}
/**
 * @return double[] an array containing the parameters of
 * the distribution.
 */
public double[] parameters()

```

```

{
    double[] answer = new double[2];
    answer[0] = mu;
    answer[1] = beta;
    return answer;
}
/**
 * This method assumes that the range of the argument has been checked.
 * @return double the value for which the distribution function
 *         is equal to x.
 * @param x double value of the distribution function.
 */
public double privateInverseDistributionValue ( double x)
{
    return x < 0.5
        ? mu + beta * Math.log( 2 * x)
        : mu - beta * Math.log( 2 - 2 * x);
}
/**
 * @return double a random number distributed according to the receiver.
 */
public double random( )
{
    double r = -beta * Math.log(generator().nextDouble());
    return generator().nextDouble() > 0.5 ? mu + r : mu - r;
}
/**
 * @param p double[] assigns the parameters
 */
public void setParameters( double[] params)
{
    defineParameters( params[0], params[1]);
}
/**
 * @return double skewness of the distribution.
 */
public double skewness( )
{
    return 0;
}
/**
 * @return double standard deviation of the distribution
 */
public double standardDeviation( )
{
    return beta / Math.sqrt( 2);
}

```

```

    }
    /**
     * @return java.lang.String
     */
    public String toString()
    {
        StringBuffer sb = new StringBuffer();
        java.text.DecimalFormat fmt = new java.text.DecimalFormat(
            "####0.00000");
        sb.append("Laplace distribution (");
        sb.append(fmt.format(mu));
        sb.append(', ');
        sb.append(fmt.format(beta));
        sb.append(')');
        return sb.toString();
    }
    /**
     * @return double probability density function
     * @param x double random variable
     */
    public double value( double x)
    {
        return Math.exp( -Math.abs( x - mu) / beta) / ( 2 * beta);
    }
    /**
     * Evaluate the distribution and the gradient of the distribution with respect
     * to the parameters.
     * @return double[] 0: distribution's value, 1,2,...,n distribution's gradient
     * @param x double
     */
    public double[] valueAndGradient( double x)
    {
        double[] answer = new double[3];
        answer[0] = value( x);
        double y = x - mu;
        if ( y >= 0 )
        {
            answer[1] = answer[0] / beta;
            answer[2] = (y / beta - 1) * answer[0] / beta;
        }
        else
        {
            answer[1] = - answer[0] / beta;
            answer[2] = - (y / beta + 1) * answer[0] / beta;
        }
        return answer;
    }

```

}
}

D.6 Log normal distribution

Table D.6 shows the properties of the log normal distribution. The log normal distribution

Table D.6: Properties of the log normal distribution

Range of random variable	$[0, +\infty[$
Probability density function	$P(x) = \frac{1}{x\sqrt{2\pi\sigma^2}} e^{-\frac{(\ln x - \mu)^2}{2\sigma^2}}$
Parameters	$-\infty < \mu < +\infty$ $0 < \sigma < +\infty$
Distribution function	(no closed expression)
Average	$e^{\mu + \frac{\sigma^2}{2}}$
Variance	$e^{2\mu + \sigma^2} (e^{\sigma^2} - 1)$
Skewness	$\sqrt{e^{\sigma^2} - 1} (e^{\sigma^2} + 2)$
Kurtosis	$e^{4\sigma^2} + 2e^{3\sigma^2} + 3e^{2\sigma^2} - 6$

is used to describe quantities that are the product of a large number of other quantities. It is an ad-hoc distribution whose shape is similar to that of gamma distributions with $\alpha > 1$. Figure D.6 shows the shapes taken by the log normal distribution for a few values of the parameters.

D.6.1 Log normal distribution — Smalltalk implementation

Listing D.11 shows the implementation of the log normal distribution in Smalltalk.

Listing D.11 Smalltalk implementation of the log normal distribution

```

Class                DhbLogNormalDistribution
Subclass of          DhbProbabilityDensityWithUnknownDistribution
Instance variable names: normalDistribution

```

Class methods




Figure D.6: Log normal distribution for a few parameters

distributionName

^'Log normal distribution'

fromHistogram: aHistogram

| average variance sigma2 |

aHistogram minimum < 0

ifTrue: [^nil].

average := aHistogram average.

average > 0

ifFalse: [^nil].

variance := aHistogram variance.

sigma2 := (variance / average squared + 1) ln.

sigma2 > 0

ifFalse: [^nil].

^self new: (average ln - (sigma2 * 0.5)) sigma: sigma2 sqrt

new

^self new: 0 sigma: 1

new: aNumber1 sigma: aNumber2

^super new initialize: aNumber1 sigma: aNumber2

Instance methods

average

$$\frac{1}{\sqrt{2\pi \text{normalDistribution variance}}} \exp\left(-\frac{\text{normalDistribution average}^2}{2 \text{normalDistribution variance}}\right)$$

changeParametersBy: aVector

normalDistribution changeParametersBy: aVector.

fourthCentralMoment

$$\frac{3}{8} \frac{\text{normalDistribution average}^4}{\text{normalDistribution variance}^2} + \frac{3}{2} \frac{\text{normalDistribution average}^2}{\text{normalDistribution variance}} + \frac{15}{8}$$

initialize: aNumber1 sigma: aNumber2

normalDistribution := DhbNormalDistribution new: aNumber1 sigma:
aNumber2.
^self

kurtosis

$$\frac{3}{5} \frac{\text{normalDistribution variance}^2}{\text{normalDistribution average}^4} + \frac{6}{5} \frac{\text{normalDistribution average}^2}{\text{normalDistribution variance}} + \frac{6}{5}$$

parameters

^normalDistribution parameters

random

^normalDistribution random exp

skewness

$$\frac{3}{2} \frac{\text{normalDistribution average}^3}{\text{normalDistribution variance}^{3/2}} + \frac{3}{2} \frac{\text{normalDistribution average}}{\text{normalDistribution variance}^{1/2}}$$

thirdCentralMoment

```

| y x |
y := normalDistribution average exp.
x := normalDistribution variance exp.
^( y squared * y) * ( x raisedTo: (3/2))
  * ( ( x squared negated + 3) * x - 2)

```

value: aNumber

```

^aNumber > 0
  ifTrue: [ ( normalDistribution value: aNumber ln) / aNumber]
  ifFalse:[ 0]

```

variance

```

^( normalDistribution average * 2 + normalDistribution variance)
  exp * ( normalDistribution variance exp - 1)

```

D.6.2 Log normal distribution — Java implementation

Listing D.12 shows the implementation of the log normal distribution in Java.

Listing D.12 Java implementation of the log normal distribution

```

package DhbStatistics;

import DhbScientificCurves.Histogram;
import DhbInterfaces.ParametrizedOneVariableFunction;
/**
 * Log normal distribution
 *
 * @author Didier H. Besset
 *
 */
public final class LogNormalDistribution
    extends ProbabilityDensityWithUnknownDistribution
{
    /**
     * Normal distribution with the same parameters.
     */
    private NormalDistribution normalDistr;

```



```

/**
 * Defines a Log Normal distribution with known parameters.
 * @param mu double
 * @param sigma double
 * @exception java.lang.IllegalArgumentException
 *         when sigma is non-positive
 */
public LogNormalDistribution ( double mu, double sigma)
    throws IllegalArgumentException
{
    normalDistr = new NormalDistribution( mu, sigma);
}
/**
 * Create an instance of the receiver with parameters estimated from
 * the given histogram using best guesses. This method can be used to
 * find the initial values for a fit.
 * @param h DhbScientificCurves.Histogram
 * @exception java.lang.IllegalArgumentException
 *         when no suitable parameter can be found.
 */
public LogNormalDistribution( Histogram h)
    throws IllegalArgumentException
{
    if ( h.getMinimum() < 0 )
        throw new IllegalArgumentException(
            "Log normal distribution is only defined for non-negative values");
    double average = h.average();
    if ( average <= 0 )
        throw new IllegalArgumentException(
            "Log normal distribution is only defined for positive average");
    double variance = h.variance();
    double sigma2 = Math.log( variance / (average * average) + 1);
    if ( sigma2 <= 0 )
        throw new IllegalArgumentException(
            "Log normal distribution is only defined for positive sigma");
    normalDistr = new NormalDistribution( Math.log( average),
        Math.sqrt( sigma2));
}
/**
 * @return double average of the distribution.
 */
public double average ( )
{
    return Math.exp( normalDistr.variance() / 2
        + normalDistr.average());
}

```

```
/**
 * @return double the lowest value of the random variable
 */
protected double lowValue()
{
    return 0;
}
/**
 * @return java.lang.String name of the distribution
 */
public String name ( )
{
    return "Log normal distribution";
}
/**
 * @return double[] an array containing the parameters of
 * the distribution.
 */
public double[] parameters ( )
{
    return normalDistr.parameters();
}
/**
 * @return double a random number distributed according to the receiver.
 */
public double random( )
{
    return Math.exp( normalDistr.random());
}
/**
 * @param m double
 */
public void setMu( double mu)
{
    normalDistr.setAverage( mu);
}
/**
 * @param p double[] assigns the parameters
 */
public void setParameters( double[] params)
{
    setMu( params[0]);
    setSigma( params[1]);
}
/**
 * @param sigma double
```

```

*/
public void setSigma( double sigma)
{
    normalDistr.setStandardDeviation( sigma);
}
/**
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    java.text.DecimalFormat fmt = new java.text.DecimalFormat(
        "####0.00000");
    sb.append("Log normal distribution (");
    sb.append(fmt.format(normalDistr.average()));
    sb.append(', ');
    sb.append(fmt.format(normalDistr.standardDeviation()));
    sb.append(')');
    return sb.toString();
}
/**
 * @return double probability density function
 * @param x double random variable
 */
public double value ( double x)
{
    return x > 0 ? normalDistr.value( Math.log(x)) / x : 0;
}
/**
 * @return double variance of the distribution.
 */
public double variance ( )
{
    double variance = normalDistr.variance();
    return Math.exp( variance + 2 * normalDistr.average())
        * ( Math.exp( variance) - 1);
}
}

```

D.7 Triangular distribution

Table D.7 shows the properties of the triangular distribution. The triangular distribution is ad-hoc distribution used when a variable is limited to an interval.

Table D.7: Properties of the triangular distribution

Range of random variable	$[a, b]$
Probability density function	$P(x) = \begin{cases} \frac{2(x-a)}{(b-a)(c-a)} & \text{if } a \leq x \leq c, \\ \frac{2(b-x)}{(b-a)(b-c)} & \text{if } c \leq x \leq b. \end{cases}$
Parameters	$\begin{aligned} -\infty < a \leq c \leq b < +\infty \\ a < b \end{aligned}$
Distribution function	$F(x) = \begin{cases} \frac{(x-a)^2}{(b-a)(c-a)} & \text{if } a \leq x \leq c, \\ 1 - \frac{(b-x)^2}{(b-a)(b-c)} & \text{if } c \leq x \leq b. \end{cases}$
Average	$\frac{a + b + c}{3}$
Variance	$\frac{a^2 + b^2 + c^2 - ab - ac - bc}{18}$
Skewness	$\frac{a^3 + b^3 + c^3}{135} + \dots$
Kurtosis	\dots

D.7.1 Triangular distribution — Smalltalk implementation

Listing D.13 shows the implementation of the triangular distribution in Smalltalk.

Listing D.13 Smalltalk implementation of the triangular distribution

Class **DhbTriangularDistribution**

Subclass of DhbProbabilityDensity

Instance variable names: lowLimit highLimit peak

Class methods

distributionName

^'Triangular distribution'

fromHistogram: aHistogram

| b c|

```
b := aHistogram standardDeviation * 1.73205080756888
```

```
new
```

```
^self new: (1 / 2) from: 0 to: 1
```

```
new: aNumber1 from: aNumber2 to: aNumber3
```

```
^super new initialize: aNumber1 from: aNumber2 to: aNumber3
```

```
Instance methods
```

```
acceptanceBetween: aNumber1 and: aNumber2
```

```
^self privateAcceptanceBetween: aNumber1 and: aNumber2
```

```
average
```

```
^(lowLimit + peak + highLimit) / 3
```

```
changeParametersBy: aVector
```

```
lowLimit := lowLimit + ( aVector at: 1).
```

```
highLimit := highLimit + ( aVector at: 2).
```

```
peak := peak + ( aVector at: 3).
```

```
distributionValue: aNumber
```

```
| norm |
```

```
^( aNumber between: lowLimit and: highLimit)
```

```
  ifTrue: [ aNumber < peak
```

```
    ifTrue: [ norm := ( highLimit - lowLimit) * (
```

```
      peak - lowLimit).
```

```
      ( aNumber - lowLimit) squared /
```

```
      norm
```

```
    ]
```

```
  ifFalse: [ aNumber > peak
```

```
    ifTrue: [ norm := ( highLimit
```

```
      - lowLimit) * ( highLimit - peak).
```

```
      1 - ( (
```

```
      highLimit - aNumber) squared / norm)
```

```
    ]
```

```
  ifFalse: [ ( peak - lowLimit)
```

```
    / ( highLimit - lowLimit)]
```

```
  ]
```

```

    ]
    ifFalse:[ 0]

initialize: aNumber1 from: aNumber2 to: aNumber3
    ( aNumber2 < aNumber3 and: [ aNumber1 between: aNumber2 and:
                                aNumber3])
    ifFalse: [ self error: 'Illegal distribution parameters' ].
    peak := aNumber1.
    lowLimit := aNumber2.
    highLimit := aNumber3.
    ^self

inverseAcceptanceAfterPeak: aNumber
    ^ highLimit - ( ( ( 1 - aNumber) * ( highLimit - lowLimit) * (
                                highLimit - peak)) sqrt)

inverseAcceptanceBeforePeak: aNumber
    ^ ( aNumber * ( highLimit - lowLimit) * ( peak - lowLimit)) sqrt
      + lowLimit

kurtosis
    ^(-6/10)

parameters
    ^Array with: lowLimit with: highLimit with: peak

privateInverseDistributionValue: aNumber
    ^ ( peak - lowLimit) / ( highLimit - lowLimit) > aNumber
      ifTrue: [ self inverseAcceptanceBeforePeak: aNumber]
      ifFalse: [ self inverseAcceptanceAfterPeak: aNumber]

skewness
    ^(((lowLimit squared * lowLimit + ( peak squared * peak) + (
                                highLimit squared * highLimit) ) / 135)
      -(((lowLimit squared * peak) + (lowLimit squared * highLimit) +
        (peak squared * lowLimit) + (peak squared * highLimit) + (highLimit
        squared * lowLimit) + (highLimit squared * peak))/90)
      + ( 2 * lowLimit * peak * highLimit / 45)) / ( self
        standardDeviation raisedToInteger: 3)

```

value: aNumber

```

| norm |
^( aNumber between: lowLimit and: highLimit)
  ifTrue: [ aNumber < peak
    ifTrue: [ norm := ( highLimit - lowLimit) * (
      peak - lowLimit).
      2 * ( aNumber - lowLimit) / norm
    ]
    ifFalse: [ aNumber > peak
      ifTrue: [ norm := ( highLimit
        - lowLimit) * ( highLimit - peak).
        2 * ( highLimit
        - aNumber) / norm
      ]
      ifFalse: [ 2 / ( highLimit -
        lowLimit)]
    ]
  ]
  ifFalse: [ 0]

```

variance

```

^(lowLimit squared + peak squared + highLimit squared - (
lowLimit * peak) - ( lowLimit * highLimit) - ( peak * highLimit)) /
18

```

D.7.2 Triangular distribution — Java implementation

Listing D.14 shows the implementation of the triangular distribution in Java.

Listing D.14 Java implementation of the triangular distribution

```

package DhbStatistics;

import DhbScientificCurves.Histogram;
import DhbInterfaces.ParametrizedOneVariableFunction;
/**
 * Triangular distribution.
 *
 * @author Didier H. Besset
 */
public final class TriangularDistribution
    extends ProbabilityDensityFunction

```

```

{
    /**
     * Low limit.
     */
    private double a;
    /**
     * High limit.
     */
    private double b;
    /**
     * peak location.
     */
    private double c;

    /**
     * Constructor method.
     * @param low double    low limit
     * @param high double   high limit
     * @param peak double    peak of the distribution
     * @exception java.lang.IllegalArgumentException
     *             if the limits are inverted or
     *             if the peak is outside the limits.
     */
    public TriangularDistribution ( double low, double high, double peak) throws IllegalArgumentException
    {
        if ( low >= high )
            throw new IllegalArgumentException(
                "Limits of distribution are equal or reversed");
        if ( peak < low || peak > high )
            throw new IllegalArgumentException(
                "Peak of distribution lies outside the limits");
        a = low;
        b = high;
        c = peak;
    }

    /**
     * Create an instance of the receiver with parameters estimated from
     * the given histogram using best guesses. This method can be used to
     * find the initial values for a fit.
     * @param h DhbScientificCurves.Histogram
     */
    public TriangularDistribution( Histogram h)
    {
        b = h.standardDeviation() * 1.73205080756888; // sqrt(12)/2
        c = h.average();
        a = c - b;
    }
}

```



```

        b += c;
    }
    /**
     * @return double average of the distribution.
     */
    public double average()
    {
        return ( a + b + c ) / 3;
    }
    /**
     * Returns the probability of finding a random variable smaller
     * than or equal to x.
     * @return integral of the probability density function from a to x.
     * @param x double upper limit of integral.
     */
    public double distributionValue ( double x)
    {
        if ( x < a )
            return 0;
        else if ( x < c )
            return ( x - a ) * ( x - a ) / ( ( b - a ) * ( c - a ) );
        else if ( x < b )
            return 1 - ( b - x ) * ( b - x ) / ( ( b - a ) * ( b - c ) );
        else
            return 1;
    }
    /**
     * @return java.lang.String    name of the distribution
     */
    public String name()
    {
        return "Triangular distribution";
    }
    /**
     * @return double[] an array containing the parameters of
     *                  the distribution.
     */
    public double[] parameters()
    {
        double[] answer = new double[3];
        answer[0] = a;
        answer[1] = b;
        answer[2] = c;
        return answer;
    }
    /**

```

```

* This method assumes that the range of the argument has been checked.
* @return double the value for which the distribution function
*           is equal to x.
* @param x double value of the distribution function.
*/
protected double privateInverseDistributionValue ( double x)
{
    return ( x < ( c - a ) / ( b - a )
        ? Math.sqrt( x * ( b - a ) * ( c - a ) ) + a
        : b - Math.sqrt( ( 1 - x ) * ( b - a ) * ( b - c ) ) );
}
/**
* @param p double[] assigns the parameters
*/
public void setParameters( double[] params)
{
    a = params[0];
    b = params[1];
    c = params[2];
}
/**
* @return double probability density function
* @param x double random variable
*/
public double value( double x)
{
    if ( x < a )
        return 0;
    else if ( x < c )
        return 2 * ( x - a ) / ( ( b - a ) * ( c - a ) );
    else if ( x < b )
        return 2 * ( b - x ) / ( ( b - a ) * ( b - c ) );
    else
        return 0;
}
/**
* @return double variance of the distribution
*/
public double variance()
{
    return ( a * a + b * b + c * c - a * b - a * c - b * c ) / 18;
}
}

```

D.8 Uniform distribution

Table D.8 shows the properties of the uniform distribution. The uniform distribution is

Table D.8: Properties of the uniform distribution

Range of random variable	$[a, b]$
Probability density function	$P(x) = \frac{1}{b-a}$
Parameters	$-\infty < a < b < +\infty$
Distribution function	$F(x) = \frac{x-a}{b-a}$
Average	$\frac{a+b}{2}$
Variance	$\frac{(b-a)^2}{12}$
Skewness	0
Kurtosis	-1.2

another ad-hoc distribution used when a variable is limited to an interval.

D.8.1 Uniform distribution — Smalltalk implementation

Listing D.15 shows the implementation of the uniform distribution in Smalltalk.

Listing D.15 Smalltalk implementation of the uniform distribution

```

Class                DhbUniformDistribution
Subclass of          DhbProbabilityDensity
Instance variable names: lowLimit highLimit

```

```

    Class methods
distributionName
    ^'Uniform distribution'

from: aNumber1 to: aNumber2
    ^super new initialize: aNumber1 to: aNumber2

fromHistogram: aHistogram

```

```

| b c|
b := aHistogram standardDeviation * 1.73205080756888

```

```

new

```

```

^self from: 0 to: 1

```

```

Instance methods

```

```

acceptanceBetween: aNumber1 and: aNumber2

```

```

^self privateAcceptanceBetween: aNumber1 and: aNumber2

```

```

average

```

```

^( highLimit + lowLimit) / 2

```

```

changeParametersBy: aVector

```

```

lowLimit := lowLimit + ( aVector at: 1).
highLimit := highLimit + ( aVector at: 2).

```

```

distributionValue: aNumber

```

```

aNumber < lowLimit
  ifTrue: [ ^0].
^aNumber < highLimit
  ifTrue: [ (aNumber - lowLimit) / ( highLimit - lowLimit)]
  ifFalse:[ 1]

```

```

initialize: aNumber1 to: aNumber2

```

```

aNumber1 < aNumber2
  ifFalse: [ self error: 'Illegal distribution parameters'].
lowLimit := aNumber1.
highLimit := aNumber2.
^self

```

```

kurtosis

```

```

^-12 / 10

```

```

parameters

```

```

^Array with: lowLimit with: highLimit

privateInverseDistributionValue: aNumber
    ^(highLimit - lowLimit) * aNumber + lowLimit

skewness
    ^0

standardDeviation
    ^( highLimit - lowLimit) / 3.46410161513775

value: aNumber
    ^( aNumber between: lowLimit and: highLimit)
    ifTrue: [ 1/( highLimit - lowLimit)]
    ifFalse:[ 0]

variance
    ^( highLimit - lowLimit) squared / 12

```

D.8.2 Uniform distribution — Java implementation

Listing D.16 shows the implementation of the uniform distribution in Java.

Listing D.16 Java implementation of the uniform distribution

```

package DhbStatistics;

import DhbScientificCurves.Histogram;
/**
 * Uniform distribution over a given interval.
 *
 * @author Didier H. Besset
 */
public final class UniformDistribution
    extends ProbabilityDensityFunction
{
    /**

```

```

    * Low limit.
    */
    private double a;
    /**
    * High limit.
    */
    private double b;

    /**
    * Constructor method.
    * @param low double    low limit
    * @param high double    high limit
    * @exception java.lang.IllegalArgumentException
    *             if the limits are inverted.
    */
    public UniformDistribution ( double low, double high)
        throws IllegalArgumentException
    {
        if ( low >= high )
            throw new IllegalArgumentException(
                "Limits of distribution are equal or reversed");
        a = low;
        b = high;
    }
    /**
    * Create an instance of the receiver with parameters estimated from
    * the given histogram using best guesses. This method can be used to
    * find the initial values for a fit.
    * @param h DhbScientificCurves.Histogram
    */
    public UniformDistribution( Histogram h)
    {
        b = h.standardDeviation() * 1.73205080756888; // sqrt(12)/2
        double c = h.average();
        a = c - b;
        b += c;
    }
    /**
    * @return double average of the distribution.
    */
    public double average()
    {
        return ( a + b ) * 0.5;
    }
    /**
    * Returns the probability of finding a random variable smaller

```

```

* than or equal to x.
* @return integral of the probability density function from a to x.
* @param x double upper limit of integral.
*/
public double distributionValue ( double x)
{
    if ( x < a )
        return 0;
    else if ( x < b )
        return ( x - a ) / ( b - a );
    else
        return 1;
}
/**
 * @return double kurtosis of the distribution.
 */
public double kurtosis( )
{
    return -1.2;
}
/**
 * @return java.lang.String    name of the distribution
 */
public String name ( )
{
    return "Uniform distribution";
}
/**
 * @return double[] an array containing the parameters of
 *                  the distribution.
 */
public double[] parameters ( )
{
    double[] answer = new double[2];
    answer[0] = a;
    answer[1] = b;
    return answer;
}
/**
 * This method assumes that the range of the argument has been checked.
 * @return double the value for which the distribution function
 *                is equal to x.
 * @param x double value of the distribution function.
 */
protected double privateInverseDistributionValue ( double x)
{

```

```

        return (b - a) * x + a;
    }
    /**
     * @param p double[]  assigns the parameters
     */
    public void setParameters( double[] params)
    {
        a = params[0];
        b = params[1];
    }
    /**
     * @return double skewness of the distribution.
     */
    public double skewness( )
    {
        return 0;
    }
    /**
     * @return double probability density function
     * @param x double random variable
     */
    public double value( double x)
    {
        if ( x < a )
            return 0;
        else if ( x < b )
            return 1 / (b - a);
        else
            return 0;
    }
    /**
     * @return double variance of the distribution
     */
    public double variance()
    {
        double range = b - a;
        return range * range / 12;
    }
}

```

D.9 Weibull distribution

Table D.9 shows the properties of the Weibull distribution. The Weibull distribution is used to model the behavior of reliability. It is defined by its acceptance function. Its

Table D.9: Properties of the Weibull distribution

Range of random variable	$[0, +\infty[$
Probability density function	$P(x) = \frac{\alpha x^{\alpha-1}}{\beta^\alpha} e^{-\left(\frac{x}{\beta}\right)^\alpha}$
Parameters	$0 < \alpha < \infty$ $0 < \beta < \infty$
Distribution function	$F(x) = 1 - e^{-\left(\frac{x}{\beta}\right)^\alpha}$
Average	$\frac{\beta}{\alpha} \Gamma\left(\frac{1}{\alpha}\right)$
Variance	$\frac{\beta^2}{\alpha} \left[2\Gamma\left(\frac{2}{\alpha}\right) - \frac{1}{\alpha} \Gamma\left(\frac{1}{\alpha}\right)^2 \right]$
Skewness	
Kurtosis	

shape is similar to that of the gamma distribution and, thus, can be applied to the same types of problems. Figure D.7 shows the shapes taken by the Weibull distribution for a few values of the parameters.

Because the Weibull distribution is defined by its distribution function, the estimation of the initial values of the parameters from a histogram is made by computing the distribution function at 2 positions. These positions are determined using the histogram limits and the average so that the estimation of the distribution function using the histogram has enough significance.

D.9.1 Weibull distribution — Smalltalk implementation

Listing D.17 shows the implementation of the Weibull distribution in Smalltalk.

Listing D.17 Smalltalk implementation of the Weibull distribution

```

Class                DhbWeibullDistribution
Subclass of          DhbProbabilityDensity
Instance variable names:  alpha beta norm

Class methods
distributionName
^^Weibull distribution'
```




Figure D.7: Weibull distribution for a few parameters

fromHistogram: aHistogram

```

| average xMin xMax accMin accMax |
aHistogram minimum < 0
  ifTrue: [ ^nil].
average := aHistogram average.
xMin := ( aHistogram minimum + average) / 2.
accMin := ( aHistogram countsUpTo: xMin) / aHistogram totalCount.
xMax := ( aHistogram maximum + average) / 2.
accMax := ( aHistogram countsUpTo: xMax) / aHistogram totalCount.
^[self solve: xMin acc: accMin upper: xMax acc: accMax]
  when: ExAll do: [ :signal | signal exitWith: nil]

```

new

```

^self error: 'Illegal creation message for this class'

```

shape: aNumber1 scale: aNumber2

```

^super new initialize: aNumber1 scale: aNumber2

```

solve: lowX acc: lowAcc upper: highX acc: highAcc

```

| lowLnAcc highLnAcc deltaLnAcc lowLnX highLnX |
lowLnAcc := (1 - lowAcc) ln negated ln.

```

```

highLnAcc := (1 - highAcc) ln negated ln.
deltaLnAcc := highLnAcc - lowLnAcc.
lowLnX := lowX ln.
highLnX := highX ln.
^self shape: deltaLnAcc / (highLnX - lowLnX)
  scale: ((highLnAcc * lowLnX - (lowLnAcc * highLnX)) /
          deltaLnAcc) exp

```

Instance methods

acceptanceBetween: aNumber1 and: aNumber2

```
^self privateAcceptanceBetween: aNumber1 and: aNumber2
```

average

```
^(1 / alpha) gamma * beta / alpha
```

changeParametersBy: aVector

```

alpha := alpha + ( aVector at: 1).
beta := beta + ( aVector at: 2).
self computeNorm.

```

computeNorm

```
norm := alpha / ( beta raisedTo: alpha).
```

distributionValue: aNumber

```

^aNumber > 0
  ifTrue: [ 1 - ( ( aNumber / beta) raisedTo: alpha) negated
            exp)]
  ifFalse:[ 0]

```

initialize: aNumber1 scale: aNumber2

```

( aNumber1 > 0 and: [ aNumber2 > 0])
  ifFalse: [ self error: 'Illegal distribution parameters'].
alpha := aNumber1.
beta := aNumber2.
self computeNorm.
^self

```

parameters

```

^Array with: alpha with: beta

privateInverseDistributionValue: aNumber
    ^((1 - aNumber) ln negated raisedTo: (1 / alpha)) * beta

value: aNumber
    ^((aNumber / beta) raisedTo: alpha) negated exp * (aNumber
        raisedTo: (alpha - 1)) * norm

variance
    ^((beta squared / alpha) * ((2 / alpha) gamma * 2 - ((1 / alpha)
        gamma squared / alpha))

```

D.9.2 Weibull distribution — Java implementation

Listing D.18 shows the implementation of the Weibull distribution in Java.

Listing D.18 Java implementation of the Weibull distribution

```

package DhbStatistics;

import DhbScientificCurves.Histogram;
import DhbFunctionEvaluation.GammaFunction;
import DhbInterfaces.ParametrizedOneVariableFunction;
/**
 * Weibull distribution.
 *
 * @author Didier H. Besset
 */
public final class WeibullDistribution
    extends ProbabilityDensityFunction
{
    /**
     * Shape parameter of the distribution.
     */
    protected double alpha;
    /**
     * Scale parameter of the distribution.
     */
    private double beta;

```

```

/**
 * Norm of the distribution (cached for efficiency).
 */
private double norm;

/**
 * Create a new instance of the Weibull distribution with given shape and scale.
 * @param shape double shape parameter of the distribution (alpha).
 * @param scale double scale parameter of the distribution (beta).
 * @exception java.lang.IllegalArgumentException
 *         if any of the parameters is non-positive.
 */
public WeibullDistribution ( double shape, double scale)
    throws IllegalArgumentException
{
    if ( shape <= 0 )
        throw new IllegalArgumentException(
            "Shape parameter must be positive");
    if ( scale <= 0 )
        throw new IllegalArgumentException(
            "Scale parameter must be positive");
    defineParameters( shape, scale);
}

/**
 * Create an instance of the receiver with parameters estimated from
 * the given histogram using best guesses. This method can be used to
 * find the initial values for a fit.
 * @param h DhbScientificCurves.Histogram
 * @exception java.lang.IllegalArgumentException
 *         when no suitable parameter can be found.
 */
public WeibullDistribution( Histogram h)
    throws IllegalArgumentException
{
    if ( h.getMinimum() < 0 )
        throw new IllegalArgumentException(
            "Weibull distribution is only defined for non-negative values");
    double average = h.average();
    if ( average <= 0 )
        throw new IllegalArgumentException(
            "Weibull distribution must have a non-negative average");
    double xMin = ( h.getMinimum() + average ) * 0.5;
    double accMin = Math.log( -Math.log( 1 - h.getCountsUpTo(xMin)
        / h.totalCount()));
    double xMax = ( h.getMaximum() + average ) * 0.5;
    double accMax = Math.log( -Math.log( 1 - h.getCountsUpTo(xMax)

```

```

        / h.totalCount()));
    double delta = accMax - accMin;
    xMin = Math.log( xMin);
    xMax = Math.log( xMax);
    defineParameters( delta / ( xMax - xMin),
        Math.exp( ( accMax * xMin - accMin * xMax)
            / delta));
}
/**
 * @return double average of the distribution.
 */
public double average()
{
    return GammaFunction.gamma( 1 / alpha) * beta / alpha;
}
/**
 * Assigns new values to the parameters.
 * This method assumes that the parameters have been already checked.
 */
public void defineParameters ( double shape, double scale)
{
    alpha = shape;
    beta = scale;
    norm = alpha / Math.pow( beta, alpha);
    return;
}
/**
 * Returns the probability of finding a random variable smaller
 * than or equal to x.
 * @return integral of the probability density function from 0 to x.
 * @param x double upper limit of integral.
 */
public double distributionValue ( double x)
{
    return 1.0 - Math.exp(-Math.pow( x / beta, alpha));
}
/**
 * @return java.lang.String the name of the distribution.
 */
public String name()
{
    return "Weibull distribution";
}
/**
 * @return double[] an array containing the parameters of
 * the distribution.

```

```

*/
public double[] parameters()
{
    double[] answer = new double[2];
    answer[0] = alpha;
    answer[1] = beta;
    return answer;
}
/**
 * This method assumes that the range of the argument has been checked.
 * @return double the value for which the distribution function
 *         is equal to x.
 * @param x double value of the distribution function.
 */
protected double privateInverseDistributionValue ( double x)
{
    return Math.pow( -Math.log( 1 - x), 1. / alpha ) * beta;
}
/**
 * @param p double[]  assigns the parameters
 */
public void setParameters( double[] params)
{
    defineParameters( params[0], params[1]);
}
/**
 * This method was created in VisualAge.
 * @return java.lang.String
 */
public String toString()
{
    StringBuffer sb = new StringBuffer();
    java.text.DecimalFormat fmt = new java.text.DecimalFormat(
        "####0.00000");
    sb.append("Weibull distribution (");
    sb.append(fmt.format(alpha));
    sb.append(',');
    sb.append(fmt.format(beta));
    sb.append(')');
    return sb.toString();
}
/**
 * @return double probability density function
 * @param x double random variable
 */
public double value( double x)

```

```
{
    return norm * Math.pow( x, alpha - 1)
        * Math.exp( -Math.pow( x / beta, alpha));
}
/**
 * @return double variance of the distribution.
 */
public double variance ( )
{
    double s = GammaFunction.gamma( 1 / alpha);
    return beta * beta * ( 2 * GammaFunction.gamma( 2 / alpha)
        - s * s / alpha) / alpha;
}
}
```

Appendix E

Accurate accumulation of expectation values

E.1 Accurate accumulation of central moments

This section shows the detailed derivation of equation 9.13 of section 9.2. The aim of this demonstration is to express the central moment of order k estimated over a sample of $n + 1$ measurements as a function of the central moments of order lower or equal to k estimated over a sample of n measurements. The estimator of the central moment of order k is defined by:

$$\left\langle (x - \bar{x})^k \right\rangle_{n+1} = \frac{1}{n+1} \sum_{i=1}^{n+1} (x_i - \langle x \rangle_{n+1})^k. \quad (\text{E.1})$$

We shall now concentrate on changing the sum of equation E.1 in such way as to bring quantities which are already computed. The sum of E.1 is equal to

$$\begin{aligned} S &= (n+1) \left\langle (x - \bar{x})^k \right\rangle_{n+1} \\ &= \sum_{i=1}^{n+1} (x_i - \langle x \rangle_{n+1})^k \\ &= \sum_{i=1}^{n+1} (x_i - \langle x \rangle_n + \langle x \rangle_n - \langle x \rangle_{n+1})^k \\ &= \sum_{i=1}^{n+1} [(x_i - \langle x \rangle_n) + \Delta_{n+1}]^k, \end{aligned} \quad (\text{E.2})$$

where we have introduced the correction defined in equation 9.12. We can now transform the expression inside the sum using the binomial expansion:

$$\begin{aligned} S &= \sum_{i=1}^{n+1} \sum_{l=0}^k \binom{l}{k} (x_i - \langle x \rangle_n)^l \Delta_{n+1}^{k-l} \\ &= \sum_{l=0}^k \binom{l}{k} \sum_{i=1}^{n+1} (x_i - \langle x \rangle_n)^l \Delta_{n+1}^{k-l} \end{aligned} \quad (\text{E.3})$$

In the second part of equation E.3 the two sums have been permuted. Like in the case of the average, we now make the last term of the inner sum explicit. The remaining sum can then be expressed as a function of the estimators of the central moments over n measurements. The term containing the $(n+1)^{\text{th}}$ measurement can be rewritten as a function of the correction defined in equation 9.12. We have:

$$\begin{aligned} S &= \sum_{l=0}^k \binom{l}{k} \left[(x_{n+1} - \langle x \rangle_n)^l \Delta_{n+1}^{k-l} + \Delta_{n+1}^{k-l} \sum_{i=1}^{n+1} (x_i - \langle x \rangle_n)^l \right] \\ &= \sum_{l=0}^k \binom{l}{k} \left[(n+1)^l \left(\frac{x_{n+1} - \langle x \rangle_n}{n+1} \right)^l \Delta_{n+1}^{k-l} + n \langle (x - \bar{x})^l \rangle_n \Delta_{n+1}^{k-l} \right] \\ &= \sum_{l=0}^k \binom{l}{k} \left[(-n-1)^l \Delta_{n+1}^{k-l} + n \langle (x - \bar{x})^l \rangle_n \Delta_{n+1}^{k-l} \right] \\ &= \sum_{l=0}^k \binom{l}{k} (-n-1)^l \Delta_{n+1}^{k-l} + \sum_{l=0}^k \binom{l}{k} n \langle (x - \bar{x})^l \rangle_n \Delta_{n+1}^{k-l}. \end{aligned} \quad (\text{E.4})$$

In the last line of equation E.4 the first term contains the binomial expansion of the following expression

$$\sum_{l=0}^k \binom{l}{k} (-n-1)^l = [1 + (-n-1)]^k = (-n)^k. \quad (\text{E.5})$$

Thus, we have:

$$S = (-n \Delta_{n+1})^k + n \sum_{l=0}^k \binom{l}{k} \langle (x - \bar{x})^l \rangle_n \Delta_{n+1}^{k-l}. \quad (\text{E.6})$$

In this last equation, the first term of the sum is just Δ_{n+1}^k and the second term of the sum vanishes by definition of the average \bar{x} . This gives us the final expression to compute the estimator of the central moment computed over $n+1$ measurements as a function of the estimator of the central moment computed over n measurements

$$\langle (x - \bar{x})^k \rangle_{n+1} = \frac{n}{n+1} \left\{ [1 - (-n)^{k-1}] \Delta_{n+1}^k + \sum_{l=2}^k \binom{l}{k} \langle (x - \bar{x})^l \rangle_n \Delta_{n+1}^{k-l} \right\}. \quad (\text{E.7})$$

Quod erat demonstrandum. . .

E.2 Accurate accumulation of the covariance

This section shows the detailed derivation of equation 12.6. To simplify notation, the components x_i and x_j have been renamed x and y respectively.

The estimator of the covariance of two random variables x and y over n measurements is defined by:

$$\text{cov}_n(x, y) = \langle (x_i - \langle x \rangle_n) (y_i - \langle y \rangle_n) \rangle_n = \frac{1}{n} \sum_{i=1}^n (x_i - \langle x \rangle_n) (y_i - \langle y \rangle_n). \quad (\text{E.8})$$

The estimator of the covariance of x and y over $n + 1$ measurements is then given by:

$$\text{cov}_{n+1}(x, y) = \frac{1}{n+1} \sum_{i=1}^{n+1} (x_i - \langle x \rangle_{n+1}) (y_i - \langle y \rangle_{n+1}). \quad (\text{E.9})$$

The sum in the equation above can then be expressed as:

$$\begin{aligned} C_{n+1} &= (n+1) \langle (x_i - \langle x \rangle_{n+1}) (y_i - \langle y \rangle_{n+1}) \rangle_{n+1} \\ &= \sum_{i=1}^{n+1} (x_i - \langle x \rangle_{n+1}) (y_i - \langle y \rangle_{n+1}) \\ &= \sum_{i=1}^{n+1} (x_i - \langle x \rangle_n + \langle x \rangle_n - \langle x \rangle_{n+1}) (y_i - \langle y \rangle_n + \langle y \rangle_n - \langle y \rangle_{n+1}) \\ &= \sum_{i=1}^{n+1} (x_i - \langle x \rangle_n + \Delta_{x,n+1}) (y_i - \langle y \rangle_n + \Delta_{y,n+1}), \end{aligned} \quad (\text{E.10})$$

where we have introduced the corrections to the estimation of the expectation value of x and y as follow:

$$\left\{ \begin{aligned} \Delta_{x,n+1} &= \langle x \rangle_n - \langle x \rangle_{n+1} \\ &= \frac{\langle x \rangle_n - x_{n+1}}{n+1}, \\ \Delta_{y,n+1} &= \langle y \rangle_n - \langle y \rangle_{n+1} \\ &= \frac{\langle y \rangle_n - y_{n+1}}{n+1}. \end{aligned} \right. \quad (\text{E.11})$$

Thus, we have:

$$\begin{aligned} C_{n+1} &= \sum_{i=1}^{n+1} [(x_i - \langle x \rangle_n) (y_i - \langle y \rangle_n) + \Delta_{y,n+1} (x_i - \langle x \rangle_n) \\ &\quad + \Delta_{x,n+1} (y_i - \langle y \rangle_n) + \Delta_{x,n+1} \Delta_{y,n+1}] \\ &= \sum_{i=1}^{n+1} (x_i - \langle x \rangle_n) (y_i - \langle y \rangle_n) + (n+1) \Delta_{x,n+1} \Delta_{y,n+1} \\ &\quad + \Delta_{y,n+1} (x_{n+1} - \langle x \rangle_n) + \Delta_{x,n+1} (y_{n+1} - \langle y \rangle_n). \end{aligned} \quad (\text{E.12})$$

The last line is obtained from the definition of the expectation values $\langle x \rangle_n$ and $\langle y \rangle_n$. Using the definitions of $\Delta_{x,n+1}$ and $\Delta_{y,n+1}$ we have:

$$\begin{aligned}
 C_{n+1} &= \sum_{i=1}^{n+1} (x_i - \langle x \rangle_n) (y_i - \langle y \rangle_n) - (n+1) \Delta_{x,n+1} \Delta_{y,n+1} \\
 &= \sum_{i=1}^n (x_i - \langle x \rangle_n) (y_i - \langle y \rangle_n) - (n+1) \Delta_{x,n+1} \Delta_{y,n+1} \quad (\text{E.13}) \\
 &\quad + (x_{n+1} - \langle x \rangle_n) (y_{n+1} - \langle y \rangle_n) \\
 &= n \text{cov}_n(x, y) + n(n+1) \Delta_{x,n+1} \Delta_{y,n+1}.
 \end{aligned}$$

Now, we obtain the expression for the estimator of the covariance over $n+1$ measurements as a function of the estimator of the covariance over n measurements:

$$\text{cov}_{n+1}(x, y) = \frac{n}{n+1} \text{cov}_n(x, y) + n \Delta_{x,n+1} \Delta_{y,n+1}. \quad (\text{E.14})$$

Note that this equation yields equation 9.14 if one put $y = x$.

Appendix F

Accompanying CD ROM

All the code presented in this book is available in the accompanying CD ROM. The contents of the CD ROM is described in the sections of this chapter.

Permission to use the code for any purpose is granted to the purchaser of the book. The code is distributed on an "AS IS" basis, without warranty. Neither the author nor Morgan Kaufmann, Inc. make any representation, or warranty, either expressed or implied, with respect to the code, its quality or accuracy, or fitness for a specific purpose. Therefore, neither the author nor Morgan Kaufmann, Inc. shall have any liability to you or any other person or entity with respect to any liability, loss or damage caused or alleged to have been caused directly or indirectly by the code contained on the CD ROM. This includes, but is not limited to, interruption of service, loss of data, loss of classroom time, loss of consulting or anticipatory profits, or consequential damages from the use of the code contained on the CD ROM.

F.1 Smalltalk code

The Smalltalk code presented in this book was created with Visual Age for Smalltalk under ENVY. The entire code is contained in a single application. The application is split into several subapplications: these subapplications correspond more or less to the chapters of this book. To be honest, they correspond to the chapters of the early version.

F.1.1 non-ENVY users

Readers using Smalltalk without ENVY must load the classes in so-called chunk format. These files are located in the directory Smalltlk and all subdirectories. The names of the subdirectories and files correspond to the subapplications of the ENVY repository. These names have been modified to be used on systems using an 8.3 file name convention.

The classes of early chapters are referenced by classes of subsequent chapters. The reader should refer to figure 1.3 of chapter 1 for an explicit description of the dependency between classes.

Here is a description of the files:

Smalltlk/DataMing/dataming.app All classes defined in chapter 12.

Smalltlk/Estim/genestim.app All classes defined in chapter 10 and extensions of classes DhbHistogram and DhbStatisticalMoments related to the t -test.

Smalltlk/Estim/estimChi.app Extensions of classes DhbHistogram and DhbStatisticalMoments related to the χ^2 -test.

Smalltlk/Estim/estimF.app Extensions of classes DhbHistogram and DhbStatisticalMoments related to the F -test.

Smalltlk/Functs/general.app The class DhbPolynomial defined in chapter 2 and all classes defined in chapter 3.

Smalltlk/Functs/beta.app Extensions of class Number defined in chapter 2 for the beta function.

Smalltlk/Functs/errfunct.app The class DhbErfApproximation defined in chapter 2 and extensions of class Number defined in chapter 2 for the error function.

Smalltlk/Functs/gamma.app The class DhbLanczosFormula defined in chapter 2 and extensions of classes Number and Integer defined in chapter 2 for the gamma function.

Smalltlk/Iterativ/iterativ.app All classes defined in chapters 4, 5, 6 and 7.

Smalltlk/LinAlgeb/linalgeb.app The classes defined in chapter 8, except for classes defined in the subapplications of this one.

Smalltlk/LinAlgeb/eigenv.app The classes DhbLargestEigenValueFinder and DhbJacobiTransformation defined in chapter 8.

Smalltlk/LinAlgeb/determin.app Extensions of classes DhbMatrix and DhbLUPDecomposition related to the computation of the matrix determinant.

Smalltlk/LinAlgeb/inverse.app Extensions of classes DhbMatrix, DhbSymmetricMatrix and DhbLUPDecomposition related to the computation of matrix inversion.

Smalltlk/Minimize/minimize.app All classes defined in chapter 11.

Smalltlk/NumPrec/numprec.app The class DhbFloatingPointMachine defined in chapter 1 and the class DhbDecimalFloatingNumber described in appendix A.

Smalltlk/Statist/stats.app All classes defined in chapter 9 as well as all classes describing a probability distribution (defined in chapters 9, 10 and appendix D).

F.1.2 ENVY users

Readers using ENVY can load the ENVY repository contained in the file Vast/OONumS.dat. This repository contains a single application, DhbNumericalMethods, containing the entire code presented in this book. Here is a brief description of the contents of each subapplication.

DhbDataMining All classes defined in chapter 12.

DhbEstimation All classes defined in chapter 10 and extensions of classes DhbHistogram and DhbStatisticalMoments related to the t test.

DhbEstimationChi2 Extensions of classes DhbHistogram and DhbStatisticalMoments related to the χ^2 test.

DhbEstimationF Extensions of classes DhbHistogram and DhbStatisticalMoments related to the F test.

DhbFunctionEvaluation The class DhbPolynomial defined in chapter 2 and all classes defined in chapter 3.

DhbBetaFunction Extensions of class Number defined in chapter 2 for the beta function.

DhbErrorFunction The class DhbErfApproximation defined in chapter 2 and extensions of class Number defined in chapter 2 for the error function.

DhbGammaFunction The class DhbLanczosFormula defined in chapter 2 and extensions of classes Number and Integer defined in chapter 2 for the gamma function.

DhbIterativeAlgorithms All classes defined in chapters 4, 5, 6 and 7.

DhbLinearAlgebra The classes defined in chapter 8, except for classes defined in the subapplications of this one.

DhbEigenValuesAndVectors The classes DhbLargestEigenValueFinder and DhbJacobiTransformation defined in chapter 8.

DhbMatrixDeterminant Extensions of classes DhbMatrix and DhbLUPDecomposition related to the computation of the matrix determinant.

DhbMatrixInversion Extensions of classes DhbMatrix, DhbSymmetricMatrix and DhbLUPDecomposition related to the computation of matrix inversion.

DhbMinimization All classes defined in chapter 11.

DhbNumericalPrecision The class DhbFloatingPointMachine defined in chapter 1 and the class DhbDecimalFloatingNumber described in appendix A.

DhbStatistics All classes defined in chapter 9 as well as all classes describing a probability distribution (defined in chapters 9, 10 and appendix D).

F.2 Java code

The Java code presented in this book was created with Visual Age for Java. The entire code is split into several packages: these packages correspond more or less to the chapters of this book. They differ slightly from the Smalltalk applications because the Java code was generated independently from the Smalltalk code. Also due to the fact that Java does not allow to extend a class in another file, all files contain only one class.

The following packages have been defined:

`DhbDataMining` All classes defined in chapter 12.

`DhbEstimation` All classes defined in chapter 10.

`DhbFunctionEvaluation` All classes defined in chapters 1 and 2.

`DhbInterfaces` All interfaces defined in this book.

`DhbInterpolation` All classes defined in chapter 3.

`DhbIterations` All classes defined in chapters 4, 5, 6 and 7.

`DhbMatrixAlgebra` All classes defined in chapter 8.

`DhbOptimizing` All classes defined in chapter 11.

`DhbScientificCurves` The classes `Histogram` and `Curve`.

`DhbStatistics` The classes defined in chapter 9, except class `Histogram` and the classes defined in appendix D.

F.2.1 Visual Age for Java

The file `CDRom/Vaj/OONumJ.dat` contains a Visual Age repository. Users of Visual for Java can load the entire code by loading this file using the import facility of Visual Age.

F.2.2 Other Java systems

Classes must be loaded from the hierarchy of directories rooted at `CDRom/Java`. Each subdirectory corresponds to one of the packages described earlier.

Note: Because some Java system requires that the name of the class be identical to the name of the file, the original names of the classes have been kept. Similarly, the name of the directory must be equal to the name of the package. If you have transferred the contents of the CD ROM over a system limiting the size of the file name, you must first rename each file and directory before loading the classes.

F.3 Utilities

The directory Utility can only be used on a 32-bit Windows systems. It contains two programs to study some numerical effects experimentally. The directory must be copied as is onto a 32-bit Windows systems (Windows NT 4.0 or higher, Windows 98 or higher). The executable files must be run from this directory because it contains all the resources needed to execute them.

F.3.1 Distribution demo

The program contained in the file DistributionDemo.exe allows the reader to study all probability density functions described in this book. After starting the program, a window containing a note book will open. Each page of the notebook corresponds to a distribution. The parameters can be modified.

Then, clicking on the button labeled Generate will generate random values distributed according to the distribution. A window allows to modify the number of generated values and the parameters of the histogram into which the values are accumulated. When the generating is completed the resulting histogram is displayed. The reader can then attempt a least square fit or a maximum likelihood fit and view the result. The button labeled Integral allows to view the values as a distribution function.

F.3.2 File reader

This small utility program allows to display data read from a text file, either in tab or coma delimited format, in graphical form. Once data are read, the user can investigate interpolation, linear regression and polynomial least square fit. This program was used to generate the figures of chapters 3 and 10.

After starting the program, a window containing the parameters of the file containing the data will open. Here the reader can select the name of the file to read, the format used (blank or coma delimited) and which column contains which data. For error bars, there are three cases:

1. no error;
2. the error is given explicitly;
3. the error is computed from the standard deviation and the number of values used to compute the standard deviation (c.f. equation 9.6 in chapter 9).

If the data were read without error, the reader can investigate the three interpolation algorithms described in chapter 3. In addition linear regression and polynomial least square fit can be tried, each point having an identical weight.

If the data were read with errors, the reader can investigate only Lagrange and Bulirsch-Stoer interpolation. Linear regression and polynomial least square fit are performed using the supplied error to compute the weight of each point as described in chapter 10.

Bibliography

- [Abramovitz & Stegun] Milton Abramovitz and Irene A. Stegun, Handbook of Mathematical Functions, Dover publications, Inc., 1964.
- [Achtley & Bryant] William R. Achtley and Edwin H. Bryant editors, Benchmark Papers in Systematic and Evolutionary Biology, Vol. 1, Dowden, Hutchinson & Ross, Inc., Stroudsburg, Pa.; distributed by Halsted Press [John Wiley & Sons, Inc.], New York, 1975.
- [Bass] J. Bass, Cours de Mathématiques, Tome II, Masson, 1968.
- [Beck] Kent Beck, Smalltalk Best Practice Patterns, Prentice Hall, 1997.
- [Berry & Linoff] Michael J.A. Berry and Gordon Linoff, Data mining for marketing, sales and customer support, John Wiley & Sons, Inc., 1997.
- [Cormen et al.] Thomas H. Cormen, Charles E. Leiserson and Ronald L. Rivest, Introduction to Algorithms, McGraw-Hill, 1990.
- [Gamma et al.] Erich Gamma, Richard Helm, Ralph Johnson and John Vlissides, Design Patterns, Addison-Wesley, 1995.
- [Gullberg] Jan Gullberg, Mathematics From the Birth of the Numbers, W.W. Norton & Company, 1997.
- [Ifrah] Georges Ifrah, Histoire Universelle des Chiffres, Robert Laffont, 1994.
- [Knudth 1] Donald E. Knuth, The Art of Computer Programming Vol. 1, Addison-Wesley, 1973.
- [Knudth 2] Donald E. Knuth, The Art of Computer Programming Vol. 2, Addison-Wesley, 1981.
- [Knudth 3] Donald E. Knuth, The Art of Computer Programming Vol. 3, Addison-Wesley, 1973.

- [Koza et al.] John R. Koza, Forrest H. Bennett III, David Andre and Martin A. Keane, Genetic Programming III, Morgan Kaufmann, 1999.
- [Law & Kelton] Averill M. Law and W. David Kelton, Simulation Modeling and Analysis, McGraw-Hill, 1982.
- [Phillips & Taylor] G.M. Phillips and P.J. Taylor, Theory and Applications of Numerical Analysis, Academic Press: London and New York, 1973 .
- [Press et al.] William H. Press, Saul A. Teukolsky, William T. Vetterling and Brian P. Flannery, Numerical recipes for C : the art of scientific computing, Cambridge University Press, 1992.
- [Alpert et al.] Sherman R. Alpert, Kyle Brown and Bobby Woolf, Design Pattern Smalltalk Companion, Addison-Wesley, 1998.
- [Smith] David N. Smith, IBM Smalltalk, The language, Addison-Wesley, 1995.
- [Flanagan] David Flanagan, Java in a nutshell, O'Reilly, 1996.